BUSINESS DEVELOPMENT · SALES

CRM, Sales Metrics & Sales Tools

The quant's edge in business development: pipelines, unit economics, the modern sales stack, and analytics that get you hired and promoted.

Study Guide 04 · Business Development & Sales Track
Audience: MBA-Finance fresher pivoting into BD/Sales (Excel · Python · SQL · Power BI · NISM Series-XV)
Context: Indian B2B / SaaS / fintech · figures in INR (₹)

Contents

1. Why This Guide Is Your Unfair Advantage

Most people who walk into a business-development interview can talk. They are personable, they tell stories, they "love talking to customers." What almost none of them can do is open a spreadsheet and tell you *why* the pipeline will or will not close, or defend a forecast number, or explain what a healthy LTV:CAC ratio looks like for a fintech at Series B. That is your gap in the market.

You are an MBA-Finance fresher who is comfortable in Excel, Python, SQL and Power BI, and who has cleared NISM Series-XV. In sales, that is not a side skill — it is a superpower. The whole discipline of **revenue operations (RevOps)** exists precisely because most salespeople are allergic to data. A Sales Development Representative (SDR) or Business Development Executive (BDE) who can build a clean funnel dashboard, calculate conversion rates by stage, and speak fluently about unit economics is on a fast track to Sales Ops, Sales Strategy, or a quota-carrying Account Executive (AE) role with a data reputation.

This guide teaches three things at once: (1) the **CRM** — the system of record where all selling is logged; (2) the **metrics** — the numbers that measure whether selling is working; and (3) the **tools** — the software stack that modern sales teams run on. Master these and you can walk into any BD interview and out-quant the room.

> **Key takeaways:** Sales is increasingly a numbers game run on software. Your finance/analytics background is a differentiator, not a handicap. Learn the CRM data model, memorise the metric formulas cold, and always tie numbers back to a business decision (hire, forecast, quota, budget).

2. CRM Fundamentals: What It Is and Why It Matters

2.1 What a CRM actually is

CRM stands for **Customer Relationship Management**. In practice, a CRM is a database with a friendly interface that stores every person, company, deal and interaction relevant to selling. Think of it as the "general ledger of the revenue team" — just as an accounting system is the single source of truth for money, the CRM is the single source of truth for customers and deals.

A CRM does four jobs:

- **System of record** — who are our customers and prospects, and what do we know about them?
- **System of workflow** — what needs to happen next, by whom, by when? (tasks, reminders, follow-ups)
- **System of reporting** — how much is in the pipeline, what is our win rate, will we hit the number?
- **System of accountability** — who did what, when? (activity logging, so managers can coach)

2.2 Why it matters — for the company and for you

For the **company**, a CRM turns selling from a set of private conversations trapped in individual reps' heads and inboxes into a shared, queryable asset. If a rep quits, their accounts and history stay. Forecasts become possible. Marketing can see which campaigns produced revenue. Leadership can see where deals stall.

For **you as a candidate or new hire**, being "good in the CRM" is a reputation you can build in your first 30 days. Reps who log diligently and keep clean data get better territories, cleaner leads and more trust from managers. In interviews, saying "I live in the CRM and I trust the pipeline because I keep it clean" signals professionalism far beyond a fresher's years.

Common trap — the CRM is not a formality. Many junior reps treat the CRM as bureaucratic box-ticking they do at 6 pm on Friday from memory. That produces garbage data, and garbage data produces garbage forecasts. Treat CRM updates as part of the sale, done in real time. "If it isn't in the CRM, it didn't happen" is the professional standard.

3. The CRM Data Model: Leads, Contacts, Accounts, Opportunities

This is the single most important conceptual section for an interview. If you can draw the CRM object model on a whiteboard, you instantly look experienced. There are four core objects.

Object	What it represents	Indian fintech example	Key point
Lead	An unqualified individual who *might* become a customer. Raw, unverified interest.	Someone downloads your "UPI reconciliation" whitepaper and leaves an email.	Leads are "not yet qualified." They get worked, then either converted or disqualified.
Contact	A qualified person you have a real relationship with, usually linked to a company.	Priya Sharma, Head of Finance at a Bengaluru D2C brand, a verified buyer.	A Contact is a person. Multiple contacts can belong to one account.
Account	A company / organisation you sell to.	"Zestco Retail Pvt Ltd" — the company, not any one person.	An Account holds many contacts and many deals over time.
Opportunity / Deal	A specific, qualified revenue possibility with a value and a close date.	"Zestco — Payment Gateway annual contract, ₹18,00,000, close Q3."	Opportunities move through pipeline stages and are what you forecast.

3.1 Lead vs Contact — the distinction that confuses everyone

In the classic Salesforce model, a **Lead** is a separate object that has not yet been qualified. Once you qualify it (confirm it is a real, relevant potential buyer), you **convert** the lead, which creates a Contact (the person), often an Account (their company), and optionally an Opportunity (a deal). HubSpot simplifies this: everyone is a "Contact" with a *lifecycle stage* property (Subscriber → Lead → MQL → SQL → Opportunity → Customer). Both models are trying to capture the same idea: a funnel from "cold, unverified" to "real buyer."

3.2 MQL vs SQL — marketing and sales handoff

Two acronyms you will hear constantly:

- **MQL (Marketing Qualified Lead)** — a lead that marketing believes is good enough to pass to sales, based on behaviour or fit (e.g., downloaded three resources, works at a target-size company).
- **SQL (Sales Qualified Lead / Sales Accepted Lead)** — a lead that sales has looked at, accepted, and is now actively working, because it meets qualification criteria (budget, need, timing).

The MQL → SQL conversion rate is a critical health metric for the marketing-sales relationship. A low rate means marketing is sending junk; sales rejects it; friction follows.

Lead-to-Opportunity Conversion Rate

= (Opportunities Created ÷ Leads Received) × 100

where: measured over the same cohort/time window. Track it per source (paid, referral, event) to see which channels produce real deals, not just clicks.

3.3 Qualification frameworks: BANT and MEDDIC

How do you decide a lead is "qualified"? You use a framework. Two are worth knowing:

- **BANT** — **B**udget, **A**uthority, **N**eed, **T**imeline. Simple, classic. Does the prospect have money, decision power, a real problem, and a timeframe?
- **MEDDIC** — **M**etrics, **E**conomic buyer, **D**ecision criteria, **D**ecision process, **I**dentify pain, **C**hampion. Used for larger, complex B2B deals; more rigorous.

> **Key takeaways:** Lead (unqualified person) → Contact (qualified person) inside an Account (company), with Opportunities (deals) attached. Qualify with BANT/MEDDIC. Know the MQL→SQL handoff. If you can whiteboard this model, you look senior.

4. Pipeline Stages, Probability & Activities

4.1 What a pipeline is

The **pipeline** is the ordered set of stages a deal passes through from first contact to closed. Each stage is a milestone in the buyer's journey, not just an internal task. A typical B2B SaaS pipeline in India might look like:

Stage	Meaning (buyer-defined milestone)	Typical win probability
1. Prospecting / Lead	Identified, not yet contacted or unengaged	5–10%
2. Qualification	Confirmed fit, budget, need, timeline (BANT passed)	15–20%
3. Discovery / Needs Analysis	Deep-dived pain, mapped requirements	25–30%
4. Demo / Proposal	Solution presented, value shown	40–50%
5. Negotiation	Pricing, terms, procurement, legal in motion	60–75%
6. Closed Won	Contract signed, deal booked	100%
—. Closed Lost	Prospect declined or went dark	0%

4.2 Stage probability and weighting

Each stage carries a historical **win probability** — the share of deals at that stage that eventually close. This is what lets you compute a **weighted pipeline** (covered in Section 7). Crucially, these probabilities should be *derived from your own historical conversion data*, not made up. This is exactly where your SQL/Excel skills shine: query closed deals, count how many were at each stage 90 days before close, and back out empirical probabilities.

4.3 Activities and logging

Activities are the atomic units of selling work logged against contacts and deals: calls, emails, meetings, demos, LinkedIn touches, tasks. They matter for two reasons: (1) they are the **leading indicators** that predict future pipeline; (2) they create the paper trail managers use to coach and to reconstruct a deal's history.

An SDR's day is measured heavily in activity metrics — dials, connects, emails sent, meetings booked. An AE's day is measured more in pipeline and closed revenue. Both must log.

> **Common trap — stages are not tasks.** A rookie mistake is treating stages as internal to-dos ("I sent the proposal, so we're in Proposal stage"). Correct discipline: a stage advances only when the *buyer* takes

a corresponding action (they agreed to a demo, they asked for pricing). Stage = buyer commitment, not seller effort. This keeps your probabilities honest and your forecast trustworthy.

5. Data Hygiene: The Discipline That Separates Pros from Amateurs

"Dirty CRM data" is the silent killer of sales orgs. Duplicate accounts, stale close dates, deals rotting in a stage for 90 days, missing amounts, contacts with no email — each is small, but together they make every report and forecast a lie. As the resident analyst, being the person who champions clean data will make you disproportionately valuable.

What good hygiene looks like

- **Deduplicate** accounts and contacts (one "Zestco Retail," not five spellings).
- **Update stages and close dates in real time** — never let a deal show a close date in the past.
- **Every open deal has a clear next step** with a date.
- **Fill required fields**: amount, stage, source, owner, close date.
- **Disqualify honestly** — mark dead leads Closed Lost with a reason, don't leave zombies inflating the pipeline.
- **Log activities the same day**, not from memory on Friday.

> Data Hygiene / Pipeline Freshness signal
>
> % of open deals with (a) a future close date AND (b) an activity in the last 14 days
>
> where: a healthy team keeps this high (say >80%). A low number means the pipeline is full of stale, untrustworthy deals — a red flag you can surface in a dashboard.

> **Common trap — vanity data volume.** A pipeline stuffed with 200 "opportunities" that nobody has touched in two months looks impressive and forecasts nothing. Quantity of records is not quality of pipeline. Clean, current, honestly-staged data beats a big dirty database every time.

6. The Major CRMs & How to Talk About Them

You do not need to be an admin of any of these. You need to (a) know what category each occupies, (b) be able to say a sentence about strengths, and (c) show you can learn any of them fast. Below is the cheat table.

CRM	Positioning	Sweet spot	One-line to say in interview
Salesforce	The market-leading enterprise CRM; hugely customisable, vast ecosystem (AppExchange).	Mid-market to large enterprise; complex processes.	"The industry standard — deeply customisable, and I'm working through Trailhead to get hands-on."
HubSpot	Inbound-marketing-born, all-in-one, famously easy to use; strong free tier.	SMB to mid-market; marketing + sales aligned.	"Great UX and reporting out of the box; I've done HubSpot Academy certifications."
Zoho CRM	Affordable, feature-rich, part of the broad Zoho suite; strong in India.	Indian SMBs, cost-sensitive teams.	"Very popular with Indian SMEs for value; flexible and integrates with the Zoho stack."
Freshsales (Freshworks)	Indian-origin (Chennai/SaaS unicorn), clean UI, built-in phone/email, AI (Freddy).	SMB to mid-market; sales teams wanting simplicity.	"An Indian product with strong built-in engagement features and a friendly UI."
Pipedrive	Lightweight, visual pipeline-first CRM built for salespeople.	Small sales teams who want a simple visual pipeline.	"Pipeline-centric and dead simple — great for small, fast-moving teams."

6.1 How to talk about a CRM you have not used

Interviewers know a fresher may not have enterprise Salesforce experience. What they are testing is your **mental model** and **learning agility**. A strong answer: *"I haven't administered Salesforce in production, but I understand the object model — leads, contacts, accounts, opportunities, activities — and the pipeline/forecasting logic, which transfers across any CRM. I've been doing Trailhead modules and I'm confident I can be productive in the tool within a week."* That reframes inexperience with a specific tool into competence with the underlying concepts.

6.2 Free ways to learn (do these — they are resume-worthy)

- **HubSpot Academy** — free certifications: "Sales Software," "Inbound Sales," "Sales Management," "Revenue Operations." Each is a badge you can list on LinkedIn and your resume. Highly credible for BD/SDR roles.
- **Salesforce Trailhead** — free, gamified, hands-on with a real Salesforce org. Earn badges and "Superbadges." Look for the "Sales Cloud Basics," "CRM for Lightning Experience," and admin beginner trails.
- **Zoho / Freshworks** — both offer free tiers and academies; spin up a free account and build a mock pipeline.

Key takeaways: Know the five names, their category and one differentiator each. For India, flag Zoho and Freshsales specifically. Emphasise that the CRM *concepts* transfer across tools, and back it up with free HubSpot Academy and Salesforce Trailhead badges — concrete proof of self-driven learning.

7. Pipeline Management & Forecasting

This is the heart of a BD/AE's analytical job. Forecasting is answering the question every leader asks: *"Are we going to hit the number this quarter?"* Get comfortable here and you become indispensable.

7.1 Pipeline coverage ratio

The single most important pipeline-health number. It asks: do we have enough pipeline to hit quota, given that most deals won't close?

> Pipeline Coverage Ratio
>
> = Total Open Pipeline Value ÷ Sales Target (Quota) for the period
>
> where: rule of thumb is 3× to 4× coverage. If your win rate is ~30%, you need roughly 3× your quota in open pipeline to have a realistic shot at hitting it.

Worked Example 1 — Pipeline Coverage

Aditya is an AE at a Mumbai payments-SaaS firm. His quarterly quota is **₹60,00,000** in new ARR. His open pipeline totals **₹1,80,00,000** across all stages.

Coverage = ₹1,80,00,000 ÷ ₹60,00,000 = **3.0×**.

His historical win rate is ~33%. Expected closes ≈ 33% × ₹1,80,00,000 = ₹59,40,000 — just under quota. **Verdict:** 3× is the bare minimum; he is on the edge. If his win rate were only 25%, 3× coverage would yield just ₹45,00,000 and he'd need ~4× (₹2,40,00,000) of pipeline. **Action:** generate more pipeline now, don't wait.

7.2 Weighted vs unweighted pipeline

Unweighted pipeline is the raw sum of all open deal values — optimistic, treats a ₹10L deal in Qualification the same as one in Negotiation. **Weighted pipeline** multiplies each deal by its stage probability, giving a risk-adjusted, more realistic figure.

> Weighted Pipeline (Expected Value)
>
> = Σ (Deal Value × Stage Win Probability)
>
> where: Σ means sum across all open deals. Also called "expected pipeline." It is a forecasting input, not the forecast itself.

Worked Example 2 — Weighted Pipeline

Aditya's four open deals:

Deal	Value (₹)	Stage	Probability	Weighted (₹)
Zestco	18,00,000	Negotiation	70%	12,60,000
NimbusPay	40,00,000	Proposal	45%	18,00,000
Kirana360	72,00,000	Discovery	25%	18,00,000
FinLoop	50,00,000	Qualification	20%	10,00,000

Unweighted pipeline = 18 + 40 + 72 + 50 = **₹1,80,00,000**.
Weighted pipeline = 12.6 + 18 + 18 + 10 = **₹58,60,000**.

Insight: the raw ₹1.8 crore looks like triple his ₹60L quota, but risk-adjusted he's expected to land ₹58.6L — just short. This gap between unweighted and weighted is exactly what a forecast conversation is about.

7.3 Forecast categories: Commit, Best Case, Pipeline

Beyond math, sales teams tag each deal with a **forecast category** reflecting the rep's judgement, not just the stage probability:

Category	Meaning	Confidence
Closed	Already won and booked this period.	100%
Commit	Rep is confident it will close this period; would "put their name on it."	~90%+
Best Case (Upside)	Could close if things go well; not banked on.	~50%
Pipeline	Real deals, but not expected to close this specific period.	low
Omitted / Lost	Excluded from the forecast.	0%

7.4 Three forecasting methods

- **Stage-weighted (pipeline) forecast** — sum of (deal value × stage probability). Objective, data-driven, but can be skewed by a couple of huge early-stage deals. Best when you have many deals.
- **Category / commit forecast** — the rep/manager's judgemental roll-up (Closed + Commit, plus a slice of Best Case). Captures human insight the model misses; risk is bias (sandbagging or happy ears).
- **Historical / run-rate forecast** — extrapolate from past performance (e.g., "last 4 quarters we closed ~₹55L on average, growing 10%"). Good sanity check.

Serious orgs triangulate: they compare the weighted-pipeline number, the commit roll-up, and the historical run-rate, and investigate when they diverge.

Worked Example 3 — Building a Quarter Forecast

Continuing with Aditya. Already Closed Won this quarter: **₹22,00,000**. His judgemental tags on the four open deals:

- Zestco (₹18L) — **Commit** (contract in legal).
- NimbusPay (₹40L) — **Best Case** (champion strong, budget unconfirmed).
- Kirana360 (₹72L) — **Pipeline** (won't close this quarter).
- FinLoop (₹50L) — **Best Case**.

Method A — Category forecast: Closed 22 + Commit 18 = **₹40,00,000** committed. Add a conservative 50% of Best Case (0.5 × [40+50] = 45) → forecast range **₹40L (floor) to ₹85L (ceiling)**, most-likely call ≈ **₹62L**.

Method B — Weighted pipeline + closed: 22 (closed) + 58.6 (weighted open, from Ex. 2) = **₹80.6L** — but note Kirana360's ₹18L weighted contribution is a Discovery deal he says won't close this quarter, so strip it: 80.6 − 18 = **₹62.6L**.

Reconciliation: both methods land near **₹62L** against a ₹60L quota. That agreement gives confidence. The honest call to his manager: "I'm forecasting ₹62L, with a firm floor of ₹40L already in Commit-or-better."

Common trap — sandbagging & happy ears. Two opposite forecasting sins. **Sandbagging** = deliberately lowballing your forecast (hiding deals in "Pipeline") so you can "beat" it and look like a hero — it destroys planning and trust. **Happy ears** = optimistically calling deals as Commit that aren't real, blowing the forecast. Good forecasting is honest and calibrated: your Commit deals should close ~90% of the time, and you should track your own forecast accuracy over time.

Forecast Accuracy

$$= 1 - |\text{Forecast} - \text{Actual}| \div \text{Actual}$$

where: expressed as a %. If you forecast ₹62L and land ₹58L, accuracy = 1 − |62−58|/58 = 1 − 6.9% = 93.1%. Tracking your own accuracy each quarter is a pro move that impresses managers.

Key takeaways: Coverage ratio (aim 3–4×) tells you if there's enough pipeline; weighted pipeline risk-adjusts it; forecast categories (Commit/Best Case/Pipeline) add human judgement; triangulate methods and reconcile. Be the rep whose forecast is boringly accurate.

8. The Core Sales & SaaS Metrics

This is the section to memorise cold. Every metric below is interview fair game. Learn the formula, what it measures, a healthy benchmark, and one trap. Note the units carefully — mixing monthly and annual, or gross and net, is the most common error.

8.1 Revenue metrics: MRR, ARR, ACV, TCV

MRR — Monthly Recurring Revenue

MRR = Σ (monthly subscription value of all active customers)

Also: MRR = Number of customers × ARPA (avg revenue per account per month)

where: only **recurring** subscription revenue counts — exclude one-time setup fees, hardware, or professional services.

ARR — Annual Recurring Revenue

ARR = MRR × 12

where: ARR is the annualised run-rate of recurring revenue. For annual contracts, ARR = total annual contract value of recurring components.

ACV & TCV

ACV (Annual Contract Value) = contract value per year

TCV (Total Contract Value) = ACV × contract length + one-time fees

where: a 3-year ₹15L/yr deal with ₹2L setup has ACV = ₹15L, TCV = (15×3)+2 = ₹47L. ACV is used for annualised comparisons; TCV for total deal size.

Common trap — confusing MRR and ARR (and mixing in non-recurring revenue). ARR is simply MRR × 12; never quote a monthly figure as if it were annual or vice-versa. And never inflate MRR/ARR with one-time implementation fees or services revenue — investors and CFOs will catch it instantly, and it makes every downstream metric (growth rate, retention) wrong.

MRR movement (the growth story)

MRR changes each month via five components — know these; they show up in board decks:

- **New MRR** — from brand-new customers.
- **Expansion MRR** — upsell/cross-sell to existing customers (upgrades, more seats).
- **Reactivation MRR** — churned customers who came back.
- **Contraction MRR** — existing customers who downgraded (negative).
- **Churned MRR** — customers who cancelled (negative).

Net New MRR

= New + Expansion + Reactivation − Contraction − Churn

where: this is the month-over-month change in MRR. Positive and accelerating is the goal.

8.2 CAC — Customer Acquisition Cost

CAC — Customer Acquisition Cost

CAC = Total Sales & Marketing Spend in a period ÷ Number of New Customers acquired in that period

where: include salaries, commissions, ad spend, tools, and overhead for sales & marketing. Fully-loaded CAC is the honest version; "ad-spend-only CAC" understates the truth.

Worked Example 4 — CAC

A Pune B2B SaaS spent, in Q1: sales salaries ₹24,00,000, marketing salaries ₹9,00,000, ad spend ₹12,00,000, tools ₹3,00,000. Total S&M = **₹48,00,000**. They acquired **40 new customers**.

CAC = ₹48,00,000 ÷ 40 = **₹1,20,000 per customer**.

Interpretation: it costs ₹1.2L fully-loaded to win one customer. Whether that's good depends entirely on what a customer is worth (LTV) and how fast you recover it (payback) — see below.

8.3 LTV — Customer Lifetime Value

LTV (a.k.a. CLV / CLTV)
LTV = (ARPA × Gross Margin %) ÷ Customer Churn Rate
Simpler form: LTV = Average Revenue per Customer per period × Average Customer Lifetime (in periods) × Gross Margin %

where: ARPA = avg revenue per account per period; Churn Rate is per the same period; Average Lifetime (in periods) = 1 ÷ churn rate. Always apply gross margin — LTV should be on profit, not revenue.

Worked Example 5 — LTV

Same SaaS: ARPA = **₹8,000/month**, gross margin = **80%**, monthly customer churn = **2%**.

Average lifetime = 1 ÷ 0.02 = **50 months**.
LTV = ₹8,000 × 80% × 50 = **₹3,20,000 per customer**.
(Equivalently: (8,000 × 0.80) ÷ 0.02 = 6,400 ÷ 0.02 = ₹3,20,000.)

Interpretation: each customer generates ₹3.2L of gross profit over their lifetime. Note how sensitive this is to churn: if churn were 4% not 2%, lifetime halves to 25 months and LTV drops to ₹1.6L. **Churn is the master variable.**

8.4 LTV:CAC ratio — the unit-economics headline

LTV : CAC Ratio
= LTV ÷ CAC

where: benchmark ≈ **3:1 or higher** is healthy for SaaS. Below 1:1 you lose money on every customer. Above ~5:1 can signal you are *under-investing* in growth (leaving market share on the table).

Worked Example 6 — LTV:CAC

From Ex. 4 and 5: LTV = ₹3,20,000, CAC = ₹1,20,000.

LTV:CAC = 3,20,000 ÷ 1,20,000 = **2.67 : 1**.

Interpretation: below the 3:1 gold standard but above the 1:1 danger line — the business is viable but not yet efficient. Two levers to improve it: cut CAC (better targeting, more inbound, higher conversion) or raise LTV (reduce churn, upsell, raise prices). Because LTV is so churn-sensitive, cutting churn from 2% to 1.5% would push LTV to ₹4.27L and the ratio to 3.6:1 — the highest-leverage fix.

8.5 CAC Payback Period

CAC Payback Period

= CAC ÷ (ARPA per month × Gross Margin %)

where: answer is in months — how long until a customer's gross profit repays their acquisition cost. Benchmark: <12 months is good for SMB SaaS; 12–18 acceptable; enterprise can stretch longer.

Worked Example 7 — CAC Payback

CAC = ₹1,20,000; ARPA = ₹8,000/mo; gross margin = 80%.

Monthly gross profit per customer = 8,000 × 0.80 = ₹6,400.
Payback = 1,20,000 ÷ 6,400 = **18.75 months**.

Interpretation: nearly 19 months to recover CAC — on the long side. Since average lifetime is 50 months, the customer is still profitable overall, but the company ties up cash for a year and a half per customer. This is why cash-strapped startups obsess over payback: shorter payback = less funding needed to grow.

8.6 Churn: logo vs revenue, gross vs net

Logo (Customer) Churn Rate

= Customers Lost in period ÷ Customers at Start of period × 100

where: counts *accounts*, treating a ₹5,000 and a ₹5,00,000 customer equally.

Gross Revenue Churn Rate

= MRR Lost (churn + contraction) in period ÷ MRR at Start × 100

where: weights by money. Cannot be negative. This is the "leaky bucket" rate before any expansion.

Worked Example 8 — Churn

Start of month: 500 customers, MRR ₹40,00,000. During the month you lose 15 customers worth ₹1,80,000 of MRR, and existing customers downgrade by ₹40,000, while others upgrade (expansion) by ₹2,60,000.

Logo churn = 15 ÷ 500 = **3.0%**.
Gross revenue churn = (1,80,000 + 40,000) ÷ 40,00,000 = 2,20,000 ÷ 40,00,000 = **5.5%**.
Net revenue churn = (2,20,000 − 2,60,000) ÷ 40,00,000 = −40,000 ÷ 40,00,000 = **−1.0%** (i.e., negative churn — expansion outran losses).

Insight: logo churn (3%) < revenue churn (5.5%) means you lost slightly larger-than-average accounts. But *net* churn is negative, the holy grail — the existing base grew even before new sales.

8.7 NRR and GRR — retention as a growth engine

NRR — Net Revenue Retention

= (Starting MRR + Expansion − Contraction − Churn) ÷ Starting MRR × 100

where: measured on an existing cohort, excluding brand-new customers. **>100% is excellent** — the base grows on its own. Best-in-class SaaS: 120%+. This is the single metric investors scrutinise most.

GRR — Gross Revenue Retention

= (Starting MRR − Contraction − Churn) ÷ Starting MRR × 100

where: expansion is *excluded*, so GRR ≤ 100% always. It shows how "leaky" the base is. 90%+ is strong for SMB; 95%+ for enterprise.

Worked Example 9 — NRR & GRR

Cohort starting MRR = ₹40,00,000. Over 12 months: expansion +₹9,00,000, contraction −₹1,50,000, churn −₹3,50,000.

NRR = (40 + 9 − 1.5 − 3.5) ÷ 40 = 44 ÷ 40 = **110%**.
GRR = (40 − 1.5 − 3.5) ÷ 40 = 35 ÷ 40 = **87.5%**.

Interpretation: 110% NRR means this cohort spends 10% more a year later despite some losses — expansion is the engine. 87.5% GRR says the underlying leak (12.5%/yr lost) is a bit high; there's a retention problem masked by strong upselling. A sharp analyst reports *both* so the leak isn't hidden.

Common trap — NRR hiding a churn problem. A shiny 115% NRR can conceal a leaky bucket if a few whale accounts are expanding fast while many small accounts churn. Always pair NRR with GRR (and logo churn). If NRR is great but GRR is weak, you are papering over churn with upsell — unsustainable if expansion slows.

8.8 Sales-effectiveness metrics: win rate, deal size, cycle length, quota attainment

Win Rate

= Deals Won ÷ (Deals Won + Deals Lost) × 100   [close/won rate]
or = Deals Won ÷ Total Opportunities Created × 100   [opportunity win rate]

where: be explicit which denominator you mean — decided deals (won+lost) vs all opportunities. They differ a lot; state your definition.

Average Deal Size (ASP)

= Total Revenue Won in period ÷ Number of Deals Won

where: watch mix — a few big deals skew the average; also report the median.

Sales Cycle Length

= Average number of days from Opportunity Created (or first contact) to Closed Won

where: shorter is more capital-efficient. Segment by deal size — enterprise cycles are far longer than SMB.

Quota Attainment

= Actual Sales Achieved ÷ Quota (Target) × 100

where: per rep or team. Healthy orgs see the majority of reps at or above ~80–100%. If almost no one hits quota, the quota (or the model) is broken.

Worked Example 10 — Win Rate & Sales Velocity

Last quarter Neha created 50 opportunities; 12 are Won, 18 Lost, 20 still open.

Close/won rate = 12 ÷ (12+18) = 12/30 = **40%**.
Opportunity win rate = 12 ÷ 50 = **24%**.

Now **Sales Velocity** — how much revenue the pipeline generates per day:

> Sales Velocity
> = (Number of Opportunities × Win Rate × Average Deal Size) ÷ Sales Cycle Length (days)

With 20 active opps, 40% win rate, ₹4,00,000 avg deal, 60-day cycle:
Velocity = (20 × 0.40 × 4,00,000) ÷ 60 = 32,00,000 ÷ 60 = ≈ **₹53,333 per day**.

Why it matters: the four levers of the formula are literally the four ways to sell more — more opps, higher win rate, bigger deals, or faster cycles. Improving any one improves velocity. This is a fantastic interview framework to recite.

8.9 Conversion rates by stage (funnel math)

> Stage Conversion Rate
> = Deals advancing to next stage ÷ Deals entering the stage × 100
>
> where: chain these to model the full funnel. The product of all stage conversions = overall lead-to-customer rate.

Worked Example 11 — Funnel Conversion

A month's funnel: 1,000 Leads → 300 MQL → 120 SQL → 60 Opportunities → 18 Closed Won.

Transition	Rate
Lead → MQL	300/1000 = 30%
MQL → SQL	120/300 = 40%
SQL → Opp	60/120 = 50%
Opp → Won	18/60 = 30%
Lead → Won (overall)	**18/1000 = 1.8%**

Use: to hit 30 deals next month you'd need ~1,667 leads at these rates. And the biggest leak (lowest rate) is where coaching/effort pays off most — here the Opp→Won 30% and Lead→MQL 30% are the constraints. This is how a data-driven SDR reasons about their own targets.

8.10 Activity metrics (leading indicators)

These are the inputs an SDR/BDE controls directly. They predict tomorrow's pipeline:

- **Dials / calls made**, **connect rate** (conversations ÷ dials), **emails sent**, **open/reply rates**.
- **Meetings booked** and **meetings held (show rate)** — the SDR's core output.
- **Opportunities created** — the handoff metric to AEs.

Connect Rate = Conversations ÷ Dials × 100 | **Meeting Show Rate** = Meetings Held ÷ Meetings Booked × 100

where: these diagnose *where* outbound breaks — low connect rate = list/timing problem; low show rate = qualification/confirmation problem.

8.11 Efficiency metrics (brief): Magic Number & Rule of 40

SaaS Magic Number

= (Net New ARR in a quarter × 4, i.e. annualised) ÷ Sales & Marketing spend in the prior quarter

where: measures S&M efficiency. **>0.75 (roughly 1.0)** = efficient, invest more; <0.5 = inefficient, fix before scaling.

Rule of 40

Revenue Growth Rate (%) + Profit Margin (%) ≥ 40

where: a health check for SaaS balancing growth and profitability. A company growing 30% at 15% margin scores 45 — healthy. Growing 60% at −25% can still pass (35 fails, 60−15=45 passes). It's a trade-off frame, not a precise law.

8.12 Metric summary table — what each one measures

Metric	Measures	Healthy benchmark
MRR / ARR	Recurring revenue scale	Growing MoM/YoY
CAC	Cost to win a customer	Lower is better (context-dependent)
LTV	Lifetime gross profit per customer	Higher; churn-driven
LTV:CAC	Return on acquisition spend	≥ 3:1
CAC Payback	Months to recover CAC	< 12–18 months
Logo Churn	% of customers lost	SMB <3–5%/mo; enterprise <1%
GRR	Revenue kept (no expansion)	90%+ (95%+ enterprise)
NRR	Revenue kept + expanded	>100%; 120%+ elite
Win Rate	Deals won vs decided	~20–40% (varies)
Sales Cycle	Days to close	Shorter is better
Quota Attainment	Actual vs target	Majority of reps ≥80%
Pipeline Coverage	Pipeline vs quota	3–4×
Magic Number	S&M efficiency	>0.75
Rule of 40	Growth + profit balance	≥ 40

Common trap — vanity metrics. Impressive-sounding numbers that don't drive decisions: total emails sent, LinkedIn impressions, "leads in the database," website visits, gross pipeline with no weighting. They feel like progress. The test: does the metric change what you *do*? If not, it's vanity. Pair every activity metric with an outcome metric (emails → replies → meetings → opportunities → revenue).

Key takeaways: Memorise CAC, LTV, LTV:CAC, payback, churn (logo vs revenue), GRR/NRR, win rate, cycle, quota attainment, coverage. Watch units (monthly vs annual, gross vs net, revenue vs profit). Churn is the master variable behind LTV. Report NRR with GRR so churn can't hide.

9. Unit Economics: Why the Numbers Drive Hiring & Quota

Unit economics = the profit-and-loss of a *single customer*. If one customer is profitable (LTV > CAC by a healthy multiple, with reasonable payback), then scaling customers scales profit; if one customer loses money, scaling just loses money faster. Everything in a sales org — how many reps to hire, how big quotas are, how much to spend on marketing — flows from this.

9.1 From unit economics to quota

Here is the logic chain a VP of Sales actually runs, and why your grasp of it will impress:

1. **Company needs** ₹12 crore new ARR next year (board target).
2. A ramped AE reliably closes **₹1.5 crore** of new ARR/year → **quota** = ₹1.5 cr.
3. To get ₹12 cr you need ~8 productive AEs (₹12cr ÷ ₹1.5cr). Accounting for ramp and attrition, hire ~11.
4. Each AE needs 3–4× pipeline coverage → ₹4.5–6 cr pipeline each → the org needs ₹40–50 cr total pipeline.
5. Given lead→opp and opp→win conversion rates, that pipeline requires X leads → sets the **marketing budget** and **SDR headcount**.
6. All of this only makes sense if **LTV:CAC ≥ 3** and payback is acceptable — otherwise growth burns cash unsustainably.

When an interviewer asks "why do we track LTV:CAC?", the killer answer connects it to this chain: *"Because it decides whether we should hire more reps and spend more on marketing at all. If unit economics are healthy, every rupee into sales returns more than a rupee, so we scale. If they're broken, hiring just accelerates losses."*

9.2 Fintech nuance

In Indian fintech, unit economics carry extra weight: interchange caps and the zero-MDR regime on UPI mean revenue-per-transaction can be thin, so **volume, retention and cross-sell (NRR)** matter enormously. A lending or wealth-tech firm also layers in credit risk and AUM economics. Your NISM Series-XV (Research Analyst) background lets you speak credibly about how a fintech actually earns — a genuine edge over a generic sales candidate.

> **Key takeaways:** Unit economics (per-customer P&L) drive every macro decision — headcount, quota, marketing budget. Quota is derived top-down from targets and bottom-up from rep capacity. LTV:CAC and payback are the gatekeepers of "should we scale?" Tie your metrics answers back to these decisions and you sound like an operator, not a fresher.

10. The Modern Sales Tech Stack

Modern B2B sales runs on a layered software stack. You should be able to name a category, its job, and 1–2 example tools. Do not claim deep expertise — claim conceptual fluency and eagerness to learn.

Layer	Job	Example tools	Who uses it
CRM (system of record)	Store accounts, contacts, deals, forecast	Salesforce, HubSpot, Zoho, Freshsales	Everyone
Sales Engagement / Cadence	Automate multi-touch outbound sequences (email + call + LinkedIn), track opens/replies	Outreach, Salesloft, Apollo.io, HubSpot Sequences	SDRs, AEs
Data / Enrichment / Prospecting	Find and verify contacts, emails, phone numbers, firmographics	ZoomInfo, Apollo.io, Lusha, Clearbit	SDRs, RevOps
Conversation Intelligence	Record, transcribe and analyse sales calls; coach; surface risk	Gong, Chorus (ZoomInfo), Fireflies	AEs, managers
Scheduling	Frictionless meeting booking, kill email ping-pong	Calendly, Chili Piper	SDRs, AEs
E-Signature / CPQ	Quotes, contracts, e-signing, deal closing	DocuSign, PandaDoc, Salesforce CPQ	AEs, ops
Social / Research	Prospect research, warm outreach, social selling	LinkedIn Sales Navigator	SDRs, AEs
Analytics / BI	Dashboards, forecasting, cohort/funnel analysis	Power BI, Tableau, Looker, native CRM reports	RevOps, you

10.1 How the stack works together (a day in the life)

An SDR pulls a target list from **Sales Navigator**, enriches contact details in **Apollo/ZoomInfo**, loads them into a **Salesloft/Outreach** sequence, and works the cadence. Booked meetings are set via **Calendly**, logged automatically to the **CRM**. The AE runs the demo on Zoom, recorded by **Gong**, which flags whether next steps were set. The proposal goes out via **DocuSign/PandaDoc**. Every step syncs to the CRM, and **Power BI** dashboards read the CRM to report funnel health. Notice: the **CRM is the hub**; everything else feeds it.

10.2 Note on Apollo

Apollo.io appears in three categories (data, engagement, and light CRM) because it bundles them — a popular, affordable all-in-one for startups and a great tool to learn on a free plan to demonstrate hands-on prospecting skill.

> **Key takeaways:** CRM = hub; around it sit engagement (Outreach/Salesloft/Apollo), data/enrichment (ZoomInfo/Lusha/Apollo), conversation intelligence (Gong/Chorus), scheduling (Calendly), e-sign (DocuSign), and research (Sales Navigator). Claim conceptual fluency, name examples, show eagerness. Apollo is a good free tool to get real hands-on reps.

11. Sales Analytics & Reporting: Your Differentiator

This is where you win the job. Most SDR/BDE candidates cannot build a dashboard or run a cohort analysis. You can. This section is your pitch.

11.1 Leading vs lagging indicators

The most important analytical distinction in sales:

Type	Definition	Examples	Why it matters
Leading	Inputs you control now that predict future outcomes	Calls, emails, meetings booked, new opps created, pipeline generated	You can act on them *today* to change tomorrow's revenue
Lagging	Outcomes that report what already happened	Revenue closed, quota attainment, win rate, churn	Report cards — too late to change, but validate the leading indicators

The pro insight: *you manage leading indicators to move lagging ones*. If revenue (lagging) is down, you diagnose which leading indicator broke — fewer meetings? lower conversion? — and fix that. Reciting this framing signals real understanding.

11.2 Dashboards a BDE/SDR should build

- **Activity dashboard** — dials, connects, emails, meetings booked/held vs targets (daily/weekly).
- **Funnel/conversion dashboard** — volume and conversion rate at each stage; spot the leak.
- **Pipeline dashboard** — open pipeline by stage, weighted pipeline, coverage ratio vs quota, deals with no recent activity (hygiene).
- **Forecast dashboard** — Commit/Best Case/Pipeline roll-up, forecast vs actual, forecast accuracy trend.
- **Cohort retention dashboard** — NRR/GRR by signup cohort over time.
- **Source/channel ROI** — CAC and win rate by lead source, so spend goes where deals come from.

11.3 Cohort and funnel analysis

Cohort analysis groups customers by a shared start point (e.g., "customers who signed up in Jan-2026") and tracks a metric (retention, revenue) over their lifetime. It answers "is retention getting better for newer cohorts?" — invisible in a single blended average. **Funnel analysis** tracks how a cohort of leads progresses through stages, exposing where and why they drop.

11.4 How YOUR skills apply (the pitch to rehearse)

- **SQL** — query the CRM/warehouse directly: "SELECT stage, COUNT(*), AVG(amount) FROM opportunities GROUP BY stage" to build empirical stage probabilities; pull raw win/loss for analysis.
- **Excel** — pivot tables for funnel and cohort views, weighted-pipeline models, forecast scenarios, quota tracking. A clean Excel forecast model is a fantastic interview leave-behind.
- **Power BI / Tableau** — live dashboards refreshing from the CRM; DAX measures for coverage, win rate, NRR; drill-downs by rep/region/source.
- **Python (pandas)** — automate messy data cleaning (dedup, standardise), cohort retention curves, and even simple predictive lead-scoring or churn models. Mentioning a logistic-regression lead score signals you can go beyond reporting.

- **Finance/NISM lens** — you naturally think in unit economics, margins and NPV, so LTV, payback and Rule of 40 come intuitively.

Worked Example 12 — A Forecast/Coverage model you could build in Excel

Columns: Deal | Amount | Stage | Stage Probability (VLOOKUP to a probability table) | Weighted (= Amount × Prob) | Forecast Category | Close Month. Then:

- SUM(Amount) = unweighted pipeline; SUM(Weighted) = weighted pipeline.
- Weighted ÷ Quota = coverage ratio (conditional-format red if <3).
- SUMIF(Category,"Commit",Amount) + Closed = forecast floor.
- A pivot of Weighted by Close Month = a monthly forecast curve.
- Flag rows where TODAY() – LastActivityDate > 14 → the hygiene report.

Walking an interviewer through this model, live or on paper, is the single most differentiating thing a quant-minded BD candidate can do.

Key takeaways: Manage leading indicators to move lagging ones. Build activity, funnel, pipeline, forecast, cohort and source-ROI dashboards. Your SQL/Excel/Power BI/Python turn you from a rep who reports numbers into one who *engineers* them — lead with this in every interview.

12. Sales Operations & Process

Sales Ops (part of RevOps) is the function that makes selling scalable and measurable — process, tools, data, territories, quotas, comp. As an analytically strong hire, this is a natural adjacency and often a faster-growth path than pure quota-carrying.

12.1 Lead routing and lead scoring

Lead routing = the rules that assign incoming leads to the right rep automatically (by geography, company size, industry, or round-robin). Good routing means fast follow-up; slow routing kills conversion. **Lead scoring** ranks leads by likelihood/value to convert, using fit (firmographics) and behaviour (engagement) — often a points model or an ML model. You could build one in Python.

12.2 SLAs between marketing and sales

An **SLA (Service Level Agreement)** here is a mutual promise: marketing commits to deliver *N* qualified leads of a defined quality per month; sales commits to follow up on each within, say, **5 minutes to 1 hour** and to work it a minimum number of times before disqualifying. Speed-to-lead is decisive — response within 5 minutes can lift qualification dramatically versus an hour later. SLAs reduce the classic "marketing leads are junk / sales ignores our leads" finger-pointing.

12.3 Territories, quotas and playbooks

- **Territory** — the slice of the market a rep owns (geography, industry vertical, account tier, or named accounts). Good territory design balances opportunity fairly so quotas are attainable.
- **Quota** — the target a rep must hit (see Section 9 for how it's derived). Often ~4–5× the rep's on-target earnings as a rule of thumb.
- **OTE (On-Target Earnings)** — base + variable (commission) a rep earns at 100% quota; a core comp concept to know.
- **Playbook** — the documented, repeatable "how we sell": ICP definition, qualification criteria, discovery questions, objection handling, email templates, stage exit criteria. Playbooks turn one rep's success into the whole team's.
- **ICP (Ideal Customer Profile)** — the precise definition of the best-fit customer (size, industry, geography, pain). Selling to ICP raises win rate and lowers CAC.

> **Common trap — chasing non-ICP deals.** Freshers often celebrate any logo. But deals outside the ICP close slower, churn faster, and consume time you could spend on high-fit prospects — hurting win rate, cycle and CAC. Disciplined qualification against the ICP is a metric-improving behaviour, not gatekeeping.

> **Key takeaways:** Sales Ops = process + data + tooling that scales selling. Know lead routing, lead scoring, marketing↔sales SLAs (speed-to-lead matters), territories, quota/OTE, ICP and playbooks. Your analytics skills make Sales Ops a strong, fast-growth landing spot.

13. Interview-Ready: 35+ Q&A with Model Answers

Rehearse these aloud. Answers are compact on purpose — expand with your own examples. Numbers assume the fictional fintech-SaaS context used throughout.

CRM & data model

Q1. What is a CRM and why does a sales team need one?

A: A CRM is the single source of truth for customers and deals — a database of accounts, contacts, opportunities and activities. It matters because it turns private, in-the-head selling into a shared, queryable asset: it enables forecasting, accountability, coaching, and continuity when reps leave. Simply: "if it isn't in the CRM, it didn't happen."

Q2. Explain the difference between a lead, a contact, an account and an opportunity.

A: A lead is an unqualified individual who might buy. A contact is a qualified person you have a real relationship with. An account is the company they belong to. An opportunity is a specific deal — a value and a close date — attached to an account and moving through pipeline stages. Leads convert into contacts (plus account, plus optionally an opportunity).

Q3. What's the difference between an MQL and an SQL?

A: An MQL is Marketing Qualified — marketing judges it good enough to pass to sales based on fit/behaviour. An SQL is Sales Qualified/Accepted — sales has vetted it and is actively working it. The MQL→SQL conversion rate is a key health check on lead quality and the marketing-sales handoff.

Q4. How do you qualify a lead?

A: With a framework like BANT — Budget, Authority, Need, Timeline — or MEDDIC for complex deals (Metrics, Economic buyer, Decision criteria, Decision process, Identify pain, Champion). I confirm there's a real problem, money, a decision-maker involved, and a timeframe before I invest heavily.

Q5. Why is CRM data hygiene important, and how do you maintain it?

A: Dirty data — duplicates, stale close dates, zombie deals — makes every report and forecast a lie. I keep it clean by logging activities the same day, updating stages/close dates in real time, ensuring every open deal has a dated next step, filling required fields, and honestly marking dead deals Closed Lost with a reason.

Q6. Which CRMs are you familiar with?

A: I understand the object model and pipeline logic that's common to all of them. Salesforce is the enterprise standard; HubSpot is the easy all-in-one; in India, Zoho and Freshsales (an Indian product) are very popular with SMBs. I'm doing HubSpot Academy and Salesforce Trailhead to get hands-on, and I'm confident I can be productive in any CRM within a week because the concepts transfer.

Q7. You've never used Salesforce. Is that a problem?

A: The tool is learnable in days; the concepts are what matter and I have them — leads, contacts, accounts, opportunities, activities, stages, forecasting. I'm actively doing Trailhead. I'd rather you hired someone who understands why the pipeline behaves the way it does than someone who can click around Salesforce but can't read a funnel.

Pipeline & forecasting

Q8. How do you prioritise your pipeline?

A: I score deals on three axes: (1) probability — stage and real buyer signals; (2) value — deal size and strategic fit to our ICP; (3) timing — close date and momentum. I focus first on high-value deals that are late-stage and moving, keep mid-stage deals progressing with clear next steps, and ruthlessly disqualify stale, non-ICP deals so they don't eat time or pollute the forecast. Concretely I'd sort by weighted value and recency of activity and work top-down.

Q9. What is pipeline coverage and what ratio is healthy?

A: Coverage = open pipeline value ÷ quota. Rule of thumb is 3–4×, because most deals don't close. If my win rate is ~30% I need roughly 3× to have a realistic shot; if it's 25% I need closer to 4×. Below 3× I'd flag a pipeline-generation problem early rather than miss quota silently.

Q10. What's the difference between weighted and unweighted pipeline?

A: Unweighted is the raw sum of open deal values — optimistic. Weighted multiplies each deal by its stage win probability to give a risk-adjusted expected value. A ₹1.8 crore unweighted pipeline might be only ₹58 lakh weighted, and the weighted number is a far better forecasting input.

Q11. How would you forecast your quarter?

A: I triangulate three ways. First, a category roll-up: Closed + Commit gives my floor, plus a conservative share (say 50%) of Best Case. Second, a stage-weighted pipeline number, stripping out deals I know won't close this period. Third, a historical run-rate as a sanity check. When they agree I'm confident; when they diverge I dig in. I'd commit a specific number with a floor — e.g., "₹62 lakh forecast, ₹40 lakh already in Commit-or-better" — and I track my own forecast accuracy each quarter.

Q12. What are Commit, Best Case and Pipeline categories?

A: They're confidence tags on deals. Commit — I'd put my name on it closing this period (~90%+). Best Case — could close if things go well (~50%), upside not banked. Pipeline — real deals, but not expected this period. Closed is already won. They add human judgement on top of mechanical stage probabilities.

Q13. What is sandbagging and why is it bad?

A: Sandbagging is deliberately lowballing your forecast — hiding deals in "Pipeline" — so you beat it and look good. It breaks capacity and cash planning as badly as over-forecasting does, and erodes trust. The goal is a calibrated, honest forecast, not a beatable one. The opposite sin, "happy ears," is over-calling weak deals as Commit.

Q14. A deal has been stuck in one stage for 60 days. What do you do?

A: First, is it real? I'd re-test qualification — is there still budget, authority, need, timeline? If yes, diagnose the specific blocker (no economic buyer engaged, no clear next step, a competing priority) and create a concrete dated next step to advance it. If it's gone dark and fails re-qualification, I'd mark it Closed Lost with a reason rather than let it inflate the pipeline. Zombie deals distort the forecast.

Metrics & unit economics

Q15. Which metrics do you track, and why those?

A: I'd separate leading from lagging. Leading (I control today): activities — calls, meetings booked/held, new opps created, pipeline generated. Lagging (outcomes): win rate, average deal size, sales cycle, quota attainment, and at the business level CAC, LTV, LTV:CAC, payback, and NRR/GRR. I track leading indicators to move the lagging ones — if revenue dips, I find which input broke and fix that.

Q16. Define CAC and how you'd calculate it properly.

A: Customer Acquisition Cost = total sales & marketing spend in a period ÷ new customers acquired. "Properly" means fully-loaded — salaries, commissions, ad spend, tools, overhead — not just ad spend, which understates it. E.g., ₹48 lakh S&M ÷ 40 customers = ₹1.2 lakh CAC.

Q17. How do you calculate LTV?

A: LTV = (ARPA × gross margin %) ÷ churn rate, or equivalently average revenue × average lifetime × gross margin. Lifetime = 1 ÷ churn. Example: ₹8,000/mo ARPA, 80% margin, 2% monthly churn → lifetime 50 months → LTV = 8,000 × 0.8 × 50 = ₹3.2 lakh. Always apply gross margin so it's profit-based, and note churn is the dominant lever.

Q18. What's a good LTV:CAC ratio and why?

A: Around 3:1 or higher is healthy. Below 1:1 you lose money per customer. Interestingly, much above ~5:1 can mean you're under-investing in growth and leaving the market to competitors. With LTV ₹3.2 lakh and CAC ₹1.2 lakh, the ratio is 2.67:1 — viable but sub-optimal; I'd attack churn to lift LTV since that's the highest-leverage fix.

Q19. What is CAC payback period?

A: Months to recover CAC from a customer's gross profit = CAC ÷ (monthly ARPA × gross margin). ₹1.2 lakh ÷ (₹8,000 × 0.8 = ₹6,400) = ~19 months. Shorter payback means less cash tied up per customer, which is why cash-constrained startups care about it as much as LTV:CAC.

Q20. Explain the difference between logo churn and revenue churn.

A: Logo churn counts customers lost ÷ customers at start — every account weighted equally. Revenue churn counts MRR lost ÷ starting MRR — weighted by money. If revenue churn exceeds logo churn, you're losing bigger-than-average accounts. Gross revenue churn excludes expansion (always ≥0); net revenue churn includes expansion and can go negative — the good kind.

Q21. What are NRR and GRR, and what's a good NRR?

A: NRR (Net Revenue Retention) = (starting MRR + expansion − contraction − churn) ÷ starting MRR, on an existing cohort. Above 100% is excellent — the base grows by itself; 120%+ is elite. GRR (Gross) excludes expansion, so it's capped at 100% and shows the raw leak. I always report both, because a great NRR can hide a weak GRR — i.e., churn masked by upselling.

Q22. What's the difference between MRR and ARR? Common mistakes?

A: ARR = MRR × 12 — annual vs monthly recurring revenue. The classic mistakes are quoting one as the other, and polluting either with non-recurring revenue like one-time setup fees or services. Only recurring subscription value counts; otherwise growth and retention metrics all break.

Q23. What is win rate — and be careful how you define it.

A: Two definitions: close/won rate = won ÷ (won + lost) — of decided deals; opportunity win rate = won ÷ all opportunities created. They differ a lot — e.g., 12 won of 30 decided is 40%, but 12 of 50 created is 24%. I always state which denominator I mean.

Q24. What is sales velocity?

A: Velocity = (number of opportunities × win rate × average deal size) ÷ sales cycle length in days — revenue generated per day. Its four inputs are literally the four levers to grow: more opps, higher win rate, bigger deals, shorter cycles. It's my favourite framework because it tells you exactly where to push.

Q25. How do unit economics drive hiring and quota decisions?

A: Unit economics are the P&L of one customer. If LTV:CAC is healthy and payback acceptable, every rupee into sales returns more than a rupee, so we scale — hire reps, raise marketing spend. Quota is derived: top-down from the revenue target and bottom-up from what a ramped rep can close; pipeline coverage then sets required lead volume, which sets marketing budget and SDR headcount. Broken unit economics mean hiring just loses money faster.

Q26. What's the Rule of 40?

A: For SaaS, revenue growth rate % + profit margin % should be ≥ 40. It's a balance check between growth and profitability — a company can grow fast and burn, or grow modestly and profit, but the sum should clear 40. It's a heuristic, not a law.

Q27. What's the Magic Number?

A: It measures sales & marketing efficiency = annualised net-new ARR from a quarter ÷ the prior quarter's S&M spend. Above ~0.75–1.0 means the go-to-market is efficient enough to invest more; below ~0.5 means fix efficiency before scaling spend.

Tools, analytics & process

Q28. Walk me through a modern sales tech stack.

A: The CRM (Salesforce/HubSpot) is the hub. Around it: sales engagement for cadences (Outreach, Salesloft, Apollo); data/enrichment for finding verified contacts (ZoomInfo, Lusha, Apollo); conversation intelligence to record and coach calls (Gong, Chorus); scheduling (Calendly); e-signature (DocuSign); research and social selling (LinkedIn Sales Navigator); and BI on top for reporting (Power BI/Tableau). Everything syncs back to the CRM.

Q29. What is a sales engagement platform?

A: A tool like Outreach or Salesloft that automates and tracks multi-touch outbound sequences — a scheduled mix of emails, calls and LinkedIn touches — with analytics on opens, replies and meetings. It lets an SDR run hundreds of personalised touches systematically instead of manually.

Q30. What is conversation intelligence and why is it useful?

A: Tools like Gong or Chorus record and transcribe sales calls, then analyse them — talk-to-listen ratios, competitor mentions, whether next steps were set, deal-risk signals. They're gold for coaching and for spotting at-risk deals, and they capture real voice-of-customer data.

Q31. What dashboards would you build as an SDR/BDE?

A: An activity dashboard (calls/emails/meetings vs target), a funnel/conversion dashboard to find the leak, a pipeline dashboard (open and weighted pipeline, coverage, stale-deal flags), a forecast dashboard (category roll-up and forecast-vs-actual), a cohort retention view (NRR/GRR), and a source-ROI view (CAC and win rate by channel). I'd build them in Power BI reading live from the CRM.

Q32. Leading vs lagging indicators — give examples.

A: Leading indicators are controllable inputs that predict outcomes — calls made, meetings booked, opps created, pipeline generated. Lagging indicators are the results — revenue, win rate, quota attainment, churn. I manage leading indicators daily to move lagging ones; if a lagging number slips, I diagnose which leading input broke.

Q33. What is a cohort analysis and when would you use one?

A: It groups customers by a shared start point — say, everyone who signed up in a given month — and tracks a metric like retention or revenue over their lifetime. It reveals trends a blended average hides, e.g., whether newer cohorts retain better after a product change. I'd build it in Excel pivots or Python.

Q34. How do your finance and analytics skills make you a better SDR/BDE?

A: I can build the funnel and pipeline dashboards most reps can't, query the CRM in SQL to derive real stage probabilities, model forecasts and cohorts in Excel/Power BI, and even prototype a lead-scoring model in Python. My NISM/finance background means unit economics — LTV, payback, margins — are second nature, so I connect my daily activity to the numbers leadership cares about. I don't just report metrics, I engineer them.

Q35. What is an SLA between marketing and sales?

A: A mutual agreement: marketing commits to a volume and quality of qualified leads; sales commits to a response time — ideally minutes, not hours — and a minimum number of follow-ups before disqualifying. Speed-to-lead is decisive; responding within 5 minutes hugely improves qualification. SLAs cut the "leads are junk / sales ignores them" conflict.

Q36. What is an ICP and why does it matter to metrics?

A: The Ideal Customer Profile precisely defines the best-fit customer — size, industry, geography, pain. Selling to ICP raises win rate, shortens cycle, lowers CAC and reduces churn. Chasing non-ICP logos flatters the logo count but hurts every real metric, so disciplined qualification against the ICP is itself a metric-improving behaviour.

Q37. Give an example of a vanity metric and a better alternative.

A: "Total emails sent" or "leads in the database" are vanity — they feel like progress but don't drive decisions. Better: replies, meetings held, and opportunities created — outcome metrics that connect activity to revenue. My test is: does the metric change what I do tomorrow? If not, it's vanity.

Q38. It's mid-quarter and you're at 40% of quota with 3 weeks left. What do you do?

A: Diagnose with data first. Check coverage — do I even have enough weighted pipeline to close the gap? Prioritise late-stage, high-value, in-ICP deals I can realistically pull in, and identify specific blockers on each with dated next steps. Pursue expansion/upsell in existing accounts for quicker wins. Be honest in my forecast about the floor and the gap, and if pipeline is structurally short, start generating now for next quarter rather than sandbagging or over-promising.

14. One-Page Cheat Sheet

Formulas at a glance

ARR = MRR × 12 | **CAC** = S&M spend ÷ new customers
LTV = (ARPA × GM%) ÷ churn | **LTV:CAC** = LTV ÷ CAC (aim ≥3)
CAC Payback = CAC ÷ (ARPA × GM%) (months)
Logo churn = lost customers ÷ start customers
Gross rev churn = MRR lost ÷ start MRR
NRR = (start + expansion − contraction − churn) ÷ start (aim >100%)
GRR = (start − contraction − churn) ÷ start (≤100%)
Win rate = won ÷ (won + lost) | **Quota attainment** = actual ÷ quota
Coverage = open pipeline ÷ quota (aim 3–4×)
Weighted pipeline = Σ(deal × stage prob)
Sales velocity = (opps × win rate × avg deal) ÷ cycle days
Magic number = annualised net-new ARR ÷ prior-Q S&M | **Rule of 40** = growth% + margin% ≥ 40

Mental models to recite

- CRM is the hub; stage = buyer commitment, not seller effort.
- Manage **leading** indicators to move **lagging** ones.
- Churn is the master variable behind LTV.
- Report NRR *and* GRR so churn can't hide.
- Unit economics decide whether to hire and raise quota.
- Triangulate forecasts; be boringly accurate, never sandbag.
- Sell to the ICP; disqualify honestly; keep data clean.
- My edge: SQL + Excel + Power BI + Python + finance intuition = I engineer the numbers, not just report them.

> **Final word:** Walk into the interview and be the candidate who can whiteboard the CRM object model, recite the unit-economics chain from LTV:CAC to quota to headcount, and offer to build a live pipeline/forecast dashboard. In a field where most people avoid numbers, your quantitative fluency is the fastest path from SDR/BDE to Sales Ops, Sales Strategy, or a data-respected AE seat. Learn the tools, memorise the formulas, and always connect every metric to a decision.