

INTERVIEW PREPARATION

Finance Interview Q&A Master Guide

1064 questions & model answers spanning Core Finance & Accounting, Equity Research, Risk, FP&A, Quant & Derivatives, Investment Banking / M&A / LBO, Portfolio Management, Markets & Instruments, the Economy & Financial Institutions, and Corporate Finance

India & Global context • formulas & worked examples • basic to advanced • 10 sections • Core Concepts primer

Prepared for **Saikumar Kaleru** • June 2026

Table of Contents

How to use this guide. Each question is a real interview prompt with a concise model answer; key formulas are written in plain text and many answers include a quick worked number. Topics run easy to advanced within each section. Read end to end once, then revisit the sections matching the role you are interviewing for.

★ Core Concepts Every Finance Interview Tests (start here — universal must-knows)

A. Core Finance & Accounting (Q1-Q249)

- A1. Accounting & the Three Financial Statements Q1-Q40
- A7. Advanced Financial Accounting Q41-Q73
- A2. Financial Ratios & Analysis Q74-Q98
- A8. Earnings Quality & Forensic Accounting Q99-Q128
- A3. Corporate Finance Q129-Q169
- A4. Time Value of Money & Fixed Income Q170-Q189
- A5. Economics & Statistics for Finance Q190-Q204
- A6. Indian Markets, Structure & Regulation Q205-Q224
- A9. Indian Financial System & Products (Advanced) Q225-Q249

B. Equity Research (Q250-Q339)

- B1. DCF & Intrinsic Valuation Q250-Q278
- B2. Relative Valuation & Multiples Q279-Q308
- B3. Research Process, Modeling & Technical Analysis Q309-Q339

C. Risk - Market & Credit (Q340-Q435)

- C1. Market Risk: VaR, Volatility & Backtesting Q340-Q374
- C2. Credit Risk Q375-Q404
- C3. Risk Frameworks, Regulation & Other Risks Q405-Q435

D. Financial Analyst / FP&A (Q436-Q515)

- D1. Budgeting, Forecasting & Variance Analysis Q436-Q465
- D2. 3-Statement Financial Modeling Q466-Q490
- D3. Cost Accounting, Unit Economics & Business Partnering Q491-Q515

E. Quant / Derivatives (Q516-Q646)

- E1. Options & Futures Fundamentals Q516-Q546
- E2. Black-Scholes-Merton & the Greeks Q547-Q581
- E3. Implied Volatility, the Smile & Vol Trading Q582-Q601
- E4. Option Strategies & Quant Math Q602-Q616
- E5. Advanced Derivatives & Quantitative Finance Q617-Q646

F. Investment Banking, M&A & LBO (Q647-Q738)

- F1. Mergers & Acquisitions Q647-Q676
- F2. Leveraged Buyouts & Private Equity Q677-Q707
- F3. Deal Valuation & Capital Raising Q708-Q738

G. Portfolio Management & Investments (Q739-Q856)

- G1. Modern Portfolio Theory Q739-Q768
- G2. CAPM, Factor Models & Performance Q769-Q798
- G3. Asset & Investment Management Q799-Q831
- G4. Market Efficiency & Behavioral Finance Q832-Q856

H. Markets & Instruments (Q857-Q948)

- H1. Fixed Income (Advanced) Q857-Q888
- H2. Swaps, Interest-Rate & FX Derivatives Q889-Q918
- H3. Foreign Exchange, Commodities & Money Markets Q919-Q948

I. Economy & Financial Institutions (Q949-Q1008)

- I1. Macroeconomics, Monetary & Fiscal Policy Q949-Q978
- I2. Banking & Financial Institutions Q979-Q1008

J. Advanced Corporate Finance & Contemporary Finance (Q1009-Q1064)

- J1. Advanced Corporate Finance Q1009-Q1039
- J2. ESG, Sustainable Finance & Fintech Q1040-Q1064

★ Core Concepts Every Finance Interview Tests

Start here. These are the concepts that come up in almost every finance interview, whatever the role, analyst to associate, equity research to FP&A to risk. Learn these cold and you can reason your way through most questions. The 1,064 Q&A that follow go deeper by specialty; this section is the spine that ties them together.

1. The Three Financial Statements and How They Link

This is the single most asked interview topic. You must be able to say it in 60 seconds.

- **Income Statement (P&L)** shows profitability over a period: Revenue minus expenses gives Net Income. It uses accrual accounting, so it records revenue when earned and expenses when incurred, not when cash moves.
- **Balance Sheet** is a snapshot at a point in time: Assets = Liabilities + Shareholders' Equity. It always balances.
- **Cash Flow Statement** reconciles Net Income to the actual cash that moved, in three buckets: Operating, Investing, Financing. It exists because profit is not cash.

How they link (say this): Net Income from the P&L flows to the top of Cash Flow from Operations and into Retained Earnings on the Balance Sheet. The Cash Flow Statement adjusts Net Income for non-cash items (depreciation) and working-capital changes; its bottom line, the net change in cash, updates the Cash line on the Balance Sheet. Depreciation on the P&L reduces the PP&E asset on the Balance Sheet. Debt raised or repaid hits Financing on the Cash Flow and the debt line on the Balance Sheet.

Interview angle: "Walk me through the three statements." Answer: name each, its purpose and time dimension, then the one sentence on how Net Income and cash tie them together. Then be ready for the follow-up in concept 2.

2. The Classic Flow-Through: Depreciation of 100 (tax rate 30%)

The most common follow-up. Walk it slowly and it is easy.

- **Income Statement:** Depreciation of 100 is an expense, so pre-tax income falls 100. At a 30% tax rate, taxes fall 30, so Net Income falls **70**.
- **Cash Flow:** Start from Net Income down 70, then add back the 100 of depreciation because it is non-cash. Cash from Operations rises **+30** (the tax saving). Net change in cash is +30.
- **Balance Sheet:** Cash is up 30; PP&E is down 100 (the asset depreciated). The asset side nets to **-70**. On the other side, Retained Earnings is down 70 from lower Net Income. Both sides fall 70, so it **balances**.

Why it matters: Depreciation is a non-cash expense that actually creates cash through the tax shield. This one example proves you understand accrual vs cash and how the statements interlock.

3. Working Capital and the Cash Conversion Cycle

Working Capital = Current Assets - Current Liabilities. In modeling we usually mean *operating* working capital: receivables + inventory - payables.

Cash Conversion Cycle (days)

$$CCC = DSO + DIO - DPO$$

where: DSO = days sales outstanding (receivables), DIO = days inventory outstanding, DPO = days payables outstanding

An **increase** in working capital (more receivables or inventory) is a **use of cash**; a decrease is a source of cash. Growing companies often consume cash even while profitable because they fund receivables and inventory. A shorter CCC frees up cash.

Interview angle: "If working capital goes up, what happens to cash?" It goes down. Retailers negotiate long payables and fast inventory turns to run a negative cycle (Amazon), effectively financing growth with suppliers' cash.

4. Accrual vs Cash Accounting (why profit is not cash)

Accrual accounting matches revenue to when it is earned and expenses to when they are incurred, regardless of cash timing. This gives a truer picture of performance but means **Net Income can be high while cash is negative** (a company booking sales on credit) or the reverse. This gap is exactly why the Cash Flow Statement exists, and why analysts trust cash flow over reported earnings when the two diverge.

5. Enterprise Value vs Equity Value (the bridge)

The bridge

$$\text{Enterprise Value} = \text{Equity Value} + \text{Total Debt} + \text{Preferred} + \text{Minority Interest} - \text{Cash \& Equivalents}$$

$$\text{Equity Value (market cap)} = \text{share price} \times \text{shares outstanding}$$

- **Equity Value** is what belongs to shareholders.
- **Enterprise Value** is the value of the whole operating business, what it would cost to buy the operations free of capital structure. You add debt (an acquirer assumes it) and subtract cash (an acquirer gets it, reducing net cost).

Match the metric to the value: EV pairs with pre-interest, pre-capital-structure metrics (EBIT, EBITDA, Revenue). Equity Value pairs with after-interest metrics (Net Income, EPS). So EV/EBITDA and P/E, never EV/Net Income.

Interview angle: "Why subtract cash?" Because cash could immediately be used to pay down the purchase price; you are buying the business, not its bank balance. "Why does EV use EBITDA and P/E use Net Income?" Capital-structure neutrality.

6. The Three Valuation Methods and When to Use Each

Method	What it is	Tends to give
DCF	Intrinsic: present value of future free cash flows	An intrinsic value, sensitive to assumptions
Comparable Companies	Relative: peer trading multiples (EV/EBITDA, P/E)	Current market view, a "trading" value
Precedent Transactions	Relative: multiples paid in past M&A deals	The highest, includes a control premium

DCF is theoretically purest but assumption-driven. Comps reflect where the market prices peers today. Precedents reflect what acquirers actually paid (usually highest because of control premiums and synergies). Good analysts triangulate all three on a **football field** chart.

7. DCF Mechanics

Value a business as the present value of the cash it will generate.

Unlevered Free Cash Flow (FCFF)

$$\text{FCFF} = \text{EBIT} \times (1 - \text{tax}) + \text{D\&A} - \text{Capex} - \text{Change in Net Working Capital}$$

discount FCFF at WACC to get Enterprise Value

Terminal Value (Gordon growth)

$$\text{TV} = \text{FCF final} \times (1 + g) / (\text{WACC} - g)$$

g = perpetual growth, must be below long-run GDP (2-3%). Alternative: exit-multiple method, $\text{TV} = \text{EBITDA} \times \text{an EV/EBITDA multiple}$

Steps: project 5 years of FCFF, compute a terminal value (Gordon or exit multiple), discount everything to today at WACC to get Enterprise Value, then bridge EV to Equity Value (concept 5) and divide by shares for intrinsic price. **Terminal value is usually 60-80% of the total**, so its assumptions dominate; always run a sensitivity on WACC and g .

- **FCFF** (unlevered) is discounted at WACC and gives Enterprise Value.
- **FCFE** (levered, after interest and debt flows) is discounted at cost of equity and gives Equity Value directly.

8. WACC and CAPM (the cost of capital)

WACC

$$\text{WACC} = (E/V) \times \text{Re} + (D/V) \times \text{Rd} \times (1 - \text{tax})$$

E = equity value, D = debt, $V = E + D$; Re = cost of equity, Rd = cost of debt

CAPM (cost of equity)

$$\text{Re} = \text{Rf} + \text{Beta} \times (\text{Rm} - \text{Rf})$$

Rf = risk-free rate, Beta = sensitivity to the market, $(\text{Rm} - \text{Rf})$ = equity risk premium

WACC is the blended required return of all capital providers, and the discount rate for unlevered cash flows. Debt is multiplied by $(1 - \text{tax})$ because **interest is tax-deductible** (the tax shield), which makes debt cheaper than equity. Cost of equity comes from CAPM: a stock with beta above 1 is more volatile than the market and demands a higher return.

Interview angle: "Why is debt cheaper than equity?" Two reasons: debt holders are senior and bear less risk, so demand a lower return; and interest is tax-deductible. "What happens to WACC as you add debt?" It falls at first (cheap debt), then rises as financial-distress and bankruptcy risk push up both R_d and R_e .

9. Time Value of Money, NPV and IRR

A rupee today is worth more than a rupee tomorrow because it can be invested. Everything in finance discounts future cash to the present.

Net Present Value

$$\text{NPV} = \text{sum of } [CF_t / (1 + r)^t] - \text{initial investment}$$

accept a project if $\text{NPV} > 0$

- **IRR** is the discount rate that makes $\text{NPV} = 0$, the project's implied annualized return. Accept if $\text{IRR} >$ hurdle rate (cost of capital).
- **NPV vs IRR conflict:** for mutually exclusive projects of different size or timing they can disagree; **trust NPV**, because it measures value created in absolute terms and assumes reinvestment at the cost of capital, whereas IRR assumes reinvestment at the IRR itself.

10. Key Ratios and the DuPont Identity

Group	Ratios	Reads
Liquidity	Current, Quick (acid-test)	Can it pay short-term bills
Leverage	Debt/Equity, Debt/EBITDA, Interest Coverage	How much debt, can it service it
Profitability	Gross, Operating, Net margin; ROE, ROA, ROIC	How efficiently it earns
Efficiency	Asset turnover, inventory turns, DSO	How hard assets work

DuPont: decomposing ROE

$$\text{ROE} = \text{Net Margin} \times \text{Asset Turnover} \times \text{Equity Multiplier}$$

$= (\text{profitability}) \times (\text{efficiency}) \times (\text{leverage})$. Shows whether a high ROE comes from real margins or just from debt

Interview angle: "Two firms have the same ROE, are they equal?" Not necessarily, use DuPont: one may earn it through strong margins, the other by loading on leverage, which is far riskier.

11. Capital Structure: Debt vs Equity

Debt is cheaper (senior claim, tax-deductible interest) but adds fixed obligations and bankruptcy risk. Equity is flexible (no repayment, no mandatory dividend) but dilutive and more expensive. **Modigliani-Miller** says that in a perfect, tax-free world capital structure does not affect firm value; adding the real-world

frictions of the **tax shield** (favors debt) and **financial-distress costs** (limits debt) gives the trade-off theory: an optimal structure balances the tax benefit of debt against the rising cost of distress.

12. Multiples: P/E vs EV/EBITDA

- **P/E** = Price / EPS: simple, equity-level, but distorted by different debt loads, tax rates, and depreciation policies.
- **EV/EBITDA**: capital-structure and depreciation neutral, so it is better for comparing companies with different leverage or for M&A. Preferred in most cross-company comparison.
- **PEG** = P/E / growth: adjusts the P/E for how fast earnings grow, so a high P/E can still be cheap if growth is high.

13. Risk and Return

- **Total risk = systematic + unsystematic.** Unsystematic (company-specific) risk is diversifiable, spread across many stocks and it washes out. Systematic (market) risk cannot be diversified away and is what investors are paid to bear, captured by **beta**.
- **Diversification** works because assets are not perfectly correlated; combining them lowers portfolio volatility for a given return (Markowitz).

Sharpe Ratio

$\text{Sharpe} = (\text{Return of portfolio} - R_f) / \text{Standard deviation of portfolio}$

excess return per unit of total risk; higher is better risk-adjusted performance

14. Bonds: the Essentials

- **Price and yield move inversely.** When market rates rise, existing bonds with lower coupons become less attractive, so their price falls.
- **Duration** measures interest-rate sensitivity: a duration of 5 means roughly a 5% price fall for a 1% rise in yield. Longer maturity and lower coupon mean higher duration.
- **Credit spread** is the extra yield over the risk-free rate that compensates for default risk; it widens when the economy weakens.

Price sensitivity

Percentage price change is approximately $-\text{Duration} \times \text{change in yield}$

convexity is the second-order correction for large yield moves

15. Accretion / Dilution (M&A basics)

A deal is **accretive** if the acquirer's post-deal EPS rises, **dilutive** if it falls. Quick rule for an all-stock deal: if the acquirer's **P/E is higher than the target's P/E, the deal is accretive** (you are buying cheaper earnings with expensive shares), and dilutive if lower. Cash and debt financing are usually more accretive than stock because cash earns little and debt interest is cheap and tax-deductible, though they raise risk.

16. Non-Cash Items and Deferred Taxes

Depreciation and amortization, stock-based compensation, and impairments reduce accounting profit but are **added back** in cash flow because no cash left. **Deferred tax** arises when book and tax accounting differ in timing (for example, accelerated depreciation for tax): the company reports a tax expense it has not yet paid, creating a deferred tax liability that reverses later.

17. How to Pitch a Stock (structure to have ready)

Interviewers in any markets-facing role ask "pitch me a stock." Use this skeleton: **(1) Recommendation** (buy/sell and target), **(2) Business** in one line, **(3) Thesis**, two or three reasons the market is mispricing it, **(4) Valuation** (your DCF or multiple vs the market), **(5) Catalysts** (what closes the gap), **(6) Risks** and what would make you wrong. Have one long and one short idea prepared.

18. The Behavioral Core

Every interview also tests fit. Prepare crisp, structured answers for: **"Walk me through your resume"** (a 90-second narrative, not a list), **"Why finance / why this role"**, **"Tell me about a time you..."** (use STAR: Situation, Task, Action, Result), and a few smart **questions for them**. Know your own resume and projects cold, every number on them is fair game.

19. EBITDA: What It Is and Its Limits

EBITDA = Earnings Before Interest, Taxes, Depreciation and Amortization. It approximates core operating profitability while stripping out capital structure (interest), tax regimes, and non-cash D&A, so you can compare the raw operating performance of companies with different debt and asset bases. It drives EV/EBITDA and the leverage ratio Debt/EBITDA.

Its limits (say these): EBITDA is **not** cash flow. It ignores capex, working-capital needs, and cash taxes, all real costs. Capital-intensive businesses look deceptively strong on EBITDA because their heavy capex is invisible.

Interview angle: "Is EBITDA a good proxy for cash flow?" No, it excludes capex, working-capital changes, and taxes. That is the whole reason free cash flow exists.

20. The Free Cash Flow Family

FCFF (unlevered) = $EBIT \times (1 - t) + D\&A - Capex - Change\ in\ NWC$ → discount at **WACC** → Enterprise Value
FCFE (levered) = $Net\ Income + D\&A - Capex - Change\ in\ NWC + Net\ Borrowing$ → discount at **cost of equity** → Equity Value

The one rule that trips people up: **match the cash flow to the discount rate.** Unlevered FCFF pairs with WACC and gives Enterprise Value; levered FCFE pairs with cost of equity and gives Equity Value directly. And FCF is not EBITDA, FCF subtracts the capex, working capital, and taxes that EBITDA leaves out.

21. Return on Capital and Value Creation

Metric	Formula	Reads
ROE	Net Income / Equity	Return to shareholders (flattered by leverage)
ROA	Net Income / Assets	Asset efficiency, ignores capital structure
ROIC	NOPAT / Invested Capital	Return on all invested capital

A company **creates value only when ROIC > WACC**, it earns more than the cost of the capital it deploys. If ROIC is below WACC, growth actually **destroys** value.

Interview angle: "How do you know a company is creating value?" The ROIC-minus-WACC spread. Positive and widening is the mark of a good business.

22. Goodwill and Intangibles

Goodwill = purchase price - fair value of net identifiable assets acquired in an acquisition. It captures brand, customer relationships, and expected synergies that are not separately recorded. Under IFRS and US GAAP it is **not amortized** but tested annually for **impairment**; a write-down is a non-cash hit to the P&L. Other intangibles (patents, software, licences) are typically amortized over their useful life.

Interview angle: "A firm pays 100 for a business whose net assets are worth 70. What goes on the balance sheet?" 30 of goodwill.

23. Revenue Recognition and Deferred Revenue

Revenue is recognized when it is **earned** (goods or services delivered), not when cash arrives, that is accrual accounting. Cash collected **before** delivery (an annual software subscription paid upfront) is **deferred revenue, a liability**, released to the P&L as the service is delivered. This is a classic divergence between cash and reported revenue and a common area for manipulation (premature recognition, channel stuffing).

Interview angle: "Is deferred revenue an asset or a liability?" A liability, the company still owes the service it was paid for.

24. Dividends vs Share Buybacks

Both return cash to shareholders. A **dividend** pays cash per share (cash and retained earnings fall). A **buyback** repurchases shares (cash and equity fall, share count shrinks, so EPS rises). Buybacks are more tax-efficient and flexible; dividends signal a stable, ongoing commitment. Neither creates value on its own, it depends on whether the company had higher-return uses for the cash.

Interview angle: "How does a buyback flow through the statements?" Cash down (financing outflow), treasury stock up so equity down, fewer shares outstanding so EPS up.

25. LBO Basics

In a **leveraged buyout** a private-equity firm buys a company funded mostly with **debt** (often 60-70%), and the target's own cash flows service that debt. Returns come from three levers: **deleveraging** (cash flow pays down debt, growing the equity slice), **EBITDA growth**, and **multiple expansion** (exiting at a higher multiple). Measured by **IRR** and **MOIC** (money on money). A "Sources and Uses" table must balance.

Interview angle: "What makes a good LBO candidate?" Stable, predictable cash flows, low existing debt, modest capex, and strong assets to borrow against.

26. Diluted Shares, the Treasury Stock Method and Net Debt

Diluted shares add the effect of in-the-money options, warrants, and convertibles. The **Treasury Stock Method** assumes option holders exercise and the company uses the proceeds to buy back shares at the market price; net new shares = options minus shares repurchased. **Net Debt = Total Debt - Cash**. Equity value uses the diluted share count; enterprise value uses net debt.

27. Levered vs Unlevered Beta

Observed (equity/levered) beta includes the company's financial risk from debt. To compare pure **business risk** across peers or to reweight for a target capital structure, you **unlever** it and then **relever**.

Hamada

$$\text{Unlevered beta} = \text{Levered beta} / [1 + (1 - t) \times (D/E)]$$

strip out leverage to get asset beta, then relever at the target D/E for CAPM

Interview angle: "Why unlever beta?" To remove capital-structure differences so peers are comparable, then relever to the company you are valuing.

28. Operating Leverage, Financial Leverage and Break-Even

Operating leverage is the fixed-vs-variable cost mix: high fixed costs make profit swing more for a given revenue change. **Financial leverage** is the use of debt, which amplifies both returns and risk.

$$\text{Break-even (units)} = \text{Fixed Costs} / \text{Contribution Margin per unit}$$

$$\text{Contribution Margin} = \text{Price} - \text{Variable Cost per unit}$$

Interview angle: "Two firms, same revenue, one with high fixed costs?" The high-fixed-cost firm has higher operating leverage, more profit upside when sales rise and more pain when they fall.

29. Interest Rates and the Central Bank

The central bank (RBI in India, the Fed in the US) sets the **policy rate** to balance inflation and growth. **Higher rates raise discount rates (WACC), which lowers the present value of future cash flows**, so valuations fall, most sharply for long-duration growth stocks whose value sits far in the future. Higher rates also raise borrowing costs and cool demand. Rate cuts do the reverse.

Interview angle: "Why do rising rates hit growth stocks hardest?" Their cash flows are far out, so a higher discount rate compresses them most.

30. Inflation: Real vs Nominal

Nominal figures include inflation; **real** figures strip it out.

Fisher relation

Real rate is approximately Nominal rate - Inflation

more precisely: $(1 + \text{nominal}) = (1 + \text{real}) \times (1 + \text{inflation})$

Inflation erodes purchasing power, pushes up nominal interest rates, helps borrowers (they repay in cheaper money) and hurts lenders and fixed-rate bondholders. Rising inflation expectations push bond prices down.

31. Credit Analysis Essentials

Lenders weigh the **5 Cs**: Character, Capacity, Capital, Collateral, Conditions. The core numbers:

Coverage & serviceability

Interest Coverage = $\text{EBIT} / \text{Interest}$ • Leverage = $\text{Debt} / \text{EBITDA}$ • DSCR = $\text{Cash available for debt service} / \text{Debt service}$

DSCR above 1 means the borrower generates enough to cover principal + interest

Interview angle: "How do you judge if a firm can take on more debt?" Compare leverage and coverage to peers and covenants, and test debt capacity under a stressed cash-flow case.

32. Liquidity vs Solvency

Liquidity is the ability to meet short-term obligations (current and quick ratios, cash on hand). **Solvency** is the ability to meet long-term obligations (leverage, the equity cushion). They are different: a **profitable company can still go bankrupt** if it runs out of cash to pay bills as they fall due, a liquidity crisis, even while solvent on paper.

Interview angle: "Can a profitable company go bankrupt?" Yes, if it is illiquid. Profit is accrual; bankruptcy is about cash.

33. Financial Modeling Best Practices

A clean, auditable model beats a clever one. The principles interviewers look for: separate **inputs/assumptions** from **calculations**; drive the model off a few clear **drivers**; keep each row's formula consistent across columns; build in **checks** (the balance sheet must balance, cash must tie to the cash-flow statement); handle **circularity** (interest depends on debt, which depends on cash, which depends on interest) with an iterative switch or a circularity breaker; label units; and make scenarios toggle-able.

Interview angle: "Where is the circular reference in a 3-statement model?" Interest expense ↔ debt balance ↔ cash. Solved with iterative calculation or a copy-paste breaker.

34. Performance Metrics (Portfolio and Investing)

Metric	Formula	Measures
Alpha	Return above CAPM / benchmark	Skill, value added
Beta	Sensitivity to the market	Systematic risk
Sharpe	(Return - Rf) / total volatility	Return per unit of total risk
Treynor	(Return - Rf) / beta	Return per unit of market risk
Information Ratio	Active return / tracking error	Consistency of outperformance

Interview angle: Sharpe uses total risk (standard deviation); Treynor uses only systematic risk (beta). Use Treynor when the portfolio is one part of a diversified whole.

Accounting deep-dive. Accounting is the language every finance interview is conducted in. These are the accounting mechanics that come up again and again, whatever the role, from a fund accountant to an equity analyst.

35. The Accounting Equation and Double-Entry

The identity behind the balance sheet
 $\text{Assets} = \text{Liabilities} + \text{Shareholders' Equity}$

Every transaction has two equal sides (double-entry): a **debit** and a matching **credit**, which keeps the equation in balance. Debits increase assets and expenses; credits increase liabilities, equity, and revenue.

Interview angle: "A company buys 100 of inventory on credit, what happens?" Inventory (asset) up 100 and Accounts Payable (liability) up 100. Both sides rise, the equation still balances.

36. Adjusting Entries: Accruals vs Deferrals

At period end, accrual accounting needs adjusting entries so revenue and expense fall in the right period. Four types:

Type	Situation	Creates
Accrued revenue	Earned, not yet billed	An asset (receivable)
Accrued expense	Incurred, not yet paid (e.g. wages)	A liability (payable)
Prepaid expense	Paid in advance (e.g. insurance)	An asset, expensed over time
Unearned (deferred) revenue	Cash received before delivery	A liability

37. Inventory: FIFO vs LIFO vs Weighted Average

The cost-flow assumption decides which costs go to COGS.

- **FIFO** (first in, first out): oldest costs to COGS. In inflation this gives **lower COGS, higher profit, higher tax**, and a higher ending inventory value.

- **LIFO** (last in, first out): newest costs to COGS. In inflation this gives **higher COGS, lower profit, lower tax**. LIFO is allowed under **US GAAP but banned under IFRS**.
- **Weighted average**: blends all unit costs.

Interview angle: "In an inflationary period, which method shows higher profit?" FIFO. "Which is not allowed under IFRS?" LIFO.

38. Depreciation and Amortization Methods

Depreciation spreads a tangible asset's cost over its useful life; amortization does the same for intangibles; both are **non-cash**. **Straight-line** charges an equal amount each year. **Accelerated** methods (written-down-value, double-declining) charge more early. Firms often use straight-line for the books and an accelerated method for tax.

Interview angle: "Why use accelerated depreciation for tax but straight-line for books?" Higher tax depreciation early lowers taxable income now, deferring tax, a time-value benefit. The gap creates a deferred tax liability (topic 41).

39. Capitalize vs Expense (Capex vs Opex)

A cost whose benefit extends beyond the current period is **capitalized** (recorded as an asset and depreciated over time). A cost consumed now is **expensed** immediately. Capitalizing lifts current profit (the cost is spread out) and creates an asset plus future depreciation. **Aggressively capitalizing costs that should be expensed is a classic fraud red flag** (WorldCom).

Interview angle: "A company capitalizes a cost it should have expensed, what happens?" Current profit and assets are overstated; cash flow classification shifts from operating to investing.

40. Accounts Receivable and the Allowance for Doubtful Accounts

Receivables are amounts customers owe. Because some will not pay, firms record an **allowance for doubtful accounts** (a contra-asset) and a matching **bad-debt expense**, recognizing the likely loss in the same period as the sale (prudence). **Net receivables = gross receivables - allowance**. Writing off a specific bad debt later reduces both gross AR and the allowance, with no further P&L hit because the expense was already taken.

41. Deferred Tax Assets and Liabilities

Because book accounting and tax rules differ in **timing**, the tax expense shown on the P&L differs from tax actually paid; the gap sits on the balance sheet.

- **Deferred Tax Liability (DTL):** book profit exceeds taxable profit now (e.g. accelerated tax depreciation), so more tax will be paid later.
- **Deferred Tax Asset (DTA):** more tax is paid now than the books imply (e.g. carryforward losses, warranty provisions), a future tax saving.

Interview angle: "Accelerated depreciation for tax creates what?" A deferred tax liability that reverses in later years.

42. Provisions, Contingent Liabilities and Reserves

- A **provision** is a liability of uncertain timing or amount, recognized when an outflow is **probable and can be estimated** (warranty costs, restructuring).
- A **contingent liability** is possible but not probable or not estimable, so it is **disclosed in the notes, not booked** (a pending lawsuit).
- A **reserve** is an appropriation of retained earnings within equity, not an expense.

Interview angle: "A lawsuit you might lose, provision or disclosure?" If loss is probable and estimable, provide for it; if only possible, disclose it.

43. Lease Accounting: Operating vs Finance

A **finance (capital) lease** transfers substantially all the risks and rewards of ownership, so the lessee records an asset and a liability and books depreciation plus interest. An **operating lease** was historically off balance sheet (just rent expense). Under **IFRS 16 / ASC 842**, lessees now bring almost all leases onto the balance sheet as a **right-of-use asset and a lease liability**.

Interview angle: "What changed with the new lease standards?" Operating leases moved onto the balance sheet, raising reported assets and debt-like liabilities.

44. Stock-Based Compensation

Stock-based comp (options, RSUs) is a **real, non-cash expense** on the P&L, employees are paid in equity. On the cash flow statement it is **added back** (no cash left); on the balance sheet it increases paid-in capital; and it increases the **diluted share count**, so it dilutes existing shareholders.

Interview angle: "Is SBC a real expense?" Yes, it transfers value to employees and dilutes owners, even though no cash goes out. Ignoring it flatters adjusted earnings.

45. Consolidation, Minority Interest and the Equity Method

Ownership drives the accounting:

Stake	Treatment
Below 20%	Cost / fair value
20-50% (significant influence)	Equity method: book your share of the investee's profit
Above 50% (control)	Full consolidation: combine 100% of the subsidiary, then subtract the share you do not own as minority (non-controlling) interest

Interview angle: "You own 30% of a company, how do you account for it?" Equity method. "Why does enterprise value add minority interest?" Because consolidated EBITDA includes 100% of the subsidiary, so EV must reflect the whole.

46. Basic vs Diluted EPS

Basic EPS = (Net Income - Preferred Dividends) / Weighted-average shares

Diluted EPS also adds shares from in-the-money options, warrants and convertibles (via the treasury stock method)

so diluted EPS is always less than or equal to basic EPS; analysts use diluted

47. Impairment of Assets

When an asset's **carrying value exceeds its recoverable amount** (what it can earn or be sold for), it is written down. The impairment is a **non-cash expense** on the P&L that also reduces the asset and equity. Goodwill is tested for impairment annually. Under US GAAP most impairments are **not reversed**; IFRS allows reversal for some assets but never for goodwill.

Interview angle: "How does a goodwill impairment flow through?" Net income falls (P&L expense), the goodwill asset and equity fall, and it is added back on the cash flow as non-cash.

48. GAAP vs IFRS: Key Differences

US GAAP is **rules-based**; IFRS is **principles-based**. The differences interviewers expect you to know:

Area	US GAAP	IFRS
Inventory	LIFO allowed	LIFO banned
PP&E	Historical cost	Revaluation to fair value allowed
Development costs	Mostly expensed	Can be capitalized if criteria met
Impairment reversal	Not allowed	Allowed (except goodwill)

49. Direct vs Indirect Cash Flow Statement

The operating section can be built two ways. The **indirect method** (used almost everywhere) starts from **Net Income** and adjusts for non-cash items and working-capital changes. The **direct method** lists actual cash receipts and payments. Both produce the **same operating cash figure**; indirect is preferred because it reconciles profit to cash and is easy to build from the ledger.

Interview angle: "Which cash-flow method is more common and why?" Indirect, because it ties net income to cash and is simpler to prepare.

Markets, macro and institutions. Beyond accounting and valuation, interviewers test whether you understand how markets, derivatives, the economy, and banks actually work. These are the broad must-knows for any markets-facing or finance role.

50. Primary vs Secondary Markets and the IPO Process

The **primary market** is where securities are issued for the first time and the **company raises capital** (IPO, follow-on offer, rights issue, bond issuance). The **secondary market** is where investors trade existing securities among themselves on an exchange, the company receives no new money. An IPO runs: appoint bankers, file the prospectus (a DRHP with SEBI in India), price the shares via book-building, then list.

Interview angle: "When Reliance shares trade on the NSE, does Reliance get the cash?" No, that is the secondary market. The company only raised money at issuance.

51. How Exchanges Work: Bid-Ask, Liquidity and Order Types

Buyers post **bids**, sellers post **asks**; the **bid-ask spread** between them is a cost of trading, a narrower spread means more **liquidity** (how easily you can trade without moving the price). Order types: a **market order** executes now at the best available price, a **limit order** only at a set price or better, and a **stop-loss** triggers once a level is hit. **Market makers** quote both sides to supply liquidity and earn the spread.

Interview angle: "What does a wide bid-ask spread tell you?" The security is illiquid and expensive to trade.

52. Long, Short and Short Selling

Going **long** means buying, you profit if the price rises. Going **short** means selling borrowed shares hoping to buy them back cheaper; you profit if the price falls. A long's maximum loss is what you invested, but a **short's loss is theoretically unlimited** because the price can rise without bound. Shorts pay a borrow fee and can be caught in a short squeeze.

Interview angle: "What is the maximum loss on a short position?" Unlimited, versus a capped loss on a long.

53. Futures vs Forwards

Both lock in a price today for a transaction in the future.

	Forward	Future
Traded	Private, OTC	Exchange-traded
Terms	Customized	Standardized
Counterparty risk	Yes	Minimal (clearinghouse)
Settlement	At maturity	Marked-to-market daily, margin

Interview angle: The three differences to name: standardization, daily mark-to-market, and a clearinghouse that removes counterparty risk.

54. Options: Calls, Puts, Intrinsic vs Time Value

A **call** is the right (not the obligation) to **buy** at the strike; a **put** is the right to **sell**. The buyer pays a premium and can only lose that premium. Option value has two parts: **intrinsic value** (how in-the-money it is now, $\max(S - K, 0)$ for a call) plus **time value** (the extra paid for the chance it moves further before expiry). Time value decays to zero at expiration, that decay is **theta**.

Interview angle: "Most a call buyer can lose?" The premium. "An option with no intrinsic value still has a price, why?" Time value, the chance it finishes in the money.

55. Hedging, Speculation and Arbitrage

- **Hedging:** taking an offsetting position to reduce risk (a farmer shorting wheat futures to lock a price).
- **Speculation:** taking on risk deliberately for expected profit (betting on direction).
- **Arbitrage:** a near risk-free profit from the same asset being priced differently in two markets, quickly competed away.

56. Interest Rate Swaps

Two parties exchange interest streams on a **notional** amount: typically one pays **fixed** and receives **floating**, the other the reverse. It is used to convert floating-rate debt to fixed (or vice versa) to manage interest-rate risk. Only the **net interest** changes hands; the notional itself is never exchanged.

Interview angle: "Why pay fixed in a swap?" To hedge floating-rate borrowing against rising rates by locking in a fixed cost.

57. The Yield Curve

The yield curve plots yield against maturity for one issuer, usually government bonds. **Normal** is upward-sloping (longer maturities pay more for time and risk). An **inverted** curve, where short rates exceed long rates, has historically been a **recession signal**: the market expects the central bank to cut rates as growth slows.

Interview angle: "What does an inverted yield curve suggest?" A likely economic slowdown or recession ahead.

58. Monetary vs Fiscal Policy

Monetary policy is run by the central bank, managing interest rates, open-market operations, and reserve requirements (CRR/SLR) to steer inflation and growth. **Fiscal policy** is run by the government, using taxation and public spending. The two can reinforce or offset each other.

Interview angle: Know the owners: rates and money supply belong to the central bank; taxes and spending to the government.

59. GDP, the Business Cycle and Leading Indicators

$$\text{GDP (expenditure)} = \text{Consumption} + \text{Investment} + \text{Government spending} + \text{Net exports}$$

Economies move through **expansion, peak, contraction (recession), trough**. **Leading indicators** turn before the economy (yield curve, PMI, new orders, building permits, equities); **lagging indicators** turn after (unemployment). The rule of thumb for recession is two consecutive quarters of falling real GDP.

60. Exchange Rates and the Current Account

An exchange rate is the price of one currency in another, driven by interest-rate differentials, inflation, and trade and capital flows. A **current-account deficit** (importing more than exporting) tends to weaken a currency. **Depreciation** helps exporters and hurts importers, and it raises imported inflation.

Interview angle: "A central bank hikes rates, what happens to its currency?" It usually appreciates, higher yields attract foreign capital.

61. The Four Types of Financial Risk

Risk	What it is
Market	Losses from moving prices: equity, interest-rate, FX, commodity
Credit	A borrower or counterparty defaults
Liquidity	Cannot trade or fund a position without a large loss
Operational	Failures of people, process, systems, or external events

Interview angle: "A bank cannot sell a bond without a steep price cut, which risk?" Liquidity risk.

62. Working Capital Management

Managing the short-term operating cycle to stay liquid without trapping cash: collect receivables faster (lower DSO), turn inventory quicker (lower DIO), and stretch payables sensibly (higher DPO), which shortens the cash conversion cycle. Too little working capital risks a liquidity crunch; too much wastes capital that could earn a return.

Interview angle: "Levers to free up cash?" Cut DSO and DIO, extend DPO.

63. Dividend Policy, Signaling and the Agency Problem

Firms decide how much profit to return versus reinvest. Starting or raising a dividend **signals** confidence in stable future earnings, and cutting it signals distress, which is why the market often punishes a cut sharply. The **agency problem** is the conflict between managers (agents) and shareholders (principals); returning cash through dividends and buybacks reduces the free cash managers might waste, helping align their interests alongside governance and incentive pay.

Interview angle: "Why does a dividend cut hit the stock hard?" It signals management expects weaker earnings, not just less cash returned.

64. How Banks Make Money (and Their Key Ratios)

A bank's core profit is the **Net Interest Margin (NIM)** = interest earned on loans minus interest paid on deposits, plus fee income. **CASA** (current and savings account) deposits are cheap funding, so a high CASA ratio lifts margins. Health is judged by the **Capital Adequacy Ratio (CAR)** under Basel and by asset quality through **non-performing assets (the gross NPA ratio)**.

Interview angle: "How does a bank make money?" Mainly the net interest margin; a high CASA ratio and low NPAs mark a strong bank.

Rounding out the toolkit. Statement analysis, deals, funds, statistics, behavioral finance, ESG, fintech and the Indian market, plus how to handle behavioral and guesstimate questions. This block covers the remaining ground interviewers probe across roles.

65. Financial Statement Analysis

Three lenses beyond the raw numbers: **common-size (vertical)** analysis expresses each line as a percent of a base (revenue on the P&L, total assets on the balance sheet) so firms of different size compare; **horizontal (trend)** analysis tracks growth across years; and **ratio analysis** (topic 10). **Earnings-quality red flags:** profit rising while operating cash flow falls, receivables growing faster than sales, recurring "one-off" adjustments, and a rising DSO.

Interview angle: "Net income is up but operating cash flow is down, should you worry?" Yes, a classic red flag for aggressive revenue recognition or a working-capital build.

66. Types of M&A, Synergies and Why Deals Fail

Deals are **horizontal** (same industry), **vertical** (up or down the supply chain), or **conglomerate** (unrelated). **Synergies** are cost (remove duplication, usually credible) or revenue (cross-sell, usually optimistic). Most acquisitions disappoint because of **overpayment** (a big control premium), **overestimated synergies**, and **integration and culture** problems.

Interview angle: "Why do most acquisitions destroy value?" Overpaying and over-promising synergies, then failing to integrate.

67. Private Equity vs Venture Capital

	Private Equity	Venture Capital
Stage	Mature, cash-generative firms	Early-stage startups
Stake	Majority / control	Minority
Leverage	Heavy (LBO)	None
Return driver	Cash flow, deleveraging, operations	A few big winners fund the losses

68. Mutual Funds, ETFs and NAV

A **mutual fund** pools investors' money into a managed portfolio; you buy or redeem at the end-of-day **NAV** = **(assets - liabilities) / units**. An **ETF** is a similar basket but **trades intraday on an exchange** like a stock. **Active** funds try to beat a benchmark (higher fees); **passive / index** funds and most ETFs simply track it (low fees). A **SIP** invests a fixed amount regularly, averaging the entry price over time.

Interview angle: "ETF vs mutual fund?" ETFs trade intraday at market price; open-end mutual funds transact once a day at NAV.

69. Credit Ratings and the Agencies

Agencies (S&P, Moody's, Fitch globally; CRISIL, ICRA in India) grade creditworthiness. **Investment grade** is BBB-/Baa3 and above; **speculative or "junk"** is below. A **downgrade raises borrowing costs and widens spreads**, and can force funds with rating mandates to sell.

Interview angle: Know the investment-grade cutoff (BBB-) and that a downgrade widens credit spreads.

70. Statistics for Finance

Mean is the average; **standard deviation** measures dispersion and is the standard proxy for **risk/volatility**; **variance** is its square. Under a **normal distribution** about 68% of outcomes fall within 1 standard deviation and 95% within 2. **Correlation** (-1 to +1) measures co-movement but is **not causation**. A **regression** fits a line to data, beta is the slope of a stock's returns on the market's. Real markets have **skew and fat tails**, so they are riskier than a normal assumes.

Interview angle: "What does standard deviation measure in finance?" Volatility, the usual proxy for total risk.

71. Behavioral Finance: Common Biases

Markets are not fully rational. Key biases: **overconfidence**, **anchoring** (fixating on a reference number), **loss aversion** (losses hurt roughly twice as much as equal gains), **herding**, **recency bias**, **confirmation bias**, and **mental accounting**. They help explain bubbles, momentum, and why investors sell winners too early and cling to losers.

Interview angle: "Give an example of a bias in action." Loss aversion, holding a losing stock hoping to break even instead of cutting it.

72. ESG and Sustainable Finance

ESG = **Environmental, Social, Governance** factors used to judge non-financial risk and impact. Investors integrate ESG to manage long-term risks (regulation, stranded assets, reputation) that standard financials miss. Instruments include **green bonds** and **sustainability-linked loans**. Live debates: **greenwashing** and inconsistent ratings. Disclosure is increasingly mandated (BRSR in India).

Interview angle: "Why do investors care about ESG?" It flags long-term risks not visible in the financials and increasingly affects the cost of capital.

73. Fintech, Digital Payments and Blockchain/Crypto

Fintech is technology reshaping finance: UPI and digital payments, neobanks, robo-advisors, digital lending, and open banking. **Blockchain** is a shared, tamper-evident ledger; **cryptocurrencies** (Bitcoin) are assets recorded on one; **stablecoins** peg to a fiat currency; **CBDCs** are central-bank digital currencies (the digital rupee). Know the promise (speed, access, lower cost) and the risks (volatility, fraud, unclear regulation).

Interview angle: Explain a blockchain in one line, "a shared, tamper-evident ledger," and UPI's impact on Indian payments.

74. Indian Market Structure and Regulation

Regulators: **SEBI** (securities), **RBI** (banking and monetary policy), **IRDAI** (insurance). Exchanges: **NSE and BSE**, with benchmark indices **Nifty 50 and Sensex**. Shares are held in **demat** form through depositories (NSDL, CDSL). India settles trades at **T+1**, among the fastest in the world. Segments: cash, F&O, currency, and commodities.

Interview angle: Know SEBI's role, the two exchanges and indices, demat, and T+1 settlement.

75. Market Sizing and Guesstimates

Case-style questions ("how many credit cards are there in India?") test **structured thinking, not a precise number**. Approach: state your assumptions, break the problem down **top-down** (population to segments to penetration to per-capita) or bottom-up, do round-number arithmetic, then **sanity-check** the answer. Talk through the structure out loud.

Interview angle: They grade the framework and assumptions, not the exact figure. Show clean structure and reasonable numbers.

76. The Behavioral Interview: Fit and STAR

Every process tests fit. Prepare tight answers for: **walk me through your resume** (a 90-second narrative), **why finance / why this firm / why this role**, **strengths and weaknesses**, a **failure or conflict**, and **where you see yourself**. Use **STAR** (Situation, Task, Action, Result) for "tell me about a time" questions, and bring thoughtful questions to ask them.

Interview angle: Structure and specificity beat generic enthusiasm. Tie every answer back to the role and know your own resume cold.

77. Bond Types and Features

Core features: **face value**, **coupon** (fixed, floating, or zero), **maturity**, and **issuer**. Types: **government** (sovereign, safest), **corporate** (credit risk), **zero-coupon** (no coupon, issued at a discount), **convertible** (can convert to equity), **inflation-linked** (principal indexed), and **callable/putable**. Seniority and whether the bond is secured drive recovery in a default.

Interview angle: "What is a zero-coupon bond?" It pays no coupon, is issued below par, and redeems at par, the return is the discount.

78. Corporate Governance

The system of rules and incentives that aligns management with shareholders and other stakeholders: an **independent board**, **audit and remuneration committees**, separation of chairman and CEO, transparent disclosure, and credible external auditors. Weak governance (dominant promoters, related-party abuse, rubber-stamp boards) is a major risk, especially in emerging markets. It is the practical answer to the agency problem (topic 63).

Interview angle: Name mechanisms that protect minority shareholders, independent directors, an audit committee, and full disclosure.

79. Securitization and Asset-Backed Securities

Securitization pools illiquid loans (mortgages, auto loans, receivables) and sells **tradable securities backed by their cash flows**, usually sliced into **tranches** by risk (senior down to equity). It frees the originator's capital and transfers risk to investors. Opaque, mispriced mortgage-backed securities and CDOs were at the center of the **2008 financial crisis**.

Interview angle: Securitization turns illiquid loans into tradable, tranching securities, know its role in 2008.

Institutional and specialist topics. Money markets, credit derivatives, alternatives, insurance, bank regulation, compliance, and corporate actions, the deeper, institution-specific ground that comes up for treasury, banking, risk, and markets roles.

80. Money Market Instruments

Short-term (under one year), highly liquid, low-risk debt used to fund operations and park cash:

- **Treasury bills:** government paper, issued at a discount with no coupon.
- **Commercial paper:** unsecured short-term corporate note.
- **Certificates of deposit:** issued by banks.
- **Call / notice money:** interbank overnight lending.

Yields track the policy rate.

Interview angle: "What is a T-bill?" A short-term government security issued below face value; the return is the discount, not a coupon.

81. Repo, Reverse Repo and the Central Bank's Liquidity Tools

A **repo** is a sale of securities with an agreement to buy them back shortly, effectively a **collateralized short-term loan**. The **repo rate** is the rate at which the RBI lends to banks (injecting liquidity); **reverse repo** absorbs it. **CRR** (cash reserve ratio) and **SLR** (statutory liquidity ratio) are reserves banks must hold. Together they steer money supply and short-term rates.

Interview angle: "The central bank raises the repo rate, what happens?" Borrowing gets costlier, liquidity tightens, aiming to cool inflation.

82. Credit Default Swaps and Credit Derivatives

A **CDS** is insurance against a borrower defaulting: the buyer pays a periodic premium (the **spread**) and receives a payout if a credit event occurs. It is used to **hedge** credit exposure or to **speculate** on a name's creditworthiness. CDS spreads are a live market read on default risk. Naked CDS positions amplified the 2008 crisis (AIG).

Interview angle: "What does a widening CDS spread tell you?" The market sees rising default risk for that reference entity.

83. Hedge Funds and Their Strategies

Lightly regulated pooled vehicles for sophisticated investors, typically charging "**2 and 20**" (2% management + 20% performance) and often using leverage, short-selling, and derivatives. Strategies: **long/short equity, global macro, event-driven / merger arbitrage, relative-value arbitrage, and quant**. They target **absolute** returns, not benchmark-relative ones.

Interview angle: Know the "2 and 20" fee model and that hedge funds can go short and chase absolute returns.

84. Alternative Investments

Assets beyond stocks, bonds, and cash: **private equity, venture capital, hedge funds, real estate, infrastructure, commodities, private credit, and collectibles**. The appeal is **diversification** (low correlation with public markets) and higher return potential; the trade-offs are **illiquidity, high fees, opacity, and hard-to-verify valuations**.

85. REITs, InvITs and Real Estate / Infrastructure Finance

A **REIT** (real estate investment trust) pools capital to own income-producing property and must distribute most of its income as dividends, giving retail investors liquid access to real estate. An **InvIT** does the same for infrastructure (roads, power lines). **Project finance** funds large infrastructure on a **non-recourse** basis, repaid only from the project's own cash flows.

Interview angle: "What is a REIT?" A listed vehicle that owns income property and passes rental income to investors as dividends.

86. Insurance: How Insurers Make Money

Two profit engines: **underwriting profit** (premiums minus claims and expenses) and **investment income on the float** (premiums held before claims are paid). Core concepts: pooling and the **law of large numbers, adverse selection, moral hazard**, and **reinsurance** (insurers insuring their own risk). Life and general (non-life) insurance differ in duration and reserving.

Interview angle: "What is insurance float?" Premiums the insurer holds and invests before paying claims, a major source of investment income.

87. Treasury and Asset-Liability Management (ALM)

A **treasury** manages a firm's or bank's cash, funding, liquidity, and financial risk (rate and FX). **ALM** matches the maturity and rate profile of assets and liabilities to control interest-rate and liquidity risk. A bank funding long-dated loans with short-term deposits runs a **maturity mismatch**. Tools: gap analysis, duration matching, and hedging.

Interview angle: "What risk does a bank run borrowing short and lending long?" Interest-rate and liquidity risk from the maturity mismatch, the classic bank failure mode.

88. Trade Finance

Instruments that de-risk cross-border trade. A **letter of credit (LC)** is a bank's guarantee that the seller is paid once correct shipping documents are presented. **Factoring** sells receivables to a financier for immediate cash; **forfaiting** does the same for medium-term export receivables. Bank guarantees and export credit round it out.

Interview angle: "What is a letter of credit?" A bank's promise to pay the seller on the buyer's behalf once documents are in order, bridging trust between distant parties.

89. Basel III and Bank Capital / Liquidity

The global bank rulebook after 2008. **Capital:** banks hold minimum capital against **risk-weighted assets** (the CAR), with a core **CET1** (common equity tier 1) floor. **Liquidity:** the **LCR** (liquidity coverage ratio) ensures 30 days of high-quality liquid assets under stress; the **NSFR** (net stable funding ratio) pushes for stable funding. The goal is safer, better-capitalized banks.

Interview angle: "What does the LCR ensure?" Enough high-quality liquid assets to survive a 30-day funding stress.

90. AML, KYC and Compliance

Anti-money-laundering rules require firms to **Know Your Customer** (verify identity), monitor for suspicious activity, and file suspicious-transaction reports. Laundering runs in three stages: **placement, layering, integration**. Failures bring heavy fines and reputational damage, so compliance is central in banking and operations roles.

Interview angle: Know the three stages of money laundering and that KYC is the foundation of AML.

91. Capital Budgeting Techniques

Choosing which long-term projects to fund. Beyond **NPV** (best, absolute value added) and **IRR** (topic 9):

- **Payback period:** years to recover the outlay, simple but **ignores the time value of money** and cash flows after payback.
- **Discounted payback:** payback using discounted cash flows.
- **Profitability index:** PV of inflows / initial outlay; accept if above 1.

Interview angle: "Weakness of the payback method?" It ignores the time value of money and everything that happens after the cutoff.

92. Corporate Actions

Events that change a company's securities: **dividends**, **stock splits** (more shares at a lower price, no change in value), **bonus / scrip issues** (free shares from reserves), **rights issues** (existing holders can buy more, usually at a discount), and **buybacks**. They move the share count and per-share metrics but often leave total value unchanged.

Interview angle: "Does a 2-for-1 split create value?" No, twice the shares at half the price; total value is unchanged. It improves liquidity and affordability.

93. Foreign Exchange, Deeper

Spot settles in two days; a **forward** locks a rate for a future date. The forward rate reflects the **interest-rate differential** (covered interest parity). A **carry trade** borrows a low-yield currency to invest in a high-yield one, earning the spread but bearing FX risk. Currencies can **float, be pegged, or be managed**; a **non-deliverable forward (NDF)** cash-settles restricted currencies.

Interview angle: "What drives the forward FX rate?" The interest-rate differential between the two currencies (interest-rate parity), not a forecast of the spot.

94. Commodities and Gold

Real assets traded via **spot and futures**: energy (oil, gas), metals (gold, copper), and agriculture. Prices reflect supply and demand, inventories, and the dollar. **Gold** is a store of value and an inflation and uncertainty hedge, but it pays no cash flow, so it competes poorly when real rates are high. **Contango** (futures above spot) versus **backwardation** (futures below spot) drives the roll return on commodity futures.

Interview angle: "Why hold gold?" As a hedge against inflation and uncertainty. Its weakness is no yield, so high real rates hurt it.

Advanced and technical topics. Deeper valuation, capital-structure theory, market efficiency, factor models, risk measurement, and options and fixed-income math, the technical ground probed for quant, risk, research, and banking interviews.

95. The Dividend Discount Model (DDM)

Values a stock as the present value of its future dividends.

Gordon growth (constant-growth DDM)

$$\text{Price} = D1 / (r - g)$$

D1 = next year's dividend, r = cost of equity, g = perpetual dividend growth (g must be below r)

A multi-stage DDM models a high-growth phase then a terminal Gordon phase. It suits **stable dividend payers** (banks, utilities, consumer staples) and is useless for firms that pay no dividend.

Interview angle: "When is DDM the right tool?" For mature companies with predictable dividends; use DCF or multiples otherwise.

96. Economic Value Added (EVA) and Residual Income

$$\text{EVA} = \text{NOPAT} - (\text{Invested Capital} \times \text{WACC})$$

the profit left after charging for the cost of ALL capital employed

Positive EVA means the firm earns **above its cost of capital**, the definition of value creation (the same idea as $\text{ROIC} > \text{WACC}$, topic 21). Residual-income valuation adds the present value of future residual income to today's book value.

Interview angle: "What does positive EVA tell you?" The company is earning more than the cost of the capital it deploys, so it is creating value.

97. Capital Structure Theories

Building on Modigliani-Miller (topic 11):

- **Trade-off theory:** an optimal debt level balances the **tax shield** benefit against rising **financial-distress** costs.
- **Pecking order theory:** firms prefer **internal funds first, then debt, then equity last**, because issuing equity signals the stock may be overvalued.
- **Market timing:** firms issue equity when it looks richly priced.

Interview angle: "Why might a healthy firm avoid issuing equity?" Pecking-order signaling, an equity raise can suggest management thinks shares are overvalued.

98. The Efficient Market Hypothesis (EMH)

Prices reflect available information, so you cannot consistently beat the market on a risk-adjusted basis. Three forms:

- **Weak:** prices reflect all past prices, so technical analysis should not work.
- **Semi-strong:** prices reflect all public information, so fundamental analysis on public data should not give an edge.
- **Strong:** prices reflect even private information.

Anomalies and behavioral biases (topic 71) are the standard challenges.

Interview angle: "If markets are semi-strong efficient, what stops working?" Trading on public news, it is already in the price.

99. Multi-Factor Models: Fama-French and APT

CAPM uses one factor (the market). **Fama-French** adds **size** (small minus big) and **value** (high minus low book-to-market); the five-factor version adds profitability and investment. **Arbitrage Pricing Theory (APT)** allows several macro factors. These explain returns better than beta alone and underpin factor and smart-beta investing.

Interview angle: Name the three Fama-French factors, market, size, and value, and note they extend CAPM.

100. Value at Risk and Expected Shortfall

VaR is the maximum loss over a horizon at a confidence level, e.g. a 1-day 99% VaR of \$10m means a 1% chance of losing more than \$10m in a day. Methods: **historical, parametric (variance-covariance), and Monte Carlo**. Its flaw: VaR says nothing about **how bad** losses get beyond the threshold. **Expected Shortfall (CVaR)** fixes this by averaging the losses in the tail beyond VaR.

Interview angle: "What is the main weakness of VaR?" It ignores the severity of tail losses beyond the cutoff; expected shortfall captures them.

101. Stress Testing, Scenario Analysis and Backtesting

Stress testing pushes a portfolio through extreme but plausible shocks (2008, a rate spike) to reveal losses beyond normal VaR. **Scenario analysis** models specific what-ifs. **Backtesting** checks whether a VaR model's actual breaches match its stated confidence over history, too many breaches means the model understates risk. Regulators require all three.

Interview angle: "Why stress test if you already compute VaR?" VaR assumes normal markets; stress tests capture the tail events it misses.

102. Put-Call Parity and Black-Scholes Intuition

Put-call parity (European, no dividends)

$$\text{Call} - \text{Put} = \text{Spot} - \text{PV}(\text{Strike})$$

a no-arbitrage identity: break it and a riskless profit exists

Black-Scholes prices a European option from spot, strike, time, the risk-free rate, and **volatility**. Volatility is the only unobservable input, so the market effectively quotes options in **implied-vol** terms.

Interview angle: "Which Black-Scholes input can't you observe?" Volatility, which is exactly why traders talk in implied volatility.

103. Implied vs Realized Volatility and the Skew

Realized volatility is what actually happened; **implied** volatility is what option prices imply about the future. Selling options when implied exceeds realized is a common edge. The **volatility skew/smile** is that implied vol varies by strike, out-of-the-money puts usually cost more (crash protection), which contradicts Black-Scholes' constant-vol assumption.

Interview angle: "What is the volatility skew?" Implied vol differs across strikes, typically higher for downside puts, reflecting demand for crash protection.

104. Duration, Convexity and Immunization

Duration (topic 14) is the first-order interest-rate sensitivity; **convexity** is the second-order correction, the curvature of the price-yield curve. Convexity **helps the bondholder**, prices rise more when yields fall than they drop when yields rise by the same amount. **Immunization** matches asset and liability duration so a portfolio is insulated from small rate moves (used by pensions and insurers).

Interview angle: "Why is positive convexity good?" Gains from falling yields exceed losses from equal rising yields.

105. The Yield Curve Deeper: Spot, Forward and Bootstrapping

The **spot (zero) rate** is the yield on a single cash flow at each maturity. **Forward rates** are the future rates between two dates implied by today's curve, via no-arbitrage. **Bootstrapping** derives spot rates step

by step from coupon-bond prices. This zero curve is what actually discounts fixed-income and swap cash flows.

Interview angle: "What is a forward rate?" The rate for a future period implied today by the spot curve, locked in by no-arbitrage, not a forecast.

106. Sum-of-the-Parts and Adjusted Present Value

Sum-of-the-parts (SOTP) values a diversified company by valuing each segment separately (its own multiple or DCF) and adding them, useful for conglomerates trading at a discount. **Adjusted Present Value (APV)** values the business as if all-equity, then adds the present value of financing effects (mainly the tax shield); it is cleaner than WACC when leverage changes over time.

Interview angle: "When use APV instead of a WACC-based DCF?" When capital structure shifts materially over the forecast, such as an LBO paying down debt.

107. Counterparty Risk, CVA and Central Clearing

Counterparty risk is that the other side of a derivative defaults before settlement. **CVA (credit valuation adjustment)** is the market price of that risk, subtracted from a derivative's value. After 2008, reforms route standardized derivatives through **central counterparties (CCPs)** with margin, which mutualizes and reduces counterparty risk.

Interview angle: "What is CVA?" The adjustment to a derivative's value for the chance the counterparty defaults.

108. Expected Credit Loss (IFRS 9) and Provisioning

IFRS 9 replaced the old "incurred loss" approach with a **forward-looking expected credit loss (ECL)** model: banks provision for expected losses from day one and move exposures through three stages as credit quality deteriorates.

Expected Loss = PD x LGD x EAD

probability of default x loss given default x exposure at default

Interview angle: "How do banks size loan provisions now?" Forward-looking expected credit loss (PD x LGD x EAD), not just losses already incurred.

109. Algorithmic Trading, HFT and Market Microstructure

Algorithmic trading executes orders by rules and code. **High-frequency trading (HFT)** is a subset competing on **speed** (co-location, ultra-low latency) for tiny per-trade edges; it adds liquidity but raises fairness and flash-crash concerns. **Market microstructure** studies how orders become prices, the order book, the bid-ask spread, market impact, and adverse selection.

Interview angle: Know that HFT competes on latency, and that microstructure explains how order flow actually moves prices.

Bottom line: concepts 1, 2, 5, 6, 7, 8 are the backbone of technical finance interviews. If time is short, drill those six until you can explain each in under a minute, then use the 1,064 Q&A below to go deep in your target area.

A. Core Finance & Accounting

A1. Accounting & the Three Financial Statements (Q1-Q40)

Q1. What are the three primary financial statements and what does each one show?

The income statement (P&L) shows profitability over a period: revenue minus expenses equals net income. The balance sheet is a snapshot at a point in time showing $\text{Assets} = \text{Liabilities} + \text{Shareholders' Equity}$. The cash flow statement reconciles the change in cash over a period across operating, investing and financing activities. Together they answer: did the firm make money, what does it own/owe, and where did the cash actually go.

Q2. What is the fundamental accounting equation, and why must it always balance?

$\text{Assets} = \text{Liabilities} + \text{Shareholders' Equity}$. It must balance because every economic transaction has two sides under double-entry bookkeeping: anything a company owns (assets) is financed either by what it owes (liabilities) or by owners' capital (equity). If you buy a \$100 machine with cash, one asset rises and another falls; if you buy it on credit, an asset and a liability both rise by \$100, keeping the equation intact.

Q3. What is the difference between accrual accounting and cash accounting?

Accrual accounting records revenue when it is earned and expenses when they are incurred, regardless of when cash moves. Cash accounting records transactions only when cash is received or paid. Accrual gives a truer economic picture (and is required under IFRS, US GAAP and Ind-AS for most entities), while cash accounting is simpler and shows actual liquidity. The gap between the two is exactly what working-capital and accrual adjustments on the cash flow statement capture.

Q4. What is double-entry bookkeeping?

Double-entry means every transaction is recorded in at least two accounts with equal debits and credits, so the books always balance. For example, paying \$500 rent debits rent expense (increasing expenses) and credits cash (decreasing an asset) by \$500. This system underpins the accounting equation and makes errors detectable because total debits must equal total credits.

Q5. On the income statement, walk through the line items from revenue to net income.

Start with Revenue (sales), subtract Cost of Goods Sold (COGS) to get Gross Profit. Subtract operating expenses (SG&A, R&D, depreciation/amortization) to get Operating Income (EBIT). Subtract interest to get Pre-Tax Income (EBT), then subtract taxes to arrive at Net Income. Net income is the "bottom line" available to equity holders, from which dividends are paid and retained earnings are built.

Q6. What is the revenue recognition principle?

Revenue is recognized when it is earned and realizable, not necessarily when cash is received. Under the modern five-step model (IFRS 15 / Ind-AS 115 / ASC 606), you identify the contract and performance obligations, determine and allocate the transaction price, and recognize revenue as obligations are satisfied. For example, a SaaS firm billing \$1,200 upfront for an annual subscription recognizes \$100 per month, holding the rest as deferred (unearned) revenue, a liability.

Q7. What is the matching principle?

The matching principle says expenses should be recognized in the same period as the revenues they help generate. For instance, COGS for a product is recorded when the product is sold (matching the sale revenue), not when the inventory was purchased. Depreciation likewise spreads an asset's cost over the periods it generates revenue. This prevents profit distortion from timing mismatches between cost and benefit.

Q8. How are accounts receivable and accounts payable defined, and where do they sit?

Accounts receivable (AR) is money owed to the company by customers for goods/services already delivered on credit; it is a current asset. Accounts payable (AP) is money the company owes to suppliers for goods/services received; it is a current liability. Rising AR means revenue booked but cash not yet collected (a use of cash); rising AP means costs incurred but not yet paid (a source of cash).

Q9. What is the difference between a prepaid expense and an accrued expense?

A prepaid expense is cash paid in advance for a future benefit (e.g., paying 12 months' insurance upfront); it is a current asset that is expensed over time. An accrued expense is a cost incurred but not yet paid (e.g., salaries earned but unpaid at month-end); it is a current liability. Prepaid = pay now, expense later; accrued = expense now, pay later.

Q10. What is deferred (unearned) revenue and how does it differ from accrued revenue?

Deferred revenue is cash collected before the good/service is delivered, so it is a liability until earned (e.g., an airline ticket sold for a future flight). Accrued revenue is revenue earned but not yet billed or collected, so it is an asset (e.g., interest earned but not yet received). They are mirror images: one is "cash before work," the other is "work before cash."

Q11. What is depreciation, and why does it appear on the income statement?

Depreciation allocates the cost of a tangible fixed asset (plant, machinery, vehicles) over its useful life to match the expense against the revenue it helps produce. It is a non-cash expense on the income statement that reduces taxable income without an immediate cash outflow (the cash went out when the asset was bought). On the cash flow statement it is added back to net income in operating activities.

Q12. Explain straight-line depreciation versus the written-down value (declining balance) method.

Straight-line charges an equal amount each year: $(\text{Cost} - \text{Salvage}) / \text{Useful Life}$. Written-down value (WDV) / declining balance applies a fixed percentage to the reducing book value each year, so charges are higher early and lower later. Example: a \$10,000 asset, 10-year life, no salvage, gives \$1,000/year straight-line; at 20% WDV it gives \$2,000 in year 1, \$1,600 in year 2, and so on. WDV is common for tax in India (Income Tax Act uses block-of-assets WDV), while companies often use straight-line for book reporting.

Q13. What is amortization and how does it differ from depreciation?

Amortization spreads the cost of intangible assets (patents, software, licenses, customer lists) over their useful life, while depreciation does the same for tangible assets. Mechanically they are identical (usually straight-line for intangibles) and both are non-cash. A key nuance: indefinite-life intangibles like goodwill are not amortized under IFRS/Ind-AS/US GAAP but are tested for impairment instead.

Q14. How does a \$100 increase in depreciation flow through the three statements (assume a 25% tax rate)?

Income statement: pre-tax income falls \$100, tax falls \$25, so net income falls \$75. Cash flow: start with net income down \$75, add back the \$100 non-cash depreciation, so cash from operations rises \$25 (the tax shield). Balance sheet: cash up \$25 (assets), PP&E down \$100 (assets), so total assets down \$75; retained earnings down \$75 (equity), and the balance sheet balances. This is the classic interview link question.

Q15. What is the capitalize-versus-expense decision and why does it matter?

If a cost provides benefits beyond the current period, it is capitalized as an asset and depreciated/amortized over time; if it only benefits the current period, it is expensed immediately. Capitalizing boosts near-term profit (cost is spread out) and creates an asset on the balance sheet, while expensing reduces current profit. The choice affects EBITDA, net income timing and cash flow classification (capex is investing, opex is operating), which is why aggressive capitalization is a common earnings-quality red flag.

Q16. What is working capital and how is it read off the balance sheet?

Working capital = Current Assets - Current Liabilities. Operating (net) working capital usually excludes cash and debt: $(AR + Inventory + Prepaids) - (AP + Accrued liabilities)$. It measures the short-term capital tied up in running the business; rising operating working capital consumes cash, while falling working capital releases cash. It is the bridge between accrual profit and operating cash flow.

Q17. How do changes in working capital affect the cash flow statement?

In the indirect method, an increase in an operating current asset (e.g., AR or inventory) is a use of cash (subtracted), and an increase in an operating current liability (e.g., AP) is a source of cash (added). Intuition: if AR grows, you booked revenue but did not collect cash; if AP grows, you booked an expense but did not pay cash. So fast-growing firms often show strong profit but weak operating cash flow because working capital absorbs cash.

Q18. What are the three sections of the cash flow statement?

Operating activities: cash from core business operations (net income adjusted for non-cash items and working-capital changes). Investing activities: cash for buying/selling long-term assets and investments (capex, acquisitions, asset sales). Financing activities: cash from/to capital providers (issuing or repaying debt, issuing or buying back equity, paying dividends). The three net to the change in cash for the period, which ties to the balance sheet cash line.

Q19. What is the difference between the direct and indirect methods of presenting operating cash flow?

The direct method lists actual cash inflows and outflows (cash from customers, cash paid to suppliers and employees). The indirect method starts from net income and adjusts for non-cash items (depreciation, etc.) and working-capital changes to arrive at operating cash flow. Both yield the same operating cash figure; the indirect method is far more common in practice because it reconciles easily to the income statement, though IFRS/Ind-AS and US GAAP encourage (but rarely see) the direct method.

Q20. What is EBITDA and how does it differ from net income?

EBITDA = Earnings Before Interest, Taxes, Depreciation and Amortization, roughly Operating Income + D&A. It strips out financing (interest), tax regime and non-cash D&A to approximate operating cash generation and enable cross-company comparison. Net income is the bottom line after all of those, reflecting actual returns to equity. EBITDA is not a GAAP/Ind-AS line item and ignores capex and working-capital needs, so it can overstate true cash flow.

Q21. Why is EBITDA a poor proxy for free cash flow?

EBITDA excludes three real cash drains: capital expenditure (capex), changes in working capital, and cash taxes and interest. A capital-intensive business can show high EBITDA yet generate little free cash flow once it reinvests in PP&E. Free cash flow to firm is roughly $EBIT \times (1 - \text{tax}) + D\&A - \text{capex} - \text{change in working capital}$, which is why analysts move from EBITDA toward FCF for valuation.

Q22. What are the main inventory costing methods?

FIFO (First-In, First-Out) assumes the oldest inventory is sold first. LIFO (Last-In, First-Out) assumes the newest is sold first. Weighted average cost assigns the average per-unit cost to all units. Each yields different COGS and ending-inventory values when prices change, which directly affects gross profit, taxes and the balance sheet inventory figure.

Q23. In a period of rising prices, how do FIFO and LIFO affect profit and taxes?

Under rising prices, FIFO charges older (cheaper) costs to COGS, giving lower COGS, higher gross profit, higher taxes, and a higher (closer-to-current) inventory value on the balance sheet. LIFO charges newer (more expensive) costs to COGS, giving higher COGS, lower profit, lower taxes, and an understated inventory balance. That tax advantage is the main reason US firms use LIFO.

Q24. Do IFRS and Ind-AS permit LIFO?

No. IFRS (IAS 2) and Indian Ind-AS 2 prohibit LIFO; only FIFO and weighted-average cost are allowed. LIFO is permitted under US GAAP, which is a notable GAAP-vs-IFRS difference. Consequently Indian and most non-US companies report inventory on FIFO or weighted-average, improving comparability and reducing balance-sheet understatement.

Q25. How is inventory valued under "lower of cost or net realizable value"?

Under IFRS/Ind-AS (IAS 2 / Ind-AS 2), inventory is carried at the lower of cost and net realizable value (NRV = estimated selling price minus costs to complete and sell). If NRV falls below cost, a write-down is recognized as an expense. US GAAP uses "lower of cost or market" for LIFO/retail method and lower of cost or NRV otherwise. IFRS/Ind-AS also allow reversal of prior write-downs if NRV recovers, which US GAAP generally does not.

Q26. What is goodwill and when does it arise?

Goodwill is the premium paid in an acquisition over the fair value of the acquired company's identifiable net assets. It is calculated as $\text{Purchase Price} - \text{Fair Value of Identifiable Net Assets Acquired}$ and is recorded as an intangible asset on the acquirer's consolidated balance sheet. It represents unidentifiable value such as brand, customer relationships, synergies and workforce. Goodwill only appears via acquisitions, never for internally generated goodwill.

Q27. How is goodwill treated after acquisition under IFRS/Ind-AS and US GAAP?

Goodwill is not amortized; instead it is tested for impairment at least annually (and whenever there are indicators). If its recoverable/fair value falls below carrying value, an impairment loss is recognized on the income statement and goodwill is written down. Under IFRS/Ind-AS impairment testing is at the cash-generating-unit level using a one-step recoverable-amount test; US GAAP tests at the reporting-unit level. Goodwill impairments are non-cash and cannot be reversed once recognized.

Q28. What is asset impairment and how does it differ from depreciation?

Impairment is a sudden, often one-time write-down when an asset's carrying value exceeds its recoverable amount (the higher of fair value less costs to sell and value in use). Depreciation is a planned, systematic allocation of cost over an asset's life. Impairment reflects an unexpected loss of value (technology obsolescence, market decline), is recognized as an expense, and under IFRS/Ind-AS can be reversed for assets other than goodwill if conditions improve, whereas US GAAP prohibits reversals.

Q29. What is a deferred tax liability (DTL) and what commonly creates one?

A DTL is tax that a company will owe in the future due to a temporary timing difference between book and tax accounting, where book profit currently exceeds taxable profit. The classic cause is using straight-line depreciation for books but accelerated/WDV depreciation for tax: tax depreciation is higher early, so taxable income (and tax paid) is lower now but will be higher later. The deferred amount sits as a liability until the timing difference reverses.

Q30. What is a deferred tax asset (DTA) and what commonly creates one?

A DTA is a future tax benefit arising when a company has paid more tax (or recognized more taxable income) now than book accounting implies, so it will pay less later. Common sources are carried-forward tax losses, provisions/warranty expenses recognized for books but deductible only when paid, and revenue taxed before it is booked. A DTA is only recognized to the extent future taxable profits are probable; otherwise a valuation allowance reduces it.

Q31. Why might effective tax rate differ from the statutory tax rate?

The effective tax rate ($\text{tax expense} / \text{pre-tax income}$) differs from the statutory rate because of permanent differences (non-deductible expenses, tax-exempt income), tax credits, different jurisdictional rates, changes in deferred tax balances, and valuation allowances on DTAs. For example, in India the headline corporate

rate may be 25% but generous exemptions, MAT credits or non-deductible items can push the effective rate above or below that. Reconciling the two is a standard earnings-quality check.

Q32. What is a provision and how does it differ from an accrual?

A provision is a liability of uncertain timing or amount, recognized when a present obligation exists, an outflow is probable, and the amount can be reliably estimated (IAS 37 / Ind-AS 37), e.g., warranty provisions, restructuring or litigation provisions. An accrual is a known, near-certain liability with only minor estimation (e.g., accrued wages). Provisions involve genuine measurement uncertainty; accruals are essentially timing entries for known amounts.

Q33. What is a contingent liability and how is it disclosed?

A contingent liability is a possible obligation depending on an uncertain future event, or a present obligation that is not probable or cannot be reliably measured (e.g., a pending lawsuit or guarantee). Under IAS 37 / Ind-AS 37 it is not recognized on the balance sheet but is disclosed in the notes; it is recorded as a provision only if an outflow becomes probable and estimable. If the outflow is remote, no disclosure is required.

Q34. How do you distinguish recognizing a provision from disclosing a contingent liability?

It hinges on probability and measurability. If a present obligation is probable (more likely than not) and reliably estimable, you recognize a provision as a liability and an expense. If the obligation is only possible, or probable but not reliably measurable, you disclose it as a contingent liability in the notes. If remote, you do nothing. This decision tree from IAS 37 / Ind-AS 37 prevents both premature liability recognition and hidden obligations.

Q35. What are the key differences between US GAAP and IFRS?

US GAAP is rules-based; IFRS is principles-based. Major differences: LIFO is allowed under US GAAP but banned under IFRS; IFRS permits revaluation of PP&E and reversal of impairments (except goodwill), US GAAP generally does not; development costs can be capitalized under IFRS (IAS 38) but are mostly expensed under US GAAP; and presentation/terminology differs. These cause real comparability issues for cross-listed firms and analysts.

Q36. What is Ind-AS and how does it relate to IFRS?

Ind-AS (Indian Accounting Standards) is India's set of accounting standards converged with IFRS, notified under the Companies Act and applicable to larger listed and unlisted companies based on net-worth thresholds. It is "converged with" rather than identical to IFRS, with carve-outs (e.g., certain treatments for foreign-exchange differences and specific exemptions). Smaller Indian companies still follow the older Indian GAAP (AS). Like IFRS, Ind-AS bans LIFO and does not amortize goodwill.

Q37. How do interest paid and dividends paid get classified on the cash flow statement under IFRS/Ind-AS versus US GAAP?

Under US GAAP, interest paid and interest/dividends received are operating, and only dividends paid are financing. Under IFRS/Ind-AS there is a policy choice: interest and dividends paid can be classified as either operating or financing, and interest/dividends received as either operating or investing, applied consistently. This flexibility means operating cash flow is not always directly comparable across IFRS and US GAAP reporters.

Q38. A company reports strong net income but negative operating cash flow. What could explain it, and is it a red flag?

It often signals a working-capital drain: surging accounts receivable (sales booked but uncollected), inventory build-up, or shrinking payables. It can also reflect aggressive revenue recognition or large non-cash gains inflating profit without cash backing. While rapid growth can legitimately consume working

capital, a persistent gap between accrual profit and operating cash flow is a classic earnings-quality red flag worth investigating in the cash flow reconciliation.

Q39. How does issuing \$100 of new equity flow through the three statements?

Cash flow: financing activities rise \$100, increasing total cash by \$100. Balance sheet: cash (asset) rises \$100, and on the other side common stock/share capital and additional paid-in capital (equity) rise \$100, so it balances. The income statement is unaffected because raising capital is not revenue or expense. This illustrates that financing transactions hit the balance sheet and cash flow but bypass the P&L.

Q40. How do the three statements link together into one connected model?

Net income from the income statement flows into the top of the cash flow statement (operating section) and into retained earnings on the balance sheet. The cash flow statement reconciles net income to the change in cash, and its ending cash becomes the balance sheet's cash line. Non-cash items (D&A) and working-capital changes bridge the income statement and balance sheet through the cash flow statement. If all links are correct, the balance sheet balances: this interconnection is the backbone of any three-statement financial model.

A7. Advanced Financial Accounting (Q41-Q73)

Q41. Under the old standards, what was the difference between an operating lease and a finance (capital) lease, and why did it matter?

A finance lease transferred substantially all risks and rewards of ownership, so the lessee capitalized the asset and recognized a corresponding liability on the balance sheet, splitting payments into depreciation plus interest. An operating lease was treated as a simple rental: the entire payment hit the income statement as a single lease/rent expense (above EBIT) and nothing appeared on the balance sheet. The classification mattered because operating leases kept large lease obligations off-balance-sheet, understating leverage and overstating returns on assets.

Q42. What changed under IFRS 16 / Ind-AS 116, and why was the old model considered abused?

IFRS 16 (and Ind-AS 116, effective 1 April 2019 in India) introduced a single lessee model that puts almost all leases on the balance sheet, eliminating the operating-vs-finance distinction for lessees. The old split-model was criticized because companies (airlines, retailers, telecoms) structured leases to qualify as "operating" to hide debt; investors had to manually capitalize lease commitments from the footnotes. Now the lessee recognizes a right-of-use (ROU) asset and a lease liability for nearly every lease, with only short-term (≤ 12 months) and low-value leases exempted.

Q43. Walk through the initial recognition of a lease under IFRS 16 / Ind-AS 116.

The lease liability is recognized at the present value of remaining lease payments, discounted at the interest rate implicit in the lease (or the lessee's incremental borrowing rate if that is not readily determinable). The ROU asset starts equal to the lease liability, plus initial direct costs, prepaid lease payments, and estimated restoration costs, less any lease incentives received. Example: 5 annual payments of INR 100 discounted at 8% give a liability of roughly INR 399, so both the ROU asset and lease liability are booked at INR 399 on day one.

Q44. How does on-balance-sheet lease accounting affect EBITDA and key ratios?

Under IFRS 16/Ind-AS 116 the old single rent expense is replaced by depreciation of the ROU asset plus interest on the liability, both of which sit below EBITDA, so EBITDA rises (the rent that used to depress it is gone). Total assets and total debt both increase, so debt/EBITDA can move either way but leverage ratios and gross debt clearly rise, while ROA and asset turnover fall because the asset base is larger. Note US GAAP

(ASC 842) also capitalizes operating leases but keeps a straight-line single lease expense in operating costs, so under US GAAP EBITDA is NOT boosted the way it is under IFRS.

Q45. Why is the income statement expense front-loaded for a lease under IFRS 16?

Because depreciation of the ROU asset is typically straight-line while interest is calculated on a declining liability balance, total expense (depreciation + interest) is higher in early years and lower later, even though the cash payment is level. This contrasts with the old operating lease's flat rent expense. Over the full lease term total expense is identical; only the timing differs.

Q46. State the IFRS 15 / Ind-AS 115 / ASC 606 five-step revenue recognition model.

(1) Identify the contract with a customer; (2) identify the distinct performance obligations in the contract; (3) determine the transaction price; (4) allocate the transaction price to the performance obligations based on relative standalone selling prices; (5) recognize revenue as (or when) each performance obligation is satisfied, i.e., when control transfers to the customer. The core principle is that revenue reflects the consideration the entity expects to be entitled to in exchange for transferring goods or services.

Q47. What is the difference between recognizing revenue "over time" versus "at a point in time"?

Revenue is recognized over time if any of three criteria are met: the customer simultaneously receives and consumes the benefits (e.g., a SaaS subscription), the entity creates/enhances an asset the customer controls, or the asset has no alternative use and the entity has an enforceable right to payment for work done. If none apply, revenue is recognized at a point in time, typically when control transfers (delivery, title, acceptance). A construction contract usually qualifies as over-time using a percentage-of-completion measure of progress.

Q48. How does the 5-step model handle a bundled contract, e.g., a phone sold with a 24-month plan?

The transaction price is allocated to each distinct performance obligation (the handset and the network service) based on their relative standalone selling prices, not the contractual invoicing pattern. So even if the handset is "free" up front, a portion of the monthly fees is reallocated to the handset and recognized as revenue at delivery, with the rest recognized over 24 months. This often shifts revenue earlier than the cash-based view and creates a contract asset on the balance sheet.

Q49. What are contract assets and contract liabilities under IFRS 15?

A contract asset arises when the entity has satisfied a performance obligation (earned revenue) but does not yet have an unconditional right to payment (which would be a receivable) — for example, unbilled revenue. A contract liability (deferred/unearned revenue) arises when the customer has paid, or payment is due, before the entity has performed. Example: an annual SaaS plan billed upfront for INR 1,200 creates a contract liability that is released to revenue at INR 100 per month.

Q50. How is stock-based compensation (SBC) accounted for, and why is it a real expense?

Under IFRS 2 / Ind-AS 102 (and ASC 718), equity-settled SBC is measured at the grant-date fair value of the award (often via Black-Scholes or a binomial model for options) and expensed over the vesting period, with a credit to equity (additional paid-in capital). It is a genuine expense because the company is transferring value to employees in lieu of cash compensation, diluting existing shareholders. It is non-cash, so it is added back in the operating section of the cash flow statement.

Q51. Why do analysts often dislike adding SBC back to free cash flow even though it is "non-cash"?

Because SBC is a recurring economic cost: it compensates employees and, if the company wants to avoid dilution, it must spend cash buying back shares to offset the new grants. Treating it as a costless add-back (as in "adjusted EBITDA" or FCF) flatters cash generation and ignores the dilution borne by shareholders. A rigorous valuation either expenses SBC or models the share-count growth it causes.

Q52. Explain the treasury stock method (TSM) for diluted EPS.

The TSM assumes in-the-money options and warrants are exercised, the company receives the exercise proceeds, and uses those proceeds to buy back shares at the average market price. The net new shares added to the diluted count = options outstanding – (options * strike / average price). Example: 1,000 options at a INR 50 strike with a INR 100 average price raise INR 50,000, which repurchases 500 shares, so net dilution is $1,000 - 500 = 500$ shares.

Q53. What is the difference between a defined benefit (DB) and a defined contribution (DC) pension plan, and who bears the risk?

In a DC plan the employer pays a fixed contribution (e.g., into a provident/pension fund) and the employee bears all investment and longevity risk; the expense equals the contribution and there is no ongoing balance-sheet obligation. In a DB plan the employer promises a defined future benefit (e.g., a formula on final salary and years of service), so the employer bears investment and actuarial risk and must recognize a pension obligation. DB plans are far more complex because they require actuarial assumptions about discount rates, salary growth, and mortality.

Q54. What is the PBO and how is funded status determined?

The Projected Benefit Obligation (PBO) is the actuarial present value of pension benefits earned to date, measured using expected future salary levels and discounted at a high-quality corporate bond rate. Funded status = fair value of plan assets – PBO; a negative figure is an underfunded plan, shown as a net pension liability on the balance sheet. Under IAS 19 the net interest cost is computed by applying the discount rate to this net liability/asset.

Q55. Under IAS 19 / Ind-AS 19, how are the components of pension cost split between P&L and OCI?

Service cost (current and past) and net interest on the net defined benefit liability go through profit or loss. Remeasurements — actuarial gains and losses from changes in assumptions and the difference between actual and expected return on plan assets — are recognized in Other Comprehensive Income (OCI) and are not recycled to P&L. Routing volatile remeasurements through OCI keeps earnings from swinging with market and actuarial noise.

Q56. Distinguish temporary from permanent differences in deferred tax accounting.

A temporary difference is a timing gap between the book carrying value and the tax base of an asset or liability that will reverse in the future, creating a deferred tax asset (DTA) or liability (DTL) — e.g., accelerated tax depreciation vs straight-line book depreciation. A permanent difference never reverses and so creates no deferred tax — e.g., tax-exempt income, fines/penalties that are not deductible, or expenses permanently disallowed. Only temporary differences hit deferred taxes; permanent differences instead change the effective tax rate.

Q57. How does accelerated tax depreciation create a deferred tax liability? Give a quick number.

Faster depreciation for tax purposes means lower taxable income (and lower tax paid) now, but higher taxable income later when book depreciation exceeds tax depreciation, so the company owes more tax in the future — a DTL. Example: an asset is depreciated INR 300 for tax but only INR 100 for books in year 1; the INR 200 temporary difference at a 25% tax rate creates a DTL of INR 50 that reverses in later years. The DTL represents tax deferred, effectively an interest-free loan from the government.

Q58. What is a deferred tax valuation allowance, and when is it recorded?

A valuation allowance (US GAAP term) reduces a DTA to the amount that is "more likely than not" to be realized, recognizing that future taxable profits may be insufficient to use the asset. It is recorded when there is doubt — recent losses, expiring carryforwards, or a weak earnings outlook. Under IFRS/Ind-AS there is no separate allowance account; instead the DTA is recognized only to the extent realization is probable, achieving the same conservatism through a recognition test rather than a contra account.

Q59. How do net operating losses (NOLs) create deferred tax assets?

An NOL can be carried forward to offset future taxable income, reducing future cash taxes, so it is recognized as a DTA equal to the NOL times the enacted tax rate (subject to recognition/valuation-allowance tests). Example: a USD 100m NOL at a 21% rate is a USD 21m DTA. If the company is not expected to generate enough future profit before the NOL expires (where limits apply), the DTA is written down or only partially recognized.

Q60. In purchase accounting, how is goodwill created? Give the mechanics.

Goodwill = purchase consideration – fair value of identifiable net assets acquired (assets at fair value minus liabilities assumed, including any newly recognized intangibles and the related deferred taxes). It is the premium paid over the fair value of what is separately identifiable, capturing synergies, assembled workforce, and going-concern value. Example: pay INR 1,000 for a target whose identifiable net assets are fair-valued at INR 700, and goodwill is INR 300.

Q61. What is purchase price allocation (PPA) and what intangibles get recognized that weren't on the target's books?

PPA steps up the acquired assets and liabilities to fair value and recognizes identifiable intangibles the target never capitalized — customer relationships, brands/trademarks, patents, developed technology, and order backlog — with the residual going to goodwill. These intangibles are then amortized (except indefinite-lived ones like some brands), creating extra post-deal amortization that depresses GAAP earnings. The write-up also creates a deferred tax liability when the step-up is not deductible for tax.

Q62. What is a bargain purchase and how is it treated?

A bargain purchase (negative goodwill) occurs when the consideration paid is less than the fair value of identifiable net assets acquired, e.g., a forced or distressed sale. Under IFRS 3 / Ind-AS 103 the acquirer first reassesses whether all assets and liabilities were correctly identified and measured, then recognizes any remaining excess immediately as a gain in profit or loss. This contrasts with goodwill, which is capitalized; negative goodwill is never carried as a liability.

Q63. Is goodwill amortized or tested for impairment? Note the India nuance.

Under IFRS 3 / Ind-AS 103 and US GAAP, goodwill is not amortized but tested for impairment at least annually (or on a triggering event); impairment is recognized when the recoverable amount of the cash-generating unit falls below its carrying amount and cannot be reversed. Note that Indian AS (the old standard for unlisted/SME entities) did amortize goodwill, and US private companies may elect amortization — but under Ind-AS, listed Indian companies follow the impairment-only model.

Q64. When do you consolidate a subsidiary, use the equity method, or just hold it at fair value?

Control (generally >50% of voting rights, or power to direct relevant activities) triggers full consolidation, where 100% of the subsidiary's assets, liabilities, revenue and expenses are combined and a non-controlling interest is shown. Significant influence (typically 20–50%, no control) triggers the equity method, where the investment is carried at cost plus the investor's share of post-acquisition profits as a one-line item. Below significant influence, the stake is a financial asset measured at fair value (FVTPL or FVOCI) under IFRS 9 / Ind-AS 109.

Q65. What is non-controlling (minority) interest and how does it flow through the statements?

When a parent owns less than 100% of a consolidated subsidiary, non-controlling interest (NCI) represents the equity in that subsidiary not owned by the parent. On the balance sheet NCI appears within equity (separate from parent equity); on the income statement, total consolidated net income is split into the portion attributable to the parent and the portion attributable to NCI. Example: own 80% of a sub earning INR 100, and INR 20 of consolidated net income is allocated to NCI.

Q66. Contrast the equity method with proportionate consolidation for a joint arrangement.

Under the equity method (required by IFRS 11 / Ind-AS 28 for joint ventures), the investor records a single line "investment in associate/JV" and picks up its share of profit as one line, with no gross-up of the investee's assets and liabilities. Proportionate consolidation (no longer permitted under IFRS for joint ventures) would instead bring the investor's percentage share of each individual asset, liability, revenue and expense onto the face of the statements. The equity method keeps leverage and revenue understated relative to the economic exposure, which is why analysts sometimes "look through" to the gross figures.

Q67. Explain functional vs presentation currency and the resulting cumulative translation adjustment (CTA).

The functional currency is the currency of the primary economic environment in which an entity operates; the presentation (reporting) currency is the one the financials are presented in. Under IAS 21 / Ind-AS 21, a foreign subsidiary's assets and liabilities are translated at the closing rate and its income at average rates, and the resulting exchange differences accumulate in OCI as a cumulative translation adjustment (CTA) within equity. The CTA is recycled to profit or loss only when the foreign operation is disposed of.

Q68. What is the difference between the current-rate (translation) method and the temporal (remeasurement) method?

The current-rate method (used when the subsidiary's functional currency differs from the parent's) translates all assets/liabilities at the closing rate, equity at historical rates, and income at average rates, with differences going to CTA in OCI. The temporal/remeasurement method (used when the books are kept in a currency other than the functional currency, e.g., a hyperinflationary or integrated operation) translates monetary items at the closing rate but non-monetary items at historical rates, and the resulting remeasurement gain/loss flows through net income. The key distinction is OCI (translation) vs P&L (remeasurement).

Q69. What are off-balance-sheet items and special purpose entities (SPEs), and why are they a red flag?

Off-balance-sheet items are obligations or assets not shown on the face of the balance sheet — historically operating leases, certain guarantees, factored receivables, and interests in unconsolidated SPEs/VIEs. SPEs (special purpose vehicles) are legal entities created to isolate assets/liabilities, often to move debt off the sponsor's books, as infamously abused by Enron. They are a red flag because they can hide leverage and risk; modern standards (IFRS 10/Ind-AS 110's control model and IFRS 16's lease capitalization) have pulled much of this back on-balance-sheet.

Q70. How is R&D treated under US GAAP versus IFRS / Ind-AS, and what is the effect on the financials?

Under US GAAP, R&D is generally expensed as incurred (with narrow exceptions like certain software costs). Under IAS 38 / Ind-AS 38, research is expensed but development costs MUST be capitalized once technical and commercial feasibility, intention, and ability to use/sell are demonstrated, then amortized over the asset's useful life. The result is that an IFRS firm can show higher near-term earnings and an intangible asset, while an otherwise identical US GAAP firm expenses the same spend immediately — a real comparability issue when benchmarking R&D-heavy companies.

Q71. How does Ind-AS 37 / IAS 37 distinguish a provision, a contingent liability, and a contingent asset?

A provision is recognized (a liability is booked) when there is a present obligation from a past event, an outflow of resources is probable (>50% likely), and the amount can be reliably estimated — e.g., warranty costs or restructuring. A contingent liability is only disclosed, not recognized, when the outflow is possible but not probable, or the amount cannot be reliably measured. A contingent asset is never recognized and is only disclosed when an inflow is probable, recognized only when realization becomes virtually certain — reflecting accounting conservatism's asymmetry.

Q72. Explain the fair value hierarchy under IFRS 13 / Ind-AS 113 (Level 1, 2, 3).

Level 1 inputs are unadjusted quoted prices in active markets for identical assets/liabilities (e.g., a listed equity's closing price) — the most reliable. Level 2 inputs are observable but not Level 1, such as quoted prices for similar items or model inputs like yield curves and benchmark rates (e.g., an interest-rate swap valued off market curves). Level 3 inputs are unobservable, relying on the entity's own assumptions and models (e.g., a private equity stake or a bespoke illiquid derivative), and carry the highest measurement uncertainty and disclosure burden.

Q73. What is the difference between basic and diluted EPS, and how are convertibles handled?

Basic EPS = (net income – preferred dividends) / weighted average shares outstanding. Diluted EPS adjusts the denominator (and sometimes the numerator) for all dilutive potential shares — options/warrants via the treasury stock method, and convertible bonds/preferreds via the "if-converted" method, which adds the conversion shares to the denominator and adds back the after-tax interest (or preferred dividend) to the numerator. Anti-dilutive securities (those that would raise EPS) are excluded, so diluted EPS is always less than or equal to basic EPS.

A2. Financial Ratios & Analysis (Q74-Q98)

Q74. What is the current ratio and what does it measure?

The current ratio measures short-term liquidity: a firm's ability to cover liabilities due within a year using assets convertible to cash within a year. $\text{Current Ratio} = \text{Current Assets} / \text{Current Liabilities}$. Example: current assets of INR 600 cr and current liabilities of INR 300 cr give a ratio of 2.0, meaning INR 2 of current assets per INR 1 of near-term obligations. A ratio above 1 is generally safe, though the "ideal" level is industry-specific.

Q75. How does the quick (acid-test) ratio differ from the current ratio?

The quick ratio strips out inventory (and often prepaid expenses) because inventory is the least liquid current asset and may not convert to cash quickly. $\text{Quick Ratio} = (\text{Current Assets} - \text{Inventory}) / \text{Current Liabilities}$, or equivalently $(\text{Cash} + \text{Marketable Securities} + \text{Receivables}) / \text{Current Liabilities}$. If current assets are 600 with 200 of inventory and current liabilities are 300, the quick ratio is $(600-200)/300 = 1.33$ versus a current ratio of 2.0. It is a stricter test favored for inventory-heavy businesses.

Q76. What is the cash ratio and when is it most relevant?

The cash ratio is the most conservative liquidity measure, counting only cash and cash equivalents (plus marketable securities) against current liabilities. $\text{Cash Ratio} = (\text{Cash} + \text{Marketable Securities}) / \text{Current Liabilities}$. It assumes receivables and inventory may not materialize, so it matters most in distress, banking stress tests, or for assessing a firm that cannot rely on collections. A cash ratio of 0.5 means half of short-term debts could be paid from cash on hand immediately.

Q77. Define gross, operating, and net profit margins.

These are profitability ratios scaled to revenue. $\text{Gross Margin} = \text{Gross Profit} / \text{Revenue}$ (revenue minus COGS), $\text{Operating Margin} = \text{Operating Income (EBIT)} / \text{Revenue}$, and $\text{Net Margin} = \text{Net Income} / \text{Revenue}$. Each strips out progressively more costs: gross margin reflects production economics, operating margin adds overheads/SG&A and D&A, and net margin includes interest and taxes. For a firm with revenue 1000, COGS 600, EBIT 200, net income 120: margins are 40%, 20%, and 12% respectively.

Q78. What is Return on Assets (ROA) and what does it tell you?

ROA measures how efficiently a company uses its total assets to generate profit. $ROA = \text{Net Income} / \text{Total Assets}$ (often using average total assets). A 10% ROA means each INR 100 of assets produces INR 10 of net income. It is most useful for comparing asset-intensive businesses and is independent of the financing mix only when computed on a net-income basis; some analysts use $ROA = \text{EBIT} / \text{Total Assets}$ to remove financing effects.

Q79. What is Return on Equity (ROE) and why do investors watch it?

ROE measures the return generated on shareholders' invested capital. $ROE = \text{Net Income} / \text{Shareholders' Equity}$ (often average equity). An ROE of 18% means the firm earns INR 18 per INR 100 of equity. Investors watch it because, over the long run, a company compounding book value at a high ROE creates substantial value, but ROE can be inflated by leverage, so it must be read alongside the debt load.

Q80. What is ROCE and how does it differ from ROE?

Return on Capital Employed measures pre-financing operating returns on all long-term capital. $ROCE = \text{EBIT} / \text{Capital Employed}$, where $\text{Capital Employed} = \text{Total Assets} - \text{Current Liabilities}$ (equivalently equity plus long-term debt). Unlike ROE, which is an after-tax, after-interest equity return, ROCE uses EBIT and the full capital base, so it is unaffected by capital structure and is a cleaner gauge of operating efficiency. It is widely cited in India for comparing capital-intensive sectors like cement, metals, and infrastructure.

Q81. What is ROIC and why is it considered a high-quality metric?

Return on Invested Capital measures the after-tax operating return on capital actively invested in the business. $ROIC = \text{NOPAT} / \text{Invested Capital}$, where $\text{NOPAT} = \text{EBIT} * (1 - \text{tax rate})$ and $\text{Invested Capital} = \text{Total Debt} + \text{Equity} - \text{Cash}$ (or net working capital plus net fixed assets). It is considered high-quality because it isolates the productivity of operating capital from financing and excess cash, making it the right lens for judging whether management allocates capital well.

Q82. Why does comparing ROIC to WACC matter?

A firm creates economic value only when its return on invested capital exceeds its cost of that capital: value is created when $ROIC > WACC$ and destroyed when $ROIC < WACC$. The spread $ROIC - WACC$ times invested capital approximates economic profit (EVA). For example, a company earning 14% ROIC against a 10% WACC earns a 4% spread; growth then adds value, whereas if ROIC were 8%, growth would actually destroy shareholder value.

Q83. What is the debt-to-equity ratio and how is it interpreted?

D/E is a leverage/solvency ratio comparing borrowed funds to owners' funds. $\text{Debt-to-Equity} = \text{Total Debt} / \text{Shareholders' Equity}$. A D/E of 1.5 means INR 1.50 of debt per INR 1 of equity. Higher ratios mean greater financial risk and fixed interest obligations but can amplify equity returns; acceptable levels vary by industry, with utilities and banks tolerating much higher leverage than asset-light tech firms.

Q84. What does the debt-to-EBITDA ratio measure?

Debt-to-EBITDA gauges how many years of current operating cash earnings would be needed to repay total debt, and is a core metric in credit analysis and loan covenants. $\text{Debt-to-EBITDA} = \text{Total Debt} / \text{EBITDA}$. A ratio of 3.0 means debt equals three years of EBITDA; lenders often set covenants around 3.0-4.0x, and rating agencies use it heavily. It is preferred over D/E for credit because EBITDA proxies the cash available to service debt.

Q85. What is the interest coverage ratio (TIE) and what level signals risk?

Times Interest Earned measures how comfortably operating profit covers interest expense. $\text{Interest Coverage (TIE)} = \text{EBIT} / \text{Interest Expense}$. A ratio of 6 means EBIT covers interest six times over. A

coverage below roughly 1.5-2.0x is a warning sign that a downturn could leave the firm unable to service debt; some analysts use EBITDA or (EBIT + lease)/interest variants for a fuller picture.

Q86. What is the financial leverage (equity) multiplier?

The equity multiplier captures how much of a firm's assets are financed by equity versus debt. $\text{Equity Multiplier} = \text{Total Assets} / \text{Shareholders' Equity}$. A multiplier of 2.5 means assets are 2.5 times equity, implying 60% of assets are debt-financed. It is the leverage component of the DuPont decomposition, and a higher multiplier mechanically raises ROE while also raising financial risk.

Q87. What is asset turnover and what does it indicate?

Asset turnover measures how efficiently a company converts its asset base into sales. $\text{Asset Turnover} = \text{Revenue} / \text{Total Assets}$ (often average total assets). A turnover of 1.5 means each INR 1 of assets generates INR 1.50 of revenue. Low-margin, high-volume businesses (retailers, distributors) typically run high asset turnover, while capital-intensive firms run low turnover; it pairs with margin to explain ROA.

Q88. What is inventory turnover and how do you convert it to Days Inventory Outstanding (DIO)?

Inventory turnover shows how many times a firm sells and replaces its inventory in a period. $\text{Inventory Turnover} = \text{COGS} / \text{Average Inventory}$. To express it in days: $\text{DIO} = 365 / \text{Inventory Turnover} = (\text{Average Inventory} / \text{COGS}) * 365$. If COGS is 600 and average inventory is 100, turnover is 6x and DIO is about 61 days, meaning inventory sits roughly two months before being sold. Higher turnover (lower DIO) generally signals efficient inventory management.

Q89. What is receivables turnover and Days Sales Outstanding (DSO)?

Receivables turnover measures how quickly a firm collects credit sales. $\text{Receivables Turnover} = \text{Revenue} / \text{Average Accounts Receivable}$, and $\text{DSO} = 365 / \text{Receivables Turnover} = (\text{Average Receivables} / \text{Revenue}) * 365$. With revenue of 1000 and average receivables of 150, turnover is 6.7x and DSO is about 55 days, the average time to collect from customers. Rising DSO can signal weaker collections or looser credit terms and ties up working capital.

Q90. What is Days Payable Outstanding (DPO) and is a higher value good or bad?

DPO measures how long a firm takes to pay its suppliers. $\text{DPO} = (\text{Average Accounts Payable} / \text{COGS}) * 365$, or $365 / \text{Payables Turnover}$ where payables turnover is $\text{COGS} / \text{Average Payables}$. A higher DPO means the firm holds onto cash longer, effectively obtaining free supplier financing, which benefits working capital, but stretching it too far can strain supplier relationships or forfeit early-payment discounts. With payables of 80 and COGS 600, DPO is about 49 days.

Q91. What is the cash conversion cycle and how is it computed?

The cash conversion cycle (CCC) measures the number of days cash is tied up in operations between paying suppliers and collecting from customers. $\text{CCC} = \text{DIO} + \text{DSO} - \text{DPO}$. Using earlier figures (DIO 61, DSO 55, DPO 49), CCC is about 67 days. A shorter or even negative CCC is ideal: firms like large retailers and e-commerce platforms can collect from customers before paying suppliers, funding growth with supplier capital.

Q92. Walk through the 3-step DuPont decomposition of ROE.

The 3-step DuPont breaks ROE into profitability, efficiency, and leverage: $\text{ROE} = \text{Net Margin} * \text{Asset Turnover} * \text{Equity Multiplier}$, i.e. $(\text{Net Income} / \text{Revenue}) * (\text{Revenue} / \text{Assets}) * (\text{Assets} / \text{Equity})$, which telescopes back to $\text{Net Income} / \text{Equity}$. For example, 10% net margin * 1.5 asset turnover * 2.0 equity multiplier = 30% ROE. It reveals whether a firm's ROE is driven by fat margins, efficient asset use, or heavy leverage.

Q93. What does the 5-step (extended) DuPont decomposition add?

The 5-step DuPont further splits net margin to isolate the drags from interest and taxes: $ROE = (EBIT/Revenue) * (Revenue/Assets) * (Assets/Equity) * (EBT/EBIT) * (Net\ Income/EBT)$. Here the first term is operating margin, the fourth (EBT/EBIT) is the interest burden, and the fifth (Net Income/EBT) is the tax burden. It lets an analyst see whether a falling ROE stems from operations, rising interest costs from leverage, or a higher effective tax rate.

Q94. Using DuPont, explain how leverage inflates ROE.

In the DuPont identity, $ROE = ROA * Equity\ Multiplier$, so for a fixed ROA, raising the equity multiplier (more debt-funded assets) mechanically raises ROE. If ROA is 8% and the firm levers up from a multiplier of 1.5 to 3.0, ROE rises from 12% to 24% with no operational improvement. The catch: leverage also magnifies downside, raises interest expense (cutting the interest-burden term), and increases bankruptcy risk, so a high ROE built purely on debt is lower quality.

Q95. When is a high ROE actually a red flag rather than a strength?

A high ROE is a red flag when it is manufactured by excessive leverage, a shrinking equity base (e.g., large buybacks or accumulated losses), or one-off gains rather than durable operating returns. Because $ROE = ROA * Equity\ Multiplier$, two firms can show identical 20% ROE while one earns it through 18% ROA and the other through 6% ROA plus 3.3x leverage. DuPont decomposition and a check on ROIC versus WACC separate genuine quality from financial engineering.

Q96. What is common-size analysis and why is it useful?

Common-size analysis restates financial statements as percentages of a base figure so firms of different sizes and time periods can be compared. On the income statement every line is expressed as a percent of revenue; on the balance sheet every line is a percent of total assets. For example, showing COGS at 60% of revenue this year versus 55% last year instantly flags margin compression. It is the backbone of trend (vertical) analysis and peer benchmarking.

Q97. What are the main limitations of ratio analysis?

Ratios are only as good as the accounting behind them and lack meaning in isolation. Key limitations: they rely on historical book values that may not reflect market reality; differing accounting policies (e.g., Ind-AS vs IFRS vs US GAAP, FIFO vs weighted-average inventory, lease capitalization) hurt comparability; window-dressing and one-off items distort them; and "good" levels are industry-specific, so cross-sector comparison misleads. Ratios should be read as trends and against peers, not as standalone verdicts.

Q98. How do accounting standards and seasonality complicate ratio comparisons in practice, with an India angle?

Two firms can look very different purely due to standards and timing rather than economics. Under Ind-AS 116 / IFRS 16, operating leases are capitalized onto the balance sheet, inflating both assets and debt and thus changing D/E, asset turnover, and EBITDA versus older or US-GAAP-treated peers; depreciation methods and revaluations similarly shift ROA and ROCE. Seasonality matters too: a retailer's quarter-end inventory and receivables around Diwali can distort turnover and liquidity ratios, which is why analysts use average (not period-end) balances and compare same-quarter year-over-year.

A8. Earnings Quality & Forensic Accounting (Q99-Q128)

Q99. What is "quality of earnings" and why do analysts care more about it than the headline EPS?

Quality of earnings measures how repeatable, cash-backed and conservatively-stated a company's reported profit is, versus how much it depends on aggressive accruals, one-offs or accounting choices. High-quality earnings are supported by operating cash flow, recur period after period, and are not flattered by reserve

releases or capitalization tricks. Two firms can report identical EPS but the one whose profit converts cleanly into cash is worth more because the number is more likely to persist.

Q100. Why is a persistent gap between operating cash flow (OCF) and net income a warning sign?

Over time, accrual accounting profit should largely convert into cash; if net income keeps rising while OCF stagnates or falls, profit is being driven by non-cash accruals (booked revenue not yet collected, deferred costs, etc.) rather than real cash generation. A simple screen is the cash conversion ratio $OCF / \text{Net Income}$ — sustainably below ~0.8-1.0 invites scrutiny. The divergence often precedes write-downs or restatements because accruals eventually reverse.

Q101. Define the accruals ratio and explain how to read it.

The balance-sheet-free version is $\text{Accruals Ratio} = (\text{Net Income} - \text{CFO}) / \text{Total Assets}$, or using both operating and investing cash flow $(\text{NI} - \text{CFO} - \text{CFI}) / \text{Average Total Assets}$. It captures the share of earnings NOT backed by cash. A high positive value means earnings are accrual-heavy (lower quality and, per Sloan, likely to mean-revert); a low or negative value means earnings are cash-rich. Example: $\text{NI} = 100$, $\text{CFO} = 40$, $\text{Total Assets} = 1,000$, so $\text{accruals ratio} = (100-40)/1000 = 6\%$, signalling 60 of the 100 profit is non-cash.

Q102. What is the Sloan accrual anomaly?

Richard Sloan (1996) showed that firms with high accruals (low cash content of earnings) systematically underperform, while firms with low accruals outperform, because investors naively fixate on bottom-line earnings and fail to discount the lower persistence of the accrual component. A long-low-accruals / short-high-accruals strategy historically earned abnormal returns. The intuition: the accrual part of earnings reverses faster than the cash part, so high-accrual earnings are "borrowed from the future."

Q103. What is the Beneish M-score and what is it used for?

The Beneish M-score is an eight-variable statistical model that estimates the probability that a company has manipulated its earnings. A score greater than -1.78 flags a likely manipulator (the model is built so that less negative / higher is worse). It was notably the model a group of students used to flag Enron before its collapse. It is a screening tool, not proof of fraud — it raises a red flag for deeper forensic work.

Q104. Name the eight inputs to the Beneish M-score and the intuition behind the key ones.

The eight ratios are: DSRI (Days Sales in Receivables Index), GMI (Gross Margin Index), AQI (Asset Quality Index), SGI (Sales Growth Index), DEPI (Depreciation Index), SGAI (SG&A Index), LVGI (Leverage Index), and TATA (Total Accruals to Total Assets). Intuitively, a jump in DSRI suggests inflated/uncollected revenue; deteriorating GMI pressures management to manage earnings; rising AQI hints at capitalizing costs; high SGI creates pressure; falling DEPI suggests slowing depreciation to lift profit; and TATA captures the degree to which earnings are accrual-driven.

Q105. In the M-score, which two variables typically carry the most weight and why?

DSRI and TATA (and SGI) tend to be the most influential. DSRI spikes when receivables grow faster than sales, the classic signature of channel stuffing or premature/fictitious revenue; TATA directly measures how much of earnings is accruals rather than cash, the core of the accrual-manipulation story. A manipulator usually shows up as inflated receivables plus high non-cash accruals simultaneously.

Q106. What is channel stuffing and how does it show up in the financials?

Channel stuffing is shipping more product to distributors/retailers than they can sell, recognizing the revenue now to hit targets, while real end-demand is unchanged. It inflates current-period revenue and receivables and borrows from future periods, so it shows up as a sharp rise in DSO and receivables outpacing sales, often followed by elevated returns/allowances and a revenue "air pocket" next period. Auto, pharma and FMCG companies near quarter-end are classic venues.

Q107. Explain "cookie-jar reserves."

Cookie-jar reserves are over-provisions (for warranties, bad debts, restructuring, litigation) deliberately built up in good years and released into income in weak years to smooth earnings. The reserve acts as a savings jar that management dips into to manufacture steady, beating-expectations profit. Watch for reserves rising when business is strong and falling exactly when an earnings miss looms — the smoothing pattern is the tell.

Q108. What is a "big bath" and why do managers take one?

A big bath is deliberately overstating losses/write-offs in an already-bad period (often around a CEO change, recession, or restructuring) to clear the decks. By kitchen-sinking every conceivable charge now, future periods are relieved of those costs and depreciation, making the subsequent recovery look stronger and easier to beat. It is the mirror image of cookie-jar reserves — both move earnings across time.

Q109. How does capitalizing operating costs inflate earnings, and how do you detect it?

Recording a cost that should be an expense (e.g., routine maintenance, software, or in WorldCom's case line/network costs) as a capital asset moves it off the income statement into the balance sheet, boosting current profit and only bleeding through slowly as depreciation. Detect it via rising AQI (more "other assets" or capitalized costs), capex growing far faster than revenue, a widening gap between net income and CFO, and unusually long or lengthening asset lives. WorldCom's USD 3.8bn fraud was exactly this.

Q110. What is bill-and-hold revenue and when is it legitimate vs abusive?

Bill-and-hold is invoicing a customer and booking revenue while the goods remain physically in the seller's warehouse. It is legitimate only under strict conditions: the buyer (not seller) requested it, has a substantive business reason, the goods are segregated and complete, and risks/rewards have genuinely transferred. It is abusive when used to pull future sales into the current quarter without real transfer of control — a frequent revenue-recognition fraud (Sunbeam under "Chainsaw Al" Dunlap).

Q111. What is round-tripping (or "lazy Susan") revenue?

Round-tripping is two companies selling roughly equal amounts of goods/services to each other, or one selling and simultaneously buying back, to inflate reported revenue (and sometimes both sides' top lines) with no real economic substance and no net cash or profit. It was rampant among dot-com and energy-trading firms (e.g., swaps of network capacity). Red flag: revenue grows but margins, cash flow and net income do not, and there are large offsetting purchases from the same counterparties.

Q112. Why is a rising Days Sales Outstanding (DSO) a revenue red flag?

$DSO = (\text{Accounts Receivable} / \text{Revenue}) * 365$. Rising DSO means receivables are growing faster than sales — cash is being collected more slowly, which can signal channel stuffing, premature/fictitious revenue, looser credit terms to buy sales, or deteriorating customer quality. Example: revenue up 10% but receivables up 40% is a strong signal that reported sales are not converting to cash. Always compare DSO trend against peers since some seasonality and industry norms differ.

Q113. What do unbilled receivables (contract assets) tell you about earnings quality?

Unbilled receivables arise when revenue is recognized (e.g., percentage-of-completion or over-time contracts) before the customer can be invoiced. A fast-growing unbilled balance means profit is being booked on management's estimates of progress, not on billings or cash — highly subjective and easy to inflate. IT services and EPC/construction firms (in India, infra and IT) warrant extra scrutiny here; spiking unbilled revenue relative to sales is a classic aggressive-recognition flag.

Q114. How can swings in deferred revenue signal earnings management?

Deferred (unearned) revenue is a liability for cash collected but not yet earned; its change is a leading indicator of future revenue and of recognition policy. A sudden drop in deferred revenue that boosts current sales can mean the company is accelerating recognition or that bookings are slowing. Conversely,

manipulators may also use it as a reserve to smooth — releasing or building it to hit targets. Track the deferred-revenue balance and its growth alongside billings, not just recognized revenue.

Q115. Describe a cash-flow classification game and why it matters.

Companies can flatter operating cash flow (the metric investors prize) by shifting outflows out of the operating section: e.g., classifying recurring expenditures as investing (capex), securitizing/factoring receivables to pull collections into CFO, stretching payables, or capitalizing interest. Because OCF is the headline "real cash" number, dressing it up while reported earnings hold can mask deteriorating quality. Always reconcile OCF to net income, read the cash-flow notes, and check whether free cash flow (OCF minus capex) tells a worse story.

Q116. Under IFRS/Ind-AS vs US GAAP, how does interest classification create comparability issues in the cash-flow statement?

Under US GAAP, interest paid and received and dividends received are all operating; dividends paid are financing. Under IFRS/Ind-AS there is policy choice — interest paid can be operating or financing, interest/dividends received operating or investing — so two otherwise-identical firms can report different OCF. This is a classification game that can legitimately or aggressively boost reported operating cash flow, so for cross-border comparison reclassify everyone to a consistent definition before comparing.

Q117. Why is segment and geographic analysis useful for spotting earnings-quality issues?

Consolidated numbers can hide a struggling core business behind a strong segment or a one-off. Segment disclosures (Ind-AS 108 / IFRS 8) let you see which units drive margin, where growth is slowing, where capital is being parked, and whether profit is concentrated in opaque geographies (often low-tax or hard-to-audit jurisdictions). Sudden margin divergence between segments, profits booked in tax havens, or "corporate/unallocated" line items that absorb losses are red flags worth probing.

Q118. What is the abuse of non-GAAP / adjusted EBITDA, and what should you watch for?

Companies present "adjusted" metrics excluding items they deem non-recurring — but serial add-backs of stock-based compensation, "restructuring" that happens every year, "one-off" legal costs, or capitalized software amortization inflate the number that gets multiplied in valuations. The tell is a persistently widening gap between adjusted and GAAP earnings, and add-backs that are clearly recurring cash costs (especially stock comp, which is a real expense). Always rebuild EBITDA yourself and judge each add-back's legitimacy rather than trusting the company's bridge.

Q119. Distinguish pro-forma figures from reported (GAAP/Ind-AS) figures.

Reported figures follow the accounting standards and are audited; pro-forma figures are management-constructed hypotheticals — e.g., "as if" an acquisition had occurred all year, or earnings excluding chosen items. Pro-forma can aid genuine comparability (combining merged entities) but is unregulated and easily slanted to flatter. Treat pro-forma as a starting point, always reconcile it back to reported numbers, and never value a company purely off management-defined metrics.

Q120. Why are related-party transactions (RPTs) a forensic red flag, and what governs them in India?

RPTs are dealings with entities connected to promoters/management (subsidiaries, promoter firms, key-personnel relatives) where prices may not be arm's-length, enabling tunneling of cash/profits out of the listed company or fabricating revenue. They were central to Satyam and to many promoter-driven frauds. In India, the Companies Act 2013 (Sec 188) and SEBI LODR Regulation 23 mandate audit-committee approval, shareholder approval for material RPTs, and disclosure; a heavy or growing RPT load relative to revenue warrants deep scrutiny.

Q121. What are the main types of audit opinion and what does each imply?

Unqualified/clean: statements fairly presented. Qualified ("except for"): a specific, material but non-pervasive issue. Adverse: misstatements are material AND pervasive — the statements as a whole are unreliable. Disclaimer: the auditor could not obtain enough evidence to form an opinion. Anything other than a clean opinion is a serious signal; an adverse or disclaimer opinion effectively means you cannot trust the reported numbers.

Q122. What is a going-concern qualification and how should an analyst react?

A going-concern emphasis/qualification is the auditor flagging material uncertainty about the company's ability to continue operating for the next ~12 months (e.g., recurring losses, negative working capital, debt covenant breaches, inability to refinance). It does not mean imminent bankruptcy but materially raises the risk and usually triggers covenant, rating and liquidity concerns. Treat it as a trigger to stress-test liquidity, debt maturities and refinancing capacity before any investment.

Q123. What is a restatement and why does it matter for earnings quality?

A restatement is a formal correction of previously issued financial statements because they contained material errors or were not prepared per GAAP/Ind-AS. It directly destroys the reliability of historical trends, often signals weak internal controls or fraud, and is empirically associated with sharp negative stock reactions and elevated future risk. Distinguish "little r" (immaterial, revised in current filings) from "Big R" (material, requiring reissuance and an 8-K-type disclosure) — the latter is far more serious.

Q124. How does off-balance-sheet leverage distort the picture, and how do you adjust?

Obligations kept off the balance sheet — historically operating leases, SPEs/VIEs, factored receivables, take-or-pay contracts, and guarantees — understate true leverage and overstate returns on capital. To compare fairly, capitalize them: bring the debt-like obligation and a corresponding asset onto an adjusted balance sheet and recompute leverage and coverage. Enron's special-purpose entities were the textbook abuse, hiding billions of debt off-book until the structures collapsed.

Q125. How did IFRS 16 / Ind-AS 116 change lease accounting, and how do you adjust pre-2019 comparables?

IFRS 16 / Ind-AS 116 (effective ~2019) require lessees to capitalize nearly all leases — recognizing a right-of-use asset and a lease liability — so rent that was once a single operating expense now splits into depreciation plus interest, lifting reported EBITDA and gross debt. To compare a post-2019 company with older data or with US GAAP (which keeps operating leases largely off-balance-sheet in P&L geography), either capitalize old leases (a common rule of thumb is debt-equivalent $\approx 8\times$ annual rent, or PV of lease commitments) or strip the lease effect out, so the comparison is apples-to-apples.

Q126. Why must pensions and other post-employment obligations be adjusted for comparability?

A material unfunded defined-benefit pension or gratuity obligation is economically debt, and pension expense can be flattered by aggressive assumptions — a high expected return on plan assets, a high discount rate, or low wage-growth assumptions all reduce reported expense and the liability. For comparison, add the net unfunded obligation to debt, and normalize the assumed return/discount rate to a realistic level. In India, gratuity and leave-encashment provisions under Ind-AS 19 deserve the same scrutiny, especially the discount-rate choice.

Q127. What does it mean to "normalize" earnings, and how do you do it?

Normalizing strips out non-recurring, non-operating and accounting-distortion items to estimate sustainable, repeatable earnings power for valuation. Steps: remove genuine one-offs (asset-sale gains, litigation settlements, restructuring), reverse obvious accrual manipulation, adjust for cyclicity (use mid-cycle margins), put leases/pensions on a consistent basis, normalize the tax rate to a sustainable level, and

exclude unsustainable reserve releases. Example: reported PBT 120 including a 30 one-time land-sale gain and a 10 unusual litigation charge normalizes toward $120 - 30 + 10 = 100$ of recurring pre-tax profit.

Q128. Give a quick worked example of working-capital manipulation and tie together the cash-fraud lesson from Satyam.

Near year-end a firm can stretch payables, factor receivables and defer inventory purchases to flatter both reported profit and operating cash flow, then reverse it next period — so always look at multi-quarter working-capital trends, not a single snapshot. The deeper lesson from Satyam (2009), where ~INR 7,000 crore of "cash and bank balances" were simply fictitious, is that even the asset analysts trust most can be fabricated: independently verify cash via bank confirmations, reconcile interest income to the cash balance (fake cash earns no real interest), and treat unverifiable cash, related-party flows and auditor red flags as central to forensic analysis rather than back-office detail.

A3. Corporate Finance (Q129-Q169)

Q129. What is the Net Present Value (NPV) of a project and what is the decision rule?

NPV is the sum of all cash flows discounted to present value at the project's required rate of return, minus the initial outlay: $NPV = -CF_0 + \sum (CF_t / (1+r)^t)$. Accept if $NPV > 0$ (it adds value), reject if $NPV < 0$. Example: outlay 100, cash flows of 60 and 70 at $r=10\%$ give $NPV = -100 + 60/1.1 + 70/1.21 = -100 + 54.5 + 57.9 = +12.4$, so accept.

Q130. What is the Internal Rate of Return (IRR)?

IRR is the discount rate that makes NPV exactly zero: it solves $0 = -CF_0 + \sum (CF_t / (1+IRR)^t)$. The rule is to accept a project if $IRR > \text{cost of capital (hurdle rate)}$. It is a percentage measure of return, which managers find intuitive, but it implicitly assumes interim cash flows are reinvested at the IRR itself.

Q131. What is the payback period and its main weakness?

Payback period is the time taken for cumulative cash flows to recover the initial investment. If a project costs 100 and returns 40 per year, payback is 2.5 years. Its weaknesses are that it ignores the time value of money and ignores all cash flows occurring after the payback point, so it can reject high-NPV long-dated projects.

Q132. How does discounted payback differ from ordinary payback?

Discounted payback first discounts each cash flow to present value, then measures how long the cumulative discounted cash flows take to recover the outlay. It fixes the time-value flaw of simple payback but still ignores cash flows beyond the cutoff. It is always longer than ordinary payback because discounting shrinks each inflow.

Q133. What is the Profitability Index (PI) and its decision rule?

PI is the ratio of the present value of future cash flows to the initial investment: $PI = PV(\text{future cash flows}) / \text{Initial Investment} = 1 + NPV/\text{Investment}$. Accept if $PI > 1$. It is useful for ranking projects under capital rationing because it measures value created per rupee (or dollar) invested.

Q134. What is the difference between independent and mutually exclusive projects?

Independent projects can all be accepted if each meets the criterion; the decision on one does not affect the others. Mutually exclusive projects compete, so accepting one means rejecting the rest (e.g. two designs for the same factory). For independent projects NPV and IRR usually agree; for mutually exclusive projects they can conflict, and NPV should be used to rank.

Q135. Why is NPV generally preferred over IRR?

NPV measures the actual rupee/dollar value added to shareholders, assumes reinvestment at the realistic cost of capital, and works correctly even with non-conventional cash flows. IRR can mislead because it

ignores project scale, assumes reinvestment at the IRR, and can produce multiple or no solutions. When NPV and IRR conflict on mutually exclusive projects, NPV wins.

Q136. What causes the NPV-versus-IRR conflict on mutually exclusive projects?

Conflicts arise from differences in scale (size of investment) or in the timing pattern of cash flows. IRR favours projects with early cash flows or smaller scale because it is a rate, whereas NPV captures total value. A large project may have a lower IRR but a much higher NPV; since the goal is to maximise value, NPV should decide.

Q137. What is the crossover rate?

The crossover rate is the discount rate at which two projects have equal NPV — i.e. the IRR of the differential cash flows (Project A minus Project B). Below the crossover rate one project has higher NPV; above it, the ranking flips. It tells you the cost of capital at which your preference between two mutually exclusive projects changes.

Q138. When can a project have multiple IRRs?

Multiple IRRs occur with non-conventional (unconventional) cash flows — when the sign of cash flows changes more than once (e.g. outflow, inflows, then a large outflow for cleanup/decommissioning). By Descartes' rule of signs there can be as many IRRs as sign changes. In such cases IRR is unreliable and NPV or MIRR should be used instead.

Q139. What is MIRR and how does it fix IRR's reinvestment problem?

Modified IRR (MIRR) compounds all positive cash flows forward to the terminal date at the cost of capital (the reinvestment rate) and discounts all negative flows to time zero at a finance rate, then finds the single rate linking them: $MIRR = (FV \text{ of inflows} / PV \text{ of outflows})^{(1/n)} - 1$. Because reinvestment is at the realistic cost of capital and there is one unique solution, MIRR avoids both the optimistic-reinvestment and multiple-IRR problems.

Q140. A project costs 1000, returns 600 and 700 over two years; cost of capital is 12%. Compute MIRR.

FV of inflows at year 2 = $600 \times (1.12) + 700 = 672 + 700 = 1372$. PV of outflow = 1000. $MIRR = (1372/1000)^{(1/2)} - 1 = 1.1713 - 1 = 0.1713$, about 17.1%. Since MIRR (17.1%) exceeds the 12% hurdle, accept.

Q141. How do you choose between mutually exclusive projects with unequal lives?

Use the Equivalent Annual Annuity (EAA) or the replacement-chain method to put projects on a comparable annual basis. EAA converts each project's NPV into a level annual cash flow: $EAA = NPV / \text{annuity_factor}(r, n)$. Choose the project with the higher EAA, since over repeated cycles it generates more value per year.

Q142. Why should sunk costs be ignored but opportunity costs included in capital budgeting?

Sunk costs are already incurred and unrecoverable, so they do not change with the accept/reject decision and must be excluded from incremental cash flows. Opportunity costs — the value forgone by using an asset you already own (e.g. land that could be sold) — are real economic costs of the project and must be included. Capital budgeting evaluates only incremental, forward-looking cash flows.

Q143. What is CAPM and what does each term mean?

The Capital Asset Pricing Model gives the required return on equity: $Re = Rf + \text{beta} \times (Rm - Rf)$. Rf is the risk-free rate, beta measures the stock's sensitivity to market moves, and $(Rm - Rf)$ is the equity risk premium. Example: $Rf=7\%$ (Indian 10-yr G-sec), $\text{beta}=1.2$, $\text{premium}=6\%$, so $Re = 7\% + 1.2 \times 6\% = 14.2\%$.

Q144. What does beta measure and what does a beta of 1.5 mean?

Beta measures the systematic (non-diversifiable) risk of a stock relative to the overall market; it is the covariance of the stock with the market divided by the market variance. A beta of 1.5 means the stock tends

to move 1.5% for every 1% move in the market — more volatile than the index (Nifty or S&P 500). The market portfolio has beta 1.0 and the risk-free asset has beta 0.

Q145. What is the after-tax cost of debt and why use the after-tax figure?

After-tax cost of debt is $R_d \cdot (1 - T_c)$, where R_d is the yield/interest rate on debt and T_c is the marginal corporate tax rate. Interest is tax-deductible, so debt provides a tax shield that lowers its effective cost. Example: $R_d=9\%$, $T_c=25\%$, after-tax cost = $9\% \cdot (1-0.25) = 6.75\%$.

Q146. Write the WACC formula and explain why it is the discount rate for a project.

$WACC = (E/V) \cdot R_e + (D/V) \cdot R_d \cdot (1 - T_c)$, where E and D are the market values of equity and debt, $V = E + D$, R_e is cost of equity and $R_d \cdot (1 - T_c)$ is after-tax cost of debt. It is the blended required return of all capital providers, so it is the minimum return a project of average risk must earn to satisfy both shareholders and lenders. Discounting at WACC values cash flows at the firm's true opportunity cost of capital.

Q147. Compute WACC: equity 600, debt 400, $R_e=15\%$, $R_d=10\%$, tax 30%.

$V = 600 + 400 = 1000$, so $E/V = 0.6$ and $D/V = 0.4$. $WACC = 0.6 \cdot 15\% + 0.4 \cdot 10\% \cdot (1 - 0.30) = 9.0\% + 0.47\% = 9.47\%$. This 9.47% is the hurdle rate for projects of similar risk to the firm.

Q148. Why should WACC use market values and target weights rather than book values?

Market values reflect the actual current opportunity cost of capital that investors demand, whereas book values are historical accounting figures that can be badly out of date. Using market (or the firm's target) capital-structure weights gives a forward-looking cost consistent with how the firm intends to finance itself. Book weights can materially distort WACC, especially for firms whose equity trades far above book value.

Q149. What is the equity risk premium and roughly what value is used in India versus the US?

The equity risk premium (ERP) is the extra return investors require for holding the market portfolio over the risk-free asset, $(R_m - R_f)$. In the US it is commonly estimated around 4.5-5.5%; for India practitioners typically use a higher mature-market premium plus an additional country risk premium, giving an effective ERP often in the 7-9% range. The exact figure depends on whether a historical, implied, or survey approach is used.

Q150. What is country risk premium and how is it applied for valuing Indian companies?

Country risk premium (CRP) is extra return demanded for the political, economic and currency risk of investing in a given country relative to a mature market like the US. A common approach (Damodaran) is $CRP = \text{sovereign default spread} \cdot (\text{equity volatility} / \text{bond volatility})$, using India's rating-based spread. It is added into the cost of equity, e.g. $R_e = R_f + \text{beta} \cdot (\text{mature ERP}) + CRP$, raising the discount rate for Indian cash flows.

Q151. What is the difference between levered and unlevered beta?

Levered beta (equity beta) reflects both business risk and the extra risk from financial leverage, so it rises with debt. Unlevered beta (asset beta) strips out leverage to capture pure business risk. You unlever a comparable firm's beta to isolate business risk, then relever it at your own firm's capital structure to estimate your equity beta.

Q152. State the Hamada equation for levering and unlevering beta.

Hamada: $\text{beta}_{\text{levered}} = \text{beta}_{\text{unlevered}} \cdot [1 + (1 - T_c) \cdot (D/E)]$. To unlever, rearrange: $\text{beta}_{\text{unlevered}} = \text{beta}_{\text{levered}} / [1 + (1 - T_c) \cdot (D/E)]$. Example: comparable $\text{beta}_L=1.2$, its $D/E=0.5$, $T_c=25\%$ gives $\text{beta}_U = 1.2 / (1 + 0.75 \cdot 0.5) = 1.2 / 1.375 = 0.87$ — the pure business-risk beta.

Q153. Why does a firm's equity beta increase as it takes on more debt?

More debt means fixed interest obligations that must be paid before equity holders, magnifying the volatility of residual cash flows to shareholders. This financial leverage adds financial risk on top of business risk, so

the equity beta rises. The Hamada relation shows β_{levered} increasing with D/E , holding business risk (asset beta) constant.

Q154. State Modigliani-Miller Proposition I without taxes and its core intuition.

MM Proposition I (no taxes) states that in perfect markets the value of a firm is independent of its capital structure: $V_{\text{levered}} = V_{\text{unlevered}}$. The "pie" of firm value is fixed by its assets and cash flows; slicing it differently between debt and equity does not change the total. Investors can replicate or undo corporate leverage through homemade leverage, so they will not pay a premium for it.

Q155. State MM Proposition II without taxes.

MM Proposition II says the cost of equity rises linearly with leverage: $R_e = R_a + (R_a - R_d) \cdot (D/E)$, where R_a is the unlevered (asset) cost of capital. As cheaper debt is added, the increasing financial risk raises the required equity return by exactly enough to keep WACC constant. So WACC stays flat and firm value is unchanged.

Q156. How does MM Proposition I change when corporate taxes are introduced?

With taxes, interest is deductible, creating a tax shield that adds value: $V_{\text{levered}} = V_{\text{unlevered}} + T_c \cdot D$, where $T_c \cdot D$ is the present value of the perpetual debt tax shield. This implies firm value rises with leverage and, taken literally, that 100% debt is optimal. The unrealistic all-debt conclusion is what motivates the trade-off theory.

Q157. What is the trade-off theory of capital structure?

Trade-off theory says firms balance the tax-shield benefit of debt against the rising present value of financial distress and bankruptcy costs. The optimal capital structure is where the marginal tax benefit of an extra rupee of debt equals the marginal expected distress cost: $V_{\text{levered}} = V_{\text{unlevered}} + PV(\text{tax shield}) - PV(\text{distress costs})$. It explains why profitable, stable firms borrow more and why an interior optimum exists rather than 100% debt.

Q158. What is the pecking order theory and how does it differ from trade-off theory?

Pecking order theory (Myers-Majluf) says, because of asymmetric information, firms prefer internal funds first, then debt, and issue equity only as a last resort, since issuing equity signals managers think shares are overvalued. Unlike trade-off theory, it implies no well-defined target leverage ratio — observed leverage simply reflects the cumulative gap between internal cash and investment needs. It explains why highly profitable firms often carry little debt (they fund internally).

Q159. What are the costs of financial distress?

Direct costs are legal, accounting and administrative fees of bankruptcy/restructuring. Indirect costs are larger and include lost sales, supplier and customer flight, key-employee departures, fire-sale asset disposals, and agency problems such as underinvestment and asset substitution. These costs rise with leverage and cap how much debt a firm can usefully take on.

Q160. How does WACC behave as a firm increases leverage under the trade-off view?

Initially WACC falls as cheap, tax-shielded debt replaces costly equity. Beyond an optimal point, rising costs of debt and equity (from financial distress and default risk) outweigh the tax benefit, so WACC turns up, giving a U-shaped curve. The optimal capital structure is the leverage that minimises WACC, which also maximises firm value.

Q161. What is the agency cost of debt and what conflicts does it create?

Agency costs of debt arise from conflicts between shareholders and bondholders once debt is in place. Examples are asset substitution (taking on riskier projects that benefit equity at lenders' expense) and the underinvestment/debt-overhang problem (skipping positive-NPV projects because gains accrue mainly to

bondholders). Lenders anticipate this and impose covenants and higher rates, which are a real cost of leverage.

Q162. State the Dividend Irrelevance theorem and its assumptions.

Miller and Modigliani's dividend irrelevance theorem says that, in perfect markets (no taxes, no transaction costs, fixed investment policy), dividend policy does not affect firm value; only the firm's investment and earning power matter. A shareholder wanting cash can sell shares (homemade dividends), and one wanting to reinvest can buy more, so the payout split is irrelevant. Real-world frictions — taxes, signaling, transaction costs — are what make dividend policy matter.

Q163. What is dividend signaling?

Dividend signaling theory holds that, because managers have inside information, changes in dividends convey information about expected future earnings. Raising the dividend signals management's confidence in sustainable cash flows; cutting it signals trouble and usually triggers a sharp price drop. This explains why firms are reluctant to cut dividends and tend to smooth them over time.

Q164. Compare a cash dividend with a share buyback.

Both return cash to shareholders, but a dividend pays everyone proportionally and tends to be sticky, whereas a buyback reduces share count, boosts EPS, and is more flexible/one-off. Buybacks let shareholders self-select on timing and (in many regimes) offer tax advantages by converting income into capital gains. In an idealised no-tax world the two are equivalent in value; differences arise from taxes, signaling and flexibility.

Q165. What is the Gordon Growth (dividend discount) model and how does it link price to payout?

The Gordon Growth Model values a share as the present value of dividends growing at a constant rate g : $P_0 = D_1 / (R_e - g)$, where D_1 is next year's dividend and R_e the cost of equity. It links payout, growth and value because $g = \text{retention ratio} * ROE$ and $D_1 = EPS * (1 - \text{retention})$. Example: $D_1=10$, $R_e=12\%$, $g=4\%$ gives $P_0 = 10 / (0.12 - 0.04) = 125$.

Q166. Does paying a higher dividend (lower retention) always raise the share price in the Gordon model?

No. Raising the payout increases D_1 but lowers g (since $g = \text{retention} * ROE$), so the net effect depends on whether the firm's ROE exceeds its cost of equity R_e . If $ROE > R_e$, retaining and reinvesting creates value and a higher payout lowers price; if $ROE < R_e$, paying out is better. Only when $ROE = R_e$ is dividend policy value-neutral, consistent with MM irrelevance.

Q167. What is the clientele effect in dividend policy?

The clientele effect is the tendency for different investor groups to prefer different payout policies based on their tax status and income needs — e.g. tax-exempt pension funds and retirees favour high dividends, while high-tax investors prefer low-payout, growth/buyback stocks. Firms attract a clientele matched to their policy, so changing policy can disturb that base. At the margin this makes a stable, predictable policy valuable.

Q168. Explain the history of the Dividend Distribution Tax (DDT) in India and the current dividend tax regime.

Historically India levied a Dividend Distribution Tax (DDT) paid by the company on dividends declared (effectively around 20.36% grossed-up), making dividends tax-free in investors' hands up to limits. The Finance Act 2020 abolished DDT effective 1 April 2020 and shifted to the classical system: dividends are now taxed in the shareholder's hands at their applicable slab/rate, with companies deducting TDS (10% above the threshold for residents). This made dividends less tax-efficient for high-income investors and increased the relative appeal of buybacks.

Q169. How are share buybacks taxed in India and how does that compare with dividends?

After DDT's abolition, buybacks of listed shares were subject to a buyback distribution tax paid by the company (around 20% under Section 115QA), making proceeds tax-free for shareholders. From 1 October 2024 the regime changed so that buyback proceeds are taxed in the shareholder's hands as deemed dividend income at their slab rate, aligning buyback and dividend taxation. This removed the earlier tax arbitrage that had made buybacks markedly more attractive than dividends for Indian investors.

A4. Time Value of Money & Fixed Income (Q170-Q189)

Q170. What is the time value of money, and how do you compute present value and future value?

Money today is worth more than the same amount later because it can be invested to earn a return. Future value compounds a present sum forward: $FV = PV \cdot (1+r)^n$, while present value discounts a future sum back: $PV = FV / (1+r)^n$. For example, \$100 invested at $r=8\%$ for $n=3$ years grows to $FV = 100 \cdot 1.08^3 = \125.97 ; conversely, \$125.97 received in 3 years is worth \$100 today at an 8% discount rate.

Q171. What is the difference between an ordinary annuity and an annuity due, and how do you value each?

An ordinary annuity pays at the end of each period (most bonds, loans), while an annuity due pays at the beginning (rent, leases, some insurance premiums). The PV of an ordinary annuity is $PV = CF \cdot [1 - (1+r)^{-n}] / r$; an annuity due is simply that multiplied by $(1+r)$ because every cash flow arrives one period earlier. So an annuity due is always worth more than an otherwise identical ordinary annuity.

Q172. Work out the present value of a 5-year annuity paying INR 10,000 per year at 9%.

Using $PV = CF \cdot [1 - (1+r)^{-n}] / r = 10,000[1 - 1.09^{-5}] / 0.09$. Here $1.09^{-5} = 0.6499$, so the annuity factor $= (1 - 0.6499) / 0.09 = 3.8897$. $PV = 10,000 \cdot 3.8897 = \text{INR } 38,897$. If it were an annuity due, multiply by 1.09 to get INR 42,398.

Q173. What is a perpetuity and how is it valued?

A perpetuity is a stream of equal cash flows that continues forever, valued as $PV = CF / r$. For example, a console-style bond paying \$50 per year at a 5% discount rate is worth $50 / 0.05 = \$1,000$. The intuition: even though payments never stop, distant cash flows discount to near-zero, so the sum converges.

Q174. What is a growing perpetuity and where is the formula used?

A growing perpetuity has cash flows that grow at a constant rate g forever: $PV = CF_1 / (r - g)$, where CF_1 is next period's cash flow. It requires $r > g$ or the value is meaningless (infinite). This is the engine behind the Gordon Growth dividend discount model and the terminal value in a DCF; e.g., a dividend of \$2 growing at 3% with $r=8\%$ is worth $2 / (0.08 - 0.03) = \$40$.

Q175. What is the effective annual rate and why does compounding frequency matter?

The stated (nominal) rate understates the true return when interest compounds more than once a year. The effective annual rate is $EAR = (1+r/m)^m - 1$, where r is the nominal annual rate and m is compounding periods per year. A 12% nominal rate compounded monthly gives $EAR = (1+0.12/12)^{12} - 1 = 12.68\%$, so the borrower actually pays more than 12%.

Q176. How does continuous compounding work and what is its formula?

Continuous compounding is the limit as compounding frequency goes to infinity, using the exponential: $FV = PV \cdot e^{(rT)}$ and $EAR = e^r - 1$. For $r=12\%$, the continuously compounded $EAR = e^{0.12} - 1 = 12.75\%$, the maximum possible for a 12% nominal rate. Continuous compounding is standard in derivatives pricing (e.g., Black-Scholes discounting) for mathematical convenience.

Q177. Define the key terms of a bond: par/face value, coupon, and maturity.

Face (par) value is the principal repaid at maturity, typically \$1,000 in the US or INR 100/INR 1,000 for Indian G-Secs. The coupon is the periodic interest, quoted as an annual coupon rate times face value (e.g., a

6% coupon on \$1,000 pays \$60/year, often \$30 semi-annually). Maturity is the date the principal is returned and coupons stop.

Q178. What is the difference between the clean price and the dirty price of a bond?

The dirty (full) price is what the buyer actually pays and includes accrued interest since the last coupon; the clean price is the quoted price excluding accrued interest. The relationship is $\text{Dirty} = \text{Clean} + \text{Accrued Interest}$, where $\text{accrued} = \text{coupon} \times (\text{days since last coupon} / \text{days in period})$. Markets quote clean prices so the figure does not saw-tooth up and down purely due to coupon accrual.

Q179. Why do bond prices and yields move inversely?

A bond's cash flows are fixed, so its price is just the PV of those flows; when market yields rise, the discount rate goes up and the PV (price) falls, and vice versa. If you hold a 6% bond and new bonds yield 8%, yours is less attractive and must trade below par to compensate. This inverse, convex relationship is the foundation of all fixed income risk.

Q180. What is yield to maturity (YTM)?

YTM is the single discount rate that sets the PV of all a bond's future coupons and principal equal to its current market price; it is the bond's internal rate of return if held to maturity and all coupons are reinvested at that same rate. It is solved iteratively from $\text{Price} = \sum (\text{Coupon} / (1+y)^t) + \text{Face} / (1+y)^n$. A bond trading at a discount has YTM above its coupon rate; at a premium, YTM is below the coupon.

Q181. Distinguish current yield, YTM, and yield to call.

Current yield = annual coupon / current price, a simple income measure that ignores capital gain/loss and reinvestment. YTM captures total return including the pull-to-par and the timing of all cash flows. Yield to call (YTC) is the same calculation but assuming the issuer redeems the bond at the first call date and call price; for a callable premium bond, prudent investors quote the lower of YTM and YTC (yield-to-worst).

Q182. What is a zero-coupon bond and how is it priced and taxed?

A zero-coupon bond pays no coupons; it is issued at a deep discount and redeems at face value, with all return coming from price appreciation. $\text{Price} = \text{Face} / (1+y)^n$; e.g., a 10-year zero at 7% yield is worth $100 / 1.07^{10} = \text{INR } 50.83$ per 100 face. Note that even with no cash coupon, imputed interest is typically taxed annually as accrued (phantom income) in the US and India.

Q183. What is Macaulay duration?

Macaulay duration is the weighted-average time, in years, until a bond's cash flows are received, where each weight is that cash flow's PV as a fraction of price. It is expressed as $\text{MacDur} = \sum (t \times \text{PV}(\text{CF}_t)) / \text{Price}$. A zero-coupon bond's Macaulay duration equals its maturity exactly, while coupon bonds have duration less than maturity because some value is returned earlier.

Q184. What is modified duration and how does it estimate price changes?

Modified duration converts Macaulay duration into a price-sensitivity measure: $\text{ModDur} = \text{MacDur} / (1+y/m)$, where y/m is the periodic yield. The first-order price approximation is $\% \text{change in Price} = -\text{ModDur} \times \text{change in yield}$. For example, a bond with $\text{ModDur} = 7$ facing a +50bp yield move ($\text{delta}_y = 0.005$) loses roughly $-7 \times 0.005 = -3.5\%$ of its price.

Q185. What is convexity and why does it matter?

Convexity measures the curvature of the price-yield relationship that linear duration misses; because the curve is convex, duration overestimates price falls when yields rise and underestimates price gains when yields drop. The improved estimate is $\% \text{change in Price} = -\text{ModDur} \times \text{delta}_y + 0.5 \times \text{Convexity} \times \text{delta}_y^2$. Higher convexity is desirable (more upside, less downside for a given yield move), which matters most for large rate shifts and barbell portfolios.

Q186. What are DVo1 and PVo1 and how are they used?

DVo1 (dollar value of a basis point), also called PVo1, is the absolute price change in money terms for a 1bp (0.01%) move in yield: $DV01 = \text{ModDur} \times \text{Price} \times 0.0001$. A bond with $\text{ModDur} = 7$ priced at \$1,000 has $DV01 = 710000.0001 = \$0.70$ per basis point. Traders use DVo1 to size hedges so that the DVo1s of the long and short positions offset, neutralizing rate risk.

Q187. What do the shapes of the yield curve signal, and what does inversion mean?

The yield curve plots yield against maturity: normal (upward-sloping) reflects expectations of growth and a term premium; flat suggests uncertainty or transition; inverted (short rates above long rates) means investors expect rate cuts amid slowing growth. An inverted curve, especially the 10y-2y or 10y-3m spread, has historically been one of the most reliable recession predictors in the US. In India the focus is on the G-Sec curve, e.g., the 10-year benchmark yield versus shorter T-bills.

Q188. What is a credit spread and what drives it?

A credit spread is the extra yield a corporate (or other risky) bond pays over a comparable-maturity risk-free government bond: $\text{Credit Spread} = \text{Corporate Yield} - \text{Risk-free Yield}$. It compensates investors for default risk, liquidity risk, and uncertainty; e.g., a BBB bond yielding 9.5% against a 7% G-Sec has a 250bp spread. Spreads widen sharply in stress (flight to quality) and narrow in risk-on conditions, making them a key gauge of credit-market sentiment.

Q189. What are G-Secs and what is the RBI's role in the Indian bond market?

Government Securities (G-Secs) are sovereign debt issued by the Government of India, including dated bonds and short-term Treasury bills, considered risk-free in INR terms and used as the benchmark for the rupee yield curve. The RBI acts as the government's debt manager and banker, conducting primary auctions, managing the borrowing program, and using open market operations (OMO) and the repo rate to steer liquidity and yields. Unlike the US where the Treasury issues and the Fed conducts monetary policy separately, in India the RBI both manages public debt and sets policy rates, and banks must hold G-Secs to meet their SLR (Statutory Liquidity Ratio).

A5. Economics & Statistics for Finance (Q190-Q204)

Q190. What is GDP and what are its components under the expenditure approach?

GDP (Gross Domestic Product) is the total market value of all final goods and services produced within a country's borders in a given period. Under the expenditure approach: $GDP = C + I + G + (X - M)$, where C is private consumption, I is investment (capex + inventory change), G is government spending, and (X - M) is net exports. For example, if consumption is INR 140 lakh cr, investment 60, government 35, exports 25 and imports 30, then $GDP = 140 + 60 + 35 + (25 - 30) = \text{INR } 230 \text{ lakh cr}$.

Q191. What is the difference between real and nominal GDP, and what is the GDP deflator?

Nominal GDP is measured at current prices, so it rises with both output and inflation; real GDP holds prices fixed at a base year, isolating actual changes in output. The link is the GDP deflator: $\text{GDP Deflator} = (\text{Nominal GDP} / \text{Real GDP}) \times 100$. If nominal GDP grows 11% but the deflator shows 5% inflation, real growth is roughly $11\% - 5\% = 6\%$, which is the figure that reflects genuine economic expansion.

Q192. How is inflation measured via CPI and WPI, and how do they differ in India?

CPI (Consumer Price Index) tracks a basket of goods and services bought by households (heavily weighted to food and fuel), while WPI (Wholesale Price Index) tracks prices at the wholesale/producer level and excludes services. Inflation is the percent change: $\text{Inflation} = (\text{Index}_{\text{now}} - \text{Index}_{\text{year_ago}}) /$

$\text{Index}_{\text{year_ago}} * 100$. In India the RBI's official inflation target is based on CPI (Combined), not WPI, because CPI better reflects the cost of living faced by consumers.

Q193. Explain India's inflation-targeting framework and the RBI's mandate.

Since 2016 the RBI follows a flexible inflation-targeting regime set jointly with the government, aiming for CPI inflation of 4% with a tolerance band of $\pm 2\%$ (i.e., 2% to 6%). The Monetary Policy Committee (MPC), a six-member body, sets the policy repo rate to keep inflation within this band while supporting growth. If CPI breaches the band for three consecutive quarters, the RBI must formally explain the failure and remedial steps to the government.

Q194. What are the RBI's main monetary policy tools: repo rate, CRR, SLR, and OMO?

The repo rate is the rate at which the RBI lends short-term funds to banks; raising it makes credit costlier and cools demand/inflation. CRR (Cash Reserve Ratio) is the share of deposits banks must park with the RBI as cash (no interest), and SLR (Statutory Liquidity Ratio) is the share they must hold in approved liquid assets like G-secs; both control how much banks can lend. OMO (Open Market Operations) is the RBI buying or selling government securities to inject or absorb durable liquidity, e.g., buying G-secs pushes money into the system and softens rates.

Q195. How does monetary policy transmit through interest rates to bonds and equities?

When a central bank raises policy rates, bond yields rise and existing bond prices fall (inverse relationship), because $\text{Bond Price} = \sum \text{CF} / (1+y)^t$ falls as y increases. Equities typically fall too: higher rates raise the discount rate in DCF valuations and increase borrowing costs, compressing the present value of future cash flows. Rate-sensitive sectors (banks, real estate, autos) and high-duration growth/tech stocks react most sharply to rate changes.

Q196. What is fiscal policy, and how does the fiscal deficit differ from a budget surplus?

Fiscal policy is the government's use of taxation and spending to influence aggregate demand; expansionary policy (more spending or tax cuts) stimulates growth, while contractionary policy restrains an overheating economy. The fiscal deficit is the shortfall when total expenditure exceeds total revenue (excluding borrowings) and is usually expressed as a percentage of GDP. India targets a glide path toward a fiscal deficit around 4.5% of GDP; large deficits are financed by borrowing, which can crowd out private investment and pressure bond yields.

Q197. What moves the USD/INR exchange rate, and what do appreciation and depreciation mean?

USD/INR is the number of rupees per US dollar; if it rises from 83 to 86, the rupee has depreciated (weakened) and the dollar appreciated. Key drivers include interest-rate differentials, inflation gaps, the current account/trade deficit, FII/FPI capital flows, crude oil prices (India is a big importer), and RBI intervention via forex reserves. A weaker rupee helps exporters and IT/pharma firms but raises the cost of imports and dollar-denominated debt, feeding into imported inflation.

Q198. What is the business cycle, and what are its phases?

The business cycle is the recurring pattern of fluctuations in economic activity around a long-run growth trend, with four phases: expansion, peak, contraction (recession), and trough. A common rule-of-thumb defines a recession as two consecutive quarters of negative real GDP growth. Policymakers use countercyclical tools, cutting rates and boosting spending in downturns and tightening during overheating booms, to smooth the cycle.

Q199. Explain demand-supply equilibrium and price elasticity of demand.

Market equilibrium occurs where the quantity demanded equals quantity supplied, setting the clearing price; a shift in either curve moves price and quantity. Price elasticity of demand measures responsiveness: $\text{PED} = (\% \text{ change in quantity demanded}) / (\% \text{ change in price})$. If a 10% price rise cuts quantity

demanding by 20%, PED = -2 (elastic); essentials like staples or fuel tend to be inelastic ($|PED| < 1$), so revenue rises when their price increases.

Q200. Define mean, variance, and standard deviation, and why standard deviation matters in finance.

The mean is the average, $\text{mean} = (\text{sum of } x_i) / n$; variance is the average squared deviation from the mean, $\text{var} = \text{sum}((x_i - \text{mean})^2) / n$; and standard deviation is $\text{sigma} = \sqrt{\text{var}}$. Standard deviation is the workhorse measure of risk/volatility in finance because it is in the same units as returns. For returns of 2%, -1%, 4%, 3%, the mean is 2%, and sigma quantifies how widely actual returns scatter around that 2%.

Q201. What is the normal distribution and how do z-scores work?

The normal (Gaussian) distribution is the symmetric bell curve fully described by its mean and standard deviation, where about 68% of outcomes fall within ± 1 sigma, 95% within ± 2 sigma, and 99.7% within ± 3 sigma. A z-score standardizes a value: $z = (x - \text{mean}) / \text{sigma}$, telling you how many standard deviations x lies from the mean. If a stock's monthly return is 8% with mean 2% and sigma 3%, then $z = (8 - 2) / 3 = 2$, an unusually high outcome (top ~2.5% of the distribution).

Q202. Explain covariance and correlation, including the correlation formula.

Covariance measures whether two variables move together: $\text{cov}(x,y) = \text{sum}((x_i - \text{mean}_x)(y_i - \text{mean}_y)) / n$; its sign indicates direction but its magnitude is unbounded and scale-dependent. Correlation normalizes it to a -1 to +1 range: $\text{corr} = \text{cov}(x,y) / (\text{sigma}_x * \text{sigma}_y)$. A correlation of +1 means perfect co-movement, 0 means no linear relationship, and -1 means they move exactly opposite; low or negative correlations between assets are what make diversification reduce portfolio risk.

Q203. How are regression slope, beta, and R-squared defined and interpreted?

In a simple linear regression $y = \alpha + \beta x + \text{error}$, the slope is $\beta = \text{cov}(x,y) / \text{var}(x)$. In finance, beta from regressing a stock's returns on the market's returns measures systematic risk: $\beta = 1.2$ means the stock tends to move 1.2% for every 1% market move. R-squared is the fraction of the dependent variable's variance explained by the model ($R^2 = 1 - \text{SS}_{\text{residual}} / \text{SS}_{\text{total}}$); for a stock-vs-market regression, R-squared shows how much of the stock's movement is driven by the market versus stock-specific factors.

Q204. What is hypothesis testing with a p-value, and why do skewness and kurtosis matter for fat tails in finance?

In hypothesis testing you assume a null hypothesis (e.g., a strategy's true alpha is zero), compute a test statistic, and the p-value is the probability of seeing a result this extreme if the null were true; $p < 0.05$ conventionally lets you reject the null as statistically significant. Skewness measures asymmetry of returns (negative skew = frequent small gains but occasional large losses), and kurtosis measures tail-heaviness, with the normal distribution having kurtosis of 3 (excess kurtosis 0). Financial returns typically show negative skew and high (leptokurtic) kurtosis, meaning extreme crashes occur far more often than a normal distribution predicts, which is exactly why VaR models that assume normality understate true tail risk.

A6. Indian Markets, Structure & Regulation (Q205-Q224)

Q205. What is SEBI and what are its core functions?

SEBI (Securities and Exchange Board of India), established as a statutory body under the SEBI Act 1992, is the regulator of the Indian securities market. Its mandate is threefold: protect the interests of investors, promote development of the market, and regulate it. In practice it registers and supervises intermediaries (brokers, mutual funds, RAs, merchant bankers), frames regulations (disclosure norms, insider-trading

rules, takeover code), and has quasi-legislative, quasi-executive and quasi-judicial powers including investigation, search-and-seizure, penalties and disgorgement.

Q206. What is the difference between the primary market and the secondary market?

The primary market is where securities are issued for the first time and capital flows directly from investors to the issuer — e.g., an IPO, FPO, rights issue, or fresh bond issue. The secondary market is where already-issued securities trade between investors (on exchanges like NSE/BSE), providing liquidity and price discovery but no new capital to the issuer. A company raises money only in the primary market; subsequent share-price moves on the exchange do not affect its cash balance.

Q207. What is the difference between NSE and BSE?

Both are SEBI-regulated stock exchanges. The BSE (Bombay Stock Exchange), founded 1875, is Asia's oldest exchange and lists the most companies; its benchmark is the Sensex (30 stocks). The NSE (National Stock Exchange), launched in 1994, pioneered fully electronic screen-based trading and dominates cash-equity and especially derivatives volumes in India; its benchmark is the Nifty 50. For most liquid stocks the NSE has higher turnover, so traders default to it, but listed shares are fungible across both.

Q208. What are the Nifty 50 and the Sensex, and how do they differ?

Both are large-cap benchmark indices. The Sensex (BSE) tracks 30 large, financially-sound companies; the Nifty 50 (NSE) tracks 50 large-caps across sectors. Both are free-float market-capitalization weighted, so a stock's weight reflects only the shares freely available for public trading, not promoter/locked holdings. They move closely together since they share most heavyweight constituents (e.g., Reliance, HDFC Bank, ICICI Bank).

Q209. What is free-float market-cap weighting and why is it used?

Free-float market cap = price * (shares outstanding * free-float factor), where the free-float factor excludes promoter holdings, strategic/government stakes, and locked-in shares. An index weights each stock by its free-float market cap divided by the index's total, so larger and more freely-traded companies move the index more. It is preferred over full-market-cap because it reflects only investable shares, reduces distortion from closely-held companies, and makes the index easier to replicate by index funds.

Q210. What are Bank Nifty and Nifty Next 50?

Bank Nifty (Nifty Bank) is a sectoral index of the most liquid and large-cap banking stocks on the NSE (around 12 banks), heavily used for trading and derivatives due to its volatility. Nifty Next 50 comprises the 50 stocks ranked just below the Nifty 50 by free-float market cap — effectively the "on-deck" large-caps from which future Nifty 50 entrants are often drawn. Both are free-float market-cap weighted like the headline indices.

Q211. Walk through the book-building IPO process in India.

In book building, the issuer and merchant bankers (lead managers) set a price band, and investors bid for quantities at prices within that band during the offer period. The cut-off (final issue) price is determined from the demand book, and shares are allotted at that price. Key documents are the DRHP (Draft Red Herring Prospectus, filed with SEBI for review) and the RHP (Red Herring Prospectus, filed with the RoC before the issue opens). Retail and other categories have reserved quotas, and listing follows on a T+3 basis from issue close under current SEBI norms.

Q212. What is a Red Herring Prospectus?

A Red Herring Prospectus (RHP) is the offer document filed by an issuer ahead of a book-built IPO that contains most disclosures — business, risk factors, financials, use of proceeds — but omits the final price and exact number of shares (hence "red herring," because price is not yet fixed). It is filed with the Registrar of

Companies before the issue opens; the final prospectus with price and size is filed after the book closes. The earlier draft version reviewed by SEBI is the DRHP.

Q213. What is ASBA and how does it work in IPO applications?

ASBA (Application Supported by Blocked Amount) is the mandatory mechanism for IPO applications in India whereby the application money is merely blocked (lien-marked) in the investor's own bank account rather than debited upfront. The funds keep earning interest and are debited only to the extent of actual allotment; the blocked amount on unallotted shares is released. This protects investors and improves cash efficiency versus the old system of paying the full amount in advance. UPI-based ASBA is used by retail investors.

Q214. What are NSDL and CDSL, and what is dematerialization?

NSDL (National Securities Depository Limited) and CDSL (Central Depository Services Limited) are India's two depositories — institutions that hold securities in electronic (book-entry) form, analogous to a "bank for shares." Dematerialization is the process of converting physical paper share certificates into electronic holdings in a demat account, which is accessed via a Depository Participant (DP) such as a broker or bank. SEBI mandates demat form for trading listed securities, which eliminated forgery, bad deliveries and physical transfer delays.

Q215. What is the NISM Series-XV exam and what do the SEBI RA Regulations require?

The SEBI (Research Analysts) Regulations, 2014 require anyone issuing research reports or making buy/sell recommendations on securities to register with SEBI as a Research Analyst and meet certification and qualification norms. The relevant certification is the NISM Series-XV: Research Analyst certification (NISM = National Institute of Securities Markets, SEBI's educational arm), which must be obtained and kept valid. The rules also mandate disclosure of conflicts of interest, separation of research from other business lines, and limits on the analyst trading against their own recommendations.

Q216. What is a clearing corporation, what does NSCCL do, and what is T+1 settlement?

A clearing corporation novates and guarantees settlement, acting as the central counterparty (CCP) — buyer to every seller and seller to every buyer — so neither side bears the other's default risk; NSCCL (National Securities Clearing Corporation, now NSE Clearing) does this for the NSE, netting obligations, running the settlement guarantee fund, collecting margins and managing pay-in/pay-out. Settlement happens on a T+1 cycle, meaning a trade on day T settles the next business day. India fully adopted T+1 in early 2022 — ahead of the US, which moved only in 2024 — shortening the trade-to-settlement gap, freeing capital faster and cutting counterparty risk; SEBI has also piloted optional same-day (T+0) settlement.

Q217. How do index options differ from stock options in India regarding settlement and style?

Indian index options (Nifty, Bank Nifty) are European-style — exercisable only at expiry — and are cash-settled against the index level, since you cannot deliver an index. Single-stock options are also European-style but are physically settled: open positions at expiry result in delivery/receipt of the underlying shares rather than cash. This physical-settlement rule for stock F&O, introduced by SEBI, was designed to curb excessive speculation and align derivative prices with the cash market.

Q218. What are SPAN and exposure margins in F&O?

SPAN (Standard Portfolio Analysis of Risk) margin is the core initial margin computed by simulating the portfolio's worst-case loss across a grid of price and volatility scenarios; it covers the largest likely one-day move. The exposure (or extreme-loss) margin is an additional buffer levied on top of SPAN to cover tail events beyond SPAN's scenarios. Together SPAN + exposure form the upfront margin a trader must keep; mark-to-market losses are settled daily on top of this.

Q219. Explain lot sizes and weekly vs monthly expiry in Indian derivatives.

Indian F&O trades in standardized "lots" (a fixed number of underlying units per contract), with lot sizes set by the exchange so that contract value stays around a SEBI-prescribed minimum — you cannot trade a single share in F&O. Monthly contracts expire on a designated day of the last week of each month, while index options also have weekly expiries that settle every week, giving short-dated trading opportunities. SEBI has been rationalizing the number of weekly expiries (limiting each exchange to one weekly-expiry index) to temper retail speculation.

Q220. What is Securities Transaction Tax (STT)?

STT is a tax levied on the value of transactions in listed securities executed on Indian exchanges, collected at source by the broker. Rates differ by segment — e.g., it applies to both buy and sell legs of delivery equity, only to the sell side of intraday and futures, and on the sell (premium/settlement) side of options — and is generally low (fractions of a percent). Because equity trades are subject to STT, they qualify for the concessional equity capital-gains tax regime under Sections 111A/112A.

Q221. How are LTCG and STCG taxed on listed equity in India?

For listed equity shares and equity mutual funds (which are subject to STT), the holding-period threshold is 12 months: gains on holdings sold within 12 months are Short-Term Capital Gains (STCG) taxed at a flat concessional rate, while gains held beyond 12 months are Long-Term Capital Gains (LTCG) taxed at a lower rate with an annual exemption threshold on aggregate LTCG. (Following the 2024 Budget, equity STCG is taxed at 20 percent and LTCG at 12.5 percent above the exemption, but the candidate should state the rule conceptually since rates change.) This contrasts with debt and other assets, which carry longer holding thresholds and different rates.

Q222. How is income from F&O trading taxed in India?

Profits and losses from futures and options are treated as non-speculative business income (not capital gains), reported under "Profits and Gains from Business or Profession" and taxed at the individual's slab rate. Because it is business income, the trader can deduct related expenses (brokerage, internet, data subscriptions), and losses can be set off and carried forward under business-loss rules. Significant turnover may trigger tax-audit requirements under the Income Tax Act.

Q223. What are circuit breakers and price bands on Indian exchanges?

Price bands are stock-level limits (e.g., 5/10/20 percent) beyond which an individual security cannot move in a day, preventing runaway single-stock moves. Market-wide circuit breakers halt trading across the entire market when a benchmark index (Nifty/Sensex) moves by 10, 15 or 20 percent, with halt durations escalating by the trigger level and time of day — coordinated across exchanges. Their purpose is to give the market a cooling-off period, curb panic-driven cascades, and let information disseminate during extreme volatility.

Q224. How do the jurisdictions of RBI and SEBI differ, and what are FPI/FII flows?

SEBI regulates the securities market — exchanges, listed companies, brokers, mutual funds, and capital-market intermediaries — while the RBI (Reserve Bank of India) is the central bank governing monetary policy, banks, the rupee, government securities, and foreign-exchange management under FEMA; there is overlap on areas like corporate bonds and the G-sec market. Mutual funds are SEBI-regulated AMC's (Asset Management Companies) whose industry body is AMFI (Association of Mutual Funds in India), which sets self-regulatory standards. FPI/FII (Foreign Portfolio/Institutional Investor) flows are overseas capital invested in Indian listed securities — registered with SEBI and routed under RBI/FEMA limits — and large net inflows or outflows are a major driver of Nifty/Sensex movements and the rupee.

A9. Indian Financial System & Products (Advanced) (Q225-Q249)

Q225. Map the four main financial regulators in India and what each one supervises.

RBI regulates banks, NBFCs, the money market, government securities, payment systems and monetary policy. SEBI regulates the securities markets — stock exchanges, listed companies, mutual funds, AIFs, brokers, RIAs and research analysts. IRDAI regulates the insurance sector (life, general, reinsurance), and PFRDA regulates the National Pension System (NPS) and pension funds. A useful boundary case: ULIPs (insurance-linked investment) sit with IRDAI, mutual funds with SEBI, and corporate/government bonds straddle RBI (issuance, primary dealers for G-Secs) and SEBI (listed corporate bonds, secondary trading on exchanges).

Q226. What are the main money market instruments in India and their typical tenors?

The money market handles short-term debt up to one year. Key instruments: Treasury Bills (T-bills, issued by RBI for the government at 91, 182 and 364 days, sold at a discount, zero coupon); Commercial Paper (CP, unsecured promissory notes by corporates/PDs, 7 days to 1 year); Certificates of Deposit (CDs, issued by banks/FIs, 7 days to 1 year); and call/notice/term money (interbank borrowing — call money is overnight, notice is 2-14 days). All are designed for liquidity management rather than long-term financing.

Q227. How is a 91-day T-bill priced and what yield does it imply if it is issued at Rs 98.50 per Rs 100 face value?

T-bills are zero-coupon and sold at a discount to face value; the return is the difference between purchase price and the Rs 100 redemption. The implied yield (simple, money-market convention) is $\text{Yield} = ((\text{Face} - \text{Price})/\text{Price}) * (365/\text{Days})$. For Rs 98.50 over 91 days: $((100 - 98.50)/98.50) * (365/91) = 0.01523 * 4.011 = 6.11\%$ annualised. There is no periodic coupon — all return comes from the discount.

Q228. Explain repo, reverse repo, the LAF corridor and the MSF.

Repo is RBI lending to banks against G-Sec collateral at the repo rate (the policy rate); reverse repo / the Standing Deposit Facility (SDF) is banks parking surplus funds with RBI at a lower rate. The Liquidity Adjustment Facility (LAF) corridor is the band: the repo rate sits in the middle, the SDF rate forms the floor, and the Marginal Standing Facility (MSF) rate forms the ceiling. MSF lets banks borrow overnight against SLR securities (even dipping into SLR) at a penal rate above repo — an emergency liquidity backstop. The corridor width (e.g. SDF and MSF each typically 25 bps either side of repo) keeps overnight rates anchored near the policy rate.

Q229. What is the G-Sec market and what role do Primary Dealers play?

Government Securities (G-Secs) are sovereign debt — dated securities (long tenor, coupon-bearing) and T-bills — issued by RBI on the government's behalf through auctions on the E-Kuber/NDS-OM platform. Primary Dealers (PDs) are RBI-registered entities obligated to bid in every primary auction (underwriting/market-making), ensuring the auction is fully subscribed, and to make two-way quotes in the secondary market to provide liquidity. They are the backbone of price discovery in the sovereign curve; retail access has been opened via the RBI Retail Direct scheme.

Q230. Why is India's corporate bond market considered underdeveloped relative to its equity market?

Despite a deep sovereign market, the corporate bond market is shallow because issuance is dominated by private placements to a few institutions (concentrated in AAA/AA names), secondary-market liquidity is thin, and a large share is held to maturity by insurers/pension funds and banks. Demand-side constraints (limited retail participation, FPI investment limits in debt) and supply concentration in top-rated issuers crowd out lower-rated paper, so smaller/riskier firms rely on banks instead. SEBI reforms — electronic

bidding platforms, a Request-for-Quote (RFQ) platform, and a backstop/limited-purpose Corporate Debt Market Development Fund — aim to deepen liquidity.

Q231. What are AIFs and how do SEBI's Category I, II and III differ?

Alternative Investment Funds (AIFs) are privately pooled vehicles registered under SEBI's AIF Regulations, 2012, with a minimum investor commitment of Rs 1 crore and capped at 1,000 investors per scheme (250 for angel funds). Category I invests in socially/economically desirable areas — VC funds, SME funds, infrastructure, angel funds (incentivised, no leverage). Category II is the residual catch-all — private equity, debt funds, fund-of-funds — that uses no leverage except for day-to-day needs. Category III employs complex/diverse strategies including leverage and derivatives — hedge funds and long-short equity funds — and is the only category that can trade actively for short-term returns.

Q232. How are Category III AIFs taxed differently from Category I and II?

Category I and II AIFs enjoy "pass-through" status under Section 115UB — income (other than business income) is taxed in the investors' hands as if they invested directly, so the fund itself is not taxed on it. Category III AIFs do NOT have full pass-through status; the fund is generally taxed at the fund level (often at maximum marginal rate on certain income), which is a key structural consideration when choosing between categories. This is why long-short/hedge strategies in Cat III carry a heavier tax-administration burden.

Q233. What is a REIT, why was it created, and how is it structured in India?

A Real Estate Investment Trust (REIT) is a SEBI-regulated trust that pools investor money to own income-generating commercial real estate (offices, malls), letting retail investors access real estate without buying property directly and giving developers an exit/monetisation route. Structure: Sponsor, Trustee, Manager and Special Purpose Vehicles (SPVs) holding the assets. SEBI mandates at least 80% of assets in completed, rent-generating property and requires distribution of at least 90% of net distributable cash flows to unitholders, paid at least semi-annually — making it a yield instrument akin to a property-backed dividend stock.

Q234. What is an InvIT and how does it differ from a REIT?

An Infrastructure Investment Trust (InvIT) is the infrastructure analogue of a REIT — it pools capital to own operating infrastructure assets (roads, power transmission lines, gas pipelines, telecom towers) and is used by developers/governments to recycle capital out of completed projects. The structure (Sponsor, Trustee, Investment Manager, Project SPVs) and the 90% distribution rule mirror REITs; the difference is the underlying asset class (infrastructure vs commercial real estate) and InvITs can be publicly listed or privately placed. Both distribute via a mix of interest, dividend and capital-return components, each taxed differently in the unitholder's hands.

Q235. How is income from REIT/InvIT units taxed for an Indian investor?

Distributions come in components, each taxed differently: the interest component is taxable at the investor's slab rate; the dividend component is generally tax-free to the investor if the SPV did not opt for the concessional corporate tax regime (otherwise taxable at slab); and the return-of-capital/amortisation component reduces the cost of acquisition (deferring tax until sale). Capital gains on the units follow listed-security rules — short-term if held under the threshold, long-term beyond it. This component split is why a single REIT payout can have multiple tax treatments in one financial year.

Q236. Explain SEBI's mutual fund scheme categorization and why it was introduced.

In 2017-18 SEBI introduced scheme categorization to stop AMC's running dozens of overlapping schemes; it defined five broad groups (Equity, Debt, Hybrid, Solution-oriented, Other/Index-FoF) with standardised sub-categories and only one scheme per category per AMC (with exceptions like index funds/ETFs). Equity sub-categories are defined by market-cap rules: large-cap = top 100 stocks by full market cap, mid-cap =

101st-250th, small-cap = 251st onwards. The goal was comparability and true-to-label investing, so a "large-cap fund" must actually hold a minimum quota (e.g. 80%) in large-caps.

Q237. What is the Total Expense Ratio (TER), and how does it affect returns?

TER is the annual percentage of a fund's assets charged to investors for management, administration, distribution and other costs; it is deducted daily and already reflected in the NAV, so investors never pay it separately. SEBI caps TER on a sliding scale that falls as AUM rises (e.g. higher slabs for the first few hundred crores, lower for larger funds), and equity funds are allowed higher TER than debt/index funds. On a 12% gross return, a 1.5% TER nets the investor ~10.5%; compounded over decades the drag is large, which is the core argument for low-cost index funds.

Q238. What is the difference between a direct plan and a regular plan?

Both invest in the identical portfolio of the same scheme, but a regular plan pays a trail commission to the distributor/agent, which is embedded in a higher TER, while a direct plan (bought straight from the AMC/RTA/SEBI-registered platform) carries no distribution commission and thus a lower TER and higher NAV. The difference (often 0.5-1.0% per year) compounds significantly over time — a direct plan can leave an investor materially richer over 15-20 years purely from the cost gap, with no difference in the underlying holdings.

Q239. Define NAV and explain how SIP rupee-cost averaging works.

Net Asset Value (NAV) is the per-unit value of a fund:
$$\text{NAV} = (\text{Total assets} - \text{liabilities}) / \text{Units outstanding}$$
 struck once daily (for most funds) on a forward-pricing basis. A Systematic Investment Plan (SIP) invests a fixed rupee amount at regular intervals, so it automatically buys more units when NAV is low and fewer when NAV is high — "rupee-cost averaging" — lowering the average cost per unit over a volatile period versus a single lump-sum timed badly. It does not guarantee profit or protect against loss in a falling market, but it removes timing risk and instils discipline.

Q240. How do index funds and ETFs differ, especially in the Indian context?

Both passively track an index (e.g. Nifty 50, Sensex), but an index fund is a mutual fund bought/sold at end-of-day NAV directly from the AMC, while an ETF trades intraday on the exchange like a stock at a live price that can deviate from its iNAV. ETFs need a demat account and a broker, and rely on Authorised Participants/market makers and a creation-redemption mechanism to keep price near NAV; liquidity and tracking error (and bid-ask spread) are the practical concerns. In India, EPFO inflows and government disinvestment (e.g. CPSE/Bharat 22 ETFs) have boosted ETF AUM, though many retail investors prefer index funds via SIP for simplicity.

Q241. What happens to share price on a bonus issue versus a stock split?

In a bonus issue the company capitalises reserves to issue free shares (e.g. 1:1), increasing the share count while face value stays the same; in a stock split the face value itself is divided (e.g. Rs 10 to Rs 2, a 5-for-1 split), also increasing share count. In both the share price drops proportionally on the ex-date so total market cap and an investor's wealth are unchanged — a 1:1 bonus roughly halves the price, doubling units. Neither creates value; they improve liquidity/affordability and reset the per-share optics.

Q242. How does a rights issue work and what is the theoretical ex-rights price (TERP)?

A rights issue lets existing shareholders buy new shares in proportion to holdings at a discount to market price, raising fresh capital while protecting them from dilution if they subscribe. The post-issue price settles at the theoretical ex-rights price:
$$\text{TERP} = (\text{Old shares} \times \text{Cum-price} + \text{New shares} \times \text{Issue price}) / (\text{Old} + \text{New shares})$$
 Example: holding at Rs 100 cum-rights, a 1:1 rights at Rs 60 gives $\text{TERP} = (1 \times 100 + 1 \times 60) / 2 = \text{Rs } 80$; the right itself is renounceable and trades on the exchange during the offer.

Q243. Explain a buyback and how it differs from a dividend, including the Indian tax angle.

A buyback is the company repurchasing its own shares (via tender offer or open market), reducing share count and returning cash; it boosts EPS and signals that management sees shares as undervalued, whereas a dividend distributes cash without changing share count. Tax treatment in India has shifted: previously the company paid buyback tax and proceeds were tax-free to shareholders, but post the October 2024 change, buyback proceeds are taxed in the shareholder's hands as deemed dividend at slab rates — narrowing the historic tax arbitrage that made buybacks attractive over dividends for promoters.

Q244. What are ASBA and UPI in the IPO application process?

ASBA (Application Supported by Blocked Amount) means the application money stays in the investor's bank account and is only blocked (not debited) until allotment — it is debited only for shares actually allotted, and the block is released otherwise, so the investor keeps earning interest and avoids refund delays. ASBA is mandatory for IPO applications. Retail investors up to a value limit (Rs 5 lakh) can apply via the UPI mechanism, where the bank block is triggered by approving a UPI mandate in their payment app, integrating the broker, exchange, sponsor bank and RTA into one fast workflow.

Q245. What are NSDL and CDSL, and what does "demat" really involve?

NSDL and CDSL are India's two depositories — institutions that hold securities in electronic (dematerialised) form, eliminating physical certificates. Investors don't deal with them directly; they open a demat account with a Depository Participant (DP), typically a broker or bank, which is the agent of the depository. Demat converts paper shares into book entries so that ownership transfer on trade settlement is a simple electronic debit/credit; corporate actions (dividends, bonus, splits), pledges and ISIN-level identification all flow through the depository, making settlement fast, safe and irreversible.

Q246. Explain India's T+1 settlement cycle and the role of clearing corporations.

India moved to a T+1 rolling settlement (full rollout completed January 2023), meaning a trade executed today settles — funds and securities exchanged — the next business day, faster than the global T+2 norm; SEBI has since piloted optional T+0 same-day settlement. The clearing corporations — NSE Clearing (NSCCL/NCL) for NSE and Indian Clearing Corporation (ICCL) for BSE — step in as the central counterparty (CCP), novating each trade so they become buyer to every seller and seller to every buyer, guaranteeing settlement even if one side defaults. This novation removes counterparty risk between the two original parties.

Q247. What is the Settlement Guarantee Fund and why does it matter?

The Settlement Guarantee Fund (SGF, also called the Core SGF) is a pool of capital maintained by the clearing corporation, funded by contributions from the CC, the exchange and clearing members, that is used to complete settlement if a member defaults. It is the financial backbone of the CCP guarantee — when a clearing member fails to pay or deliver, the CC draws on the defaulter's margins and contributions first, then the SGF, in a defined "default waterfall." This is what makes the settlement guarantee credible rather than just a promise, underpinning systemic trust in exchange-traded markets.

Q248. Describe India's margin framework: SPAN + exposure, upfront and peak margin.

For F&O, total initial margin is SPAN margin (a portfolio risk-based margin computed by simulating worst-case price/volatility moves across a scenario grid) plus an Exposure margin (an additional ad-hoc buffer over SPAN). SEBI mandates upfront margin collection — brokers must collect the full initial margin before/at the time of the trade, not at end of day. The peak margin rule requires margins to be met based on the highest of multiple intraday snapshots (clearing corporations take random snapshots through the day), curbing the old practice of brokers offering excessive intraday leverage; shortfalls attract penalties.

Q249. Summarise SEBI's PIT (Prohibition of Insider Trading) Regulations and the NISM Research Analyst obligations.

Under SEBI (PIT) Regulations 2015, an "insider" (connected person or anyone with Unpublished Price Sensitive Information, UPSI) is barred from trading in securities while in possession of UPSI or from communicating it (tipping), unless for legitimate purposes; companies must maintain structured digital databases, codes of conduct, and trading windows/closure around results. A SEBI-registered Research Analyst (passing NISM Series XV) must disclose conflicts of interest, any financial interest or holding in the subject company, and whether they/relatives received compensation; they cannot trade contrary to their own recommendation within prescribed windows, must keep records, separate research from other business, and present price targets with assumptions and risk factors. FPIs (Foreign Portfolio Investors) bring portfolio/secondary-market money into listed securities subject to SEBI registration and investment limits, whereas FDI (Foreign Direct Investment) is a lasting controlling stake (typically 10%+) in a business with sector caps and RBI/government routes — and for the same investor, debt is taxed largely at slab/applicable rates with no STCG concession while listed equity enjoys concessional STCG (15%, with grandfathered LTCG taxed at 10%/12.5% above a threshold), driving very different after-tax outcomes.

B. Equity Research

B1. DCF & Intrinsic Valuation (Q250-Q278)

Q250. What is a discounted cash flow (DCF) valuation and why is it used?

A DCF values a business as the present value of the cash flows it is expected to generate, discounted back at a rate that reflects their risk and the time value of money: $\text{Value} = \sum \text{of } CF_t / (1+r)^t$. It is used because it is the most direct way to estimate intrinsic value from fundamentals (cash, growth, risk) rather than relying on what the market is paying for comparable companies. Its strength is that it forces explicit assumptions; its weakness is that it is highly sensitive to those assumptions, especially the discount rate and terminal value.

Q251. What is the core difference between FCFF and FCFE?

Free cash flow to the firm (FCFF) is the cash available to all capital providers (both debt and equity holders) before financing flows; free cash flow to equity (FCFE) is the cash available only to equity holders after interest and net debt repayments. FCFF is "unlevered" (independent of capital structure), while FCFE is "levered" because it is affected by interest and borrowing. Roughly, $\text{FCFE} = \text{FCFF} - \text{Interest} \cdot (1-t) + \text{Net Borrowing}$.

Q252. Write out the formula for FCFF starting from EBIT.

$\text{FCFF} = \text{EBIT} \cdot (1-t) + \text{D\&A} - \text{CapEx} - \text{change in NWC}$. You start with EBIT (operating profit), tax it at the marginal rate to get unlevered after-tax operating profit (NOPAT), add back non-cash D&A, subtract reinvestment in fixed assets (CapEx) and subtract the increase in net working capital. For example, EBIT 200, t 25%, D&A 40, CapEx 60, change in NWC 20 gives $\text{FCFF} = 150 + 40 - 60 - 20 = 110$.

Q253. Why do we add back D&A and subtract CapEx and change in NWC in the FCFF bridge?

D&A is a non-cash expense already deducted to get EBIT, so it is added back to convert accounting profit to cash. CapEx is a real cash outflow for long-lived assets that the income statement only recognises gradually, so it is subtracted in full. An increase in net working capital (more receivables/inventory than payables) ties up cash, so a rise in NWC reduces free cash flow while a fall releases cash.

Q254. What is the difference between an unlevered DCF and a levered DCF?

An unlevered DCF discounts FCFF (pre-financing cash flows) at the WACC to arrive at enterprise value, and is the standard approach because it isolates operating performance from capital structure. A levered DCF discounts FCFE (post-financing cash flows) at the cost of equity to arrive directly at equity value. The unlevered version is preferred in practice because WACC changes less than FCFE does when leverage shifts, making it more stable.

Q255. Which discount rate matches FCFF, and what value does it produce?

FCFF is discounted at the WACC because FCFF belongs to all capital providers and WACC is the blended cost of all capital. Discounting FCFF at WACC produces enterprise value (the value of the operating business). Matching the cash flow to the right rate is the cardinal rule: an all-investor cash flow needs an all-investor discount rate.

Q256. Which discount rate matches FCFE, and what value does it produce?

FCFE is discounted at the cost of equity (R_e), typically from CAPM $R_e = R_f + \text{beta} \cdot (\text{equity risk premium})$, because FCFE is the residual cash belonging only to shareholders. Discounting FCFE at the cost of equity produces equity value directly, with no enterprise-to-equity bridge needed. Mixing them up (e.g., discounting FCFF at the cost of equity) is a classic and serious error.

Q257. Walk through the enterprise-to-equity value bridge.

Start with enterprise value (EV) from the FCFF/WACC DCF, then $\text{Equity Value} = \text{EV} - \text{Total Debt} - \text{Preferred} - \text{Minority Interest} + \text{Cash \& equivalents}$. You subtract claims that rank ahead of common equity and add non-operating assets like cash. For example, EV 1,000, debt 300, cash 50 gives equity value = $1,000 - 300 + 50 = 750$; dividing by shares outstanding gives intrinsic value per share.

Q258. What is terminal value and why is it needed?

Terminal value (TV) captures the value of all cash flows beyond the explicit forecast horizon, because you cannot forecast individual years to infinity. It is estimated at the end of the explicit period and then discounted back to today like any other cash flow. TV typically accounts for 60-80% of total DCF value, which is why its assumptions deserve heavy scrutiny.

Q259. State the Gordon growth (perpetuity growth) terminal value formula.

$$\text{TV} = \text{FCF}_{(n+1)} / (\text{WACC} - g) = \text{FCF}_n * (1+g) / (\text{WACC} - g)$$
, where FCF_n is the final-year free cash flow, g is the perpetual growth rate, and WACC is the discount rate. For example, final-year FCFF 110, g 3%, WACC 9% gives $\text{TV} = 110 * 1.03 / (0.09 - 0.03) = 113.3 / 0.06 = 1,888$. This TV is then discounted back from year n to the present.

Q260. What is the exit multiple method for terminal value?

The exit multiple method estimates TV by applying a market-based multiple (most commonly EV/EBITDA) to the final-year metric: $\text{TV} = \text{exit multiple} * \text{final-year EBITDA}$. It anchors the terminal value to where comparable businesses currently trade rather than to a growth assumption. The drawback is that it imports current market sentiment into an intrinsic valuation, so it is less "pure" than perpetuity growth.

Q261. How do you sanity-check the Gordon growth TV against the exit multiple TV?

Compute TV both ways and back out the implied figure each method produces: derive the implied exit multiple from your perpetuity-growth TV, and back out the implied perpetuity growth rate from your exit-multiple TV. If the perpetuity method implies a 30x EV/EBITDA exit (versus 8x for peers) or the exit multiple implies 6% perpetual growth (above GDP), your assumptions are inconsistent and need revisiting. The two methods should land in a similar range for a credible valuation.

Q262. What is the maximum sensible value for the perpetual growth rate g , and why?

The perpetual growth rate must be at or below the long-run nominal GDP growth rate of the economy, because no company can grow faster than the whole economy forever (it would eventually become the entire economy). For developed markets this is roughly 2-3%; for India a higher 4-6% nominal can be justified given faster trend growth, but should still be tied to long-run GDP and inflation. A g that exceeds long-run GDP is the single most common terminal-value abuse.

Q263. What is the mid-year convention and why does it raise value?

The mid-year convention assumes cash flows arrive on average in the middle of each year rather than at year-end, so each period is discounted using exponents like 0.5, 1.5, 2.5 instead of 1, 2, 3. This shortens the effective discounting period, raising the present value modestly (typically a few percent). It is more realistic because firms generate cash continuously through the year rather than in a single year-end lump.

Q264. How do you choose the length of the explicit forecast horizon?

The explicit period should run until the business reaches a steady state where growth, margins and reinvestment have stabilised to sustainable levels, so the perpetuity assumption is reasonable; this is usually 5-10 years. High-growth or cyclical firms need longer horizons (10-15 years) so the company can mature before terminal value kicks in. Too short a horizon forces an unrealistically high terminal growth onto a company still growing fast, distorting the valuation.

Q265. What does a WACC-versus-g sensitivity table show and why is it standard?

A sensitivity table arrays implied value (or per-share price) across a grid of WACC on one axis and terminal growth g on the other, showing how value responds to the two assumptions that drive it most. It is standard because point estimates give false precision; the grid communicates a value range and reveals how fragile the conclusion is. If a 0.5% change in WACC or g swings value by 30%, the model is telling you the answer is highly assumption-dependent.

Q266. What is a reverse DCF and what does it tell you?

A reverse DCF takes the current market price as given and solves for the assumptions (usually the growth rate or operating margin) that the market is implicitly pricing in. Instead of asking "what is the company worth," it asks "what must you believe for today's price to be fair." For example, if a stock's price only makes sense at 18% revenue growth for 10 years, you can judge whether that expectation is plausible, which is often more useful than a forward DCF's single output.

Q267. Why is a standard DCF inappropriate for banks and insurance companies?

For banks and insurers, debt (deposits, policy reserves) is raw material in their operations rather than just financing, so FCFF, CapEx and net working capital are not meaningful and enterprise value is ill-defined. The interest spread is the core business, so you cannot separate "operating" from "financing." Instead, analysts use equity-based methods: a dividend discount model, an FCFE/excess-capital model, or relative valuation on P/B and return-on-equity.

Q268. What does it mean to normalise or use through-cycle FCF, and when does it matter?

Normalising means basing the forecast (and especially the terminal year) on mid-cycle or sustainable cash flows rather than a peak or trough year, smoothing out temporary distortions. It matters most for cyclical businesses (commodities, metals, autos, real estate) where a peak-year FCF capitalised into perpetuity wildly overstates value. The fix is to use average margins/returns over a full cycle, or to ensure the terminal year reflects a normalised steady state.

Q269. How are share buybacks treated in a DCF?

In an enterprise-value DCF, buybacks are a financing/distribution decision and do not change FCFF or enterprise value; they reroute cash that is already captured in the model to shareholders. They matter at the equity stage: a buyback reduces both cash and the share count, so per-share intrinsic value (equity value / shares) can rise even though total equity value is unchanged. Double-counting buybacks as both reinvestment and a return source is a mistake.

Q270. How should stock-based compensation (SBC) be handled in a DCF?

SBC is a real economic cost that transfers value from shareholders to employees, so it should not simply be added back as a "non-cash" item the way naive cash-flow templates do. Best practice is either to treat SBC as a cash operating expense (do not add it back) or to add back the cash effect but then increase the diluted share count for the new shares issued. Ignoring SBC overstates free cash flow and double-favours the company by also inflating value before dilution.

Q271. List several of the most common DCF mistakes.

Common errors include: mismatching cash flow and discount rate (FCFF with cost of equity, or FCFE with WACC); a terminal growth rate above long-run GDP; inconsistent forecasts (revenue growth without matching CapEx/working-capital investment); double-counting items like SBC or buybacks; using a too-short horizon that dumps unstabilised growth into the terminal value; and forgetting to adjust the EV-to-equity bridge for items like minority interest, leases or pensions. Each can swing the valuation materially.

Q272. How does a change in WACC move the valuation?

Value moves inversely to WACC: a higher discount rate lowers every present value and, because WACC appears in the denominator ($WACC - g$) of the terminal value, the effect is amplified in the terminal portion. As WACC approaches g , the terminal value explodes toward infinity, so small WACC changes produce large value swings when the spread ($WACC - g$) is small. This convexity is exactly why the sensitivity grid and a sensible WACC are critical.

Q273. How does a change in the terminal growth rate g move the valuation?

A higher g raises value both by increasing the numerator $FCF \cdot (1+g)$ and by shrinking the denominator ($WACC - g$), so the impact is non-linear and concentrated in the terminal value. The closer g is to WACC, the more violently value responds to small changes in g . This is why a disciplined cap on g (at or below long-run GDP) protects the model from runaway terminal values.

Q274. How can you value a company that currently has negative free cash flow?

Forecast forward to the point where the business turns FCF-positive as it scales and margins improve, so the explicit period captures the path from cash burn to maturity; the early negative years simply enter as negative present values. The value then rests heavily on the later years and terminal value, so the model must justify the route to profitability (unit economics, margin trajectory, reinvestment). A longer explicit horizon is essential so that terminal value is anchored on a normalised, profitable steady state rather than on a loss-making year.

Q275. What special care does a DCF for a high-growth firm require?

High-growth firms need a longer explicit horizon (often 10-15 years), an explicit fade of growth and margins toward mature, sustainable levels by the terminal year, and reinvestment (CapEx, working capital, R&D) that genuinely supports the assumed growth. The terminal growth must still respect the GDP ceiling no matter how fast the firm grows today. The biggest risk is capitalising a high growth/high margin terminal year that the company will never sustain, so a sanity check against the implied exit multiple is vital.

Q276. Why is enterprise value (not equity value) the natural output of an unlevered DCF?

Because FCFF is the cash available to all capital providers, discounting it at WACC values the entire operating enterprise, independent of how it is financed, so the result is enterprise value. Equity is then derived by subtracting net debt and other senior claims in the bridge. This separation is useful because it lets you value the operating business cleanly and then handle the capital structure explicitly.

Q277. How do you build WACC for use in an FCFF DCF, and is it constant?

$WACC = (E/V) \cdot Re + (D/V) \cdot Rd \cdot (1-t)$, where E and D are the market values of equity and debt, $V = E+D$, Re is the cost of equity (often CAPM), Rd is the pre-tax cost of debt, and t is the marginal tax rate; the $(1-t)$ captures the tax deductibility of interest. WACC is not necessarily constant: if leverage is expected to change over the forecast (e.g., an LBO deleveraging), WACC should be re-estimated each period, or you should switch to APV (adjusted present value). For Indian companies, use the local risk-free rate (10-year G-Sec) and an India equity risk premium so the discount rate is currency-consistent with INR cash flows.

Q278. Provide a compact end-to-end worked DCF example.

Assume FCFF of 100, 110, 121 over years 1-3 (10% growth), WACC 10%, terminal g 3%. Terminal value at year 3 = $121 \cdot 1.03 / (0.10 - 0.03) = 124.6 / 0.07 = 1,781$; discount the three FCFFs and the TV at 10%: $PV(FCFF) = 100/1.1 + 110/1.21 + 121/1.331 = 90.9 + 90.9 + 90.9 = 272.7$, $PV(TV) = 1,781/1.331 = 1,338$. Enterprise value = $272.7 + 1,338 = 1,611$; if net debt is 300, equity value = 1,311, and dividing by shares gives intrinsic value per share. Note the terminal value is ~83% of EV here, underscoring how dominant the terminal assumptions are.

B2. Relative Valuation & Multiples (Q279-Q308)

Q279. What is relative valuation, and what is its core underlying assumption?

Relative valuation prices an asset by comparing it to similar assets using standardized multiples (e.g., P/E, EV/EBITDA), rather than by discounting its own cash flows. The core assumption is that the market prices comparable companies correctly on average, so applying a peer multiple to the target's metric yields a fair value. If the entire peer set is mispriced (a bubble or a depressed sector), relative valuation simply inherits that mispricing.

Q280. What is a "multiple" and why are multiples standardized?

A multiple is a ratio of value (price or enterprise value) to a fundamental driver (earnings, EBITDA, sales, book value), e.g., $P/E = \text{Price per share} / \text{EPS}$. Standardizing by a per-unit driver lets you compare a 500 crore company with a 50,000 crore company on the same scale. The implicit logic is that two firms with the same fundamentals (growth, risk, returns) should trade at the same multiple.

Q281. What is the difference between trading comps and precedent transaction comps?

Trading comps (comparable company analysis) use multiples of currently listed peers trading in the public market, reflecting standalone, minority-stake values. Precedent transaction comps use multiples paid in actual M&A deals, which embed a control premium and expected synergies. Transaction comps therefore usually give higher multiples and set a ceiling/acquisition value, while trading comps give a "no-control" market value.

Q282. What is the difference between an enterprise-value multiple and an equity multiple?

EV multiples (EV/EBITDA, EV/EBIT, EV/Sales) value the whole firm available to all capital providers (debt + equity), so the denominator must be a pre-interest, pre-financing metric. Equity multiples (P/E, P/B, P/S) value only the shareholders' claim, so the denominator must be post-interest (net income, book equity). The cardinal rule is consistency: never put an equity-side numerator over a firm-side denominator (e.g., Price/EBITDA is wrong; EV/Net income is wrong).

Q283. Write out the EV bridge: how do you get from equity market cap to enterprise value?

$$EV = \text{Equity market cap} + \text{Total debt} + \text{Minority (non-controlling) interest} + \text{Preferred stock} - \text{Cash and cash equivalents}$$

The logic: an acquirer buying the equity must also assume the debt, preferred, and minority claims, but gets the cash which can repay debt. Example: market cap 800, debt 300, minority 20, preferred 0, cash 100 gives $EV = 800 + 300 + 20 + 0 - 100 = 1,020$.

Q284. Why do we add minority interest and subtract cash when computing enterprise value?

Minority interest is added because the income statement consolidates 100% of a partly-owned subsidiary's EBITDA, so EV must include 100% of the claim against it for consistency. Cash is subtracted because it is a non-operating asset that the acquirer can use to pay down the purchase price, and EV is meant to capture only the operating (core) business value.

Q285. Why is EV/EBITDA called "capital-structure neutral" and often preferred over P/E?

EBITDA is computed before interest expense, so it is unaffected by how the firm splits its financing between debt and equity; EV (the numerator) also covers all capital providers, so the ratio is independent of leverage. P/E, by contrast, uses net income (after interest) and price (equity only), so two identical businesses with different debt loads show different P/Es. This makes EV/EBITDA better for comparing peers with very different leverage and for cross-border peers.

Q286. Beyond leverage neutrality, what other advantages does EV/EBITDA have over P/E?

EBITDA strips out depreciation and amortization, reducing distortion from different depreciation policies and accounting choices, and it is usually positive even when net income is negative, so loss-making but cash-

generative firms remain comparable. EV/EBITDA is also less affected by one-off below-the-line items and differing tax rates. It is the workhorse multiple in M&A, LBOs, and capital-intensive sectors.

Q287. What are the main weaknesses of EV/EBITDA?

EBITDA ignores capital expenditure and changes in working capital, so it overstates the cash generation of capex-heavy businesses (a steel plant and a software firm at the same EV/EBITDA are not equally cheap). It also ignores the depreciation that proxies for real asset consumption — Warren Buffett's critique that "EBITDA pretends capex doesn't exist." For capex-heavy firms, EV/EBIT or EV/(EBITDA - capex) is often more honest.

Q288. When is P/E a perfectly fine multiple, and when does it mislead?

P/E works well for mature, profitable companies with similar capital structures and stable, comparable accounting — for example comparing two FMCG majors. It misleads when earnings are negative or near zero (the ratio becomes meaningless or huge), when leverage differs sharply across peers, or when one-off items, different tax regimes, or aggressive accounting distort net income. In cyclical firms, a low trailing P/E at the peak of the cycle can be a value trap.

Q289. What is P/B (price-to-book) and why does it suit banks and financials?

P/B = Market cap / Book value of equity (or Price per share / Book value per share). It suits banks, insurers, and NBFCs because their balance sheets are marked largely to financial assets, so book value is a meaningful, near-economic measure of net worth, and their "earnings power" is closely tied to equity capital. For these firms earnings are volatile and EBITDA is not meaningful (interest is their core revenue), so P/B and P/E on normalized earnings dominate. In India, banks like HDFC Bank are routinely valued on P/B.

Q290. How does P/B link to ROE, and why does that matter for justified valuation?

A firm's justified P/B is driven by how its return on equity compares to its cost of equity: $P/B = (ROE - g) / (r - g)$ in a Gordon-style framework. A bank earning ROE above its cost of equity should trade above 1x book; one earning below should trade below book. So when comparing two banks at the same P/B, the one with higher sustainable ROE is the cheaper, better value.

Q291. What is P/S (price-to-sales), and when is it useful?

P/S = Market cap / Revenue (the equity-side version; EV/Sales is the firm-side analogue). It is useful for early-stage or loss-making companies with no meaningful earnings (e.g., young SaaS or e-commerce firms), and because revenue is harder to manipulate than earnings. Its weakness is that it ignores profitability entirely — a high-revenue, low-margin firm can look cheap on P/S while being a poor business.

Q292. What is the PEG ratio and how do you interpret it?

PEG = (P/E) / (annual EPS growth rate in %), designed to adjust P/E for growth so high-growth stocks don't automatically look expensive. A PEG around 1.0 is often taken as fairly valued; below 1.0 may indicate undervaluation relative to growth, above 1.0 overvaluation. Example: a stock with P/E of 30 growing earnings at 30% has PEG = $30/30 = 1.0$, while P/E 30 growing at 15% has PEG = 2.0.

Q293. What are the pitfalls of the PEG ratio?

PEG assumes a linear, proportional relationship between P/E and growth, which is theoretically wrong — value rises faster than linearly with growth and also depends on risk and payout. It relies on a single forecast growth number that is easily gamed or optimistic, ignores differences in risk/quality between firms, and breaks down for low-growth or negative-growth companies. It is a quick screen, not a rigorous valuation.

Q294. What is the difference between trailing and forward multiples, and which is generally preferred?

A trailing multiple uses the last reported (LTM/historical) metric; a forward multiple uses a forecast (e.g., next-twelve-months or next-fiscal-year EBITDA/EPS). Forward multiples are generally preferred because

valuation is about the future and they reflect expected growth, but they depend on the reliability of estimates. For a growing company, the forward P/E is lower than the trailing P/E; mixing trailing for one peer and forward for another is a classic comparability error.

Q295. How do you choose a good peer set for trading comps?

Select companies similar on the value-relevant dimensions: same industry/business model, comparable size, similar growth and margin profile, similar geography/regulatory regime, and similar capital intensity. The goal is firms whose fundamentals (growth, risk, returns on capital) are close enough that they "deserve" the same multiple. A tight set of 5-8 true comparables beats a long list of loosely related names; you then screen for outliers and accounting comparability.

Q296. Why might a stock legitimately trade at a premium or discount to its peers?

Differences in fundamentals justify multiple differences: higher expected growth, higher and more sustainable ROIC/ROE, lower risk (beta, leverage, earnings volatility), better management/governance, or stronger competitive moat all warrant a premium. Discounts arise from cyclicalities, governance concerns, customer concentration, lower liquidity, or holding-company/conglomerate structures. A premium is only "expensive" if it exceeds what the superior fundamentals justify.

Q297. A peer trades at 12x EV/EBITDA. Your target has EBITDA of 500 crore, debt of 800 crore, and cash of 200 crore. What is the implied equity value?

Implied EV = $12 \times 500 = 6,000$ crore. Then equity value = EV - net debt (and any preferred/minority): $6,000 - (800 - 200) = 6,000 - 600 = 5,400$ crore. Divide by shares outstanding to get the implied price per share. This is the standard "apply multiple, then walk down the EV bridge to equity" workflow.

Q298. State the Gordon Growth (single-stage DDM) formula and the meaning of each term.

$P_0 = D_1 / (r - g)$, where P_0 is the fair price today, D_1 is next year's expected dividend ($D_1 = D_0 \times (1+g)$), r is the cost of equity, and g is the constant perpetual dividend growth rate. It assumes dividends grow forever at a constant rate g , with $r > g$ required for a finite value. Example: $D_1 = 20$, $r = 12\%$, $g = 4\%$ gives $P_0 = 20 / (0.12 - 0.04) = 250$.

Q299. Why is the Gordon Growth model so sensitive, and when is it appropriate?

Because the denominator is $(r - g)$, small changes in either input swing the value dramatically — if g moves from 4% to 5% with $r = 12\%$, value jumps from 250 to 285 (+14%) on the same D_1 . It is appropriate only for stable, mature firms paying predictable dividends with growth at or below the long-run economy rate (g cannot exceed r , and realistically not the GDP growth rate forever). It is widely used for regulated utilities and large banks.

Q300. What is a multi-stage DDM and why use it?

A multi-stage DDM models an explicit high-growth period with year-by-year dividends, then applies a Gordon-growth terminal value for the stable phase: $P_0 = \sum D_t / (1+r)^t$ for the explicit years + $[D_{(n+1)} / (r - g_{\text{stable}})] / (1+r)^n$. It is used when a firm cannot grow at a constant rate forever — e.g., a fast-growing company that will mature. The two-stage (and H-model) variants let growth fade from high to sustainable, which is more realistic than assuming a single perpetual g .

Q301. What is sum-of-the-parts (SOTP) valuation and when is it used?

SOTP values a diversified company by valuing each business segment separately with the multiple or method appropriate to that segment, then summing the parts and adjusting for net debt, corporate costs, and any holding-company discount. It is used for conglomerates and multi-segment firms where a single blended multiple hides value — e.g., valuing a group's IT arm on EV/EBIT, its NBFC on P/B, and its telecom on EV/EBITDA. Indian conglomerates like Reliance or the Tata/Aditya Birla groups are common SOTP candidates.

Q302. What is a holding-company discount and why does it appear in SOTP?

A holding-company (or conglomerate) discount is the gap between the market value of a holding company and the sum of its stakes in listed/unlisted subsidiaries. It arises from double taxation on dividends passing up, lack of control over subsidiary cash flows, capital-allocation concerns, and reduced transparency. In SOTP for Indian holdcos, analysts often apply a 20-50% discount to the underlying NAV to reflect this.

Q303. What is a control premium, and why do precedent transaction multiples include it?

A control premium is the extra amount an acquirer pays above the unaffected market price to obtain control — the ability to change management, capital structure, and strategy, and to capture synergies. Because precedent transactions are completed acquisitions, their multiples bake in this premium (often 20-40% over the pre-deal share price). That is why transaction comps typically exceed trading comps and are used to estimate acquisition (takeover) value rather than standalone value.

Q304. How do synergies affect transaction multiples, and should you pay for them?

Synergies (cost savings + revenue gains) raise the post-deal earnings/EBITDA the acquirer expects, so a buyer may justify a higher multiple on the target's standalone metric. But standard practice is that the seller should not capture all synergy value — the buyer pays a control premium partly funded by synergies but ideally retains a share, otherwise the deal destroys value for the acquirer. When using transaction comps you should be aware their multiples are inflated by deal-specific synergies that may not apply to your target.

Q305. What is a "football field" chart and what is it used for?

A football field is a horizontal bar chart that displays the valuation range for a company from each method side by side — 52-week trading range, trading comps, precedent transactions, DCF, LBO analysis — so you can see where the ranges overlap. It is the standard summary exhibit in valuation/M&A decks because it conveys that valuation is a range, not a point, and shows which methods agree. The overlap zone often guides the negotiation or fairness-opinion range.

Q306. How do you reconcile a DCF value that is much higher (or lower) than what comps imply?

First check that both are internally consistent — same forecast assumptions, leverage, and metric definitions — then identify the source of divergence: comps reflect current market sentiment for the sector, while DCF reflects your explicit cash-flow and terminal-value assumptions. If DCF is far above comps, either the market is pricing the sector pessimistically (an opportunity) or your DCF growth/terminal assumptions are too aggressive; reverse-engineering the implied growth/multiple from each helps locate the disagreement. Analysts typically present both and triangulate rather than forcing them to match.

Q307. Why are sector-specific multiples used, and give examples across sectors?

Each industry has a value driver that standard accounting multiples capture poorly, so practitioners use a metric tied to the real economics. Examples: banks and insurers on P/B (and P/E on normalized earnings); telecom and cable on EV/EBITDA and EV per subscriber; hotels on EV per room; oil & gas on EV/EBITDAX or EV per barrel of reserves; real estate/REITs on price-to-NAV and EV per square foot; internet/SaaS on EV/Sales, EV/ARR, and EV per user/MAU. The chosen multiple should correlate with the cash the asset ultimately produces.

Q308. Why are EV/Subscriber and similar "per-unit operating" multiples sometimes used despite ignoring profitability?

Per-unit metrics (EV per subscriber, per user, per room, per ton of capacity) are used when current earnings are negative or volatile but the operational base is the true long-run value driver, allowing comparison of firms not yet profitable — common for early-stage telecom, streaming, and internet companies. They normalize by scale of the customer/asset base, which is more stable than fluctuating early earnings. The danger is that they ignore the quality, monetization, and cost of serving each unit, so a high EV/subscriber

can flatter a business with poor ARPU or churn; they are best used alongside profitability and lifetime-value analysis.

B3. Research Process, Modeling & Technical Analysis (Q309-Q339)

Q309. What is the difference between sell-side and buy-side equity research?

Sell-side research is produced by brokerages and investment banks (e.g., Morgan Stanley, Kotak Institutional Equities) and is distributed to clients to generate trading commissions and corporate-access revenue; it publishes ratings, target prices and detailed models. Buy-side research sits inside asset managers, mutual funds, hedge funds and insurers (e.g., a fund house's internal analyst) and is kept private to drive the firm's own portfolio decisions. Sell-side analysts cover more names and write more publicly; buy-side analysts go deeper on fewer names and answer only to their portfolio managers.

Q310. What are the core deliverables of a sell-side equity research analyst?

The four staples are: (1) an earnings/financial model forecasting revenue, margins, EPS and cash flow; (2) a rating (Buy/Hold/Sell or Overweight/Neutral/Underweight); (3) a 12-month target price derived from a valuation method; and (4) written notes — an initiating-coverage report plus ongoing updates on results, management meetings and news. The analyst also fields client calls and arranges corporate access. In India, sell-side research must carry SEBI Research Analyst (RA) registration and disclosures.

Q311. What does it mean to "initiate coverage" on a stock?

Initiating coverage is the first formal research report an analyst publishes on a company, establishing a rating, target price and full investment thesis. It is typically long and comprehensive — business overview, industry context, model, valuation, risks — because it sets the baseline against which all later updates are measured. After initiation the analyst is expected to maintain the stock with regular updates.

Q312. How is a target price typically derived and what does it represent?

A target price is the analyst's estimate of fair value over a stated horizon, usually 12 months, and the rating is set by the implied upside/downside versus the current price. It is commonly derived from a forward multiple (e.g., target P/E times next-year EPS), a DCF, a sum-of-the-parts, or a blend. Example: if FY+1 EPS is INR 50 and the justified P/E is 20x, the target price is $50 * 20 = \text{INR } 1,000$.

Q313. What is the difference between top-down and bottom-up revenue forecasting?

Top-down starts from a large aggregate and narrows: total addressable market (TAM) times an assumed market share, or GDP/industry growth scaled to the company. Bottom-up builds revenue from the company's own operational drivers — units * price, stores * sales-per-store, subscribers * ARPU — and sums them. Bottom-up is usually more defensible and auditable; analysts often build bottom-up and sanity-check it against a top-down TAM number.

Q314. What are "revenue drivers" and why model revenue by drivers instead of a single growth rate?

Revenue drivers are the underlying volume and price levers that mechanically produce sales, e.g., $\text{Revenue} = \text{Volume} * \text{Price}$, or for a telco $\text{Revenue} = \text{Subscribers} * \text{ARPU}$, for retail $\text{Stores} * \text{Avg sales per store}$, for SaaS $\text{Customers} * \text{ACV with churn}$. Driver-based modeling makes assumptions explicit and testable, lets you flex one lever at a time for scenarios, and ties forecasts to operating reality rather than an arbitrary single growth percentage that hides the logic.

Q315. How do you forecast operating margins in a model?

Forecast each cost line either as a percentage of revenue or via its own driver, then derive margin as the residual. Distinguish variable costs (scale with volume — raw materials, freight) from largely fixed costs (rent, core salaries, depreciation) so the model captures operating leverage: as revenue grows, fixed costs

spread over more units and EBIT margin expands. Anchor assumptions to historical trends, management guidance and capacity, and check that gross, EBITDA and EBIT margins move coherently.

Q316. What makes a good investment thesis?

A good thesis is a clear, differentiated and falsifiable view on why a stock will outperform: what the market is mispricing, why you are right (variant perception), the catalyst and timeframe that closes the gap, and the valuation upside if you are correct. It states explicit assumptions and the specific data points that would prove it wrong. "Good company" is not a thesis; "good company whose margin recovery the Street is underestimating, re-rating on H2 results" is.

Q317. What is a catalyst, and why does it matter to a research call?

A catalyst is a specific, datable event expected to force the market to re-price a stock toward fair value — earnings beats, new product launch, regulatory approval, capacity coming online, margin inflection, M&A, or index inclusion. A cheap stock can stay cheap indefinitely (a "value trap") without a catalyst, so identifying one gives the call a timeframe and converts a valuation gap into an actionable trade. Catalysts are central to differentiating a Buy from a "fairly valued."

Q318. What is a bull case and bear case in a research report?

The bull case and bear case are scenario valuations bracketing the base-case target: the bull case stacks optimistic but plausible assumptions (faster growth, margin expansion, higher multiple) and the bear case the pessimistic ones (share loss, margin compression, de-rating). Together they frame the risk/reward — e.g., base INR 1,000, bull INR 1,400, bear INR 700 against a price of 950 gives roughly +47% / -26% skew. A favorable skew (more upside than downside) strengthens a Buy.

Q319. What is "scuttlebutt" or primary research in equity analysis?

Coined by Philip Fisher, scuttlebutt is gathering insight from primary, ground-level sources rather than filings alone — talking to customers, suppliers, distributors, ex-employees, competitors and channel checks at stores. It helps verify demand, pricing power and management claims, and often surfaces variant perception before it shows in reported numbers. The constraint is legality and ethics: it must not cross into material non-public information (insider trading), a SEBI/SEC violation.

Q320. Outline Porter's Five Forces and their use in industry analysis.

The five forces that determine industry profitability are: (1) threat of new entrants, (2) bargaining power of suppliers, (3) bargaining power of buyers, (4) threat of substitutes, and (5) rivalry among existing competitors. An industry where all five are weak (high entry barriers, fragmented suppliers and buyers, few substitutes, benign rivalry) tends to earn structurally high returns. Analysts use the framework to judge whether a company's margins are durable or will be competed away.

Q321. What is an economic moat, and what are Morningstar's moat sources?

An economic moat is a durable structural advantage that lets a firm earn returns above its cost of capital for an extended period while fending off competitors. Morningstar identifies five sources: (1) intangible assets (brands, patents, regulatory licenses), (2) switching costs, (3) the network effect, (4) cost advantages (scale, process, location), and (5) efficient scale (a niche market profitably served by a few). Morningstar rates moats as none, narrow, or wide.

Q322. What is the Competitive Advantage Period (CAP) and why does it matter for valuation?

CAP (also called the fade period) is the length of time a company is expected to earn returns on invested capital above its WACC before competition erodes them to the cost of capital. A wider moat implies a longer CAP, which justifies a higher multiple because value-creating growth persists longer. In a DCF it shows up as how many years you let ROIC exceed WACC before fading spreads to zero in the terminal value.

Q323. What are the main red flags of poor earnings quality?

Key red flags: net income growing while operating cash flow stagnates or falls (earnings not converting to cash), rising Days Sales Outstanding (DSO) signaling aggressive or uncollectible sales, rising inventory days, recurring "one-time" or exceptional items, frequent restatements, large gaps between reported and tax accounting, capitalizing costs that should be expensed, and bill-and-hold or end-of-quarter revenue spikes. Each suggests reported profits may be overstated or unsustainable.

Q324. Why is rising Days Sales Outstanding (DSO) a warning sign?

DSO measures how long it takes to collect receivables: $DSO = (\text{Accounts Receivable} / \text{Revenue}) * 365$. If DSO rises faster than revenue, the company may be booking sales it hasn't collected — through looser credit terms, channel stuffing, or premature revenue recognition — inflating reported revenue and earnings while cash lags. Example: receivables up 40% on revenue up 10% means DSO is climbing sharply and the quality of that growth is suspect.

Q325. What is the cash conversion ratio and what does a falling ratio indicate?

Cash conversion compares cash generation to accounting profit, commonly $OCF / \text{Net Income}$ (or $FCF / \text{Net Income}$), with a healthy figure near or above 1.0 over time. A persistently falling ratio means profits are increasingly tied up in working capital or driven by non-cash accruals rather than real cash — a classic earnings-quality concern. It can flag aggressive revenue recognition, ballooning receivables/inventory, or capitalized expenses propping up reported income.

Q326. What is channel stuffing and how would you detect it?

Channel stuffing is shipping more product to distributors/retailers than end-demand warrants, often near quarter-end, to inflate current-period revenue at the expense of future periods. Detect it via spiking DSO and receivables, rising distributor inventory, a sales surge followed by returns or a sharp next-quarter drop, and lengthening payment terms or generous incentives. Channel checks (scuttlebutt with distributors) are often the fastest way to confirm it.

Q327. What is the accruals ratio and how is it used to assess earnings quality?

The accruals ratio isolates the non-cash portion of earnings; the balance-sheet version is $\text{Accruals Ratio} = (\text{NOA}_{\text{end}} - \text{NOA}_{\text{begin}}) / \text{Avg NOA}$, where NOA is net operating assets, or equivalently $(\text{Net Income} - \text{OCF} - \text{FCF} \dots) / \text{Avg NOA}$ via the cash-flow version $(\text{NI} - \text{CFO} - \text{CFI}) / \text{Avg NOA}$. High positive accruals mean earnings are driven by accounting estimates rather than cash, and academic evidence (Sloan, 1996) shows high-accrual firms tend to see earnings reverse and underperform. Lower, stable accruals signal higher-quality, cash-backed earnings.

Q328. How is ESG incorporated into modern equity research?

ESG analysis assesses environmental, social and governance factors as drivers of risk and long-term value rather than just ethics screens — e.g., carbon-transition risk to a fossil producer, labor and supply-chain risk, and especially governance (board independence, related-party transactions, promoter pledging). Analysts integrate material ESG factors into the discount rate, growth/margin assumptions, or the valuation multiple, and may flag controversies as thesis risks. In India, SEBI mandates BRSR (Business Responsibility and Sustainability Reporting) for the top 1,000 listed firms, giving analysts standardized ESG data.

Q329. How do you value a loss-making, high-growth company?

Because near-term earnings are negative, anchor on revenue and the path to profitability: use EV/Sales or EV/Gross Profit versus growth-adjusted peers, a DCF built to a normalized steady-state margin in the terminal year, or unit economics (LTV/CAC, contribution margin, cohort retention) to prove the model works at scale. Test whether the implied terminal margin and market share are realistic, and stress the cash runway and dilution risk. Example: a SaaS firm at INR 500 cr revenue growing 60% with 80% gross margin and clear operating leverage can justify a high EV/Sales if the path to ~25% EBIT margin is credible.

Q330. What defines an uptrend and a downtrend in technical analysis?

An uptrend is a sequence of higher highs and higher lows — each rally peak and each pullback trough is above the prior one — reflecting net buying pressure; a downtrend is the mirror, lower highs and lower lows. The trend is presumed intact until that structure breaks (e.g., an uptrend makes a lower low). "The trend is your friend" captures the bias to trade in the direction of the prevailing structure rather than against it.

Q331. What are support and resistance levels?

Support is a price zone where buying interest has repeatedly halted declines; resistance is a zone where selling has repeatedly capped advances. They form at prior swing highs/lows, round numbers, and high-volume nodes, and reflect clustered orders and trader memory. A key principle is polarity: once decisively broken, support tends to become resistance and vice versa, which is why breakouts and retests are watched closely.

Q332. What is the difference between a Simple and an Exponential Moving Average?

A Simple Moving Average (SMA) is the equal-weighted mean of the last N closes: $SMA = (P1 + P2 + \dots + PN) / N$. An Exponential Moving Average (EMA) weights recent prices more heavily: $EMA_{today} = Price * k + EMA_{yesterday} * (1 - k)$, where $k = 2 / (N + 1)$. The EMA therefore reacts faster to new price action (less lag) but is more prone to whipsaws; SMA is smoother but slower.

Q333. What are the golden cross and death cross?

A golden cross is a bullish signal where a shorter moving average crosses above a longer one — classically the 50-day SMA crossing above the 200-day SMA — suggesting upward momentum is building. A death cross is the bearish opposite: the 50-day crossing below the 200-day. Both are lagging signals confirming a trend already underway rather than predicting turns, and work best filtered with volume and broader context.

Q334. How is the RSI calculated and how is it interpreted?

The Relative Strength Index is a momentum oscillator bounded 0-100: $RSI = 100 - 100 / (1 + RS)$, where $RS = \text{average gain} / \text{average loss}$ over a lookback (default 14 periods). Readings above 70 are conventionally overbought and below 30 oversold, though in strong trends RSI can stay extended. A widely watched signal is divergence — price makes a new high but RSI does not — hinting the move is losing momentum.

Q335. How does the MACD indicator work?

MACD measures momentum via the gap between two EMAs: $MACD\ line = EMA(12) - EMA(26)$, the $signal\ line = EMA(9)$ of the MACD line, and the $histogram = MACD\ line - signal\ line$. A bullish trigger is the MACD line crossing above the signal line; bearish is crossing below; crossing the zero line marks a shift in the underlying trend. As with RSI, MACD/price divergence is used to anticipate weakening momentum.

Q336. What are Bollinger Bands and how are they used?

Bollinger Bands plot a middle band (typically a 20-period SMA) with upper and lower bands at plus/minus two standard deviations of price: $Upper = SMA(20) + 2 * \sigma$, $Lower = SMA(20) - 2 * \sigma$. The bands widen when volatility rises and contract ("squeeze") when it falls, often before a large move. Price tagging a band is not itself a buy/sell signal — in a trend price can "walk the band" — so it is read alongside trend and volume.

Q337. Why is volume important for confirming price moves?

Volume gauges conviction behind a move: a breakout or trend advance on rising volume is more credible, while a move on thin volume is suspect and prone to failure. Classic interpretation — rising price on rising volume confirms an uptrend; rising price on falling volume warns of weakening demand and possible reversal. Volume confirmation is also key for validating breakouts from support/resistance or chart patterns.

Q338. Describe the doji, engulfing, and hammer candlestick patterns.

A doji has open and close nearly equal (a tiny body), signaling indecision and a possible turning point. A bullish engulfing is a large up-candle whose body fully engulfs the prior down-candle, indicating a momentum shift to buyers (bearish engulfing is the inverse). A hammer is a small body with a long lower wick appearing after a decline, showing buyers rejected lower prices — a potential bottom. All gain reliability when confirmed by the next candle, support/resistance, and volume.

Q339. What are the head-and-shoulders and double top/bottom patterns, and how does technical analysis complement fundamental analysis?

Head-and-shoulders is a bearish reversal pattern — three peaks with a higher middle "head" — that signals a top when price breaks the neckline; the inverse pattern signals a bottom. A double top (two failed highs at resistance) is bearish and a double bottom (two failed lows at support) is bullish, confirmed on a break of the intervening level. TA does not replace fundamentals: fundamentals (valuation, thesis, catalysts) answer *what* and *why* to own, while TA helps with *when* — timing entries/exits, sizing, setting stops near support, and gauging crowd momentum around the fundamental call.

C. Risk - Market & Credit

C1. Market Risk: VaR, Volatility & Backtesting (Q340-Q374)

Q340. What is Value at Risk (VaR) and what three parameters define it?

VaR is the maximum loss on a portfolio that is not expected to be exceeded over a given holding period, at a given confidence level, under normal market conditions. It is defined by three parameters: the horizon (e.g. 1 day, 10 days), the confidence level (e.g. 95%, 99%), and the base currency/value of the position. Formally, VaR is the loss quantile: $P(\text{Loss} > \text{VaR}) = 1 - \text{confidence}$.

Q341. How do you read the statement "1-day 99% VaR = \$1m"?

It means that on a normal trading day there is a 99% probability the loss will be no more than \$1m, or equivalently a 1% chance (roughly 1 day in 100, about 2-3 days a year) the loss will exceed \$1m. Crucially, VaR tells you the threshold but says nothing about how bad the loss is on those breach days — it could be \$1.1m or \$50m.

Q342. What is the difference between absolute VaR and relative VaR?

Absolute VaR measures the loss relative to zero (the worst loss in money terms), while relative VaR measures the loss relative to the expected return or a benchmark. Over short horizons like 1 day the expected return (drift) is tiny, so the two are nearly identical and drift is usually ignored; over longer horizons the distinction matters more.

Q343. What are the three main methods for calculating VaR?

(1) Historical simulation — re-price the portfolio using actual historical returns and read off the loss quantile; (2) Parametric / variance-covariance — assume a distribution (usually normal) and compute $\text{VaR} = z\sigma V$ from estimated volatility and correlations; (3) Monte Carlo simulation — generate thousands of random scenarios from an assumed model and read off the quantile. Each trades off realism, speed, and ease of implementation.

Q344. What is the parametric (variance-covariance) VaR formula and what are the standard z-values?

$\text{VaR} = z * \sigma * V$, where z is the standard-normal quantile, σ is the return volatility over the horizon, and V is the portfolio value. The z -multipliers are 1.65 for 95% and 2.33 for 99% (one-tailed). Example: a \$10m portfolio with 1-day $\sigma = 2\%$ has 99% $\text{VaR} = 2.33 * 0.02 * 10\text{m} = \$466,000$.

Q345. Work a quick parametric VaR: 95% 1-day VaR on a 50 crore INR equity book with daily volatility of 1.5%.

$\text{VaR} = z * \sigma * V = 1.65 * 0.015 * 50,00,00,000 = \text{roughly } 1.24 \text{ crore INR}$. So on 95 days out of 100 the one-day loss should stay below about 1.24 crore. Bumping to 99% replaces 1.65 with 2.33, giving about 1.75 crore.

Q346. What are the pros and cons of historical simulation VaR?

Pros: no distributional assumption (it captures fat tails and actual correlations as they occurred), easy to explain, and handles non-linear instruments like options by full revaluation. Cons: it assumes the past window represents the future, the result is bounded by the worst event in the window, it reacts slowly to regime changes, and a short window gives a noisy tail while a long window mixes stale data.

Q347. What are the pros and cons of parametric / variance-covariance VaR?

Pros: fast, analytic, low data needs (just volatilities and a correlation matrix), and easy to decompose into marginal/component contributions. Cons: it assumes returns are normal (or elliptical), so it understates tail

risk for fat-tailed assets, and it handles option-like non-linearity (gamma, vega) poorly unless you add delta-gamma approximations.

Q348. What are the pros and cons of Monte Carlo VaR?

Pros: very flexible — it can use any distribution, capture non-linear payoffs through full revaluation, and model path dependence and complex correlations. Cons: it is computationally expensive (thousands of full repricings), results depend entirely on the chosen model and parameters (model risk), and convergence in the tail is slow, requiring many simulations for a stable 99% number.

Q349. What is Expected Shortfall (ES) / Conditional VaR (CVaR)?

Expected Shortfall is the average loss conditional on being in the tail beyond the VaR level — i.e. the mean of losses worse than VaR. ES at 97.5% answers "if we breach VaR, how bad is it on average?" Formally $ES_{\alpha} = E[\text{Loss} \mid \text{Loss} \geq \text{VaR}_{\alpha}]$, so ES is always at least as large as the corresponding VaR.

Q350. Why does Expected Shortfall give more information than VaR?

VaR is a single threshold and is blind to the shape and severity of the tail beyond it — two portfolios with identical VaR can have wildly different worst-case losses. ES averages over the entire tail, so it "sees" how heavy the tail is and penalises portfolios with catastrophic but rare payoffs (like short deep-out-of-the-money options) that VaR can hide.

Q351. What are the four properties of a coherent risk measure?

A risk measure ρ is coherent (Artzner et al., 1999) if it satisfies: (1) Monotonicity — if portfolio A always loses more than B, $\rho(A) \geq \rho(B)$; (2) Translation invariance — adding cash c reduces risk by c , $\rho(X + c) = \rho(X) - c$; (3) Positive homogeneity — $\rho(\lambda X) = \lambda \rho(X)$ for $\lambda \geq 0$ (scaling); (4) Subadditivity — $\rho(A + B) \leq \rho(A) + \rho(B)$, so merging never increases risk.

Q352. Why is VaR not a coherent risk measure?

VaR satisfies monotonicity, homogeneity, and translation invariance but can violate subadditivity: for non-elliptical or fat-tailed/discrete distributions, the VaR of a combined portfolio can exceed the sum of the standalone VaRs. This is perverse because it implies diversification "increases" measured risk, and it can be gamed by portfolios with payoffs concentrated just beyond the quantile.

Q353. Why is Expected Shortfall coherent while VaR is not?

ES satisfies all four axioms including subadditivity for any distribution, because it averages over the whole tail rather than picking a single quantile point. By integrating the tail, ES properly rewards diversification (ES of a merged book is never more than the sum) and cannot be hidden by lumping losses just beyond the VaR cutoff.

Q354. Give a simple example where VaR violates subadditivity.

Take two independent corporate bonds each with a 4% default probability and a loss of 100 on default, zero otherwise. Each bond's 95% VaR is 0 (default sits in the 4% tail, beyond the 5% point). Combined, the chance at least one defaults is about 7.84%, so the 95% VaR of the two-bond portfolio is 100 — greater than $0 + 0$. Diversifying made VaR worse, violating subadditivity.

Q355. What is the square-root-of-time rule for scaling VaR?

To convert a 1-day VaR to an h -day VaR you scale volatility by $\sqrt{h}$: $VaR_h = VaR_1 * \sqrt{h}$. Example: a 1-day 99% VaR of \$1m becomes a 10-day VaR of $1m * \sqrt{10}$ = about \$3.16m. Basel's old market-risk framework used exactly this to scale a 1-day figure to the 10-day regulatory horizon.

Q356. What assumptions underlie square-root-of-time scaling, and when does it break down?

It assumes returns are independent and identically distributed (i.i.d.) with zero autocorrelation and constant volatility, so variances add linearly across days. It breaks down under volatility clustering (GARCH effects),

mean reversion (which makes it overstate risk), or momentum/autocorrelation and fat tails (which make it understate risk). For options and path-dependent positions it is also unreliable because risk is non-linear in time.

Q357. What is backtesting of VaR and why is it required?

Backtesting compares the number of days actual losses exceeded the predicted VaR (exceptions/breaches) against the number expected from the confidence level. For a 99% 1-day VaR over 250 trading days you expect about 2.5 breaches; far more suggests the model understates risk, far fewer suggests it is too conservative. Regulators (Basel) require banks to backtest internal VaR models to validate them.

Q358. What is the Kupiec POF (proportion of failures) test?

Kupiec's POF test checks whether the observed breach frequency matches the expected rate $p = 1 - \text{confidence}$ using a likelihood-ratio statistic. $LR_POF = -2 \ln \left[\frac{((1-p)^{(N-x)} * p^x)}{((1-x/N)^{(N-x)} * (x/N)^x)} \right]$, where N is the number of observations and x the number of breaches; it is compared to a chi-squared with 1 degree of freedom. It tests unconditional coverage only — whether the total count of breaches is right.

Q359. What is the Christoffersen independence test and conditional coverage test?

Christoffersen's independence test checks whether breaches are clustered in time (a breach today predicting a breach tomorrow) versus being independent — clustering signals the model fails to react to volatility regimes. His conditional coverage test combines Kupiec's coverage test with the independence test: $LR_cc = LR_POF + LR_ind$, compared to chi-squared with 2 degrees of freedom. A good model needs both the right number of breaches and breaches that are independent over time.

Q360. Describe Basel's traffic-light approach to backtesting.

Basel evaluates breaches of the 99% 1-day VaR over the most recent 250 days and assigns a zone: Green (0-4 exceptions) — model fine, no penalty; Yellow (5-9) — possible problem, a graduated capital multiplier add-on applies; Red (10+) — model rejected, maximum add-on and supervisory action. The multiplier scales the capital charge from a base of 3 up to 4 as breaches rise.

Q361. What is realized (historical) volatility and how is it computed?

Realized/simple volatility is the standard deviation of past returns, $\sigma = \sqrt{\frac{1}{(n-1)} * \sum (r_t - \text{mean})^2}$, often computed over a rolling window. It is then annualised by multiplying the daily figure by $\sqrt{252}$. Equal-weighting all observations is its weakness — an old extreme return stays in and then drops out abruptly, causing "ghosting" jumps in the estimate.

Q362. What is EWMA / RiskMetrics volatility and why is it preferred over a simple rolling window?

EWMA (Exponentially Weighted Moving Average) weights recent returns more heavily: $\sigma_t^2 = \lambda \sigma_{t-1}^2 + (1-\lambda)r_{t-1}^2$. RiskMetrics uses $\lambda = 0.94$ for daily data. It responds faster to volatility changes and avoids the abrupt ghosting jumps of equal-weighted windows, while needing only the previous variance and return to update.

Q363. In EWMA, what does the decay factor λ control?

λ sets how fast old observations lose weight: a high λ (e.g. 0.97) gives a smooth, slowly-adapting estimate with a long memory, while a low λ (e.g. 0.90) reacts quickly but is noisier. The weight on the return k days ago is $(1-\lambda)^k$, so with $\lambda = 0.94$ the effective memory is roughly $1/(1-\lambda) = \text{about } 17 \text{ days}$.

Q364. What is the GARCH(1,1) model for volatility?

GARCH(1,1) models conditional variance as $\sigma_t^2 = \omega + \alpha r_{t-1}^2 + \beta \sigma_{t-1}^2$, where ω is a constant, α weights the latest squared return (the "news"), and β weights yesterday's

variance (persistence). It captures volatility clustering and mean reversion to a long-run level, and EWMA is just the special case $\omega = 0$, $\alpha + \beta = 1$.

Q365. In GARCH(1,1), what are persistence and the long-run variance?

Persistence is $\alpha + \beta$; the closer it is to 1, the longer volatility shocks take to decay (typical equity estimates are around 0.98-0.99). The long-run (unconditional) variance is $\omega / (1 - \alpha - \beta)$, to which the model mean-reverts. Example: $\omega = 0.000002$, $\alpha = 0.08$, $\beta = 0.91$ gives long-run variance $0.000002/0.01 = 0.0002$, i.e. a long-run daily vol of about 1.41%.

Q366. What is the GJR-GARCH model and the leverage effect?

GJR-GARCH adds an asymmetry term so negative returns raise future volatility more than positive ones: $\sigma_t^2 = \omega + (\alpha + \gamma I) r_{t-1}^2 + \beta \sigma_{t-1}^2$, where the indicator $I = 1$ when $r_{t-1} < 0$. This captures the leverage effect — empirically, equity markets become more volatile after falls than after equal-sized rallies, which symmetric GARCH and EWMA miss.

Q367. What is volatility clustering and why does it matter for VaR?

Volatility clustering is the empirical fact that large moves tend to be followed by large moves and calm by calm — high and low volatility come in bursts (Mandelbrot's observation). It violates the i.i.d. assumption, so static volatility estimates lag reality; this is exactly why EWMA and GARCH models, which let volatility evolve, give more responsive and accurate VaR than equal-weighted methods.

Q368. What are fat tails (leptokurtosis) and why do they matter for VaR?

Financial return distributions have more probability mass in the extremes than the normal distribution (excess kurtosis > 0), meaning large moves happen far more often than a bell curve predicts. This is why normal-based parametric VaR systematically understates the probability and size of extreme losses, and why ES, historical simulation, or t-distribution assumptions are used to capture the tail better.

Q369. Why does normal-distribution VaR understate tail risk?

The normal distribution assigns vanishingly small probability to large moves, but real markets show fat tails and clustering, so 4-5 sigma events occur far more often than "once in millennia." A 99% normal VaR uses $z = 2.33$, but the true 1% loss quantile of a fat-tailed series is larger, so the model under-reserves precisely in the crashes that matter. This is a core reason regulators supplement VaR with stressed measures and ES.

Q370. What is Stressed VaR (sVaR) and why was it introduced?

Stressed VaR computes VaR using model inputs (volatilities, correlations) calibrated to a continuous 12-month period of significant financial stress, such as 2007-2009, rather than the recent calm window. Basel 2.5 introduced it after 2008 because ordinary VaR estimated on benign recent data badly understated risk going into the crisis; total capital became roughly $\text{VaR} + \text{sVaR}$, sharply raising charges.

Q371. Define marginal, incremental, and component VaR.

Marginal VaR is the change in portfolio VaR per unit change in a position's size (the sensitivity/derivative). Incremental VaR is the actual change in total VaR from adding or removing an entire position or trade. Component VaR allocates the total VaR across positions such that the components sum exactly to portfolio VaR, showing each position's true risk contribution including diversification — useful for risk budgeting.

Q372. What is the diversification benefit in VaR and how is it measured?

Because assets are not perfectly correlated, portfolio VaR is less than the sum of standalone VaRs; the gap is the diversification benefit. Diversification benefit = $\text{sum}(\text{standalone VaRs}) - \text{portfolio VaR}$. For two positions, $\text{portfolio sigma} = \sqrt{w_1^2 s_1^2 + w_2^2 s_2^2 + 2w_1 w_2 \rho s_1 s_2}$, so the lower the correlation ρ , the bigger the benefit — and it can vanish in a crisis when correlations spike toward 1.

Q373. What are basis risk and gap risk in a market-risk context?

Basis risk is the risk that a hedge and the hedged item do not move perfectly together — e.g. hedging a Nifty cash portfolio with Nifty futures still leaves the cash-futures basis, or hedging a corporate bond with a government bond leaves spread risk. Gap risk is the risk of a sudden discontinuous price jump (overnight or on a news shock) that prevents trading out at intermediate levels, so a "delta-neutral" book or stop-loss is overrun — VaR built on continuous moves understates both.

Q374. What are the key limitations of VaR, and why did Basel move to Expected Shortfall under FRTB?

VaR ignores tail severity beyond the quantile, is not subadditive (penalising diversification), can be gamed, assumes stable distributions, and gives a false sense of a single number. Under the Fundamental Review of the Trading Book (FRTB), Basel replaced 99% VaR with a 97.5% Expected Shortfall measure — coherent, tail-sensitive, and calibrated to a stress period with liquidity-horizon scaling — to better capture extreme and illiquid risks exposed in 2008.

C2. Credit Risk (Q375-Q404)

Q375. What is credit risk, and what are its main forms?

Credit risk is the risk of loss arising because a borrower or counterparty fails to meet its contractual obligations. It has three broad forms: default risk (the borrower fails to pay interest or principal), downgrade/migration risk (the borrower's credit rating deteriorates, raising the required yield and lowering bond price even without default), and spread risk (the market-wide credit spread widens, repricing the exposure). All three reduce the value of a credit-exposed position.

Q376. What are the three key parameters of credit risk: PD, LGD and EAD?

PD (Probability of Default) is the likelihood the borrower defaults over a given horizon, usually one year, expressed as a percentage. LGD (Loss Given Default) is the fraction of exposure lost if default occurs, after recoveries (e.g., 45%). EAD (Exposure at Default) is the expected amount owed at the time of default, in currency (e.g., the drawn loan plus a portion of undrawn commitments). Together they drive the Expected Loss of an exposure.

Q377. Define recovery rate and its relationship to LGD.

Recovery rate is the fraction of the exposure that the lender recovers after a default, through collateral sale, restructuring or bankruptcy proceeds. It is simply the complement of LGD: $\text{Recovery Rate} = 1 - \text{LGD}$. So if LGD is 40%, the recovery rate is 60%; a senior secured loan with strong collateral has high recovery (low LGD), while subordinated unsecured debt has low recovery (high LGD).

Q378. What is Expected Loss and how is it calculated? Give a worked number.

Expected Loss (EL) is the average credit loss a lender anticipates over a horizon and is treated as a cost of doing business (covered by provisions/pricing). The formula is $\text{EL} = \text{PD} * \text{LGD} * \text{EAD}$. Example: a loan with $\text{EAD} = \text{INR } 100 \text{ crore}$, $\text{PD} = 2\%$ and $\text{LGD} = 45\%$ gives $\text{EL} = 0.02 * 0.45 * 100 = \text{INR } 0.9 \text{ crore}$ (INR 90 lakh) of expected loss per year.

Q379. What is Unexpected Loss and how does it differ from Expected Loss?

Unexpected Loss (UL) is the volatility of credit loss around the expected level, i.e., the potential for losses to exceed EL in a bad year. EL is covered by provisions and loan pricing, whereas UL must be absorbed by economic capital. For a single exposure, $\text{UL} = \text{EAD} * \sqrt{\text{PD} * (1 - \text{PD}) * \text{LGD}^2 + \text{PD} * \sigma_{\text{LGD}}^2}$; capital is sized to cover UL up to a high confidence level (e.g., 99.9%), which is the basis of credit VaR.

Q380. How do PD, LGD and EAD map to the Basel regulatory capital framework?

Under the Basel Internal Ratings-Based (IRB) approach, banks estimate PD, LGD and EAD as inputs to a supervisory risk-weight formula that computes capital for Unexpected Loss. The Foundation IRB lets banks

estimate PD only (LGD/EAD prescribed by the regulator), while Advanced IRB lets them estimate all three. Expected Loss is covered separately by provisions, and any shortfall of provisions versus EL is deducted from capital.

Q381. What is the Altman Z-score and what does it measure?

The Altman Z-score is a multivariate discriminant model (Edward Altman, 1968) that predicts the probability of corporate bankruptcy within about two years from five financial ratios. It blends profitability, leverage, liquidity, solvency and activity into a single score. A higher Z indicates greater financial health; it remains a popular quick screen for distress despite being decades old.

Q382. State the five-ratio Altman Z-score formula and its cut-offs.

For public manufacturers: $Z = 1.2 \cdot X_1 + 1.4 \cdot X_2 + 3.3 \cdot X_3 + 0.6 \cdot X_4 + 1.0 \cdot X_5$, where X_1 = working capital/total assets, X_2 = retained earnings/total assets, X_3 = EBIT/total assets, X_4 = market value of equity/total liabilities, X_5 = sales/total assets. Interpretation: $Z > 2.99$ = "safe" zone, $Z < 1.81$ = "distress" zone, and $1.81-2.99$ = "grey" zone. A score below 1.81 flags a high bankruptcy risk.

Q383. What are the Z' and Z'' variants of the Altman model?

The Z' model is for private manufacturers: it replaces market value of equity in X_4 with book value of equity and re-estimates the coefficients (safe > 2.90 , distress < 1.23). The Z'' model drops X_5 (sales/total assets) to suit non-manufacturers and is also applied to emerging-market corporates: $Z'' = 6.56 \cdot X_1 + 3.26 \cdot X_2 + 6.72 \cdot X_3 + 1.05 \cdot X_4$, with a constant of +3.25 added for the EM scoring (EMS) version to align it with US bond rating equivalents.

Q384. Walk through a quick Altman Z-score calculation.

Take a public manufacturer with $X_1 = 0.15$, $X_2 = 0.20$, $X_3 = 0.10$, $X_4 = 0.8$, $X_5 = 1.5$. Then $Z = 1.2 \cdot 0.15 + 1.4 \cdot 0.20 + 3.3 \cdot 0.10 + 0.6 \cdot 0.8 + 1.0 \cdot 1.5 = 0.18 + 0.28 + 0.33 + 0.48 + 1.5 = 2.77$. A Z of 2.77 sits in the grey zone (1.81-2.99), suggesting caution but not imminent distress.

Q385. What is the Merton structural model of default?

The Merton (1974) model treats a firm's equity as a call option on the value of its assets, struck at the face value of debt D maturing at time T . If at maturity asset value V exceeds D , equityholders repay debt and keep the residual; if $V < D$, they default and hand the firm to creditors. Default is therefore an endogenous event driven by the firm's asset value falling below its debt, and equity = $\max(V - D, 0)$.

Q386. Why is equity modelled as a call option in the Merton framework?

Equityholders have limited liability: their downside is capped at zero (they walk away in bankruptcy) but their upside in the firm's assets is unlimited, exactly the payoff of a call option on assets V with strike D . This lets you price equity with Black-Scholes and, crucially, back out the unobservable asset value V and asset volatility σ from the observable equity value and equity volatility. Debt is correspondingly a risk-free bond minus a put option (the default option) held by equityholders.

Q387. Define distance-to-default and how PD is derived in Merton.

Distance-to-default (DD) measures how many standard deviations the firm's asset value is above its default point: $DD = (\ln(V/D) + (\mu - 0.5 \cdot \sigma^2) \cdot T) / (\sigma \cdot \sqrt{T})$, where V is asset value, D the debt/default point, μ the asset drift and σ the asset volatility. The model's default probability is $PD = N(-DD)$, the normal CDF of the negative distance-to-default. A larger DD means the firm is further from the default boundary and has a lower PD.

Q388. Work through a simple distance-to-default and PD calculation.

Suppose $V = 120$, $D = 100$, $\mu = 5\%$, $\sigma = 25\%$, $T = 1$. Then $\ln(120/100) = 0.1823$, and the drift term $(0.05 - 0.5 \cdot 0.25^2) \cdot 1 = 0.0188$, so numerator = 0.2011; denominator = $0.25 \cdot \sqrt{1} = 0.25$; $DD =$

$0.2011/0.25 = 0.80$. Hence $PD = N(-0.80) = 21.2\%$. A higher asset cushion or lower volatility would push DD up and PD sharply down.

Q389. Why do structural-model PDs for investment-grade firms collapse to nearly zero?

Because $PD = N(-DD)$ and the normal distribution has very thin tails, a large distance-to-default maps to an almost vanishing probability. For a well-capitalised IG firm with DD around 5-7 standard deviations, $N(-DD)$ is on the order of $1e-7$ or smaller, far below observed IG default rates. This is a known weakness of basic Merton: real short-horizon credit spreads imply non-trivial default risk, so practitioners use jumps, stochastic volatility, or calibrate to market spreads (e.g., Moody's KMV's empirical EDF mapping) rather than the raw normal CDF.

Q390. What are reduced-form (intensity) models and how do they differ from structural models?

Reduced-form models treat default as an unpredictable event governed by a stochastic hazard rate (intensity) λ , without modelling the firm's assets or capital structure. Default arrives as the first jump of a Poisson process, so the survival probability to time T is $S(T) = \exp(-\int_0^T \lambda dt)$, and for constant intensity $PD \text{ over } T = 1 - \exp(-\lambda T)$. Unlike structural models (where default is a predictable boundary crossing), default here is a "surprise," which produces realistic non-zero short-term spreads and lets parameters be calibrated directly to market credit spreads and CDS quotes.

Q391. What is a credit rating and what scale do agencies use?

A credit rating is an independent opinion of a borrower's or instrument's creditworthiness, i.e., its relative probability of default and expected loss. S&P and Fitch use AAA (highest) down through AA, A, BBB, BB, B, CCC, CC, C to D (default), with +/- modifiers; Moody's uses Aaa, Aa, A, Baa, Ba, B, Caa, Ca, C with 1/2/3 modifiers. Ratings drive borrowing cost, investor eligibility and regulatory capital.

Q392. What distinguishes investment grade from speculative grade?

Investment grade (IG) is BBB-/Baa3 and above: relatively low default risk, eligible for many regulated/institutional portfolios. Speculative grade (also "high yield" or "junk") is BB+/Ba1 and below: meaningfully higher default risk and yields. The IG/HY boundary matters because many mandates, indices and capital rules force selling when an issuer is downgraded from BBB- to BB+ (a "fallen angel"), which can trigger forced selling and spread widening.

Q393. Who are the major credit rating agencies globally and in India?

Globally the "Big Three" are S&P Global Ratings, Moody's Investors Service and Fitch Ratings, which dominate sovereign and corporate ratings. In India, SEBI-registered domestic agencies include CRISIL (majority-owned by S&P), ICRA (associate of Moody's), CARE Ratings, India Ratings (a Fitch group company), Acuité and Brickwork. Indian issuers map to national-scale symbols such as CRISIL AAA, "[ICRA]AA", or "CARE A" with similar gradations.

Q394. What is the difference between through-the-cycle (TTC) and point-in-time (PIT) ratings?

A through-the-cycle rating reflects a borrower's creditworthiness across an entire economic cycle, deliberately stable and "stressed," so it does not jump with every macro swing; agency ratings are broadly TTC. A point-in-time rating/PD reflects current conditions at the assessment date and so rises and falls with the cycle. TTC ratings give stable capital and lending decisions; PIT estimates (used in IFRS 9 / Ind-AS 109 expected loss) are more responsive and react quickly to downturns.

Q395. What is a rating transition matrix and what does it tell you?

A rating transition (migration) matrix gives the probabilities that an issuer in a given rating moves to each other rating (or defaults) over a horizon, typically one year. Rows are the starting rating, columns the ending rating, each row summing to 100%; the diagonal (staying put) is usually the largest entry, and the last column is the PD for that grade. For example, a one-year AAA-to-default probability is essentially 0%, while

a B-rated issuer might show a 4-6% default column, so the matrix feeds migration-based credit VaR and multi-year cumulative PD estimates.

Q396. What is a credit spread and what does it compensate investors for?

A credit spread is the extra yield a risky bond pays over a comparable risk-free benchmark (e.g., a US Treasury or, in India, a G-Sec of similar tenor): $\text{Credit Spread} = \text{Corporate Yield} - \text{Risk-Free Yield}$. It compensates investors for expected default loss plus several premia: a risk premium for uncertain/unexpected losses, a liquidity premium, and tax/structural factors. As a rough first approximation, Spread is approximately $\text{PD} * \text{LGD}$, the annualised expected loss.

Q397. How is a credit spread decomposed?

A credit spread can be decomposed into: (1) an expected-loss component, approximately $\text{PD} * \text{LGD}$, the compensation for average default losses; (2) a credit-risk premium for unexpected losses and the systematic, non-diversifiable nature of default risk; (3) a liquidity premium for the difficulty of trading the bond; and (4) tax and other technical effects. Empirically the expected-loss piece explains only a minority of IG spreads, a puzzle known as the "credit spread puzzle," with risk and liquidity premia making up the rest.

Q398. What is a credit default swap (CDS)?

A CDS is a bilateral derivative that transfers the credit risk of a reference entity from a protection buyer to a protection seller. The buyer pays a periodic premium (the CDS spread, in basis points on notional) and, if a defined credit event occurs (default, failure to pay, restructuring), the seller compensates the buyer for the loss, typically $(1 - \text{Recovery}) * \text{Notional}$. It functions like insurance on a bond and lets investors hedge or take a synthetic long/short credit position without owning the underlying.

Q399. Explain the premium leg versus the protection leg of a CDS.

The premium (or fee) leg is the stream of periodic spread payments the protection buyer makes to the seller until maturity or a credit event, plus any accrued premium at default. The protection (or contingent) leg is the single payment the seller makes to the buyer if a credit event occurs, equal to $(1 - \text{Recovery Rate}) * \text{Notional}$. A CDS is fair when the present value of the premium leg equals the present value of the expected protection leg; solving that equality gives the par CDS spread.

Q400. Why is the CDS spread approximately equal to PD times LGD?

Setting the PV of the protection leg equal to the PV of the premium leg, the par spread is approximately $\text{Spread} \approx \lambda * (1 - \text{Recovery}) = \text{PD} * \text{LGD}$ per year, where λ is the default intensity. Intuitively, the annual premium must cover the annual expected loss: probability of default times loss given default. Example: if the market quotes a 5-year CDS at 300 bps (3%) and recovery is assumed at 40% (LGD 60%), the implied annual PD is roughly $0.03 / 0.60 = 5\%$.

Q401. What is credit VaR and how does it differ from market VaR?

Credit VaR is the worst credit loss over a horizon (often 1 year) at a given confidence level (e.g., 99.9%), measured as the difference between that quantile loss and the expected loss; capital is held against this unexpected loss. Unlike market VaR, the credit loss distribution is highly skewed and fat-tailed: most years see small losses near EL, but rare clustered defaults create a long right tail. Models like CreditMetrics (rating migration) and CreditRisk+ (actuarial intensity) are used to generate this distribution.

Q402. Why do default correlation and concentration matter in a credit portfolio?

Because credit losses are driven by common factors (the economy, sector, region), defaults cluster rather than occur independently, fattening the tail of the loss distribution and raising the capital needed for a given confidence level. Higher default correlation means diversification benefits shrink and unexpected loss rises. Concentration, large single-name exposures or heavy sector/geographic tilts, further amplifies tail risk,

which is why regulators impose large-exposure limits and why portfolio managers monitor name and sector concentration (e.g., via a Herfindahl index).

Q403. What are the RBI IRAC norms and how do NPAs get classified in India?

RBI's Income Recognition and Asset Classification (IRAC) norms govern how Indian banks recognise bad loans. A loan becomes a Non-Performing Asset (NPA) when interest/principal is overdue for more than 90 days. NPAs are then graded: Sub-Standard (NPA up to 12 months), Doubtful (NPA beyond 12 months), and Loss assets (uncollectible), with rising minimum provisioning at each stage (e.g., 15% for secured sub-standard, scaling up to 100% for loss assets). Accounts overdue 31-90 days are flagged as Special Mention Accounts (SMA-1/SMA-2) as an early-warning stage.

Q404. How does Expected Credit Loss under Ind-AS 109 / IFRS 9 work, and how does it relate to RBI provisioning?

Ind-AS 109 (India's version of IFRS 9) uses a forward-looking Expected Credit Loss (ECL) model with three stages: Stage 1 (performing) provides 12-month ECL; Stage 2 (significant increase in credit risk) and Stage 3 (credit-impaired/default) provide lifetime ECL. The core computation is $ECL = PD * LGD * EAD$, discounted, and incorporating macro forecasts, so provisions are recognised before an account actually turns NPA. This differs from RBI's rule-based IRAC provisioning (fixed percentages by classification); Indian banks still follow IRAC, while NBFCs report under Ind-AS ECL, and the RBI has proposed moving banks to an expected-loss framework, with the higher of the two often maintained.

C3. Risk Frameworks, Regulation & Other Risks (Q405-Q435)

Q405. What are the main categories of financial risk a bank or trading firm manages?

The standard taxonomy is: market risk (losses from moves in prices, rates, FX, equities), credit risk (a borrower or issuer defaulting), liquidity risk (inability to fund or to sell without large price concessions), operational risk (failures of people, process, systems or external events), counterparty credit risk (a trading counterparty defaulting before settlement), model risk (losses from using a wrong or misused model), and reputational risk (loss of franchise value from damaged perception). Regulators (Basel) explicitly capitalize market, credit and operational risk; the others are managed but harder to capital-charge.

Q406. What is the difference between market risk and credit risk?

Market risk is the risk of loss from adverse movements in market prices or rates (equity, interest rate, FX, commodity, volatility) on positions you hold. Credit risk is the risk that a counterparty or borrower fails to meet contractual obligations, causing loss of principal and/or interest. A bond portfolio has both: market risk if yields move and the price falls, and credit risk if the issuer defaults.

Q407. What is liquidity risk, and what are its two main types?

Liquidity risk is the risk of being unable to meet obligations or unwind positions in a timely way without incurring unacceptable losses. The two types are funding liquidity risk (you cannot raise cash to meet liabilities as they fall due) and market (asset) liquidity risk (you cannot sell an asset quickly near its fair value because the market is thin). In a crisis the two reinforce each other: fire-sales widen bid-ask spreads, which worsens both.

Q408. What is operational risk? Give examples.

Operational risk is the risk of loss from inadequate or failed internal processes, people and systems, or from external events (Basel definition); it explicitly includes legal risk but excludes strategic and reputational risk. Examples: rogue trading (Nick Leeson at Barings, Kerviel at Societe Generale), internal/external fraud, settlement and system failures, cyber-attacks, and natural disasters. Unlike market and credit risk, it is not taken deliberately to earn return.

Q409. What is model risk?

Model risk is the risk of adverse consequences from decisions based on models that are incorrect, misspecified, fed bad data, or used outside their intended scope. Examples include mispricing exotic options with a flawed volatility surface, or VaR understating tail risk because it assumes normality. It is managed through independent model validation, back-testing, and governance (e.g. the US SR 11-7 guidance).

Q410. What is reputational risk?

Reputational risk is the risk that negative perception by clients, counterparties, investors or regulators erodes a firm's revenue, funding access or franchise value. It is often a second-order consequence of another failure (a fraud, a mis-selling scandal, a data breach) and is hard to quantify or capitalize. Banks manage it through conduct standards, disclosure and crisis communication rather than a capital charge.

Q411. Briefly trace the evolution from Basel I to Basel II to Basel III.

Basel I (1988) introduced a simple risk-weighted capital framework focused almost entirely on credit risk, requiring 8% capital against risk-weighted assets. Basel II (2004) added a three-pillar structure, more risk-sensitive credit approaches (standardized and internal ratings-based), and an explicit operational risk charge. Basel III (post-2008, phased from 2013) kept the pillars but sharply raised the quality and quantity of capital, and added liquidity standards (LCR, NSFR), a leverage ratio, and capital buffers.

Q412. What are the three pillars of Basel II?

Pillar 1 is minimum capital requirements (the quantitative rules for credit, market and operational risk capital). Pillar 2 is the supervisory review process, where regulators assess a bank's internal capital adequacy (ICAAP) and can require capital above the minimum for risks not fully captured in Pillar 1. Pillar 3 is market discipline, achieved through public disclosure of risk and capital so the market can assess the bank.

Q413. What is the Capital Adequacy Ratio (CAR), and what is the minimum?

CAR (also called CRAR in India) measures a bank's capital against its risk-weighted assets: $CAR = \text{Eligible Capital} / \text{Risk-Weighted Assets}$. The Basel minimum total CAR is 8%; the RBI requires Indian banks to hold a minimum of 9%, plus buffers on top. Example: a bank with INR 90 crore of eligible capital and INR 1,000 crore of RWA has $CAR = 90/1000 = 9\%$.

Q414. What are risk-weighted assets (RWA)?

RWA scale each asset by a risk weight reflecting its credit (and other) riskiness, so safer assets need less capital. For example a sovereign exposure might carry 0% weight, a residential mortgage 35%, and an unrated corporate loan 100%: $RWA = \sum(\text{exposure}_i * \text{risk_weight}_i)$. A USD 100 loan at 100% weight is USD 100 of RWA needing USD 8 of capital at 8%; the same USD 100 in a 20% asset is only USD 20 RWA needing USD 1.60.

Q415. What is the difference between Tier 1 and Tier 2 capital?

Tier 1 is going-concern capital that absorbs losses while the bank operates: it comprises Common Equity Tier 1 (CET1 - common shares, retained earnings) plus Additional Tier 1 (AT1 - perpetual instruments like contingent convertibles). Tier 2 is gone-concern capital that absorbs losses in liquidation, such as subordinated debt and certain loan-loss reserves. Tier 1 is higher quality and more loss-absorbing than Tier 2.

Q416. What is CET1 capital and what is its minimum requirement under Basel III?

Common Equity Tier 1 is the highest-quality capital: common shares, share premium, retained earnings and disclosed reserves, net of regulatory deductions like goodwill. Under Basel III the minimum CET1 ratio is 4.5% of RWA, the minimum Tier 1 is 6%, and minimum total capital is 8% - before buffers. CET1 is the first line of loss absorption and the focus of stress tests.

Q417. What is the capital conservation buffer?

The capital conservation buffer is an extra 2.5% of CET1 above the minimum, designed to be built up in good times and drawn down in stress. With it, effective CET1 rises to 7% and total capital to 10.5%. If a bank's capital falls into the buffer, it faces automatic restrictions on dividends, buybacks and bonuses until capital is rebuilt.

Q418. What is the countercyclical capital buffer (CCyB)?

The CCyB is a discretionary buffer of 0-2.5% CET1 that national regulators switch on when they judge that excessive system-wide credit growth is building up risk, and release in downturns. Its aim is macroprudential - to lean against the credit cycle and protect the banking system, not just the individual bank. In India the RBI has the framework in place but has so far kept the CCyB at 0%.

Q419. What is the leverage ratio and why was it added in Basel III?

The leverage ratio is a non-risk-based backstop: $\text{Leverage Ratio} = \text{Tier 1 Capital} / \text{Total Exposure (on- plus off-balance-sheet)}$, with a 3% minimum. It was added because risk-weighting can be gamed or mis-calibrated (pre-2008 banks looked well-capitalized on RWA but were hugely levered), so a simple, hard cap on total leverage constrains balance-sheet growth regardless of model-assigned risk weights.

Q420. What is the Liquidity Coverage Ratio (LCR)?

The LCR ensures a bank holds enough high-quality liquid assets (HQLA) to survive a 30-day stressed cash outflow: $\text{LCR} = \text{HQLA} / \text{Total Net Cash Outflows over 30 days} \geq 100\%$. HQLA are assets that can be sold or pledged quickly with little loss (cash, central-bank reserves, high-grade government bonds). It is a short-term, acute-stress liquidity measure introduced by Basel III.

Q421. What is the Net Stable Funding Ratio (NSFR)?

The NSFR addresses structural, longer-term liquidity over a one-year horizon: $\text{NSFR} = \text{Available Stable Funding} / \text{Required Stable Funding} \geq 100\%$. Available stable funding weights liabilities and capital by how reliable they are (equity and long-term deposits score high; overnight wholesale funding scores low), while required stable funding reflects the illiquidity and tenor of assets. It discourages over-reliance on short-term wholesale funding to hold illiquid assets.

Q422. What is stress testing, and how does it differ from VaR?

Stress testing estimates losses under specific severe-but-plausible scenarios (e.g. a 30% equity crash, a 200 bps rate shock, a 2008-style event), whether historical or hypothetical. Unlike VaR, which gives a probabilistic loss threshold under normal conditions and tends to underrepresent tails, stress tests are scenario-driven and deliberately probe extreme conditions. They are central to Pillar 2 / ICAAP and to supervisory exercises like the Fed's CCAR and the RBI's macro stress tests.

Q423. What is reverse stress testing?

Reverse stress testing starts from the outcome and works backwards: it asks "what scenarios would cause the firm to fail or breach viability?" rather than "what is the loss under scenario X?" This identifies hidden vulnerabilities and concentrations that forward scenarios might miss, and informs recovery and resolution planning. Regulators increasingly require it as a complement to conventional stress tests.

Q424. What are the Basel approaches for operational risk capital?

Under Basel II banks could use the Basic Indicator Approach (capital as a fixed percentage, 15%, of average gross income), the Standardized Approach (different percentages by business line), or the Advanced Measurement Approach (AMA, internal loss models). Basel III replaces all of these with a single Standardized Measurement Approach (SMA), which combines a Business Indicator (a proxy for size from income statement items) with an Internal Loss Multiplier based on the bank's own historical operational losses.

Q425. What is counterparty credit risk (CCR) in derivatives?

CCR is the risk that the counterparty to a derivative or securities-financing trade defaults before the final settlement of cash flows, when the contract has positive value to you (so the default creates a replacement loss). Unlike a loan, the exposure is uncertain and bilateral - it changes with market moves and can be positive or negative. It is measured through metrics like Expected Exposure (EE), Potential Future Exposure (PFE) and Exposure at Default (EAD).

Q426. What are CVA and DVA?

Credit Valuation Adjustment (CVA) is the adjustment to a derivative's value reflecting the expected loss from the counterparty defaulting - effectively the market price of counterparty credit risk: roughly $CVA \sim LGD * \sum (EE(t) * PD(t))$ over future times t . Debit Valuation Adjustment (DVA) is the mirror image: the gain from your own default risk, which reduces the value of your liabilities. DVA is controversial because a firm's own credit deterioration produces an accounting gain.

Q427. What is wrong-way risk?

Wrong-way risk arises when a counterparty's credit quality is positively correlated with the exposure to it, so that exposure rises just as the counterparty becomes more likely to default. A classic example: buying credit protection on a company from a counterparty closely tied to that company. Right-way risk is the favorable opposite, where exposure falls as default risk rises.

Q428. How do netting and collateral/margin reduce counterparty exposure?

Close-out netting under an ISDA Master Agreement lets you offset positive and negative mark-to-market across all trades with a counterparty, so exposure is the net rather than gross amount. Collateral (margin) further reduces it: variation margin is posted to cover current mark-to-market moves, while initial margin covers potential future exposure over the close-out period. Together they sharply cut the loss-given-default, which is why a CSA (Credit Support Annex) governs collateral terms.

Q429. What is a central counterparty (CCP), and why did central clearing expand after 2008?

A CCP interposes itself between the two sides of a trade through novation, becoming buyer to every seller and seller to every buyer, and manages risk via multilateral netting, daily margining and a default fund/waterfall. After the 2008 crisis (where opaque bilateral OTC exposures, e.g. around AIG and Lehman, amplified contagion), the G20 mandated central clearing of standardized OTC derivatives to increase transparency and reduce systemic interconnectedness. The trade-off is concentration of risk in the CCPs themselves, making them systemically critical.

Q430. What is hedging, and why do firms hedge?

Hedging is taking an offsetting position whose value moves opposite to an existing exposure, so that gains on one roughly cancel losses on the other, reducing the variability of outcomes. Firms hedge to protect margins and cash flows from risks outside their core competence (FX, rates, commodity prices) - for example an Indian IT exporter selling USD forward to lock in the INR value of dollar receivables. The goal is to reduce risk, not to make a profit, and a perfect hedge eliminates the targeted risk entirely.

Q431. What is basis risk in hedging?

Basis risk is the residual risk that the hedge and the underlying exposure do not move in perfect lockstep, because they differ in instrument, maturity, location or quality. For instance, hedging jet fuel with crude oil futures leaves the crude-to-jet-fuel spread (the basis) unhedged. The basis is $\text{Basis} = \text{Spot price of asset hedged} - \text{Futures price used}$, and unexpected changes in it are the source of basis risk.

Q432. What is the difference between over-hedging and under-hedging?

Under-hedging means the hedge covers less than the full exposure (hedge ratio below 1), leaving part of the original risk live. Over-hedging means the hedge notional exceeds the exposure (hedge ratio above 1), so the

"hedge" itself creates a net speculative position in the opposite direction. The optimal minimum-variance hedge ratio is $h = \text{correlation} * (\sigma_{\text{spot}} / \sigma_{\text{futures}})$, and deviations from it leave residual or added risk.

Q433. What is the difference between economic capital and regulatory capital?

Regulatory capital is the minimum capital prescribed by supervisors (Basel/RBI rules) using standardized formulas, regardless of the bank's own view. Economic capital is the bank's internal estimate of the capital needed to absorb unexpected losses to a chosen confidence level and horizon (e.g. 99.9% over one year), tailored to its actual risk profile and used for risk-based pricing and performance (RAROC). They can diverge: a bank may hold economic capital above the regulatory floor where it judges its risks to be greater than the standardized rules assume.

Q434. What is risk appetite, and how do limits operationalize it?

Risk appetite is the board-approved statement of the amount and type of risk a firm is willing to take to meet its objectives, expressed through a risk appetite framework (RAF). It cascades into concrete, monitored limits - VaR limits, notional or position limits, stop-losses, concentration limits, stress-loss limits - that traders and desks must operate within. Limits convert a high-level appetite into enforceable controls, with escalation and breach procedures when they are exceeded.

Q435. How are risk and capital oversight split between the RBI and SEBI in India?

The RBI is the banking regulator: it sets capital adequacy (CRAR minimum 9% plus buffers), implements Basel III, prescribes liquidity (LCR/NSFR), runs supervisory stress tests, and oversees bank credit, market and operational risk. SEBI regulates securities markets - exchanges, brokers, mutual funds, clearing corporations - and governs market risk controls such as margining, position limits, circuit breakers and disclosure. So a bank's solvency sits with the RBI, while conduct and risk in the capital markets (including F&O margins and the clearing corporations that act as CCPs) sit with SEBI.

D. Financial Analyst / FP&A

D1. Budgeting, Forecasting & Variance Analysis (Q436-Q465)

Q436. What is FP&A, and what is the purpose of a budget?

FP&A (Financial Planning & Analysis) is the corporate finance function that builds budgets and forecasts, analyzes performance against plan, and advises management on resource allocation. A budget is a quantified plan for a period (usually a fiscal year) that sets revenue and cost targets, allocates resources, and creates a benchmark to measure actual results against. Its purposes are planning, coordination across departments, communication of targets, motivation, and control (variance analysis).

Q437. Walk through the typical annual budgeting cycle.

A typical cycle: (1) management sets high-level assumptions and targets (growth, margins, headcount), (2) departments build bottom-up budgets, (3) FP&A consolidates into a master budget, (4) iterate and negotiate to close gaps to targets, (5) board/CFO approval, (6) load the approved budget into the GL/planning system, then (7) monitor actuals vs budget monthly and reforecast. In India the fiscal year runs Apr-Mar, so the cycle typically finalizes in Jan-Mar; US/global firms commonly use a Jan-Dec calendar.

Q438. What is incremental budgeting and what is its main drawback?

Incremental budgeting starts from the prior year's actuals or budget and adjusts by a percentage or fixed amount for the new period (e.g., last year's INR 100 lakh spend +5% = INR 105 lakh). It is fast and simple but tends to perpetuate past inefficiencies ("budget creep") because existing spend is never re-justified, and it can encourage "use it or lose it" year-end spending.

Q439. What is zero-based budgeting (ZBB) and when is it used?

In ZBB every line item starts from a "zero base" and must be justified from scratch each cycle, rather than carrying forward prior amounts. It forces managers to defend every cost against its value, which can strip out legacy waste, but it is time-consuming and analytically heavy. It is often deployed for cost-transformation programs (e.g., consumer-goods firms cutting SG&A) rather than every year.

Q440. What is activity-based budgeting (ABB)?

ABB budgets costs based on the activities that drive them and their expected volumes, rather than on departmental line items. You identify cost drivers (e.g., number of purchase orders, machine setups, customer calls), estimate the activity volume, then multiply by the cost per activity. Budget = sum over activities of (expected driver volume * cost per driver unit). It links spending to operational demand and is the budgeting counterpart of activity-based costing.

Q441. What is a rolling (continuous) budget and how does it differ from a static annual budget?

A rolling budget always maintains a fixed horizon (e.g., 12 or 18 months) by adding a new period as each one elapses, so a month-end re-plan continuously extends the view. A static annual budget is set once and remains the fixed benchmark for the whole year. Rolling budgets keep the plan current and reduce the "January is detailed, December is stale" problem, at the cost of more frequent planning effort.

Q442. Compare top-down and bottom-up budgeting.

Top-down: senior management sets targets and pushes them down to departments; it is fast and strategically aligned but can be unrealistic and demotivating if disconnected from operations. Bottom-up: departments build estimates from operational detail and roll them up; it improves ownership and accuracy but risks padding and may not hit corporate targets. Most firms use a hybrid ("goal-setting top-down, build bottom-up, then reconcile").

Q443. Distinguish the operating budget, capital budget, and cash budget.

The operating budget covers day-to-day revenues and expenses (sales, COGS, opex) and feeds the budgeted P&L. The capital budget covers long-term asset investments (plant, equipment, IT) evaluated via NPV/IRR/payback. The cash budget projects cash inflows and outflows and timing to ensure liquidity; it is distinct from the P&L because of accruals, capex timing, and working-capital changes.

Q444. What is the master budget and what are its main components?

The master budget is the consolidated, integrated set of all sub-budgets producing the budgeted financial statements. It starts with the sales budget, then production/purchases, direct materials, direct labour, overhead, and SG&A budgets (the operating budgets), plus the capital and cash budgets, culminating in the budgeted income statement, balance sheet, and cash flow statement. The sales forecast is the keystone because nearly everything else depends on volume.

Q445. Define forecast, budget, and actual, and how they relate.

The budget is the fixed target approved at the start of the period; the forecast is the current best estimate of how the period will actually end (updated through the year); the actual is what was realized. FP&A reports all three: actuals vs budget show performance against the original plan, while forecast vs budget signals whether the year is tracking to land on, above, or below target.

Q446. What is run-rate, and how do you annualize it?

Run-rate annualizes recent performance as if it continued unchanged. Simple form: Annualized = latest period figure * number of periods per year (e.g., monthly revenue of INR 50 lakh => INR 6 crore run-rate). A smoother version uses a trailing average, e.g., Annualized = (sum of last 3 months / 3) * 12. It is a quick sanity check but ignores seasonality and one-offs, so it can mislead for cyclical businesses.

Q447. Describe trend-based forecasting and its limitation.

Trend forecasting extrapolates the historical growth pattern forward, e.g., next period = current * (1 + historical growth rate), or fitting a line/CAGR through past data. $CAGR = (Ending / Beginning)^{(1/n)} - 1$. It is simple and works when the past is a good guide to the future, but it breaks at inflection points (new products, recessions, regime changes) because it assumes the trend persists.

Q448. How is regression used in forecasting?

Regression fits a statistical relationship between a dependent variable (e.g., sales) and one or more drivers, e.g., $y = a + bx$, where b is the slope (sales per unit of driver) and a the intercept. For example, $Sales = 200 + 1.5 \text{MarketingSpend}$ predicts INR 1.5 of sales per INR 1 of marketing. Quality is judged by R-squared and the significance of coefficients; the risk is over-fitting and assuming correlation implies a stable causal driver.

Q449. What is driver-based forecasting and why is it preferred for planning?

Driver-based forecasting models outputs as a function of operational drivers rather than as a single top-line growth number, e.g., Revenue = units * average price, or = traffic * conversion rate * average order value. It makes assumptions explicit and lets managers test "what if conversion improves 1 point," improving transparency and scenario flexibility versus a blunt percentage-growth forecast.

Q450. What does "favorable" vs "unfavorable" mean in variance analysis, and what is the sign convention?

A variance is favorable (F) when it increases profit versus plan, and unfavorable (U) when it decreases profit. For revenue, Actual > Budget is favorable; for costs, Actual > Budget is unfavorable (you spent more). A common convention reports cost variance as Budget - Actual so a positive number is favorable, while revenue variance is Actual - Budget so positive is favorable; always state direction explicitly to avoid sign confusion.

Q451. Define a flexible (flex) budget and explain why it is useful.

A flexible budget recalculates budgeted revenues and costs at the actual activity (volume) level, using the original per-unit standards and budgeted fixed costs. For example, if budgeted variable cost was INR 100/unit and actual volume was 1,200 units, the flex budget variable cost is INR 1,20,000 regardless of the 1,000-unit static budget. It separates volume effects from price/efficiency effects so you compare like-for-like at the achieved volume.

Q452. How do the flex-budget variance and the sales-volume variance decompose the total static-budget variance?

Total static-budget variance = flexible-budget variance + sales-volume variance. The sales-volume variance = (actual units - budgeted units) * budgeted contribution margin per unit, isolating the effect of selling more/fewer units. The flexible-budget variance = actual results - flexible budget (at actual volume), isolating price and efficiency effects. Example: static op income variance of INR 50,000 might split into INR 30,000 volume + INR 20,000 flex (price/cost) effects.

Q453. Decompose a revenue variance into price and volume variances with a worked example.

Price variance = (Actual price - Budget price) * Actual quantity; Volume variance = (Actual quantity - Budget quantity) * Budget price. Example: budget 1,000 units at INR 100; actual 1,100 units at INR 105. Price variance = $(105-100)1,100 = +\text{INR } 5,500 \text{ F}$; Volume variance = $(1,100-1,000)100 = +\text{INR } 10,000 \text{ F}$. Total revenue variance = $1,100 \times 105 - 1,000 \times 100 = 1,15,500 - 1,00,000 = +\text{INR } 15,500 \text{ F} = 5,500 + 10,000$. (The cross term is conventionally absorbed into the price variance by using actual quantity.)

Q454. What is mix variance, and how does total variance split into price, volume (quantity), and mix?

With multiple products, total variance = price variance + pure quantity (volume) variance + mix variance. Mix variance captures the effect of selling a different proportion of products than budgeted, at constant total units: Mix variance = sum over products of [(actual mix % - budget mix %) * total actual units * budgeted contribution per unit]. The quantity variance uses budgeted mix on the change in total units, so mix isolates the shift toward higher- or lower-margin products independent of total volume.

Q455. Work a price/volume/mix decomposition for a two-product example.

Budget: A 600u at CM 10, B 400u at CM 20 (total 1,000u; budget mix 60/40; avg CM = $(600 \times 10 + 400 \times 20) / 1000 = 14$). Actual total 1,100u, mix 50/50 (550 A, 550 B), CMs as budget. Quantity variance = $(1,100-1,000)14 = +1,400 \text{ F}$. Mix variance = $\text{sum}[(\text{actual mix}\% - \text{budget mix}\%) \times \text{actual total units} \times \text{budget CM}] = (0.50-0.60)1,100 \times 10 + (0.50-0.40)1,100 \times 20 = -1,100 + 2,200 = +1,100 \text{ F}$. Total volume+mix = 2,500 F (matches: actual CM $550 \times 10 + 550 \times 20 = 16,500$ vs budget 14,000). Price variance would add on top if CMs changed.

Q456. Explain the direct material price and usage (quantity) variances.

Material price variance = (Actual price - Standard price) * Actual quantity purchased; it measures buying efficiency. Material usage variance = (Actual quantity used - Standard quantity for actual output) * Standard price; it measures consumption efficiency. Example: standard 2 kg at INR 50/kg for output; actual 2.2 kg at INR 48/kg for one unit. Price = $(48-50)2.2 = -4.4 \text{ (F)}$; Usage = $(2.2-2.0)50 = +10 \text{ (U)}$.

Q457. Explain the direct labour rate and efficiency variances.

Labour rate variance = (Actual rate - Standard rate) * Actual hours; it reflects wage differences. Labour efficiency variance = (Actual hours - Standard hours for actual output) * Standard rate; it reflects productivity. Example: standard 5 hrs at INR 200/hr; actual 5.4 hrs at INR 210/hr. Rate = $(210-200)5.4 = +54 \text{ (U)}$; Efficiency = $(5.4-5.0)200 = +80 \text{ (U)}$.

Q458. What are the overhead spending and volume variances?

For variable overhead: spending variance = actual VOH - (actual hours * standard VOH rate); efficiency variance = (actual hours - standard hours) * standard VOH rate. For fixed overhead: spending (budget)

variance = actual fixed OH - budgeted fixed OH; volume variance = budgeted fixed OH - (standard hours allowed * fixed OH rate), arising because fixed costs are absorbed on a per-unit basis and actual volume differs from the level used to set the rate. The fixed-OH volume variance is purely an absorption artifact, not a cash over/underspend.

Q459. How do you set a materiality threshold for investigating variances?

Firms set a rule combining absolute size and percentage, e.g., investigate any variance over INR 5 lakh OR over 10% of the line, whichever is breached, to avoid chasing tiny noise or ignoring large small-percentage items. Thresholds may be tighter for controllable/high-value lines and looser for volatile ones. The principle is management by exception: focus analyst time where the variance is both material and actionable.

Q460. What is a rolling reforecast and how often is it done?

A rolling reforecast updates the full-year (or rolling-horizon) estimate using actuals to date plus a revised view of the remaining periods, typically monthly or quarterly. A common structure: Full-year forecast = actuals YTD + reforecast of remaining months. It keeps guidance and resource decisions current and is the practical bridge between the static budget and reality, often paired with run-rate checks for the back half.

Q461. What KPIs does FP&A typically track in budgeting and performance reporting?

Common KPIs: revenue growth, gross and operating margin, EBITDA, budget vs actual variance %, forecast accuracy (e.g., absolute % error), cost per unit, headcount and cost per FTE, working-capital metrics (DSO, DPO, DIO), cash burn/runway (for startups), and operational drivers like units, ARPU, churn, or utilization. Forecast accuracy = $1 - |\text{Actual} - \text{Forecast}| / \text{Actual}$, tracked over time to gauge planning quality.

Q462. What is budget padding (sandbagging) and how do you detect and mitigate it?

Padding/sandbagging is deliberately overstating costs or understating revenue so targets are easy to beat. Detection: persistent favorable variances in the same departments, hockey-stick spending near year-end, and benchmarking against drivers/peers. Mitigation: zero-based or driver-based reviews, holding managers to forecast accuracy (not just hitting target), stretch goals, and decoupling the budget used for incentives from the forecast used for planning.

Q463. Describe the FP&A monthly close and reporting loop.

After accounting closes the books (typically business days 3-7), FP&A pulls actuals, computes variances vs budget and prior forecast, investigates material items with business owners, prepares the management/board reporting pack with commentary, and updates the rolling reforecast and KPIs. The loop is: close -> analyze variances -> explain -> reforecast -> report -> act, repeating each month so the organization steers off current information rather than the stale annual budget.

Q464. What is scenario and sensitivity analysis in planning, and how do they differ?

Sensitivity analysis flexes one variable at a time to see its impact on an output (e.g., "a 1% price cut reduces EBITDA by INR 20 lakh"), revealing which assumptions matter most. Scenario analysis changes a coherent set of assumptions together to describe a whole state of the world (base, bull, bear), each with internally consistent revenue, cost, and cash figures. Sensitivity isolates drivers; scenarios stress-test the plan against plausible futures and inform contingency planning.

Q465. A static budget showed operating income of INR 1,00,000 (1,000 units at INR 1,000, variable cost INR 600, fixed INR 3,00,000); actuals were 1,100 units, price INR 980, variable cost INR 620/unit, fixed INR 3,10,000. Reconcile via the flexible budget.

Budget CM/unit = $1,000 - 600 = 400$. Sales-volume variance = $(1,100 - 1,000)400 = +40,000$ F. Flexible budget (at 1,100u, budget prices): CM = $1,100 \times 400 = 4,40,000$, less fixed 3,00,000 = 1,40,000 op income. Actual: revenue $1,100 \times 980 = 10,78,000$; variable $1,100 \times 620 = 6,82,000$; CM = 3,96,000; less fixed 3,10,000 =

86,000. Flex-budget variance = $86,000 - 1,40,000 = -54,000$ U (price -22,000, variable cost -22,000, fixed -10,000). Total vs static (1,00,000) = $86,000 - 1,00,000 = -14,000$ U = +40,000 volume - 54,000 flex.

D2. 3-Statement Financial Modeling (Q466-Q490)

Q466. What are the three financial statements and what does each one tell you?

The income statement shows profitability over a period (revenue minus expenses down to net income). The balance sheet is a snapshot at a point in time of what the company owns and owes (Assets = Liabilities + Equity). The cash flow statement reconciles net income to actual cash movement across operating (CFO), investing (CFI), and financing (CFF) activities, explaining the change in the cash balance between two balance sheet dates. Under Ind-AS the structure mirrors IFRS, so the same three-statement logic applies for Indian listed companies.

Q467. At a high level, how do the three statements link together?

Net income from the bottom of the income statement flows to the top of the cash flow statement (as the starting point of CFO) and into retained earnings on the balance sheet. CapEx sits in CFI, debt and equity issuance/repayment and dividends sit in CFF, and changes in working capital sit in CFO. The single ending cash figure produced by the cash flow statement (CFO + CFI + CFF + beginning cash) becomes the cash line on the balance sheet, and depreciation links the IS expense, the CF add-back, and the reduction of PP&E on the balance sheet.

Q468. Where does net income show up across all three statements?

Net income is the last line of the income statement, the first line of the cash flow statement (CFO), and it increases retained earnings within equity on the balance sheet (Ending RE = Beginning RE + Net Income - Dividends). This is why a single net income number ties the whole model together. If you change one assumption that moves net income, all three statements should move consistently.

Q469. Walk me through how a \$10 increase in depreciation flows through all three statements (assume a 25% tax rate).

Income statement: pre-tax income falls by 10, so taxes fall by 2.50 ($10 * 25\%$) and net income falls by 7.50. Cash flow statement: net income is down 7.50 but you add back the full 10 of non-cash depreciation, so cash from operations rises by 2.50; total cash increases by 2.50. Balance sheet: cash up 2.50 on the asset side, PP&E down 10 (net assets down 7.50), while retained earnings fall by 7.50 on the equity side, so the balance sheet stays balanced. The net 2.50 cash gain is exactly the tax shield from depreciation.

Q470. Why does extra depreciation actually increase a company's cash, even though it is an expense?

Depreciation is a non-cash expense, so the only real cash effect is that it lowers taxable income and therefore the cash taxes paid. With a 25% tax rate, 10 of extra depreciation saves 2.50 of taxes ($10 * 0.25$), and that 2.50 is the "depreciation tax shield." Without taxes, extra depreciation would have zero cash effect because you subtract it on the IS and add the identical amount back on the CF.

Q471. Walk me through a \$10 increase in inventory across the three statements (purchased with cash, no tax effect).

Income statement: no immediate impact, because buying inventory is not an expense until it is sold (it sits in current assets). Cash flow statement: inventory is a use of cash, so CFO falls by 10 and total cash falls by 10. Balance sheet: inventory (asset) up 10, cash (asset) down 10, so total assets are unchanged and the statement still balances with no change to liabilities or equity.

Q472. Walk me through \$10 of CapEx across all three statements.

Income statement: no immediate impact, since CapEx is capitalized, not expensed (only the future depreciation hits the IS). Cash flow statement: CapEx is an outflow in CFI of 10, so total cash falls by 10. Balance sheet: cash (asset) down 10 and gross PP&E (asset) up 10, so total assets are unchanged and the balance sheet balances. In later periods the depreciation of that asset then flows through the IS.

Q473. Walk me through what happens when a company raises \$100 of debt.

Income statement: no immediate impact at the moment of issuance (interest expense only accrues over time). Cash flow statement: a financing inflow (CFF) of 100, so cash rises by 100. Balance sheet: cash (asset) up 100 and debt (liability) up 100, so Assets = Liabilities + Equity still holds. In subsequent periods, interest expense reduces net income and the debt schedule tracks principal repayments.

Q474. Walk me through a \$10 inventory write-down across the three statements (25% tax rate).

Income statement: the 10 write-down is an expense (often in COGS or as an impairment), so pre-tax income falls 10, taxes fall 2.50, and net income falls 7.50. Cash flow statement: net income down 7.50 but the write-down is non-cash, so add back 10; CFO and total cash rise by 2.50 (the tax shield). Balance sheet: inventory down 10, cash up 2.50 (net assets down 7.50), retained earnings down 7.50 - balanced. The mechanics are identical to extra depreciation because both are non-cash charges.

Q475. How does the balance sheet "balance," and what do we mean by the plug?

The balance sheet balances when total assets equal total liabilities plus equity, and in a working model this happens automatically only if every line item is correctly linked to the IS and CF. The "plug" is the line that absorbs the residual flow of cash so the model stays internally consistent - typically the cash balance (and a revolver/short-term debt line when cash would otherwise go negative). If your statements are all correctly linked, the balance sheet balances by construction; you never manually force it.

Q476. What is the revolver (cash-sweep) mechanism and why is it used in a model?

The revolver is a short-term line of credit that acts as the financing plug: if the model projects a cash shortfall (cash below a minimum required balance), the revolver is automatically drawn to cover it; if there is excess cash above the minimum, the model uses it to pay the revolver down (the "cash sweep"). This keeps cash from going negative and keeps the balance sheet balanced under any operating scenario. The revolver balance sits in CFF and on the liability side, and its interest feeds back into the income statement.

Q477. Explain the circular reference created by interest on debt and why it arises.

Interest expense depends on the debt balance, but interest reduces net income, which reduces cash, which (via the cash sweep) changes how much revolver is drawn or repaid - which changes the debt balance and therefore interest again. This loop (interest -> net income -> cash -> debt -> interest) is a genuine circularity because each variable depends on the others within the same period. It is most acute when you calculate interest on the average of beginning and ending debt balances.

Q478. How do you resolve a circular reference from interest in a model?

Three common approaches: (1) enable iterative calculation in Excel (File > Options > Formulas > Enable iterative calculation) so the spreadsheet converges on a solution; (2) use a "circularity switch" - a toggle cell that, when flipped, replaces the live interest formula with a hardcoded zero or value to break the loop if a #REF or #DIV/o error corrupts the model, then flip it back; or (3) compute interest on the beginning balance only, which removes the circularity entirely at the cost of slight accuracy. Copy-paste-special-values of the interest figure is a quick manual break but it freezes the link.

Q479. Why is a circularity switch considered best practice when using iterative calc?

Iterative calculation can "blow up" - if any single cell in the loop briefly returns an error (#DIV/o, #REF), the error propagates around the circle and the whole model fills with errors that iteration cannot clear on its own. A circularity switch lets you set interest to 0 (breaking the loop), let the model recalculate cleanly, then

turn the switch back on to restore the live calculation. It is essentially a circuit breaker that makes the circular model recoverable.

Q480. How is a working capital schedule built using days-based drivers?

You forecast the major working capital items from turnover ratios expressed in days: DSO (Days Sales Outstanding) = Accounts Receivable / Revenue * 365 drives receivables; DIO (Days Inventory Outstanding) = Inventory / COGS * 365 drives inventory; DPO (Days Payable Outstanding) = Accounts Payable / COGS * 365 drives payables. You set forecast assumptions for each days metric, then back into the balance, e.g. forecast AR = DSO / 365 * forecast Revenue. The period-over-period change in each item then flows into CFO.

Q481. How does a change in working capital flow into the cash flow statement?

An increase in an operating asset (receivables, inventory) is a use of cash and reduces CFO; a decrease is a source of cash. An increase in an operating liability (payables, accrued expenses) is a source of cash and increases CFO; a decrease is a use. The cash flow change equals the negative of the change in net working capital (excluding cash and debt). For example, if receivables rise 10 and payables rise 4, the net WC use of cash is 6.

Q482. Walk me through a debt schedule and how interest is calculated.

A debt schedule tracks, for each tranche: beginning balance, mandatory amortization/repayments, optional prepayments or drawdowns, and ending balance (Ending = Beginning - Repayments + Drawdowns). Interest expense is typically calculated as the interest rate times the average of beginning and ending balance, or sometimes on the beginning balance to avoid circularity. The ending balances feed the balance sheet liabilities, the interest feeds the income statement, and principal flows feed CFF. The cash sweep determines optional revolver paydown from excess free cash.

Q483. How is a PP&E and depreciation schedule constructed?

Roll forward the net PP&E balance: Ending PP&E = Beginning PP&E + CapEx - Depreciation - Disposals. CapEx is often forecast as a percentage of revenue, and depreciation as a percentage of beginning PP&E or of revenue, or built bottom-up from a depreciation waterfall on each asset vintage. The ending PP&E goes to the balance sheet, CapEx is an outflow in CFI, and depreciation is a non-cash add-back in CFO and an expense on the IS. This single schedule keeps the three depreciation/CapEx touchpoints consistent.

Q484. What are deferred taxes and how do they arise in a model?

Deferred taxes arise from timing differences between book (financial-reporting) income and taxable income - most commonly because tax depreciation (e.g. accelerated/written-down value under the Income Tax Act in India, or MACRS in the US) differs from straight-line book depreciation. When tax depreciation exceeds book depreciation, the company pays less cash tax now and books a Deferred Tax Liability (DTL), which reverses later. The book tax expense stays on the IS, while the DTL change reconciles book vs cash taxes.

Q485. How do deferred taxes flow through the cash flow statement and balance sheet?

The income statement shows the book tax provision, but cash taxes actually paid differ; the difference is the change in deferred tax balances. An increase in a Deferred Tax Liability is a non-cash add-back in CFO (you expensed more tax on the books than you paid in cash), and it grows the DTL on the balance sheet liability side. A Deferred Tax Asset works in reverse. Modeling DTL/DTA changes is how you reconcile book tax expense to actual cash tax outflow.

Q486. What are the core modeling best practices around inputs and formulas?

No hardcoded numbers buried inside formulas - every assumption should live in a clearly labeled input cell so it can be flexed and audited. Convention is to color inputs/assumptions blue and formulas/calculations black (and links to other sheets often green), so anyone can instantly see what is an assumption versus a

calculation. Build error checks (e.g. a balance check row showing Assets - Liabilities - Equity = 0), avoid circular links except where intended, and keep one consistent sign convention. Clean structure and consistency matter as much as the math.

Q487. What is the balance check and how should it be implemented?

The balance check is a dedicated row that computes Total Assets minus (Total Liabilities + Total Equity) for every period; it should equal exactly 0 in all columns. You usually highlight it (e.g. conditional formatting that flags any non-zero, or a single master cell that sums the absolute checks). If it shows any non-zero value, the model is internally inconsistent and you must fix the linkage before trusting any output. It is the single most important integrity check in a three-statement model.

Q488. What are the most common reasons a model does not balance?

Typical culprits: net income not flowing to both retained earnings and the top of the cash flow; a sign error on a cash flow line (treating a use of cash as a source); CapEx or depreciation linked in only two of the three statements; the change in a balance sheet item not reflected in CFO; dividends deducted from retained earnings but not from CFF (or vice versa); and the ending cash from the CF not being linked into the balance sheet cash line. A disciplined approach is to isolate the period where the check first breaks and trace each balance sheet item's change to its cash flow counterpart.

Q489. What is driver-based forecasting and why is it preferred over straight-line growth?

Driver-based forecasting builds line items from underlying operational drivers (units * price, customers * ARPU, capacity * utilization, store count * sales per store) rather than just applying a flat growth rate to last year's number. It makes assumptions transparent and testable, lets you run sensitivities on the real levers, and ties the financials to business reality. For example, revenue = number of stores * average revenue per store, so you can flex each driver independently for scenario analysis.

Q490. At a high level, how do you model an acquisition or a share buyback in a three-statement model?

For an acquisition, the cash (or debt) used is an outflow in CFI/CFF, the target's assets and liabilities are added to the balance sheet, any premium over net assets creates goodwill (an asset), and you consolidate the target's IS going forward; new acquisition debt adds interest expense. For a buyback, cash leaves via CFF (a financing outflow), treasury stock/equity is reduced on the balance sheet, share count falls (boosting EPS), and there is no income statement impact except the lower share count in the EPS denominator. Both must keep the balance check at zero - cash down on the asset side matches the offsetting reduction or increase in equity/liabilities.

D3. Cost Accounting, Unit Economics & Business Partnering (Q491-Q515)

Q491. What is the difference between fixed, variable and semi-variable costs?

Fixed costs stay constant in total within a relevant range regardless of output (rent, salaries, depreciation), so per-unit fixed cost falls as volume rises. Variable costs change in total in proportion to output (raw materials, packaging, sales commissions) but stay constant per unit. Semi-variable (mixed) costs have both a fixed base and a variable component, e.g. an electricity bill with a fixed connection charge plus a per-kWh usage charge, or a mobile plan with a base rental plus per-GB overage.

Q492. How do you split a semi-variable cost into its fixed and variable parts?

The simplest method is high-low: variable cost per unit = (cost at highest activity - cost at lowest activity) / (highest units - lowest units), then fixed cost = total cost - (variable per unit * units) at either point. Example: at 10,000 units cost is 70,000 and at 4,000 units cost is 40,000, so variable per unit = (70,000 - 40,000) / (10,000 - 4,000) = 5, and fixed = 70,000 - 5 * 10,000 = 20,000. Regression analysis is the more accurate alternative because it uses all data points, not just two extremes.

Q493. What is the difference between direct and indirect costs?

Direct costs can be traced economically to a single cost object (the wood in a chair, the wages of the worker building it). Indirect costs (overheads) are shared across multiple cost objects and must be allocated, such as factory rent, the plant manager's salary, or electricity. The same cost can be direct or indirect depending on the chosen cost object: a supervisor's salary is indirect to a single unit but direct to the department they run.

Q494. What is the difference between product costs and period costs?

Product (inventoriable) costs are incurred to make the product (direct materials, direct labour, manufacturing overhead) and sit on the balance sheet as inventory until the goods are sold, when they flow to COGS. Period costs are expensed in the period incurred regardless of sales, and include selling, general and administrative (SG&A) costs like marketing, office rent and the CFO's salary. Under both Ind-AS 2 and IAS 2, inventory is valued at production cost; selling and admin costs are excluded from inventory and expensed as period costs.

Q495. Define contribution margin and the contribution margin ratio.

Contribution margin per unit = selling price - variable cost per unit; it is the amount each unit contributes toward covering fixed costs and then profit. The CM ratio = CM per unit / price = total CM / total sales, expressed as a percentage. Example: price 100, variable cost 60, so CM = 40 per unit and CM ratio = 40/100 = 40%, meaning 40 paise of every rupee of sales is available to cover fixed costs and profit.

Q496. How do you compute the break-even point in units and in revenue? Give a worked example.

BEP (units) = fixed costs / CM per unit; BEP (revenue) = fixed costs / CM ratio. Example: fixed costs 200,000, price 100, variable cost 60, so CM = 40 and CM ratio = 40%. BEP units = 200,000/40 = 5,000 units, and BEP revenue = 200,000/0.40 = 500,000 (which equals 5,000 * 100, confirming consistency).

Q497. How do you find the volume needed to hit a target profit?

Treat the desired profit like an extra fixed cost: target units = (fixed costs + target profit) / CM per unit. Using fixed costs 200,000, CM 40 and a target profit of 100,000, required units = (200,000+100,000)/40 = 7,500 units. For an after-tax target, gross it up first: pre-tax profit = after-tax profit / (1 - tax rate).

Q498. What is the margin of safety and why does it matter?

Margin of safety = actual (or budgeted) sales - break-even sales, often expressed as a percentage: MoS% = (actual sales - BEP sales) / actual sales. It measures how far sales can fall before the firm hits a loss, so it is a cushion against demand risk. Example: if BEP is 5,000 units and you sell 8,000, MoS = 3,000 units or 3,000/8,000 = 37.5%.

Q499. What is operating leverage and how is the degree of operating leverage (DOL) calculated?

Operating leverage measures how sensitive operating income is to a change in sales, driven by the proportion of fixed costs in the cost structure. DOL = contribution margin / operating income (EBIT); a DOL of 4 means a 10% rise in sales produces a roughly 40% rise in operating income. Example: CM 400,000 and operating income 100,000 gives DOL = 4.

Q500. Why do high fixed-cost firms have more volatile profits?

Because fixed costs do not move with volume, a larger fixed-cost base means a higher DOL, so each percentage change in sales is magnified into a larger percentage change in profit, in both directions. Airlines, telecom and software firms have high fixed costs and low variable costs, so they swing sharply between large losses below break-even and high margins above it. The trade-off is operational: high operating leverage rewards growth but punishes downturns.

Q501. What is the difference between variable (marginal) costing and absorption costing?

Under variable (marginal) costing, only variable manufacturing costs are inventoried and all fixed manufacturing overhead is expensed as a period cost. Under absorption (full) costing, fixed manufacturing overhead is allocated to each unit and stays in inventory until sold. Absorption costing is required for external reporting under Ind-AS 2 / IAS 2 / US GAAP, while variable costing is used internally for CVP and decision-making because it isolates contribution margin.

Q502. How can the choice between absorption and variable costing change reported profit in a period?

The difference equals the fixed manufacturing overhead trapped in or released from inventory: when production exceeds sales, absorption costing defers some fixed overhead into ending inventory, so it reports higher profit than variable costing. When sales exceed production, the reverse happens and absorption shows lower profit. Example: 1,000 units produced, 800 sold, fixed overhead 50 per unit; absorption defers $200 \times 50 = 10,000$ of fixed cost into inventory, so absorption profit is 10,000 higher than variable that period. This is also why absorption costing can incentivise overproduction to inflate short-term profit.

Q503. What is activity-based costing (ABC) and how does it differ from traditional overhead allocation?

ABC assigns overhead to products based on the activities that actually drive cost, using multiple cost pools and cost drivers, rather than spreading all overhead on one volume measure like direct labour hours. It gives more accurate product costs when overhead is large and products consume resources unevenly, exposing that low-volume, complex products are often under-costed by traditional methods. Steps: identify activities, pool their costs, pick a cost driver for each pool, compute a driver rate, then assign cost to products by their driver usage.

Q504. What is a cost driver? Give examples.

A cost driver is the factor that causes the cost of an activity to change, used as the allocation base in ABC. Examples: number of machine setups drives setup cost, number of purchase orders drives procurement cost, number of inspections drives quality cost, and number of customer support tickets drives support cost. Choosing a driver that genuinely correlates with resource consumption is what makes ABC more accurate than a single volume-based rate.

Q505. What is the difference between standard costing and actual costing, and what is a variance?

Standard costing records inventory and production at predetermined (standard) costs set from engineering estimates and budgets, while actual costing uses costs actually incurred. The variance is the difference between standard and actual, split into price/rate variances (paid more or less per unit of input) and quantity/efficiency variances (used more or less input than standard). Variance analysis turns the gap into management signals; for example an unfavourable material price variance points to procurement, while an unfavourable efficiency variance points to the shop floor.

Q506. What is cost-volume-profit (CVP) analysis and what assumptions underlie it?

CVP analysis models how profit changes with volume using the relationship $\text{Profit} = (\text{price} - \text{variable cost}) \times \text{units} - \text{fixed costs}$, and underpins break-even, target-profit and margin-of-safety calculations. Key assumptions: selling price and variable cost per unit are constant, total fixed costs are constant within the relevant range, costs split cleanly into fixed and variable, and product mix is stable for multi-product firms. These assumptions make CVP a short-run planning tool rather than a precise long-run forecast.

Q507. What is a relevant cost for decision-making, and what makes a cost a sunk cost?

A relevant cost is a future cash flow that differs between the alternatives being considered; only relevant costs and revenues belong in a decision. A sunk cost is a cost already incurred that cannot be changed by any future choice, so it is never relevant, no matter how large. For example, money already spent on R&D for a product is sunk and should be ignored when deciding whether to launch; only incremental future costs and revenues matter.

Q508. How do you make a make-vs-buy (outsourcing) decision on a relevant-cost basis?

Compare the incremental cost to buy from a supplier against the costs you would actually avoid by not making it in-house, which are the variable production costs plus any avoidable (escapable) fixed costs. Unavoidable fixed costs like factory rent that continues regardless are irrelevant. Example: making costs 50 variable + 20 fixed per unit but only 8 of that fixed is avoidable; relevant make cost is 58, so buying at 55 is cheaper, but at 60 it is not. Also weigh qualitative factors: quality control, supply reliability and any opportunity use of freed capacity.

Q509. How should a company evaluate a one-off special order priced below normal price?

If there is spare capacity, accept the order when the price exceeds the incremental (mostly variable) cost per unit, because fixed costs are already covered by regular business and any contribution is extra profit. Example: normal price 100, variable cost 60; a special order at 70 still adds 10 of contribution per unit if capacity is idle. Be careful that the order does not cannibalise full-price sales, breach price-discrimination rules, or set a precedent that erodes the regular price.

Q510. How do you decide whether to drop a product line or segment?

Drop it only if the contribution margin lost is less than the fixed costs actually saved (avoidable fixed costs); a segment showing an accounting loss after allocated common costs can still be worth keeping if it covers its own variable costs and contributes to shared overhead. Example: a line with 80,000 contribution but 100,000 of allocated fixed cost looks like a 20,000 loss, yet if only 30,000 of those fixed costs disappear on closure, dropping it worsens profit by $80,000 - 30,000 = 50,000$. Allocated common costs that continue after closure are irrelevant to the decision.

Q511. What is transfer pricing and why does the chosen method matter?

A transfer price is the internal price one division charges another for goods or services, and it shifts measured profit between divisions, affecting performance evaluation, manager incentives and, across borders, taxable income in each jurisdiction. Common methods: market-based (use the external market price), cost-based (variable or full cost, sometimes plus a markup), and negotiated. Tax authorities require cross-border transfers to follow the arm's-length principle; in India this is enforced under the Income-tax Act transfer-pricing rules (Sections 92 to 92F) aligned to OECD guidelines, with documentation and penalties for non-compliance.

Q512. Define the core SaaS unit-economics KPIs: CAC, LTV and the LTV/CAC ratio.

CAC (customer acquisition cost) = total sales and marketing spend / number of new customers acquired in the period. LTV (lifetime value) = average revenue per account * gross margin % / churn rate, i.e. the gross profit a customer generates over their lifetime. LTV/CAC compares value created to cost to acquire; a ratio around 3x or higher is considered healthy, while below 1x means you lose money on every customer and well above 5x may signal underinvestment in growth.

Q513. What are CAC payback period, MRR/ARR and gross margin, and how do they fit together?

CAC payback = CAC / (monthly gross-margin revenue per customer); it tells you how many months of margin recover the acquisition cost, with under 12 months generally considered good. MRR is monthly recurring revenue and ARR = MRR * 12; both count only recurring subscription revenue, excluding one-off services. SaaS gross margin = (recurring revenue - cost of delivering the service, e.g. hosting and support) / recurring revenue, typically 70-85% for software; high gross margin is what makes LTV and payback attractive.

Q514. Explain churn and net revenue retention (NRR), and why NRR can exceed 100%.

Churn is the rate at which customers or revenue are lost; logo churn counts customers, revenue churn counts lost MRR. Net revenue retention = (starting MRR + expansion - contraction - churn) / starting MRR,

measured for an existing cohort and excluding new customers. NRR can exceed 100% when expansion (upsells, seat growth, price increases) outweighs churn and downgrades, meaning the existing base grows even without new logos; best-in-class SaaS firms run NRR of 120%+, a powerful compounding engine.

Q515. What is the Rule of 40, and what is the FP&A business-partnering role including driver-based planning?

The Rule of 40 says a healthy software company's revenue growth rate plus profit margin (often EBITDA or free-cash-flow margin) should be at least 40%, e.g. 30% growth + 12% margin = 42% passes; it captures the growth-versus-profitability trade-off. FP&A business partnering means finance works alongside operating teams to turn strategy into forecasts, challenge assumptions and inform decisions rather than just report numbers. Driver-based planning builds the model from operational drivers (units, price, headcount, conversion rate, churn) rather than simply growing last year's totals by a percentage; a simple cost model links volume drivers to variable costs, adds step-fixed costs as capacity thresholds are crossed, and layers in fixed overhead, so a change in a driver flows transparently through to cost, margin and cash.

E. Quant / Derivatives

E1. Options & Futures Fundamentals (Q516-Q546)

Q516. What is a call option versus a put option, and who holds the right versus the obligation?

A call gives the buyer the RIGHT (not obligation) to BUY the underlying at the strike K; a put gives the buyer the right to SELL at K. The option BUYER holds the right and pays a premium; the option SELLER (writer) receives the premium and takes on the OBLIGATION to deliver (call) or buy (put) if exercised. So buyers have limited loss (premium) and sellers have limited gain (premium) but potentially large loss.

Q517. What does "long" versus "short" a position mean, and what are the four basic option positions?

Long means you bought/own it; short means you sold/wrote it. The four positions are: long call (bullish, pay premium), short call (bearish/neutral, receive premium), long put (bearish, pay premium), short put (bullish/neutral, receive premium). Long positions have the rights; short positions carry the obligations and the margin requirement.

Q518. What is the payoff of a long call at expiry, and write the formula?

Payoff = $\max(S_T - K, 0)$, where S_T is the spot at expiry and K the strike. Profit = $\max(S_T - K, 0) - \text{premium}$. Example: Nifty call K = 22000 bought for 150; if Nifty expires at 22300, payoff = 300, profit = $300 - 150 = 150$ per unit.

Q519. What is the payoff of a long put at expiry?

Payoff = $\max(K - S_T, 0)$. Profit = $\max(K - S_T, 0) - \text{premium}$. Example: a TCS 3800 put bought for 60; if TCS expires at 3700, payoff = 100 and profit = $100 - 60 = 40$ per share. A put gains value as the underlying falls, capped because price cannot go below zero.

Q520. What are the payoffs to the option WRITER (short call and short put)?

Short call payoff = $-\max(S_T - K, 0)$; short put payoff = $-\max(K - S_T, 0)$. The writer keeps the premium if the option expires worthless. Profit (short call) = premium - $\max(S_T - K, 0)$; a naked short call has theoretically unlimited loss, while a short put's max loss is $K - \text{premium}$ (if the stock goes to zero).

Q521. Define intrinsic value and time value of an option.

Intrinsic value is the in-the-money amount if exercised now: call = $\max(S - K, 0)$, put = $\max(K - S, 0)$. Time value = option premium - intrinsic value, reflecting the chance of moving further ITM before expiry. Example: stock 105, call K = 100 priced 8: intrinsic = 5, time value = 3. Time value decays to zero at expiry (theta decay).

Q522. Explain moneyness (ITM/ATM/OTM) for calls and puts.

For a call: ITM when $S > K$, ATM when $S \sim K$, OTM when $S < K$. For a put it reverses: ITM when $S < K$, OTM when $S > K$. ITM options have intrinsic value; ATM/OTM options are pure time value. Example: spot 100 — the 95 call is ITM, the 100 call ATM, the 105 call OTM; the 105 put is ITM and the 95 put OTM.

Q523. What is the difference between a payoff diagram and a profit diagram?

A payoff diagram plots the value at expiry (ignoring cost), so a long call's payoff is a hockey stick that is flat at zero up to K then slopes up at 45 degrees. A profit diagram subtracts the premium paid (or adds premium received), shifting the line down by the premium, so breakeven for a long call is at $K + \text{premium}$, not K. Profit accounts for the upfront cost; payoff does not.

Q524. What is the breakeven point for a long call and a long put?

Long call breakeven = $K + \text{premium}$ (stock must rise enough to recover the premium). Long put breakeven = $K - \text{premium}$. Example: 22000 Nifty call at 150 breaks even at 22150; 22000 put at 130 breaks even at 21870. Below/above breakeven respectively the buyer is still net negative even if the option is ITM.

Q525. What are the main drivers of an option premium?

Six drivers (Black-Scholes inputs): underlying price S , strike K , time to expiry T , volatility σ , risk-free rate r , and dividends/yield q . Calls rise with S , T , σ and r , and fall with q ; puts rise with K , T , σ and q , and fall with r . Volatility and time are the key "time value" drivers — higher σ or more time means more premium for both calls and puts.

Q526. Why does higher volatility increase BOTH call and put premiums?

Options have asymmetric payoffs (limited downside, open upside), so a wider distribution of S_T raises the expected payoff of the optionality without a matching penalty. Higher σ fattens the tails, increasing the probability of finishing deep ITM, while the OTM side is already floored at zero. This is captured by positive vega for both calls and puts.

Q527. What is the difference between American and European exercise style?

American options can be exercised any time up to expiry; European options only at expiry. American options are worth at least as much as European (more flexibility). In India, both index options (Nifty, Bank Nifty) AND single-stock options on NSE are European-style; US equity options are typically American, while US index options like SPX are European.

Q528. When is it ever optimal to early-exercise an American option?

For a non-dividend-paying stock, an American CALL is never optimally exercised early (you forgo time value and interest on K) — so it equals a European call. American PUTS can be optimal to exercise early (deep ITM, to receive K and earn interest). Early call exercise becomes rational just before a large dividend, to capture the dividend.

Q529. State the put-call parity relationship and explain each term.

For European options: $C - P = S - Ke^{-rT}$, where C and P are call and put prices on the same K and T , S is spot, r the risk-free rate. It says a long call plus short put (same strike) replicates a forward on the stock. Rearranged: $C + Ke^{-rT} = P + S$ (fiduciary call = protective put). With a dividend yield q , replace S with Se^{-qT} .

Q530. Use put-call parity to value a missing option: $S = 100$, $K = 100$, $r = 6\%$, $T = 0.5$, call = 7. Find the put.

$P = C - S + Ke^{-rT} = 7 - 100 + 100e^{-(0.06 \times 0.5)} = 7 - 100 + 100 \times 0.9704 = 7 - 100 + 97.04 = 4.04$. The put is worth about 4.04. The call exceeds the put here because the present-valued strike (97.04) is below spot, reflecting the interest carry benefit to the call holder.

Q531. How do you build a synthetic forward (or synthetic long stock) using options?

Synthetic long forward = long call + short put at the same strike K and expiry; its payoff is $S_T - K$, identical to a forward. Synthetic long stock today = long call + short put + invest Ke^{-rT} in bonds. Conversely synthetic short stock = short call + long put. These follow directly from put-call parity $C - P = S - Ke^{-rT}$.

Q532. What does Indian regulation say about exercise style and how does it affect parity?

SEBI/NSE moved all equity (single-stock) options to European style years ago, and index options were always European, so standard European put-call parity $C - P = S - Ke^{-rT}$ holds cleanly in India. Because Indian stock options are physically settled, the parity should also account for any expected dividend (use Se^{-qT}) and the cost of holding/borrowing the stock to deliver.

Q533. What is a covered call versus a naked call?

A covered call is writing a call while owning the underlying stock — the stock covers the delivery obligation, capping upside but earning premium income. A naked (uncovered) call is writing a call without owning the stock, exposing the writer to theoretically unlimited loss and requiring large margin. Covered calls are an income/yield strategy; naked calls are pure short-volatility bets.

Q534. What is the difference between exercise and assignment?

Exercise is the option holder (long) choosing to invoke their right to buy/sell at K. Assignment is when the clearing corporation (e.g., NSE Clearing / OCC in the US) selects a short option holder to fulfil that obligation. Assignment is usually random across writers; an in-the-money option at expiry is auto-exercised, triggering assignment to a writer who must then deliver or take delivery.

Q535. What is the difference between open interest and volume?

Volume is the number of contracts traded during the day (resets daily). Open interest (OI) is the total number of contracts currently open/outstanding (not yet closed or expired). OI rises when a new buyer and new seller create a contract and falls when both close. Rising price with rising OI suggests a strengthening trend; rising price with falling OI suggests short covering.

Q536. How do futures differ from forwards?

Futures are exchange-traded, standardized, marked-to-market daily, and guaranteed by a clearinghouse with margin — minimizing counterparty risk. Forwards are OTC, customized, settled at maturity, and carry counterparty/credit risk with no daily cash flows. Because of daily settlement, futures have small convexity differences from forwards, but for short maturities and low rate-correlation they are treated as equal in price.

Q537. What is the cost-of-carry fair value of a future/forward? Give the formula.

$F = Se^{((r - q)T)}$ for continuous compounding, where r is the risk-free rate, q the dividend yield (or convenience yield for commodities), and T the time to expiry. For commodities with storage cost u : $F = Se^{((r + u - y)T)}$, y = convenience yield. Example: $S = 22000$ Nifty, $r = 6\%$, $q = 1.5\%$, $T = 0.25$: $F = 22000e^{(0.045 \times 0.25)} = 22000 \times 1.0113 = 22249$.

Q538. What is contango versus backwardation?

Contango is when futures price is above spot ($F > S$), the normal state when carry cost (r) exceeds yield (q) — the curve slopes upward. Backwardation is $F < S$, common when convenience yield or dividends exceed financing cost, or with supply tightness. As expiry approaches, the basis converges to zero, so futures pull toward spot.

Q539. Define basis and basis risk.

Basis = spot price - futures price (some texts use futures - spot; be consistent). At expiry basis converges to zero. Basis risk is the risk that the basis changes unpredictably before you close a hedge, so the hedge does not perfectly offset the cash position — it arises when the hedging instrument, location, or maturity does not exactly match the exposure. Example: hedging jet fuel with crude futures leaves basis risk from the fuel-crude spread.

Q540. How do you hedge an equity portfolio with index futures using beta? Give the contract formula.

Number of contracts = $\text{beta} \times (\text{Portfolio value}) / (\text{Futures price} \times \text{lot multiplier})$. Going SHORT that many index futures neutralizes market risk. Example: INR 1 crore portfolio, beta 1.2, Nifty future 22000, lot 50, contract value = $22000 \times 50 = 11,00,000$; $\text{contracts} = 1.2 \times 1,00,00,000 / 11,00,000 = 10.9$, so short about 11 contracts.

Q541. How do you change a portfolio's beta (not fully hedge) using index futures?

Contracts = $(\text{beta}_{\text{target}} - \text{beta}_{\text{current}}) \times \text{Portfolio value} / (\text{Futures price} \times \text{lot multiplier})$. A positive result means buy futures (raise beta), negative means sell (lower beta). Example: move beta from 1.0 to 0.5 on a

INR 1 crore book, Nifty future contract value 11,00,000: contracts = $(0.5 - 1.0) * 1,00,00,000 / 11,00,000 = -4.5$, so short ~5 contracts.

Q542. What is the minimum-variance hedge ratio for a futures hedge?

$h = \text{correlation}(\text{spot}, \text{futures}) * (\sigma_{\text{spot}} / \sigma_{\text{futures}}) = \text{covariance}(\text{spot returns}, \text{futures returns}) / \text{variance}(\text{futures returns})$. It is the slope of regressing spot price changes on futures price changes, minimizing variance of the hedged position. Number of contracts = $h * (\text{size of exposure}) / (\text{futures contract size})$; for a portfolio this h^* effectively equals the beta versus the index.

Q543. What are initial margin and maintenance margin, and what triggers a margin call?

Initial margin is the deposit required to open a futures position; maintenance margin is the lower threshold the balance must stay above. As positions are marked-to-market daily, losses reduce the margin account; if it falls below maintenance, a margin call requires topping up back to the initial margin. Example: initial 1,00,000, maintenance 75,000 — if MTM losses drop the balance to 70,000, you must add 30,000.

Q544. What is SPAN margin in the Indian F&O market?

SPAN (Standard Portfolio Analysis of Risk) is the system used by NSE/exchanges to compute the worst-case one-day loss of a portfolio across price and volatility scenarios, setting the initial margin. Total upfront margin = SPAN margin + Exposure margin (an additional buffer). SPAN nets risk across a portfolio (e.g., a hedged spread needs far less margin than two outright positions).

Q545. What are typical lot sizes and expiries for Nifty and Bank Nifty derivatives?

Lots are exchange-defined and revised periodically — Nifty is commonly 25 (revised from 50/75 over time) and Bank Nifty around 15-35; always check the current NSE circular. Index options had weekly and monthly expiries, though SEBI has rationalized to fewer weekly expiries (one benchmark weekly per exchange) plus monthly contracts that expire on a designated weekday. Stock F&O are monthly only.

Q546. What is physical versus cash settlement, and how does India apply it?

Cash settlement pays the net difference in cash at expiry; physical settlement requires actual delivery of the underlying. In India, all single-stock F&O (futures and options) are PHYSICALLY settled — ITM stock options and stock futures held to expiry result in delivery obligations and full contract-value funds/shares. Index derivatives (Nifty, Bank Nifty) are CASH settled against the settlement value, since you cannot deliver an index.

E2. Black-Scholes-Merton & the Greeks (Q547-Q581)

Q547. What does the Black-Scholes-Merton (BSM) model do, and who is it for?

BSM gives a closed-form fair value for a European option on a non-dividend-paying stock, derived in 1973 by Black, Scholes and Merton. Its central insight is that an option can be perfectly replicated by a continuously rebalanced portfolio of the stock and a risk-free bond, so by no-arbitrage the option must cost the same as that replicating portfolio. This removes any need to know investor risk preferences or the stock's expected return, which is why it became the foundation of modern derivatives pricing.

Q548. List the core assumptions behind BSM.

The key assumptions are: (1) stock prices follow geometric Brownian motion, so log returns are normally distributed and prices are lognormal; (2) volatility σ is constant and known; (3) the risk-free rate r is constant and known; (4) no arbitrage opportunities exist; (5) trading is continuous and the stock is perfectly divisible; (6) no transaction costs or taxes; (7) the option is European (exercise only at expiry); and (8) no dividends in the base model (Merton later relaxed this). Most real-world deviations from BSM trace back to one of these assumptions failing.

Q549. Write the BSM formula for a European call.

$C = SN(d_1) - Ke^{(-rT)}N(d_2)$, where $d_1 = (\ln(S/K) + (r + 0.5\sigma^2)T) / (\sigma\sqrt{T})$ and $d_2 = d_1 - \sigma\sqrt{T}$. Here S is spot, K is strike, T is time to expiry in years, r is the continuously-compounded risk-free rate, sigma is annualized volatility, and N() is the standard normal cumulative distribution function (CDF). The call value is the discounted expected payoff under the risk-neutral measure.

Q550. What is the BSM formula for a European put?

$P = Ke^{(-rT)}N(-d_2) - SN(-d_1)$, using the same d_1 and d_2 as the call. Equivalently you can get it from put-call parity: $P = C - S + Ke^{(-rT)}$. Both routes give the identical price, since BSM is internally consistent with no-arbitrage parity.

Q551. What does "lognormal prices" mean and why does BSM assume it?

It means the stock price S_T at any future date is lognormally distributed, i.e. $\ln(S_T)$ is normally distributed. This follows from modeling returns as $dS/S = \mu dt + \sigma dW$, geometric Brownian motion, where percentage changes (not absolute changes) are random and normally distributed. Lognormality conveniently keeps prices strictly positive and makes the expected payoff integral solvable in closed form.

Q552. Interpret $N(d_2)$ in the BSM call formula.

$N(d_2)$ is the risk-neutral probability that the call finishes in-the-money, i.e. the probability that $S_T > K$ under the risk-neutral measure. The term $Ke^{(-rT)}N(d_2)$ is therefore the present value of paying the strike, weighted by the chance you actually exercise. It is a risk-neutral probability, not the real-world probability, because it uses r rather than the true expected return.

Q553. Interpret $N(d_1)$ in the BSM call formula.

$N(d_1)$ is the option's delta — the sensitivity of the call price to the stock — and the term $S*N(d_1)$ is the present value of receiving the stock conditional on exercise. More technically $N(d_1)$ is the risk-neutral expected value of S_T given finishing ITM, discounted and divided by S. Loosely, $N(d_1)$ tells you how many shares to hold to hedge, while $N(d_2)$ is the exercise probability.

Q554. Why does BSM use the risk-free rate r instead of the stock's real expected return?

Because the option is priced by replication/hedging, not by forecasting. A delta-hedged option position is locally riskless, so it must earn the risk-free rate to avoid arbitrage; the stock's true drift μ cancels out of the pricing equation entirely. This is the essence of risk-neutral valuation: we discount the expected payoff computed under a measure where every asset drifts at r.

Q555. What is risk-neutral valuation in one sentence?

Price any derivative as the expected value of its payoff computed under the risk-neutral probability measure (where all assets grow at r), discounted back at the risk-free rate: $\text{Price} = e^{(-rT)} * E_Q[\text{payoff}]$. It works because of no-arbitrage and replication, and it sidesteps the need to estimate risk premia or real-world drift.

Q556. How does Merton's dividend adjustment change the formula?

For a stock paying a continuous dividend yield q, replace S with $Se^{(-qT)}$ everywhere. So $C = Se^{(-qT)}N(d_1) - Ke^{(-rT)}N(d_2)$, with $d_1 = (\ln(S/K) + (r - q + 0.5\sigma^2)T) / (\sigma\sqrt{T})$. Intuitively dividends leak value out of the stock you would receive on exercise, so the call is worth less and the put more. The same q-adjustment is used to price index options (where q is the index dividend yield) and, via cost-of-carry, FX and futures options.

Q557. Define delta and give the BSM delta for a call and a put.

Delta is the rate of change of option value per unit change in the underlying: $\text{delta} = dV/dS$. For a (non-dividend) European call, $\text{delta} = N(d_1)$, ranging from 0 (deep OTM) to 1 (deep ITM); for a put, $\text{delta} = N(d_1)$

-1, ranging from -1 to 0. With a dividend yield q the call delta is $e^{-qT}N(d_1)$. An ATM call has delta around 0.5.

Q558. Why is delta also called the hedge ratio?

Because to neutralize the directional risk of one short call you hold delta shares of the underlying (delta shares per option), making the combined position locally insensitive to small spot moves. If you are short a call with delta 0.55 on a lot, you buy $0.55 \times \text{lot_size}$ shares to be delta-neutral. As spot moves, delta changes, so the hedge must be rebalanced — that is dynamic delta hedging.

Q559. Why is delta a rough approximation of the probability of finishing ITM?

For a call, $\text{delta} = N(d_1)$ while the true risk-neutral ITM probability is $N(d_2)$, and since $d_1 > d_2$ these differ only by the $\sigma\sqrt{T}$ gap. For short-dated, low-vol options that gap is small, so traders use delta as a quick proxy for "chance of expiring ITM." A 25-delta option is loosely treated as having roughly a 25% chance of finishing ITM, though strictly the probability is $N(d_2)$, a bit lower.

Q560. Define gamma and state where it is largest.

Gamma is the second derivative of value with respect to spot, $\text{gamma} = d(\text{delta})/dS = d^2V/dS^2$; it measures how fast delta changes. Gamma is identical for a call and the corresponding put (parity), is always positive for long options, and is largest for at-the-money options close to expiry, where delta swings most violently around the strike. The BSM formula is $\text{gamma} = N'(d_1) / (S\sigma\sqrt{T})$, where $N'()$ is the normal PDF.

Q561. Why do traders care about gamma even though they hedge delta?

Because delta hedging is only exact for infinitesimal moves; gamma measures the hedging error from large moves. High positive gamma means your delta hedge becomes stale quickly, forcing frequent rebalancing, but it also means a delta-neutral long-option book profits from big moves in either direction. Gamma is the curvature that a single delta number cannot capture, so a full hedge tracks both.

Q562. Define vega and its units.

Vega is the sensitivity of option value to a change in volatility: $\text{vega} = dV/d\sigma$. By convention it is quoted per 1 percentage-point (0.01) change in vol. The BSM formula is $\text{vega} = S\sqrt{T}N'(d_1)$ (times e^{-qT} with dividends); it is the same for calls and puts, always positive for long options. Note "vega" is not a real Greek letter, but the name stuck.

Q563. Where is vega largest, and how does it differ from gamma's profile?

Vega is largest for at-the-money options that are longer-dated, because more time means more room for volatility to move the terminal price distribution. This contrasts with gamma, which is also largest ATM but peaks for short-dated options. So a long-dated ATM option is a big vega / small gamma bet, while a near-expiry ATM option is a big gamma / small vega bet.

Q564. Define theta and explain its usual sign.

Theta is the rate of change of value with the passage of time: $\text{theta} = dV/dt$, usually quoted per day (per calendar or trading day). For most long options theta is negative — the option loses time value as expiry approaches ("time decay") — and decay is fastest for at-the-money options near expiry. Deep ITM European puts (and some dividend cases) can have positive theta, but the typical long-option holder is fighting theta.

Q565. What is the relationship between theta, gamma and vega for an option's time value?

They are linked through the BSM PDE: $\text{theta} + 0.5\sigma^2S^2\text{gamma} + rS\text{delta} - rV = 0$. The key intuition is the gamma-theta tradeoff: positive gamma (good — you profit from moves) is paid for by negative theta (bad — you bleed time value). A delta-neutral long-gamma book earns from realized volatility but pays theta; it is profitable only if realized vol exceeds the implied vol it paid.

Q566. Define rho and say when it matters.

Rho is the sensitivity of value to the risk-free rate: $\rho = dV/dr$, quoted per 1 percentage-point change. For a call $\rho = KTe^{(-rT)}N(d_2) > 0$; for a put $\rho = -KTe^{(-rT)}N(-d_2) < 0$. Rho matters most for long-dated options (LEAPS) and in high-rate environments; for typical short-dated equity options it is the least important Greek.

Q567. Walk through delta hedging and why it must be dynamic.

To hedge a sold call you buy delta shares; the position is then insensitive to small spot moves. But because gamma is non-zero, delta changes as spot moves and as time passes, so you must continuously rebalance — buy more shares as spot rises, sell as it falls. In a perfect BSM world this continuous re-hedging exactly replicates the option, and the cost of replication equals the BSM premium.

Q568. What is gamma scalping?

Gamma scalping is the trading strategy of holding a long-gamma, delta-neutral position (e.g. long a straddle, delta-hedged) and profiting from realized volatility. As spot moves, your delta drifts; you rebalance by selling shares after rallies and buying after dips — mechanically "buy low, sell high" — locking in small profits. These scalping gains must beat the theta you pay; you win when realized vol > implied vol.

Q569. Spell out the gamma-theta tradeoff with a quick example.

Long gamma earns money from movement but bleeds theta; short gamma collects theta but loses on big moves. Roughly, the daily P&L of a delta-hedged book is about $0.5\gamma(dS)^2 - \theta_{\text{per_day}}$. If a stock with gamma giving a $0.5\gamma S^2$ of, say, a 0.01 daily variance equivalent earns more from the squared move than the day's theta cost, the long-gamma trader profits; otherwise theta wins.

Q570. Summarize the sign of each Greek for a long call versus a long put.

Long call: delta + (0 to 1), gamma +, vega +, theta - (usually), rho +. Long put: delta - (-1 to 0), gamma +, vega +, theta - (usually), rho -. Gamma and vega are positive for any long option (call or put); shorting any option flips all signs. The asymmetry is in delta and rho, which carry the directional and rate-direction information.

Q571. What is vanna?

Vanna is a second-order Greek measuring how delta changes when volatility changes (or equivalently how vega changes when spot changes): $\text{vanna} = d^2V/(dS d\sigma)$. It is important for risk-managing skew, because it links your directional exposure to vol moves. Books with significant vanna can see their delta shift purely from a change in implied vol, even with spot unchanged.

Q572. What is volga (vomma)?

Volga, also called vomma, is the second derivative of value with respect to volatility: $\text{volga} = d^2V/d\sigma^2$, i.e. how vega changes as vol changes. It captures the convexity of the option in vol and is central to pricing volatility-of-volatility and the smile. Long options that are away-from-the-money tend to have high positive volga, which benefits from large swings in implied vol.

Q573. What is charm?

Charm (delta decay) measures how delta changes with the passage of time: $\text{charm} = d(\text{delta})/dt = d^2V/(dS dt)$. It tells a hedger how their delta will drift overnight or over a weekend even if spot does not move, which matters for setting hedges before non-trading periods. ATM options near expiry have the largest charm, so hedges can go stale quickly as time passes.

Q574. How do the main Greeks change as an option approaches expiry?

Gamma and theta both spike sharply for at-the-money options as $T \rightarrow 0$ (delta flips between 0 and 1 over a tiny spot range, and time value decays fastest). Vega shrinks toward zero because there is little time left for

vol to act. Delta tends toward 0 or 1 (a step function), and rho also shrinks as the discounting horizon collapses.

Q575. How do Greeks change as volatility rises?

Higher sigma widens the terminal price distribution: vega is positive so option values rise, and OTM options gain the most in relative terms. Higher vol pushes deltas of OTM/ITM options back toward 0.5 (more uncertainty about ending side of the strike), spreads gamma over a wider spot range (lowering the ATM peak), and increases time value, making theta (in absolute terms) larger. In short, vol "smears" the Greek profiles across strikes.

Q576. Describe the Cox-Ross-Rubinstein (CRR) binomial model.

CRR discretizes time into N steps where each step the stock moves up by factor $u = e^{(\sigma\sqrt{dt})}$ or down by $d = 1/u$, with risk-neutral up-probability $p = (e^{(r\Delta t)} - d)/(u - d)$. You build the price tree forward, then discount expected payoffs backward step by step at the risk-free rate. It is flexible enough to price American options and discrete dividends, which closed-form BSM cannot handle directly.

Q577. How does the binomial model relate to BSM?

As the number of steps N goes to infinity ($dt \rightarrow 0$), the CRR binomial price converges to the BSM price for a European option — the discrete lognormal tree approaches continuous geometric Brownian motion. CRR is essentially a numerical scheme that, in the limit, reproduces the closed-form BSM value. This convergence is why binomial trees are taught as the intuitive bridge to BSM.

Q578. Why can't BSM price American options directly, and how is it handled?

BSM assumes European exercise (only at expiry), but American options can be exercised early, adding an optimal-stopping decision BSM's formula ignores. American puts (and calls on dividend-paying stocks) can be worth exercising early, so they need numerical methods — binomial/trinomial trees, finite differences, or least-squares Monte Carlo. Note an American call on a non-dividend stock is never optimally exercised early, so it equals its European value (and BSM applies).

Q579. What is the volatility smile/skew, and why does it contradict BSM?

BSM assumes a single constant sigma, which would imply identical implied vol across all strikes. In reality, plotting implied vol against strike shows a smile or (for equity indices) a downward skew, with OTM puts priced at higher implied vols. This reflects fat tails and crash fear — markets price more probability into extreme down-moves than the lognormal assumption allows — directly violating BSM's constant-vol and lognormal assumptions.

Q580. Name the main situations where BSM breaks down in practice.

BSM struggles with: (1) American early exercise (needs trees/numerical methods); (2) fat tails and jumps — real returns have more extreme moves than lognormal predicts; (3) the volatility smile/skew — vol is not constant across strikes; (4) stochastic and time-varying volatility; (5) transaction costs and discrete (not continuous) hedging; and (6) discrete dividends and changing rates. Practitioners patch these with local/stochastic-vol models, jump-diffusion, and implied-vol surfaces.

Q581. State put-call parity and how it constrains the Greeks of calls and puts.

Put-call parity is $C - P = Se^{(-qT)} - Ke^{(-rT)}$ (or $C - P = S - Ke^{(-rT)}$ with no dividends). Differentiating gives Greek relationships: $\text{call_delta} - \text{put_delta} = e^{(-qT)}$ (so for $q=0$, $\text{delta_call} - \text{delta_put} = 1$); gamma and vega are identical for a matched call and put; and theta/rho differ only by the derivatives of the deterministic parity terms. This is why a call and put on the same strike/expiry share their gamma and vega exposures exactly.

E3. Implied Volatility, the Smile & Vol Trading (Q582-Q601)

Q582. What is the difference between implied volatility and realized (historical) volatility?

Realized (historical) volatility is backward-looking: it is the actual standard deviation of past returns, e.g. computed as the annualized stdev of daily log returns. Implied volatility (IV) is forward-looking: it is the volatility input embedded in the current market price of an option, i.e. the market's expectation of future volatility over the option's life. They often differ — IV typically trades above realized vol (a "variance risk premium") because option sellers demand compensation for taking on risk.

Q583. What exactly is implied volatility?

Implied volatility is the value of sigma that, when plugged into the Black-Scholes-Merton (BSM) formula, makes the model price equal the observed market price of the option. All other BSM inputs (spot S, strike K, time T, rate r, dividend q) are observable, so the option's market price can be "inverted" to back out a single unknown — sigma. It is the market's consensus forecast of annualized volatility expressed in BSM's language, not a directly observed quantity.

Q584. Why can't you solve for implied volatility with a closed-form formula?

Because the BSM price is a non-linear, monotonic-but-not-invertible function of sigma — the option price depends on sigma through the cumulative normal $N(d_1)$ and $N(d_2)$, which cannot be algebraically rearranged to isolate sigma. So IV must be found numerically using iterative root-finding such as Newton-Raphson or bisection. The good news is that price is strictly increasing in sigma ($\text{vega} > 0$), so a unique solution exists for any arbitrage-free price.

Q585. How does Newton-Raphson solve for implied volatility?

You iterate $IV_{\{n+1\}} = IV_n - (\text{model_price}(IV_n) - \text{market_price}) / \text{vega}(IV_n)$, where vega is the derivative of price with respect to sigma. Starting from a guess (say 20%), each step uses the pricing error divided by vega to correct the estimate, and it converges quadratically — usually 3-5 iterations to high precision. Example: if model price is 4.80 vs market 5.00 and vega is 0.40 (per 1.00 = 100% vol), then $IV_{\{n+1\}} = IV_n + 0.20/0.40 = IV_n + 0.50$ vol points adjustment.

Q586. What is the volatility smile?

The volatility smile is the empirical pattern that implied volatility varies with strike (or moneyness) rather than being constant as BSM assumes. When you plot IV against strike, out-of-the-money puts and calls often show higher IV than at-the-money options, producing a U or "smile" shape. It exists because real return distributions have fatter tails and are not lognormal, so the market prices tail strikes richer than a flat-vol BSM would.

Q587. What is volatility skew and how does it differ from a smile?

Skew refers to an asymmetric smile where IV is systematically higher on one side of the strike range than the other. Equity indices show a downward (negative) skew: low strikes (downside puts) carry much higher IV than high strikes. A symmetric smile implies the market sees roughly equal fat-tail risk on both sides, whereas skew implies asymmetric crash/tail expectations.

Q588. Why do equity index options exhibit a pronounced downward skew?

Three reinforcing reasons: (1) crash-o-phobia — after 1987, markets price a higher probability of sudden large down moves than up moves; (2) the leverage effect — as equity prices fall, firms' debt/equity ratios rise, increasing volatility, so down moves and higher vol are correlated; and (3) supply/demand — institutions persistently buy downside puts for portfolio insurance, bidding up low-strike IV. The Nifty and Bank Nifty options on the NSE show this same negative skew.

Q589. Why do some assets (like FX or certain commodities) show a more symmetric smile?

For currencies, a large move in either direction is roughly equally likely and equally feared — a sharp appreciation or depreciation both create tail risk — so demand for OTM puts and calls is more balanced, producing a near-symmetric smile. Equities skew because the downside is the feared, leverage-amplified direction; FX has no such asymmetry between "up" and "down." Some commodities can even show a reverse (upward) skew when supply shocks make spikes the feared tail.

Q590. What is the term structure of volatility?

It is how implied volatility varies with time to maturity for a fixed strike or moneyness (typically ATM). In calm markets it usually slopes upward (longer maturities price more vol as uncertainty accumulates and mean-reverts toward a long-run level); in stress it inverts, with near-term IV spiking above long-dated IV because the shock is expected to be acute but temporary. It reflects the market's expected path of volatility over time.

Q591. What is the implied volatility surface?

The IV surface is a 3D plot of implied volatility as a function of two axes — strike (or moneyness/delta) and time to maturity. A cross-section at fixed maturity gives the smile/skew; a cross-section at fixed strike gives the term structure. Traders use the whole surface to price and risk-manage exotic and vanilla books consistently, and a well-behaved surface must be free of arbitrage in both the strike and time directions.

Q592. What is a 25-delta risk reversal and what does it measure?

A 25-delta risk reversal is the IV of the 25-delta call minus the IV of the 25-delta put ($RR = IV(25d \text{ call}) - IV(25d \text{ put})$). It quantifies the skew/asymmetry of the smile: a negative risk reversal means puts are bid over calls (typical for equities), signalling downside fear. It is the standard FX-market measure of skew and is quoted directly as a tradable structure (long call / short put at equal delta).

Q593. What is a 25-delta butterfly and what does it measure?

The 25-delta butterfly is $BF = 0.5 \cdot (IV(25d \text{ call}) + IV(25d \text{ put})) - IV(ATM)$, the average of the two 25-delta wing vols minus the at-the-money vol. It measures the convexity (curvature) of the smile — how much the wings are bid relative to the center, i.e. how fat the market prices both tails. Together, ATM vol, the risk reversal, and the butterfly fully parameterize the smile at a given maturity.

Q594. What does the VIX measure and why is it called the "fear gauge"?

The VIX is the CBOE Volatility Index: the market's expected 30-day implied volatility of the S&P 500, expressed in annualized percentage points. It is computed model-free from a strip of OTM SPX option prices (a variance-swap-style weighted sum across strikes), not from a single BSM inversion. It is called the fear gauge because it spikes when investors rush to buy protection during sell-offs — VIX above ~30 signals stress, below ~15 signals calm.

Q595. What is India VIX and how does it relate to the VIX?

India VIX is the NSE's volatility index, measuring the expected 30-day implied volatility of the Nifty 50 from near- and next-month Nifty options, using the same CBOE-style model-free methodology. It is widely watched as India's fear gauge: it jumps around events like Union Budget, RBI policy, and election results. Like the VIX, it is annualized in percent and tends to be negatively correlated with the underlying index.

Q596. What is the difference between IV rank and IV percentile?

IV rank measures where current IV sits between its 52-week low and high: $IV \text{ Rank} = (IV_{\text{now}} - IV_{\text{low}}) / (IV_{\text{high}} - IV_{\text{low}}) \cdot 100$. IV percentile measures the fraction of days over the lookback during which IV was below the current level. IV rank is sensitive to a single extreme spike (it anchors on the min/max), while IV percentile better reflects the typical distribution; an IV rank of 50 with a percentile of 80 tells you IV is mid-range but has rarely been this high day-to-day.

Q597. What is vega and how do volatility traders use it to express views?

Vega is the option's sensitivity to a 1-percentage-point change in implied volatility ($d\text{Price}/d\sigma$); for example a vega of 0.12 means the option gains about 0.12 if IV rises by 1 point. A trader who thinks vol is too cheap buys options (long vega); one who thinks vol is too rich sells options (short vega). Vega is largest for at-the-money, longer-dated options, so vol views are often expressed with ATM straddles or variance swaps to get clean vega exposure with minimal directional delta.

Q598. Explain sticky-strike versus sticky-delta (sticky-moneyness) regimes.

Sticky-strike assumes each strike keeps its IV as the spot moves, so an option's IV stays put while its moneyness changes — common in quiet, range-bound markets. Sticky-delta (sticky-moneyness) assumes the IV attached to a given delta/moneyness stays fixed, so the whole smile slides with the spot — common in trending markets. The assumption matters for hedging because it changes the effective delta: under sticky-delta the skew adds to the BSM delta, altering how many shares you hedge with.

Q599. What no-arbitrage constraints must a volatility surface satisfy?

Two key families: (1) the calendar-spread (term) constraint — total implied variance $\sigma^2 T$ must be non-decreasing in maturity for a fixed moneyness, otherwise a calendar arbitrage exists; and (2) the butterfly constraint — the implied risk-neutral density from the strike dimension must be non-negative, which requires call prices to be convex and monotonic in strike. Violations (e.g. IV curving too sharply, or near-dated variance exceeding far-dated) imply free lunches, so quants fit arbitrage-free parameterizations like SVI.

Q600. Why is selling options described as "short vol" and "short gamma"?

When you sell an option you are short vega (you lose if IV rises) — hence short vol. You are also short gamma: a delta-hedged short option position must buy as the underlying rises and sell as it falls (buy high, sell low), so realized movement bleeds money via negative gamma P&L. The short option earns positive theta (time decay) to compensate; the position profits when realized volatility comes in below the IV you sold, and loses when the market moves more than priced.

Q601. What is the difference between variance and volatility, and why does volatility mean-revert?

Volatility is σ (standard deviation of returns); variance is σ^2 . The distinction matters because variance is additive over time and is what variance swaps and the VIX actually trade — payoffs are linear in variance but non-linear (concave) in vol, so a vol position has convexity. Implied (and realized) volatility mean-reverts: it oscillates around a long-run average, spiking in crises and decaying back, which is why models like Heston and GARCH include a mean-reversion term and why traders fade extreme IV rank readings expecting normalization.

E4. Option Strategies & Quant Math (Q602-Q616)

Q602. What is a covered call, and what market view does it express?

A covered call means you hold the underlying (long stock or futures) and sell one call against it, usually slightly out-of-the-money. The view is mildly bullish to neutral: you expect the stock to stay flat or rise modestly, and you collect the premium as extra yield. Max profit = (strike - purchase price) + premium received, capped if the stock rallies above the strike; max loss is large (stock can fall to zero, cushioned only by the premium); breakeven = purchase price - premium. On NSE this is a common Nifty/Bank Nifty F&O overlay where investors write monthly calls on a held position to earn premium income.

Q603. What is a protective put and when would you use it?

A protective put is long underlying plus a long put, acting like an insurance policy on your position. The view is bullish but nervous: you want upside exposure while capping downside, so you pay a premium to floor your losses at the strike. Max loss = (purchase price - put strike) + premium paid; max profit is unlimited

(stock can keep rising) minus the premium cost; breakeven = purchase price + premium. Indian portfolio managers buy Nifty puts as a hedge before events like the Union Budget or RBI policy to protect against a market drop.

Q604. Construct a bull call spread and state its risk profile.

A bull call spread buys a lower-strike call and sells a higher-strike call of the same expiry, financing part of the cost with the short call. The view is moderately bullish with defined risk: you profit if the underlying rises toward the higher strike but you cap both gain and loss. Max profit = $(K_{\text{high}} - K_{\text{low}}) - \text{net premium paid}$; max loss = net premium paid; breakeven = $K_{\text{low}} + \text{net premium}$. Example on Nifty: buy 24000 call at 200, sell 24200 call at 120, net debit 80; max profit = $200 - 80 = 120$, max loss = 80, breakeven = 24080.

Q605. How does a bear put spread work and what is its payoff?

A bear put spread buys a higher-strike put and sells a lower-strike put of the same expiry, a defined-risk way to play a moderate fall. The view is moderately bearish: you profit as the underlying declines toward the lower strike, with capped gain and capped loss. Max profit = $(K_{\text{high}} - K_{\text{low}}) - \text{net premium paid}$; max loss = net premium paid; breakeven = $K_{\text{high}} - \text{net premium}$. It is the mirror image of the bull call spread and is cheaper than buying a naked put because the short put offsets premium.

Q606. What is a long straddle, and when is it profitable?

A long straddle buys a call and a put at the same strike (usually at-the-money) and expiry. The view is long volatility / direction-agnostic: you expect a big move either way but are unsure of direction, common around earnings or binary events. Max loss = total premium paid (if the stock pins the strike at expiry); max profit is unlimited on the upside and large on the downside; there are two breakevens at strike $\pm$ total premium. You need the realized move to exceed the combined premium, so a long straddle loses if implied volatility was already rich and the stock stays quiet.

Q607. How does a long strangle differ from a long straddle?

A long strangle buys an out-of-the-money call and an out-of-the-money put (different strikes, same expiry), so it is cheaper than a straddle but needs a larger move to pay off. The view is still long volatility expecting a big swing. Max loss = total premium paid; max profit is unlimited up / large down; breakevens are at $K_{\text{call}} + \text{total premium}$ and $K_{\text{put}} - \text{total premium}$. Because both legs start OTM the upfront cost (and max loss) is lower, but the no-profit zone between the strikes is wider than a straddle's.

Q608. Describe a short straddle/strangle and its risk.

Selling a straddle (short ATM call + short ATM put) or a strangle (short OTM call + short OTM put) collects premium and profits when the underlying stays range-bound and implied volatility falls. The view is short volatility / neutral. Max profit = premium received (when the stock pins the strikes); max loss is theoretically unlimited on a large move in either direction; breakevens are strike $\pm$ premium (straddle) or each strike $\pm$ premium (strangle). This is high-risk: a sharp gap (e.g., a Bank Nifty move on RBI surprise) can blow past the premium, which is why SEBI margins for short option positions are heavy.

Q609. What is a long butterfly spread and what view does it express?

A long call butterfly buys one lower-strike call, sells two middle-strike calls, and buys one higher-strike call (equally spaced strikes, same expiry) for a small net debit. The view is range-bound / short volatility: you expect the underlying to expire near the middle strike. Max profit = (wing width) - net premium, achieved at the middle strike; max loss = net premium paid; two breakevens sit at $K_{\text{low}} + \text{net premium}$ and $K_{\text{high}} - \text{net premium}$. It is a cheap, fully defined-risk way to bet on a pin, with the lost premium being the only downside.

Q610. Explain the iron condor and the risk reversal as two range / skew strategies.

An iron condor sells an OTM put spread and an OTM call spread simultaneously (four legs), collecting net premium; the view is range-bound short volatility, with max profit = net premium received if the underlying stays between the short strikes, and max loss = (spread width - net premium) capped by the long wings. A risk reversal is directional-plus-skew: sell an OTM put and buy an OTM call (bullish version), often for near-zero cost, giving long-stock-like exposure financed by the put premium; max loss is large if the stock falls below the put strike, max profit is unlimited above the call strike. A collar is the conservative cousin of the risk reversal applied to a held position: long stock + long protective put + short covered call, which caps both downside and upside for little or no net premium and is widely used by Indian PMS/HNI clients to lock in a band around a concentrated holding.

Q611. What is geometric Brownian motion and why is it the standard model for stock prices?

Geometric Brownian motion (GBM) models the stock as the SDE $dS = \mu S dt + \sigma S dW$, where μ is the drift (expected return), σ the volatility, and dW the increment of a Wiener process. It is "geometric" because returns rather than absolute price levels are random: dividing by S gives $dS/S = \mu dt + \sigma dW$, so percentage moves are stationary. This is preferred because it keeps prices strictly positive and produces lognormal prices with proportional volatility, matching how equity returns actually behave; it is the engine behind Black-Scholes.

Q612. What are the defining properties of a Wiener process (standard Brownian motion)?

A Wiener process $W(t)$ satisfies: $W(0) = 0$; it has independent increments; each increment is normally distributed with $W(t) - W(s) \sim \text{Normal}(0, t-s)$ (variance grows linearly with time); and its paths are continuous but nowhere differentiable. A key scaling consequence is that dW has standard deviation $\sqrt{dt}$, i.e., $E[dW] = 0$ and $E[dW^2] = dt$, which is why volatility scales with $\sqrt{T}$ (the square-root-of-time rule). These properties make W the canonical source of continuous-time randomness in finance.

Q613. Apply Ito's lemma to S and explain why $d(\ln S)$ carries a $-0.5\sigma^2$ term.

Ito's lemma is the stochastic chain rule: for $f(S)$ it adds a second-order term, $df = (df/dS) dS + 0.5(d^2f/dS^2) (dS)^2$, where crucially $(dW)^2 = dt$. Applying it to $f = \ln S$ with $dS = \mu S dt + \sigma S dW$ gives $d(\ln S) = (\mu - 0.5\sigma^2) dt + \sigma dW$. The $-0.5\sigma^2$ (the Ito correction) appears because $\ln$ is concave and $(dS)^2 = \sigma^2 S^2 dt$ is non-zero in stochastic calculus; intuitively, volatility drags down the growth rate of the log because gains and losses are not symmetric in compounded returns.

Q614. Why are stock prices lognormal while log-returns are normal?

From $d(\ln S) = (\mu - 0.5\sigma^2) dt + \sigma dW$, the log of the price is normally distributed, so the continuously-compounded return $\ln(S_T/S_0)$ is $\text{Normal}((\mu - 0.5\sigma^2)T, \sigma^2 T)$. Because the price itself is the exponential of a normal variable, $S_T = S_0 \cdot \exp((\mu - 0.5\sigma^2)T + \sigma\sqrt{T}Z)$, S_T follows a lognormal distribution: skewed right and bounded below at zero. This is economically sensible, prices cannot go negative and a stock can multiply many-fold but only fall to zero, which is exactly the asymmetry a lognormal captures.

Q615. How does Monte-Carlo simulation price a derivative, and how fast does it converge?

Monte-Carlo simulates many price paths under the risk-neutral GBM $S_T = S_0 \cdot \exp((r - 0.5\sigma^2)T + \sigma\sqrt{T}Z)$, computes the payoff on each path (e.g., $\max(S_T - K, 0)$ for a call), averages them, and discounts: $\text{price} = \exp(-rT) \cdot (1/N) \cdot \sum(\text{payoff}_i)$. Its statistical error shrinks as $1/\sqrt{N}$, so cutting the error in half requires roughly 4x the paths, slow but dimension-independent, which makes it the go-to method for path-dependent and high-dimensional options (Asian, basket, barrier) where closed forms do not exist. Variance-reduction tricks like antithetic variates and control variates are used to tighten the estimate without simply quadrupling N .

Q616. What is the difference between the real-world and risk-neutral measures, and why price under risk-neutral?

Under the real-world (physical) measure P the stock drifts at its true expected return μ , reflecting investors' risk premium; under the risk-neutral measure Q the drift is replaced by the risk-free rate r , so $dS = rSdt + \sigma SdW$ and all assets earn r on average. We price derivatives under Q because, by the fundamental theorem of asset pricing, the absence of arbitrage means a discounted price is a martingale under Q , giving $\text{price} = \exp(-rT) * E_Q[\text{payoff}]$, a value that does not depend on anyone's risk appetite. The real-world measure is used for risk management, forecasting and VaR, while the risk-neutral measure is used purely for consistent, arbitrage-free valuation; only the drift changes between them, the volatility σ is the same.

E5. Advanced Derivatives & Quantitative Finance (Q617-Q646)

Q617. What distinguishes a path-dependent option from a path-independent one, and give examples of each?

A path-independent option's payoff depends only on the underlying's level at expiry (or exercise), not the route taken - a vanilla European call paying $\max(S_T - K, 0)$ is the canonical example. A path-dependent option's payoff depends on the trajectory: Asian options use the average price over the life, barrier options depend on whether a level was touched, and lookbacks use the running max/min. Path dependence usually rules out simple closed forms and pushes you toward trees, PDEs, or Monte Carlo.

Q618. What is an Asian (average) option, why is it cheaper than a vanilla, and who uses it?

An Asian option's payoff uses the average underlying over a set of dates rather than the single expiry price - an average-rate call pays $\max(A - K, 0)$ where A is the arithmetic mean of observed prices. Averaging lowers effective volatility, so the option is cheaper than a vanilla and is robust to a single-day spike or expiry-day manipulation. Treasurers and commodity desks use them to hedge a stream of regular transactions (e.g. a fuel buyer hedging monthly purchases) where the average cost is what they actually pay; note the geometric-average version has a closed form while the arithmetic one needs Monte Carlo.

Q619. Define knock-in versus knock-out barrier options, the parity linking them to vanillas, and why their Greeks are dangerous.

A knock-out starts alive and dies if the underlying touches the barrier (an up-and-out call dies if S hits H above spot); a knock-in only springs into existence once the barrier is touched, and for the same terms $\text{knock-in} + \text{knock-out} = \text{the vanilla}$. Barriers are cheaper than vanillas because you forgo payoff in some scenarios, which is why FX desks and structured notes love them. Near the barrier the value can jump to zero over a tiny move, so delta and gamma spike and flip sign, making delta-hedging extremely costly - dealers manage this with barrier shifting or vanilla over-hedges.

Q620. What is a lookback option and why is it expensive?

A lookback's payoff references the best price reached over the life: a floating-strike lookback call pays $S_T - \min(S)$, effectively letting you buy at the low, while a fixed-strike version pays $\max(\max(S) - K, 0)$. Because you always capture the optimal historical price, lookbacks embed perfect market-timing you could never achieve, so they are expensive and mostly a structuring/teaching tool rather than a liquid product.

Q621. Explain a digital (binary) option and how desks actually hedge its blow-up risk near the strike.

A cash-or-nothing digital pays a fixed amount (say 1) if $S_T > K$ and zero otherwise, so its payoff is a step function. As expiry nears, delta becomes a tall narrow spike at K and gamma flips sign violently around it. Dealers therefore replicate and hedge a digital with a tight vanilla call spread - long a call at $K - \epsilon$, short at $K + \epsilon$, scaled by $1/(2\epsilon)$ - which caps the local risk and is what really gets traded.

Q622. What is a chooser option and how does it decompose?

A chooser lets the holder decide at a future date t whether the option becomes a call or a put (same strike K , expiry T). At the choice date a rational holder takes $\max(C, P)$, so via put-call parity a simple chooser equals a call expiring at T plus a put expiring at the earlier choice date t . It suits a buyer expecting a big move of uncertain direction and is cheaper than buying an outright straddle.

Q623. What is a compound option and where do compound options arise naturally?

A compound option is an option on an option - e.g. a call-on-a-call gives the right to buy a call for a preset premium at the first expiry, so it has two strikes and two dates. They appear in real-options analysis (a phased capital project is a sequence of options on options), in hedging contingent deals where the underlying hedge may or may not be needed, and in installment-premium structures. They are very sensitive to volatility, exhibiting strong gamma-of-vega.

Q624. What is a basket option and why is correlation its key risk?

A basket option pays off on a weighted portfolio of underlyings, e.g. $\max(w_1S_1 + w_2S_2 + \dots - K, 0)$. Its price hinges on pairwise correlations because basket variance = sum over i,j of $w_iw_j\sigma_i\sigma_j\rho_{ij}$; higher correlation raises basket vol and the price. Buyers use baskets to express a view on a group (an index, an FX basket) more cheaply than buying options on each name, precisely because diversification lowers basket vol.

Q625. Derive the risk-neutral up-probability in a binomial tree and state the CRR choice of u and d .

Risk-neutral pricing requires the discounted expected stock to equal today's price: $S = e^{(-r\Delta t)}(pSu + (1-p)Sd)$, giving $p = (e^{(r\Delta t)} - d) / (u - d)$, which must satisfy $d < e^{(r\Delta t)} < u$ so $0 < p < 1$. Cox-Ross-Rubinstein sets $u = e^{(\sigma\sqrt{\Delta t})}$ and $d = 1/u$ so the tree recombines and matches lognormal variance as $\Delta t \rightarrow 0$. Example: $S=100$, $\sigma=20\%$, $\Delta t=1y$, $r=5\%$ gives $u=1.221$, $d=0.819$, $p=(1.0513-0.819)/(1.221-0.819)=0.578$; p is a pricing measure, not the real-world probability.

Q626. How does a binomial tree price an American option, and when is early exercise actually optimal?

You roll backward from expiry taking value = $\max(\text{immediate exercise payoff}, e^{(-r\Delta t)}(pV_{up} + (1-p)V_{down}))$ at every node, comparing intrinsic value to continuation value - backward induction is required because a node's continuation value needs its children's already-computed values. This early-exercise test is exactly what Black-Scholes cannot capture. Early exercise is never optimal for an American call on a non-dividend stock (you would forfeit time value and interest on the strike), but can pay for a deep in-the-money American put or a call just before a large dividend.

Q627. What does a trinomial tree add over a binomial tree?

A trinomial tree allows three moves per step (up, unchanged, down), giving an extra degree of freedom that improves stability and convergence and maps directly onto an explicit finite-difference scheme. The middle branch lets you align nodes with specific levels - invaluable for barrier options and for interest-rate models like Hull-White - and fit the local drift/vol more flexibly than a binomial lattice.

Q628. Why is Monte Carlo natural for exotics, and why does its error shrink only like $1/\sqrt{N}$?

Monte Carlo simulates many risk-neutral paths, averages the payoffs, and discounts: price = $e^{(-rT)}\text{mean}(\text{payoff})$; it handles path dependence and high dimensions (baskets, many dates) trivially where trees suffer the curse of dimensionality. By the Central Limit Theorem the standard error is $\sigma_{\text{payoff}}/\sqrt{N}$, so error falls as $1/\sqrt{N}$ - halving it needs 4x the paths and one extra decimal needs 100x. Its weakness is early exercise (needs Longstaff-Schwartz regression) and this slow convergence, which is why variance reduction matters.

Q629. Explain antithetic and control variates as variance-reduction techniques.

Antithetic variates pair each random draw Z with $-Z$; the two payoffs are negatively correlated, so $\text{Var}((X+Y)/2) = (\text{Var}(X) + \text{Var}(Y) + 2\text{Cov}(X,Y))/4$ shrinks when Cov is negative - essentially free and best for monotonic payoffs. A control variate uses a correlated quantity with a known analytic price: estimator = $\text{MC}(\text{target}) - b(\text{MC}(\text{control}) - \text{TrueValue}(\text{control}))$, with optimal $b = \text{Cov}(\text{target}, \text{control}) / \text{Var}(\text{control})$. For an arithmetic Asian the geometric Asian is the classic control because it is highly correlated and has a closed form, often cutting variance by an order of magnitude.

Q630. What are finite-difference methods conceptually, and how do explicit and implicit schemes differ?

Finite-difference methods solve the Black-Scholes PDE directly by discretizing price and time onto a grid, approximating derivatives with differences, and stepping backward from the payoff boundary. The explicit scheme is simple but only stable under a tight time-step condition (like a trinomial tree); implicit and Crank-Nicolson schemes solve a linear system each step but are unconditionally stable and more accurate. FDM excels for low-dimensional American and barrier problems where you want the full price surface and clean Greeks.

Q631. Why do models beyond Black-Scholes exist - what does the volatility smile prove?

Black-Scholes assumes one constant volatility, which would imply a flat implied-vol curve across strikes; the market instead shows a smile/skew (out-of-the-money equity puts priced richer). This proves the real return distribution has fatter tails and negative skew than lognormal, so a single sigma cannot price all strikes consistently. Hence local-vol, stochastic-vol, and jump models exist to reproduce the observed smile and price exotics consistently with vanillas.

Q632. Contrast local volatility (Dupire) with stochastic volatility models.

Local volatility makes sigma a deterministic function $\sigma(S,t)$ calibrated to fit today's entire vanilla surface exactly via the Dupire equation, but it predicts unrealistic, flattening future smile dynamics. Stochastic volatility makes variance itself a random process with its own vol-of-vol and spot correlation, capturing the realistic forward skew that matters for exotics and forward-starting products. Desks often use local-stochastic-vol (LSV) hybrids to get both an exact fit and sensible dynamics.

Q633. What do the Heston and SABR stochastic-volatility models add, and what do their parameters control?

Heston models variance as a mean-reverting CIR process $dv = \kappa(\theta - v)dt + \xi\sqrt{v}dW_2$ correlated (rho) with spot: kappa is reversion speed, theta the long-run variance, xi the vol-of-vol setting smile convexity, and rho (negative for equities) the skew - it has a semi-closed form for vanillas making calibration fast. SABR models a forward as $dF = \alpha F^{\beta} dW_1$, $d(\alpha) = \nu \alpha^{\beta} dW_2$: alpha sets the vol level, beta the backbone (0 normal, 1 lognormal), rho the skew, nu the curvature. SABR is the market standard for swaption smiles because Hagan's formula gives implied vol per strike directly.

Q634. What does Merton's jump-diffusion model capture that pure diffusion cannot?

Merton adds a compound-Poisson jump term to geometric Brownian motion, so the stock can gap: $dS/S = (\mu - \lambda k)dt + \sigma dW + (J-1)dN$ with jump intensity lambda. Jumps generate the steep short-dated skew and fat tails that pure diffusion cannot (diffusion smiles flatten at short maturities), and they make perfect delta-hedging impossible, so markets become incomplete. This matters for crash risk, gap risk on barriers, and overnight/earnings moves.

Q635. Explain delta-hedging P&L, the gamma-theta relationship and the breakeven daily move, and how transaction costs alter it.

For a delta-hedged book the P&L over a small interval is roughly $0.5\gamma(dS)^2 + \theta dt - \text{long gamma earns on big moves while theta bleeds time value} - \text{so the two break even at } dS \sim \sqrt{2S\gamma/\theta} \sqrt{dt}$, about one implied standard deviation. Example: $S=20000$, $\sigma=15\%$, $dt=1/252$ gives breakeven $\sim$

$200000.15 \cdot \sqrt{1/252} \sim 189$ points/day; move more and long gamma wins, less and theta wins. In reality you re hedge discretely, so transaction costs act like a volatility add-on (Leland) that scales with $\sqrt{\text{rehedge frequency}}$, and bid-ask plus impact can erode the theoretical realized-minus-implied-vol edge entirely.

Q636. Define vega, vanna, and volga and why a book must watch all three.

Vega is sensitivity to volatility ($dV/d\sigma$); vanna is the cross term $dVega/dSpot$ (equivalently $dDelta/d\sigma$); and volga (vomma) is $dVega/d\sigma$, the convexity of value in vol. A vega-neutral book can still bleed when spot moves (vanna) or when vol moves a lot (volga), so smile-risk management requires hedging all three - this underpins the vanna-volga pricing method in FX. Ignoring vanna/volga is why naive vega hedges fail in skewed, moving markets.

Q637. How does a variance swap pay off and how is it replicated with a strip of options?

A variance swap pays $\text{notional} \cdot (\text{realized_variance} - K_var)$, where realized variance is the annualized sum of squared log returns, giving pure variance exposure with no path-dependent delta. It is replicated statically by a portfolio of out-of-the-money calls and puts weighted by $1/K^2$ across all strikes, plus a dynamic delta hedge; the fair strike is essentially the $1/K^2$ -weighted integral of option prices (the model-free basis for the VIX). The $1/K^2$ weighting is precisely what makes the payoff track variance regardless of where spot goes.

Q638. How does a volatility swap differ from a variance swap, and why is the variance swap easier to hedge?

A variance swap settles on realized variance (vol squared) while a volatility swap settles on realized volatility itself, $\text{notional} \cdot (\text{realized_vol} - K_vol)$. The variance swap has clean static replication via the $1/K^2$ option strip, but a vol swap's payoff is the concave square root of variance, so by Jensen's inequality its fair strike sits below $\sqrt{\text{variance strike}}$ by a convexity adjustment and it cannot be statically replicated. That vol-of-vol convexity is an extra risk that makes vol swaps harder to hedge.

Q639. What is dispersion trading and what view does it express?

Dispersion trading sells index volatility/variance and buys volatility on the constituent single names (or vice versa), profiting when realized correlation differs from the implied correlation embedded in index vol. Since $\text{index variance} = \text{sum of weighted single-name variances} + \text{correlation terms}$, the trade is essentially short implied correlation - it makes money when stocks move idiosyncratically while index vol stays subdued. The main risk is a correlation spike in a sell-off, when everything moves together and the short-index-vol leg loses badly.

Q640. Decompose a capital-guaranteed (principal-protected) note into options.

A capital-guaranteed note is a zero-coupon bond plus a long call: the issuer puts most of the principal into a discount bond that accretes to par (guaranteeing capital) and spends the residual buying upside via a call, often capped or with a participation rate. Example: at 5% rates, ~78 of 100 buys a 5y zero maturing at 100, leaving ~22 to buy index calls; the participation rate is whatever option exposure that budget funds. This is why low rates make these notes far less attractive - less premium is left for the call.

Q641. Decompose a reverse convertible and an autocallable into options.

A reverse convertible is a high-coupon note where the investor is effectively short a put: you receive an enhanced coupon for selling downside protection and are repaid in shares (taking the loss) if the underlying breaches a level at maturity. An autocallable embeds a short down-and-in put plus a series of digital/barrier coupons with an automatic early-redemption feature - it pays attractive coupons and redeems early if the underlying is above a trigger on observation dates, but exposes the buyer to sharp losses through the protection barrier. Both are yield-enhancement products that sell tail risk to coupon-hungry investors.

Q642. What is a quanto option and what extra risk does it introduce?

A quanto pays off on a foreign-currency underlying but settles in the domestic currency at a fixed exchange rate - e.g. an option on the S&P 500 paying in INR at a pre-agreed rate, removing FX risk for the rupee

investor. Pricing requires a quanto drift adjustment: the foreign asset's drift shifts by $-\rho\sigma_{FX}$ because the payoff is exposed to the asset-FX covariance. The key extra risk is therefore the unobservable asset-FX correlation, which must be estimated and hedged.

Q643. What is CVA and how does counterparty risk enter OTC derivative pricing?

Credit Valuation Adjustment is the market value of counterparty default risk - the amount you mark down a derivative's risk-free value because the counterparty might default when you are in-the-money. Conceptually $CVA = \int \text{loss-given-default} \times \text{expected positive exposure} \times \text{marginal default probability}$, discounted, so it depends on the counterparty's credit spread and your future exposure profile. Post-2008, XVA desks price CVA (plus DVA, FVA, MVA, KVA) into every OTC trade, which is why uncollateralized trades cost more than cleared/collateralized ones.

Q644. How do you compute and manage the Greeks of a whole derivatives portfolio?

Because the Greeks are partial derivatives of value, they add across positions on a common underlying: portfolio delta = sum of (position delta * quantity), and likewise for gamma, vega, theta. A desk nets these to a few aggregate numbers and hedges the residual - trade the underlying to zero delta, then use options to flatten gamma and vega within limits. The subtlety is that vega across strikes/expiries is not truly fungible (smile and term-structure risk), so books bucket vega by tenor and strike rather than summing one number.

Q645. Distinguish risk-neutral from real-world probabilities and explain why pricing uses the risk-neutral measure.

Real-world (physical) probabilities include a risk premium and are used for forecasting, P&L distributions, and risk metrics like VaR. Risk-neutral probabilities are an artificial measure under which all assets drift at the risk-free rate, so a derivative's price equals the discounted risk-neutral expected payoff - this holds because it enforces no-arbitrage via replication, not because investors are actually risk-neutral. The two are linked by the market price of risk (Girsanov): you price with risk-neutral but forecast and size risk with real-world.

Q646. Why does gamma explode near expiry for at-the-money index options, and what does it mean for weekly Nifty/Bank Nifty expiries?

For an at-the-money option, gamma scales roughly as $1/(\sigma\sqrt{T})$, so as $T \rightarrow 0$ gamma blows up and tiny spot moves cause huge delta swings. On NSE weekly expiries, dealers short these cheap near-dated options carry enormous gamma, and their delta-hedging forces large index buying/selling that can pin price near big open-interest strikes or amplify intraday moves on expiry day. This is why expiry sessions see sharp, mean-reverting moves around key strikes and why short-gamma sellers can be badly hurt by a single large move.

F. Investment Banking, M&A & LBO

F1. Mergers & Acquisitions (Q647-Q676)

Q647. What is the difference between a merger, an acquisition, and a takeover?

A merger is the combination of two companies into one new or surviving entity, often framed as a "merger of equals" with mutual agreement. An acquisition is one company (the acquirer) buying another (the target), with the target ceasing to exist independently or becoming a subsidiary. A takeover is simply an acquisition viewed from the angle of control changing hands; it can be friendly (board-approved) or hostile (pursued over the board's objection, e.g., via a tender offer directly to shareholders).

Q648. What are horizontal, vertical, and conglomerate mergers?

A horizontal merger combines two firms in the same industry and stage (e.g., two banks merging) to gain market share and scale. A vertical merger combines firms at different stages of the same supply chain (e.g., a car maker buying a parts supplier) to secure inputs or distribution. A conglomerate merger combines firms in unrelated businesses (e.g., a cement firm buying a telecom company), driven by diversification rather than operational overlap.

Q649. What are the main strategic rationales for doing an M&A deal?

The classic motives are: synergies (cost savings and revenue gains from combining), increased market share and pricing power, vertical integration to control the supply chain, geographic or product diversification, acquiring technology/talent/IP faster than building it, and tax benefits (using a target's losses or a more favorable structure). A buyer may also acquire simply to deploy excess cash or because organic growth has stalled.

Q650. What are synergies in M&A, and what are the two main types?

Synergies are the incremental value created when the combined firm is worth more than the sum of the two standalone firms. Cost synergies come from eliminating duplicate costs (overlapping staff, facilities, IT, procurement scale). Revenue synergies come from selling more (cross-selling, new geographies, pricing power). Total deal value justification often hinges on $\text{Synergy value} = \text{PV of after-tax synergy cash flows} - \text{integration costs}$.

Q651. Why are cost synergies considered more credible than revenue synergies?

Cost synergies are largely within management's control: you can decide to close a plant, cut a duplicate finance team, or consolidate vendors, and the savings are quantifiable and time-bound. Revenue synergies depend on customer behavior, competitive response, and sales execution, all of which are uncertain and slow to materialize. Experienced acquirers and analysts therefore discount or ignore revenue synergies and underwrite deals primarily on cost synergies.

Q652. Walk me through the typical M&A deal process from start to close.

On the sell-side, a banker pitches the engagement, then a teaser (anonymous one-pager) goes to buyers, followed by an NDA, then a detailed CIM (Confidential Information Memorandum). Interested buyers submit indicative bids, sign a Letter of Intent (LOI), and conduct due diligence in a data room. The parties negotiate and sign the SPA (Share/Stock Purchase Agreement), obtain regulatory and shareholder approvals, and then close (consideration paid, ownership transfers). Post-close comes integration.

Q653. What is a teaser versus a CIM in a sell-side process?

A teaser is a short, anonymous 1-2 page document describing the business at a high level (sector, financial highlights, investment thesis) sent to gauge interest before any confidentiality agreement. A CIM

(Confidential Information Memorandum), also called an Information Memorandum (IM), is a detailed 30-100 page document with the company named, full financials, management, operations, and projections, shared only after an NDA is signed.

Q654. What is a Letter of Intent (LOI) and is it binding?

An LOI is a preliminary document outlining the key proposed deal terms: indicative price, structure (asset vs stock), exclusivity, and timeline. Most of the LOI is non-binding (price can still change after due diligence), but certain clauses, typically exclusivity/no-shop, confidentiality, and break fees, are binding. It signals serious intent and frames the subsequent definitive agreement (SPA).

Q655. What is the difference between an asset purchase and a stock (share) purchase?

In a stock purchase, the buyer acquires the target's shares and takes on the entire legal entity, including all assets, liabilities, and contracts (known and unknown). In an asset purchase, the buyer cherry-picks specific assets and assumed liabilities, leaving unwanted liabilities behind, and gets a stepped-up tax basis on the assets. Sellers usually prefer stock deals (clean exit, often better tax treatment); buyers often prefer asset deals (liability protection, tax step-up).

Q656. What is a control premium and why does an acquirer pay it?

A control premium is the amount paid above a target's pre-announcement market price to gain control of the company. $\text{Control premium \%} = (\text{Offer price per share} - \text{Unaffected share price}) / \text{Unaffected share price}$. Acquirers pay it because control lets them direct strategy, capture synergies, replace management, and access cash flows that minority shareholders cannot; typical premiums run roughly 20-40%, e.g., offering INR 120 for a stock trading at INR 100 is a 20% premium.

Q657. How do you choose between cash, stock, and mixed consideration in a deal?

Cash is preferred when the acquirer has the funds or cheap debt, is confident in the deal, and wants to keep 100% of the upside (and avoid diluting EPS). Stock is used when cash/debt is constrained, the acquirer's shares are richly valued (an "expensive currency"), or the buyer wants the seller to share deal risk. Mixed consideration balances the two; the choice also reflects tax (stock deals can be tax-deferred for sellers) and signaling (all-stock can signal the buyer thinks its stock is overvalued).

Q658. What does it mean for an acquisition to be accretive or dilutive to EPS?

A deal is accretive if the acquirer's pro-forma EPS (post-deal) is higher than its standalone EPS, and dilutive if it is lower. $\text{Pro-forma EPS} = (\text{Acquirer NI} + \text{Target NI} + \text{after-tax synergies} - \text{after-tax financing cost}) / \text{pro-forma share count}$. Accretion/dilution is a quick first-pass screen of whether a deal "adds to earnings," though it says nothing about whether the deal actually creates value.

Q659. What is the rule of thumb for accretion/dilution based on P/E in an all-stock deal?

In an all-stock deal, if the acquirer's P/E is higher than the target's P/E (the price paid), the deal is accretive; if the acquirer's P/E is lower, it is dilutive. Intuitively, you are issuing your high-P/E (cheap-to-issue) shares to buy cheaper low-P/E earnings, which boosts EPS. The break-even is when the two P/Es are equal.

Q660. Walk me through a simple all-stock accretion/dilution example.

Acquirer: net income USD 100m, 100m shares, EPS USD 1.00, price USD 20 (P/E 20x). Target: net income USD 20m, bought for USD 240m all-stock (target P/E = $240/20 = 12x$). Acquirer issues $240/20 = 12m$ new shares. Pro-forma NI = $100 + 20 = 120m$; pro-forma shares = 112m; pro-forma EPS = $120/112 = \text{USD } 1.071$. EPS rose from 1.00 to 1.07, so the deal is accretive, consistent with acquirer P/E (20x) > target P/E (12x).

Q661. How does financing with cash or debt change accretion/dilution versus stock?

Cash/debt deals are accretive when the after-tax yield given up or paid is lower than the target's earnings yield (E/P). Using cash, you forgo interest income (e.g., 2% after-tax) to buy earnings yielding, say, 8% (12x

P/E), so it is usually accretive. Using debt, compare the after-tax cost of debt to the target's earnings yield: accretive if after-tax cost of debt < target E/P. Because cash and debt rates are typically far below equity earnings yields, cash/debt deals are more often accretive than stock deals.

Q662. Give a worked cash-financed accretion/dilution example with debt.

Acquirer EPS USD 1.00 (NI 100m, 100m shares). It buys a target's net income of USD 20m for USD 240m, funded with debt at 5% pre-tax, tax rate 25%. After-tax interest = $240 * 5\% * (1 - 0.25) = \text{USD } 9\text{m}$. Pro-forma NI = $100 + 20 - 9 = 111\text{m}$; shares unchanged at 100m; pro-forma EPS = USD 1.11. Accretive, because the target's earnings yield ($20/240 = 8.3\%$) exceeds the after-tax cost of debt ($5\% * 0.75 = 3.75\%$).

Q663. How is goodwill created in an acquisition?

Goodwill is the excess of the purchase price (equity purchase price / consideration) over the fair value of the target's identifiable net assets acquired. $\text{Goodwill} = \text{Purchase price} - \text{Fair value of net identifiable assets (incl. identified intangibles)}$. For example, paying INR 1,000 cr for a company whose net identifiable assets are fairly valued at INR 700 cr (after writing up assets and booking intangibles) creates INR 300 cr of goodwill on the consolidated balance sheet.

Q664. Is goodwill amortized, and how is it treated under IFRS/Ind-AS vs US GAAP?

Under IFRS and Ind-AS, goodwill is not amortized; it is tested for impairment at least annually and written down if impaired. US GAAP also requires impairment testing (private companies may elect amortization). So goodwill typically sits on the balance sheet indefinitely until an impairment charge reduces it, which hits the income statement as a non-cash expense. The old practice of routine goodwill amortization was eliminated in the early 2000s.

Q665. How do you value synergies in a deal?

You estimate the incremental after-tax cash flows from synergies (cost savings plus credible revenue gains, minus the costs to achieve them), then discount them to present value, usually at the acquirer's or combined firm's WACC. $\text{Synergy value} = \text{sum of } [\text{after-tax synergy CF}_t / (1 + \text{WACC})^t] - \text{PV of integration/restructuring costs}$. The maximum justifiable premium is bounded by this synergy value, since paying more than the synergies are worth transfers all the value to the target's shareholders.

Q666. What is the difference between a friendly and a hostile takeover?

In a friendly takeover, the target's board negotiates and recommends the deal to its shareholders. In a hostile takeover, the acquirer bypasses or overrides the board, going directly to shareholders, typically via a tender offer (public bid to buy shares) or a proxy fight (to replace directors). Hostile deals are riskier, costlier, and rarer because of due-diligence limitations and entrenched defenses.

Q667. What is a poison pill and how does it work?

A poison pill (shareholder rights plan) is a defense that makes a hostile takeover prohibitively expensive by diluting the acquirer. A "flip-in" pill lets all shareholders except the hostile bidder buy new shares at a steep discount once the bidder crosses a trigger threshold (e.g., 15% ownership), massively diluting the bidder's stake. It is designed not to be triggered but to force the bidder to negotiate with the board.

Q668. Explain white knight, golden parachute, and staggered board defenses.

A white knight is a friendlier third-party acquirer the target invites to outbid a hostile bidder. A golden parachute is a large severance package for executives triggered by a change of control, raising the deal's cost and aligning management. A staggered (classified) board elects only a fraction of directors each year, so a hostile bidder cannot replace the whole board in one proxy contest, slowing any takeover.

Q669. What is an earnout and when is it used?

An earnout is deferred contingent consideration: part of the purchase price is paid later only if the target hits agreed performance targets (revenue, EBITDA, milestones) post-close. It is used to bridge a valuation gap when buyer and seller disagree on future performance, and to retain and motivate selling founders/management. Risks include disputes over how targets are measured and misaligned incentives during the earnout period.

Q670. What are the main areas covered in due diligence?

Due diligence verifies the target across multiple dimensions: financial (quality of earnings, working capital, debt), legal (contracts, litigation, title), tax, commercial/market (customers, competition), operational, HR (key staff, pensions), IT/technology, environmental, and regulatory/compliance. The goal is to confirm the investment thesis, uncover hidden liabilities, and inform price, representations/warranties, and indemnities in the SPA.

Q671. What is a fairness opinion and who provides it?

A fairness opinion is a formal letter from an independent investment bank stating whether the consideration in a deal is fair, from a financial point of view, to a company's shareholders. The target's board obtains it to support its recommendation and to discharge its fiduciary duty (reducing litigation risk). It is an opinion on price fairness, not a guarantee the deal is the best possible outcome, and the bank is typically paid a fixed fee for it.

Q672. What is a minority squeeze-out (freeze-out)?

A squeeze-out is a transaction that compels remaining minority shareholders to sell their shares so the acquirer can reach 100% ownership, usually after crossing a high ownership threshold. It is done via a short-form merger, scheme of arrangement, or delisting, with minority holders receiving cash (often a court- or regulator-supervised fair price). It lets the acquirer fully consolidate, remove public-company costs, and avoid minority veto rights.

Q673. Why do studies suggest most M&A deals destroy value for the acquirer?

Acquirers frequently overpay (winner's curse in competitive auctions), pay control premiums that exceed realized synergies, and overestimate revenue synergies while underestimating integration difficulty, culture clash, and customer/employee attrition. The value created by synergies is often handed entirely to target shareholders via the premium, leaving acquirer shareholders worse off. Empirically, the target's stock usually pops on announcement while the acquirer's often falls.

Q674. Under India's SEBI Takeover Code (SAST), when is a mandatory open offer triggered?

Under the SEBI (Substantial Acquisition of Shares and Takeovers) Regulations, 2011, an acquirer crossing 25% of voting rights in a listed company must make a mandatory open offer to public shareholders for at least an additional 26% of shares. A separate "creeping acquisition" trigger applies between 25% and the maximum permissible non-public shareholding: acquiring more than 5% in a financial year also triggers an open offer. The open-offer price is set by a formula based on highest negotiated/market prices over defined look-back periods.

Q675. What is the role of the CCI in Indian M&A?

The Competition Commission of India (CCI) reviews combinations (mergers/acquisitions) above prescribed asset/turnover thresholds to ensure they do not cause an Appreciable Adverse Effect on Competition (AAEC) in India. Qualifying deals require prior CCI approval before closing; the CCI can approve, approve with modifications/remedies, or block them. India also added a "deal value threshold" (transactions above INR 2,000 cr with substantial Indian operations) bringing large digital/tech acquisitions into the review net.

Q676. What is an NCLT scheme of arrangement in Indian M&A?

A scheme of arrangement under Sections 230-232 of the Companies Act, 2013 is a court-supervised process, adjudicated by the National Company Law Tribunal (NCLT), used for mergers, demergers, amalgamations, and capital restructurings. It requires approval by the requisite majority of shareholders and creditors (75% in value), plus sanction from the NCLT, and clearances from regulators (SEBI/stock exchanges for listed firms, RBI, CCI as applicable). Once sanctioned, the scheme is binding on all stakeholders, which is why it is favored for complex restructurings and squeeze-outs.

F2. Leveraged Buyouts & Private Equity (Q677-Q707)

Q677. What is a leveraged buyout (LBO)?

An LBO is the acquisition of a company using a large amount of borrowed money (debt) to fund the purchase price, with a relatively small slice of equity from a private equity (PE) sponsor. The acquired company's own assets and cash flows secure and service the debt, so the target effectively pays for its own purchase over time. A typical structure might be 60-70% debt and 30-40% equity, versus an all-equity acquisition.

Q678. Why do PE firms use leverage instead of buying companies with all equity?

Leverage amplifies equity returns: by putting in less of their own money and borrowing the rest, the sponsor's percentage return on the equity invested is magnified when the deal goes well. Debt is also cheaper than equity (interest is tax-deductible, and lenders accept lower returns because they are senior and secured), which lowers the blended cost of capital. The trade-off is higher financial risk, since fixed interest and principal payments must be met regardless of business performance.

Q679. What are the three levers of value creation in an LBO?

The three levers are: (1) debt paydown / deleveraging, where the company's cash flow repays debt and transfers enterprise value from lenders to equity; (2) EBITDA growth, via revenue growth, margin expansion or cost cuts; and (3) multiple expansion, selling at a higher EV/EBITDA multiple than the entry multiple. Returns come from the combination of growing the numerator (EBITDA), shrinking net debt, and ideally re-rating the multiple.

Q680. Walk through the intuition of how an LBO generates returns.

You buy a company for an enterprise value funded mostly with debt and a thin equity check. Over the hold period the company's free cash flow pays down debt, so even if enterprise value stays flat, equity value (EV minus net debt) grows simply because debt shrank. Layer on EBITDA growth and a higher exit multiple, and the equity value at exit can be several times the original investment. It is analogous to buying a house with a mortgage: as you pay down the loan and the house appreciates, your home equity grows much faster than the house price.

Q681. Explain a "Sources & Uses" of funds table in an LBO.

Uses lists what the cash is needed for: the purchase of equity, refinancing of existing debt, and transaction fees (advisory, financing, legal). Sources lists where the cash comes from: new debt (term loans, bonds), sponsor equity, and sometimes rollover equity from management or cash on the target's balance sheet. Sources must equal Uses; the sponsor equity is typically the plug that makes the two sides balance.

Q682. Give a simple Sources & Uses example.

Suppose you buy a company for an enterprise value of 1,000 (USD millions). Uses: 1,000 for the purchase plus 30 of fees = 1,030. Sources: 600 of new debt, 30 of target cash used, and the remaining 400 as sponsor equity (the plug). Equity = Uses minus other sources = $1,030 - 600 - 30 = 400$.

Q683. What are the main tranches of debt in an LBO capital structure, from most to least senior?

In rough order of seniority: revolver (revolving credit facility), Term Loan A, Term Loan B, senior secured notes, senior unsecured / high-yield bonds, mezzanine debt, and PIK instruments. Each lower tranche is riskier, more expensive, and sits further down the repayment waterfall. Below all debt sits the sponsor's common equity, which is the residual claim.

Q684. What is the difference between Term Loan A and Term Loan B?

Term Loan A (TLA) amortizes more heavily (meaningful annual principal repayment), has a shorter maturity (typically 5-6 years), a lower spread, and is usually held by banks. Term Loan B (TLB) has minimal amortization (often 1% per year with a bullet at maturity), a longer maturity (6-7 years), a higher spread, and is held by institutional investors (CLOs, credit funds). TLB is the workhorse of most US LBO structures.

Q685. What is a revolver and why is it part of the structure?

A revolving credit facility is a flexible line of credit the company can draw down and repay as needed, primarily to fund working capital and seasonal cash swings. It is usually left undrawn at close (so it is often not a "source" of funds) and carries an unused commitment fee on the undrawn portion. It provides liquidity insurance, ensuring the business can operate even if cash flow is temporarily tight.

Q686. What is mezzanine debt and where does it sit?

Mezzanine is subordinated debt that sits between senior debt and equity in the capital structure, often unsecured and carrying high coupons (frequently part cash, part PIK) plus sometimes equity warrants. It fills the financing gap when senior lenders will not lend the full amount needed, taking more risk for a higher return (often low-to-mid teens). It is junior to senior debt but senior to common equity in the waterfall.

Q687. What is PIK (payment-in-kind) debt?

PIK debt does not pay cash interest; instead the interest accrues and is added to the principal balance (it "pays in kind" with more debt). This preserves cash for the operating business or for debt paydown in early years, but the balance compounds, so the total payoff at maturity is larger. PIK toggles let the borrower choose between cash and PIK each period, useful when cash flow is uncertain.

Q688. What is the debt waterfall in an LBO?

The waterfall is the priority order in which cash flows and, in a default, liquidation proceeds are paid out. Senior secured lenders (revolver, term loans) are paid first, then senior unsecured/high-yield, then mezzanine and PIK, and finally common equity holders receive whatever residual is left. Higher tranches have lower risk and lower returns; equity is last in line but captures all the upside.

Q689. What characteristics make a good LBO candidate?

Ideal targets have stable, predictable cash flows (to service debt reliably), low and discretionary capex (so cash goes to debt paydown rather than reinvestment), strong tangible assets (to serve as collateral), a defensible market position with low cyclicality, and a reasonable/undervalued entry price. Opportunities for operational improvement and a clear exit path also help. Mature, cash-generative businesses (e.g., consumer staples, services) fit better than high-growth, capital-hungry tech.

Q690. Why is low capex important for an LBO target?

Free cash flow available for debt repayment equals operating cash flow minus capex. High capex consumes the cash that would otherwise pay down debt, slowing deleveraging and increasing default risk under a heavy debt load. Low-capex businesses convert more EBITDA into free cash flow, which is exactly what is needed to service and retire the LBO debt.

Q691. What is MOIC (also called MoM) and how is it calculated?

MOIC (Multiple on Invested Capital), also called MoM (Multiple of Money) or cash-on-cash, measures total value returned relative to equity invested: $MOIC = \text{Total Equity Proceeds} / \text{Equity Invested}$. If you

invest 100 of equity and exit for 300, $MOIC = 300/100 = 3.0x$. It ignores time, so it is always paired with IRR.

Q692. What is IRR and how does it relate to MOIC?

IRR (Internal Rate of Return) is the annualized discount rate that sets the net present value of all cash flows to zero; it is the time-weighted return. For a single in-and-out investment, $IRR = MOIC^{(1/n)} - 1$, where n is the holding period in years. For a 3.0x MOIC over 5 years: $IRR = 3^{(1/5)} - 1 = 1.2457 - 1 = 0.2457$, about 24.6%.

Q693. Work through a returns example: invest 100, exit at 300 in 5 years.

$MOIC = 300 / 100 = 3.0x$. $IRR = 3^{(1/5)} - 1$. Since $3^{(0.2)} \sim 1.2457$, $IRR \sim 24.6\%$. A useful rule of thumb: roughly 3x in 5 years approximates a mid-20s percent IRR, and 2x in 5 years is about 15% ($2^{(1/5)} - 1 = 14.9\%$).

Q694. Why can MOIC and IRR tell different stories?

MOIC measures absolute multiple of capital but ignores timing; IRR measures annualized speed but rewards early cash returns. A deal returning 2.0x in 2 years has a strong IRR (41%) but modest MOIC, while 2.5x in 7 years has a higher MOIC but a lower IRR (~14%). LPs care about both: IRR for return rate and MOIC for absolute dollars (and to avoid being misled by quick flips that flatter IRR).

Q695. Walk through a conceptual "paper LBO."

A paper LBO is a back-of-envelope deal you compute mentally in an interview. Steps: (1) Entry — compute purchase EV = entry multiple x EBITDA, split into debt and equity. (2) Project EBITDA growth over ~5 years. (3) Build free cash flow each year and use it to pay down debt, tracking the falling net debt balance. (4) Exit — compute exit EV = exit multiple x final-year EBITDA, subtract remaining net debt to get exit equity value. (5) Returns — $MOIC = \text{exit equity} / \text{entry equity}$, then $IRR = MOIC^{(1/n)} - 1$.

Q696. Paper LBO: buy at 10x EBITDA of 100 with 5x debt, EBITDA grows to 150, exit at 10x, 50 debt repaid over 5 years. What are the returns?

Entry EV = $10 \times 100 = 1,000$; debt = $5 \times 100 = 500$; entry equity = $1,000 - 500 = 500$. Exit EV = $10 \times 150 = 1,500$; net debt at exit = $500 - 50 = 450$; exit equity = $1,500 - 450 = 1,050$. $MOIC = 1,050 / 500 = 2.1x$; $IRR = 2.1^{(1/5)} - 1 \sim 16\%$. Here all the gain came from EBITDA growth and debt paydown, with no multiple expansion.

Q697. How exactly does leverage amplify equity returns? Show with numbers.

Suppose EV grows from 1,000 to 1,300 (a 30% EV gain). All-equity: equity goes 1,000 to 1,300 = 1.3x. With 700 debt and 300 equity, if debt stays at 700, exit equity = $1,300 - 700 = 600 = 2.0x$ on the 300 invested. The same 30% asset move becomes a 100% equity gain because the fixed debt claim does not share in the upside, leaving more for equity, the essence of financial leverage (and the same mechanism magnifies losses on the downside).

Q698. What is the difference between maintenance and incurrence covenants?

Maintenance covenants require the borrower to satisfy financial tests on an ongoing basis (e.g., every quarter), such as keeping leverage below a ceiling, tested whether or not the company takes any action. Incurrence covenants are tested only when the company takes a specific action (incurring new debt, paying a dividend, making an acquisition). "Covenant-lite" (cov-lite) loans have only incurrence covenants, giving borrowers more flexibility, common in TLB markets.

Q699. What are the two most common financial covenants and how are they defined?

A leverage covenant caps $\text{Net Debt} / \text{EBITDA}$ below some maximum (e.g., must stay under 6.0x). A coverage covenant requires $\text{EBITDA} / \text{Interest Expense}$ (interest coverage) or $(\text{EBITDA} - \text{Capex}) /$

(Interest + Principal) (fixed-charge coverage) to stay above a minimum (e.g., above 2.0x). They protect lenders by ensuring debt remains serviceable and proportionate to cash flow; breaching one is an event of default.

Q700. What is the J-curve in private equity returns?

The J-curve describes how a PE fund's net returns (or net cash flow to LPs) dip negative in early years before turning positive, tracing a "J" shape. Early on, management fees and deal costs are drawn while investments are immature and held at cost, so reported IRR is negative; later, as portfolio companies are improved and exited, distributions flow back and IRR climbs. It explains why judging a fund's performance in its first 2-3 years is misleading.

Q701. What is the difference between entry multiple and exit multiple, and why does it matter?

The entry multiple is the EV/EBITDA paid at acquisition; the exit multiple is the EV/EBITDA received at sale. Multiple expansion (exit > entry) boosts returns; multiple contraction destroys them. Conservative LBO models often assume exit multiple equals entry multiple, so returns rely only on EBITDA growth and debt paydown rather than betting on a market re-rating.

Q702. What is a dividend recapitalization?

A dividend recap is when a PE-owned company raises new debt and uses the proceeds to pay a special dividend to the sponsor, rather than reinvesting it. It lets the sponsor extract cash and de-risk the investment before a full exit, locking in returns early and improving IRR (cash returned sooner). The downside is it re-levers the company, increasing its financial risk and reducing cushion for the remaining equity.

Q703. How do PE firms exit their investments?

The main exit routes are: (1) strategic sale (trade sale) to a corporate buyer in the same industry, often the highest price due to synergies; (2) secondary buyout, selling to another financial sponsor (PE firm); (3) IPO, listing the company on a public exchange, which may be a partial exit with staged selldown; and occasionally (4) dividend recaps as a partial return. The chosen route depends on market conditions, company size, and which yields the best risk-adjusted proceeds.

Q704. Explain the "2 and 20" fee structure, GP, and LP.

In a PE fund, Limited Partners (LPs, e.g., pension funds, endowments) provide the capital; the General Partner (GP, the PE firm) manages it. "2 and 20" means a 2% annual management fee on committed/invested capital plus 20% carried interest (carry), the GP's share of profits. On a 1,000 fund, that is ~20/year in fees and 20% of gains above the return of capital (and typically above a hurdle). Carry aligns GP incentives with LP returns.

Q705. What is a hurdle rate and a GP catch-up?

The hurdle rate (preferred return, often 8%) is the minimum annualized return LPs must receive before the GP earns any carried interest. The catch-up is a provision after the hurdle is met that pays the GP a larger share (often 100%) of subsequent profits until the GP has received its full 20% of total profits-above-cost, then both revert to the standard 80/20 split. Together they ensure LPs get their preferred return first, but the GP eventually earns carry on all the profit, not just the excess above the hurdle.

Q706. What is the difference between private equity and venture capital?

PE typically buys mature, cash-generative companies using significant leverage, often acquiring control (majority/100%) and improving operations; VC invests minority equity in early-stage, high-growth, often unprofitable startups with little or no debt. PE returns come from cash flow, leverage and operational improvement with moderate failure rates; VC follows a power-law where a few big winners offset many write-offs. PE deal sizes and certainty are higher; VC upside per deal is far larger but riskier.

Q707. What are add-on (bolt-on) acquisitions and the buy-and-build strategy, and how does PE look in India?

An add-on (bolt-on) is a smaller acquisition made by an existing PE portfolio company ("platform") to expand scale, geography or capabilities; buy-and-build strings together many add-ons to create a larger, more valuable group. Because the larger combined entity often commands a higher exit multiple than the small targets were bought at, this creates "multiple arbitrage." In India, classic leveraged buyouts are less common because RBI regulations historically restrict banks from funding share acquisitions and acquisition financing is limited; consequently Indian PE skews toward growth capital and minority stakes (e.g., funding expansion in tech, financials and consumer), with control buyouts and leverage growing but still modest versus the US/Europe, all under SEBI's AIF (Alternative Investment Fund) framework.

F3. Deal Valuation & Capital Raising (Q708-Q738)

Q708. What is precedent transaction analysis and how does it differ from trading comps?

Precedent transaction analysis (also called "deal comps" or "M&A comps") values a company based on multiples actually paid in past acquisitions of similar companies, e.g. EV/EBITDA paid in completed deals. The key difference from trading comps is that trading comps use current market prices of standalone public peers (a minority, non-control value), while precedent transactions reflect prices paid to acquire whole companies, which include a control premium and often synergy expectations. As a result, transaction multiples are usually higher than trading-comp multiples for the same sector.

Q709. Why do transaction multiples include a control premium?

In a trading comp, an investor buys a small minority stake at the prevailing share price and gets no control over the business. In a transaction, the acquirer buys 100% (or a controlling stake) and gains the ability to change management, strategy, capital structure and capture synergies, so it pays extra for that control. This excess over the unaffected trading price is the control premium, typically around 20-40% over the target's pre-announcement share price, which is why EV/EBITDA paid in deals exceeds standalone trading EV/EBITDA.

Q710. How do you select comparable deals for a precedent transactions analysis?

You screen for deals involving targets in the same industry and of broadly similar size, business model and geography, ideally completed (announced) recently so the multiples reflect comparable market and financing conditions. You prefer deals with disclosed financials so the multiple can actually be computed, and you note the deal type (strategic vs financial buyer, distressed vs competitive auction) because that drives the premium. Older deals (more than ~3-5 years) and deals from very different market cycles are usually excluded or footnoted.

Q711. Why might a strategic buyer's deal multiple be a poor comp for a financial-sponsor deal?

A strategic (corporate) acquirer can realize cost and revenue synergies by combining the target with its existing operations, so it can justify paying a higher multiple. A financial sponsor (PE firm) generally has no operating synergies and is constrained by leverage and target IRR, so it typically pays less. Using a synergy-rich strategic deal to value a sponsor-led transaction will overstate value, so bankers separate the two buyer types when presenting the range.

Q712. What is a "football field" chart in a valuation?

A football field is a horizontal bar chart that summarizes the valuation ranges produced by each methodology (DCF, trading comps, precedent transactions, LBO analysis, 52-week trading range, analyst price targets) on a single page, so you can see where the ranges overlap. Each method is one bar showing its low-to-high value

(e.g. enterprise value or per-share price), and the overlapping zone suggests a defensible valuation range. It is the standard one-slide "valuation summary" that bankers present to a board or client.

Q713. What are the three main valuation methodologies banks present, and how do they typically rank?

The three core methods are discounted cash flow (DCF), comparable companies (trading comps), and precedent transactions. As a rough generalization, DCF often produces the highest value (because it can capture full long-term growth and an optimistic terminal value), precedent transactions sit in the middle (they bake in a control premium and synergies), and trading comps tend to be lowest (minority, no-control market value). This ordering is not a rule, but it is the common pattern interviewers expect, and a banker should always be ready to explain why their specific case deviates.

Q714. Why does a DCF often give the highest valuation?

A DCF captures the full intrinsic value of all projected future cash flows, including years of growth and a terminal value that often represents 60-80% of total value, so small changes in the growth or discount-rate assumptions move value a lot. Because management projections tend to be optimistic and the terminal value is sensitive to the perpetuity growth or exit multiple chosen, the DCF range can run above market-based methods. That said, it is also the most assumption-dependent method, so a disciplined banker stress-tests it with sensitivity tables.

Q715. Why would an acquirer pay a premium over the target's current share price?

The acquirer pays a premium because it expects to be worth more than the standalone trading value: it gains control, can capture synergies (cost savings and revenue uplift), and must offer enough to persuade shareholders to give up their stock. In a competitive auction the premium also reflects the price needed to beat rival bidders. The premium is justified only if the present value of synergies plus the standalone value exceeds the price paid; otherwise the acquirer overpays and destroys value.

Q716. How do you justify an acquisition premium using synergies?

You estimate the after-tax run-rate synergies (e.g. cost savings from headcount, procurement, facility consolidation, plus any revenue synergies), discount their present value net of integration costs, and compare that to the dollar premium being paid. If PV of synergies exceeds the premium, the deal creates value for the acquirer's shareholders even at a premium price. Example: paying a \$200m premium is justified if \$40m/year of after-tax synergies discounted as a perpetuity at 10% is worth \$400m in PV.

Q717. What is the difference between cost synergies and revenue synergies, and which do investors trust more?

Cost synergies are savings from eliminating duplicate functions, consolidating procurement, and shutting overlapping facilities; revenue synergies are incremental sales from cross-selling, broader distribution, or pricing power. Investors and analysts give far more credit to cost synergies because they are more controllable and quicker to realize, whereas revenue synergies are harder to deliver and often fail to materialize. Bankers typically present cost synergies as the credible base and treat revenue synergies as upside.

Q718. Walk through the IPO process end to end.

The company first runs a "beauty parade" and awards a mandate to lead underwriters (bookrunners), who then conduct due diligence and draft the registration document (S-1 in the US, DRHP in India). After regulator review (SEC / SEBI comments), the company markets the deal via a roadshow to institutional investors, the bankers build a book of demand at various prices (book-building), set the final offer price, allocate shares, list and trade, and may stabilize the price using a greenshoe over-allotment option. Insiders are usually subject to a lock-up (often 180 days) before they can sell.

Q719. What is an S-1 versus a DRHP?

The S-1 is the registration statement filed with the US SEC for an IPO; it contains the business description, risk factors, MD&A, audited financials and use-of-proceeds disclosure. The DRHP (Draft Red Herring Prospectus) is the Indian equivalent filed with SEBI and the stock exchanges; "red herring" refers to the fact that the price/number of shares is not finalized at filing. Both are the core disclosure documents that let investors evaluate the offering; the final versions are the "424B" prospectus (US) and RHP/Prospectus (India).

Q720. What happens during an IPO roadshow and book-building?

In the roadshow, management and bankers present the investment case to institutional investors over one to two weeks, while the syndicate desk collects indications of interest at different price levels. Book-building is the process of aggregating that demand into a "book" that shows how many shares investors will buy at each price, which lets the bankers gauge the demand curve. The final price is then set based on where the book is well-covered, balancing maximizing proceeds against leaving enough demand for a healthy aftermarket.

Q721. What is a greenshoe / over-allotment option?

The greenshoe (over-allotment option) lets the underwriters sell up to ~15% more shares than originally offered and then buy them back (or exercise the option) depending on aftermarket demand. If the stock trades up, underwriters exercise the option and the issuer sells the extra shares; if it trades down, they buy shares in the open market to cover the short, which supports (stabilizes) the price. It is named after the Green Shoe Manufacturing Company, the first issuer to use it.

Q722. What is an IPO lock-up period and why does it exist?

A lock-up is a contractual restriction (commonly 180 days in the US; for Indian IPOs promoter and pre-IPO shareholder lock-ins are set by SEBI rules) that prevents insiders and pre-IPO holders from selling their shares immediately after listing. It exists to prevent a flood of insider selling that would crater the price and signal a lack of confidence right after the offering. Lock-up expiry dates are watched because they can create selling pressure when large blocks become free to trade.

Q723. Why are IPOs often underpriced, and what is the evidence?

Underpricing means the offer price is set below the price the stock trades at on day one, so the stock "pops"; on average first-day returns are positive (historically high-single to double-digit percent). It happens because of information asymmetry (issuers and bankers set a price that ensures the book is covered and rewards institutional investors for revealing demand), risk aversion by underwriters who want a successful deal, and the desire to create positive momentum and goodwill. The cost to the issuer is "money left on the table" — the difference between proceeds at the offer price and what fuller pricing would have raised.

Q724. What is the difference between primary and secondary shares in an IPO?

Primary shares are newly issued by the company, so the cash raised goes to the company's balance sheet (the share count increases, funding growth or debt paydown). Secondary shares are existing shares sold by current holders (founders, PE/VC), so the proceeds go to those selling shareholders, not the company, and there is no new dilution from those shares. A single IPO can include both; an offer that is mostly secondary (in India, an "Offer for Sale") signals insiders cashing out rather than the company raising capital.

Q725. What is the difference between book-building and fixed-price issues in India?

In a fixed-price issue the company sets a single price in the prospectus before the offer opens, and investors apply at that price. In book-building the issuer specifies a price band (floor and cap), institutions and other investors bid within the band, and the final cut-off price is discovered from the book; most large Indian IPOs use the 100% book-built route. SEBI requires demand details to be disclosed, and retail investors can typically bid "at cut-off" to accept whatever final price is set.

Q726. Who are anchor investors in an Indian IPO?

Anchor investors are large institutional investors (QIBs) who are allotted shares one day before the IPO opens to the public, at a price within the band, to lend credibility and demand confidence to the offering. SEBI requires anchors to commit for a lock-in (currently 50% of anchor shares for 90 days and the remainder for 30 days from allotment), so they cannot immediately flip. Their participation, and the marquee names involved, is a signal retail and other investors watch closely.

Q727. What are the QIB, NII and retail categories in an Indian book-built IPO?

SEBI mandates reservation buckets: QIBs (Qualified Institutional Buyers — mutual funds, banks, FIIs, insurers) get up to 50%, Non-Institutional Investors (NIIs/HNIs applying for more than Rs 2 lakh) get at least 15%, and Retail Individual Investors (applications up to Rs 2 lakh) get at least 35% (these are the proportions for the standard profitable-company route). Retail and NII over-subscription is allotted by lottery/proportionate rules when demand exceeds supply. The split ensures broad public participation rather than the deal going entirely to institutions.

Q728. What is a follow-on public offering (FPO) and how does it differ from an IPO?

A follow-on (or secondary) public offering is a sale of shares by a company that is already listed, as opposed to an IPO, which is the first time shares are sold to the public. An FPO can be dilutive (new primary shares issued, raising cash for the company) or non-dilutive (existing holders selling). Because there is already a market price, an FPO is usually priced at a small discount to the prevailing trading price rather than discovered from scratch.

Q729. What is a rights issue and what is the theoretical ex-rights price (TERP)?

A rights issue offers existing shareholders the right to buy new shares pro-rata, usually at a discount to the market price, preserving their ownership percentage if they subscribe. The theoretical ex-rights price is the blended price after the new shares are added: $TERP = (N_{old} * P_{old} + N_{new} * P_{sub}) / (N_{old} + N_{new})$. Example: 1-for-1 rights at Rs 80 when the stock is Rs 120 gives $TERP = (1120 + 180)/2 = Rs 100$, and the value of one right = market price minus TERP = $120 - 100 = Rs 20$.

Q730. How does a rights issue cause dilution, and does subscribing protect a shareholder?

A rights issue increases the share count, so EPS and each share's claim on the company fall (the price drops toward TERP), and any shareholder who does not subscribe is diluted both in ownership percentage and in value as the share price re-bases below the old price. A shareholder who fully subscribes maintains their ownership percentage and is economically neutral — the loss on existing shares from the lower TERP is offset by buying new shares cheaply. Those who cannot subscribe may sell their rights (where renounceable) to recover the rights' value.

Q731. What is a QIP (Qualified Institutions Placement) in India and why is it popular?

A QIP is a SEBI mechanism that lets a listed company raise capital quickly by issuing shares (or convertibles) only to Qualified Institutional Buyers, without the lengthy documentation of a public offering. It is popular because it is fast, avoids extensive regulatory filing and roadshows, and the floor price is formula-based (broadly the average of the last two weeks' weekly high-low of closing prices), with a permitted discount of up to 5%. It is the Indian analog of a US private placement / registered direct to institutions.

Q732. What is the difference between Equity Capital Markets (ECM) and Debt Capital Markets (DCM)?

ECM teams help companies raise equity — IPOs, follow-ons, rights issues, QIPs, convertibles and block trades — which gives up ownership but adds no repayment obligation. DCM teams help companies raise debt — bonds, notes, and structured debt — which must be repaid with interest but does not dilute owners and the interest is usually tax-deductible. Equity is more sensitive to market sentiment and valuation, while debt issuance hinges on credit ratings, interest-rate conditions, and credit spreads.

Q733. Walk through the bond issuance process and the role of the credit rating.

The issuer mandates underwriters, who help structure the bond (size, tenor, seniority, covenants), obtain a credit rating from agencies (S&P/Moody's/Fitch globally; CRISIL/ICRA/CARE in India), and prepare an offering memorandum/prospectus. The rating is central because it largely determines the credit spread over the benchmark rate (e.g. US Treasuries or the G-sec yield), and thus the coupon the issuer must pay; a higher rating means a lower spread and cheaper funding. Bankers then market the bond, build a book of orders at a given spread, set the coupon/price, and allocate the bonds.

Q734. How is a bond's coupon set, and what does "yield to issuer" mean?

The coupon is set so the bond can be sold at (or near) par given prevailing yields: coupon roughly equals the benchmark risk-free yield plus the credit spread appropriate for the issuer's rating and tenor. If demand is strong, the book lets bankers tighten the spread (lower yield), reducing the issuer's cost; the all-in yield the company pays — including the spread and any fees — is effectively its cost of debt for that issue. For example, a BBB issuer pricing 150 bps over a 4.0% government benchmark would carry roughly a 5.5% coupon if issued at par.

Q735. What is a convertible bond and why do issuers use them?

A convertible bond is debt that the holder can convert into a set number of the issuer's shares, so it combines a bond with an embedded equity call option. Issuers use them because the equity option lets them pay a lower coupon than on straight debt (the investor accepts less yield for the upside), and they can effectively sell equity at a premium to the current price (the conversion price is set above today's stock price). The trade-off is potential future dilution if the stock rises and bonds convert; issuers often pair them with call-spread overlays to raise the effective conversion price.

Q736. What is the underwriting spread (gross spread) in an offering?

The gross spread is the underwriters' total fee: the difference between the price investors pay (offer price) and the net proceeds the issuer receives per share. For example, a \$20.00 offer with a 7% gross spread nets the issuer \$18.60 per share, with \$1.40 going to the syndicate, split into the management fee, underwriting fee, and selling concession. US IPOs have historically clustered around a 7% gross spread; investment-grade bond and large follow-on fees are much smaller (often well under 1%).

Q737. What is the difference between a bought deal and a best-efforts offering, and what are SPACs?

In a bought deal the underwriter buys the entire issue from the company at a set price and resells it, taking the placement risk (common for follow-ons and bonds); in a best-efforts deal the bank only agrees to use its best efforts to sell shares and bears no risk for unsold stock, which is typical for riskier or smaller IPOs. A SPAC (Special Purpose Acquisition Company) is a shell company that IPOs to raise a cash trust, then merges with a private target to take it public; pros are speed, price certainty and the ability to share forward projections, while cons are sponsor dilution (the ~20% "promote"), redemptions that can drain the trust, and a poor average post-merger return record.

Q738. How does a banker pitch a financing, and what are league tables?

In a financing pitch the banker shows the client the optimal structure (equity vs debt vs convertible), the likely amount and pricing, market windows and investor demand, comparable recent deals, and the bank's relevant track record and distribution capability. To prove credibility the pitch cites league tables — published rankings (e.g. by Bloomberg, Dealogic, Refinitiv, Prime Database in India) that rank banks by deal volume or value in a product/region over a period. A top league-table position is used as a marketing credential to win the mandate, even though rankings can be sliced to favor the pitching bank.

G. Portfolio Management & Investments

G1. Modern Portfolio Theory (Q739-Q768)

Q739. What is the expected return of a portfolio, and how is it calculated?

A portfolio's expected return is the weighted average of the expected returns of its holdings: $E[R_p] = \sum(w_i * E[R_i])$, where w_i are the portfolio weights (summing to 1). For example, 60% in equities at 12% and 40% in bonds at 6% gives $E[R_p] = 0.6*12\% + 0.4*6\% = 9.6\%$. Note that return aggregates linearly, but risk does not.

Q740. Write the formula for the variance of a two-asset portfolio.

$\text{var}(R_p) = w_1^2*s_1^2 + w_2^2*s_2^2 + 2*w_1*w_2*s_1*s_2*\text{corr}$, where s_1, s_2 are the standard deviations and corr is the correlation between the two assets. The portfolio standard deviation is $\sigma_p = \sqrt{\text{var}(R_p)}$. The cross term $2*w_1*w_2*s_1*s_2*\text{corr}$ is the covariance contribution, and it is what makes portfolio risk lower than the weighted average of individual risks whenever $\text{corr} < 1$.

Q741. Work a quick two-asset risk/return example with $\text{corr} = 0.3$.

Assets: A has $E[R]=10\%$, $s=20\%$; B has $E[R]=6\%$, $s=10\%$; weights 50/50; $\text{corr}=0.3$. Return: $E[R_p] = 0.5*10\% + 0.5*6\% = 8\%$. Variance: $0.25*0.04 + 0.25*0.01 + 2*0.5*0.5*0.20*0.10*0.3 = 0.01 + 0.0025 + 0.003 = 0.0155$, so $\sigma_p = \sqrt{0.0155} = 12.45\%$. The simple weighted-average risk would have been 15%, so diversification saved about 2.55 percentage points.

Q742. What is the difference between covariance and correlation?

Covariance measures how two assets move together in absolute (unit-dependent) terms: $\text{cov}(A,B) = \text{corr} * s_A * s_B$, and can range over any value. Correlation is the standardized version, $\text{corr} = \text{cov}(A,B) / (s_A*s_B)$, always bounded between -1 and +1, which makes it comparable across asset pairs. Covariance enters the variance formula directly, but correlation is easier to interpret because it strips out the scale of each asset's volatility.

Q743. How does correlation drive the benefit of diversification?

The lower the correlation between assets, the smaller the cross (covariance) term in the variance formula, so portfolio volatility falls below the weighted average of individual volatilities. As long as $\text{corr} < 1$, combining assets reduces risk; the benefit grows as correlation approaches -1. Diversification is therefore "the only free lunch in finance" because it can lower risk without necessarily lowering expected return.

Q744. What happens to two-asset portfolio risk in the special cases $\text{corr} = +1$, 0, and -1?

If $\text{corr} = +1$, risk is just the weighted average, $\sigma_p = w_1*s_1 + w_2*s_2$ (no diversification benefit). If $\text{corr} = 0$, the cross term vanishes and $\sigma_p = \sqrt{w_1^2*s_1^2 + w_2^2*s_2^2}$, giving partial diversification. If $\text{corr} = -1$, risk can be driven to zero with the right weights, because the assets perfectly offset each other. Real-world correlations sit between these extremes, so diversification helps but cannot fully eliminate risk.

Q745. With $\text{corr} = -1$, what weights produce a zero-variance two-asset portfolio?

When $\text{corr} = -1$, set the weights inversely to volatility: $w_1 = s_2/(s_1+s_2)$ and $w_2 = s_1/(s_1+s_2)$. For example with $s_1=20\%$ and $s_2=10\%$, $w_1 = 10/30 = 33.3\%$ in asset 1 and $w_2 = 66.7\%$ in asset 2 gives $\sigma_p = 0$. This is a theoretical edge case; perfect negative correlation rarely exists, but it illustrates the upper bound of diversification.

Q746. Distinguish systematic from unsystematic risk.

Systematic (market or non-diversifiable) risk affects all assets and stems from economy-wide factors like interest rates, inflation, recessions, or geopolitics. Unsystematic (specific, idiosyncratic, or diversifiable) risk is unique to a company or sector, such as a product recall, management fraud, or a plant fire. Total risk = systematic + unsystematic, and only systematic risk is compensated with a higher expected return because the rest can be diversified away.

Q747. Why can diversification not remove systematic risk?

Diversification works by averaging away the independent, asset-specific shocks that partially cancel out across many holdings. Systematic risk is common to every asset, so it does not cancel; adding more stocks just leaves you exposed to the same market-wide factor. As the number of holdings rises, portfolio variance falls toward the average covariance between assets, not toward zero, and that residual floor is the systematic risk.

Q748. How does the number of assets affect portfolio risk?

For an equal-weighted portfolio of N assets, variance is $\text{var}(R_p) = (1/N) \cdot \text{avg_variance} + (1 - 1/N) \cdot \text{avg_covariance}$. As N grows, the first term shrinks toward zero and the variance converges to the average covariance, the systematic floor. Empirically much of the diversifiable risk is removed with roughly 20-30 stocks; adding more gives diminishing benefit. In India, a Nifty 50 index fund captures most of this diversification within a single instrument.

Q749. Write portfolio variance in matrix form and explain the terms.

$\text{var}(R_p) = \mathbf{w}' \cdot \mathbf{\Sigma} \cdot \mathbf{w}$, where $\mathbf{w}$ is the column vector of weights, $\mathbf{w}'$ is its transpose, and $\mathbf{\Sigma}$ is the covariance matrix of asset returns. The diagonal of $\mathbf{\Sigma}$ holds the variances and the off-diagonal entries hold the covariances $\text{cov}(i, j) = \text{corr}_{ij} \cdot s_i \cdot s_j$. This compact form generalizes the two-asset formula to any number of assets and is what optimizers actually compute.

Q750. What is the efficient frontier and what does it represent?

The efficient frontier is the set of portfolios that deliver the maximum expected return for each level of risk (or the minimum risk for each level of return). Plotted in risk-return space it is the upward-sloping upper edge of the feasible region of all portfolio combinations. Any portfolio below the frontier is sub-optimal because you could earn more return for the same risk; rational mean-variance investors hold only frontier portfolios.

Q751. What is the minimum-variance portfolio (MVP)?

The MVP is the single portfolio on the efficient frontier with the lowest possible total risk, sitting at the leftmost tip of the frontier curve. For two assets its weight in asset 1 is $w_1 = (s_2^2 - \text{cov}_{12}) / (s_1^2 + s_2^2 - 2 \cdot \text{cov}_{12})$. It is the starting point of the efficient frontier; portfolios below the MVP (the lower curve) are inefficient because another portfolio offers the same risk with higher return.

Q752. Compute the minimum-variance weights for a simple two-asset case.

With $s_1=20\%$, $s_2=10\%$, $\text{corr}=0.3$, covariance is $\text{cov}_{12} = 0.3 \cdot 0.20 \cdot 0.10 = 0.006$. Then $w_1 = (0.01 - 0.006) / (0.04 + 0.01 - 2 \cdot 0.006) = 0.004 / 0.038 = 10.5\%$, so $w_2 = 89.5\%$. The portfolio tilts heavily toward the lower-volatility asset, which is the intuition behind the MVP minimizing total risk.

Q753. What changes when you add a risk-free asset to the opportunity set?

A risk-free asset (with zero variance and return R_f) lets you form straight-line combinations between R_f and any risky portfolio. The best such line is the one tangent to the efficient frontier, which dominates the curved frontier of risky-only portfolios at every risk level. Investors now hold a mix of the risk-free asset and that single tangency portfolio, achieving better risk-return trade-offs than with risky assets alone.

Q754. What is the Capital Market Line (CML)?

The CML is the straight line from the risk-free rate, tangent to the efficient frontier, representing the best attainable combinations of the risk-free asset and the market (tangency) portfolio. Its equation is $E[R_p] = R_f + ((E[R_m] - R_f)/\sigma_m) * \sigma_p$. The slope is the market price of risk, and only efficient (fully diversified) portfolios lie on the CML. It differs from the SML, which prices individual assets by beta.

Q755. What is the tangency (market) portfolio?

The tangency portfolio is the risky portfolio where the line from R_f just touches the efficient frontier; it has the highest Sharpe ratio of any risky portfolio. Under CAPM assumptions, in equilibrium every investor holds this same portfolio of risky assets, which must therefore be the market portfolio, value-weighted across all risky assets. All efficient portfolios are then a blend of R_f and this single portfolio.

Q756. What is the Sharpe ratio and how does it relate to the CML?

The Sharpe ratio is excess return per unit of total risk: $\text{Sharpe} = (R_p - R_f) / \sigma_p$. It is exactly the slope of the line connecting the risk-free asset to a portfolio, so the CML is the line with the maximum Sharpe ratio, achieved by the tangency portfolio. For example, a fund returning 14% with $R_f=6\%$ and $\sigma=16\%$ has $\text{Sharpe} = (14-6)/16 = 0.5$. Higher Sharpe means better risk-adjusted performance.

Q757. State the two-fund separation theorem.

Two-fund separation says that when a risk-free asset exists, every investor's optimal portfolio is a combination of just two funds: the risk-free asset and the single tangency (market) portfolio. Risk tolerance only changes the split between these two, not the composition of the risky fund. A conservative investor lends (holds more R_f), while an aggressive one can borrow at R_f to lever up the tangency portfolio beyond 100%.

Q758. How do investors with different risk appetites use the CML?

They all hold the same tangency portfolio but vary the weight placed on it versus the risk-free asset. A point left of the tangency portfolio means lending part of the money at R_f (lower risk and return); a point to the right means borrowing at R_f to invest more than 100% in the tangency portfolio (leverage, higher risk and return). This is the practical consequence of two-fund separation.

Q759. Why is equal-weighting sometimes preferred over mean-variance optimization?

Equal-weighting ($1/N$) ignores the estimated return and covariance inputs entirely, so it sidesteps the estimation error that plagues optimizers. Empirical studies (e.g., DeMiguel, Garlappi, Uppal) found $1/N$ often matches or beats "optimal" mean-variance portfolios out of sample. It also gives a natural small-cap tilt and forces contrarian rebalancing, though it can over-concentrate in volatile names since it ignores risk.

Q760. Contrast market-cap weighting with equal weighting.

Market-cap weighting holds each asset in proportion to its market value, so it is self-rebalancing, low-turnover, and reflects the consensus "market portfolio" (e.g., Nifty 50, S&P 500). Equal weighting holds the same amount in each, which tilts toward smaller and value names and requires periodic rebalancing as prices drift. Cap-weighting can become top-heavy in a few mega-caps, while equal-weight spreads exposure but trades more.

Q761. What is risk parity and how does it differ from cap-weighting?

Risk parity allocates so that each asset (or asset class) contributes equally to total portfolio risk, rather than equal capital or cap weights. Because bonds are less volatile than equities, this typically means a large bond allocation, often levered up to reach a target return (as in Bridgewater's All Weather). It aims for more balanced risk exposure across regimes, but depends on stable correlations and can suffer when leverage and rates move against it.

Q762. Why is rebalancing necessary, and what does it accomplish?

Over time, winners grow and losers shrink, so weights drift away from targets and the portfolio's risk profile changes (often becoming more equity-heavy). Rebalancing sells outperformers and buys underperformers back to target, enforcing a disciplined "sell high, buy low" and controlling risk. It can be calendar-based (e.g., quarterly) or threshold-based (e.g., when a weight drifts 5% off target); in India, watch STT and capital-gains tax as rebalancing costs.

Q763. What is estimation error in MPT, and why are the inputs fragile?

Mean-variance optimization needs expected returns, variances, and the full covariance matrix, all estimated from noisy historical data. Expected returns are especially hard to estimate and small input changes produce wildly different "optimal" weights. The optimizer treats estimation noise as if it were signal, piling into assets that merely look attractive in-sample, which is why MPT is fragile and often disappoints out of sample.

Q764. What did Michael Best and Robert Grauer (and Richard Michaud) mean by "error maximization"?

Michaud famously called mean-variance optimizers "estimation-error maximizers" because they overweight assets with overestimated returns or underestimated risk and short those with the opposite errors. Tiny changes in inputs can flip the recommended portfolio dramatically, producing extreme, unstable, and often un-investable weights. Practitioners counter this with constraints, shrinkage estimators (e.g., Ledoit-Wolf), Bayesian methods like Black-Litterman, and resampled efficiency.

Q765. How does the Black-Litterman model address MPT's input fragility?

Black-Litterman starts from a neutral equilibrium (the market-cap portfolio implied by reverse-optimization) rather than raw historical means, then blends in the investor's explicit views with confidence levels. This Bayesian shrinkage produces more stable, intuitive, and diversified weights that don't swing violently with small input changes. It directly fixes the "garbage in, garbage out" problem of feeding noisy expected returns into a naive optimizer.

Q766. Why does diversifying across asset classes help more than within one class?

Assets within a single class (e.g., Indian large-cap equities) tend to be highly correlated, so their idiosyncratic risks cancel but they still rise and fall together with that market. Combining low-correlation classes such as equities, bonds, gold, and real estate lowers the average covariance and pushes the efficient frontier up and to the left. Gold, for instance, often has low or negative correlation with equities and acts as a portfolio hedge during crises.

Q767. What is the case for international diversification, and what is its main limitation?

Holding foreign equities adds assets driven partly by different economies, currencies, and cycles, lowering portfolio correlation and risk (e.g., a global investor pairing US, European, and emerging-market stocks). The main limitation is that correlations tend to spike toward 1 during global crises, exactly when diversification is most needed, as seen in 2008. Currency exposure and home-country constraints (for Indian investors, LRS limits and taxation of foreign funds) also temper the benefit.

Q768. Put it together: build the efficient set for a two-asset stock-bond mix and identify the key portfolios.

Take equities ($E[R]=12\%$, $s=18\%$) and bonds ($E[R]=6\%$, $s=7\%$) with $\text{corr}=0.2$. Varying the equity weight from 0 to 100% traces a curved frontier; the MVP weight in equities is $w_{\text{eq}} = (s_{\text{bond}}^2 - \text{cov}) / (s_{\text{eq}}^2 + s_{\text{bond}}^2 - 2 \cdot \text{cov})$ with $\text{cov} = 0.2 \cdot 0.18 \cdot 0.07 = 0.00252$, giving $w_{\text{eq}} = (0.0049 - 0.00252) / (0.0324 + 0.0049 - 0.00504) = 0.00238 / 0.03226 = 7.4\%$. Adding a risk-free asset at $R_f=5\%$ and drawing the tangent line locates the tangency portfolio (highest Sharpe), and all efficient holdings become blends of R_f and that tangency mix.

G2. CAPM, Factor Models & Performance (Q769-Q798)

Q769. What is the CAPM equation, and what does it say in words?

The Capital Asset Pricing Model states $E[R_i] = R_f + \beta_i * (E[R_m] - R_f)$, where R_f is the risk-free rate, $E[R_m] - R_f$ is the equity (market) risk premium, and β_i measures the asset's sensitivity to market moves. In words: the expected return on an asset equals the risk-free rate plus a reward for the systematic risk it adds to a diversified portfolio. For example, with $R_f = 7\%$ (India 10Y G-Sec proxy), $ERP = 6\%$, and $\beta = 1.2$, $E[R] = 7\% + 1.2 * 6\% = 14.2\%$. Only non-diversifiable (systematic) risk is priced; idiosyncratic risk earns no premium.

Q770. What are the key assumptions behind CAPM?

CAPM assumes investors are rational mean-variance optimizers with homogeneous expectations, can borrow and lend freely at a single risk-free rate, face no taxes or transaction costs, hold the same single-period horizon, and that all assets are perfectly divisible and markets are frictionless and competitive. Under these conditions everyone holds the same risky "market portfolio" combined with the risk-free asset, which is what makes beta the single relevant risk measure. The assumptions are unrealistic, but the model is valued for the clean intuition that only systematic risk is compensated.

Q771. Define beta and give its formula.

Beta measures an asset's systematic risk: how much its return moves with the overall market. $\beta = \text{Cov}(R_i, R_m) / \text{Var}(R_m)$, which equals $\rho(i,m) * \sigma_i / \sigma_m$, where ρ is the correlation. $\beta = 1$ means the asset moves in line with the market; $\beta > 1$ is more volatile/cyclical (e.g., a high-beta auto or metals stock), $\beta < 1$ is defensive (e.g., FMCG/pharma). For instance, if $\text{Cov}(R_i, R_m) = 0.024$ and $\text{Var}(R_m) = 0.020$, $\beta = 1.2$.

Q772. How do you interpret a stock with $\beta = 0.6$ versus $\beta = 1.5$?

A beta of 0.6 means that for a 1% market move the stock tends to move about 0.6% in the same direction — it is defensive and dampens portfolio volatility (typical of utilities, consumer staples). A beta of 1.5 means roughly a 1.5% move per 1% market move — it amplifies both gains and losses and is cyclical (typical of high-leverage or discretionary names). Higher beta implies a higher required return under CAPM, not necessarily a "better" stock.

Q773. What is the beta of cash and of the risk-free asset?

Cash and the risk-free asset have a beta of zero because their returns have essentially no covariance with the market. Under CAPM their required return is just R_f . This matters in portfolio aggregation: holding cash mechanically lowers a portfolio's beta, since portfolio beta is the value-weighted average of component betas, and the cash slice contributes zero.

Q774. How do you estimate beta in practice?

The standard approach regresses the stock's excess returns on the market's excess returns: $R_i - R_f = \alpha + \beta * (R_m - R_f) + e$, and the slope is beta. Practitioners commonly use 2–5 years of monthly data (or ~1–2 years of weekly) against a broad index (S&P 500 globally, Nifty 50 in India). Raw beta is often "adjusted" toward 1 using the Blume formula $\text{adjusted } \beta = 0.67 * \text{raw } \beta + 0.33 * 1$, since betas tend to mean-revert. Estimates are sensitive to the index, frequency, and window chosen.

Q775. What is the difference between the Security Market Line (SML) and the Capital Market Line (CML)?

The CML plots expected return against total risk (standard deviation, σ) and applies only to efficient portfolios — combinations of the risk-free asset and the market portfolio; its slope is the Sharpe ratio of the market. The SML plots expected return against systematic risk (beta) and applies to all assets and portfolios,

efficient or not; its slope is the equity risk premium. The key contrast: CML uses total risk and holds only for efficient portfolios, while SML uses beta and prices every individual security.

Q776. What does it mean for a stock to plot above or below the SML?

The SML gives the CAPM-required return for a given beta. A stock plotting above the SML offers more return than required for its systematic risk — it is undervalued and has positive alpha. A stock below the line offers too little for its beta — overvalued, negative alpha. In equilibrium CAPM says all assets should lie exactly on the SML; deviations are the basis for active stock selection.

Q777. What is Jensen's alpha and what does a positive alpha indicate?

Jensen's alpha is the intercept from the CAPM regression — the return earned above (or below) what CAPM predicts for the portfolio's beta: $\alpha = R_p - [R_f + \beta(R_m - R_f)]$. A positive alpha means the manager or asset delivered excess risk-adjusted return — skill, mispricing, or exposure to factors CAPM omits. Example: a fund returns 16% with beta 1.1 when $R_f = 7\%$ and $R_m = 13\%$; CAPM expects $7\% + 1.1 \times 6\% = 13.6\%$, so $\alpha = +2.4\%$.

Q778. What is the equity risk premium (ERP), and roughly what range is used?

The ERP is the extra return investors demand for holding equities over the risk-free asset: $ERP = E[R_m] - R_f$. It can be estimated historically (average realized equity return minus risk-free), via a forward-looking implied approach (e.g., dividend discount / Gordon model solved for the discount rate), or from surveys. Mature-market ERPs are commonly assumed around 4.5–6%; for India a country/equity risk premium is typically added, pushing total ERP higher (often ~7–9% depending on the source and date).

Q779. What is the difference between levered (equity) beta and unlevered (asset) beta?

Equity beta reflects both business risk and the amplification from financial leverage, while asset (unlevered) beta strips out leverage to capture pure business risk. They are linked by $\beta_{equity} = \beta_{asset} * [1 + (1 - \tau_c) * (D/E)]$ (assuming debt beta ~ 0). More debt raises equity beta because fixed interest claims make residual equity cash flows more volatile. Asset beta lets you compare the underlying riskiness of firms with different capital structures.

Q780. Walk through unlevering and relevering beta for a comparable-company valuation.

To value a private or target firm you take listed comparables' equity betas, unlever each using $\beta_{asset} = \beta_{equity} / [1 + (1 - \tau_c) * (D/E)]$, average the asset betas, then relever to the target's own capital structure with $\beta_{equity} = \beta_{asset} * [1 + (1 - \tau_c) * (D/E_{target})]$. Example: comp beta = 1.3, $D/E = 0.5$, $\tau_c = 25\%$ gives $\beta_{asset} = 1.3 / [1 + 0.75 \times 0.5] = 1.3 / 1.375 = 0.945$; *relevering to a target D/E of 1.0 gives $0.945[1 + 0.75 \times 1.0] = 1.65$* . This isolates business risk before reapplying the target's leverage.

Q781. Should you assume debt beta is zero, and when does it matter?

Assuming debt beta = 0 is a common simplification (the Hamada relation), reasonable for investment-grade firms whose debt is nearly risk-free. But for highly levered or distressed companies, debt carries real market risk (debt beta > 0), and ignoring it overstates the asset beta and miscounts how risk shifts between debt and equity. A more complete relever uses $\beta_{asset} = [E/(E+D)] * \beta_{equity} + [D/(E+D)] * \beta_{debt}$, allocating systematic risk across both claims.

Q782. What is the single-index (market) model, and how does it differ from CAPM?

The market model is the empirical regression $R_i = \alpha_i + \beta_i R_m + e_i$, decomposing a security's return into a market-driven component and a firm-specific residual. It is a statistical/return-generating model that simplifies covariance estimation (only betas and one market variance are needed). CAPM is an equilibrium model predicting expected returns; it implies the intercept (in excess-return form) should be zero. The market model describes how returns behave; CAPM asserts what expected returns must be in equilibrium.

Q783. Under the single-index model, how is total risk decomposed?

Total variance splits into systematic and unsystematic parts: $\text{Var}(R_i) = \beta_i^2 * \text{Var}(R_m) + \text{Var}(e_i)$. The first term is market (systematic) risk, the second is firm-specific (diversifiable) risk. As you add securities, the residual terms wash out and portfolio risk converges toward the systematic component. This is why diversification eliminates idiosyncratic risk but not market risk — and why only beta is priced.

Q784. What is Arbitrage Pricing Theory (APT) and how does it differ from CAPM?

APT, developed by Ross, says expected return is a linear function of multiple systematic risk factors: $E[R_i] = R_f + \beta_{i1}RP_1 + \beta_{i2}RP_2 + \dots + \beta_{ik}RP_k$, where each β is the sensitivity to a factor and RP is that factor's risk premium. Unlike CAPM, it does not require the market portfolio, mean-variance investors, or normal returns — it rests on a no-arbitrage condition. Its weakness is that the theory does not specify which factors matter; they must be chosen empirically (e.g., inflation, industrial production, term and credit spreads).

Q785. What is the Fama-French three-factor model?

It extends CAPM with two empirically motivated factors: $E[R_i] - R_f = \beta_m(R_m - R_f) + s*SMB + h*HML$. SMB ("small minus big") captures the size premium — small caps historically outperforming large caps — and HML ("high minus low") captures the value premium — high book-to-market (value) stocks outperforming low book-to-market (growth) stocks. The model explains the cross-section of returns far better than CAPM alone, suggesting size and value proxy for systematic risks CAPM misses.

Q786. What did Fama and French add in the five-factor model?

The 2015 five-factor model adds profitability and investment to the three factors: $R_m - R_f$, SMB, HML, RMW, and CMA. RMW ("robust minus weak") is the return spread of firms with robust versus weak operating profitability; CMA ("conservative minus aggressive") is the spread of firms that invest conservatively versus aggressively. The intuition is that high-profitability, low-investment firms tend to earn higher returns. Notably, once RMW and CMA are included, HML often becomes redundant in explaining average returns.

Q787. What is the momentum factor and the Carhart four-factor model?

Momentum is the tendency of recent winners (past ~3–12 month returns) to keep outperforming recent losers over the next several months. Carhart (1997) added this as a fourth factor (UMD or WML, "up minus down"/"winners minus losers") to the Fama-French three-factor model: $\dots + \beta_m(R_m - R_f) + s*SMB + h*HML + w*WML$. It is widely used to evaluate fund performance because much apparent "skill" can be explained by passive exposure to the momentum factor.

Q788. What is Roll's critique of CAPM?

Roll (1977) argued that CAPM is not genuinely testable because the true "market portfolio" must include all risky assets — equities, bonds, real estate, human capital, etc. — which is unobservable. Any test uses a proxy (like the S&P 500 or Nifty), so a rejection of CAPM may simply mean the proxy is mean-variance inefficient, not that the model is wrong; conversely, if the proxy is efficient by construction, the SML relation holds tautologically. In short, tests of CAPM are really joint tests of the model and the chosen market proxy.

Q789. What is the low-beta (low-volatility) anomaly?

Empirically, low-beta stocks have delivered higher risk-adjusted returns than CAPM predicts, while high-beta stocks have underperformed — the realized SML is flatter than theory says. Explanations include leverage constraints (investors who cannot borrow bid up high-beta stocks to chase returns, depressing their future returns) and behavioral preferences for "lottery"-like volatile stocks. This anomaly underpins low-volatility and minimum-variance "smart beta" strategies.

Q790. Beyond Roll's critique, what are the main empirical failures of CAPM?

Studies find the relationship between beta and average returns is too flat (high-beta stocks earn less than predicted), and several characteristics — size, value (book-to-market), momentum, profitability, and low volatility — predict returns even after controlling for beta, which CAPM says should not happen. Fama and French's work showed beta adds little explanatory power once size and value are included. These findings motivated multifactor models as more complete descriptions of expected returns.

Q791. Define the Sharpe ratio and explain what it measures.

The Sharpe ratio is excess return per unit of total risk: $\text{Sharpe} = (R_p - R_f) / \sigma_p$, where σ_p is the portfolio's standard deviation. It tells you how much reward you earn for each unit of total volatility, making it useful for comparing standalone portfolios or for an investor's entire wealth. Example: a fund returning 14% with $\sigma = 12\%$ when $R_f = 7\%$ has $\text{Sharpe} = (14 - 7)/12 = 0.58$. Higher is better; it is the slope of the line from the risk-free asset through the portfolio.

Q792. Define the Treynor ratio and contrast it with Sharpe.

The Treynor ratio is excess return per unit of systematic risk: $\text{Treynor} = (R_p - R_f) / \beta_p$. It uses beta rather than total volatility, so it rewards return relative to market risk only. Use Treynor when the portfolio is one of several well-diversified sleeves within a larger fund (idiosyncratic risk already diversified away, so only beta matters); use Sharpe when the portfolio represents the investor's total holdings and total risk is what they bear. Example: $R_p = 14\%$, $R_f = 7\%$, $\beta = 1.1$ gives $\text{Treynor} = 7/1.1 = 6.36\%$.

Q793. What is the Information Ratio and how is it used?

The Information Ratio measures active return per unit of active risk: $\text{IR} = (R_p - R_b) / \text{TE}$, where R_b is the benchmark return and TE is the tracking error (the standard deviation of active returns $R_p - R_b$). It gauges the consistency and quality of a manager's outperformance relative to a benchmark. An IR around 0.5 is good and 1.0 is excellent. Example: active return of 3% with tracking error of 5% gives $\text{IR} = 0.6$.

Q794. What is tracking error, and what drives it?

Tracking error is the standard deviation of the difference between a portfolio's returns and its benchmark's returns: $\text{TE} = \text{stdev}(R_p - R_b)$. It quantifies how closely a portfolio follows its benchmark — a pure index fund targets near-zero TE, while a high-conviction active fund may run 6–10%. It is driven by active bets: sector tilts, factor exposures, position concentration, and cash holdings. Tracking error is the denominator of the Information Ratio and a core constraint in active mandates.

Q795. What is the Sortino ratio and why use it instead of Sharpe?

The Sortino ratio replaces total volatility with downside deviation, penalizing only returns below a target (often R_f or zero): $\text{Sortino} = (R_p - \text{target}) / \text{downside_deviation}$. The rationale is that investors dislike downside volatility, not upside, so penalizing both (as Sharpe does) can unfairly mark down strategies with large favorable swings. It is especially relevant for asymmetric or option-like return profiles where upside and downside risk differ materially.

Q796. What is M-squared (Modigliani-Modigliani) and what advantage does it offer?

M-squared restates the Sharpe ratio as a return figure by levering or de-levering the portfolio with the risk-free asset until its volatility matches the market's, then comparing that adjusted return to the market's: $M2 = R_f + \text{Sharpe}_p * \sigma_m$. Because the answer is in percentage points (e.g., "the fund is 1.8% better than the market on a risk-adjusted basis"), it is more intuitive to communicate than a raw Sharpe number, while preserving the same ranking.

Q797. In active management, what is the distinction between alpha and beta?

Beta is the return earned from passive exposure to systematic risk factors (the market, or size/value/momentum tilts) — cheap and replicable via index or factor funds. Alpha is the genuine value a manager adds beyond those exposures — security selection or timing skill that cannot be captured passively.

Performance attribution aims to separate the two, since investors should not pay active fees for returns that are really just packaged beta. The Information Ratio and Jensen's alpha are the standard tools for isolating true alpha.

Q798. What is "smart beta," and how does it relate to factor models and benchmark selection?

Smart beta refers to rules-based strategies that systematically tilt toward priced factors — value, size, momentum, quality, low volatility — instead of weighting by market capitalization, aiming to harvest factor premia at low cost. It sits between passive and active: the exposures come from research like Fama-French and Carhart, but they are delivered via transparent, index-like rules (e.g., the Nifty Alpha Low-Volatility 30 or value-weighted indices). This blurs benchmark selection: choosing an appropriate benchmark matters because a "smart beta" fund's apparent alpha against a cap-weighted index may simply be its deliberate factor exposure, which a factor-aware benchmark would reveal as beta, not skill.

G3. Asset & Investment Management (Q799-Q831)

Q799. What is the difference between active and passive investment management?

Passive management seeks to replicate a benchmark index (e.g. S&P 500, Nifty 50) at minimum cost, accepting market returns. Active management uses research, security selection and timing to try to beat the benchmark after fees. The trade-off: active offers the possibility of outperformance (alpha) but charges higher fees and incurs higher turnover/taxes, while passive locks in low cost and near-market returns.

Q800. Explain the "arithmetic of active management" and why most active funds underperform after fees.

William Sharpe's argument: before costs, the average actively managed dollar must earn exactly the market return, because all investors collectively hold the market. Since active investing costs more (fees, trading, taxes) than passive, the average actively managed dollar must underperform the average passive dollar AFTER costs. This is an accounting identity, not an empirical claim: e.g. SPIVA reports consistently show 70-90% of active large-cap funds trailing their index over 10-15 year horizons.

Q801. What does the SPIVA scorecard measure, and what is the typical finding?

SPIVA (S&P Indices Versus Active) tracks the percentage of actively managed funds that underperform their benchmark over various horizons, in the US, India and elsewhere. The consistent finding is that the majority underperform, and the share underperforming rises with the time horizon. Survivorship matters: many laggard funds are merged or closed, so reported active performance is flattered unless you correct for it.

Q802. Define strategic asset allocation versus tactical asset allocation.

Strategic asset allocation (SAA) sets the long-term target mix across asset classes (e.g. 60% equity / 40% bonds) based on the investor's objectives, risk tolerance and capital market expectations; it is the policy anchor. Tactical asset allocation (TAA) makes deliberate short-to-medium-term deviations from those targets to exploit perceived mispricings or cyclical views (e.g. overweighting equities when valuations are cheap). SAA drives most long-run risk and return; TAA is an active overlay that can add or subtract value.

Q803. What did the Brinson, Hood and Beebower studies conclude about asset allocation?

The 1986/1991 Brinson studies of US pension plans found that the policy (strategic) asset allocation explained about 90% of the variability of returns OVER TIME for a given fund. This is widely (and often loosely) cited as "asset allocation determines 90% of returns." Later work (Ibbotson & Kaplan, 2000) clarified the distinction: ~90% of the time-series variability, but only ~40% of the variation in returns ACROSS funds, and roughly 100% of the average return level.

Q804. Is the claim "asset allocation determines 90% of returns" accurate? Clarify the nuance.

It is a misinterpretation. Brinson's 90% refers to the variance of a single portfolio's returns over time being driven by its policy mix, not to the level of return. Across different funds, security selection and timing explain a larger share of the differences. The correct takeaway: asset allocation dominates the volatility experience and is the single most important decision, but it does not literally "determine 90% of the dollars you make."

Q805. What is an Investment Policy Statement (IPS) and why does it matter?

An IPS is the governing document that formalizes a client's objectives and constraints and the manager's responsibilities, serving as the blueprint for portfolio construction and a discipline against emotional decisions. It anchors the strategic asset allocation, defines the benchmark, and sets rebalancing and review procedures. In stressed markets it is the reference that prevents panic selling or style drift.

Q806. List the standard objectives and constraints captured in an IPS.

Objectives: return requirement and risk tolerance (ability + willingness to bear risk). Constraints, often remembered as "RR-TTLU": Risk, Return, then Time horizon, Taxes, Liquidity, Legal/regulatory, and Unique circumstances (e.g. ESG screens, concentrated stock, religious restrictions like Shariah). A common mnemonic for the five constraints is TTLU.

Q807. How do "ability" and "willingness" to take risk differ in an IPS, and which governs?

Ability to take risk is objective: time horizon, wealth relative to needs, liquidity demands, and income stability. Willingness is psychological and subjective: the client's stated comfort with volatility. When they conflict, the prudent rule is to adopt the LOWER of the two (the more conservative), while educating the client; you never let high willingness override low ability.

Q808. What is calendar rebalancing versus threshold (percentage-of-portfolio) rebalancing?

Calendar rebalancing restores target weights on a fixed schedule (e.g. quarterly or annually) regardless of market moves—simple but ignores intra-period drift. Threshold (or tolerance-band) rebalancing acts only when an asset's weight breaches a set corridor, e.g. a 5% absolute band so a 60% equity target triggers a trade at 55% or 65%. Threshold responds to actual risk drift; a hybrid (check on calendar, trade only if outside the band) is common in practice.

Q809. Why does disciplined rebalancing tend to add value, and what is its hidden cost?

Rebalancing forces "sell high, buy low" by trimming outperformers and topping up laggards, controlling risk and harvesting a small "rebalancing bonus" from mean-reverting, volatile assets. The hidden costs are transaction costs and, in taxable accounts, realized capital gains—so wider bands are used for taxable or illiquid portfolios. In a sustained trending market, frequent rebalancing can hurt by cutting winners too early.

Q810. What is a style box and how is it constructed?

Popularized by Morningstar, the equity style box is a 3x3 grid: the vertical axis is market capitalization (large / mid / small) and the horizontal axis is style (value / blend / growth). It lets investors see a fund's footprint and avoid unintended overlap or factor bets. The fixed-income version uses credit quality (high/medium/low) by interest-rate sensitivity (short/intermediate/long duration).

Q811. Distinguish value and growth investing styles.

Value investing buys stocks trading cheap relative to fundamentals—low P/E, low P/B, high dividend yield—betting on mean reversion or mispricing (Graham/Buffett lineage). Growth investing buys companies with high expected earnings/revenue growth, accepting rich multiples in anticipation of compounding. Value has a long-run premium in the Fama-French framework (the HML factor), though it underperformed growth through much of the 2010s.

Q812. What are the quality and momentum equity factors?

Quality favours firms with strong, stable profitability (high ROE/ROIC), low leverage, stable earnings and good accruals—it tends to be defensive. Momentum buys recent winners and sells recent losers, typically using the trailing 12-month return excluding the most recent month (12-1), exploiting return persistence. Momentum delivers high returns but suffers occasional sharp "crashes" during sharp market reversals, so it is often paired with low-volatility or quality screens.

Q813. What is GARP and the dividend strategy?

GARP (Growth At a Reasonable Price) blends growth and value—seeking above-average growth but only at undemanding valuations, often screened with the PEG ratio (P/E divided by earnings growth rate; PEG below ~1 is attractive). Dividend strategies target high or, better, growing and sustainable dividends as a source of income and a quality signal; "dividend aristocrats" are firms with decades of consecutive dividend increases. Watch payout ratios—an unsustainably high yield can signal distress (a value trap).

Q814. Explain bond ladder, barbell, and bullet strategies.

A ladder spreads bonds evenly across maturities (e.g. equal amounts maturing each year for 10 years), giving steady reinvestment and reduced timing risk. A bullet concentrates maturities around a single target date—useful for funding a known future liability. A barbell combines short and long maturities while avoiding the middle, giving liquidity at the short end plus yield/convexity at the long end; it outperforms a bullet of equal duration when the yield curve flattens non-parallel.

Q815. What is duration timing and credit strategy in fixed income?

Duration timing (rate anticipation) lengthens portfolio duration when you expect yields to fall (to maximize price gains) and shortens it when you expect yields to rise. Price sensitivity is approximated by: percent price change $\approx$ -ModifiedDuration * change_in_yield + 0.5Convexity(change_in_yield)². Credit strategy seeks excess return from spread—buying lower-rated or out-of-favour issuers expected to tighten or be upgraded, accepting default and liquidity risk.

Q816. Compare mutual funds, ETFs, and index funds on structure and cost.

An index fund is any fund (mutual or ETF) that passively tracks an index. A traditional mutual fund is bought/sold once daily at NAV directly with the fund; an ETF trades intraday on an exchange like a stock at a market price that can sit at a premium/discount to NAV. ETFs generally have lower expense ratios and use an in-kind creation/redemption mechanism that makes them more tax-efficient than equivalent mutual funds.

Q817. Why are ETFs typically more tax-efficient than mutual funds?

ETF shares are created and redeemed "in kind" with authorized participants—the fund hands over low-basis securities rather than selling them, so it rarely realizes capital gains that must be distributed to holders. A traditional open-end fund must sell securities to meet redemptions, triggering taxable distributions to all remaining holders. Note: this in-kind advantage is a US feature; in India both ETFs and mutual funds distribute, and taxation is at the investor's hands on sale.

Q818. What is tracking error and what drives it in an index product?

Tracking error is the standard deviation of the difference between the fund's return and its benchmark's return—a measure of how faithfully it replicates. It is driven by the expense ratio, cash drag, sampling vs full replication, transaction costs, securities-lending income, and dividend reinvestment timing. A good large-cap index fund may have tracking error well under 0.5% annually; higher tracking error signals either skill (for active) or inefficiency (for passive).

Q819. Distinguish open-end and closed-end funds.

An open-end fund continuously issues and redeems shares at NAV, so its size fluctuates with flows and it always trades at NAV. A closed-end fund (CEF) raises a fixed pool of capital via an IPO, then trades on an

exchange; its market price can deviate—often at a persistent discount to NAV—driven by sentiment, liquidity and distribution policy. CEFs can stay fully invested in illiquid assets without redemption pressure, which is why the structure suits less-liquid strategies.

Q820. Describe long/short equity and equity market-neutral hedge fund strategies.

Long/short equity holds long positions in expected winners and short positions in expected losers, typically keeping a net long bias and aiming to profit from stock selection on both sides. Equity market-neutral tightens this to roughly zero net market exposure ($\beta \approx 0$), often dollar- and sector-neutral, so returns come almost purely from the long-short spread (alpha), not market direction. Market-neutral has lower volatility but relies heavily on leverage to generate meaningful returns.

Q821. What is global macro and managed futures (CTA) investing?

Global macro takes directional or relative-value positions across equities, rates, currencies and commodities based on top-down macroeconomic views, using cash and derivatives, often with leverage (e.g. shorting a currency on a balance-of-payments thesis). CTAs/managed futures are usually systematic trend-followers trading futures across asset classes; they tend to be long volatility and provide "crisis alpha," performing well in sustained trends like 2008 or 2022. Both can be highly diversifying because their returns have low correlation to equities.

Q822. Explain event-driven, merger arbitrage, and distressed strategies.

Event-driven strategies profit from corporate events—mergers, spin-offs, restructurings—where the catalyst, not the market, drives returns. Merger (risk) arbitrage typically buys the target and shorts the acquirer in a stock deal, capturing the spread between the deal price and current price as compensation for deal-break risk. Distressed investing buys deeply discounted debt/equity of firms near or in bankruptcy, aiming to profit from restructuring or recovery—illiquid, legally intensive, and long-horizon.

Q823. What is relative value, and how do quant strategies fit in?

Relative value seeks to profit from price discrepancies between related securities (e.g. convertible arbitrage, fixed-income arbitrage, capital-structure arb), typically hedging out market direction and earning small, leveraged spreads as mispricings converge. Quant strategies use systematic, data-driven models—statistical arbitrage, factor models, machine learning—to trade large baskets, relying on diversification across many small bets. Both are vulnerable to crowding and deleveraging shocks (e.g. the August 2007 "quant quake").

Q824. Explain the 2-and-20 fee model and a high-water mark.

"2-and-20" is the classic hedge fund/PE fee: a 2% annual management fee on assets plus a 20% performance (incentive) fee on profits. A high-water mark ensures the manager earns the 20% only on NEW net profits above the highest prior NAV—so after a loss, gains that merely recover the loss are not charged again. Example: a \$100m fund rising to \$120m earns $0.2 \times \$20\text{m} = \4m incentive; if it then falls to \$110m and recovers to \$120m, no new incentive is due until NAV exceeds \$120m.

Q825. What is a hurdle rate, and how do lock-ups and gates relate to liquidity?

A hurdle rate is a minimum return (e.g. a cash benchmark or fixed %) that must be cleared before the performance fee applies. A lock-up is a period (often 1-2 years) during which investors cannot redeem; gates limit the percentage of the fund that can be redeemed in any single period. These illiquidity terms let managers hold less-liquid, higher-return positions without forced selling—the investor is, in effect, paid an illiquidity premium for accepting them.

Q826. How do leverage and short selling amplify hedge fund returns and risks?

Leverage (borrowing or derivatives) multiplies both gains and losses on a position; gross exposure of 200/100 (200% long, 100% short) means 300% gross on 100% capital. Short selling sells borrowed securities to profit from declines, but losses are theoretically unlimited as the price can rise indefinitely, and

shorts face borrow costs and squeeze risk. The combination is why even market-neutral funds can blow up: small adverse moves are magnified, and margin calls force deleveraging at the worst time.

Q827. What is portfolio risk budgeting?

Risk budgeting allocates a portfolio's total risk (rather than its capital) across positions or strategies, recognizing that a 5% capital weight in a volatile asset may consume far more than 5% of total risk. A common metric is each position's marginal/component contribution to total volatility (sum of component contributions = total portfolio volatility). Risk parity is the limiting case—sizing assets so each contributes equal risk, which typically means leveraging low-volatility bonds up to equity-like risk.

Q828. Outline ESG integration in investment mandates and the main approaches.

ESG integration systematically incorporates Environmental, Social and Governance factors into security analysis and portfolio construction—either as risk mitigation or to capture sustainability-linked alpha. Approaches range from negative screening (excluding tobacco, weapons, coal), to positive/best-in-class tilts, thematic investing, and active stewardship/engagement and proxy voting. Key challenges are inconsistent ESG ratings, greenwashing, and regulation (EU SFDR disclosure tiers; in India SEBI's BRSR Core and ESG mutual-fund sub-category rules).

Q829. Summarize the SEBI mutual fund framework and the PMS/AIF structures in India.

Indian mutual funds are regulated by SEBI under the 1996 MF Regulations, with a three-tier sponsor-trustee-AMC structure, standardized scheme categories, a TER cap, and mandatory daily NAV disclosure; distribution is direct (no commission) or regular plans. PMS (Portfolio Management Services) offers personalized, discretionary portfolios with a minimum investment of Rs 50 lakh. AIFs (Alternative Investment Funds) are pooled vehicles in three categories—Category I (VC, infrastructure, social), Category II (PE, debt funds), Category III (hedge-fund-style, long-short)—with a minimum commitment of Rs 1 crore, aimed at sophisticated investors.

Q830. What is goal-based investing and how does it differ from traditional portfolio theory?

Goal-based investing organizes wealth into separate "buckets," each mapped to a specific goal (retirement, child's education, home) with its own time horizon, target amount and risk level, rather than optimizing one aggregate portfolio against a single risk-tolerance number. It reframes risk as "probability of missing the goal" instead of return volatility, which improves client behaviour and discipline. It draws on mental-accounting and behavioural-portfolio theory and underpins much of modern wealth management and goal-planning advice.

Q831. Explain dollar-cost averaging / SIP and its mathematical edge.

Dollar-cost averaging (in India, a Systematic Investment Plan or SIP) invests a fixed amount at regular intervals regardless of price, so you automatically buy more units when prices are low and fewer when high. Because a fixed sum buys a varying number of units, the average COST per unit is the harmonic mean of prices, which is always less than or equal to the arithmetic average price paid. Example: investing Rs 6,000 monthly at NAVs of 100 then 150 buys $60+40=100$ units for Rs 12,000—average cost Rs 120, below the Rs 125 simple-average price; SIPs also reduce timing risk and instill discipline.

G4. Market Efficiency & Behavioral Finance (Q832-Q856)

Q832. What is the Efficient Market Hypothesis (EMH) in one sentence?

The EMH states that asset prices fully and rapidly reflect all available information, so securities are always fairly priced given what is known and you cannot consistently earn risk-adjusted excess returns (alpha) by trading on that information. It implies that any new information is incorporated into prices almost instantly, leaving no predictable, exploitable patterns. Eugene Fama formalized the idea in 1970.

Q833. What are the three forms of EMH and what information does each say is already priced?

Weak form: prices reflect all past price and volume data, so technical analysis cannot add value. Semi-strong form: prices reflect all publicly available information (financials, news, analyst reports, macro data), so neither technical nor fundamental analysis on public data helps. Strong form: prices reflect all information, public and private/insider, so even insiders cannot beat the market. The forms are nested — strong implies semi-strong implies weak.

Q834. If markets are weak-form efficient, which type of analysis is rendered useless and why?

Technical analysis, because it relies solely on historical prices, volume, and chart patterns to predict future moves. If all past price information is already embedded in the current price, past patterns carry no predictive power and price changes follow a random walk. Most academic evidence supports weak-form efficiency in major liquid markets, including the NSE and BSE for large-cap stocks.

Q835. What does the random walk hypothesis say about stock prices?

It says successive price changes are independent and identically distributed, so the next move is unpredictable and prices wander randomly: $P(t) = P(t-1) + \text{drift} + \text{random error}$. A random walk is a sufficient (though stronger) condition for weak-form efficiency. The practical implication is that the best forecast of tomorrow's price is today's price plus a small expected drift, not any chart pattern.

Q836. What are the implications of EMH for active fund management?

If markets are efficient, active managers cannot reliably beat the market after fees, so the rational strategy is low-cost passive indexing. This underpins the rise of index funds and ETFs; SPIVA scorecards show the majority of active large-cap funds underperform their benchmark over 10-15 years in the US and increasingly in India. EMH does not say active managers are stupid — it says competition among them is exactly what keeps prices efficient.

Q837. Under semi-strong efficiency, can you profit by trading immediately after a company's earnings announcement?

In a strictly semi-strong efficient market, no — prices adjust to public news so fast that by the time you trade, the information is already priced in and you earn only a normal return. In reality, the post-earnings-announcement drift anomaly shows prices keep drifting in the direction of the surprise for weeks, which is evidence against perfect semi-strong efficiency. This drift is one of the most robust and persistent anomalies documented globally.

Q838. What is the joint-hypothesis problem in testing market efficiency?

Any test of market efficiency is simultaneously a test of efficiency AND of the asset-pricing model used to define "normal" expected returns. If you find abnormal returns, you cannot tell whether the market is inefficient or whether your risk model (e.g., CAPM) is simply wrong/incomplete. Fama emphasized this means market efficiency per se is not testable in isolation — a rejection could just mean a bad benchmark for risk.

Q839. What is the size effect anomaly?

Documented by Banz (1981), it is the historical tendency of small-capitalization stocks to earn higher average returns than large-cap stocks, even after adjusting for beta. It became the SMB (small-minus-big) factor in the Fama-French model. Its strength has weakened since publication, and much of the premium concentrates in tiny, illiquid stocks, raising the question of whether it is a true risk premium or a transaction-cost mirage.

Q840. What is the value effect, and is it considered a risk premium or a mispricing?

The value effect is the tendency of value stocks (high book-to-market, low P/E) to outperform growth stocks over long horizons; it is the HML (high-minus-low) factor in Fama-French. Fama-French argue it is

compensation for risk (distress/financial leverage), while behavioralists argue it reflects investors over-extrapolating growth and overpaying for glamour stocks. Value suffered a long drought in the 2010s, reviving the debate over whether it has been arbitrated away.

Q841. What is the momentum anomaly?

Momentum (Jegadeesh-Titman, 1993) is the tendency of stocks that performed well over the past 3-12 months to keep outperforming over the next 3-12 months, and losers to keep losing. It is hard to reconcile with EMH and with risk-based explanations, making it a leading exhibit for behavioral finance (under-reaction then over-reaction). It works across markets and asset classes but suffers occasional sharp "momentum crashes" during market reversals.

Q842. What is the low-volatility anomaly and why is it puzzling?

It is the empirical finding that low-volatility/low-beta stocks deliver returns as high as or higher than high-volatility stocks on a risk-adjusted basis, directly contradicting CAPM's prediction that more risk should earn more return. Explanations include leverage constraints (investors who cannot borrow bid up risky stocks) and a lottery-preference for volatile names. It has spawned "min-vol" and "low-vol" smart-beta ETFs globally.

Q843. What is the January effect and does it still survive?

The January effect is the historical tendency for small-cap stocks to earn unusually high returns in early January, often linked to tax-loss selling in December followed by repurchasing. It is a calendar anomaly; like many such effects it has largely faded or reversed once widely publicized, consistent with the idea that arbitrage erodes anomalies once known. Its decay is often cited as evidence the market adapts.

Q844. What does the disappearance of many anomalies after publication tell us about EMH?

McLean and Pontiff (2016) showed that documented anomalies' returns decline roughly 30-60% after publication, as arbitrageurs trade against them. This supports an "adaptive" view: markets are not perfectly efficient, but they self-correct as profitable patterns become known, pushing returns back toward normal. It is a middle ground — anomalies are real but often not durable enough to exploit net of costs.

Q845. What is the Grossman-Stiglitz paradox?

Grossman and Stiglitz (1980) argued that if prices perfectly reflected all information, no one would have any incentive to spend money gathering information, since they could free-ride on prices. But then prices could not reflect information no one collected — a contradiction. The resolution is that markets must be "efficiently inefficient": small persistent mispricings must exist to reward information gatherers, so perfect efficiency is impossible in equilibrium.

Q846. What is overconfidence bias and how does it harm investors?

Overconfidence is the tendency to overestimate the accuracy of one's knowledge and predictions. It leads investors to trade too much and under-diversify; Barber and Odean found the most active retail traders underperform the market by several percent a year, largely due to transaction costs. Example: a trader who, after a few winning F&O trades on the Nifty, scales up position size believing they have an edge, only to give back gains.

Q847. What is anchoring bias, with an example?

Anchoring is over-relying on an initial reference number when making estimates. Example: an investor refuses to sell a stock bought at INR 500 even after fundamentals deteriorate, because they anchor on the purchase price and "want to get back to even." Analysts also anchor target prices to current prices or to round numbers, causing sluggish revisions when new information arrives.

Q848. What is loss aversion and how does prospect theory describe it?

Loss aversion (Kahneman and Tversky) is the finding that losses hurt roughly twice as much as equivalent gains feel good — the value function is steeper for losses. Prospect theory says people evaluate outcomes as gains/losses relative to a reference point (not absolute wealth), are risk-averse in gains but risk-seeking in losses, and overweight small probabilities. Practically, a INR 1,000 loss causes about as much pain as a INR 2,000 gain gives pleasure.

Q849. What is the disposition effect and how is it linked to prospect theory?

The disposition effect is the tendency to sell winners too early and hold losers too long. It flows from prospect theory: investors are risk-averse over gains (lock in the sure win) but risk-seeking over losses (gamble to avoid realizing a loss). It is tax-inefficient (you defer gains and pay tax sooner) and empirically reduces returns; it is widely observed among Indian and global retail investors.

Q850. What is mental accounting, with an example?

Mental accounting (Thaler) is treating money differently depending on arbitrary categories rather than recognizing money is fungible. Example: an investor keeps a "safe" fixed deposit earning 6% while simultaneously carrying credit-card debt at 36%, because they mentally separate "savings" from "debt." Another: treating dividend income as spendable "income" while refusing to sell a tiny slice of principal that is economically identical.

Q851. What is herding behavior and how does it affect markets?

Herding is following the crowd rather than one's own analysis, often from a fear of missing out or safety in numbers. It amplifies bubbles (everyone piles into hot stocks) and crashes (everyone sells together), increasing volatility and correlation. Example: retail investors rushing into a hyped IPO or a meme/penny stock purely because others are buying, ignoring valuation.

Q852. Distinguish confirmation bias, recency bias, and the availability heuristic.

Confirmation bias: seeking and weighting only evidence that supports your existing view (e.g., a bull reading only bullish research on a held stock). Recency bias: overweighting recent events, like assuming a market that rose all year will keep rising. Availability heuristic: judging probability by how easily examples come to mind, e.g., overestimating crash risk right after a vivid market crash because it is fresh in memory.

Q853. What is framing bias, with an example?

Framing is letting the way information is presented — not its substance — change your decision. Example: investors react more positively to a fund described as having a "90% success rate" than one with a "10% failure rate," though they are identical. In product sales, "this insurance protects your family" (gain frame) sells better than the economically equivalent "if you die uninsured your family suffers."

Q854. What are limits to arbitrage and why can't smart money always correct mispricing?

Limits to arbitrage are the frictions that stop rational traders from eliminating mispricings: fundamental risk (the asset's value can move against you), noise-trader risk (mispricing can worsen before it corrects — "the market can stay irrational longer than you can stay solvent"), short-selling constraints and costs, capital constraints, and career/redemption risk. Shleifer and Vishny showed arbitrageurs face margin calls and client withdrawals exactly when mispricings are widest, forcing them to unwind rather than press the bet.

Q855. Who are noise traders and what role do they play in behavioral models?

Noise traders are investors who trade on sentiment, rumor, or psychological biases rather than fundamental information. In the DeLong-Shleifer-Summers-Waldmann model, their unpredictable sentiment creates "noise-trader risk" that deters rational arbitrageurs, so noise traders can persist and even earn higher returns for bearing the risk they themselves create. Their presence helps explain why prices deviate from fundamentals and why those deviations are not quickly erased.

Q856. Explain the equity premium puzzle, behavioral drivers of bubbles, and the Adaptive Markets Hypothesis together.

The equity premium puzzle (Mehra-Prescott) is that the historical excess return of stocks over bonds (~5-7% in the US) is far too large to be justified by standard risk-aversion models; behavioral explanations include loss aversion plus "myopic" frequent evaluation (Benartzi-Thaler), which makes investors demand a big premium to hold volatile stocks. The same biases — herding, overconfidence, extrapolation and weak limits to arbitrage — drove historical bubbles like Dutch tulip mania (1637), the dot-com bubble (2000), and the 2008 housing/credit crisis, where prices detached from fundamentals before crashing. Andrew Lo's Adaptive Markets Hypothesis reconciles EMH and behavioral finance by treating markets as evolving ecosystems where efficiency is not absolute but depends on competition, adaptation, and changing conditions — explaining why anomalies appear, get arbitrated away, and why contrarian and value strategies (buying what fearful crowds oversell) can pay off over full cycles.

H. Markets & Instruments

H1. Fixed Income (Advanced) (Q857-Q888)

Q857. What is the term structure of interest rates, and what does a "normal" vs "inverted" yield curve signal?

The term structure is the relationship between yield and time to maturity for bonds of identical credit quality (e.g., G-Secs or US Treasuries), plotted as the yield curve. A normal (upward-sloping) curve has long yields above short yields, reflecting term premium and expectations of stable/rising rates; an inverted curve (short above long) has historically preceded recessions because markets expect future rate cuts. A flat curve signals transition or uncertainty. The curve is the backbone for discounting, relative-value, and forward-rate analysis.

Q858. State the pure expectations theory of the term structure and its key limitation.

Pure (unbiased) expectations theory says long-term rates are purely the geometric average of expected future short-term rates, so forward rates are unbiased predictors of future spot rates: $(1+s_2)^2 = (1+s_1)^* (1+E[r_{1,2}])$. Its limitation is that it ignores risk premia — it implies investors are indifferent between holding one long bond or rolling short bonds, which contradicts the observed average upward slope and the fact that long bonds carry more price risk. Empirically forwards are biased high as predictors, motivating the liquidity-premium refinement.

Q859. Contrast the liquidity preference theory with the market segmentation and preferred habitat theories.

Liquidity preference theory adds a positive, maturity-increasing term premium to expectations: forward rate = expected future spot + liquidity premium, explaining the typical upward slope even when rates are expected to be flat. Market segmentation theory says yields at each maturity are set by independent supply/demand of distinct investor clienteles (banks short, insurers/pension funds long) with no substitution across maturities. Preferred habitat is the middle ground: investors have a preferred maturity but will move out of it if compensated by a sufficient yield premium, so the curve reflects both expectations and supply/demand imbalances.

Q860. Define spot rate, par yield, and forward rate, and explain how they differ.

A spot rate s_n is the yield on a zero-coupon bond maturing in n periods — the rate to discount a single cash flow at time n . A par yield is the coupon rate that prices a coupon bond exactly at par (100) given the spot curve; it is a weighted blend of spot rates. A forward rate $f(a,b)$ is the rate, agreed today, for borrowing/lending between two future dates a and b , implied by today's spot rates. When the spot curve is upward sloping, forwards lie above spots, which lie above par yields.

Q861. How do you bootstrap the spot (zero) curve from coupon-bond or par-yield data?

Bootstrapping solves for spot rates sequentially from shortest to longest maturity, using each newly found spot to discount earlier coupons of the next bond. For a 2-year par bond at coupon c with annual coupons: $100 = c/(1+s_1) + (100+c)/(1+s_2)^2$, solve for s_2 after s_1 is known from the 1-year. Iterating outward recovers the full zero curve consistent with observed prices. This zero curve, not the coupon yields, is what should be used to value arbitrary cash-flow streams (no-arbitrage pricing).

Q862. Given $s_1 = 5\%$ and $s_2 = 6\%$ (annual), compute the implied 1-year forward rate one year from now, $f(1,2)$.

Use $(1+s_2)^2 = (1+s_1)^* (1+f(1,2))$, so $1+f(1,2) = (1.06)^2 / 1.05 = 1.1236 / 1.05 = 1.07010$. Thus $f(1,2) = 7.01\%$. The forward sits above both spots because the spot curve is rising — the second-year rate must exceed s_2 to

pull the average up from s1.

Q863. Why are forward rates "locked in" today, and how would you arbitrage a mispriced forward?

By buying a 2-year zero and shorting (or borrowing via) a 1-year zero in the right proportions, an investor can synthetically guarantee a known reinvestment rate for year 2 — that is the implied forward, enforced by no-arbitrage. If a traded forward/FRA differs from the implied $f(1,2)$, you go long the cheap leg and short the rich leg (cash-and-carry): e.g., borrow 1-year, lend 2-year, and lock the FRA to capture the spread risk-free. In INR markets this shows up in OIS/MIBOR FRA vs G-Sec implied forwards.

Q864. What is the par yield curve and why is it not the same as the coupon yields of on-the-run bonds?

The par yield curve gives, for each maturity, the single coupon rate that makes a bond price exactly to par given the underlying zero curve: $100 = \sum [c/(1+s_t)^t] + 100/(1+s_n)^n$. It is a synthetic construct used to read off "the coupon you'd need to issue at par today" at any maturity. It differs from actual on-the-run yields because real bonds rarely price exactly at par and carry liquidity/coupon-effect differences; the par curve removes those by construction.

Q865. Distinguish modified duration from effective duration and say when each is appropriate.

Modified duration = Macaulay duration / $(1 + y/k)$ and assumes fixed cash flows, so it is exact for option-free bonds. Effective duration uses a parallel shift of the benchmark curve and reprices with a model: $\text{EffDur} = (P(-dy) - P(+dy)) / (2 * P_0 * dy)$, capturing how cash flows themselves change when rates move. Use modified duration for plain vanilla bonds and effective duration for bonds with embedded options (callable, puttable, MBS) where cash flows are rate-dependent.

Q866. Compute effective duration given $P_0 = 100.00$, $P(+25bp) = 98.30$, $P(-25bp) = 101.80$.

$\text{EffDur} = (P(-dy) - P(+dy)) / (2 * P_0 * dy) = (101.80 - 98.30) / (2 * 100.00 * 0.0025) = 3.50 / 0.50 = 7.0$. So a 100bp parallel shift moves price by roughly 7%. Because we reprice with the model on each side, this automatically reflects any optionality embedded in the bond.

Q867. What are key-rate (partial) durations and why do they matter beyond a single effective duration?

Key-rate durations measure price sensitivity to a 1bp shift at specific maturity points on the curve (e.g., 2y, 5y, 10y, 30y) while holding other points fixed, so they capture exposure to non-parallel moves (steepening, flattening, twists). A single effective duration assumes parallel shifts and can hide curve risk: two portfolios with identical total duration can react very differently to a flattener. Summing the key-rate durations recovers the overall effective duration.

Q868. How is portfolio duration computed, and what is its main assumption?

Portfolio duration is the market-value-weighted average of the component durations: $D_p = \sum [w_i * D_i]$, where w_i = market value of bond i / total portfolio value. For example, 60% in a 3-year-duration bond and 40% in an 8-year gives $0.63 + 0.48 = 5.0$. The key assumption is a parallel shift of the yield curve; it understates risk when yields move by different amounts across maturities, which is where key-rate durations are needed.

Q869. What is immunization, and what are the conditions to immunize a single liability?

Immunization structures a bond portfolio so its value moves in step with a target liability, locking in a return regardless of rate moves. The classic conditions: (1) portfolio Macaulay duration equals the liability's horizon, (2) present value of assets $\geq$ present value of the liability, and (3) minimize convexity dispersion of asset cash flows around the horizon to reduce structural risk. Duration matching balances price risk and reinvestment risk so they offset at the horizon date.

Q870. Why does duration matching alone not fully immunize against non-parallel curve shifts, and what improves it?

Duration matching protects against small parallel shifts because price and reinvestment effects offset only at first order. Non-parallel shifts (twists/butterflies) break this, since cash flows clustered far from the liability date get re-priced differently — barbell structures with high dispersion are especially exposed. Matching also higher moments (convexity/dispersion) and minimizing the spread of asset cash flow maturities around the liability, or using key-rate duration matching, improves robustness. Rebalancing is required as durations drift with time and rates.

Q871. Define dollar duration (DVo1/PVo1) and show how to hedge a position with it.

Dollar duration (a.k.a. DVo1 or PVo1) is the change in a position's value for a 1bp yield move: $DVo1 = \text{ModDur} * \text{Price} * 0.0001$ (per 100 face, then scaled by position). To hedge, set the notional of the hedge instrument so its DVo1 offsets the position's: $\text{hedge notional} = \text{portfolio DVo1} / \text{hedge-per-unit DVo1}$. E.g., a position with DVo1 of INR 50,000 hedged with a futures contract worth INR 25 DVo1 needs 2,000 contracts; in the US you'd hedge with Treasury futures or the cheapest-to-deliver.

Q872. A 10-year bond has modified duration 8.2 and price 96. What is its DVo1 per INR 100 face, and the P&L on INR 50 crore for a +15bp move?

$DVo1 = \text{ModDur} * \text{Price} * 0.0001 = 8.2 * 96 * 0.0001 = 0.07872$ per 100 face. On INR 50 crore face that scales by $50,00,00,000/100 = 50,00,000$ units, so $DVo1 = 0.07872 * 50,00,000 = \text{INR } 3,93,600$ per bp. A +15bp move loses about $15 * 3,93,600 = \text{INR } 59.0$ lakh (long bond, rates up). This linear estimate ignores convexity, which would slightly reduce the actual loss.

Q873. Write the second-order (duration + convexity) price-change formula and explain each term.

Percentage price change $\approx -\text{ModDur} * dy + 0.5 * \text{Convexity} * dy^2$, where dy is the yield change in decimal. The first term is the linear duration estimate (negative slope of the price-yield curve); the second is the convexity correction, which is always positive for option-free bonds because the price-yield relationship is curved (convex). Convexity makes the duration estimate too pessimistic for falling yields and too optimistic for rising yields, so the adjustment improves accuracy for large dy .

Q874. With $\text{ModDur} = 7.0$ and $\text{Convexity} = 65$, estimate the price change for a -100bp move, including convexity.

$dy = -0.01$. Duration term $= -7.0 * (-0.01) = +0.070$ ($= +7.0\%$). Convexity term $= 0.5 * 65 * (-0.01)^2 = 0.5 * 65 * 0.0001 = +0.00325$ ($= +0.325\%$). Total $\approx +7.325\%$, versus $+7.0\%$ from duration alone — convexity adds about 33bp of price gain on a large favorable move. For a +100bp move convexity would instead cushion the loss to about -6.675%.

Q875. Why do callable bonds exhibit negative convexity, and over what yield region?

A callable bond gives the issuer the right to redeem early, so when yields fall and the embedded call moves into the money, price appreciation is capped near the call price ("price compression") — the price-yield curve bends the wrong way, giving negative convexity. This occurs in the low-yield region where the call is likely to be exercised; at high yields the bond behaves like an option-free bond with positive convexity. The investor is short a call, so they earn a higher yield as compensation.

Q876. How does a puttable bond's convexity and value differ from a callable bond's?

A puttable bond lets the holder sell it back at a set price when yields rise, putting a floor under the price, so the holder is long a put: $\text{puttable value} = \text{straight bond} + \text{put value}$, and it exhibits more positive convexity than an option-free bond at high yields. A callable bond holder is short a call: $\text{callable value} = \text{straight bond} - \text{call value}$, producing negative convexity at low yields. Hence puttables trade at lower yields (holder pays for protection) while callables trade at higher yields.

Q877. Define nominal spread, I-spread, Z-spread, and OAS, and order them by sophistication.

Nominal spread is yield-to-maturity over the matching-maturity benchmark (single point). I-spread is the bond's yield over the interpolated swap rate of equal maturity. Z-spread (zero-volatility spread) is the constant spread added to every spot rate so discounted cash flows equal the market price — it accounts for the whole curve. OAS (option-adjusted spread) removes embedded-option value by averaging across interest-rate paths in a model: it is the Z-spread net of option cost, and is the cleanest measure for comparing bonds with optionality.

Q878. What is the relationship between Z-spread and OAS for callable vs putable bonds?

For a callable bond, $OAS = Z\text{-spread} - \text{call option cost (in bp)}$, so $OAS < Z\text{-spread}$, because part of the apparently high Z-spread is just compensation for being short the call rather than true credit/liquidity spread. For a putable bond the holder owns the option, so $OAS = Z\text{-spread} + \text{put value benefit}$ and $OAS > Z\text{-spread}$. For an option-free bond, $OAS = Z\text{-spread}$. OAS thus lets you compare a callable, a putable, and a bullet on an apples-to-apples basis.

Q879. What are mortgage-backed securities (MBS) and the prepayment risk they carry?

An MBS is a pass-through security backed by a pool of mortgages whose principal and interest are passed to investors; in the US these are issued by agencies (GNMA/FNMA/FHLMC). Prepayment risk is the homeowner's option to repay early: when rates fall, refinancing accelerates prepayments (contraction risk — cash comes back to be reinvested at low rates), and when rates rise prepayments slow (extension risk — duration lengthens just as rates rise). This embedded short call gives MBS negative convexity, modeled via measures like PSA/CPR.

Q880. Explain CPR, SMM, and the PSA benchmark used to model prepayments.

CPR (conditional prepayment rate) is the annualized fraction of the remaining pool expected to prepay; SMM (single monthly mortality) is its monthly equivalent: $SMM = 1 - (1 - CPR)^{(1/12)}$. The PSA benchmark assumes CPR ramps linearly from 0.2% in month 1 by 0.2%/month to 6% at month 30, then stays flat — "100 PSA"; faster/slower speeds are quoted as multiples (e.g., 150 PSA, 50 PSA). These drive the timing of principal cash flows and hence average life and effective duration.

Q881. What is securitization and tranching, and how do senior/mezzanine/equity tranches differ?

Securitization pools illiquid assets (auto loans, credit-card receivables, mortgages) into an SPV that issues asset-backed securities; tranching slices the SPV's cash flows into a credit waterfall. The senior tranche is paid first and absorbs losses last (highest rating, lowest yield); mezzanine sits in the middle; the equity/first-loss tranche absorbs initial losses and earns the residual/highest yield. Subordination and overcollateralization provide credit enhancement so senior tranches can be rated above the underlying pool.

Q882. How did CDO tranching contribute to the 2008 financial crisis?

Subprime mortgages were pooled into RMBS, then the risky mezzanine tranches were re-securitized into CDOs and re-tranched, with the new senior pieces rated AAA on flawed assumptions of low, uncorrelated default correlation. When house prices fell nationally, defaults became highly correlated, so losses blew through subordination and reached supposedly safe senior/super-senior tranches. Mispriced correlation, rating-agency conflicts, and leverage turned a housing downturn into a systemic credit collapse.

Q883. For a convertible bond, define conversion ratio, conversion (parity) value, and conversion premium with a quick example.

Conversion ratio is the number of shares received per bond; conversion (parity) value = conversion ratio * current share price. Conversion premium = (convertible price - parity value) / parity value, the extra paid for the bond's downside protection and optionality. Example: ratio 20, share price INR 45 gives parity = $20 * 45 = \text{INR } 900$; if the convertible trades at INR 1,000, the premium is $100/900 = 11.1\%$. The conversion price = $\text{par} / \text{ratio} = 1000/20 = \text{INR } 50$ is the breakeven share price.

Q884. How do you value a convertible bond and what embedded options does it contain?

A convertible's value has a floor at the greater of its straight-bond value (investment value) and its conversion/parity value, plus the time value of the embedded equity call option: convertible $\approx$ straight bond + call on the stock. Many are also callable by the issuer (an embedded short call to the holder) to force conversion, and some are puttable. It behaves like debt when the stock is low (bond floor dominates) and like equity when the stock is high (delta approaches the conversion ratio).

Q885. How does a floating-rate note (FRN) work, and why is its price near par with low duration?

An FRN pays a coupon that resets each period to a reference rate plus a quoted margin (e.g., 3M T-Bill / MIBOR / SOFR + 80bp), so the coupon tracks market rates. Because the coupon resets, interest-rate (duration) risk is limited to the short period until the next reset, keeping price close to par; its effective rate duration is roughly the time to next reset. Price still moves if the issuer's credit (required margin) changes — captured by spread duration rather than rate duration.

Q886. Define credit spread and spread duration, and quantify a spread move's price impact.

Credit spread is the yield premium over the risk-free/benchmark curve compensating for default and liquidity risk; spread duration measures price sensitivity to a change in that spread: $\% \Delta \text{Price} \approx -\text{SpreadDur} * \Delta \text{spread}$. For a bond with spread duration 6, a +50bp widening costs about $-6 * 0.005 = -3.0\%$. For fixed-rate bonds spread duration $\approx$ modified duration, but for FRNs rate duration is tiny while spread duration remains full-term, which is why FRNs still carry meaningful credit risk.

Q887. How do inflation-linked bonds like US TIPS and India's IIBs work, and what is the breakeven inflation rate?

Inflation-linked bonds protect real purchasing power: US TIPS index the principal to CPI and pay a fixed real coupon on the inflated principal, so investors earn a real yield. India's Inflation Indexed Bonds (IIBs) similarly link principal to WPI/CPI (the retail CGB-linked IINSS-C used CPI). Breakeven inflation = nominal Treasury yield - real TIPS yield; if 10y nominal is 4.3% and 10y TIPS real is 1.8%, breakeven $\approx 2.5\%$, and you prefer TIPS when you expect realized inflation to exceed breakeven.

Q888. Describe the Indian G-Sec yield curve and how corporate bond spreads are quoted against it.

The G-Sec curve is built from sovereign rupee bonds across tenors (benchmarked by the liquid 10-year on-the-run G-Sec), with the curve published via CCIL/FBIL and yields driven by RBI policy (repo rate), inflation, and the fiscal borrowing calendar. AAA PSU and corporate bonds are quoted as a spread over the matching-tenor G-Sec (e.g., AAA 10y at G-Sec + 60-90bp), widening down the rating ladder (AA, A) and in liquidity stress. SDLs (state development loans) trade at a spread above central G-Secs reflecting state credit and lower liquidity.

H2. Swaps, Interest-Rate & FX Derivatives (Q889-Q918)

Q889. What is an interest-rate swap (IRS)?

An interest-rate swap is an OTC contract in which two parties agree to exchange interest payments on a notional principal: typically one party pays a fixed rate while the other pays a floating rate (e.g. SOFR, EURIBOR, or in India MIBOR/MIFOR). Only interest cash flows are exchanged, and they are usually netted, so it is a tool to transform the interest character of an asset or liability without touching the underlying principal. A "plain vanilla" IRS is fixed-for-floating in a single currency.

Q890. In a plain-vanilla IRS, is the notional principal exchanged?

No. The notional is purely a reference amount used to compute the interest cash flows; it is never exchanged in a single-currency interest-rate swap. This is the key difference from a currency swap, where principal

usually is exchanged. Because principal stays put, the credit exposure of an IRS is far smaller than its notional suggests.

Q891. What does "net settlement" mean in a swap?

On each payment date both legs are calculated, and only the difference between the fixed and floating amounts changes hands, paid by whichever party owes more. For example, on a USD 10m notional, if the fixed payer owes 5% (USD 500k for the year) and the floating receiver is owed 4.5% (USD 450k), only USD 50k is paid by the fixed payer. Netting reduces settlement risk and the size of cash flows actually moving.

Q892. Why would a company use an interest-rate swap?

The classic motive is converting floating-rate debt into fixed-rate (to lock in costs and hedge against rising rates) or vice versa. A firm with a floating loan can enter a "pay-fixed, receive-floating" swap: the floating it receives offsets the floating it owes on the loan, leaving it with a net fixed cost. Swaps also let firms exploit comparative advantage in different markets, manage duration, and speculate on rate moves.

Q893. Explain the comparative-advantage argument for swaps.

If Company A borrows more cheaply in fixed and Company B more cheaply in floating, but each actually wants the other type, they can each borrow where they have an advantage and swap. The total interest saving (the difference between the two fixed spreads and the two floating spreads) is shared between them, so both end up paying less than borrowing directly in their preferred market. Critics note part of this "free lunch" reflects hidden credit-risk differences and renewal risk in floating rates.

Q894. A firm has a floating loan at MIBOR+150 bps and wants certainty. How does a swap help?

It enters a pay-fixed/receive-MIBOR swap on the same notional, say paying 7% fixed and receiving MIBOR. Its net cost becomes (MIBOR + 1.50%) on the loan minus MIBOR received plus 7% paid = 8.50% fixed. The floating exposure cancels, locking the all-in cost at 8.50% regardless of where MIBOR moves.

Q895. Why is the value of a swap approximately zero at initiation?

At inception the fixed rate (the swap rate) is set precisely so that the present value of the fixed leg equals the present value of the expected floating leg, so neither party pays the other to enter. As market rates move afterward, one leg becomes more valuable and the swap acquires positive value to one party and negative to the other. $V_{mid} = PV(\text{fixed leg}) - PV(\text{floating leg}) = 0$ at $t=0$.

Q896. How can a swap be valued as a portfolio of bonds?

To the fixed-rate receiver, a swap equals being long a fixed-coupon bond and short a floating-rate bond: $V = B_{\text{fixed}} - B_{\text{floating}}$. The fixed bond is discounted at current rates; the floating-rate bond resets to par at each reset, so right after a reset it is worth its notional. Thus immediately after a reset, V (to fixed receiver) = $B_{\text{fixed}} - \text{Notional}$.

Q897. How can a swap be valued as a series of FRAs?

Each exchange of fixed-for-floating on a future date is economically a forward rate agreement: you can value each settlement by replacing the unknown future floating rate with today's forward rate, compute the net cash flow, and discount it back. The swap value is the sum of these discounted FRA values. This "strip of FRAs" view and the bond view give the same answer.

Q898. What is the swap rate?

The swap rate is the fixed rate that makes a new swap have zero value at initiation, i.e. the fixed rate at which $PV(\text{fixed leg}) = PV(\text{floating leg})$. It equals a weighted average of the forward rates implied by the curve. Formula: $\text{swap rate} = (1 - DF_n) / (\text{sum of } DF_i * \tau_i)$, where DF are discount factors and τ the accrual fractions.

Q899. What is the swap curve and why does it matter?

The swap curve plots par swap rates against maturity (2y, 5y, 10y, etc.) and is a key benchmark term structure, often more liquid than government bonds in some tenors. It is used to discount cash flows, price other derivatives, and gauge market rate expectations and credit/liquidity conditions. In India, the INR OIS (MIBOR-based) curve serves this role for rupee rates.

Q900. What is a forward rate agreement (FRA)?

An FRA is an OTC contract that fixes an interest rate on a notional for a future period; one party locks a borrowing/lending rate today for, say, a 3-month period starting in 6 months (a "6x9 FRA"). At settlement only the difference between the agreed rate and the realized reference rate is exchanged. It is effectively a single-period building block of a swap.

Q901. What is the payoff of an FRA and how is it settled?

Payoff to the fixed-rate payer (borrower) = $\text{Notional} * (R_{\text{reference}} - R_{\text{FRA}}) * (\text{days}/360) / (1 + R_{\text{reference}} * (\text{days}/360))$. The cash flow is settled at the start of the reference period, so it is discounted back from the period end, which is why the divisor appears. If realized rates exceed the locked rate, the fixed payer (who feared rising rates) receives the payment.

Q902. Worked FRA example: 6x12 FRA, notional USD 5m, agreed 4.00%, settles at 5.00%.

The period is 6 months (~182 days, ~0.5 year). Unadjusted gain to the fixed payer = $5\text{m} * (5.00\% - 4.00\%) * 0.5 = \text{USD } 25,000$. Discounted to settlement at the start of the period: $25,000 / (1 + 0.05 * 0.5) = \text{USD } 24,390$. The borrower who locked 4% benefits because rates rose to 5%.

Q903. What was LIBOR and why is the market transitioning away from it?

LIBOR (London Interbank Offered Rate) was a survey-based estimate of unsecured interbank borrowing costs, long used as the floating benchmark for trillions in loans and derivatives. It was discredited by the 2012 rigging scandal and undermined by the collapse of actual interbank lending, making it manipulable and unrepresentative. Regulators ceased its publication (USD LIBOR's last panel rates ended June 2023), forcing a move to transaction-based risk-free rates.

Q904. What is SOFR and how does it differ from LIBOR?

SOFR (Secured Overnight Financing Rate) is the US replacement: an overnight rate based on actual secured (repo) Treasury transactions, making it robust and nearly manipulation-proof. Unlike term, forward-looking, unsecured LIBOR, SOFR is overnight, secured (so it carries almost no bank-credit premium), and is typically compounded in arrears to build term rates. Because it lacks LIBOR's credit spread, contracts add a fixed spread adjustment (e.g. ~26 bps for 3-month USD) when converting.

Q905. Name the risk-free reference rates in major markets.

US: SOFR; UK: SONIA; Euro area: ESTR (and EURIBOR still survives reformed); Switzerland: SARON; Japan: TONA. India uses MIBOR (Mumbai Interbank Offer Rate / FBIL overnight MIBOR) as the basis for the rupee OIS market, plus MIFOR for FX-implied rupee rates. These are largely overnight, transaction-anchored rates compounded for term exposure.

Q906. What is OIS and what is OIS discounting?

An Overnight Indexed Swap exchanges a fixed rate for the geometric average of an overnight index (SOFR, ESTR, or MIBOR in India) over the period. Post-2008, dealers shifted to "OIS discounting" — using the OIS curve, not LIBOR, to discount collateralized swap cash flows, because cash collateral earns the overnight rate, making OIS the true funding/discount rate. This better reflects the near-risk-free cost of funding posted collateral.

Q907. What is a currency swap?

A currency swap exchanges principal and interest payments in one currency for principal and interest in another. Unlike an IRS, the principal usually is exchanged at the start and re-exchanged at maturity (at the original spot rate). It lets a firm raise cheap funding in one currency and convert it into the currency it actually needs while hedging both interest and FX exposure.

Q908. What is cross-currency basis?

The cross-currency basis is the spread added to one leg of a cross-currency swap that deviates from what covered interest parity would predict, reflecting supply/demand imbalances for funding in a given currency. A negative USD basis means non-US borrowers pay a premium to swap into dollars, common in stress periods when dollar funding is scarce. It is a market measure of how expensive it is to obtain a currency via the FX swap market.

Q909. What is an FX forward?

An FX forward is an OTC agreement to exchange two currencies at a fixed rate on a future date, locking in the exchange rate today and removing uncertainty about the future spot. An importer owing USD in 3 months can buy USD forward to fix its INR cost. The forward rate is determined by spot and the interest-rate differential, not by a forecast of the future spot.

Q910. State covered interest rate parity (CIP).

CIP says the forward exchange rate is pinned by the interest-rate differential, otherwise arbitrage is possible: $F = S * (1 + r_{\text{domestic}}) / (1 + r_{\text{foreign}})$, where rates are for the contract horizon and S, F are quoted as domestic per foreign. The currency with the higher interest rate trades at a forward discount, exactly offsetting its rate advantage so no risk-free profit exists.

Q911. Worked CIP example: INR/USD, spot 83.00, INR rate 7%, USD rate 5%, 1 year.

Treating INR as domestic and USD as foreign: $F = 83.00 * (1 + 0.07) / (1 + 0.05) = 83.00 * 1.07/1.05 = 83.00 * 1.01905 = 84.58$. So the 1-year USD/INR forward is about 84.58, i.e. USD is at a forward premium of ~1.58 rupees because INR's higher rate must be offset.

Q912. What is a forward premium versus a forward discount?

A currency is at a forward premium if it is more expensive to buy forward than spot, and at a discount if cheaper. Annualized forward premium $\sim (F - S)/S * (12/\text{months})$. By CIP the low-interest-rate currency trades at a forward premium and the high-rate currency at a discount; in INR/USD terms the rupee, with higher rates, trades at a forward discount (you get more INR per USD forward).

Q913. How does an FX swap differ from a currency swap?

An FX swap is a single contract combining a spot (or near-date) exchange with a simultaneous reverse forward exchange of the same two currencies — it is a short-dated funding/rollover tool with no intermediate interest payments; the rate differential is baked into the forward points. A cross-currency (currency) swap is longer-dated and involves periodic interest payments on both legs throughout its life. In short: FX swap = two principal exchanges, no coupons; currency swap = principal plus a stream of interest payments.

Q914. What is an interest-rate cap, and how does it relate to caplets?

A cap is an option-based hedge that pays the holder whenever a floating reference rate rises above a strike (cap) rate, effectively setting a ceiling on borrowing cost. It is a strip (portfolio) of "caplets," one per reset period, each a call option on the interest rate. Caplet payoff = $\text{Notional} * \max(R_{\text{reference}} - R_{\text{strike}}, 0) * (\text{days}/360)$, and the cap premium is the sum of the caplet premiums.

Q915. What are floors and collars?

A floor is the mirror image of a cap: a strip of "floorlets" (puts on the rate) that pays when the floating rate falls below a strike, guaranteeing a minimum return to a lender/investor. A collar combines buying a cap and

selling a floor (for a borrower), so the rate is confined to a band; the premium received from the sold floor offsets the cost of the bought cap, and a "zero-cost collar" sets strikes so net premium is nil.

Q916. What is a swaption?

A swaption is an option granting the right, but not the obligation, to enter a swap at a preset fixed rate on a future date. A "payer swaption" gives the right to pay fixed/receive floating (valuable if rates rise); a "receiver swaption" gives the right to receive fixed (valuable if rates fall). Firms use them to hedge anticipated financing or to lock a future borrowing rate with optionality, paying an upfront premium.

Q917. What is basis risk in interest-rate/FX hedging?

Basis risk is the residual risk that a hedge does not move perfectly with the exposure because the hedge references a different rate, index, or tenor than the underlying. For example, hedging a loan tied to 1-month MIBOR with a 3-month MIBOR swap, or hedging SOFR-linked debt against a term-rate exposure, leaves a "basis" that can move against you. It is why LIBOR-to-SOFR transition and tenor mismatches matter to hedgers.

Q918. How did post-2008 reform change swap markets via clearing and netting, and what is India's framework?

After 2008, the G20 mandated that standardized OTC derivatives be centrally cleared through CCPs (which novate trades and net exposures), reported to trade repositories, and margined, sharply reducing counterparty risk; netting agreements (ISDA master agreements) collapse many trades into a single net obligation on default. In India, rupee IRS/OIS and MIFOR swaps clear through the Clearing Corporation of India (CCIL), while currency and FX derivatives are jointly regulated by RBI and SEBI — exchange-traded currency futures/options trade on NSE/BSE with RBI-set position limits, and OTC forwards require an underlying exposure under RBI's FEMA rules.

H3. Foreign Exchange, Commodities & Money Markets (Q919-Q948)

Q919. In a quote like USD/INR = 83, which is the base currency and which is the quote currency, and what does the number mean?

USD is the base currency and INR is the quote (or counter) currency; the rate tells you how many units of the quote currency buy ONE unit of the base. So USD/INR = 83 means 1 US dollar costs 83 Indian rupees. The convention is base/quote, and you always read the price as "units of quote per one base."

Q920. What is the spot FX market and what is a typical settlement convention?

The spot market is for immediate exchange of one currency for another at the current rate, but "immediate" settlement is conventionally T+2 (two business days) for most pairs, and T+1 for USD/CAD and USD/INR. It is the largest, most liquid FX segment, with daily global turnover around USD 7-8 trillion. The spot rate is the benchmark off which forwards and swaps are priced.

Q921. Explain bid, ask, and the bid-ask spread in FX.

The bid is the price at which the market maker buys the base currency from you; the ask (offer) is the price at which it sells the base to you, and ask > bid. For USD/INR quoted 82.98 / 83.02, you sell dollars at 82.98 and buy dollars at 83.02; the 0.04 spread is the dealer's compensation and a transaction cost to you. Liquid pairs like EUR/USD have spreads of a fraction of a pip; thin or emerging-market pairs are wider.

Q922. What is a pip, and what is a pip worth?

A pip ("percentage in point") is the smallest standard increment in an FX quote, conventionally the 4th decimal place for most pairs (0.0001) and the 2nd decimal (0.01) for JPY pairs. For EUR/USD moving

1.1050 to 1.1051 that is one pip; on a standard lot of 100,000 EUR the value is $100,000 * 0.0001 = \text{USD } 10$ per pip. Fractional pips (a 5th decimal, the "pipette") are now common in electronic quoting.

Q923. What is the difference between appreciation and depreciation, and how does it show in the quote?

A currency appreciates when it gains value (buys more of another currency) and depreciates when it loses value. If USD/INR moves from 83 to 85, the dollar appreciated and the rupee depreciated (it now takes more rupees per dollar). Confusingly, a RISING USD/INR number means a WEAKER rupee, because INR is the quote currency.

Q924. What is a cross rate and how do you compute one?

A cross rate is an exchange rate between two currencies derived through a common third currency (usually USD). If $\text{EUR/USD} = 1.10$ and $\text{USD/INR} = 83$, then $\text{EUR/INR} = 1.10 * 83 = 91.3$. When both quotes share the same base or quote you divide instead of multiply, so always line up the units so the common currency cancels.

Q925. What are the main drivers of exchange rates?

Key drivers are interest-rate differentials (higher real rates attract capital and strengthen a currency), inflation differentials (higher inflation tends to weaken a currency), the balance of payments (a current-account deficit creates net selling pressure that must be financed by capital inflows), risk sentiment (risk-off flows favor safe havens like USD, JPY, CHF), and central-bank intervention. Over the long run fundamentals like productivity and terms of trade dominate; over short horizons rate expectations and sentiment dominate. India's rupee is especially sensitive to its oil-driven trade deficit and to FII equity flows.

Q926. State covered interest rate parity (CIP), give the formula, and work an example.

CIP says the forward rate offsets the interest-rate differential exactly, leaving no risk-free arbitrage: $F = S * (1 + r_{\text{quote}}) / (1 + r_{\text{base}})$, for the same horizon. Example: spot $\text{USD/INR} = 83$, 1-year INR rate 7 percent, USD rate 5 percent, so $F = 83 * (1.07 / 1.05) = 84.58$. The higher-yielding currency (INR) trades at a forward DISCOUNT, compensating for the rate gap so covered arbitrage profit is zero.

Q927. What is uncovered interest rate parity (UIP), how does it differ from CIP, and is the forward a good predictor of spot?

UIP states the EXPECTED change in spot equals the interest differential, so a higher-yielding currency is expected to depreciate by about that amount: $E[S_{\text{future}}]/S = (1 + r_{\text{quote}})/(1 + r_{\text{base}})$. Unlike CIP there is no forward locking the rate, so UIP is an expectations condition that bears FX risk, not an arbitrage relation, and it fails empirically. Because of this "forward premium puzzle," the forward rate is a poor, biased predictor of future spot — it mainly reflects the rate differential, not a genuine forecast.

Q928. State the law of one price and absolute purchasing power parity (PPP).

The law of one price says an identical good should cost the same everywhere once converted to a common currency. Absolute PPP extends this to a basket: $S = P_{\text{domestic}} / P_{\text{foreign}}$, so the exchange rate equals the ratio of price levels. If a basket costs 8300 INR and USD 100, absolute PPP implies $\text{USD/INR} = 83$.

Q929. What is relative PPP and its formula?

Relative PPP focuses on CHANGES rather than levels: the percentage change in the exchange rate equals the inflation differential, so change in S is approximately $\text{inflation}_{\text{quote}} - \text{inflation}_{\text{base}}$. If India's inflation is 6 percent and US inflation 2 percent, the rupee should depreciate roughly 4 percent per year against the dollar (USD/INR rises about 4 percent). It tends to hold better over long horizons than short ones.

Q930. What is the Big Mac Index and what does it illustrate?

The Big Mac Index, published by The Economist, uses the price of a Big Mac across countries as a one-good proxy basket to gauge whether currencies look over- or under-valued versus PPP. If a Big Mac costs USD 5 in

the US and 250 INR in India, the implied PPP rate is $250/5 = 50$; with the market rate at 83 the rupee looks "undervalued" on that crude measure. It is a teaching tool, not a precise valuation, because it ignores non-tradables, taxes, and local labor costs.

Q931. What is the carry trade, why does it earn returns, and why does it blow up?

A carry trade borrows in a low-yielding currency (e.g., JPY) and invests in a high-yielding one (e.g., BRL, INR, AUD), pocketing the interest differential; it works because UIP fails, so high-yield currencies don't depreciate enough to wipe out the carry. The catch is negative skew: in risk-off episodes everyone unwinds at once, the funding currency sharply appreciates and the high-yield currency crashes, erasing months of small gains in days. It is essentially selling crash insurance, and leverage amplifies both the carry and the wipeout.

Q932. Distinguish transaction, translation, and economic FX exposure.

Transaction exposure is the risk that a specific contracted cash flow (e.g., a USD receivable) changes value in home currency before settlement. Translation (accounting) exposure arises when consolidating foreign-subsidiary financial statements into the parent's reporting currency. Economic (operating) exposure is the broader, long-run effect of currency moves on competitiveness, future cash flows, and firm value, even without explicit foreign-currency contracts.

Q933. How do corporates hedge FX risk?

Common tools are forwards (lock a rate for a future date), futures (exchange-traded, standardized), and currency options (pay a premium for downside protection while keeping upside). A money-market hedge replicates a forward: for a USD receivable, borrow its present value in USD, convert to INR at spot, and invest domestically, so the receivable repays the loan — by CIP this gives essentially the same rate as the forward. Natural hedging (matching revenue and cost currencies, netting intercompany flows) reduces exposure without derivatives; an Indian IT exporter typically sells USD forwards to protect rupee earnings.

Q934. What roles does the RBI play in managing the rupee?

The RBI runs a "managed float": it lets the rupee trade with market forces but intervenes in spot and forward markets to curb excessive volatility, buying dollars when the rupee is strong and selling them to defend it when weak. It also sets the policy (repo) rate, manages capital-flow rules, and operates a forward book. It does not target a fixed level but smooths sharp moves.

Q935. What are forex reserves and why do they matter?

Forex reserves are a central bank's holdings of foreign currencies (mostly USD), gold, SDRs, and the IMF reserve position, used to back the currency, intervene in FX markets, and ensure import and external-debt cover. India's reserves (around USD 600+ billion) provide a cushion of roughly 10-11 months of imports, supporting confidence and credit ratings. Reserves rise when the RBI buys dollars and fall when it sells to defend the rupee.

Q936. What is the NDF market and why does it matter for the rupee?

A Non-Deliverable Forward (NDF) is a cash-settled forward on a currency like INR, traded offshore in centers such as Singapore, London, and Dubai, settled in USD with no physical delivery of rupees. It exists because capital-account restrictions limit onshore access for non-residents, and it can drive or front-run the onshore rate, especially in stress. The RBI watches the NDF closely and has expanded onshore access to reduce offshore price-setting power.

Q937. Distinguish the spot and futures markets for commodities.

The spot (cash) market is for immediate delivery and payment of the physical commodity, while a futures contract is a standardized, exchange-traded agreement to buy/sell a set quantity and grade at a fixed price on a future date. Futures enable hedging and price discovery without immediate physical handling, and most are closed out or rolled before delivery. The futures price links to spot via the cost of carry.

Q938. Explain cost of carry and convenience yield for commodities.

Cost of carry is what it costs to hold a commodity to a future date — storage, insurance, and financing — which pushes the futures price above spot: $F = S * (1 + r + \text{storage})$. Convenience yield is the implicit benefit of holding the physical (availability for production, avoiding stockouts), which pulls the futures price down: $F = S * (1 + r + \text{storage} - \text{convenience yield})$. When convenience yield is high (tight supply), futures can trade below spot.

Q939. Define contango and backwardation for commodities.

Contango is when futures trade above spot (and far months above near months), typical when carrying costs dominate and supply is ample — e.g., crude in oversupply. Backwardation is when futures trade below spot, signaling tight physical supply and high convenience yield, so holders of the physical demand a premium. A long futures position earns positive "roll yield" in backwardation and suffers negative roll yield in contango.

Q940. How do producers and consumers hedge with commodity futures, and why is gold a safe haven?

A producer (e.g., a farmer or oil driller) is naturally LONG the physical, so it SELLS futures to lock a sale price; a consumer (e.g., an airline or jeweler) is short the commodity, so it BUYS futures to lock a purchase price. Gold is treated as a safe haven because it holds value in inflation, crises, and currency debasement, has low or negative correlation with risk assets, and carries no credit risk — central banks and investors buy it when confidence in fiat or markets falls.

Q941. Name the major commodity exchanges in India and globally.

In India, MCX (Multi Commodity Exchange) dominates metals and energy (gold, silver, crude, natural gas) and NCDEX leads agri commodities, both regulated by SEBI (which absorbed the former FMC). Globally, the CME Group (including NYMEX/COMEX) handles energy and metals, and the LME (London Metal Exchange) is the benchmark for industrial base metals like copper, aluminium, and zinc. ICE is key for Brent crude and soft commodities.

Q942. What is the purpose of the money market and what instruments trade there?

The money market is where short-term debt (maturities up to one year) is borrowed and lent, providing liquidity management for banks, firms, and governments at low risk. Instruments include Treasury bills, commercial paper (CP), certificates of deposit (CDs), repurchase agreements (repo), and call/notice money. It is the channel through which central-bank policy rates transmit to the rest of the economy.

Q943. Describe T-bills, commercial paper, CDs, repo, and call money.

Treasury bills are short-term sovereign debt issued at a discount (in India, 91/182/364-day tenors). Commercial paper is unsecured short-term debt issued by high-rated corporates; certificates of deposit are negotiable time deposits issued by banks. A repo is a collateralized short-term loan (sell-and-repurchase of securities); call money is uncollateralized interbank overnight lending. All price off the prevailing policy rate plus a credit/liquidity spread.

Q944. How do money-market rates price off the policy rate?

The policy rate (the RBI repo rate, or the Fed funds target) anchors the shortest, safest rate, and other money-market rates sit above it by a spread reflecting credit risk, term, and liquidity. When the central bank hikes the repo rate, interbank call money, T-bill yields, CP, and CD rates all reprice upward; the central bank uses open-market and repo operations to keep the overnight rate inside its corridor (in India, the LAF corridor with the SDF floor and MSF ceiling). This is the front end of monetary-policy transmission.

Q945. Explain repo mechanics and the haircut.

In a repo, the borrower sells securities (collateral) to the lender and agrees to repurchase them later at a slightly higher price; the price difference is the implicit interest (the repo rate). The haircut is the discount applied to the collateral's market value — e.g., a 2 percent haircut means USD 100 of bonds secures only USD

98 of cash — protecting the lender against a fall in collateral value or borrower default. Riskier or more volatile collateral carries a larger haircut.

Q946. What are LIBOR, SOFR, and MIBOR, and what changed?

LIBOR was the London Interbank Offered Rate, a survey-based benchmark for unsecured interbank lending that was discredited by manipulation and phased out (most settings ended by mid-2023). SOFR (Secured Overnight Financing Rate) replaced USD LIBOR; it is a transaction-based rate derived from overnight US Treasury repo, so it is secured and nearly risk-free. MIBOR (Mumbai Interbank Offered Rate) is India's benchmark for the rupee overnight/call market, now computed by FBIL from actual transactions.

Q947. What is the difference between the money market and the capital market?

The money market trades short-term instruments (under one year) for liquidity, with low risk and low return — T-bills, CP, CDs, repo. The capital market trades long-term instruments (over one year) — bonds and equities — to raise investment capital, with higher risk and return. Money markets are about managing cash; capital markets are about funding long-term growth.

Q948. How do you compute the discount yield on a T-bill, and how does it differ from bond-equivalent yield?

The discount yield (bank discount basis) is $(\text{Face} - \text{Price}) / \text{Face} * (360 / \text{days})$, which understates true return because it divides by face value and uses a 360-day year. The bond-equivalent yield (BEY) is $(\text{Face} - \text{Price}) / \text{Price} * (365 / \text{days})$, dividing by the price actually paid and using 365 days, giving a more accurate annualized return. Example: a 91-day bill bought at 98 (face 100) has discount yield = $2 / 100 * 360 / 91 = 7.91$ percent, but BEY = $2 / 98 * 365 / 91 = 8.18$ percent.

I. Economy & Financial Institutions

11. Macroeconomics, Monetary & Fiscal Policy (Q949-Q978)

Q949. What are the four components of GDP under the expenditure approach?

$GDP = C + I + G + NX$, where C is private consumption, I is gross investment (business capex plus inventory changes plus residential construction), G is government spending on goods and services (not transfers), and NX is net exports (exports minus imports). In most economies C is the largest slice; for India, consumption is roughly 55-60% of GDP and investment (gross fixed capital formation) is the key growth swing factor. Note G excludes transfer payments like pensions or subsidies because those are not purchases of current output.

Q950. GDP can be measured three ways. What are they and why do they give the same number?

The three approaches are expenditure (sum of all spending: $C+I+G+NX$), income (sum of all factor incomes: wages + rent + interest + profits + taxes-less-subsidies), and output/value-added (sum of value added at each production stage, avoiding double counting). They are identical because every rupee of output sold becomes someone's income and someone's expenditure — one entity's spending is another's revenue. India's CSO/NSO historically emphasized the production (gross value added, GVA) approach; $GDP = GVA + \text{net taxes on products}$.

Q951. Distinguish nominal GDP from real GDP.

Nominal GDP values output at current-year prices, so it rises with both higher quantities and higher prices. Real GDP values output at constant base-year prices, stripping out inflation so it reflects only changes in the volume of goods and services. Only real GDP growth tells you whether the economy genuinely produced more; if nominal GDP grows 11% while prices rise 6%, real growth is roughly 5%.

Q952. What is the GDP deflator and how does it differ from CPI?

$GDP \text{ deflator} = (\text{Nominal GDP} / \text{Real GDP}) * 100$; it is the broadest price index because it covers every good and service in GDP, with weights that change each year as the output mix changes. CPI tracks only a fixed basket of consumer goods, includes imports, and uses fixed weights, so it can diverge from the deflator. Example: nominal GDP 220, real GDP 200 gives a deflator of 110, implying 10% price growth versus the base year.

Q953. What is the difference between GDP and GNP (and GNI)?

GDP measures output produced within a country's borders regardless of who owns the factors; GNP (now usually called GNI) measures output produced by a country's residents/nationals regardless of location. $GNP = GDP + \text{net factor income from abroad}$ (income residents earn overseas minus income foreigners earn domestically). A country with large remittance inflows or overseas investments (e.g., India, the Philippines) has GNP slightly above GDP; one hosting heavy foreign-owned production (e.g., Ireland) has GNP below GDP.

Q954. Name the phases of the business cycle.

The classic phases are expansion (rising output, employment, incomes), peak (the cycle's high point where growth tops out), contraction/recession (falling output, often defined as two consecutive quarters of negative real GDP growth), and trough (the low point before recovery begins). A full cycle runs trough-to-trough or peak-to-peak. The NBER in the US dates US recessions using a broad set of indicators, not just the two-quarter rule.

Q955. Give examples of leading, coincident, and lagging economic indicators.

Leading indicators turn before the economy does: stock prices, the yield curve slope, building permits, new orders/PMI, and initial jobless claims. Coincident indicators move with the cycle: real GDP, industrial production, employment, and personal income. Lagging indicators turn after: unemployment rate, CPI inflation, and the average duration of unemployment. Markets watch leading indicators (e.g., India's manufacturing PMI) to anticipate turning points.

Q956. Explain demand-pull versus cost-push inflation.

Demand-pull inflation occurs when aggregate demand outpaces the economy's productive capacity — "too much money chasing too few goods" — typical of an overheating boom. Cost-push inflation arises when input costs rise (oil shocks, wage spikes, supply disruptions), shifting aggregate supply left so prices rise even as output falls. The 1970s oil crises and the 2021-22 supply-chain/energy spike were largely cost-push; a credit-fueled boom is demand-pull. Policy is harder for cost-push because curbing demand worsens the output loss.

Q957. How does CPI differ from WPI, and which does India target?

CPI measures retail price changes for a consumer basket and includes services; WPI (Wholesale Price Index) measures prices at the wholesale/producer stage and excludes services, with a heavy weight on manufactured goods and commodities. WPI is more volatile because of commodity swings. India's RBI conducts flexible inflation targeting against CPI (Combined) with a target of 4% +/- 2%; WPI is no longer the policy anchor though it is still tracked.

Q958. What is core inflation and why do central banks watch it?

Core inflation strips out food and energy (and sometimes other volatile items) from headline CPI because those prices are noisy and supply-driven rather than demand-driven. It gives a cleaner read on underlying, persistent inflation pressure that monetary policy can actually influence. The RBI and the Fed both monitor core measures; a falling headline but sticky core (say headline 4.5% but core 6%) signals broad-based price pressure that warrants caution on rate cuts.

Q959. What are the costs of inflation?

Inflation erodes the purchasing power of money and fixed incomes, creates "menu costs" and "shoe-leather costs" of constantly repricing and economizing on cash, distorts the tax system (taxing nominal gains), and arbitrarily redistributes wealth from creditors to debtors. Unanticipated inflation is worse because it injects uncertainty into long-term contracts and raises required risk premia, lifting nominal interest rates. Even moderate inflation distorts relative price signals that markets rely on to allocate resources.

Q960. What is the difference between deflation and stagflation, and why is each dangerous?

Deflation is a sustained fall in the general price level; it is dangerous because consumers defer purchases expecting lower prices, real debt burdens rise, and central banks hit the zero lower bound (Japan's "lost decades"). Stagflation is the toxic combination of stagnant/falling growth plus high inflation plus high unemployment, which defeats standard policy because fighting inflation (tightening) worsens the recession and vice versa. The 1970s OPEC-driven stagflation broke the simple Phillips-curve consensus.

Q961. Define the monetary aggregates Mo, M1, M2, M3.

Mo (reserve/base money) is currency in circulation plus bankers' deposits with the central bank. M1 (narrow money) is currency with the public plus demand/checkable deposits. M2 adds savings and small time deposits. M3 (broad money) adds larger time/term deposits and is the most common policy aggregate. In India: Mo is reserve money; M3 = currency with public + demand deposits + time deposits + other deposits with RBI, and M3 is the headline "broad money" growth figure.

Q962. What is the money multiplier?

The money multiplier links base money to broad money: in a simple model, multiplier = 1 / reserve ratio, so a 10% reserve requirement implies a 10x multiplier (a Rs 100 injection can support up to Rs 1,000 of

deposits). More generally, $\text{multiplier} = (1 + c) / (c + r)$ where c is the currency-to-deposit ratio and r the reserve ratio. Leakages into cash holdings or excess reserves shrink the actual multiplier below the theoretical maximum.

Q963. How do central banks "create money"?

Central banks create base money by crediting commercial banks' reserve accounts — typically when buying assets (government bonds in open market operations or QE) or lending to banks. Commercial banks then create the bulk of broad money endogenously by extending loans, which simultaneously create deposits; the central bank influences but does not mechanically dictate this through reserves and the policy rate. So most money in a modern economy is bank-created deposit money, not physical currency.

Q964. What are the main tools of monetary policy?

The key tools are: (1) the policy interest rate (the price of central-bank liquidity — Fed funds target, RBI repo rate), (2) reserve requirements (the fraction of deposits banks must hold, India's CRR), (3) open market operations (buying/selling government securities to adjust liquidity), and (4) unconventional tools like quantitative easing/tightening and forward guidance at the zero bound. India adds the SLR (statutory liquidity ratio) and a liquidity corridor bounded by the repo and reverse-repo/SDF rates.

Q965. What are CRR and SLR in the Indian context?

CRR (Cash Reserve Ratio) is the percentage of a bank's net demand and time liabilities (NDTL) it must keep as cash reserves with the RBI, earning no interest — a direct lever on lendable liquidity. SLR (Statutory Liquidity Ratio) is the percentage of NDTL banks must hold in approved liquid assets (mainly government securities, gold). Raising CRR drains liquidity and tightens credit; SLR both ensures solvency and creates captive demand for government bonds.

Q966. Explain quantitative easing (QE) and quantitative tightening (QT).

QE is large-scale central-bank purchase of longer-dated government and sometimes private bonds, used when the policy rate is already near zero; it expands the central bank's balance sheet, injects reserves, lowers long-term yields, and pushes investors toward riskier assets (the portfolio-balance and signaling channels). QT is the reverse — shrinking the balance sheet by letting bonds mature or selling them, draining liquidity and putting upward pressure on yields. The Fed ran massive QE in 2008-14 and 2020, then QT from 2022.

Q967. Describe the monetary policy transmission mechanism.

A policy-rate change propagates through several channels: the interest-rate channel (loan and deposit rates adjust, affecting borrowing and saving), the credit/bank-lending channel (changes in bank reserves and risk appetite alter loan supply), the asset-price/wealth channel (rates move equity and bond prices and household wealth), and the exchange-rate channel (rate differentials move the currency, affecting trade). Transmission has lags of several quarters and can be "long and variable"; in India transmission to bank lending rates improved after linking loans to external benchmarks (the EBLR regime).

Q968. What does the Taylor rule say, conceptually?

The Taylor rule is a guideline for setting the policy rate as a function of inflation and the output gap: roughly, $\text{nominal rate} = \text{neutral real rate} + \text{inflation} + 0.5(\text{inflation gap}) + 0.5(\text{output gap})$. It implies that when inflation rises above target the central bank should raise the nominal rate by more than one-for-one (the "Taylor principle") so the real rate increases and actually restrains demand. It is a descriptive benchmark, not a mechanical mandate, but it captures how systematic, rules-based policy should react.

Q969. Contrast the classical, Keynesian, and monetarist schools.

Classical economists believe markets clear quickly, prices/wages are flexible, output is supply-determined at full employment, and money is "neutral" (affects only prices, not real output) — so government intervention is largely unnecessary. Keynesians argue prices/wages are sticky, demand drives short-run output, and active

fiscal/monetary stimulus can close output gaps in recessions. Monetarists (Friedman) accept demand matters short-run but stress that "inflation is always and everywhere a monetary phenomenon" and favor stable, rules-based money-supply growth over discretionary fine-tuning.

Q970. Explain the IS-LM framework conceptually.

IS-LM models short-run equilibrium between the goods market and the money market on axes of interest rate (r) and output (Y). The IS curve (downward-sloping) shows goods-market equilibrium: lower rates spur investment and raise output. The LM curve (upward-sloping) shows money-market equilibrium: higher output raises money demand and thus the rate for a fixed money supply. Their intersection sets r and Y ; expansionary fiscal policy shifts IS right (higher Y and r), while expansionary monetary policy shifts LM right (higher Y , lower r).

Q971. What is the Phillips curve and the inflation-unemployment tradeoff?

The original Phillips curve posited a stable inverse relationship between unemployment and wage/price inflation: pushing unemployment lower stokes inflation. The 1970s stagflation broke the naive version, leading to the expectations-augmented Phillips curve, where the tradeoff exists only short-run and any attempt to hold unemployment below the "natural rate" (NAIRU) just accelerates inflation. Long-run, the curve is vertical at the natural rate — there is no permanent tradeoff, only a short-run one driven by inflation surprises.

Q972. What is the fiscal multiplier and what makes it larger or smaller?

The fiscal multiplier is the ratio of the change in output to the change in government spending or taxes; in a simple Keynesian model, spending multiplier = $1/(1 - MPC)$, so an MPC of 0.8 gives a multiplier of 5. In practice it is much smaller (often well below 1) because of leakages into savings, imports, and taxes, plus crowding out and offsetting monetary policy. Multipliers are larger in recessions, at the zero lower bound, and for spending targeted at high-MPC (low-income) households; smaller in open economies with high import propensity.

Q973. What is crowding out?

Crowding out occurs when increased government borrowing to fund a deficit raises interest rates, which discourages (crowds out) private investment and consumption, partially offsetting the fiscal stimulus. The effect is strong when the economy is near full capacity and central banks respond by tightening; it is weak in a deep recession with slack resources and accommodative monetary policy. There is also "financial" crowding out, where government bond demand absorbs savings that could have funded private projects.

Q974. What are automatic stabilizers?

Automatic stabilizers are features of the budget that dampen the cycle without any new legislation: in a downturn, tax receipts fall (progressive income tax) and transfer spending rises (unemployment benefits, welfare), automatically supporting demand; in a boom the reverse cools the economy. They make fiscal policy countercyclical with no decision lag, unlike discretionary stimulus. India relies more on discretionary measures because its safety-net coverage is narrower, though schemes like MGNREGA act as a partial stabilizer.

Q975. Why does the debt-to-GDP ratio matter, and what drives its trajectory?

Debt-to-GDP scales government debt by the economy's ability to service it; a rising ratio raises default/inflation risk premia and interest costs. Its dynamics hinge on the relationship between the average interest rate (r) on debt and the nominal growth rate (g): if g exceeds r , an economy can run modest primary deficits and still see the ratio fall, whereas r above g makes debt self-compounding. India's general government debt is roughly 80-85% of GDP; the FRBM framework targets gradual consolidation toward a debt anchor.

Q976. What is the balance of payments, and how do the current and capital accounts differ?

The balance of payments records all transactions between residents and the rest of the world and always sums to zero (every flow is financed). The current account covers trade in goods and services, primary income (interest/dividends), and secondary income (remittances/transfers); a deficit means the country imports more than it exports of these. The capital and financial account records cross-border asset flows (FDI, portfolio investment, loans, reserve changes); a current-account deficit must be financed by a capital-account surplus. India typically runs a current-account deficit funded by capital inflows.

Q977. Compare fixed, floating, and managed-float exchange-rate regimes.

Under a fixed (pegged) regime the central bank commits to a specific exchange rate and uses reserves and policy to defend it, gaining stability but losing monetary autonomy. Under a pure float the rate is set entirely by the market, giving full monetary independence but more volatility. A managed float (dirty float) lets the market set the rate but with periodic central-bank intervention to smooth volatility — India's regime, where the RBI lets the rupee move but intervenes to curb sharp swings and build reserves.

Q978. What is the impossible trinity, and how does it connect to the yield curve as a recession signal?

The impossible trinity (Mundell-Fleming trilemma) states a country can have at most two of: a fixed exchange rate, free capital mobility, and an independent monetary policy — never all three. India, for instance, chooses partial capital controls and a managed float to retain meaningful monetary autonomy. Separately, an inverted yield curve (short-term yields above long-term, e.g., the US 10-year minus 2-year going negative) has preceded most US recessions because it signals markets expect future rate cuts amid slowing growth; in India twin deficits (simultaneous fiscal and current-account deficits) similarly flag macro vulnerability, and rate hikes transmit to markets by lifting bond yields (pushing bond prices down) and compressing equity valuations through a higher discount rate.

12. Banking & Financial Institutions (Q979-Q1008)

Q979. How does a commercial bank fundamentally make money?

A bank makes money in two broad ways: net interest income (NII) and non-interest (fee) income. NII is the spread between interest earned on assets (loans, securities) and interest paid on liabilities (deposits, borrowings); fee income comes from things like account charges, card interchange, processing fees, trade finance, and wealth/distribution commissions. The core model is borrowing short and cheap (deposits) and lending long and dearer (loans), capturing the spread while managing credit, liquidity and rate risk.

Q980. What is Net Interest Income (NII) and how is it different from Net Interest Margin (NIM)?

$NII = \text{Interest Income} - \text{Interest Expense}$, an absolute rupee/dollar amount. NIM expresses that spread as a percentage of average earning assets: $NIM = (\text{Interest Income} - \text{Interest Expense}) / \text{Average Earning Assets}$. NII tells you the size of the spread; NIM normalises it so you can compare banks of different sizes. A large bank can have huge NII but a thin NIM, and vice versa.

Q981. Compute NIM for a simple example.

If a bank earns interest income of 100 and pays interest expense of 55 on average earning assets of 1,000, then $NII = 100 - 55 = 45$ and $NIM = 45 / 1,000 = 4.5\%$. Indian private banks like HDFC/Kotak historically run NIMs around 4-4.5%, while public sector banks tend to be lower (2.5-3%) due to higher cost of funds and lower-yielding asset mixes.

Q982. Why are deposits a bank's liabilities and loans its assets — isn't that backwards from how a customer sees it?

From the customer's view a deposit is "my money" (an asset to them), but to the bank it is money it owes back, so it sits on the liability side. A loan is "my debt" to the borrower but to the bank it is a claim that will be repaid with interest, hence an asset. The bank's balance sheet is essentially the mirror image of its customers': customer deposits fund customer loans.

Q983. Why is a bank's balance sheet considered "special" and highly leveraged?

Banks operate with very thin equity relative to assets — typically equity is only about 8-12% of assets, meaning leverage of roughly 8x to 12x (assets/equity). They fund mostly with other people's money (deposits and borrowings) and hold illiquid loans against liquid, on-demand liabilities. This maturity transformation plus high leverage makes banks profitable in good times but fragile to losses and runs, which is why they are heavily regulated.

Q984. What is the CASA ratio and why do banks obsess over it?

CASA = Current Account and Savings Account deposits as a share of total deposits: $\text{CASA Ratio} = (\text{CA} + \text{SA balances}) / \text{Total Deposits}$. Current accounts pay zero interest and savings accounts pay very little, so they are low-cost, sticky funds. A high CASA ratio (say 45%+) lowers the bank's overall cost of funds and structurally lifts NIM, which is why CASA-rich banks like HDFC, Kotak and SBI enjoy a funding advantage over deposit-hungry rivals.

Q985. How does a high CASA ratio translate into better margins, numerically?

Suppose a bank funds 100 of loans yielding 10%. If 50 is CASA at 1% and 50 is term deposits at 7%, blended cost = $(50 \cdot 1\% + 50 \cdot 7\%) / 100 = 4\%$, so spread = 6%. If CASA falls and the mix becomes 20 CASA / 80 term, cost = $(20 \cdot 1\% + 80 \cdot 7\%) / 100 = 5.8\%$, spread shrinks to 4.2%. Same loan yield, but the cheaper deposit mix adds nearly 2 percentage points of margin.

Q986. What is Asset-Liability Management (ALM) in a bank?

ALM is the practice of managing the structure and risk of the balance sheet so that mismatches between assets and liabilities — in maturity, repricing, currency and liquidity — stay within acceptable limits. Its main goals are to protect net interest income and economic value from interest-rate moves and to ensure the bank can always meet cash outflows. ALM is overseen by an Asset-Liability Committee (ALCO) using tools like gap analysis, duration and liquidity buckets.

Q987. What is interest-rate risk in the banking book (IRRBB) and the maturity/duration gap?

IRRBB is the risk that changes in interest rates hurt the bank's earnings or economic value because assets and liabilities reprice at different times. The repricing gap = rate-sensitive assets minus rate-sensitive liabilities in a given time bucket; a positive gap means rising rates help NII, a negative gap means rising rates hurt it. Duration gap extends this to economic value: $\text{Duration Gap} = D_{\text{assets}} - (L/A) \cdot D_{\text{liabilities}}$, and a positive duration gap means the bank's net worth falls when rates rise.

Q988. Give a quick duration-gap intuition with numbers.

If assets of 100 have duration 4 years and liabilities of 92 have duration 1.5 years, duration gap = $4 - (92/100) \cdot 1.5 = 4 - 1.38 = 2.62$ years. For a +1% (100 bps) rate rise, the change in economic value of equity is approximately $-\text{Duration Gap} \cdot (\text{rate change}) \cdot \text{Assets} = -2.62 \cdot 0.01 \cdot 100 = -2.62$, i.e. equity value drops by about 2.62. A large positive duration gap is exactly what sank some banks (e.g. SVB in 2023) when rates spiked.

Q989. What is liquidity management and how does it differ from solvency?

Liquidity management ensures the bank can meet cash outflows (deposit withdrawals, loan drawdowns) as they fall due, by holding liquid assets and maintaining funding access. Solvency is about having enough capital so assets exceed liabilities. A bank can be solvent (assets > liabilities) yet fail from a liquidity crisis if

it cannot convert assets to cash fast enough — which is precisely how bank runs kill otherwise viable banks. Basel introduced the LCR and NSFR specifically to backstop liquidity.

Q990. Walk through the typical loan/credit process at a bank.

It runs from origination and application, through credit appraisal (assessing repayment capacity, collateral and risk), to sanction with terms and pricing, documentation, disbursement, and then ongoing monitoring and collection/recovery. The pricing reflects the borrower's risk grade, cost of funds, operating cost and a risk premium. Good underwriting at appraisal is what protects the bank, because once a bad loan is disbursed, recovery is costly and uncertain.

Q991. What are the 5 Cs of credit?

Character (borrower's track record and willingness to repay), Capacity (cash flow/ability to service debt), Capital (the borrower's own skin in the game), Collateral (security pledged), and Conditions (purpose of loan and the macro/industry environment). Lenders use the 5 Cs as a structured checklist to judge both the willingness and the ability to repay. Some add a sixth C — Credit history — though it overlaps with Character.

Q992. What is a Non-Performing Asset (NPA)?

An NPA is a loan or advance where the borrower has stopped servicing it — under RBI norms, a term loan becomes an NPA when principal or interest is overdue for more than 90 days. Once classified NPA, the bank must stop booking accrued interest as income and start setting aside provisions. NPAs directly erode profitability (lost interest plus provisions) and signal asset-quality stress.

Q993. How does RBI's IRAC framework classify NPAs?

Under the Income Recognition and Asset Classification (IRAC) norms, assets are graded as Standard (performing), Sub-standard (NPA for up to 12 months), Doubtful (NPA beyond 12 months), and Loss (uncollectible, identified as a write-off). The classification drives provisioning — the worse the grade, the higher the provision. A loan also passes through SMA (Special Mention Account) stages (SMA-0/1/2 at 0-30/31-60/61-90 days overdue) as an early-warning bucket before turning NPA at 90 days.

Q994. What is provisioning, and roughly how much is required by category?

A provision is an amount charged to the P&L and held against expected loss on a loan. Indicative RBI minimums: Standard assets carry a small general provision (around 0.25-1% by sector); Sub-standard ~15% (25% if unsecured); Doubtful escalates with age from ~25% up to 100% on the secured portion plus 100% on the unsecured portion; Loss assets 100%. Banks often hold provisions above the regulatory floor, tracked via the Provision Coverage Ratio (PCR).

Q995. Distinguish Gross NPA from Net NPA.

Gross NPA is the total stock of bad loans before provisions; Net NPA = Gross NPA - provisions held against them. As ratios: Gross NPA Ratio = Gross NPAs / Gross Advances and Net NPA Ratio = Net NPAs / Net Advances. Net NPA shows the residual risk the bank's capital still carries; a bank with high Gross NPA but near-zero Net NPA has already absorbed most of the hit through provisioning.

Q996. What do "slippage" and "restructuring" mean in NPA management?

Slippage is the fresh addition of loans moving from standard (performing) into NPA during a period; the slippage ratio measures new NPAs as a fraction of opening standard advances and indicates the pace of asset-quality deterioration. Restructuring is renegotiating loan terms (lower rate, longer tenor, moratorium) to help a stressed but viable borrower keep performing. Restructuring can mask stress if overused, so RBI tightened rules (e.g. post-2015 AQR and the June 2019 stressed-asset framework) to prevent "evergreening".

Q997. What is the credit cycle and why does it matter for banks?

The credit cycle is the recurring expansion and contraction of credit availability: in upswings banks lend freely, underwriting standards loosen and asset prices rise; in downswings defaults climb, banks pull back, and provisioning spikes. Bank earnings are highly cyclical because loan losses are back-loaded — the bad loans made in the boom only surface in the bust. This is why bank valuations swing with the cycle and why "the loans of the best years go bad in the worst years".

Q998. Why is the central bank called the lender of last resort?

In a panic, even solvent banks can run out of cash if depositors withdraw en masse; the central bank lends to such banks against good collateral to stop the panic spreading — the classic Bagehot rule is to lend freely, at a penalty rate, against good collateral. In India the RBI provides this via facilities like the Marginal Standing Facility (MSF) and the LAF window. The backstop exists because a liquidity failure at one bank can cascade into a systemic crisis.

Q999. Explain fractional reserve banking and how it leads to bank runs.

Banks keep only a fraction of deposits as reserves and lend out the rest, so they cannot repay all depositors simultaneously — the money is tied up in illiquid loans. If enough depositors fear losses and rush to withdraw at once, the bank cannot meet demand and fails, even if fundamentally solvent; the fear becomes self-fulfilling. Deposit insurance and the lender-of-last-resort exist precisely to break this run dynamic.

Q1000. What is deposit insurance, and how does DICGC work in India?

Deposit insurance guarantees depositors a set amount if their bank fails, removing the incentive to run. In India the DICGC (a RBI subsidiary) insures each depositor up to INR 5 lakh per bank (raised from INR 1 lakh in 2020), covering principal plus interest across all accounts at that bank. In the US the FDIC insures up to USD 250,000 per depositor per bank. The cap is deliberate so large/sophisticated creditors still exert market discipline.

Q1001. Compare commercial, investment, and universal banks.

A commercial bank takes deposits and makes loans (retail and corporate). An investment bank does not take retail deposits; it underwrites securities, advises on M&A, and trades — earning fees and trading income rather than NII. A universal bank combines both under one roof (e.g. JPMorgan, ICICI Bank-style groups). After 1999 (Gramm-Leach-Bliley repeal of Glass-Steagall) US banks could combine these, while post-2008 rules like Volcker re-restricted proprietary trading.

Q1002. What distinguishes cooperative, small finance, and payments banks in India?

Cooperative banks are member-owned, locally focused, and dually regulated (RBI + state registrars), historically prone to governance lapses (e.g. PMC Bank). Small Finance Banks (SFBs like AU, Equitas, Ujjivan) are full banks aimed at financial inclusion, required to direct a high share of lending to priority sectors and small-ticket loans. Payments Banks (e.g. Airtel, Paytm, India Post) can take deposits up to INR 2 lakh per customer and offer payments/remittances but cannot lend or issue credit cards — they earn on float and fees, not NII.

Q1003. How does an NBFC differ from a bank, and why is that riskier on funding?

A Non-Banking Financial Company lends like a bank but cannot accept demand deposits and has no access to CASA or the central bank's liquidity window. It funds itself from banks, bonds, commercial paper and debentures — wholesale, costlier and shorter-dated money. This creates a structural asset-liability mismatch: NBFCs often lend long (infrastructure, housing) against short-term borrowings, so a sudden loss of market confidence can freeze their funding even if their loan book is sound.

Q1004. What were the IL&FS and DHFL episodes, and what lessons did they teach?

IL&FS (2018) defaulted on its debt despite high credit ratings, exposing hidden leverage, governance failures and rating-agency complacency; it triggered a liquidity crunch across NBFCs. DHFL (2019) followed with

fraud and ALM mismatch, becoming the first NBFC taken to insolvency under the IBC. The lessons: short-term borrowing to fund long-term assets is fragile, ratings can mislead, and NBFCs are vulnerable to confidence-driven funding freezes — prompting RBI's scale-based regulation and tighter liquidity (LCR) norms for large NBFCs.

Q1005. What is Priority Sector Lending (PSL) in India?

PSL is an RBI mandate requiring banks to lend a minimum share of credit to sectors deemed nationally important — agriculture, MSMEs, education, housing, weaker sections, renewables. The overall target is 40% of Adjusted Net Bank Credit for domestic banks, with sub-targets (e.g. 18% to agriculture). Banks short of target buy Priority Sector Lending Certificates (PSLCs) or park funds with NABARD/SIDBI at lower yields, so PSL acts as a directed-credit and cross-subsidy mechanism.

Q1006. What are financial intermediation and disintermediation?

Financial intermediation is banks standing between savers and borrowers — pooling deposits, assessing credit, transforming maturity and bearing risk — so neither side has to find the other directly. Disintermediation is the bypassing of banks: borrowers raise money directly from markets (bonds, commercial paper) and savers invest directly (mutual funds, equities), shrinking the bank's traditional spread business. Fintech, capital-market deepening and rising investor sophistication all accelerate disintermediation.

Q1007. Why does a bank's ROE look high while its ROA looks low?

$ROA = \text{Net Income} / \text{Average Assets}$ and $ROE = \text{Net Income} / \text{Average Equity}$, linked by leverage: $ROE = ROA * (\text{Assets}/\text{Equity})$. Banks earn a thin ROA — around 1% is considered healthy — but run leverage of ~10x, so a 1% ROA times 10x equals roughly a 10-15% ROE. The high ROE is mostly leverage, not superior profitability per dollar of assets, which is why analysts watch ROA to judge true earning quality.

Q1008. Why do banks trade on Price-to-Book (P/B) rather than P/E, and how is book value relevant; plus a word on insurance and fintech?

Banks' assets and liabilities are largely financial and marked close to fair value, so book value (equity per share) is a meaningful anchor, and P/B (often paired with ROE — high-ROE banks deserve P/B above 1) captures value better than volatile, provision-driven earnings. By analogy, insurers earn on "float" — premiums held before claims are paid — splitting profit into underwriting income (premiums minus claims/costs) and investment income on the float, with a combined ratio below 100% meaning underwriting profit. Fintechs disrupt banking by unbundling high-margin pieces (payments, lending, wealth) with lower cost-to-serve and better UX, pushing banks toward platform models while still relying on banks for the regulated deposit-and-balance-sheet core.

J. Advanced Corporate Finance & Contemporary Finance

J1. Advanced Corporate Finance (Q1009-Q1039)

Q1009. What is Economic Value Added (EVA) and how does it differ from accounting profit?

EVA is economic profit: $EVA = NOPAT - WACC \times \text{Invested Capital}$, where NOPAT is net operating profit after tax. Unlike accounting profit (net income), which only deducts interest and ignores the cost of equity, EVA charges for ALL capital — debt and equity — so a firm can report positive net income yet destroy value if it does not out-earn its full cost of capital. Example: NOPAT = INR 120 cr, invested capital = INR 1,000 cr, WACC = 11% gives $EVA = 120 - 0.11 \times 1000 = \text{INR } 10 \text{ cr}$ of true value created.

Q1010. Why can a company show accounting profit but still have negative economic profit?

Accounting net income deducts the explicit cost of debt (interest) but never charges for the opportunity cost of equity, which is real and often the largest financing cost. Economic profit subtracts a capital charge on the entire invested base. A firm earning a 9% return on capital that cost 12% has positive net income but is destroying roughly 3% of its capital base each year — the gap accounting statements hide.

Q1011. State the ROIC vs WACC value-creation test and what it tells you.

$ROIC = NOPAT / \text{Invested Capital}$; value is created only when $ROIC > WACC$. The spread ($ROIC - WACC$) times invested capital equals EVA. If $ROIC = WACC$ the firm merely earns its cost of capital and growth adds no value; if $ROIC < WACC$, growth actively destroys value. This single test reframes capital budgeting: the goal is not growth or even accounting returns, but earning a positive spread over the cost of capital.

Q1012. How does the value driver formula link growth, ROIC and value?

The key driver formula for value is $\text{Value} = NOPAT \times (1 - g/ROIC) / (WACC - g)$, where g is the growth rate of NOPAT. The term $(1 - g/ROIC)$ is the reinvestment-adjusted free cash flow. When $ROIC > WACC$, raising g increases value; when $ROIC = WACC$ the g terms cancel and value reduces to $NOPAT/WACC$ regardless of growth; when $ROIC < WACC$, faster growth lowers value. So growth and ROIC are joint drivers, not independent ones.

Q1013. "A high-growth company is always more valuable." Critique this.

Growth only creates value when invested at returns above the cost of capital. A company growing 15% at ROIC of 8% against a 12% WACC is consuming cash on negative-spread projects and is worth less than a slower-growing, high-ROIC peer. Growth is a multiplier on the ROIC-WACC spread: it amplifies value creation when the spread is positive and amplifies destruction when negative.

Q1014. Which is more valuable to improve when ROIC already exceeds WACC: growth or the ROIC spread?

It depends on diminishing returns, but mathematically improving ROIC raises value through two channels — a higher current spread AND a lower reinvestment need for any given growth (less capital tied up per unit of NOPAT). For low-ROIC firms (ROIC near WACC), lifting ROIC is far more powerful than chasing growth; for already high-ROIC firms, profitable growth becomes the dominant lever since each incremental rupee of investment earns a wide spread.

Q1015. What is a real option in capital budgeting and why does NPV understate it?

A real option is managerial flexibility to alter a project after new information arrives — the option to expand, abandon, defer, or switch. Standard NPV assumes a fixed, now-or-never cash-flow plan and ignores the asymmetric value of acting on good news while limiting downside on bad news. Because option value rises

with uncertainty, NPV systematically undervalues flexible projects: $\text{Strategic NPV} = \text{Static NPV} + \text{Value of Real Options}$.

Q1016. Describe the option to expand and how it changes a project decision.

The option to expand lets management scale up if early results are favourable — like a call option on follow-on investment. A pilot plant with a negative standalone NPV may be worth funding because it creates the right (not obligation) to build full capacity later if demand proves strong. The pilot's true value = its own NPV + value of the embedded expansion call; ignoring the latter wrongly kills strategically valuable seed investments.

Q1017. Explain the option to abandon and the option to defer.

The option to abandon is a put: if a project sours, management can sell assets or shut down to recover salvage value, truncating the downside — most valuable when assets have high resale/redeployment value. The option to defer is a call on waiting: delaying an irreversible investment until uncertainty resolves can beat investing today, even with positive NPV, because waiting preserves the chance to avoid a bad commitment. Both raise value as uncertainty increases.

Q1018. What drives the value of a real option, and which inputs matter most?

Real options are valued with option-pricing logic (Black-Scholes or binomial), so the key drivers are the underlying asset value, the exercise cost (investment), time to expiry (decision window), and especially volatility of the underlying cash flows. Higher uncertainty raises option value because it widens the favourable upside while the downside is capped by not exercising. This is why high-uncertainty R&D, mining, and pharma projects are best evaluated with a real-options lens.

Q1019. What is the profitability index and how is it used under capital rationing?

$\text{PI} = \text{PV of future cash flows} / \text{Initial Investment} = 1 + \text{NPV}/\text{Investment}$. Under capital rationing (a hard budget limit), you cannot fund every positive-NPV project, so you rank by PI to maximise NPV per rupee of scarce capital rather than total NPV. Example: a project with NPV INR 30 cr on INR 100 cr (PI 1.30) beats one with NPV INR 40 cr on INR 200 cr (PI 1.20) per unit of capital.

Q1020. When can simple PI ranking fail under capital rationing?

PI ranking is exact only for perfectly divisible projects and a single-period constraint. With indivisible (whole-or-nothing) projects, the highest-PI set may leave budget stranded, so a lower-PI combination that fully uses the budget can yield higher total NPV. Multi-period rationing requires integer/linear programming to maximise total NPV subject to each period's cash constraint, since PI cannot capture inter-temporal trade-offs.

Q1021. Distinguish sensitivity, scenario, and Monte-Carlo analysis in capital budgeting.

Sensitivity analysis flexes ONE input at a time (e.g., NPV vs unit price) to find which variable the project is most exposed to — but ignores correlations. Scenario analysis moves a consistent SET of inputs together into coherent cases (base/best/worst), capturing variable interactions but only a few discrete outcomes. Monte-Carlo assigns probability distributions to inputs, samples thousands of times, and produces a full NPV distribution with probabilities (e.g., $P(\text{NPV} < 0) = 18\%$), capturing the complete risk profile.

Q1022. What is the cash conversion cycle and how is it optimized?

$\text{CCC} = \text{DIO} + \text{DSO} - \text{DPO}$ (days inventory outstanding + days sales outstanding - days payables outstanding) measures how long cash is tied up in operations. Lowering CCC frees cash: reduce inventory (JIT, better forecasting), collect receivables faster (discounts, tighter terms), and stretch payables without harming supplier relationships. A retailer moving CCC from 60 to 40 days on INR 1,000 cr daily-equivalent sales releases roughly 20 days of working capital in cash.

Q1023. Compute the implied annual cost of forgoing a "2/10 net 30" trade-credit discount.

Forgoing the 2% discount means paying full price 20 days later (day 30 instead of day 10). Cost per period = $2/98 = 2.04\%$; annualised: $(1 + 2/98)^{(365/20)} - 1 \approx 44.6\%$ effective (or about 37.2% simple = $2.04\% \times 365/20$). Because this implicit rate dwarfs most borrowing costs, a firm with access to cheaper financing should almost always take the discount and pay on day 10.

Q1024. What is factoring and how does it affect working capital?

Factoring is selling accounts receivable to a third party (factor) at a discount for immediate cash, accelerating collections and shrinking DSO. In recourse factoring the seller bears default risk; in non-recourse the factor absorbs it for a higher fee. The cost is the discount/fee plus interest for the advance period; it suits firms needing liquidity faster than customers pay, trading some margin for cash and offloaded collection effort. In India, TReDS platforms (e.g., RXIL, M1xchange) digitise MSME receivable financing.

Q1025. Explain supply-chain finance (reverse factoring) and why both parties benefit.

In supply-chain finance the BUYER initiates the arrangement: a financier pays the supplier early at a discount based on the buyer's (typically stronger) credit rating, and the buyer repays the financier on the original due date. The supplier gets cheap early cash (priced off the buyer's low credit spread), while the buyer can extend its own DPO — improving both parties' cash conversion cycles. It differs from factoring, which is supplier-initiated and priced on the supplier's weaker credit.

Q1026. How does just-in-time inventory improve financial performance, and what is the risk?

JIT minimises inventory by receiving goods only as needed, cutting DIO, storage/obsolescence costs, and the working capital locked in stock — directly raising ROIC by shrinking invested capital. The trade-off is fragility: thin buffers leave the firm exposed to supply shocks, demand spikes, and logistics disruptions (as global supply-chain crises exposed). Optimal policy balances lower carrying cost against higher stock-out and disruption risk.

Q1027. Describe the manager-shareholder agency problem and two mechanisms that mitigate it.

Managers (agents) may pursue private benefits — empire-building, perks, excessive cash hoarding, low-risk quiet life — at the expense of shareholder (principal) value. Mechanisms include performance-linked pay/equity (aligning manager wealth with the share price), independent boards and active monitoring, debt (which disgorges free cash via fixed obligations — Jensen's free-cash-flow control), the market for corporate control (takeover threat), and disclosure. Each narrows the gap between manager incentives and owner value.

Q1028. Explain the shareholder-bondholder agency conflict and give two specific examples.

Once debt is in place, shareholders (via managers) can transfer wealth from bondholders. Asset substitution: swap safe assets for risky ones — shareholders capture the upside while bondholders bear added default risk. Debt overhang/underinvestment: shareholders skip positive-NPV projects whose gains accrue mainly to existing creditors. Other forms include excessive dividends and claim dilution. Bondholders price this in via higher yields and protective covenants (dividend limits, leverage caps).

Q1029. What governance mechanisms reduce agency costs, including India-specific ones?

Globally: independent boards, audit committees, executive equity, say-on-pay, takeover markets, and creditor covenants. India layers SEBI's LODR Regulations (mandatory independent directors, separation of chair/CEO for large firms, related-party-transaction approval), Companies Act 2013 governance norms, mandatory auditor rotation, and SEBI's stewardship code for institutional investors. These specifically target the concentrated-promoter ownership common in India, where the key conflict is controlling vs minority shareholders rather than diffuse manager-owner separation.

Q1030. State the static trade-off theory of capital structure and value the tax shield.

Trade-off theory says the optimal debt level balances the marginal value of the interest tax shield against the marginal expected cost of financial distress. For permanent debt, the tax-shield value is $T_c \times \text{Debt}$ (e.g., $T_c = 25\%$, $\text{Debt} = \text{INR } 400 \text{ cr}$ gives $\text{INR } 100 \text{ cr}$ of value). Optimal leverage is where the present value of incremental tax shields equals the present value of incremental distress costs — implying an interior, firm-specific target ratio, not maximum debt.

Q1031. What are the costs of financial distress, direct and indirect?

Direct costs are out-of-pocket: legal, advisory, court, and administrative fees of restructuring or bankruptcy (often a small % of asset value). Indirect costs are larger and subtler: lost customers and suppliers (who fear non-performance), key-employee flight, fire-sale asset disposals, distracted management, foregone positive-NPV investment (debt overhang), and tighter, costlier financing. Indirect costs scale with asset intangibility and reliance on stakeholder confidence, which is why R&D/brand-heavy firms carry less debt.

Q1032. Why do profitable, asset-light tech firms often carry little debt despite the tax shield?

Trade-off logic predicts low leverage when expected distress costs are high relative to tax benefits. Asset-light firms have few tangible assets to pledge and high indirect distress costs (value lives in people, brand, and growth options that evaporate in distress), so the marginal cost of debt is steep. Their volatile cash flows raise default probability, and many have ample internal cash, lowering the need for debt's discipline — pushing the optimum toward equity, consistent with both trade-off and pecking-order theories.

Q1033. Compare US Chapter 11 reorganisation with India's IBC/NCLT process.

US Chapter 11 is debtor-in-possession: existing management generally stays and runs the firm while proposing a reorganisation plan, with an automatic stay on creditors. India's Insolvency and Bankruptcy Code (2016) is creditor-in-control: on admission by the NCLT, an Interim/Resolution Professional displaces the board, a Committee of Creditors drives the Corporate Insolvency Resolution Process, and a resolution plan must be approved within a target $\sim 180+90$ (now up to 330) days or the firm goes to liquidation. IBC thus shifts control to creditors far more decisively than Chapter 11.

Q1034. Explain the priority of claims (the "waterfall") in liquidation, including India's IBC order.

Liquidation proceeds follow absolute priority — senior claims are paid in full before junior ones receive anything: secured creditors, then unsecured creditors, then preferred equity, then common equity last. Under IBC Section 53, the waterfall is: insolvency-resolution/liquidation costs first, then workmen's dues (24 months) and secured creditors who relinquish security *pari passu*, then other employee wages, unsecured financial creditors, government dues and remaining secured-creditor shortfalls, then any remaining debts, preference shareholders, and finally equity shareholders.

Q1035. What does a financing choice signal under information asymmetry (signaling theory)?

When managers know more than investors, the financing they choose conveys information. Issuing equity signals managers may believe the stock is overvalued (adverse selection), so announcements often push the price DOWN; issuing debt or doing a buyback signals confidence that cash flows can service obligations or that the stock is cheap, and is read positively. This underpins the pecking order: internal funds first, then debt, then equity as a last resort to avoid the negative equity-issue signal.

Q1036. When does a share buyback create value versus a dividend, and what is the EPS effect?

A buyback creates value when shares are repurchased below intrinsic value, transferring wealth to remaining holders, and offers flexibility (no implied commitment) plus tax efficiency where capital gains are taxed lighter than dividends. Mechanically, buying back shares cuts the share count, so EPS rises if the after-tax earnings yield foregone on cash is below the inverse P/E — but a buyback above fair value destroys value regardless of the cosmetic EPS lift. In India, buybacks are now taxed in the company's hands (post-2024 changes shifting the burden), altering the dividend-vs-buyback calculus.

Q1037. What is project finance with an SPV, and why is it structured that way?

Project finance funds a large standalone asset (toll road, power plant, port) through a special purpose vehicle whose debt is repaid primarily from the project's own cash flows, with limited or no recourse to the sponsors' balance sheets. The SPV ring-fences risk (sponsor failure does not sink the project and vice versa), enables high leverage against contracted cash flows, and isolates lenders' claims to project assets. It suits capital-intensive infrastructure with predictable, often contracted revenues — common in India's PPP and renewable-energy financing.

Q1038. What is a leveraged recapitalisation and why might a firm do one?

A leveraged recap raises a large slug of new debt and returns the proceeds to shareholders via a special dividend or buyback, sharply increasing leverage without changing the assets. Motives: capture more interest tax shields, return excess cash, impose debt discipline on management (reducing free-cash-flow agency costs), and sometimes deter a takeover by loading up debt. The trade-off is higher financial risk and distress probability, so it suits stable-cash-flow firms that were previously underleveraged.

Q1039. Why do firms hold cash despite the drag on ROIC, and what is the optimal cash policy?

Firms hold cash for the transactions motive (smooth operating outflows), the precautionary motive (buffer against shocks and funding-market freezes), and to preserve financial flexibility for future investment/acquisition options without issuing securities under bad conditions. Cash earns below the cost of capital and dilutes ROIC, and excess cash invites agency problems (wasteful spending), so the optimum balances the convenience/insurance value of liquidity against this carrying cost. High-uncertainty, R&D-intensive, and financially constrained firms optimally hold more cash.

J2. ESG, Sustainable Finance & Fintech (Q1040-Q1064)

Q1040. What does ESG stand for, and what kinds of factors fall under each pillar?

ESG stands for Environmental, Social, and Governance — a framework for assessing a company's non-financial risks and impacts alongside financial analysis. Environmental covers carbon emissions, energy use, water, waste, and biodiversity; Social covers labour practices, health and safety, diversity, community relations, and product safety; Governance covers board independence and structure, executive pay, shareholder rights, audit quality, and anti-corruption controls. The premise is that these factors can be financially material and that integrating them improves risk-adjusted decisions, not merely ethics.

Q1041. What is the difference between negative/exclusionary screening and positive/best-in-class screening?

Negative (exclusionary) screening removes companies or sectors that fail an ethical or risk test — for example excluding tobacco, weapons, thermal coal, or gambling from a portfolio. Positive (best-in-class) screening instead keeps a sector but selects the companies with the strongest ESG performance relative to their peers, e.g. holding the best-rated oil major rather than excluding the whole energy sector. Negative screening shrinks the investable universe by rule; best-in-class tilts within it toward leaders.

Q1042. Name the main approaches to ESG/responsible investing.

The commonly cited approaches are: (1) negative/exclusionary screening; (2) positive/best-in-class selection; (3) ESG integration — systematically factoring ESG data into traditional financial analysis and valuation; (4) thematic investing — targeting a theme like clean energy or water; (5) impact investing — seeking measurable, intentional social/environmental outcomes alongside return; and (6) active ownership/engagement and proxy voting. Many funds blend several. These categories are reflected in the UN PRI framework and surveys by bodies like the GSIA (Global Sustainable Investment Alliance).

Q1043. How does impact investing differ from ESG integration?

ESG integration is a risk-management lens: it incorporates ESG data to sharpen valuation and avoid mispriced risks, but the goal remains financial return. Impact investing is intentional — it deliberately seeks measurable, additional social or environmental outcomes (e.g. clean water access, affordable housing) and holds itself accountable for measuring that impact, while still targeting a financial return that can range from below-market (concessionary) to market rate. The distinguishing features of impact investing are intentionality, additionality, and measurement of the outcome.

Q1044. What is greenwashing, and why is it a concern in sustainable finance?

Greenwashing is making a product, fund, or company appear more environmentally or socially responsible than it actually is — for instance labelling a fund "sustainable" while it holds material fossil-fuel exposure, or touting a vague "net zero" pledge with no credible plan. It misleads investors, misallocates capital away from genuinely sustainable activities, and erodes trust in the whole ESG market. Regulators have responded with disclosure rules such as the EU's SFDR fund classifications and SEBI's disclosure norms for ESG mutual fund schemes in India, plus fund-naming requirements that mandate minimum sustainable holdings.

Q1045. Why do ESG ratings often disagree across providers like MSCI, Sustainalytics, and S&P?

Different providers use different frameworks, indicators, weights, and data sources, so the same company can score well with one and poorly with another. Studies (e.g. Berg, Koelbl and Rigobon's "Aggregate Confusion") attribute the divergence mainly to differences in scope (which issues are measured), measurement (how an indicator is proxied), and weighting (how issues are combined). Correlations between major ESG ratings are roughly 0.5-0.6, far below the ~0.99 seen across credit ratings — meaning ESG scores are noisier and less comparable, which complicates benchmarking and can mask greenwashing.

Q1046. What does "materiality" mean in an ESG context, and what is double materiality?

Materiality refers to whether an ESG factor is significant enough to affect a company's financial value or an investor's decision — the same threshold concept used in financial reporting. Single (financial) materiality asks how sustainability issues affect the company's value (outside-in). Double materiality, central to the EU's CSRD, adds the reverse view: how the company's activities affect society and the environment (inside-out). SASB/ISSB standards focus on financially material issues that are often industry-specific — water is material for a beverage maker, data privacy for a software firm.

Q1047. What is a green bond, and what distinguishes it from an ordinary bond?

A green bond is a debt instrument whose proceeds are earmarked (ring-fenced) exclusively for environmentally beneficial projects — renewable energy, clean transport, green buildings, water management. It has the same credit risk and ranking as the issuer's other senior debt; the difference is the use-of-proceeds commitment plus reporting on how the money was deployed and its environmental impact. Most follow ICMA's Green Bond Principles, often with a second-party opinion to verify alignment, and India's SEBI has its own framework for green debt securities.

Q1048. How do sustainability-linked bonds/loans differ from green or social bonds?

Use-of-proceeds bonds (green, social, sustainability bonds) restrict where the raised money is spent. Sustainability-linked bonds and loans (SLBs/SLLs) place no restriction on use of proceeds — instead the financial terms are tied to the issuer's overall performance against predefined KPIs and sustainability performance targets (SPTs), such as cutting emissions by a set percentage by a target date. If the issuer misses the target, the coupon typically steps up (a penalty); they follow ICMA's Sustainability-Linked Bond Principles. The structure rewards entity-level transformation rather than project earmarking.

Q1049. Conceptually, how does carbon pricing work, and what are the two main mechanisms?

Carbon pricing puts a cost on greenhouse-gas emissions so emitters internalise the externality they impose on society. The two main mechanisms are a carbon tax (government sets a fixed price per tonne of CO₂-

equivalent and the market determines the quantity) and a cap-and-trade emissions trading system (the government caps total allowable emissions, issues allowances, and lets firms trade them so the market sets the price). The EU ETS is the largest cap-and-trade system; India has notified a Carbon Credit Trading Scheme (CCTS) to develop a domestic compliance market.

Q1050. What is the difference between compliance carbon markets and voluntary carbon markets?

Compliance markets are created by law or regulation — emitters subject to a cap (like the EU ETS) must surrender allowances for their emissions, and trading is mandatory for covered entities. Voluntary carbon markets let companies and individuals buy carbon credits (each representing one tonne of CO₂ reduced or removed, e.g. from reforestation) on their own initiative to offset emissions toward voluntary net-zero claims. Voluntary credits face persistent concerns over additionality, permanence, double-counting, and quality, which is why integrity bodies like ICVCM have emerged to set standards.

Q1051. Distinguish physical climate risk from transition climate risk in finance.

Physical risk is the financial damage from the physical effects of climate change: acute events (floods, cyclones, wildfires) and chronic shifts (rising sea levels, heat stress, changing rainfall) that destroy assets, disrupt supply chains, and raise insurance losses. Transition risk is the financial risk from the move to a low-carbon economy — policy changes, carbon pricing, new technology, and shifting consumer preferences that can erode the value of carbon-intensive businesses. A coastal factory faces physical risk; a coal-fired power producer faces transition risk, and many firms face both.

Q1052. What are stranded assets?

Stranded assets are assets that suffer unanticipated or premature write-downs, devaluation, or conversion to liabilities because of changes such as climate policy, carbon pricing, technological disruption, or shifting demand. The classic example is fossil-fuel reserves or infrastructure (coal mines, oil fields, coal plants) that may become uneconomic to extract or operate under aggressive decarbonisation, leaving "unburnable carbon" on balance sheets. The concept matters for valuation because reserves booked today may never be monetised, implying impairment risk for energy and resource companies.

Q1053. What were the TCFD recommendations, and how do they relate to the ISSB standards?

The Task Force on Climate-related Financial Disclosures (TCFD), set up by the Financial Stability Board, recommended that companies disclose climate-related risks and opportunities across four pillars: governance, strategy (including scenario analysis), risk management, and metrics and targets (including Scope 1, 2, and where relevant Scope 3 emissions). The TCFD framework has now been consolidated into the IFRS Foundation's ISSB standards — IFRS S1 (general sustainability disclosures) and IFRS S2 (climate), which fully incorporate the TCFD recommendations and are becoming the global baseline; the TCFD itself was disbanded in 2023 after its work was handed to the ISSB.

Q1054. What is SEBI's BRSR, and who must comply with it in India?

BRSR (Business Responsibility and Sustainability Reporting) is SEBI's mandatory ESG disclosure framework for listed companies in India, replacing the earlier Business Responsibility Report. It requires reporting against the nine principles of the National Guidelines on Responsible Business Conduct (NGRBC), covering environmental, social, and governance metrics. It is mandatory for the top 1,000 listed companies by market capitalisation, and a subset of quantitative, assured "BRSR Core" KPIs is being phased in with reasonable assurance for the largest companies, alongside value-chain disclosures on a comply-or-explain basis.

Q1055. Does ESG investing help or hurt returns — what does the evidence say?

The evidence is genuinely mixed and depends on approach and period. Proponents argue ESG integration manages downside risk and that good governance correlates with quality; some meta-studies find a neutral-

to-slightly-positive link between ESG and financial performance. Critics note that exclusionary screening shrinks the universe and can reduce diversification, that much outperformance reflects sector/factor tilts (e.g. underweight energy, overweight tech), and that returns are period-dependent. The honest answer is that ESG is best framed as a risk lens rather than a guaranteed alpha source, and results vary by implementation.

Q1056. How does fiduciary duty interact with ESG considerations?

A fiduciary (e.g. a pension trustee or asset manager) must act in the best financial interest of beneficiaries. The modern interpretation — supported by reports like the UN/UNEP "Fiduciary Duty in the 21st Century" — is that considering financially material ESG factors is consistent with, and may even be required by, fiduciary duty, since ignoring material risks would be imprudent. The tension arises when ESG is used to pursue non-financial goals at the expense of returns; in the US this has become politicised, with debate over whether ERISA fiduciaries may weigh ESG, generally permitted only where it is financially relevant or as a tie-breaker.

Q1057. Give a high-level overview of the main verticals within fintech.

Fintech spans: digital payments (UPI, wallets, cards, real-time rails); lending and credit, including digital lending and BNPL (buy now, pay later); neobanks (digital-only banks, often partnering with a licensed bank); wealthtech and robo-advisors (automated, algorithm-driven portfolio management); insurtech (digital insurance distribution and underwriting); regtech (technology for compliance, KYC/AML, reporting); and the infrastructure layer — payment gateways, APIs, account aggregators, and core-banking platforms. The common thread is using technology to make financial services cheaper, faster, and more accessible.

Q1058. What is the "India Stack" and UPI, and why are they significant?

India Stack is a set of open APIs and digital public infrastructure built around identity (Aadhaar), payments, and data, enabling presence-less, paperless, cashless services. UPI (Unified Payments Interface), run by NPCI, is its flagship real-time payment system that lets users send money instantly between bank accounts using a virtual ID (VPA) across apps, with no per-transaction fee for users (P2P). UPI is significant for its massive scale — billions of transactions per month — interoperability, and role in financial inclusion, and it has become a model studied by other countries.

Q1059. What are robo-advisors, and what are their main limitations?

Robo-advisors are digital platforms that build and manage investment portfolios algorithmically based on a questionnaire about goals, risk tolerance, and horizon, typically using low-cost ETFs and automatic rebalancing at a fraction of a human advisor's fee. Their strengths are low cost, accessibility, discipline, and tax features like loss harvesting (in markets where it applies). Limitations include limited handling of complex or unusual situations, reliance on the quality of risk-profiling questions, model and herding risk in stressed markets, and weaker behavioural coaching during panics, when human advice often adds the most value.

Q1060. At a conceptual level, what are blockchain and DLT, and what finance use cases are realistic?

A blockchain is a type of distributed ledger (DLT) — a shared, append-only database replicated across many participants, where transactions are grouped into cryptographically linked blocks and validated by consensus rather than a single central authority, making records tamper-evident. Realistic finance use cases include securities settlement (potentially shortening or atomising settlement via delivery-versus-payment), tokenisation of assets, trade finance, and smart contracts — self-executing code that automatically enforces agreed terms. The benefits are reduced reconciliation, fewer intermediaries, and transparency; the trade-offs are scalability, governance, legal enforceability, and integration with existing systems.

Q1061. How should one think about cryptocurrencies, stablecoins, and CBDCs as a balanced finance topic?

Cryptocurrencies like Bitcoin are decentralised digital assets not backed by any issuer; they are highly volatile, used for speculation and some payments, and raise concerns over consumer protection, illicit use, and energy (for proof-of-work). Stablecoins aim to hold a steady value by pegging to a fiat currency or basket, usually backed by reserves — useful for crypto trading and cross-border transfers, but exposing holders to reserve-quality and de-pegging (run) risk, prompting regulation like the EU's MiCA. CBDCs (central bank digital currencies) are sovereign digital money issued by a central bank — a direct liability of the state — offering stability and policy control; India's RBI is piloting the Digital Rupee (e-rupee) in wholesale and retail forms.

Q1062. What is open banking, and how do Account Aggregators implement the idea in India?

Open banking is a framework where customers can securely share their financial data with third-party providers, via standardised APIs and explicit consent, to access better products like tailored loans or aggregated dashboards. In India this is implemented through the RBI-regulated Account Aggregator (AA) framework: licensed AAs are consent-managers (not data holders) that move data between Financial Information Providers (e.g. banks) and Financial Information Users (e.g. lenders) only with the customer's granular, revocable consent, built on a standardised consent artefact. The aim is to reduce friction in lending and advisory while keeping the customer in control of their data.

Q1063. Where is AI/ML used in finance, and what is model risk?

AI/ML is used in credit scoring (assessing default probability from richer/alternative data), algorithmic and high-frequency trading, fraud and anomaly detection, AML transaction monitoring, robo-advisory, document processing, and customer chatbots. Model risk is the risk of loss from decisions based on models that are wrong or misused — from flawed assumptions, biased or non-representative training data, overfitting, or deployment outside the conditions the model was built for. Sound practice (e.g. the US Fed/OCC SR 11-7 guidance) requires model validation, governance, explainability, monitoring for drift, and controls against discriminatory bias — especially important as opaque ML models can be hard to interpret and can entrench unfair outcomes.

Q1064. What is financial inclusion, and how have fintech and policy advanced it in India?

Financial inclusion means ensuring affordable, timely access to basic financial services — a transaction account, savings, credit, insurance, and payments — for all, especially underserved and low-income populations. In India the JAM trinity (Jan Dhan bank accounts, Aadhaar identity, Mobile) plus UPI enabled hundreds of millions of new accounts and direct benefit transfers that cut leakage, while account aggregators and alternative-data lending expand credit to thin-file borrowers. Inclusion supports economic development and is measured by indicators like the World Bank Findex and the RBI's Financial Inclusion Index, though challenges remain around account dormancy, digital literacy, and consumer protection against fraud and over-indebtedness.