

STUDY GUIDE

# Market Research — Complete Handbook

A complete, concept-first reference — prepared for Saikumar Kaleru.

certifications / study / Finance • 2026

*Prepared for Saikumar Kaleru. Read top to bottom; each part builds on the last. Bold terms are the vocabulary interviewers listen for. Indian-market context (CMIE, RBI, SEBI, NSS, industry bodies) is used throughout, matching the rest of this study library.*

---

# PART 1 — WHAT MARKET RESEARCH ACTUALLY IS

## 1.1 The core idea

**Market research** is the systematic collection, analysis, and interpretation of data about a market — its customers, competitors, and environment — to reduce uncertainty before a business decision. Where **equity research** answers "is this stock worth buying," market research answers "will this product sell, to whom, at what price, and against whom."

Every market research project ultimately answers one of four questions: 1. **Who** is the customer, and what do they need? (segmentation, needs) 2. **How big** is the opportunity? (market sizing) 3. **Who else** is competing for it? (competitive intelligence) 4. **What should we do** about price, product, or positioning? (decision support)

## 1.2 Market Research vs adjacent disciplines

- **Equity/Investment Research** studies a *company* to price its stock — outward-facing to public markets, ends in a buy/sell call.
- **Market Research** studies a *market or customer base* to guide product/business decisions — ends in a recommendation to a brand, product, or strategy team, not a trade.
- **Competitive Intelligence (CI)** is a specialised slice of market research focused specifically on competitors' moves, capabilities, and likely responses.
- **Business/Data Analytics** answers "what happened" from a company's own internal data; market research goes *outside* the company to answer "what's happening in the market."
- In practice these overlap heavily — a strong market research analyst reads company filings and does light equity-research-style work, and a strong equity analyst does light market-sizing work. The skill sets are transferable, which is the bridge worth using if coming from a finance/markets background.

## 1.3 Where market researchers sit

- **Agency side:** research/insights consultancies (Nielsen, Kantar, IMRB/Kantar IMRB, Ipsos, GfK, YouGov, Mintel) — client-facing, run studies for many brands.
- **Client/in-house side:** a consumer or B2B company's own insights/strategy team (FMCG, banking, telecom, e-commerce) — owns a category and runs research for internal stakeholders.
- **Consulting:** strategy/management consultancies run market-sizing and competitive-landscape workstreams as part of larger engagements.
- **Financial services angle** (the natural bridge from an equity/derivatives background): buy-side and consulting shops increasingly employ "market intelligence" analysts who combine public-market data with primary research to validate investment theses — this is the most direct crossover role.

## 1.4 The research process, end to end

1. **Define the business problem** — not the research question yet. ("Should we launch in Tier-2 cities?" is a business problem; "what is unaided brand awareness in Tier-2 cities?" is a research question.)
2. **Translate into research objectives** — specific, answerable questions.
3. **Choose a design** — exploratory, descriptive, or causal (Part 2).
4. **Choose methods** — secondary first (cheap, fast), then primary if gaps remain (Parts 3-4).
5. **Collect data** — fieldwork, surveys, interviews, desk research.

6. **Analyse** — statistical and qualitative analysis (Part 5).

7. **Report and recommend** — translate findings into a decision-ready recommendation (Part 8).

**Trap:** a research design that answers the *literal* question but not the *business* problem is a common interview failure point — always restate the business decision the research will inform before designing the study.

## 1.5 The research brief — the document that starts every project

A professional engagement begins with a **research brief**, written by the client/stakeholder (or, often in practice, drafted by the researcher *for* the stakeholder to sign off, since stakeholders rarely write a good brief unaided). A complete brief covers: - **Background** — why this project, why now (a competitor launch, a declining metric, a new market entry decision). - **Business objective** — the decision this research will inform, stated as a decision, not a topic ("decide whether to reformulate the product" not "understand consumer taste preferences"). - **Research objectives** — the specific, answerable questions that, once answered, let the business objective be decided. - **Target audience/population** — who the research needs to represent. - **Deliverables and timeline** — report format, presentation date, any interim readouts. - **Budget and constraints** — this shapes method choice more than researchers like to admit; a ₹5 lakh budget cannot fund a 2,000-person national survey with 20 city-tier cells. - **Stakeholder map** — who signs off, who consumes the findings, who might resist the conclusion (a market research analyst who ignores internal politics produces technically correct research that changes nothing).

**Common failure mode:** skipping straight to "let's run a survey" without a brief. A senior analyst's first move on any incoming request is to write or extract a one-page brief and get explicit sign-off on the business objective — this alone prevents the single most expensive mistake in the field: a technically well-executed study that answers a question nobody needed answered.

## 1.6 Research objectives vs research questions vs survey questions — the translation chain

Interviewers probe this distinction because junior analysts conflate the three: - **Research objective** (business-facing): "Understand whether price or convenience is the bigger barrier to online grocery adoption among non-users." - **Research question** (analysis-facing): "What proportion of non-users cite price vs convenience as the primary barrier, and does this differ by income band?" - **Survey question(s)** (respondent-facing): "Which of the following best describes why you don't currently order groceries online? [price too high / delivery too slow / prefer to see product before buying / no reliable internet / other]" plus a follow-up ranking or Likert battery on each barrier. A good analyst can walk backward from any survey question to the research question it serves to the business objective it ultimately informs — if a question can't be traced back, cut it.

---

## PART 2 — RESEARCH DESIGN

### 2.1 The three designs

- **Exploratory research** — used when the problem is vague. Goal: generate hypotheses, not test them. Methods: literature review, expert interviews, focus groups, case analysis. Small samples, qualitative, flexible.
- **Descriptive research** — used to describe a market "as it is": size, share, demographics, usage patterns, satisfaction levels. Methods: surveys, observational studies, secondary-data analysis. Larger samples, mostly quantitative, structured.
- **Causal research** — used to test whether X causes Y (e.g., does a price cut increase volume enough to grow revenue). Methods: experiments, A/B tests, controlled field trials. Requires a control group and random assignment to claim causality — correlation from a plain survey is *not* proof of causation, a point interviewers probe directly.

### 2.2 Choosing a design — the decision rule

If you don't know what you don't know → **exploratory**. If you need to size or describe something known → **descriptive**. If you need to prove a specific action will change an outcome → **causal**. Most real projects are a funnel: exploratory (a few interviews) → descriptive (a large survey) → occasionally causal (a pilot/A-B test) before a big launch decision.

### 2.3 Primary vs secondary research

- **Secondary research** uses data that already exists (someone else collected it for another purpose). Cheaper, faster, but not tailored to your exact question and may be stale.
- **Primary research** is data collected specifically for this project (surveys, interviews, experiments). Expensive and slow, but exactly fits the question and is proprietary.
- **Rule of thumb**: always exhaust secondary sources first — most "we need a survey" requests are 60-80% answerable from existing reports, company filings, and government data. This is the single biggest cost-saving move a junior researcher can make, and interviewers listen for it.

### 2.4 Qualitative vs quantitative

- **Qualitative** — depth over breadth. Answers *why* and *how*. Small samples (6-30), non-statistical, rich narrative data (interview transcripts, discussion notes).
- **Quantitative** — breadth over depth. Answers *how many*, *how much*, *what %*. Large samples (100s-1000s), statistical, numeric data (survey responses, scanner data).
- **Mixed-methods** is standard practice: qualitative first to discover the right questions and language, quantitative second to measure how widespread the finding is across the population.

### 2.5 Research design worked through six Indian industries

Interviewers frequently present a one-line business problem and ask "how would you design this?" — the fastest way to build fluency is to have already worked through several. Each example below follows the same skeleton: business problem → design choice → method mix → key risk to flag.

**1. FMCG — should a snack brand reformulate to reduce sugar content?** Business problem: will a reformulated (lower-sugar) SKU protect volume, or will core users reject the taste change? Design: causal,

via a **blind taste test** (respondents don't know which sample is reformulated) with a monadic or paired-comparison design, plus a quantitative purchase-intent survey on the winning formulation. Method mix: central-location taste test (n≈200-300, quota-controlled by current-user vs non-user), followed by a home-use test (2-week in-home trial, n≈100) to capture repeat-usage reality beyond a single sip. Key risk: a central-location "sip test" often overstates acceptance versus real repeated in-home consumption — always pair with a home-use test before a full reformulation, a classic FMCG research-design trap.

**2. Banking — will customers adopt a new UPI-based micro-lending feature?** Business problem: build or don't build; if build, what credit limit and pricing. Design: exploratory (IDIs with existing personal-loan customers and with non-borrowers) → descriptive quantitative (intent, preferred limit, acceptable interest rate via a **Gabor-Granger** or **Van Westendorp** pricing module) → light conjoint on feature bundle (limit size, interest rate, repayment tenure, processing speed). Method mix: online survey on app users (n≈800, stratified by existing credit-bureau band if available internally), 15-20 qualitative interviews for language/objections. Key risk: stated willingness-to-pay for credit products is notoriously unreliable (nobody wants to admit they'd pay high interest) — validate with a soft-launch/waitlist pilot measuring actual uptake at a real posted rate before committing to full-scale pricing.

**3. E-commerce — why is cart abandonment high on the checkout page specifically?** Business problem: diagnose and fix a conversion leak. Design: descriptive, mixed-methods — this is a case where **behavioural/analytics data should be pulled first** (internal secondary data: session recordings, funnel drop-off by step, device/payment-method cuts) before commissioning any primary research. Method mix: quant funnel analysis (internal data, free) → usability testing (5-8 moderated sessions, think-aloud, on the actual checkout flow) → optionally an intercept exit-survey ("what stopped you from completing your order today?") triggered on abandonment. Key risk: qualitative usability testing with 5-8 users finds ~80% of major usability issues (Nielsen's well-known finding) — resist the urge to commission an expensive large-n survey before this cheap, fast diagnostic step.

**4. Telecom — sizing demand for a new prepaid data-only plan for students** Business problem: is the segment large enough and price-sensitive enough to justify a dedicated plan and marketing spend. Design: descriptive, primarily secondary + light primary validation. Method mix: top-down sizing from telecom-industry reports (TRAJ data on subscriber base, student population from education-ministry data) cross-checked bottom-up via campus-intercept surveys at a sample of colleges (stratified by tier-1/tier-2 city) on current spend and switching intent. Key risk: campus intercepts oversample students physically present and willing to stop — supplement with an online panel arm to correct for this convenience-sample bias before finalising the size estimate.

**5. Healthcare/pharma — understanding doctor prescribing behaviour for a new drug category** Business problem: which messaging and which specialist segment to prioritise in field-force targeting. Design: qualitative-led, given a small, hard-to-access, expert population. Method mix: 25-30 IDIs with prescribing physicians (recruited via a medical panel/fieldwork agency, incentivised appropriately per medical-research ethics norms), segmented by specialty and prescribing volume tier; a short quantitative confirmatory survey (n≈150-200 physicians) on message resonance and stated prescribing intent. Key risk: physicians are a scarce, expensive-to-recruit population — a poorly piloted discussion guide that wastes even 5 of 30 interviews on bad questions is a real, expensive mistake; pilot rigorously with 1-2 physicians before full fielding.

**6. Auto/EV — will consumers switch from ICE two-wheelers to electric in a specific state?** Business problem: market-entry and product-spec decision (range, price, charging assumptions) for a new EV two-wheeler. Design: descriptive + causal element (conjoint) — this is a multi-attribute trade-off decision (price vs range vs charging time vs brand), the textbook conjoint-analysis use case. Method mix: quantitative survey (n≈600-1000, stratified by current vehicle type and state given large state-level policy/subsidy differences in India) with an embedded conjoint module on price/range/charging-time/brand; secondary

desk research on state EV-subsidy policy and existing charging-infrastructure density. Key risk: EV purchase intent is heavily distorted by subsidy policy that can change before launch — always model the sizing/pricing conclusion under at least two subsidy scenarios (current policy, and a "subsidy reduced/removed" downside case) rather than a single point estimate.

---

## PART 3 — SECONDARY RESEARCH (DESK RESEARCH)

---

### 3.1 Where to look, in order of cost

1. **Company's own data** — internal sales, CRM, support tickets, prior research archives (free, but easy to miss if not asked for).
2. **Government and regulatory sources** (India) — MOSPI/NSS surveys (consumer expenditure, employment), RBI (household finance, banking data), SEBI (mutual fund/market data), Census, Economic Survey, ministry reports (e.g. DPIIT for startups) — free and authoritative, but slow-moving and broad.
3. **Industry bodies and trade associations** — CII, FICCI, NASSCOM (IT), ASSOCHAM, and sector bodies (e.g. AMFI for mutual funds, IBA for banking) — free/cheap, sector-specific, often annual.
4. **Syndicated research firms** — Nielsen (retail audit, FMCG), Kantar, CMIE (Centre for Monitoring Indian Economy — the single most-used paid macro/company database in India), Euromonitor, Statista, Gartner/Forrester (tech), IDC. Paid, but standardised, comparable across companies, and update regularly.
5. **Public company filings** — annual reports, investor presentations, credit-rating rationale (CRISIL/ICRA/CARE), stock-exchange disclosures. Free, primary-sourced by the company itself, most useful for competitive benchmarking of listed players.
6. **News, trade press, and analyst notes** — free, current, but requires triangulation since a single source may be biased or wrong.

### 3.2 Desk research technique

- **Triangulate:** never trust one source for a market-size or growth number — cross-check at least two independent sources and reconcile the gap explicitly in the report (don't hide the disagreement).
  - **Check the definition, not just the number:** two "market size" figures can differ 5x because one includes adjacent categories (e.g. "EV market" including two-wheelers vs only cars) — always confirm scope before comparing.
  - **Track vintage:** a 2019 number quoted in 2026 without an update flag is a common analyst error and an easy interviewer trap ("how current is this number, and how do you know?").
  - **Build a source log** as you go — every number in the final deck should be traceable to a named, dated source; this is what separates a credible market research deliverable from a guess.
-



## PART 4 — PRIMARY RESEARCH: QUALITATIVE METHODS

---

### 4.1 In-depth interviews (IDIs)

One-on-one, semi-structured conversations (45-60 min), used to explore a topic in depth without group dynamics influencing the answer. Best for sensitive topics (income, health, B2B purchasing decisions) or expert/stakeholder input (e.g. interviewing distributors, doctors, CFOs).

### 4.2 Focus group discussions (FGDs)

6-10 participants, a trained moderator, 60-90 minutes, used to observe group dynamics, generate ideas through discussion, and test reactions to concepts/ads/packaging live. **Risk:** groupthink and dominant-participant bias — a skilled moderator actively manages this (round-robin questions, private written responses before group discussion).

### 4.3 Ethnography / observational research

Researchers observe behaviour in its natural setting (a kirana store, a bank branch, a household) rather than asking about it — used because **stated preference and revealed behaviour often diverge** (people say they'd pay for organic food, then buy the cheaper option at the shelf). Slower and more expensive, but the gold standard for behavioural truth.

### 4.4 Other qualitative tools

- **Projective techniques** — indirect prompts (word association, sentence completion, "if this brand were a person...") used to surface associations people won't state directly.
- **Diary studies** — participants log behaviour/thoughts over days/weeks, useful for infrequent or habitual behaviours a single interview can't capture.
- **Usability/concept testing** — a moderator watches a participant use a product/prototype and think aloud, common in fintech/app research.

### 4.5 Qualitative analysis

- **Coding:** tagging transcript segments with recurring themes/categories (manually or in NVivo/Atlas.ti).
  - **Thematic analysis:** grouping codes into higher-level themes and reporting prevalence ("6 of 12 respondents mentioned trust as the top barrier") — note this is *illustrative*, not statistically representative, and should never be reported as a percentage of "all customers."
-

## PART 5 — PRIMARY RESEARCH: QUANTITATIVE METHODS & SURVEY DESIGN

### 5.1 The survey design process

1. Define exactly what each question will be used to decide (no "nice to know" questions — every question earns its place).
2. Choose question types: **closed-ended** (multiple choice, Likert scale, ranking) for easy quantitative analysis; **open-ended** sparingly, for colour and to catch what closed options missed.
3. Order questions: general → specific, easy → sensitive, behavioural → attitudinal → demographic (demographics last, so people are already invested in finishing).
4. **Pilot test** on 10-20 people before full fielding — catches confusing wording, skip-logic errors, and length problems.

### 5.2 Common bias traps in question wording

- **Leading questions** — "Don't you agree our app is easy to use?" pushes agreement; ask neutrally: "How easy or difficult is our app to use?"
- **Double-barrelled questions** — "Is the app fast and reliable?" conflates two things; split into two questions.
- **Social desirability bias** — people over-report virtuous behaviour (saving, exercising) and under-report the opposite; mitigate with anonymity and neutral framing.
- **Order/anchoring effects** — the first option in a list gets chosen more often; randomise option order where the survey tool allows it.

### 5.3 Measurement scales

- **Nominal** — categories with no order (city, gender, brand used).
- **Ordinal** — ordered categories, unequal spacing (income brackets, satisfaction: poor/fair/good/excellent).
- **Interval** — ordered, equal spacing, no true zero (a 1-5 Likert scale is technically interval by convention).
- **Ratio** — ordered, equal spacing, true zero (age, spend in ₹, number of purchases) — allows the widest range of statistics (means, ratios).
- **Likert scale**: the standard 5- or 7-point agreement/satisfaction scale. Interviewers may ask why 5 vs 7 points — 7-point scales give slightly more discrimination but higher respondent fatigue; 5-point is the pragmatic default for most consumer surveys.
- **Net Promoter Score (NPS)**: "How likely are you to recommend X, 0-10?" — % Promoters (9-10) minus % Detractors (0-6), ignoring Passives (7-8). Widely used, widely criticised (a single number hides *why*), always pair with an open-ended follow-up ("what's the main reason for your score?").

### 5.4 Sampling

- **Population** — everyone the finding should generalise to (e.g. "urban Indian smartphone owners aged 18-35").
- **Sampling frame** — the actual list you can draw from (a customer database, a panel) — rarely a perfect match to the population, and the gap is a source of bias worth naming in any report.

- **Probability sampling** (every unit has a known chance of selection — allows statistically valid generalisation):
- **Simple random sampling** — pure lottery from the full frame.
- **Stratified sampling** — split the population into strata (e.g. by city tier), sample each proportionally — improves precision when subgroups matter.
- **Cluster sampling** — randomly pick clusters (e.g. a few cities), then sample fully within them — cheaper for geographically spread populations, common in India given field-cost constraints.
- **Systematic sampling** — pick every k-th unit from an ordered list.
- **Non-probability sampling** (convenience, quota, snowball, judgment) — faster/cheaper but cannot be generalised with a formal margin of error; acceptable for exploratory work, a red flag if used for a headline market-size claim.

## 5.5 Sample size — the formula interviewers actually ask for

For a proportion, at 95% confidence:

$$n = (Z^2 \times p \times (1-p)) / e^2$$

Where  $Z = 1.96$  (for 95% confidence),  $p$  = expected proportion (use 0.5 if unknown — it maximises required sample, the conservative choice),  $e$  = desired margin of error (e.g. 0.05 for  $\pm 5\%$ ).

**Worked example (India context):** A bank wants to estimate the % of urban customers who'd adopt a new UPI feature,  $\pm 5\%$  margin of error, 95% confidence,  $p$  unknown (use 0.5):  $n = (1.96^2 \times 0.5 \times 0.5) / 0.05^2 = (3.8416 \times 0.25) / 0.0025 = 0.9604 / 0.0025 \approx \mathbf{384 \text{ respondents}}$ . This is the famous "384" number that recurs across market research regardless of population size (above a few thousand) — population size barely matters once it's large, which surprises people new to the field and is a good interview talking point.

For a **finite population correction** (when the population itself is small, e.g. surveying a company's 2,000 enterprise clients):  $n_{\text{adj}} = n / (1 + (n-1)/N)$ , where  $N$  is the population size.

## 5.6 Data collection modes

- **CAPI/CATI** (computer-assisted personal/telephonic interviewing) — interviewer-led, good response rates, higher cost, still common for rural/older India respondents.
- **Online panels (CAWI)** — cheapest and fastest, but skews toward internet-literate, younger, urban respondents — a real limitation to disclose, not hide, in any Indian consumer study.
- **Intercept surveys** — stopping people in a physical location (mall, branch, store) — good for in-the-moment behaviour, biased toward that location's footfall profile.

## 5.7 A full worked questionnaire (UPI micro-lending feature, from the Part 2.5 example)

Reproducing a complete, well-formed questionnaire is one of the most useful things to have memorised for a portfolio/interview exercise — note the ordering (screening → behavioural → attitudinal → conjoint-style trade-off → demographics last), the neutral wording, and the routing logic.

**Screenener** S1. Do you currently use a UPI app for payments? [Yes → continue / No → terminate] S2. Have you ever taken a personal loan or used a "buy now pay later" product? [Yes/No — used for quota, not a terminate]

**Behavioural** Q1. In the last 3 months, has there been an occasion where you needed short-term credit (₹2,000-₹50,000) but did not borrow? [Yes/No] Q2. If yes, what did you do instead? [Borrowed from

friends/family / Used a credit card / Delayed the purchase/expense / Used a different lending app / Other — specify]

**Attitudinal / intent** Q3. How interested would you be in a feature that lets you borrow instantly through your UPI app, with the amount auto-debited from your next salary credit? [5-point: Not at all interested → Extremely interested] Q4. What would concern you most about using such a feature? [Open-ended, single response] Q5. Rate your agreement: "I trust my UPI app provider to offer fair interest rates." [5-point Likert: Strongly disagree → Strongly agree]

**Pricing (Van Westendorp Price Sensitivity Meter — four questions, same feature, asked in this order)** Q6a. At what interest rate (% per annum) would you consider this credit line too *expensive* to use? [numeric entry] Q6b. At what rate would you consider it *starting to get expensive*, but you'd still consider it? [numeric entry] Q6c. At what rate would you consider it a *bargain* — great value? [numeric entry] Q6d. At what rate would you consider it *so cheap* that you'd question the lender's legitimacy/safety? [numeric entry]

**Conjoint module** (choice-based, 6 tasks shown, each a choice among 3 hypothetical plans + "none") Example task: "Which of these would you choose?" - Plan A: ₹25,000 limit · 18% p.a. · 3-month repayment · instant approval - Plan B: ₹50,000 limit · 24% p.a. · 6-month repayment · approval in 1 hour - Plan C: ₹15,000 limit · 15% p.a. · 1-month repayment · instant approval - None of these

**Demographics** (always last) D1. Age band · D2. Gender · D3. City tier · D4. Monthly income band · D5. Occupation

**Notes an analyst should be able to explain about this instrument:** why the screener terminates non-UPI-users (population definition), why Van Westendorp uses four specific price points rather than one willingness-to-pay question (it triangulates an acceptable *range*, not a single number, and cross-checks internal consistency — a respondent whose "too cheap" answer exceeds their "too expensive" answer flags a bad/inattentive response), and why demographics sit last (reduces early-abandon risk on a survey that opens with sensitive financial behaviour).

## 5.8 Panel quality and data-cleaning for online surveys

Online panels (Section 5.6) introduce a data-quality problem beyond sampling bias: **inattentive or fraudulent respondents** rushing through for a panel incentive. Standard quality checks, expected knowledge for any quantitative role: - **Speeding**: flag/remove respondents completing in under roughly a third of the median completion time. - **Straight-lining**: selecting the same response option down an entire grid/matrix question regardless of content — a strong signal of non-engagement. - **Attention-check items**: an instruction embedded mid-survey ("select 'Somewhat agree' for this item to show you are reading carefully") — failure is grounds for exclusion. - **Open-end gibberish**: nonsensical or copy-pasted text in open-ended fields. - **Logical consistency traps**: e.g. claiming to be a heavy user of a product category while also selecting "never heard of this brand" for a brand used earlier in the survey. Reporting the **effective sample size after cleaning** (not just the raw completes) alongside the cleaning criteria applied is standard practice in a credible methodology appendix.

## PART 6 — DATA ANALYSIS FOR MARKET RESEARCH

### 6.1 Descriptive statistics (the first pass on any dataset)

Mean, median, mode, standard deviation, frequency distributions. Always report **median alongside mean** for skewed data (income, spend) — a few high spenders can make the mean misleading.

### 6.2 Cross-tabulation ("cross-tabs")

Breaking one variable by another (e.g. NPS score by age band) — the single most common output in a market research deck. Always check **base size** (n) in each cell; a finding on a base of 12 respondents is not reportable as a trend, only a footnote.

### 6.3 Segmentation

Grouping the market into distinct customer types who share needs/behaviour, so different strategies can be aimed at each. - **A-priori segmentation** — segments defined in advance by known variables (age, income, geography) — simple, easy to action, sometimes not the variable that actually drives behaviour. - **Post-hoc (cluster-based) segmentation** — statistical clustering (k-means, hierarchical clustering) on attitudinal/behavioural survey data to *discover* natural groupings — richer, but harder to explain to stakeholders and harder to identify a real customer against ("which cluster is this person?" requires a classifier built on top). - A good segmentation is: **substantial** (big enough to matter), **distinct** (segments actually differ), **stable** (doesn't shift month to month), **actionable** (marketing/product can actually reach and target it).

### 6.4 Statistical significance testing

- **t-test** — compares means between two groups (e.g. does Segment A spend more than Segment B?).
- **Chi-square test** — tests whether two categorical variables are associated (e.g. is brand preference associated with city tier?).
- **ANOVA** — compares means across 3+ groups at once.
- **p-value**: the probability of seeing a difference this large (or larger) if there were truly no difference in the population.  $p < 0.05$  is the conventional "statistically significant" threshold — but always pair with **effect size** (how big is the difference, practically) since a huge sample can make a trivial difference "significant."

### 6.5 Regression and conjoint analysis

- **Linear/logistic regression** — quantifies how much a set of variables (price, awareness, income) predicts an outcome (purchase, satisfaction) — used for driver analysis ("what actually moves NPS").
- **Conjoint analysis** — shows respondents product profiles that vary systematically on attributes (price, brand, features) and asks them to choose/rank; from choices, the model estimates how much each attribute (and each level, e.g. ₹999 vs ₹1499) contributes to preference. The standard technique for **pricing research** and feature prioritisation — a common interview topic for anyone claiming pricing-research experience.
- **MaxDiff (best-worst scaling)** — respondents repeatedly pick the "most" and "least" important item from small sets — a cleaner alternative to asking people to rate 20 things on a 1-5 scale (where everything ends up rated 4-5).

## 6.6 Tools of the trade


- **SPSS** — the long-standing standard for cross-tabs, significance testing, and basic multivariate stats in market research; still widely required in job postings.
- **Qualtrics / SurveyMonkey** — survey design, fielding, and light built-in analysis.
- **Excel** — pivot tables remain the day-to-day workhorse for cross-tabs and quick analysis.
- **Power BI / Tableau** — dashboarding for tracking studies (brand trackers, NPS trackers) that update on a cadence.
- **Python (pandas, statsmodels, scikit-learn)** — increasingly expected for segmentation, regression, and automating repetitive survey-data cleaning — a genuine differentiator for a candidate coming from a Python/SQL background, since most classically-trained market researchers don't code.
- **R** — still common in academic-leaning or advanced analytics teams for the same statistical work as Python.

## 6.7 Regression driver analysis — full worked example

A bank wants to know what drives NPS for its mobile app. Survey data on n=500 users includes NPS (0-10), and five predictor variables measured 1-5: app speed, ease of navigation, feature completeness, customer-support responsiveness, and trust in security.

A linear regression of NPS on the five predictors yields standardised coefficients (betas — comparable to each other since all variables are on the same 1-5 scale):

| DRIVER                          | STANDARDISED BETA | P-VALUE |
|---------------------------------|-------------------|---------|
| Trust in security               | 0.34              | <0.001  |
| Ease of navigation              | 0.28              | <0.001  |
| App speed                       | 0.19              | 0.002   |
| Customer-support responsiveness | 0.11              | 0.04    |
| Feature completeness            | 0.04              | 0.31    |



**Reading this table (the part interviewers actually probe):** trust in security is the single largest driver of NPS — larger than app speed or even feature completeness, which is *not* statistically significant ( $p=0.31$ , meaning we cannot conclude it drives NPS at all in this sample). The actionable insight is not "improve everything" but "prioritise security-trust communication and navigation simplicity over adding new features" — directly answering a resource-allocation question a findings-only report would not answer. Always check **R<sup>2</sup> (variance explained)** too — a driver model with  $R^2=0.15$  means the five variables jointly explain only 15% of NPS variation, a caveat that belongs in the same slide as the driver chart, not omitted.

## 6.8 Conjoint analysis — from raw choices to attribute importance


Continuing the Part 5.7 UPI lending conjoint: after  $n=600$  respondents each complete 6 choice tasks (3,600 total choice observations), a multinomial logit model estimates a **part-worth utility** for each attribute level. Illustrative output:

| ATTRIBUTE      | LEVEL    | PART-WORTH UTILITY |
|----------------|----------|--------------------|
| Interest rate  | 15% p.a. | +0.82              |
| Interest rate  | 18% p.a. | +0.21              |
| Interest rate  | 24% p.a. | -1.03              |
| Credit limit   | ₹15,000  | -0.44              |
| Credit limit   | ₹25,000  | +0.08              |
| Credit limit   | ₹50,000  | +0.36              |
| Approval speed | Instant  | +0.61              |
| Approval speed | 1 hour   | -0.61              |

**Attribute importance** is calculated from each attribute's *range* of utilities (max minus min), normalised to sum to 100% across attributes: interest rate's range (0.82 to -1.03 = 1.85) is the largest, making it the most important attribute ( $\approx 45\%$  relative importance in this illustrative case), ahead of approval speed (range



1.22,  $\approx 30\%$ ) and credit limit (range 0.80,  $\approx 20\%$  — with feature/brand making up the remainder). This is exactly how a pricing-research deliverable justifies a statement like "interest rate matters roughly 1.5x more to adoption than approval speed" to a pricing committee, and it is the calculation most likely to be asked to explain, step by step, in a technical interview for a role advertising conjoint/pricing-research experience.



**Market simulation:** once part-worths are estimated, they can simulate **share of preference** for any hypothetical product lineup (e.g. "if we launch Plan B against two named competitor plans, what share would we capture") by summing utilities per profile and converting to choice probabilities — this simulator output, not the raw utility table, is what gets presented to a business stakeholder.


## 6.10 A full worked segmentation case study, with cluster output

*A digital lending app runs a survey ( $n=800$ ) with an attitudinal/behavioural battery, then applies  $k$ -means clustering ( $k=4$ , chosen via the elbow method on within-cluster variance) on standardised variables: price sensitivity (1-5), digital comfort (1-5), loan-purpose (categorical, recoded), and monthly income (₹).*

**Illustrative cluster output:**



| SEGMENT |                            | SIZE (%) | PRICE SENSITIVITY | DIGITAL COMFORT | AVG INCOME | PRIMARY LOAN PURPOSE        |
|---------|----------------------------|----------|-------------------|-----------------|------------|-----------------------------|
| A:      | "Convenience seekers"      | 28%      | 2.1 (low)         | 4.6 (high)      | ₹95,000    | Discretionary/lifestyle     |
| B:      | "Value-driven strivers"    | 34%      | 4.4 (high)        | 3.2 (mid)       | ₹42,000    | Emergency/essential         |
| C:      | "Cautious traditionalists" | 22%      | 3.6 (mid-high)    | 1.8 (low)       | ₹58,000    | Rare, only when unavoidable |
| D:      | "Affluent optimisers"      | 16%      | 1.5 (low)         | 4.8 (high)      | ₹1,80,000  | Investment/opportunity      |



**Reading the output (the part that gets tested in interviews):** the four clusters are checked against the segmentation-quality criteria from Part 6.3 — **substantial** (all four segments are large enough to target, none is a tiny statistical artifact), **distinct** (the profiles genuinely differ across all four variables, not just one), **actionable** (Segment B, the largest at 34% and clearly price-sensitive with essential loan needs, is an obvious first product-and-messaging target; Segment D, though small, has high income and low price sensitivity — a candidate for a premium, fast-approval product tier rather than a discount play). The critical next step, per Part 6.3's caution about post-hoc clustering, is validating that these clusters actually predict a business-relevant outcome (e.g. repayment behaviour, cross-sell uptake) — a statistically clean cluster solution that doesn't differentiate on outcomes the business cares about is elegant but not useful.

**A common interview follow-up:** "how did you choose k=4 and not k=3 or k=5?" — the honest answer combines a quantitative heuristic (the elbow method: plotting within-cluster variance against k and picking the point where additional clusters stop meaningfully reducing variance) with a business-judgment check (does the resulting number of segments produce genuinely actionable, differentiable strategies, or does splitting further just fragment the target list without a corresponding difference in what the business would actually do for each group) — k-selection is not a purely mechanical decision.

## 6.9 Common statistical pitfalls specific to market research (interview favourites)

- **Confusing correlation in survey data with causation** — a survey showing heavy app users have higher NPS does not prove app usage *causes* satisfaction; satisfied customers may simply use the app more (reverse causality) — only a designed experiment (Part 2.1, causal research) can claim causation.
  - **Small-base reporting** — presenting a sub-group finding (e.g. "Gen Z respondents in Tier-3 cities") without checking the cell has a sane base size ( $n \geq 30$  as a rough floor for any percentage to be quoted,  $n \geq 100$  to report with real confidence) is one of the most common junior-analyst errors flagged by senior reviewers.
  - **Multiple comparisons** — running significance tests on dozens of cross-tab cells and only reporting the "significant" ones inflates the false-positive rate (with a 5% threshold, testing 20 unrelated things yields roughly one false "significant" result by chance alone) — pre-register the handful of hypotheses that matter rather than data-mining a whole cross-tab table for anything that clears  $p < 0.05$ .
  - **Treating a 5-point Likert mean as precise** — reporting "3.42 vs 3.38" as a meaningful difference on ordinal-by-convention data without a significance test attached is a common but unjustified precision claim.
-

## PART 7 — MARKET SIZING & COMPETITIVE INTELLIGENCE

### 7.1 TAM, SAM, SOM

- **TAM (Total Addressable Market)** — total revenue opportunity if you captured 100% of the relevant market, globally or in the defined geography.
- **SAM (Serviceable Addressable Market)** — the slice of TAM you could realistically serve given your business model/geography/channel constraints.
- **SOM (Serviceable Obtainable Market)** — the slice of SAM you can realistically capture given competition and go-to-market limits, usually over a defined time horizon (e.g. 3 years).

#### TAM → SAM → SOM: narrowing to the realistically obtainable market



### 7.2 Two ways to size a market

- **Top-down:** start from a broad, authoritative number (e.g. total Indian retail spend from MOSPI/Census) and apply a chain of narrowing percentages (% online, % in category, % in target city tier) down to the target market. Fast, defensible on data sources, but errors compound multiplicatively across each percentage applied.
- **Bottom-up:** build from unit economics (number of potential customers × price × purchase frequency, or number of outlets × average sales per outlet) — more defensible and specific, but requires more granular data and assumptions to justify.
- **Best practice:** do both and **triangulate** — if top-down and bottom-up land within roughly 20-30% of each other, confidence is high; if they diverge sharply, the discrepancy itself is the most important finding to investigate before presenting either number.

**Worked example (India, illustrative):** Sizing the market for a premium D2C protein-supplement brand targeting urban gym-goers. - *Top-down:* India population ~1.4bn → urban ~35% (~490m) → adults

18-40 ~30% of urban (~147m) → gym-goers/fitness-conscious ~8% (~11.8m) → willing to pay premium price ~15% (~1.8m potential customers) × assumed ₹6,000/year average spend ≈ **₹1,060 crore SAM**. - *Bottom-up*: ~35,000 organised gyms in India × average 150 active members × 20% buy protein supplements × ₹6,000/year ≈ **₹630 crore**. - The gap (₹1,060cr vs ₹630cr) would be flagged and investigated — likely the top-down "fitness-conscious %" assumption is too generous relative to the bottom-up gym-membership-based count, a classic and defensible finding to walk an interviewer through.

### 7.3 Competitive intelligence (CI)

- **Competitor mapping**: identify direct competitors (same product, same customer), indirect competitors (different product, same need), and potential entrants.
- **Porter's Five Forces** (also core to equity/strategy work — the crossover point with fundamental research): threat of new entrants, bargaining power of buyers, bargaining power of suppliers, threat of substitutes, competitive rivalry — used to assess how attractive/defensible a market structure is.
- **Win-loss analysis**: systematically interviewing recently won and lost deals/customers to understand *why* — one of the highest-value, most underused CI techniques in B2B contexts.
- **Benchmarking**: comparing pricing, features, service levels, and public financials (for listed competitors) side by side — this is where equity-research-style spreadsheet skills transfer directly into a market research/CI role.
- **Sources for CI in India**: competitor annual reports and investor decks, MCA filings (Ministry of Corporate Affairs — company financials for private Indian companies), app-store reviews and ratings, job postings (signal hiring direction/expansion plans), patent filings, News/PR, LinkedIn headcount trends by function.

### 7.4 Industry/sector analysis frameworks

- **PESTEL** — Political, Economic, Social, Technological, Environmental, Legal factors shaping an industry's outlook, used for the macro layer of an industry report.
- **Value chain analysis** — mapping who captures value at each stage (raw material → manufacturing → distribution → retail → customer) to identify where a new entrant or investment could capture margin.
- **Life-cycle stage** — is the industry in introduction, growth, maturity, or decline? Growth-stage industries reward market-share-gain strategies; mature industries reward cost/efficiency strategies — this framing should show up explicitly in any industry-analyst-style report or interview answer.

### 7.5 Four more worked market-sizing examples across sectors

**Health insurance — sizing the addressable market for a new digital-first health insurer, urban India, age 25-45.** Top-down: urban population ~490m → age 25-45 ~28% (~137m) → currently uninsured or under-insured ~55% (~75m) → digitally reachable/comfortable buying insurance online ~40% (~30m) → realistic 5-year penetration for a new entrant ~2% (~600,000 policies) × average premium ₹8,000/year ≈ **₹480 crore five-year SOM run-rate**. Bottom-up cross-check: existing digital insurance aggregators report roughly 2-2.5 crore health-insurance policy searches/year online in this age band; a realistic conversion-and-capture rate for a new entrant against established players (~3-5% share of searches converting to that insurer) gives a comparable order-of-magnitude figure — used here to sanity-check the top-down chain rather than as the primary method, since aggregator search volume is a proxy for intent, not guaranteed purchases.

**Quick-service restaurant (QSR) — sizing demand for a new regional (South Indian) QSR chain entering North India.** Bottom-up (the more defensible method for a physical-location business):

target 150 outlets in 5 years across 8 cities  $\times$  average footfall 300 customers/day/outlet  $\times$  average ticket size ₹250  $\times$  350 operating days/year  $\approx$  **₹394 crore annual revenue at full 150-outlet run-rate**. Each input (footfall, ticket size) benchmarked against disclosed same-store-sales metrics of listed QSR comparables (from their annual reports/investor decks) rather than invented — this is the direct equity-research-skill crossover point: reading a competitor's public disclosures to calibrate a private/hypothetical sizing model.

**B2B SaaS — sizing the Indian market for an HR-tech payroll-compliance product targeting mid-size companies (200-2,000 employees).** Bottom-up: number of Indian companies in the 200-2,000 employee band (from MCA/industry-database counts)  $\approx$  45,000  $\rightarrow$  addressable (excludes those already on adequate in-house/competitor systems, estimated via a short win-loss/market-scan exercise)  $\sim$ 60% ( $\sim$ 27,000)  $\rightarrow$  realistic 3-year penetration for a challenger product  $\sim$ 8% ( $\sim$ 2,160 customers)  $\times$  average contract value ₹4.5 lakh/year  $\approx$  **₹97 crore three-year ARR potential**. B2B sizing leans more heavily on a firmographic company count than a population-based top-down chain, since the buying unit is a company, not a consumer — an important distinction to state explicitly when asked to size any B2B market.

**Two-wheeler EV (extending the Part 2.5 conjoint example into full TAM/SAM/SOM) — a specific state, 3-year horizon.** TAM: total two-wheeler unit sales in the state per year (from SIAM/industry-body data)  $\times$  average price  $\approx$  total two-wheeler market value. SAM: the EV-eligible slice — filtered by realistic charging-infrastructure coverage (only cities/towns with adequate density) and by income bands that can afford the EV price premium even after subsidy. SOM: SAM  $\times$  assumed market share for a new entrant given 2-3 established EV players already present, informed by the conjoint-derived share-of-preference simulation (Part 6.8) rather than a flat guessed percentage — this is the cleanest example of primary research output (conjoint utilities) feeding directly into a market-sizing model rather than the two being separate exercises, a connection worth making explicitly in any portfolio writeup.

## 7.6 Competitive intelligence — a fuller worked example (win-loss analysis)

A B2B SaaS company is losing more competitive deals than expected against one named competitor. A win-loss research program:

1. **Sample:** the last 20-30 closed-lost deals against that competitor specifically, plus 10-15 closed-won for contrast, interviewed within 30-60 days of the decision (memory degrades fast beyond this window).
2. **Who to interview:** the actual economic buyer/decision-maker, not just the day-to-day product champion — different stakeholders weight different factors (a CFO cares about total cost of ownership, an IT admin cares about integration effort).
3. **Independent interviewer:** ideally not the salesperson who lost the deal, since prospects are more candid with a neutral researcher, and the salesperson's own account of "why we lost" is a biased secondary source, not primary evidence.
4. **Structured but open discussion guide:** what triggered the evaluation, who else was considered, what ultimately tipped the decision, what would have changed the outcome, unprompted then prompted on standard factors (price, features, support, brand trust, implementation timeline).
5. **Synthesis:** code responses into a small number of loss reasons and report **prevalence with the actual base** ("11 of 24 interviewed closed-lost deals cited implementation timeline as the primary factor" — not "44% of customers"), cross-tabbed against deal size and industry to check whether the loss pattern is universal or segment-specific.
6. **Output:** a ranked, evidence-backed list of the true competitive gaps (as opposed to the sales team's anecdotal, recency-biased narrative) — routinely the single most-cited-as-valuable market research deliverable in B2B software, because it replaces internal opinion with external evidence on exactly the question (why do we lose) that internal teams are least objective about.

## PART 8 — CONSUMER INSIGHTS & REPORTING

---

### 8.1 From data to insight

A **finding** is a fact from the data ("62% of respondents cited price as the top purchase barrier"). An **insight** explains *why it matters* and *what to do* ("price sensitivity is highest among first-time buyers, not repeat buyers — a starter SKU at a lower price point could convert this segment without discounting the core range"). Interviewers and hiring managers explicitly screen for candidates who default to insight, not just findings — this is the single most-repeated piece of feedback in market research interview prep material.

### 8.2 Customer journey mapping

Documenting the stages a customer goes through (awareness → consideration → purchase → onboarding → usage → advocacy/churn), the touchpoints at each stage, and the pain points/emotions at each — used to pinpoint exactly where research or product intervention should focus.

### 8.3 Personas

Composite, research-based archetypes representing a segment ("Priya, 28, first-jobber, price-sensitive, discovers brands via Instagram") — used to make segmentation memorable and actionable for product/marketing teams. Must be grounded in actual data (segmentation output, real quotes), not invented from assumption — a common quality failure point.

### 8.4 Reporting and storytelling

- **Structure:** Executive summary (the "so what" in one page) → methodology → key findings (3-5 headline insights, not 40 slides of data) → detailed appendix.
  - **One slide, one message** — a chart should support a single stated takeaway in the slide title, not force the reader to interpret it themselves.
  - **Show the uncertainty:** sample sizes, margins of error, and data vintage belong on every relevant slide, not buried in an appendix — this is what separates rigorous market research from marketing collateral dressed up as research.
-



## PART 9 — CAREERS, ROLE EXPECTATIONS & INTERVIEW PREP

### 9.1 Typical role ladder

Research Executive/Analyst (0-2 yrs, fieldwork + basic analysis) → Senior Research Analyst/Associate (2-5 yrs, owns a study end-to-end, client/stakeholder-facing) → Research Manager (5-8 yrs, owns a category/client relationship, manages analysts) → Insights Director/Head of Insights (8+ yrs, sits on the leadership team, sets the research agenda).

**"Senior Market Research Analyst" bar (typically 3-5 yrs):** can independently scope a study, choose the right method without guidance, run analysis in SPSS/Excel/Python, and present findings directly to a business stakeholder with a clear recommendation — not just execute a brief handed down from a manager.

### 9.2 What interviewers actually probe for

1. **Methodology judgment** — given a vague business problem, can you design the right study (this handbook's Parts 2-5)?
2. **Numerical fluency** — sample size, significance, basic segmentation logic, market sizing math (Parts 5-7), often via a live case/whiteboard exercise.
3. **"So what" thinking** — can you turn a data table into a decision-ready recommendation (Part 8)?
4. **Domain/business acumen** — do you understand the category (FMCG, banking, fintech, retail) well enough for the insight to be credible, not generic?

### 9.3 Sample case-style interview question (worked)

*"A bank wants to know whether to launch a co-branded credit card with an e-commerce platform. How would you research this?"* - Restate the business decision: launch or not, and if yes, target segment and card benefits. - Design: exploratory (a few interviews with existing co-brand-card holders at competitor banks, desk research on comparable co-brand card economics) → descriptive quantitative survey (interest/intent, preferred benefits, willingness to pay annual fee) on a representative sample of the platform's existing customers, stratified by spend tier. - Sizing: bottom-up from the platform's active user base × estimated card-eligible % (income/credit-score proxy) × estimated adoption rate from comparable launches. - Competitive: benchmark 2-3 existing co-brand cards in market (cashback %, fee, partner) via desk research and MITC (Most Important Terms & Conditions) disclosures. - Recommendation: a go/no-go with the target segment, expected uptake, and the single biggest risk flagged explicitly (e.g. cannibalisation of the bank's existing card base). This structure — decision, design, sizing, competitive check, risk-flagged recommendation — is the reusable skeleton for almost any market-research case interview.

### 9.4 Bridging a finance/derivatives background into market research

The honest gap (see `senior_research_skills.md` in the `job_apply_agent` repo) is direct consumer/survey experience. The honest strength is quantitative rigour most classically-trained market researchers lack: comfort with statistics, sample-size logic, regression, and — critically — **Python for automating survey-data cleaning and analysis**, which is still a differentiator in most agency and in-house insights teams. Frame the F&O/derivatives background as: "trained to read what a large, noisy dataset is actually saying, under time pressure, and to be explicit about uncertainty" — which is the exact same instinct market research analysis requires, applied to a different kind of data.


## PART 10 — BRAND & ADVERTISING RESEARCH

### 10.1 Brand equity — what it is and how it's measured

**Brand equity** is the commercial value a brand name adds beyond the functional product itself (why people pay more for a branded product than an identical unbranded one). Two widely referenced frameworks: - **Aaker's Brand Equity model** — five pillars: brand loyalty, brand awareness, perceived quality, brand associations, and other proprietary assets (patents, trademarks, channel relationships). - **Keller's Customer-Based Brand Equity (CBBE) pyramid** — builds bottom-up: salience (awareness) → performance & imagery (what it does, how it makes you feel) → judgments & feelings (rational and emotional response) → resonance (deep loyalty/identification at the top) — used heavily in brand-tracking questionnaire design, since each pyramid layer maps to a specific survey battery.

### 10.2 Awareness metrics — the standard funnel

- **Unaided (spontaneous) awareness:** "What brands of [category] can you name?" — no list shown; the strongest, most defensible awareness metric since it isn't prompted.
- **Aided awareness:** "Have you heard of [Brand X]?" with the name shown — always higher than unaided, and easy to inflate by suggestion; report both, never aided alone.
- **Top-of-mind (TOM):** the *first* brand named in an unaided-awareness question — the single strongest predictor of default/habitual purchase in low-involvement categories (FMCG especially).
- **Consideration:** "Which brands would you consider buying?" — the funnel stage between awareness and purchase; a brand can have high awareness but low consideration (a well-known but distrusted or "not for me" brand), an important diagnostic distinction.



An illustrative funnel across all the stages this Part covers: note aided awareness (61%) sits well above unaided (22%), exactly as Section 10.2 describes — always report both, never aided alone. The steep consideration-to-trial drop (38% to 24%) is a common, diagnostic pattern: it usually signals a distribution,



price, or trust barrier rather than an awareness problem, since these people already know and would consider the brand but haven't yet bought it — precisely the kind of funnel-stage diagnosis that turns a topline number into an actionable insight (Part 8.1).

### 10.3 Ad/copy testing

- **Pre-testing** (before a campaign airs): concept boards or animatics shown to a sample, measured on message takeaway (did they understand the intended message, unprompted), likeability, brand linkage (do they correctly attribute the ad to the right brand, without prompting — a real and common failure, especially in cluttered categories), and purchase-intent shift pre/post exposure.
- **Post-testing / tracking:** measuring actual campaign recall and brand-metric movement after a campaign has run in-market, via periodic (monthly/quarterly) brand-tracker waves.
- **Key diagnostic: brand linkage failure** — an ad that tests well on likeability but where a meaningful share of viewers cannot correctly name the advertised brand is, commercially, closer to a wasted spend than a "successful" ad; this single point is one of the most repeated pieces of ad-testing feedback given to junior researchers.

### 10.4 Brand tracking studies

A **tracker** is a recurring (usually monthly or quarterly), methodologically-fixed survey run continuously to monitor a standard battery of brand-health metrics (awareness, consideration, usage, NPS, key brand-association ratings) over time, so that movement is comparable wave-on-wave. Design discipline required: sample size, quotas, and exact question wording must stay fixed across waves — even a small wording change breaks comparability and is a common, costly research-governance mistake when handing a tracker between agencies or research managers.

### 10.5 Positioning and perceptual mapping

A **perceptual map** plots brands (including one's own) on two commercially meaningful axes derived from consumer ratings (e.g. "premium ↔ value" and "traditional ↔ modern"), usually built via **factor analysis** or **multi-dimensional scaling (MDS)** on a larger battery of brand-attribute ratings, collapsed to the two dimensions that explain the most variance in how consumers actually differentiate brands (not the two dimensions a marketing team assumes matter). Used to identify **white space** — an attractive, currently unoccupied position on the map — as a basis for repositioning or new-product strategy.

---

## PART 11 — B2B MARKET RESEARCH

---

### 11.1 How B2B research differs from consumer (B2C) research

- **Buying unit:** B2B involves a **buying committee** (economic buyer, technical evaluator, end user, sometimes procurement/legal), not a single consumer decision-maker — research needs to sample and weight multiple stakeholder types per account, not just "the customer."
- **Sample size and access:** B2B populations are far smaller and harder to reach (a market may have only a few hundred relevant companies globally) — research leans qualitative-heavy (Part 4) far more than B2C, and quantitative samples of  $n=100-150$  are often considered large and adequately powered in a B2B context, versus  $n=384+$  as a consumer default (Part 5.5).
- **Sales cycle length:** B2B decisions can take months to years, meaning win-loss research (Part 7.6) and longitudinal account-based tracking matter more than single-point-in-time studies.
- **Firmographics replace demographics:** company size (employees/revenue), industry (NAICS/NIC code), and tech stack/existing vendor relationships are the segmentation variables, in place of age/income/city tier.

### 11.2 Firmographic and technographic segmentation

- **Firmographic:** company size, industry, geography, growth stage (startup/scale-up/enterprise), ownership structure (private/listed/PE-backed) — directly analogous to demographic segmentation but for companies.
- **Technographic:** what software/tools/vendors a company already uses — a strong predictor of both need and switching cost, increasingly available via third-party technographic data providers (e.g. BuiltWith-style tools for web stack) as a secondary-research input before primary fielding.

### 11.3 B2B research methods in practice

- **Expert/stakeholder IDIs** dominate over surveys for market-sizing and competitive-landscape work, given small populations.
- **Account-based research (ABR):** deep-dive research on a small number of named target/strategic accounts rather than a representative sample — common for enterprise sales enablement.
- **Advisory panels / customer advisory boards:** an ongoing structured group of key customers consulted periodically — a hybrid between qualitative research and relationship management.
- **Trade-show and conference intercepts:** a practical, if biased-toward-attendees, way to reach a hard-to-access B2B population efficiently in one place.

### 11.4 B2B and the equity-research crossover

B2B market research overlaps most directly with the **industry/channel-check** work equity and credit analysts already do (calling suppliers, distributors, and customers of a company under coverage to validate a thesis) — a candidate coming from a markets background should lead with this parallel explicitly in interviews, since it is a genuinely transferable, not just adjacent, skill.

---

## PART 12 — RESEARCH OPERATIONS, VENDOR MANAGEMENT & ETHICS

### 12.1 Working with fieldwork agencies and vendors

Most in-house/client-side researchers don't run fieldwork themselves — they commission and manage agencies (for panels, moderators, translators, field interviewers). Core vendor-management skills: - **Writing a clear RFP/scope:** sample size, quotas, incidence rate (what % of the general population qualifies — a low incidence rate, e.g. finding luxury-car owners, sharply raises fieldwork cost and time), timeline, and deliverable format. - **Cost benchmarks to sanity-check quotes:** online quantitative surveys are priced roughly per completed interview (varies hugely by length, incidence, and country); qualitative IDIs are priced per interview including incentive and moderator/analyst time; always request a **per-unit cost breakdown**, not just a lump sum, to compare vendors fairly. - **Quality control of vendor delivery:** spot-checking raw data against the Part 5.8 quality-check criteria before accepting a vendor's "completes," not just trusting the topline numbers handed over.

### 12.2 Fieldwork timelines — realistic planning benchmarks

- Quick online quant survey (n≈300-500, general population, standard incidence): **1-2 weeks** fielding.
- Complex/low-incidence quant (e.g. B2B decision-makers, niche category users): **3-6 weeks**.
- Qualitative IDIs/FGDs (recruitment + fieldwork + analysis): **3-5 weeks** for a standard 15-20 interview study.
- A full mixed-methods project (qual exploratory → quant descriptive → reporting): **8-12 weeks** end to end is a realistic default planning assumption, and stakeholders who ask for this in two weeks need to be told explicitly what has to be cut (sample size, geographic spread, or the qualitative phase) to hit that timeline — a senior analyst's job includes managing this expectation up front, not silently cutting corners.

### 12.3 Research ethics and data privacy (India-specific)

- **Informed consent:** respondents must be told the purpose of the research, how data will be used, and that participation is voluntary — standard practice codified by bodies like ESOMAR internationally.
- **Anonymity/confidentiality:** individual responses should not be traceable back to a named respondent in any client-facing deliverable; only aggregated/anonymised data is reported.
- **India's Digital Personal Data Protection (DPDP) Act, 2023** — the relevant compliance framework for any Indian market research operation collecting personal data (name, phone, email, and combined with other fields, even things like precise location or purchase history): requires a clear, specific purpose for data collection, real consent (not a buried checkbox), data-minimisation (collect only what's needed for the stated research purpose), and defined data-retention limits. A market research analyst does not need to be a lawyer, but should be able to state, for any project, what personal data is collected, why, how long it's retained, and how a respondent could request deletion — increasingly a standard interview question as data-privacy regulation tightens.
- **Research vs marketing/sugging: "sugging"** (selling under the guise of research) and **"frugging"** (fundraising under the guise of research) are explicitly unethical and, per most industry codes, grounds for professional sanction — a legitimate research contact must never pivot into a sales pitch.
- **Vulnerable populations:** additional care (parental consent, simplified language, extra anonymity protection) is required when researching minors, patients, or other vulnerable groups — relevant in

healthcare/pharma and education-sector research specifically.

## 12.4 Building a research repository / insights management

Mature insights functions maintain a **searchable repository** of past studies (tagged by topic, category, date, and key findings) so that a new project starts with "what have we already learned" (Part 3's secondary-research-first principle applied to a company's *own* research history, not just external sources) rather than re-commissioning research that already exists somewhere in the organisation — a real, common inefficiency in large companies with fragmented insights functions, and a differentiator for a candidate who proactively asks "does a repository exist?" on joining a new team.

---

## PART 13 — DIGITAL & SOCIAL RESEARCH METHODS

### 13.1 Social listening

**Social listening** is the systematic monitoring and analysis of public social-media conversation (X/Twitter, Instagram, YouTube comments, Reddit, review sites, forums) about a brand, category, or competitor, using tools (Brandwatch, Sprinklr, Talkwalker, or lighter-weight tools for smaller budgets) that aggregate and code mentions at scale. - **Sentiment analysis**: classifying mentions as positive/negative/neutral, usually via NLP models — useful for trend-spotting but notoriously imperfect on sarcasm, regional slang, and code-switching (mixing English and a regional language mid-sentence, common in Indian social media) — always spot-check a sample of the automated classification against human coding before trusting the topline sentiment number in a client deliverable. - **Share of voice (SoV)**: a brand's mention volume as a % of total category mentions — a rough proxy for relative market visibility/buzz, not a direct measure of sales impact. - **Strength over survey research**: captures *unprompted*, naturally-occurring reactions (nobody is being asked a leading question), and is fast/cheap relative to fielding a survey — a genuine complement to, not full replacement for, structured survey research (Part 5), since social-media populations skew younger, more urban, and more vocal/extreme than the general population. - **Use case**: real-time crisis monitoring (a product-safety issue or PR crisis unfolding), competitive campaign tracking, and early signal detection for emerging trends before they show up in slower quarterly tracker waves (Part 10.4).

### 13.2 Web and app analytics as a research input

Behavioural digital-analytics data (page views, funnel drop-off, session recordings, heatmaps, app-event logs) is technically **secondary data** in the Part 3 sense (already being collected for operational purposes) but is one of the highest-value, most underused inputs to market research when a company has direct digital touchpoints. - **Combining with primary research**: the equity-research-style "channel checks" analogy applies directly here — analytics tells you *what* happened (a specific checkout step has high drop-off), while a follow-up qualitative usability study or exit survey tells you *why* (the Part 5.6/2.5 e-commerce cart-abandonment example is the canonical case). - **A/B testing (online experiments)**: the digital-native version of Part 2.1's causal research design — randomly assigning users to a control (existing) experience and a treatment (new) experience and measuring the difference in a defined metric (conversion rate, time-on-page, revenue-per-visitor). Requires a large enough sample and traffic volume to reach statistical power (the same significance-testing logic as Part 6.4, applied to a live product change rather than a survey question), and requires genuinely random assignment to support a causal claim — informal "before/after" comparisons without a concurrent control group are vulnerable to confounding from any other simultaneous change (seasonality, a marketing campaign, a competitor's move). - **Common pitfall**: stopping an A/B test as soon as it *looks* significant ("peeking") inflates the false-positive rate exactly like the multiple-comparisons problem in Part 6.9 — a pre-registered sample size/duration, decided before launching the test, is the correct discipline.

### 13.3 Online research communities (MROCs)

A **Market Research Online Community (MROC)** is a private, recruited panel of engaged customers/prospects who participate in ongoing discussions, polls, diary exercises, and concept reactions over weeks or months, rather than a single one-off survey or focus group. - **Advantage over a single FGD/survey**: builds longitudinal depth (the same people, tracked over time, reveal how attitudes evolve, e.g. across a product's actual usage lifecycle) and lets a research team pose follow-up questions iteratively as new findings emerge, closer to an ongoing conversation than a single snapshot. - **Trade-off**: more

expensive and operationally heavier to run than a one-off study, and the same self-selection concern as any recruited panel (Part 5.6's online-panel bias) applies — an MROC skews toward more engaged, opinionated customers, useful for depth but not for representative incidence estimates.

### 13.4 Mobile ethnography and passive metering

- **Mobile ethnography:** participants use a smartphone app to capture photos, videos, and short reflections of real moments (a grocery-store visit, a specific usage occasion) as they happen, closing the gap between the Part 4.3 ethnography ideal (observing real behaviour) and the practical cost/access constraints of sending a researcher to physically observe every participant.
  - **Passive metering:** with consent, software passively logs actual behaviour (apps used, time spent, purchase receipts scanned) rather than relying on self-report — directly addresses the stated-vs-revealed-preference gap (Part 4.3) with less participant burden than a full diary study (Part 4.4), though it raises correspondingly higher data-privacy obligations (Part 12.3/DPDP Act) given the depth of personal behavioural data collected.
-

## PART 14 — RETAIL & SHOPPER RESEARCH

### 14.1 Shopper research vs consumer research — a distinction worth stating explicitly

**Consumer research** studies why people choose a *brand/product* (need states, brand perception, usage). **Shopper research** studies behaviour specifically *at the point of purchase* — the in-store or on-page decision moment, which can diverge meaningfully from stated brand preference (a shopper who says they prefer Brand A in a survey may still buy Brand B in-store due to shelf placement, a promotion, or simple habit/convenience). A complete go-to-market research plan typically needs both.

### 14.2 Path-to-purchase and the decision journey in a physical/digital store

Mapping the sequence from store/site entry to purchase: awareness of the category need → navigation to the relevant aisle/category page → comparison among visible/available options → selection → checkout. Each step is a potential drop-off point, and — mirroring the digital funnel-analysis logic in Part 13.2 — in-store traffic patterns (via footfall sensors or, more simply, structured in-store observation) can be analysed the same way as an online conversion funnel.

### 14.3 Share of shelf and planogram research

- **Share of shelf (SOS):** a brand's proportion of shelf space (facings, linear feet, or eye-level positioning) within a category, benchmarked against the brand's actual market share — a brand that is meaningfully under-shelved relative to its market share has a quantifiable, addressable negotiating point for trade/retail relationship teams.
- **Planogram testing:** testing different shelf-layout configurations (by brand block vs by sub-category vs by price tier) via either a physical store test (a small set of "test" stores running the new planogram vs matched "control" stores on the old layout) or, increasingly, virtual-reality/3D-simulated store environments that let many layout variants be tested faster and cheaper than physical resets — a direct extension of the causal-research control-group logic from Part 2.1, applied to store layout rather than a survey question.

### 14.4 In-store audits and mystery shopping

- **Retail audits:** systematic, repeated store visits (by field staff or a syndicated provider like Nielsen's retail audit service) recording stock availability (out-of-stock rate — a chronically under-measured but high-revenue-impact metric, since a stocked-out product simply can't be bought regardless of brand preference), pricing, and promotional compliance (is the agreed promotional price/display actually live in-store).
- **Mystery shopping:** trained researchers pose as ordinary customers to evaluate service quality, staff product knowledge, and process adherence against a structured scorecard — used heavily in banking branch research, retail, and hospitality, and distinct from a standard customer-satisfaction survey because it captures *observed* service delivery rather than a customer's *recalled* impression of it.

### 14.5 E-commerce-specific shopper research

Digital-shelf equivalents of the above: **share of search** (a brand's visibility in on-platform search results for relevant category terms, the e-commerce analogue of physical share of shelf), **content/listing audits**



(image quality, review count/rating, A+ content presence — all measurable drivers of on-platform conversion), and **competitive price-tracking** (automated monitoring of competitor pricing across marketplaces, feeding directly into the pricing-research toolkit from Part 5.7's Van Westendorp/Gabor-Granger methods).

---

## PART 15 — INTERNATIONAL & CROSS-CULTURAL RESEARCH

### 15.1 Why multi-country research isn't just "run the same survey in more places"

A survey instrument validated in one market cannot be assumed to work identically elsewhere — question wording, response-scale interpretation, and even the underlying construct being measured (what "quality" or "value" *means* to a respondent) can differ meaningfully across cultures. Treating a multi-country study as a simple translation exercise is one of the most common and costly mistakes in international research.

### 15.2 Translation and back-translation

The standard rigour check: a questionnaire is translated from the source language into the target language by one translator, then **independently translated back** into the source language by a second translator who has not seen the original — discrepancies between the original and the back-translated version flag wording that likely doesn't carry the same meaning, before the survey is ever fielded. Simple, literal translation without this check risks subtly changing what's actually being measured (e.g. an idiom or brand-attribute term that has no direct equivalent).

### 15.3 Cultural response-style bias

Documented, systematic differences in how different cultural groups use rating scales, independent of their actual underlying opinion: - **Acquiescence bias**: a tendency to agree with statements regardless of content, more pronounced in some cultural contexts than others. - **Extreme response style**: a tendency to favour the endpoints of a scale (strongly agree/strongly disagree) over the middle options, again with documented cross-cultural variation. - **Practical implication**: comparing raw mean Likert scores (Part 5.3) across countries without adjusting for these response-style differences can produce misleading conclusions (e.g. "Country A loves this product more than Country B" when the gap is partly or wholly a scale-usage artifact, not a genuine preference difference) — techniques like within-respondent standardisation or focusing on rank-order/relative comparisons within each country, rather than raw cross-country score comparisons, help correct for this.

### 15.4 Sampling and fieldwork logistics across markets

Incidence rates, panel quality, and appropriate data-collection modes (Part 5.6) can vary hugely by country — a mode that works well in one market (online panels in a highly digitally-penetrated market) may badly under-represent the population in another (where CAPI/CATI or intercept methods remain necessary for a representative sample) — a multi-country study's methodology section should specify and justify the mode used *per market*, not assume one mode fits all.

### 15.5 Where this matters directly for an India-based analyst

Multinational companies operating in India (or Indian companies expanding into other emerging markets) routinely need India-specific adaptations even within a "global" study framework: language (a single "Hindi" version is often insufficient given India's genuine linguistic diversity — a truly national sample may need several regional-language instrument versions), rural/urban mode differences (CAPI still dominant in many rural fieldwork contexts, per Part 5.6), and category-specific cultural translation (e.g. "value for money" carries different weight and meaning across Indian income segments than in a Western market context) —

an analyst who can speak specifically to these adaptation points, rather than treating India as a single homogeneous market, demonstrates real category depth in an interview.

---


## PART 16 — ADVANCED PRICING RESEARCH & REVENUE MANAGEMENT

### 16.1 Beyond Van Westendorp and Gabor-Granger: choice-based pricing research

Part 5.7 and Part 6.8 covered the Van Westendorp Price Sensitivity Meter and choice-based conjoint applied to pricing. A fuller pricing-research toolkit an analyst should know: - **Gabor-Granger** (introduced briefly in Part 5's cross-references): respondents are shown a single price and asked a simple purchase-intent question ("would you buy at ₹X?"); if yes, the price is raised and asked again; if no, it's lowered — iterating to trace out a full demand curve at the individual level. Simpler to field than conjoint, but more vulnerable to anchoring on the *first* price shown, and doesn't capture trade-offs against competing products the way a full choice-based conjoint does. - **Brand-Price Trade-Off (BPTO)**: respondents repeatedly choose among branded products at varying prices across several rounds, revealing brand-price substitution patterns (at what price does a respondent switch from their preferred brand to a cheaper alternative) — useful specifically for understanding brand equity's cash value in price terms, a natural bridge to the brand-equity frameworks in Part 10.

### 16.2 Price elasticity of demand — the concept every pricing researcher must connect findings to

**Price elasticity of demand (PED)** = % change in quantity demanded ÷ % change in price. A category/product with  $PED > 1$  (in absolute value) is **elastic** — demand is highly sensitive to price, and a price cut can increase total revenue if the volume gain outweighs the per-unit price loss.  $PED < 1$  is **inelastic** — demand barely moves with price, and a price increase generally raises revenue since volume loss is proportionally smaller than the price gain. **Worked example**: a 10% price increase leads to a 4% volume decline →  $PED = -4\%/10\% = -0.4$  (inelastic, in absolute terms  $0.4 < 1$ ) → revenue impact  $\approx (1.10 \times 0.96 - 1) \approx +5.6\%$ , meaning the price increase is revenue-accretive despite losing some volume. Every pricing-research deliverable should ultimately translate a "willingness to pay" finding into this kind of revenue-impact framing, since that is what a pricing committee actually decides on.



### 16.3 Revenue management and dynamic pricing research

In categories with perishable inventory or capacity constraints (airlines, hotels, ride-hailing, event ticketing), pricing research extends into **revenue management** — understanding how willingness to pay varies by purchase timing, remaining inventory, and customer segment (business vs leisure traveller, for instance), so price can be dynamically adjusted to maximise total revenue rather than set once and left static. Market research supports this with segment-specific willingness-to-pay studies (e.g. a business traveller's WTP curve is typically far less price-elastic than a leisure traveller's for the same route) that feed the pricing algorithms revenue-management teams build — a genuine crossover point between market research and quantitative/pricing-analytics roles.

### 16.4 Common pricing-research pitfalls

- **Hypothetical bias:** stated willingness to pay in a survey setting is a well-documented overestimate of what respondents will actually pay with real money on the line — always discount stated-preference pricing findings, and validate with a real-money pilot (a limited-time actual price test) wherever feasible before a full pricing decision.
- **Ignoring competitive response:** a pricing study that estimates a company's own optimal price in isolation, without modelling how competitors might react (a price cut inviting a competitive price war, for instance), can recommend a price that looks optimal in the research but destroys value once competitive dynamics play out — game-theoretic thinking, not just a single-firm demand curve, belongs in any pricing recommendation for a competitive category.
- **Underweighting non-price value drivers:** a pricing study run in isolation from the broader value proposition (service quality, brand trust, convenience) risks over-indexing on price as the lone lever, when in reality willingness to pay is itself a function of these other factors — pricing research is most useful when integrated with the broader brand/product research (Part 10) rather than run as a standalone exercise.

## PART 17 — SYNDICATED DATA & PANEL-BASED RESEARCH

### 17.1 What syndicated research actually provides, beyond "a report you can buy"

Syndicated research firms (Nielsen, Kantar, IQVIA for pharma, GfK) run large, continuously-operating panels and retail-audit infrastructure and sell standardised access to the resulting data across many subscribing clients — the key economic idea being that the fixed cost of running a national retail audit or a household purchase panel is shared across many subscribers, making data available at a fraction of what any single company would pay to build the same infrastructure alone.

### 17.2 Retail audit panels vs household/consumer panels — a distinction worth stating precisely

- **Retail audit panels** (e.g. Nielsen's retail measurement service) track what moves *out of stores* — sales, distribution (% of stores stocking a product), pricing, and promotional activity, measured by regularly auditing a representative sample of stores' stock and sales records. This answers "how much of X sold, at what price, in what stores" — a supply-side, store-level view.
- **Household/consumer panels** (e.g. Kantar's Worldpanel) recruit a representative panel of households who scan or report every purchase they make, over time — this answers "who is buying X, how often, are they new or repeat buyers, what else do they buy" — a demand-side, individual/household-level view that retail audit data alone cannot provide (retail audit can't tell you if the same 100 units sold were bought by 100 different first-time buyers or 10 loyal repeat buyers making 10 purchases each).
- **Why an analyst needs both:** a brand can show flat retail-audit sales while its household panel data reveals it's actually losing repeat buyers but gaining an offsetting number of new triers — a materially different, more actionable diagnosis than the flat topline sales number alone would suggest, and exactly the kind of insight that separates syndicated-data fluency from just reading a topline sales report.

### 17.3 Key syndicated metrics an analyst should be able to define on the spot

- **Distribution (numeric and weighted):** numeric distribution = % of stores stocking a product; weighted distribution = % of *category sales* happening in stores that stock it (weights by store importance, not just store count) — a product could have low numeric distribution (in few stores) but high weighted distribution if it's specifically in the highest-volume stores, a much healthier position than the reverse.
- **Penetration:** % of households that bought the brand/category at least once in a period — the household-panel-derived growth lever distinct from frequency (Part 17.2).
- **Buying rate / frequency:** among buying households, how often and how much they buy — penetration and frequency together decompose total volume growth into "more people buying" vs "existing buyers buying more," a standard first-cut diagnostic for any FMCG brand-health review.
- **Source of growth / source of business:** for a new product launch, syndicated panel data can decompose incremental category sales into cannibalisation of the company's own existing products vs genuine incremental category growth vs share taken from competitors — a critical, board-level question ("is this new product actually growing the pie, or just eating our own other products' sales") that syndicated panel data is uniquely positioned to answer.

## 17.4 Limitations to disclose when using syndicated data

Panel-based estimates carry sampling error like any survey (Part 5.5's logic applies), category/channel coverage can have gaps (e.g. historically weaker coverage of small kirana-store channels relative to organised retail in some syndicated services, though this has improved), and there's often a reporting lag (syndicated data for a period is typically available some weeks after that period closes, not in real time) — an analyst presenting syndicated-data-based findings should be able to state these limitations proactively rather than presenting the numbers as a perfectly precise, real-time picture of the market.

---



## PART 18 — NEW PRODUCT DEVELOPMENT & CONCEPT TESTING

### 18.1 Where research fits in the NPD funnel

New product development typically funnels from many raw ideas down to the few that actually launch: **idea generation** → **concept screening** (quick, cheap filtering of many concepts) → **concept testing** (deeper quantitative validation of a shortlist) → **product testing** (does the actual formulation/prototype deliver) → **market testing/simulated test market** (a small-scale or simulated launch before full rollout) → **launch**. Market research has a distinct, purpose-built method at nearly every stage — using the wrong stage's method too early (e.g. running a full concept test on 50 raw, unfiltered ideas) wastes budget that a cheaper screening method would have handled.

### 18.2 Concept screening vs concept testing

- **Concept screening:** fast, low-cost, often qualitative or a simple quantitative gate (e.g. a short online survey across many concepts, n=50-100 per concept, measuring basic appeal/uniqueness/purchase-intent) — designed to cheaply eliminate weak ideas before investing in deeper research on the stronger ones.
- **Concept testing:** a fuller quantitative study (larger n, richer diagnostic battery) on the shortlisted concepts that survive screening, measuring not just overall appeal but *why* — specific attribute ratings, competitive comparison, and often price sensitivity (linking directly to Part 16's pricing-research toolkit) — used to make the actual go/no-go and refinement decisions before committing to product development investment.

### 18.3 The standard concept-test measurement battery

- **Purchase intent:** typically a 5-point scale ("definitely would buy" to "definitely would not buy") — industry norms usually apply a discount to the top-box score (Part 18.5 covers why) rather than taking it at face value.
- **Uniqueness/differentiation:** does the concept offer something meaningfully different from what's already available — a concept that scores high on appeal but low on uniqueness may simply be appealing because it resembles an already-successful competitor, raising cannibalisation and differentiation concerns.
- **Believability:** do respondents actually believe the concept's claims are achievable/credible — a concept can be highly appealing *if true* but score poorly on believability, flagging a communication or proof-point problem to fix before launch.
- **Relevance:** does the concept address a need the respondent actually has — high appeal with low relevance often signals a novelty response that won't translate into a genuine, repeatable need-driven purchase.
- **Value perception:** initial price-value read, refined later with the dedicated pricing-research toolkit (Part 16) once the concept has cleared the appeal/differentiation/believability/relevance bar.

### 18.4 Monadic vs sequential-monadic vs comparative concept testing

- **Monadic design:** each respondent evaluates only *one* concept (different respondents see different concepts) — avoids order/comparison bias and gives a clean, standalone read on each concept, but

requires a larger total sample (since each concept's score comes from a fully separate respondent group) and doesn't directly tell you which concept respondents would choose *between* the options.

- **Sequential monadic:** each respondent evaluates multiple concepts one at a time, rating each independently before moving to the next — more sample-efficient than pure monadic (each respondent contributes to multiple concepts' scores) but introduces order effects (the first concept seen can anchor subsequent ratings) that need to be controlled via rotation (randomising which concept each respondent sees first).
- **Comparative/proto-monadic:** respondents directly compare and rank/choose among concepts shown together — most efficient for a pure "which wins" question, but doesn't isolate each concept's standalone appeal the way monadic does, and doesn't reflect the real-world reality that consumers rarely see all competing concepts side by side at the actual moment of purchase.

## 18.5 Worked example — interpreting a concept test's purchase-intent score

A concept test shows 42% of respondents rate "definitely would buy" (top-box) and a further 31% rate "probably would buy" (top-two-box = 73%).

**Model answer.** Raw top-box/top-two-box scores from a concept test systematically overstate real-world trial rates — stated intent in a hypothetical, no-cost survey setting doesn't carry the friction of an actual purchase decision (the same hypothetical-bias principle from Part 16.4's pricing-research discussion applies directly to purchase-intent scores). Category-specific **norms databases** (maintained by research agencies from hundreds of prior concept tests in the same category) are the correct benchmark — a 42% top-box score should be compared against the norm for that specific category (e.g. "the FMCG snacking-category top-box norm is 35%") rather than interpreted as an absolute predictor of 42% of the market buying the product; a concept scoring meaningfully above its category norm is a genuine positive signal, even though the raw percentage overstates the actual likely trial rate.

## 18.6 Product testing and sensory research

Once a concept is validated, **product testing** evaluates whether the actual formulation/prototype delivers against the promise: - **Blind vs branded testing:** a blind test (unbranded, unlabeled) isolates the product's own performance from brand-driven expectation bias; a branded test measures the *combined* effect of the product plus the brand halo — running both and comparing reveals how much of a product's appeal comes from the brand name itself versus the underlying product performance (a well-known, striking finding in many categories: a well-loved brand's product can lose a blind taste test to a competitor yet still win decisively when branded, revealing the brand equity's real, measurable value). - **JAR (Just-About-Right) scales:** for sensory attributes (sweetness, spiciness, thickness), asking not just "how much do you like this" but "is the sweetness too little / just about right / too much" — directly actionable for reformulation, since a JAR scale pinpoints *which direction* to adjust a specific attribute, something a simple overall-liking score can't do alone.

## PART 19 — CATEGORY MANAGEMENT & TRADE MARKETING RESEARCH

### 19.1 What category management is, and where research fits

**Category management** is the retailer-and-supplier collaborative practice of managing a product category (e.g. "biscuits" or "hairecare") as a strategic business unit — deciding assortment (which SKUs to stock), shelf layout, pricing, and promotion at the category level rather than brand-by-brand — with the stated goal of growing the *category's* total sales (a genuinely shared retailer-supplier interest), not just any single brand's share within it. Market research supports this with category-level (not just brand-level) demand, segmentation, and shopper-behaviour data that a retailer needs to make assortment and space-allocation decisions.

### 19.2 Assortment research — the SKU rationalisation question

A recurring, high-stakes category-management question: "should we add/drop this SKU?" Assortment research addresses this via: - **Incrementality analysis**: does a candidate new SKU add genuinely incremental category sales, or does it primarily cannibalise the retailer's existing SKUs in the category (the same source-of-growth logic from Part 17.3, applied at the assortment-decision level rather than the new-product-launch level)? - **Trade-off/MaxDiff-based assortment optimisation**: using MaxDiff (Part 6.5) or choice-based conjoint (Part 6.8) across a full candidate SKU list to estimate which specific combination of SKUs captures the most shopper preference for a given amount of shelf space — directly informing a data-driven "keep vs drop" assortment decision rather than relying on gut feel or which supplier has the strongest existing retailer relationship.

### 19.3 Trade promotion effectiveness research

Evaluating whether a promotional spend (a temporary price cut, an end-cap display, a "buy one get one") actually pays back: - **Baseline vs incremental sales**: using syndicated panel/audit data (Part 17), decompose total sales during a promotion period into the **baseline** (what would have sold anyway, absent the promotion, estimated from non-promoted periods) and the **incremental lift** attributable to the promotion itself. - **Pull-forward/pantry-loading effect**: a common finding — a meaningful share of a promotion's apparent "incremental" volume is often existing buyers simply stocking up early (pulling forward future purchases) rather than genuinely new demand, meaning the post-promotion period frequently shows a sales *dip* below baseline that must be netted against the promotion-period lift to get a true incremental read, not just looking at the promotion week in isolation. - **ROI calculation**: (incremental revenue from the promotion – cost of the promotion, including the trade discount and any listing/display fees) ÷ cost of the promotion — trade marketing teams increasingly demand this explicit ROI framing from research rather than a simple "sales went up during the promotion" headline, which as the pull-forward point above shows, can be a misleadingly generous read of a promotion's true value.

### 19.4 The retailer relationship — research as a shared-value tool

Category-management research is somewhat unusual among market-research disciplines in that its findings are often presented jointly to, or co-developed with, the retail partner (rather than being purely internal to the manufacturer) — a supplier's category-management/insights team credible enough to bring genuinely retailer-beneficial, not just manufacturer-self-serving, analysis to a retailer relationship earns more shelf-space and assortment influence over time; research seen as one-sided advocacy for the supplier's own

brands, without genuine category-growth framing, is a fast way to lose credibility and influence with a retail partner.

---

## PART 20 — MARKETING MIX MODELING & MEDIA EFFECTIVENESS RESEARCH

### 20.1 What Marketing Mix Modeling (MMM) is, and how it differs from everything else in this handbook

**Marketing Mix Modeling** is a statistical (typically regression-based) technique that decomposes sales or another business KPI into the contribution of each marketing input — TV spend, digital spend, price, promotions, distribution, seasonality, and a base/organic level — using historical time-series data, usually at a weekly or monthly level across markets. Unlike most methods in this handbook (surveys, qualitative research, conjoint), MMM doesn't ask anyone anything — it's built entirely from observational business and marketing-spend data, closer in spirit to the regression driver analysis in Part 6.7 but applied to aggregate sales rather than individual survey responses.

### 20.2 The core MMM equation and what "decomposition" means in practice

A simplified MMM takes the form:  $\text{Sales} = \text{Base} + f(\text{TV spend}) + f(\text{Digital spend}) + f(\text{Price}) + f(\text{Promotion}) + f(\text{Distribution}) + f(\text{Seasonality}) + \text{error}$ . Each  $f()$  function typically incorporates two well-established marketing-response phenomena: - **Adstock (carryover)**: advertising's effect doesn't happen all in the week it airs — it decays over subsequent weeks (a geometric decay is common: this week's effective TV impact = this week's actual spend + a decay-rate-weighted carryover of recent weeks' spend), reflecting that consumers who saw an ad last week may still be influenced by it this week. - **Diminishing returns (saturation)**: each additional rupee of spend on a channel typically produces a smaller incremental sales lift than the rupee before it (a concave response curve) — the first ₹10 lakh of TV spend in a market might move sales meaningfully, the next ₹10 lakh on top of already-heavy spend, much less so. Once fitted, the model decomposes total historical sales into how much each input contributed — "TV drove 8% of sales, digital drove 5%, price/promotion drove 12%, and 62% was base/organic demand" — a business-critical output for budget-allocation decisions.

### 20.3 From decomposition to ROI: media-mix optimisation


The real business value of MMM isn't just decomposition — it's using the fitted response curves (Part 20.2) to calculate **marginal ROI per channel** (what does the *next* rupee of spend on each channel return, given where current spend already sits on that channel's diminishing-returns curve) and re-optimize the budget allocation across channels to maximise total sales/profit for a given total spend. A channel with high *average* ROI can still have low *marginal* ROI if it's already heavily saturated — the classic MMM insight that surprises marketing teams who look only at average ROI ("TV has always worked for us") without checking where they currently sit on that channel's saturation curve.

### 20.4 Worked example — reading an MMM output table and making a budget call

*An MMM decomposition shows: TV spend ₹40 cr/year contributing 6% of sales, with an estimated marginal ROI of 0.6 (i.e. the next rupee spent returns ₹0.60 of sales) — below breakeven after accounting for margin. Digital spend ₹15 cr/year contributing 5% of sales, with a marginal ROI of 2.1.*

**Model answer.** TV is at a marginal ROI below 1 (indeed, below what's needed to cover typical gross margin, meaning the *next* rupee spent on TV is unprofitable), strongly suggesting the channel is over-

saturated at current spend levels even though it still contributes meaningfully to total sales (past spend on TV isn't "wasted" — sunk investment already delivered its 6% contribution; the finding is specifically about the next incremental rupee). Digital's marginal ROI of 2.1 is well above breakeven, suggesting room to shift incremental budget from TV toward digital until their marginal ROIs converge (the textbook optimal-allocation rule: keep shifting budget from lower-marginal-ROI to higher-marginal-ROI channels until the marginal returns equalise across channels) — this reallocation logic, not the raw contribution percentages, is what a media/marketing budget committee actually needs from an MMM deliverable.



The chart makes the marginal-vs-average-ROI distinction visually obvious: at ₹40cr, TV's curve has nearly flattened — the tangent slope (marginal ROI) there is shallow even though the curve's height (average/total contribution) is still substantial. At ₹15cr, Digital's curve is still climbing steeply — its tangent slope is much higher, which is exactly why shifting the next rupee from TV to Digital raises total sales even without spending a rupee more in total.

## 20.5 MMM's limitations and how it complements (not replaces) survey-based research

MMM is powerful for allocation decisions across already-used channels with enough historical spend variation to estimate response curves reliably, but it cannot easily answer "why" a channel works (no attitudinal or motivational data — that's what the brand-tracking and qualitative methods in Parts 4 and 10 provide), struggles with genuinely new channels that have no historical spend history to model, and requires enough weeks/markets of spend *variation* to statistically identify each channel's effect (a channel spent at a constant % of revenue every single week gives the model little to work with). A mature insights function typically runs MMM for macro budget-allocation decisions alongside brand tracking (Part 10.4) and ad-testing (Part 10.3) for the qualitative "why" and creative-execution questions MMM can't answer.

## PART 21 — CUSTOMER RETENTION, CHURN & LIFETIME VALUE RESEARCH

### 21.1 Why retention research is treated as a distinct discipline from acquisition-focused research

Most of this handbook's methods (concept testing, brand tracking, segmentation) are oriented toward acquiring new customers or understanding the broad market. Retention research asks a different question — *why do existing customers leave, and what predicts who will leave next* — and draws on a different primary data source: the company's own behavioural/transactional data on existing customers (usage frequency, support tickets, payment history), closer in spirit to Part 13.2's digital-analytics discussion than to a traditional survey-first approach, though survey-based exit/win-back research remains a key complementary tool.

### 21.2 Defining and measuring churn

**Churn rate** = the % of customers who stop being customers over a defined period. The definition itself requires a deliberate research choice: **contractual churn** (a subscription explicitly cancelled — unambiguous to measure) versus **non-contractual/behavioural churn** (no explicit cancellation event, so "churned" must be *defined* as, e.g., "no purchase/login in the last 90 days" — a threshold choice that materially affects the measured churn rate and must be justified, not picked arbitrarily). **Voluntary churn** (customer actively leaves) is distinguished from **involuntary churn** (e.g. a payment method fails and isn't updated) since the two require completely different interventions — a win-back campaign doesn't address a failed-payment problem, which needs a payment-retry/dunning-management fix instead.

### 21.3 Predictive churn modelling — what market research contributes beyond the data-science model


A data-science/analytics team typically builds the statistical churn-prediction model itself (logistic regression or a machine-learning classifier predicting churn probability from usage and account features). Market research's distinct contribution: **exit surveys and win-back interviews** that capture the *stated reason* for leaving (which the behavioural data alone often can't reveal — a customer's usage dropping before cancelling tells you churn was coming, not *why*), and **qualitative research with at-risk-but-not-yet-churned customers** (identified by the predictive model) to understand what intervention might actually retain them before they leave, closing the loop between "the model says this customer is likely to churn" and "here's what we should actually do about it."

### 21.4 Customer Lifetime Value (CLV) — the metric that reframes retention as a revenue question

**CLV** (Customer Lifetime Value) estimates the total profit a customer will generate over their entire relationship with the company, commonly approximated as:  $CLV = (\text{Average Revenue per Customer} \times \text{Gross Margin \%}) / \text{Churn Rate}$  (a simplified perpetuity-style formula, directly analogous to the Gordon Growth perpetuity formula in the Valuation chapters — both capitalise a recurring flow using a rate that captures how long the flow persists). **Worked example:** average revenue per customer ₹2,400/year, gross margin 65%, annual churn rate 20% →  $CLV = (2,400 \times 0.65) / 0.20 = ₹7,800$ . This single number reframes retention research as a revenue question a finance-literate stakeholder immediately understands: reducing



churn from 20% to 15% in this example raises CLV to  $(2,400 \times 0.65) / 0.15 = ₹10,400$  — a 33% CLV increase from a 5-percentage-point churn reduction, exactly the kind of business-case framing that gets a retention initiative funded, and a calculation that plays directly to a candidate's finance background as a genuine crossover skill into market research.



The curve's shape is the real lesson: CLV doesn't fall linearly as churn rises — it's a  $1/\text{churn}$  relationship, so the CLV gain from cutting churn is much larger at *low* churn rates than the same percentage-point cut delivers at high churn rates. Going from 10% to 5% churn roughly doubles CLV, while going from 35% to 30% churn (the same 5-point cut) barely moves it — a genuinely important nuance when prioritising which customer segment's retention to invest in first.

## 21.5 CAC-to-CLV ratio — connecting retention research back to acquisition economics

**CAC** (Customer Acquisition Cost) compared against CLV (Part 21.4) gives the fundamental unit-economics check every subscription/recurring-revenue business must pass: a CLV:CAC ratio meaningfully above 1 (a commonly cited healthy benchmark is 3:1 or higher) indicates the business earns back its acquisition spend with room for profit over a customer's lifetime; a ratio near or below 1 signals the business is effectively losing money on each customer relationship despite what topline revenue growth might suggest. This is the single metric that most directly connects this Part's retention/CLV research back to the acquisition-focused research covered everywhere else in this handbook — a market researcher who can move fluently between "how do we get customers" and "what's a customer actually worth once we have them" operates at a genuinely senior, business-partner level rather than a narrow methodological specialist level.



## PART 22 — NEUROMARKETING & IMPLICIT MEASUREMENT TECHNIQUES

### 22.1 Why implicit methods exist — the gap between what people say and what actually drives behaviour

Every method in Parts 4-5 (interviews, focus groups, surveys) relies on respondents being willing and *able* to accurately report their own reactions and motivations. A large body of psychology and consumer-behaviour research shows people are often genuinely unaware of the true drivers of their own choices (much decision-making is fast, automatic, and emotionally-driven rather than the deliberate, rational process a direct survey question assumes) — implicit measurement techniques attempt to measure response *before* conscious rationalisation shapes the stated answer, directly addressing the stated-vs-revealed-preference gap flagged repeatedly elsewhere in this handbook (Part 4.3's ethnography rationale, Part 18.5's hypothetical-bias discussion).

### 22.2 The core implicit-measurement toolkit

- **Eye-tracking:** records exactly where and how long a respondent's gaze rests on a stimulus (a package design, an ad, a shelf layout) — used to answer concretely "did they even see the logo/price/key claim," a question self-report notoriously answers unreliably (people often can't accurately recall or report what they visually attended to).
- **Facial coding:** automated analysis of micro-expressions (via camera and computer-vision software) while a respondent views a stimulus, classifying moment-by-moment emotional response (surprise, positive affect, confusion) — used heavily in ad-testing to identify the *exact moment* in a video ad where engagement drops or a negative reaction spikes, a level of temporal precision a post-hoc survey question ("how did you feel about the ad") cannot provide.
- **Implicit Association Test (IAT):** measures response-time differences when a respondent pairs a brand with positive versus negative concepts, on the premise that faster pairing reflects a stronger, more automatic (implicit) association — used to surface brand associations respondents may not consciously report or may not even be aware they hold.
- **Galvanic skin response (GSR) / biometric arousal measures:** measures physiological arousal (skin conductance, heart rate) as a proxy for emotional intensity (though not valence — GSR shows *how strongly* someone is reacting, not whether the reaction is positive or negative, so it's typically paired with facial coding or self-report to determine direction).

### 22.3 Where implicit methods add genuine value versus where they're overused

Implicit methods are most defensible when the research question is specifically about attention, emotional intensity, or automatic association (ad pre-testing, packaging/shelf-standout testing, brand-association mapping) — situations where the gap between stated and automatic response is itself the thing being measured. They are frequently over-applied or oversold when used as a wholesale replacement for standard survey/qualitative research on questions that don't actually hinge on implicit/automatic response (e.g. using eye-tracking to answer a straightforward feature-preference question a simple MaxDiff exercise, Part 6.5, would answer more directly and cheaply) — a mature researcher matches the *method* to the *specific gap* it addresses, not to what sounds most technologically sophisticated in a pitch to a client.

## 22.4 Practical and ethical constraints

Implicit/biometric methods typically require lab-based or specialised equipment (though eye-tracking and facial coding increasingly work via standard webcams in online studies), meaningfully raising cost and sample-size constraints versus a standard online survey — sample sizes for implicit studies are often in the dozens, not hundreds, making them better suited to diagnostic, exploratory-style questions (Part 2.1) than to precise, generalisable incidence estimates. Ethically, biometric/facial data is more sensitive personal data than a standard survey response, requiring explicit informed consent (Part 12.3) specifically covering biometric data collection, storage, and retention — a heightened version of the DPDP Act obligations already discussed for standard personal data.

---

## PART 23 — FINANCIAL SERVICES & INVESTMENT PRODUCT RESEARCH

### 23.1 Why this Part exists — the most direct bridge in this handbook to a stock-market-role background

Every method covered so far applies to any consumer or B2B category. This Part applies them specifically to researching **financial products and investor behaviour** — mutual funds, insurance, broking/trading platforms, lending products, wealth management — the category most directly adjacent to an equity/markets background, and the one where a candidate's finance domain knowledge is a genuine, direct advantage over a generalist market researcher.

### 23.2 What's structurally different about researching financial-product decisions

- **Trust and risk perception dominate the decision more than in most consumer categories:** a financial product is a promise about the future (returns, protection, liquidity) rather than an immediately-verifiable experience good — research must specifically probe trust in the institution/brand, not just product features, since a technically superior product from a low-trust provider frequently loses to an inferior product from a trusted one.
- **Financial literacy varies enormously across the same target population,** more so than most product categories — a survey instrument that assumes respondents understand terms like "expense ratio" or "sum assured" will get unreliable answers from a large share of even the *target* population; research design must include either careful screening on financial literacy or plain-language, scenario-based question framing rather than jargon-based questions.
- **Regulatory disclosure requirements shape what can be researched and how:** SEBI's advertising and disclosure code for mutual funds/broking products, and IRDAI's equivalent for insurance, constrain how concept tests and ad tests for financial products can be framed (e.g. mandatory risk disclosures), and a researcher testing financial-product concepts needs at least a working awareness of these constraints to design a realistic, compliant concept stimulus.

### 23.3 Investor segmentation — a financial-services-specific application of Part 6.3

Beyond generic demographic/attitudinal segmentation, investor research typically segments on: - **Risk tolerance and investment horizon** (conservative/moderate/aggressive, short/medium/long horizon) — directly analogous to a fund's own risk-profiling questionnaire, and often validated against actual portfolio composition/behaviour data where available (revealed risk tolerance vs stated risk tolerance, echoing the stated-vs-revealed-preference theme from Part 22.1). - **Financial sophistication/self-directed vs advised:** a self-directed, high-engagement trader (checking positions daily, comfortable with derivatives) has fundamentally different research needs and product-feature priorities than an advisor-dependent, low-engagement investor — conflating the two in a single undifferentiated survey is a common and costly research design error for a broking or wealth-management client. - **Life-stage-driven need states:** a first-jobber building an initial portfolio, a mid-career investor balancing growth and children's-education goals, and a pre-retirement investor prioritising capital preservation have distinctly different product needs, messaging responsiveness, and channel preferences — segmentation frameworks in this category are frequently built around financial life stage rather than pure demographics.

## 23.4 Worked example — researching adoption barriers for a new investment product

*A brokerage wants to understand why adoption of its new options-trading feature has been slower than a comparable competitor's launch. Behavioural data shows most users view the feature's onboarding screen but don't complete activation.*

**Model answer.** This is a diagnostic research problem that should follow the same funnel-analysis-first discipline as the digital-research methods in Part 13.2: pull the drop-off data (where exactly in the activation flow do users abandon) before commissioning primary research. If most drop-off happens at a specific step (e.g. a risk-disclosure/suitability questionnaire required by regulation), a small set of moderated usability sessions (Part 4.3/13.2's usability-testing logic) watching real users attempt that specific step would likely reveal whether the barrier is comprehension (users don't understand what's being asked), friction (the step is too long/complex), or genuine hesitation (users realise at that step that options trading is riskier than they'd assumed and self-select out) — each diagnosis points to a completely different fix (simplify language, shorten the flow, or accept that this is appropriately filtering out unsuitable users and isn't actually a "problem" to solve). This mirrors the e-commerce cart-abandonment example from Part 2.5 almost exactly, applied to a financial-product activation funnel instead of a checkout funnel — the same diagnostic discipline transfers directly.

## 23.5 Tracking investor sentiment and confidence — a financial-services-specific tracker application

Beyond the standard brand-tracker metrics (Part 10.4), financial-services trackers commonly include an **investor confidence/sentiment index** — a recurring survey measure of investors' forward-looking market outlook and risk appetite, tracked over time. These indices (several syndicated versions exist in the Indian market from research/broking firms) serve a dual purpose: as a client-facing thought-leadership content asset for the firm that runs them, and as a genuine leading-indicator research signal — a sharp drop in self-reported risk appetite ahead of any actual market decline can be an early behavioural signal worth cross-referencing against the market-efficiency/behavioural-finance framework covered in the Equity & Capital Markets material, directly connecting this handbook's market-research toolkit back to the equity-research and technical-research content elsewhere in this compilation.

## PART 24 — AGILE & RAPID RESEARCH METHODS

### 24.1 Why traditional research timelines increasingly clash with business decision speed

The classic research timeline (Part 12.2: 8-12 weeks for a full mixed-methods project) was built for an era of slower product-development and marketing-decision cycles. Modern product and growth teams — especially in digital-first and startup environments — often need a directional answer in days, not months, to keep pace with sprint-based development cycles. Agile/rapid research is the discipline of deliberately trading some depth and rigour for speed, in a controlled, transparent way, rather than simply skipping research when time is short.

### 24.2 Unmoderated remote testing platforms

A major enabler of rapid research: platforms that let a researcher script a task (e.g. "find and add this product to your cart," "react to this concept") and recruit panelists who complete it independently on their own device, with the platform automatically capturing screen recordings, think-aloud audio, and survey responses — turning what used to require in-person moderated sessions and days of scheduling into a study that can field and return results within 24-48 hours. The trade-off versus moderated research (Part 4.1's IDIs): no ability to probe follow-up questions in real time based on what a participant says, so the initial task/question script must be unusually well-designed since there's no researcher present to rescue a poorly-worded task.

### 24.3 The rapid-research decision framework — when speed trade-offs are acceptable

Not every research question is a good candidate for a rapid/agile approach. A useful discipline: rapid methods are well-suited to **directional, lower-stakes, iterative** decisions (does concept A or B resonate better before a deeper dive; is this specific UI flow confusing) where being wrong is cheap and quickly correctable in the next sprint. They are poorly suited to **high-stakes, hard-to-reverse decisions** (a major market-entry decision, a pricing strategy that's expensive to unwind once launched) where the cost of an under-powered, rushed answer is much higher than the cost of waiting for a properly-scoped study — a senior researcher's real value-add is often exactly this triage judgment, not blanket enthusiasm for either "always rapid" or "always rigorous."

### 24.4 Continuous/always-on research vs the traditional project-based model

Some organisations move beyond one-off rapid studies toward a **continuous research** model — an always-available panel or ongoing lightweight survey/testing capability that product and marketing teams can tap into on demand throughout a sprint, rather than commissioning discrete projects. This requires different organisational infrastructure (a maintained, consented panel; templated, reusable study designs; a research team resourced for constant small studies rather than periodic large ones) but can dramatically shorten the time from "we have a question" to "we have an answer," at the cost of requiring more research-team bandwidth spread thinner across more, smaller studies.

## 24.5 Maintaining rigour discipline under speed pressure

The single biggest risk in rapid research is quietly abandoning the methodological discipline covered throughout this handbook (unbiased question wording, adequate base sizes even for a quick study, avoiding leading tasks) simply because the timeline is compressed — speed should change the *scope and depth* of a study (fewer questions, a smaller but still adequately powered sample, a narrower research objective), not the *rigour* applied to whatever is actually asked. A rapid study that cuts corners on wording or sampling to save an extra day produces a fast, confidently-wrong answer, which is worse than no answer at all if a team acts on it — this tension (speed vs rigour, and where genuinely to compromise) is one of the more sophisticated judgment calls a senior researcher is expected to navigate, and a strong differentiator in how a candidate discusses rapid research in an interview.

---

## PART 25 — QUALITATIVE DATA ANALYSIS: FROM TRANSCRIPTS TO THEMES

### 25.1 Why this deserves its own full Part, not just Part 4.5's two-line summary

Part 4.5 introduced coding and thematic analysis briefly. In practice, turning 20-30 hours of interview/FGD transcripts into a credible, defensible set of themes is one of the most methodologically substantial and most frequently under-explained parts of the entire research process — a candidate who can walk through this process step by step, not just name "thematic analysis" as a term, stands out sharply in an interview.

### 25.2 The coding process, step by step

1. **Familiarisation:** reading (or re-reading, and ideally re-listening to) all transcripts before any formal coding begins, to build a holistic sense of the data before fragmenting it into codes — skipping this step and coding "cold" risks anchoring too heavily on the first few transcripts read.
2. **Open/initial coding:** working through transcripts line by line, applying short descriptive labels ("codes") to meaningful segments of text — at this stage, codes are granular and numerous (often 50-100+ initial codes across a project), deliberately over-inclusive rather than prematurely narrow.
3. **Developing a codebook:** as initial coding proceeds, similar codes are consolidated and defined precisely in a shared **codebook** (a document listing each code, its definition, and an example quote) — essential both for consistency across a project's own analysis and, critically, for enabling a second coder to apply the same codes consistently (Part 25.4).
4. **Axial coding / grouping into themes:** related codes are grouped into broader **themes** — a theme is a higher-level pattern of meaning that recurs across multiple codes and, ideally, multiple respondents, not just a single vivid quote from one person.
5. **Reviewing and refining themes:** checking each candidate theme against the full dataset — does it genuinely capture a coherent pattern, or does it collapse two distinct ideas together, or is it really a sub-theme of a bigger pattern? This iterative refinement, not the initial grouping, is usually where the real analytical insight emerges.
6. **Naming, defining, and writing up themes:** each final theme gets a clear name, a one-paragraph definition, and is illustrated with a small number of well-chosen (not cherry-picked) verbatim quotes in the final deliverable.

### 25.3 Inductive vs deductive coding — a distinction worth stating explicitly

**Inductive coding** lets themes emerge from the data itself, with no predetermined framework — appropriate for genuinely exploratory research (Part 2.1) where the researcher doesn't want to constrain what might be found. **Deductive coding** applies a predetermined framework or set of codes derived from existing theory, a prior study, or specific research objectives the client has already specified — appropriate when the research has narrower, more specific objectives and comparability to a prior study or framework matters. Most real projects use a hybrid: starting with a small set of deductive codes tied directly to the research objectives, while remaining open to inductively adding new codes for genuinely unanticipated patterns that emerge.



## 25.4 Inter-coder reliability — how to defend that your themes aren't just one researcher's subjective read

When a project's credibility depends on the coding being more than one analyst's personal interpretation, a second, independent coder codes a subset of the same transcripts using the shared codebook (Part 25.2, step 3), and the two coders' results are compared for agreement — commonly quantified via **Cohen's Kappa** (a statistic that corrects raw percentage agreement for the agreement expected by chance alone; a Kappa above roughly 0.6-0.7 is generally considered acceptable agreement for qualitative coding). Where coders disagree, discussing and resolving the disagreement often improves the codebook itself, catching ambiguous code definitions the original single-coder process would never have surfaced. This is rarely done for every project (it adds real time and cost) but is a standard practice for higher-stakes qualitative research feeding a major business decision.

## 25.5 Qualitative analysis software — what it actually does (and doesn't do)

Software like **NVivo** and **Atlas.ti** doesn't perform the analytical thinking itself — it manages the *mechanics* of coding at scale: attaching codes to specific transcript segments, allowing codes to be renamed/merged/split across an entire project without manually re-editing every transcript, and generating outputs like code-frequency counts or code-co-occurrence matrices (which codes tend to appear in the same passages, often a genuine analytical lead worth investigating further). For smaller projects, some researchers code manually in a spreadsheet or even by hand on printed transcripts — perfectly legitimate for a small enough dataset, since the software's value is organisational efficiency at scale, not a shortcut to the analytical thinking the researcher must still do.

---

## PART 26 — SAMPLING WEIGHTING & POST-STRATIFICATION

### 26.1 Why even a well-designed sample usually still needs weighting

Part 5.4 covered probability sampling. In practice, even a carefully executed probability sample rarely lands with a respondent profile that exactly matches the true population on every dimension that matters — response rates differ by subgroup (younger respondents often under-respond to phone surveys, for instance), and a panel's composition drifts from the general population over time. **Weighting** corrects for this gap after data collection, adjusting each respondent's contribution to the final results so the weighted sample matches known population benchmarks (typically from census or other authoritative sources) on key variables.

### 26.2 The core weighting mechanics — a worked example

*A survey targets a population known to be 51% female / 49% male, but the achieved sample is 60% female / 40% male.*

**Model answer.** Each female respondent is downweighted and each male respondent upweighted so the weighted totals match the true 51/49 population split:

Female weight =  $51\% / 60\% = 0.85$  (each female respondent counts as 0.85 of a "person")  
 Male weight =  $49\% / 40\% = 1.225$  (each male respondent counts as 1.225 of a "person")

A female respondent's individual answers are unchanged — weighting adjusts how much each response *counts* toward the topline, not the response itself. This single-variable example generalises to **multi-variable (raking/RIM) weighting** when several variables (age, gender, region, income) all need to be simultaneously balanced — an iterative algorithm that adjusts weights across all target variables together, since correcting for gender alone can throw age balance off and vice versa if done independently.

### 26.3 Post-stratification vs simple weighting — a related but distinct technique

**Post-stratification** divides the sample into strata (e.g. age × gender × region cells) and weights each cell to match its known population share — a more granular version of the single-variable weighting in Part 26.2, correcting for imbalance across *combinations* of variables, not just each variable's marginal distribution independently. It requires the population benchmark data to exist at the cell level (not just each variable's overall split), which isn't always available, particularly for less common demographic combinations in a smaller market.

### 26.4 The honest limits and risks of weighting

Weighting corrects for known, measurable imbalances (demographics the researcher can benchmark against authoritative population data) — it cannot correct for *unmeasured* bias (e.g. if people willing to complete a 20-minute survey are systematically more opinionated or engaged than the true population on the actual topic being studied, weighting by demographics alone won't fix that). **Design effect:** heavy weighting (some respondents weighted several times more than others) inflates the *effective* margin of error beyond what the raw sample size alone would suggest — a heavily-weighted  $n=1,000$  survey can have the real statistical precision of a much smaller unweighted sample, a nuance often omitted from a topline

methodology note but one a rigorous analyst should be able to explain when asked "how much do you trust this specific weighted subgroup finding."

---

## PART 27 — DATA STORYTELLING & EXECUTIVE PRESENTATION

### 27.1 Why this is a distinct skill from Part 8.4's reporting principles

Part 8.4 covered the structural discipline of a good research report (BLUF, one-slide-one-message, showing uncertainty). This Part goes further into the craft of actually **presenting** findings live to senior stakeholders — a genuinely different skill from writing a report someone reads alone, since a live presentation must hold attention, handle real-time pushback, and land a recommendation within a fixed, often short, meeting window.

### 27.2 Structuring a live findings presentation for a senior, time-constrained audience

A senior stakeholder (a CMO, a category head, an investment committee) rarely has more than 15-20 minutes of genuine attention for a research readout, regardless of how the deck is formatted. The discipline: lead with the recommendation and the single most important finding supporting it (mirroring the research-note BLUF principle from the Equity & Capital Markets material — the same "lead with the conclusion" logic applies identically here), hold the bulk of the supporting data and methodology in an appendix available if asked, and rehearse the ability to deliver the *entire* core story in under 5 minutes in case the meeting gets cut short — a genuinely common occurrence with senior stakeholders' calendars.

### 27.3 Anticipating and handling pushback live

A senior stakeholder pushing back on a finding in real time ("that doesn't match what I'm hearing from the sales team") is not a research failure — it's a completely normal, in fact healthy, part of a live presentation, and how it's handled matters more than avoiding it. Effective handling: acknowledge the specific tension without becoming defensive, distinguish between the *finding* (what the data shows) and the *interpretation* (what it means), and be prepared to state precisely what evidence exists and doesn't exist ("we didn't specifically test that segment, so I can't speak to it directly — that's worth a follow-up" is a stronger, more credible answer than overreaching beyond what the research actually supports). A presenter who becomes visibly defensive under a reasonable challenge damages the research's credibility far more than a well-handled "I don't know, here's how we'd find out" answer.

### 27.4 Data visualisation principles specific to a live-presentation context (not a written report)

- **Fewer data points per slide than a written report would use:** a chart that's perfectly readable in a document a reader can study at their own pace is often too dense for a live audience seeing it for 20 seconds — simplify further than Part 8.4's already-lean "one slide, one message" principle for anything presented live.
- **Progressive disclosure for complex charts:** building up a complex chart in stages (animating or clicking through layers) rather than showing the full complexity at once, so the live audience's attention is guided rather than overwhelmed on first view.
- **A single, consistent visual language across the whole deck** (consistent colour meaning — e.g. always using the same colour for "our brand" vs competitors across every chart) reduces the cognitive

load of re-interpreting a new colour scheme on every slide, letting the audience focus on the finding rather than re-learning the chart each time.

## **27.5 Tools of the trade for presentation-stage work**

Beyond the analysis tools already covered (Part 6.6's SPSS/Python/R for the analysis itself), presentation-stage work typically uses PowerPoint or Google Slides for the deck itself, with charts either built natively or exported from the analysis tool and restyled for visual consistency (Part 27.4) — and increasingly, lightweight interactive dashboards (Power BI, Tableau, or simpler tools) for stakeholders who want to explore a tracking study's data themselves between formal presentation readouts, extending the standard research deliverable beyond a single static presentation moment.

---

## PART 28 — MARKET ENTRY & EXPANSION RESEARCH

### 28.1 Why market-entry research is a distinct project type, not just "market sizing plus competitive intelligence"

Parts 7 (market sizing, competitive intelligence) and 15 (international/cross-cultural research) each cover pieces of what a market-entry study needs, but a genuine market-entry decision — should we launch in this new geography or category at all — combines them into a single, higher-stakes deliverable with its own structure and a go/no-go (or how-to-enter) recommendation at the end, warranting its own dedicated treatment.

### 28.2 The market-entry research structure, end to end

1. **Market attractiveness:** sizing the opportunity (TAM/SAM/SOM, Part 7.1-7.2) combined with growth-rate and structural-trend analysis — is this a genuinely growing, structurally attractive market, or a large-but-stagnant/declining one.
2. **Competitive landscape:** mapping incumbent players, their market shares, competitive intensity, and — critically for entry decisions specifically — barriers to entry (regulatory licensing requirements, capital intensity, entrenched distribution relationships, brand loyalty) that a new entrant would face (Part 7.3-7.4's frameworks applied specifically through an entry-barrier lens).
3. **Customer/demand-side validation:** primary research (Part 4-5) validating that genuine, sufficient unmet demand exists for the entrant's specific value proposition in this market — not just that the category itself is large, but that there's a real gap the entrant can credibly fill.
4. **Regulatory and operational feasibility:** for cross-border entry specifically, licensing/approval requirements, local partnership or ownership rules, and operational feasibility (supply chain, distribution access) — often requiring legal/regulatory expertise beyond the research function itself, but the research team must know enough to flag these as gating factors early rather than only after a full market study is complete.
5. **Entry-mode recommendation:** assuming the market is attractive and feasible, is the recommended approach organic build, a joint venture/local partnership, or acquisition of an existing local player — a distinct strategic question the research should inform (e.g. primary research revealing strong incumbent brand loyalty and complex regulatory relationships might favour a partnership/acquisition entry mode over an organic build, even if the market itself sizes attractively).

### 28.3 Worked example — sizing and validating entry into a new geographic market

*A company successful in metro-market India considers expanding into Tier-2/Tier-3 cities for the same product category.*

**Model answer.** This isn't simply "the same market, smaller cities" — it requires validating several distinct assumptions that don't automatically transfer from the metro experience: does the same customer need state exist in Tier-2/3 markets (income levels, category penetration, and awareness may differ substantially, requiring primary research rather than assuming metro learnings apply); does the existing distribution/retail model work in these markets (Tier-2/3 retail structure often differs meaningfully from metro organised retail, echoing Part 14's shopper-research distinctions); and does the price point that works in metro markets hold in Tier-2/3, where price sensitivity is often structurally higher (Part 16.2's price-elasticity framework, likely showing a more elastic demand curve in these markets). A rigorous entry study

would run a small pilot in a representative sample of Tier-2/3 cities before a full rollout commitment, treating the pilot itself as a primary research instrument (Part 24's rapid-research logic — a real-world pilot resolves the hypothetical-bias problem, Part 16.4, that a pure survey-based study can't fully escape).

## 28.4 Category-entry research — the non-geographic version of the same question

The same structure applies when a company considers entering a new *product category* (not a new geography) — e.g. a personal-care brand considering entering the wellness-supplements category. The competitive-landscape and barrier-to-entry analysis shifts from geographic/regulatory barriers toward brand-permission questions specifically: does the company's existing brand equity (Part 10.1's Aaker/Keller frameworks) credibly extend into the new category in consumers' minds, or would entry require building entirely new brand credibility from scratch — a **brand-extension-fit study** (typically a quantitative concept test measuring purchase intent and credibility for the brand specifically in the new category, benchmarked against how a completely unbranded/new-brand concept performs on the same measures) is the standard research instrument for this specific question.

---



## PART 29 — ADVANCED STATISTICAL TECHNIQUES: STRUCTURAL EQUATION MODELING

### 29.1 Why SEM goes beyond the regression toolkit already covered

Part 6.7's regression driver analysis measures how a set of *observed* variables (survey ratings) predict a single observed outcome (NPS). **Structural Equation Modeling (SEM)** extends this in two ways that matter for more sophisticated brand and satisfaction research: it can model **latent variables** (unobserved constructs like "brand trust" or "perceived quality" that are only indirectly measured through several related survey items, rather than any single directly-observed question), and it can model **multiple linked relationships simultaneously** (a full causal chain — e.g. product experience → satisfaction → loyalty → advocacy — rather than a single-outcome regression).

### 29.2 Latent variables — why a single survey question rarely captures a real construct well

A construct like "brand trust" isn't well-measured by any single question ("do you trust this brand, yes/no") — real research typically asks several related items (trust in product safety, trust in the company's honesty, trust the company will resolve problems fairly) and treats "brand trust" as a **latent variable** estimated from the shared variance across those multiple observed items, filtering out each individual item's own measurement noise. This is directly related to the factor analysis mentioned in Part 10.5's perceptual mapping — SEM can be understood as extending factor analysis (extracting latent constructs from multiple items) by then also modelling how those latent constructs causally relate to each other and to outcomes.

### 29.3 Path models — visualising and testing a full causal chain

A SEM **path model** diagrams the hypothesised relationships between constructs as arrows (e.g. Product Quality → Satisfaction → Loyalty, with Satisfaction also directly influencing a separate Advocacy/NPS construct) and estimates the strength of each path simultaneously, along with **fit statistics** that indicate how well the entire hypothesised structure matches the actual data pattern (common fit indices include CFI, RMSEA, and chi-square/df ratio — a researcher doesn't need to derive these by hand, but should recognise the names and know they indicate whether the proposed causal structure is a statistically reasonable representation of the data, versus a poorly-fitting model that needs revision). This lets a researcher test not just "does X predict Y" (a regression question) but "does the full hypothesised chain of relationships hold together as a coherent structure" — a materially more rigorous test of a brand or customer-experience theory.

**SEM path model: product/service quality -> satisfaction -> loyalty -> advocacy**

Boxes = constructs (gold = latent, estimated from multiple survey items) · numbers = standardised path coefficients

The diagram makes the two SEM extensions from Part 29.1 concrete: the gold boxes (Satisfaction, Loyalty) are latent constructs, each estimated from several underlying survey items rather than any single question, and the full set of arrows is estimated *simultaneously* as one coherent model rather than as separate regressions — so the 0.52 path from Satisfaction to Loyalty and the 0.29 direct path from Satisfaction to Advocacy are estimated jointly, correctly accounting for the fact that Satisfaction's total effect on Advocacy runs both directly and indirectly through Loyalty.

## 29.4 When SEM is worth the complexity, and when it's overkill

SEM requires meaningfully larger sample sizes than simple regression (rough rules of thumb suggest at least 10-20 respondents per estimated parameter in the model, quickly requiring several hundred respondents for even a modest model), specialised software (R's `lavaan` package, or dedicated tools like AMOS/Mplus), and more advanced statistical training to run and interpret correctly. It's genuinely worth this investment when the research question specifically concerns testing a multi-step causal theory with latent constructs (e.g. validating a full customer-experience-to-loyalty model for an ongoing tracking programme that will be reused across many waves) — it's overkill for a one-off, simpler question a standard regression driver analysis (Part 6.7) would answer just as well at a fraction of the sample-size and complexity cost. Knowing when *not* to reach for the more sophisticated technique is as much a marker of genuine expertise as knowing how to run it.

## PART 30 — RESEARCH BUDGETING & ROI JUSTIFICATION

### 30.1 Why a research function itself needs an ROI story, not just individual project budgets

Part 12.1-12.2 covered vendor cost benchmarks and timelines for a single project. This Part addresses a distinct, more senior-level question: how does an insights function justify its *overall* budget and headcount to leadership — a genuinely different challenge from costing one study, since it requires demonstrating the cumulative business value the function delivers, not just that any one project came in on budget.

### 30.2 Building a research budget — the components

A research function's annual budget typically breaks into: **fieldwork/vendor costs** (the largest variable component — panel access, moderator fees, incentive payments, syndicated-data subscriptions), **software/tools licensing** (analysis software, survey platforms, qualitative-coding tools), **headcount** (usually the largest fixed component), and **discretionary/contingency** (for rapid, unplanned requests that inevitably arise mid-year, per Part 24's agile-research reality). A mature budgeting process ties planned fieldwork spend to a **research roadmap** — a prioritised list of the year's planned studies mapped to specific business decisions they'll inform — rather than requesting a lump budget with no specific link to what it will be spent investigating.

### 30.3 Demonstrating research ROI — the honest difficulty and the standard approaches

Unlike a marketing campaign with a directly trackable revenue outcome, research's value is often several steps removed from a final business result (research informs a decision, the decision drives an outcome, other factors also influence that outcome) — making a clean, single ROI number for the research function as a whole genuinely difficult to construct honestly. Standard approaches, used in combination rather than any one alone: **decision-influence tracking** (logging which specific business decisions each study informed, and, where possible, the estimated value of getting that decision right versus wrong — e.g. a concept test that prevented a low-scoring product from launching, avoiding an estimated development/launch cost); **stakeholder satisfaction/usage surveys** (do the internal teams who commission research report it's actually used and valuable — a leading indicator of value even without a hard revenue number); and **cost-avoidance framing** (research that correctly predicted a launch would underperform, preventing that specific spend, is a legitimate and often persuasive ROI argument even without a "revenue generated" number).

### 30.4 The build-vs-buy decision for research capability

A recurring leadership question: should a specific research capability be built in-house (a dedicated analyst, an owned panel) or bought externally (agency partners, syndicated subscriptions) on an as-needed basis? The framework mirrors any make-vs-buy decision: **build** when the capability is used frequently enough that in-house fixed cost beats the cumulative cost of repeated external engagements, and when the capability is close enough to the core business that institutional knowledge built up over repeated studies compounds in value; **buy** for infrequent, specialised, or highly variable-skill needs (a one-off international study needing

local-market expertise the in-house team doesn't have) where building permanent capability wouldn't be utilised often enough to justify the fixed investment.

### 30.5 Worked example — building the business case for adding a second researcher

*An insights team of one analyst is repeatedly unable to take on requests from product and marketing teams due to capacity constraints, with an estimated 40% of incoming research requests currently declined or significantly delayed.*

**Model answer.** A credible headcount business case combines several of Part 30.3's approaches: quantify the volume and nature of declined/delayed requests (not just "we're busy," but a specific count and characterisation — how many were high-stakes decisions proceeding *without* research support as a result, a genuine risk-exposure argument); estimate the cost of the current gap using cost-avoidance framing (e.g. "of the 12 declined requests this quarter, at least 2 were pricing decisions of a scale where a pricing-research miss could cost more than a second analyst's full annual salary"); and benchmark the requested headcount against a reasonable analyst-to-stakeholder-team ratio if that data exists elsewhere in the organisation or industry. This is a substantively different, more senior-level pitch than a single-project budget request — it requires the ability to characterise a *pattern* of unmet need and its business risk, not just defend one study's cost, exactly the kind of capability an analyst moving from execution-focused work toward research-leadership responsibilities needs to develop.

---

## PART 31 — COMMERCIAL DUE DILIGENCE & M&A-SUPPORT RESEARCH

### 31.1 Why this Part is the strongest direct crossover with equity/investment-banking work in this entire handbook

**Commercial due diligence (CDD)** is market research commissioned specifically to support an M&A transaction, private equity investment, or major strategic decision — validating the target company's market position, growth prospects, and competitive standing *before* a deal closes. Unlike almost anything else in this handbook, CDD is typically commissioned by the same investment banking, private equity, or corporate development teams that sit adjacent to equity research (Part on Equity & Capital Markets elsewhere in this compilation) — a candidate with a markets/finance background moving into a CDD-focused research role is making one of the most natural, directly-relevant transitions covered anywhere in this material.

### 31.2 What a CDD engagement actually covers, and how it differs from a standard market-sizing study

A CDD engagement typically combines several methods already covered in this handbook into one time-pressured, high-stakes deliverable: **market sizing and growth-rate validation** (Part 7 — critically, independently verifying the target company's own market-size claims in its investment materials, since management/sell-side data rooms have an obvious incentive to present optimistic figures); **competitive positioning** (Part 7.3-7.4 — where does the target genuinely rank against competitors, not just where its own materials claim it ranks); **customer reference calls** (a CDD-specific primary research method — structured interviews with a sample of the target's actual customers, assessing satisfaction, switching risk, and price sensitivity, functioning similarly to Part 7.6's win-loss analysis but focused on retention risk rather than competitive loss reasons); and **management-team credibility assessment**, often informed by industry-expert interviews about the target's reputation and execution track record.

### 31.3 The compressed timeline and its methodological implications

CDD engagements typically run on a **2-4 week timeline**, dramatically compressed relative to even the rapid-research methods in Part 24, driven by the deal-process clock rather than the researcher's preference — this forces heavy reliance on secondary research and expert-network interviews (fast to arrange) over large-scale primary quantitative surveys (too slow to field and analyse within the window), and demands the same speed-vs-rigour judgment from Part 24.5 applied under even tighter real-world constraints, since a CDD finding genuinely influences whether a deal proceeds and at what valuation.

### 31.4 Worked example — a CDD red flag surfaced through customer reference calls

*A private equity firm is evaluating an acquisition of a B2B software company. The target's data room shows 95% gross revenue retention. Customer reference calls (a sample of 15 customers, stratified by contract size) reveal that 4 of the 15 report they are actively evaluating switching to a competitor at their next renewal, citing pricing and a recent product-quality complaint.*

**Model answer.** A 95% headline retention rate describes *past* behaviour (customers who have renewed to date); the reference-call finding is a forward-looking signal about *future* retention risk that the historical metric alone cannot capture — if the pattern in this 15-customer sample (roughly 27% expressing switching

intent) generalises to the broader customer base, it represents a material, currently-invisible risk to the deal thesis that the historical 95% retention figure would not have surfaced on its own. This is exactly the kind of finding that justifies commercial due diligence as a distinct step beyond simply reading the data room — it demonstrates why PE/IB teams commission primary research specifically to stress-test forward-looking assumptions (future retention, in this case) that historical financial data alone cannot validate, directly informing the deal team's valuation model and negotiating position (e.g. building this risk into a downside revenue-retention scenario, or negotiating deal terms — an earnout structure, a price adjustment — that account for it).

### **31.5 Deliverable format — how CDD reporting differs from a standard research readout**

A CDD deliverable is typically structured explicitly around the transaction's key investment-thesis questions (rather than a general market overview), often organised as a direct "supports/doesn't support/mixed" verdict against each specific claim in the target's own investment materials — a structural difference from the standard research-report format (Part 8.4, Part 27) that reflects CDD's purpose: not general market understanding, but a specific, time-pressured input to a go/no-go investment decision with real capital at stake.

---

## PART 32 — SUSTAINABILITY & ESG CONSUMER RESEARCH

### 32.1 A distinct question from the equity-research ESG material elsewhere in this compilation

The Equity & Capital Markets material elsewhere in this library covers ESG from an investment-risk perspective — how governance, environmental, and social factors affect a company's cash flows and cost of capital. This Part asks a genuinely different question: what do actual *consumers* believe, feel, and how do they behave regarding sustainability and ESG-related claims — a market-research question about attitudes and purchase behaviour, not an investment-risk assessment, though the two connect (consumer sustainability attitudes are themselves one of the demand-side factors that can eventually show up as a company-level ESG-related revenue or reputational risk).

### 32.2 The well-documented attitude-behaviour gap in sustainability research specifically

Sustainability research has one of the most extensively documented **stated-vs-revealed-preference gaps** (Part 22.1's broader theme) of any research category: survey after survey shows a large majority of consumers stating they prefer sustainable/ethical products and would pay a premium for them, while actual purchase behaviour (measured via sales data, not self-report) shows a far smaller share genuinely paying that premium consistently. A market researcher covering this category must design specifically around this known gap — leaning more heavily on revealed-preference methods (actual purchase/panel data, Part 17.2's household panels; real-money pricing pilots, Part 16.4) rather than trusting stated purchase-intent survey questions on sustainability at face value, an even more pronounced caution than the general hypothetical-bias warning applied elsewhere in this handbook.

### 32.3 Segmenting sustainability attitudes rather than treating "the sustainable consumer" as one segment

Consumer sustainability attitudes are not uniform across a population — segmentation research in this space typically reveals distinct groups ranging from genuinely high-commitment, price-insensitive sustainability-driven buyers (usually a meaningful minority) through to a much larger "sustainability-aware but price-constrained" middle segment (genuinely cares, but price/convenience wins in the actual moment of purchase) to a segment largely indifferent to sustainability claims — treating the market as one undifferentiated "green consumer" segment, a common client assumption a researcher should push back on, leads to poorly-targeted sustainability-marketing spend aimed at a segment too small or too price-sensitive to justify a premium-priced sustainable offering.

### 32.4 Greenwashing perception research — a distinct, increasingly important sub-topic

As sustainability claims have proliferated, consumer skepticism toward unsubstantiated or vague claims ("greenwashing") has grown into its own researchable phenomenon — studies increasingly measure not just "does this sustainability claim appeal to you" but "does this specific claim seem credible, or does it read as marketing spin," since a claim that triggers greenwashing skepticism can actively damage brand trust rather than helping it, the opposite of the intended effect. This connects directly to Part 10.3's ad-testing

believability measure, applied specifically to sustainability claims, and to the ESG-disclosure material in the Equity & Capital Markets ESG chapter — a company making sustainability claims to consumers that aren't well-supported by its actual disclosed ESG data (BRSR filings, per that chapter) faces genuine reputational and regulatory risk if the gap becomes public, connecting this consumer-research topic back to the equity-analyst ESG-risk lens from Part 32.1.

### 32.5 Worked example — testing willingness to pay for a sustainable product reformulation

*A packaged-goods company considers reformulating a product with more sustainable (but higher-cost) packaging. A survey shows 68% of respondents say they'd be willing to pay a 10% price premium for the new packaging.*

**Model answer.** Per Part 32.2's attitude-behaviour gap, this stated 68% figure should not be taken at face value for a pricing/business-case decision — the rigorous approach discounts it substantially (informed by category norms for the gap between stated and revealed sustainability willingness-to-pay, similar to Part 18.5's norms-database logic for concept-test scores) and, wherever feasible, validates with a real-money method: a limited regional or online-only pilot actually charging the 10% premium and measuring real purchase behaviour, rather than committing to a full national reformulation based on the survey figure alone. This is precisely the kind of "don't trust stated intent alone" discipline this Part's opening theme (Part 32.2) is built around, applied to a concrete, common client request.

---



## PART 33 — PRE-IPO & INVESTOR PERCEPTION RESEARCH

### 33.1 A market-research discipline that sits directly inside the capital-raising process

The Equity & Capital Markets material elsewhere in this compilation covers the IPO process from the issuer/banker side — book-building, pricing, allotment. This Part covers a distinct but directly related activity: **investor perception research**, commissioned by a company (often via its investment bank or a specialised research/advisory firm) ahead of a significant capital-markets event — an IPO, a large follow-on offering, or even ahead of an annual investor day — to understand how institutional and analyst audiences currently perceive the company, well before the roadshow itself begins.

### 33.2 What a pre-IPO perception study actually measures

Structured interviews with a sample of institutional investors, sell-side analysts, and sometimes credit-rating analysts (drawing on the same expert-interview methodology as Part 4.1's IDIs, but with a highly specific, senior, hard-to-access respondent population) typically probe: **awareness and existing perception** of the company and its sector (does this audience already have a view, and is it accurate); **perceived strengths and concerns** about the business model, growth story, and management team, in the audience's own words rather than the company's self-description; **valuation expectations** (an informal read on what multiple/range this investor audience would consider fair, useful triangulation input alongside the formal book-building process); and **messaging effectiveness** — testing draft investor-presentation narratives and positioning statements before they're used in the actual roadshow, functionally similar to Part 10.3's ad/message testing but for an investor audience instead of a consumer one.

### 33.3 Why this research exists — the cost of a poorly-perceived roadshow

A roadshow is an extremely compressed, high-stakes communication window (management meets dozens of institutional investors over just days) — discovering during the roadshow itself that investors consistently misunderstand the growth story, or have a specific unaddressed concern (a governance question, a competitive worry), is far more costly than surfacing and addressing that same gap in a perception study conducted weeks earlier, when the investor presentation and messaging can still be revised. This is the same underlying logic as Part 18.5's concept-testing-before-launch discipline, applied to an investor-communications "product" (the equity story) instead of a consumer product.

### 33.4 Worked example — a perception study surfacing a messaging gap ahead of a roadshow

*A mid-cap company preparing for a follow-on offering commissions investor perception interviews with 15 institutional investors. A recurring theme: several investors express confusion about how the company's newer digital-services segment relates to its legacy core business, with a few explicitly stating this ambiguity makes them cautious about the growth story.*

**Model answer.** This is precisely the kind of finding a perception study is designed to surface before it becomes a roadshow-day surprise — the company's management team, deeply familiar with the business, may not realise external investors lack the same clarity on how the two segments fit together strategically. The actionable output isn't just "investors are confused" (a finding, per Part 8.1's finding-vs-insight

distinction) but the specific insight: the investor presentation needs a clearer, more explicit segment-relationship narrative (e.g. a simple diagram showing how the digital segment leverages the core business's existing customer relationships) addressing this exact confusion directly, tested again with a follow-up round of interviews if time allows before the actual roadshow — turning a research finding into a concrete presentation-design fix, the same "insight, not just finding" discipline this handbook returns to throughout.

---

## PART 34 — INFLUENCER & CREATOR ECONOMY RESEARCH

### 34.1 A genuinely new research object that didn't fit older frameworks

Influencer/creator marketing has grown into a large enough spend category that it warrants its own research treatment distinct from Part 10's traditional brand/ad research — an individual creator carries personal audience trust and relationship dynamics a traditional media placement doesn't, and measuring an influencer partnership's effectiveness requires methods adapted specifically for this format.

### 34.2 Influencer selection research — beyond follower count

A common, well-documented mistake is selecting influencer partners primarily on raw follower count, which correlates poorly with actual campaign effectiveness. More rigorous selection research examines: **audience-brand fit** (does the influencer's actual audience demographic/psychographic profile match the target segment — obtainable via platform analytics the influencer or their agency shares, or estimated via third-party social-analytics tools); **engagement quality, not just rate** (a high engagement rate inflated by bot followers or engagement-pod schemes is a well-known measurement-integrity problem, requiring audience-authenticity checks before committing spend); and **past brand-partnership performance and audience sentiment in comments** (a qualitative read on how an influencer's audience has historically responded to sponsored content, distinct from organic content — some creators see meaningful engagement/sentiment drop-off specifically on sponsored posts, a real signal worth checking before partnering).

### 34.3 Measuring campaign effectiveness — adapting the standard toolkit


Influencer campaign measurement borrows directly from Part 10's ad-effectiveness toolkit, adapted for the format: **brand lift studies** (pre/post awareness and consideration surveys among the influencer's audience, mirroring Part 10.3's ad pre/post-testing logic); **unique tracked links/codes** for direct response attribution (a more direct, e-commerce-native measurement than most traditional media allows); and, for larger, ongoing influencer programmes, incorporating influencer spend as its own channel within a Marketing Mix Model (Part 20) alongside traditional TV/digital spend, allowing the same marginal-ROI framework used for other channels to be applied to influencer spend specifically.

### 34.4 The disclosure and authenticity research dimension

Consumer research on sponsored-content disclosure (does the audience recognise a post as paid partnership, and does that recognition change their trust/purchase intent) is its own increasingly important sub-topic, connecting to Part 32.4's greenwashing-perception logic — audiences who feel a partnership disclosure was hidden or unclear respond with the same trust erosion documented for perceived-inauthentic sustainability claims, while audiences who perceive a partnership as genuinely authentic and well-disclosed show materially better response, making disclosure clarity itself a testable creative/execution variable rather than a pure compliance afterthought.

### 34.5 Worked example — diagnosing an underperforming influencer campaign

*A brand runs a campaign with 8 mid-tier influencers, achieving strong impressions and engagement-rate numbers, but tracked-link conversion is well below the brand's typical paid-media benchmark.*



**Model answer.** Strong engagement with weak conversion is a specific, diagnosable pattern (echoing Part 2.5's "diagnose before prescribing" discipline) — before concluding influencer marketing simply "doesn't work" for this brand, check: audience-brand fit (Part 34.2 — high engagement can reflect a creator's general audience appeal rather than genuine interest in this specific brand/category, especially in categories with low natural fit to the creator's usual content); whether the call-to-action and tracked link were prominent and frictionless (a common, purely executional failure point distinct from audience-fit); and whether the influencer content itself credibly conveyed the product's actual value proposition, or was engaging content that happened to feature the product without making a clear case for it — three distinct possible root causes requiring three different fixes (better influencer selection, better campaign execution mechanics, or better creative briefing), exactly the kind of specific diagnosis this handbook's research process (Part 1.4) is built to produce rather than a vague "it underperformed" conclusion.

## PART 35 — VOICE OF CUSTOMER (VoC) PROGRAMS

### 35.1 VoC as continuous listening, distinct from a project-based study

Everything covered so far in this handbook is largely project-based — a discrete study with a start and end date. **Voice of Customer (VoC)** programs are different in kind: an always-on system continuously capturing customer feedback across multiple touchpoints (post-purchase surveys, support-ticket sentiment, app-store/product reviews, social mentions), feeding a live dashboard rather than a periodic report — closer to Part 24.4's continuous-research model, but specifically structured around unsolicited and lightly-solicited feedback rather than commissioned survey waves.

### 35.2 The core VoC data sources and what each captures

- **Transactional surveys** (a short survey triggered immediately after a specific interaction — a purchase, a support call) capture in-the-moment sentiment tied to a specific, identifiable touchpoint, with response rates and recency advantages a periodic tracker (Part 10.4) can't match.
- **Support-ticket and call-transcript analysis:** applying the qualitative-coding discipline from Part 25 at much larger scale (often software-assisted, using text-analytics/NLP to auto-code recurring themes across thousands of tickets) surfaces recurring pain points customers are actively contacting the company about — a strong, high-volume signal of what's currently broken.
- **Review and social-mention mining:** unsolicited, naturally-occurring feedback (Part 13.1's social listening, applied specifically to product/service feedback rather than brand buzz generally) — valuable precisely because it's unprompted, though skewed toward more extreme (very satisfied or very dissatisfied) customers who are more likely to leave a review unprompted, a real bias to disclose when using this source.

### 35.3 Closing the loop — VoC's defining operational requirement

What distinguishes a genuine VoC program from simply collecting feedback is **closing the loop**: routing individual pieces of negative feedback to the relevant team for direct follow-up with that specific customer (a service-recovery function, distinct from the aggregate-insight function most other research in this handbook serves), while separately aggregating the same feedback for trend analysis and product/process improvement. A program that only does the aggregate-trend side misses the immediate customer-recovery opportunity; a program that only does individual follow-up misses the systemic pattern-detection value — a mature VoC program does both simultaneously from the same underlying data.

### 35.4 Worked example — triangulating a VoC signal with a formal study

*A VoC dashboard shows a rising volume of support tickets mentioning "confusing" or "can't find" language related to a recently redesigned checkout flow, alongside a modest but real dip in the post-purchase transactional-survey satisfaction score for the same period.*

**Model answer.** VoC data is well-suited to *detecting* a problem quickly (the ticket-volume spike and survey-score dip both surfaced within days of the redesign, far faster than a periodic tracking-study wave would have caught it) but is less well-suited to fully *diagnosing* the specific usability failure driving it — the appropriate next step is exactly the same escalation logic as Part 23.2's financial-services activation-funnel example: pull the specific behavioural funnel data for the new checkout flow, then run a small set of moderated usability sessions (Part 4.3/13.2) on the flagged step to pinpoint the precise design failure, using

VoC as the early-warning trigger and a targeted follow-up study as the diagnostic tool — the two research modes working together rather than either alone being sufficient.

---

## PART 36 — RETAIL MEDIA NETWORKS & E-COMMERCE ADVERTISING RESEARCH

### 36.1 A genuinely new advertising channel requiring its own measurement approach

**Retail media networks** — advertising sold directly by e-commerce platforms and retailers (sponsored product placements, on-site display ads, search-term advertising within a retailer's own app/site) — have grown into a major ad-spend category distinct from both traditional media (Part 20's MMM channels) and open-web digital advertising (Part 13.2), warranting its own research treatment given how differently it's measured and bought.

### 36.2 Why retail media measurement is unusually direct — and where that directness can mislead

Retail media's defining advantage is closed-loop measurement: because the ad and the purchase both happen within the same retailer platform, attribution from ad exposure to actual purchase is far more direct than in traditional or open-web advertising, where the Part 20 MMM/attribution modelling machinery exists specifically to bridge that gap indirectly. This directness is genuinely valuable but has a specific measurement trap: retail-media ROAS (Return on Ad Spend) figures reported by the platform itself often include **organic/incremental confusion** — some of the "attributed" sales would have happened anyway (the shopper was already searching for that product or brand) even without the sponsored placement, meaning platform-reported ROAS can overstate true incrementality unless independently validated, echoing Part 20.4's MMM-vs-attribution-model reconciliation theme (different measurement methods for the same spend can disagree, and reconciling the gap rather than picking whichever number looks best is the rigorous approach).

### 36.3 Share of search — the retail-media-specific visibility metric

**Share of search** (introduced briefly in Part 14.5's e-commerce shopper-research discussion) becomes a central, ongoing tracking metric specifically for retail media strategy: a brand's visibility in on-platform search results for its own branded terms and relevant category terms, trackable over time the same way share of shelf is tracked for physical retail (Part 14.3) — a brand that's losing share of search to competitors on its own brand-name searches (competitors bidding on and appearing for a brand's own name) is a specific, actionable, and increasingly common competitive-intelligence finding in this channel.

### 36.4 Worked example — reconciling platform-reported ROAS with true incrementality

*A brand's retail-media dashboard reports a sponsored-product campaign ROAS of 8:1. A holdout test (a matched sample of the platform's users randomly withheld from seeing the ads, compared against an exposed group) shows the true incremental sales lift from the campaign is closer to 3:1.*

**Model answer.** The gap between the platform-reported 8:1 and the holdout-measured 3:1 reflects exactly the organic/incremental confusion from Part 36.2 — a meaningful share of the "attributed" 8:1 sales would have occurred anyway from shoppers already intending to buy the product, with the sponsored placement claiming credit for a purchase it didn't actually cause. The **holdout/incrementality test** (a genuine causal-research design, Part 2.1, applied specifically to a retail-media platform's own user base) is the

rigorous way to measure true incremental ROAS rather than trusting the platform's own attribution logic at face value — a platform has some inherent incentive for its own reported ROAS to look favourable, making independent incrementality testing a genuinely important check, not just an academic nicety, whenever retail-media budget decisions are being made based on the platform's self-reported numbers.

---



## PART 37 — MOBILE APP UX RESEARCH & IN-APP ANALYTICS

### 37.1 Why mobile app research deserves its own treatment, especially for financial-services apps

For a broking/trading or wealth-management app specifically (Part 23's financial-services research territory), the mobile app is the product — the research methods covering it directly determine whether real usability barriers to activation, trading, and retention get caught before they cost the business meaningfully. This Part covers the mobile-specific extensions to the digital-research toolkit from Part 13.2.

### 37.2 In-app micro-surveys and contextual feedback capture

Short, contextually-triggered surveys (a single question appearing right after a specific in-app action — completing a trade, viewing a research report, hitting an error) capture feedback at the precise moment of an experience, avoiding the recall-decay problem of a general post-hoc survey (Part 5's broader survey-design principles apply, but timing/triggering is the mobile-specific addition) — a **Net Promoter Score (Part 5.3's NPS definition) triggered after a specific milestone** (first successful trade, a support interaction) gives a more diagnostic, moment-specific read than a single quarterly NPS survey blending all possible experiences into one number.

### 37.3 App-store review mining as a structured VoC source

Extending Part 35.2's VoC source list with a mobile-specific channel: app-store reviews (Play Store/App Store) are a rich, unsolicited, publicly-visible feedback source, but carry the same extreme-respondent skew flagged for other review-mining sources (Part 35.2) — additionally, app-store reviews are specifically valuable for **competitive benchmarking** (reading competitor apps' own reviews for recurring complaints reveals gaps a competitive product analysis, Part 7.3, might miss from feature comparison alone) and for catching **version-specific regressions** (a spike in negative reviews tightly clustered right after a specific app update is a strong, fast signal that update introduced a real usability or stability problem).

### 37.4 Mobile-specific usability testing considerations

Usability testing (Part 4.3/13.2) on a mobile app has specific considerations beyond desktop/web usability testing: screen size constraints mean information architecture and navigation patterns that work on desktop can fail on mobile even with identical underlying functionality; testing across both iOS and Android is often necessary given real platform-specific interaction-pattern differences users have learned to expect; and testing on the participant's **own device** (rather than a lab-provided device) surfaces real-world friction — an older device, a specific screen size, existing app clutter — that a controlled lab environment with a fresh, high-spec device can mask entirely.

### 37.5 Worked example — diagnosing a mobile-specific activation drop

*A broking app's activation funnel (Part 23.4's financial-services activation-funnel logic) shows materially lower completion on Android than iOS at the KYC-document-upload step specifically.*

**Model answer.** A platform-specific drop-off at one specific step is a strong, localised diagnostic signal (echoing Part 2.5's diagnose-before-prescribing discipline) — before assuming a general UX problem, the Android-specific pattern suggests checking platform-specific technical causes first: camera/file-picker

permission-handling differences between Android versions and manufacturers (a well-documented source of Android-specific friction, given Android's more fragmented device/OS-version landscape versus iOS's more controlled ecosystem), upload file-size or format handling differences, or Android-specific device performance issues during the upload step. Moderated usability sessions specifically on a sample of Android devices (Part 37.4's own-device testing principle) at this exact step would be the appropriate next research step — a materially more targeted diagnostic path than a generic "improve KYC UX" usability study that doesn't isolate the platform-specific nature of the problem from the outset.

---

## PART 38 — SURVEY PROGRAMMING & DATA QUALITY IN PRACTICE

### 38.1 The operational step between questionnaire design and fielding

Part 5's survey-design principles (question wording, ordering, scale choice) describe *what* the questionnaire should contain. This Part covers the practical step in between design and fielding: **programming** the questionnaire into a live survey tool and **testing** it thoroughly before real respondents ever see it — a genuinely distinct, often underestimated source of research errors when skipped or rushed.

### 38.2 Skip logic and routing — where good questionnaire design meets implementation risk

A well-designed questionnaire's skip logic (Part 5.1's screener/routing principles) must be correctly implemented in the survey platform — a common, costly implementation error is a broken skip pattern that either shows respondents irrelevant questions (a straight annoyance and drop-off risk) or, worse, silently skips a question that should have been asked, producing missing data that isn't discovered until analysis, by which point the field is closed and the gap can't be fixed. **Full-path testing** — walking through every possible combination of screener/routing answers, not just the "happy path" a programmer naturally tests first — is the standard discipline for catching this before launch.

### 38.3 Soft-launch / pilot fielding as the last quality gate before full launch

Extending Part 5.1's pilot-testing principle to the programming stage specifically: a **soft launch** — releasing the live, programmed survey to a small fraction (e.g. 5-10%) of the planned sample first, then pausing to review real response data (completion rates, time-to-complete, open-end response quality, any unexpected drop-off points) before releasing to the full sample — catches both wording problems (Part 5.1's original pilot purpose) and programming/implementation problems (broken logic, a scale that renders incorrectly on mobile) that only surface once real respondents interact with the live instrument, not just an internal team's own test runs.

### 38.4 Data quality checks that should be built into the programming itself, not applied only after the fact

Beyond the post-hoc data-quality checks from Part 5.8 (speeding, straight-lining, attention checks), several checks can be built directly into the survey's live logic: **minimum time-per-page thresholds** that flag or soft-block implausibly fast completion in real time; **consistency checks across related questions** (e.g. a respondent claiming to be a heavy category user in one question but "never heard of" a leading brand in that category later — Part 5.8's logical-consistency-trap example, here caught live rather than only in post-hoc cleaning); and **straight-line detection on grid questions** flagged for review or exclusion as the response comes in, rather than discovered only during analysis after the field has already closed and the opportunity to re-contact or replace a bad respondent has passed.

### 38.5 Worked example — a programming error discovered too late

*A tracking study's Wave 4 data shows an unusual jump in "don't know" responses to a specific brand-awareness question compared to Waves 1-3, only discovered during analysis after the field has closed.*

**Model answer.** Before concluding this reflects a genuine shift in the market (e.g. declining brand salience), the first step is ruling out a programming/implementation explanation (Part 38.2) — checking whether Wave 4's questionnaire was reprogrammed or migrated to a new survey platform between waves, whether the specific question's response options or randomisation setting changed, or whether a skip-logic error caused some respondents to see the question in a different context than in prior waves. If a soft-launch check (Part 38.3) wasn't run for Wave 4's specific implementation, this exact kind of undetected-until-analysis error is the predictable consequence — reinforcing why the soft-launch discipline matters most precisely for tracking studies, where comparability across waves (Part 10.4's tracker-consistency principle) depends on the instrument being implemented identically each time, not just designed identically on paper.

---

## PART 39 — GEOSPATIAL & LOCATION-BASED RESEARCH

### 39.1 Research questions that are fundamentally about location, not just demographics

Physical-network businesses — bank branches, retail stores, QSR outlets — face a distinct category of research question this handbook hasn't yet directly addressed: not just *what* customers want, but *where* to place physical points of presence, and how a location's specific catchment area shapes demand for it.

### 39.2 Trade-area analysis — defining and sizing a location's realistic catchment

A **trade area** is the geographic zone from which a specific location realistically draws the large majority of its customers, typically mapped using a combination of drive-time or distance rings (a simple, if imperfect, starting method) and, more rigorously, actual customer address/transaction data from existing nearby locations (a **gravity model** approach, where the pull of a location on a given area is weighted by distance and the location's own size/attractiveness relative to competing locations) — directly analogous to the TAM/SAM/SOM narrowing logic from Part 7.1, but applied to a specific geographic footprint rather than a demographic/product segment.

### 39.3 Site-selection research — combining trade-area sizing with competitive and demographic overlay

Site-selection research for a new physical location combines: trade-area demand sizing (Part 39.2) using demographic and spending-pattern data within the realistic catchment; a competitive-density overlay (how many existing competing locations already serve that trade area, and how saturated it already is — directly extending Part 7.3's competitive-intelligence frameworks to a spatial dimension); and increasingly, footfall/mobility data (aggregated, anonymised location data showing actual foot traffic patterns near candidate sites) as a revealed-behaviour cross-check against the demographic/gravity-model estimate, mirroring this handbook's recurring stated-vs-revealed-preference theme (Part 22.1) applied to where people actually go, not just where a model predicts they should.

### 39.4 Cannibalisation analysis — a location-specific version of the incrementality question

Before opening a new location near existing ones (same brand), **cannibalisation analysis** estimates how much of the new location's projected sales would be diverted from nearby existing locations rather than representing genuinely incremental demand — the geospatial equivalent of Part 19.2's assortment-incrementality question and Part 36.2's retail-media organic/incremental confusion, all variations on the same underlying discipline: before claiming a new investment (a product, an ad campaign, a store) drives incremental value, rigorously check how much of its apparent contribution is actually just moving existing demand around rather than growing it.

### 39.5 Worked example — evaluating a new bank-branch location

*A bank is evaluating a new branch location. The trade-area analysis (Part 39.2) shows a strong addressable population and income profile, but the site sits 1.2km from an existing branch of the same bank, in a trade area with meaningful overlap between the two catchments.*

**Model answer.** A strong standalone trade-area profile is necessary but not sufficient (Part 39.4's cannibalisation discipline) — the analysis must explicitly estimate what share of the new branch's projected business would be diverted from the existing nearby branch's current customer base within the overlapping catchment zone, versus genuinely new-to-bank or previously-underserved demand the new location would capture. If cannibalisation is estimated at, say, 40% of the new branch's projected volume, the real incremental business case is materially weaker than the standalone trade-area number alone would suggest — the same "don't claim credit for demand that would have existed anyway" discipline that runs through the retail-media (Part 36.2), assortment (Part 19.2), and CLV/retention material elsewhere in this handbook, here applied to a physical network-expansion decision specifically relevant to a financial-services branch network.

---

## PART 40 — RETAIL INVESTOR & TRADING-APP RESEARCH

### 40.1 Why retail-broking research is a distinct sub-discipline, not generic fintech UX research

Discount and full-service brokers, and the trading/investing apps built around demat accounts, sit at the intersection of financial-services research (Part 23) and mobile-app UX research (Part 37) — but retail-investor research has its own recurring, distinct research questions: onboarding/KYC funnel drop-off, trading-frequency-based segmentation, and the specific trust/risk-perception dynamics of putting real money at risk through an app, which generic e-commerce or social-app UX research frameworks don't fully capture.

### 40.2 KYC and account-activation funnel research — a regulated, multi-step drop-off problem

Opening a demat/trading account involves a regulated KYC (Know Your Customer) sequence — PAN/Aadhaar verification, bank-account linkage, a video/in-person verification step, and a segregated-nominee declaration — each an independent point of drop-off; funnel research here (extending Part 37.5's Android-vs-iOS activation-drop diagnostic) must distinguish drop-off caused by the app's own UX (a confusing screen, an unclear error message) from drop-off inherent to the regulatory step itself (a document upload that simply takes time, a step users abandon to go find a physical document) — the fix differs completely between the two, and conflating them wastes product effort on UX changes that can't move a regulatory-friction bottleneck.

### 40.3 Trading-frequency segmentation — a behavioural, not demographic, cut

Retail investors segment far more usefully by revealed trading behaviour than by demographics alone: **buy-and-hold investors** (few trades/year, primarily mutual-fund/SIP-oriented), **active investors** (regular but measured equity trading, holding periods of weeks to months), and **high-frequency/intraday traders** (many trades/day, heavily F&O-weighted) — each segment has fundamentally different feature needs (a buy-and-hold investor cares about SIP automation and portfolio-tracking clarity; an intraday trader cares about order-execution speed and real-time margin visibility), meaning a single "investor persona" for product research purposes actively obscures the different needs each segment has, extending this handbook's general a-priori-vs-behavioural segmentation caution (Part 8) to this specific, highly bimodal user base.

### 40.4 Trust and risk-perception research — the layer generic app research misses

Because trading apps involve real financial risk (not just time or attention, as with most consumer apps), research specifically on **trust signals** (visible regulatory registration/SEBI references, clear disclosure of brokerage/charges before order placement, transparent margin/leverage warnings) and **risk-perception calibration** (whether users, especially new F&O traders, correctly understand the leveraged risk they're taking on) sits alongside standard usability research as a distinct, higher-stakes research stream — closely related to, but methodologically distinct from, this handbook's investor-confidence/sentiment-index material (Part 23.4), since it's about product-level risk comprehension specifically, not market-wide sentiment.

## 40.5 Worked example — diagnosing why new F&O traders churn within 30 days

*A discount broker's data shows a large share of first-time F&O (options) traders stop trading entirely within 30 days of their first options trade, well before typical intraday/active-trader retention benchmarks. Product wants to know whether this is an app-UX problem or something else.*

**Model answer.** Before assuming a UX fix, the research should separate three distinct hypotheses (Part 40.1's discipline against treating this as generic app-UX research): (1) a genuine UX/execution-friction problem specific to options order flow (multi-leg order placement being confusing, unclear margin requirements shown at order time); (2) a risk-comprehension problem (Part 40.4) — new F&O traders taking losses larger than they understood the leveraged risk to be, and exiting the product entirely out of risk-aversion rather than dissatisfaction with the app itself; or (3) a natural, non-problematic pattern where a share of first-time options traders were always going to be one-off "curiosity" traders regardless of app quality. Distinguishing these requires combining behavioural data (loss size on the trades preceding churn, whether churned users' losses exceeded a typical "first options trade" size) with direct qualitative outreach to a churned-user sample — a UX fix only helps if hypothesis (1) is actually the dominant driver, while hypotheses (2) and (3) call for entirely different responses (better pre-trade risk education, or accepting the drop-off as a natural segment boundary rather than a problem to "fix").

---



## PART 41 — FINANCIAL LITERACY & INVESTOR EDUCATION RESEARCH

### 41.1 A distinct research stream from product-level trust/risk research

Part 40.4 covered trust signals and risk-perception calibration *within a specific trading product*. **Financial literacy and investor education research** is a related but broader stream: measuring baseline financial/market knowledge across a population or user base, and evaluating whether educational content or campaigns (broker-run investor education, exchange/regulator-run awareness programs) actually move comprehension and behaviour — a distinct research question from whether one specific app's risk disclosures are well understood.

### 41.2 Financial literacy assessment — standardised instruments over ad-hoc quizzes

Credible financial literacy research uses **standardised, validated assessment instruments** (batteries of questions covering compound interest, inflation, diversification, and risk-return tradeoffs, in the tradition of widely-used academic financial-literacy scales) rather than an ad-hoc internal quiz — standardisation matters because it allows benchmarking a given population's literacy level against published national/regional norms, and because ad-hoc questions risk conflating "financially literate" with "familiar with this specific company's product jargon," an important distinction when the research output will be compared across studies or over time.

### 41.3 Evaluating investor-education content effectiveness — a pre/post, not just satisfaction, design

Evaluating whether an investor-education module (an in-app tutorial, a webinar series, a regulator-mandated awareness module before enabling F&O trading) actually works requires a **pre/post comprehension design** — measuring the same knowledge questions before and after exposure to isolate genuine comprehension gain — rather than relying only on post-content satisfaction ratings ("how helpful was this?"), which measure perceived helpfulness, not actual knowledge change, and are prone to the same social-desirability and self-report-inflation biases covered elsewhere in this handbook's survey-design material (Part 4).

### 41.4 Behavioural follow-through — the gap between comprehension and applied behaviour

Even a well-designed pre/post study that shows genuine comprehension gain does not by itself establish that the education changed real trading/investing behaviour — a further, harder research question is whether education-module completers show different downstream behaviour (more diversified position sizing, lower incidence of the specific mistakes the module warned against) than a matched comparison group of non-completers, extending this handbook's stated-vs-revealed and intent-vs-behaviour cautions (Part 22.1) into the specific claim "education changed behaviour," which needs behavioural evidence, not just comprehension-test evidence, to support.

## 41.5 Worked example — evaluating a mandatory pre-F&O-trading education module

*A broker introduces a mandatory short education module before a user's account is enabled for F&O (options) trading. Post-module satisfaction scores are high (4.3/5), and product wants to cite this as evidence the module is "working."*

**Model answer.** High satisfaction scores (Part 41.3) are not evidence the module achieved its actual goal — improving risk comprehension and reducing preventable mistakes — since satisfaction measures how pleasant or clear the module felt, not what was retained or whether behaviour changed. The credible evaluation design pairs a pre/post knowledge assessment (using validated risk/leverage-comprehension questions, Part 41.2) around the module to establish genuine comprehension gain, and — ideally — a behavioural follow-up comparing module completers against a matched comparison group on downstream outcomes such as early account closure or outsized first-month losses (Part 41.4, tying directly back to Part 40.5's F&O-churn diagnostic elsewhere in this handbook). Citing satisfaction alone as proof of effectiveness risks the module being kept or scaled based on a metric that has no established link to the outcome it actually needs to move.

---

## PART 42 — ROBO-ADVISORY & ALGORITHMIC RECOMMENDATION RESEARCH

### 42.1 A distinct research problem — trusting a recommendation, not just an app

Robo-advisory platforms and algorithm-driven portfolio recommendations raise a research question distinct from both Part 40's broking-app trust/risk research and Part 41's financial-literacy research: whether users **trust and act on** a recommendation they did not arrive at themselves — a distinct behavioural and comprehension challenge from trusting an app's execution/security, or possessing baseline market knowledge, since a user can be financially literate and still distrust or override an algorithm's specific recommendation.

### 42.2 Risk-profiling questionnaire validation — a measurement-quality problem, not just a UX one

The risk-profiling questionnaire that feeds a robo-advisory recommendation (typically asking about risk tolerance, time horizon, and investment goals) is itself a measurement instrument, and needs the same validation discipline as any survey scale in this handbook (Part 5's scale-construction material): does it produce stable, repeatable risk-profile classifications for the same underlying person (test-retest reliability), and does the resulting classification actually predict the behaviour it's meant to predict (does a "moderate risk" classification correlate with how a user actually responds to a real portfolio drawdown)? A questionnaire that's intuitive and quick to complete (a UX win) but produces unstable or behaviourally-unvalidated risk classifications is a measurement failure hiding behind a good user experience.

### 42.3 Explainability research — how much "why" users need before they'll act

**Explainability research** tests how much justification a recommendation needs to include before users are willing to act on it — ranging from a bare recommendation with no rationale, to a brief one-line reason ("recommended based on your moderate risk profile and 10-year horizon"), to a full breakdown of the underlying logic — with the research question being where the marginal explanation stops adding trust/action and starts adding cognitive friction that reduces completion, a direct extension of this handbook's progressive-disclosure principle (Part 33) applied specifically to algorithmic-recommendation interfaces rather than general survey or content design.

### 42.4 The override/disagreement research question — what happens when users reject the recommendation

A meaningfully large share of users typically deviate from or override a robo-advisory recommendation at least once — researching *why* (a mismatch between the stated risk profile and the user's felt risk tolerance in the moment, a specific recommended holding the user distrusts for reasons the questionnaire didn't capture, herd behaviour toward a trending stock the algorithm didn't recommend) is directly diagnostic of where the risk-profiling instrument (Part 42.2) or the explanation layer (Part 42.3) is failing to fully capture or communicate what the user actually needs, making override analysis a genuinely useful diagnostic data source in its own right, not just a metric to minimise.

## 42.5 Worked example — diagnosing a spike in recommendation overrides after a market drawdown

*A robo-advisory platform sees overrides (users manually adjusting away from the algorithm's recommended allocation) spike sharply during a market drawdown, concentrated among users classified as "moderate risk" who had shown no override behaviour previously.*

**Model answer.** This pattern (Part 42.4's diagnostic framing) points toward a gap between the *stated* risk tolerance the questionnaire captured (Part 42.2) and the user's *actual, felt* risk tolerance when a real drawdown is happening — a classic instance of this handbook's broader stated-vs-revealed-preference theme (Part 22.1), here specifically exposed by market stress rather than by a hypothetical survey question. The research response should not default to blaming the explanation layer (Part 42.3) without first checking the risk-profiling questionnaire's behavioural predictive validity for this specific cohort — if "moderate risk" users are overriding en masse during their first real drawdown, this is evidence the questionnaire's classification is not capturing genuine risk tolerance for a meaningful subset of moderate-risk users, a measurement-validity finding with direct implications for how the questionnaire itself should be redesigned or supplemented (e.g. a scenario-based drawdown-simulation question, shown to more accurately reveal felt risk tolerance than an abstract self-rated question), not merely a UX or explanation-copy fix.

---

## PART 43 — ALTERNATIVE DATA AS A MARKET RESEARCH INPUT

### 43.1 Alternative data as a distinct evidence source, not a replacement for primary research

**Alternative data** — non-traditional data sources like satellite imagery (retail parking-lot traffic, agricultural yield estimates, shipping/port activity), aggregated app-download and usage rankings, web-traffic estimates, and job-posting volumes — has become a genuine input into market and competitive research, particularly for research supporting equity/investment decisions. It's a distinct evidence layer from this handbook's core primary methods (surveys, interviews, panels, Part 2-5 material): alternative data captures *revealed, aggregate, often near-real-time* activity signals, at the cost of being noisier, harder to interpret causally, and requiring careful cleaning before it's usable evidence at all.

### 43.2 The core use case — a faster, higher-frequency read on a trend primary research would only catch later

Alternative data's main value proposition is **speed and frequency**: app-download rankings can show a competitor's launch traction within days, while a syndicated tracking survey (Part 10) might not field and report for weeks; satellite-derived retail-traffic estimates can give a same-quarter read on footfall trends well before a company's own quarterly results are published. This makes alternative data particularly useful as an **early-warning or directional cross-check** ahead of a slower primary study, rather than a replacement for the deeper causal understanding (the *why*) that primary qualitative/survey research still uniquely provides.

### 43.3 Validation discipline — alternative data needs ground-truthing before it's trusted

A recurring, serious pitfall with alternative data is treating a noisy, indirect proxy as if it were a precise, validated measure — app-download rank movements, for instance, can be driven by a promotional spike, an app-store algorithm change, or a genuine usage trend, and are not inherently distinguishable from the raw numbers alone. The correct discipline (extending this handbook's triangulation theme, Part 1) is **ground-truthing**: checking a new alternative-data source against a known, already-validated outcome (e.g. does app-download rank actually correlate with the company's own subsequently-reported user-growth numbers over several past quarters?) before treating it as a leading indicator worth acting on for a new, unvalidated question.

### 43.4 Where alternative data fits in a research team's toolkit relative to primary methods

Alternative data is best positioned as a **complement that sharpens where to point primary research**, not a parallel, independent evidence stream: a sharp divergence between alternative-data signals (e.g. a competitor's app-download surge) and existing qualitative understanding of the market is itself a research trigger — the appropriate next step is a targeted qualitative or survey follow-up to understand *why* the alternative-data signal is moving, since alternative data alone can flag *that* something changed far faster than it can explain *why*, and the "why" is usually what a business decision actually needs.

### 43.5 Worked example — a competitor's app-download surge with no other visible signal

*App-download-ranking data shows a fintech competitor's trading app jumped from #45 to #8 in its category over three weeks. No corresponding news coverage, funding announcement, or major feature launch is visible from a standard secondary-research scan.*

**Model answer.** Before treating this as a validated signal of a genuine competitive shift, the ground-truthing discipline (Part 43.3) applies first — checking whether this specific data provider's rank movements have historically correlated with real user-growth outcomes for comparable apps, and ruling out a promotional/paid-acquisition spike or an app-store-side algorithm change as the driver, since a rank jump with no corresponding news is at least as consistent with a paid User Acquisition campaign (a real but temporary and less strategically significant driver) as with organic, sustained demand growth. If the alternative-data signal survives this check as plausible and significant, the appropriate next step (Part 43.4) is not to report it as a finding on its own, but to trigger a targeted piece of primary research — a quick competitive mystery-shop of the app itself, a scan of app-store reviews (Part 37's app-store review mining) for what specifically new users are reacting to — to establish the *why* the fast-but-shallow alternative-data signal cannot answer by itself.

---

## PART 44 — ADVISOR/DISTRIBUTION-CHANNEL RESEARCH FOR ASSET & WEALTH MANAGEMENT

### 44.1 Why the distribution channel is a distinct research subject from the end investor

Part 23 covered researching financial-product end investors directly. Asset and wealth management products, however, are frequently sold *through* an intermediary channel — independent financial advisors (IFAs), mutual fund distributors (MFDs), relationship managers (RMs) at banks/wealth platforms — meaning the **advisor/distributor themselves** is a distinct research subject with their own needs, incentives, and decision drivers, separate from and sitting between the asset manager and the end investor; researching only the end investor while ignoring the channel that actually influences most of their product choices leaves a significant blind spot for any AMC or wealth platform's research function.

### 44.2 Advisor satisfaction and needs research — what drives which products an advisor recommends

Advisor-facing research typically covers: **product usability** for the advisor (how easy is it to explain/pitch a fund's strategy to a client, how competitive are the disclosed costs/commissions versus alternatives), **platform/service quality** (transaction turnaround time, ease of the onboarding/KYC process for their clients, quality of wholesaler/relationship support from the AMC), and **content/enablement needs** (what sales collateral, market commentary, or client-facing materials would make the advisor more likely to recommend a given product) — each a distinct driver of advisor recommendation behaviour that end-investor research alone would never surface, since the advisor's own operational friction and incentive structure sit between the product's actual merits and what gets recommended to a client.

### 44.3 RM productivity and needs research — an internal-channel variant with different stakes

For wealth platforms and private banks with an in-house relationship manager (RM) model rather than external distributors, a closely related but internally-focused research stream studies **RM productivity drivers**: which tools, data, or workflow changes measurably reduce the time an RM spends on low-value administrative tasks versus client-facing advisory time, and what client-portfolio-review content or talking points RMs report the most demand for from clients — methodologically this overlaps with general employee/internal research, but the specific content (investment-product knowledge gaps, market-commentary needs) ties it directly to this handbook's financial-services research material (Part 23) rather than to generic internal-culture research.

### 44.4 The double-blind-spot risk — researching only the channel, or only the end investor, in isolation

The core methodological risk in this space is researching only one side: **channel-only research** risks over-indexing on what makes a product easy to *sell* rather than what genuinely serves the end investor's interests (a real conflict-of-interest risk this handbook flags explicitly, not just a research-design footnote, given the regulatory and reputational stakes in financial services specifically), while **end-investor-only research** (Part 23) misses the very real influence the advisor/RM relationship has on what actually gets recommended and purchased — the credible research program deliberately triangulates both perspectives on any given

product or feature question, explicitly checking where advisor and end-investor priorities align versus diverge.

### 44.5 Worked example — reconciling divergent advisor and end-investor feedback on a new fund

*Advisor research on a newly-launched thematic equity fund shows strong positive feedback — advisors find it easy to pitch given its clear, narrative-driven theme. End-investor research on the same fund shows advisors' clients frequently misunderstand the fund's actual risk/volatility profile, expecting it to behave like a broader diversified fund.*

**Model answer.** This is a direct instance of the double-blind-spot risk (Part 44.4) surfacing as a genuine divergence, not a research contradiction to be averaged away — the advisor-side finding (easy to pitch, well-received) and the end-investor-side finding (risk-profile misunderstanding) are both valid and both actionable, but point to different fixes. The advisor-facing "easy to pitch" quality is actually part of the problem if the narrative that makes it easy to pitch is the same narrative causing end-investor risk misunderstanding — the appropriate response is not to conclude the fund is broadly successful based on advisor sentiment alone, but to specifically test whether revised advisor-facing materials that foreground the risk/volatility profile alongside the thematic narrative change what advisors actually communicate to clients, then re-run the end-investor comprehension check (Part 41's pre/post comprehension-design discipline, applied here to advisor-mediated communication rather than direct-to-investor education) to confirm the fix actually closes the understanding gap rather than just changing the advisor-research score.

---



## PART 45 — IPO GREY-MARKET & PRE-LISTING INVESTOR SENTIMENT RESEARCH

### 45.1 A distinct research window — the pre-listing period as its own research object

Part 23's pre-IPO messaging testing covered how a company's own equity story is perceived before listing. This Part covers a related but distinct research window: the **subscription and grey-market period** between an IPO's price-band announcement and its listing day, during which several genuinely new, time-bounded research questions become answerable that don't exist for an already-listed stock — investor demand across categories, informal secondary-market pricing, and how these relate (or don't) to actual post-listing performance.

### 45.2 Reading category-wise subscription data as a segmented demand signal

IPO subscription data is reported by investor category — **Retail Individual Investors (RII)**, **Non-Institutional Investors (NII, primarily HNIs)**, and **Qualified Institutional Buyers (QIB)** — and reading these three categories *separately* rather than only the headline overall-subscription number is directly analogous to this handbook's segmentation discipline (Part 8): heavy QIB oversubscription with weak retail interest suggests institutional conviction not yet shared by (or not yet visible to) individual investors, while the reverse pattern (retail-heavy, QIB-light) can reflect retail enthusiasm around a well-marketed consumer brand that professional investors are pricing more cautiously — each pattern has different implications for likely listing-day demand and post-listing volatility.

### 45.3 Grey-market premium (GMP) — what it is, and its real, limited evidentiary value

The **grey-market premium (GMP)** — an informal, unregulated indication of what an IPO's shares might trade for above (or below) the issue price, based on unofficial pre-listing trading among a small network of market participants — is widely watched but should be treated with real caution as a research input: it reflects a thin, informal, low-participation market (not a regulated exchange), is not investable or hedgeable by most participants, and has a documented history of diverging sharply from actual listing-day price action, particularly for high-profile or heavily-retail-subscribed issues where sentiment can run ahead of fundamentals. GMP is best used as one directional, low-confidence input alongside category-wise subscription data (Part 45.2), never as a standalone predictor.

### 45.4 The listing-day-pop research question — validating (or debunking) GMP after the fact

A useful, recurring internal research exercise for a firm covering IPOs is systematically comparing GMP-implied listing expectations against **actual listing-day price action** across a sample of past IPOs, to build an evidence-based view (specific to the current market cycle, since this relationship shifts over time) of how reliable GMP has actually been recently — extending this handbook's general validation discipline (Part 43.3's ground-truthing principle, applied here to GMP specifically) rather than treating GMP's popular reputation as a substitute for checking its actual recent track record.

## 45.5 Worked example — reconciling divergent subscription and grey-market signals

*An IPO shows QIB subscription at 45x and NII at 60x, but RII subscription at only 0.8x (undersubscribed in the retail category), while the grey-market premium suggests a healthy double-digit percentage listing gain.*

**Model answer.** The stark divergence between overwhelming institutional/HNI demand (Part 45.2) and retail undersubscription is itself the most informative signal here — it suggests the issue's story or valuation may not be resonating with (or may not have reached) individual investors as effectively as it has with institutions who typically apply more rigorous, professional demand estimation, a genuinely different read than a uniformly strong subscription across all three categories would represent. The GMP's double-digit listing-gain suggestion (Part 45.3) should be weighted cautiously given its thin, informal-market basis and weighed against the specific composition of demand just identified, not treated as an independent confirming signal — a firm with a track record of validating GMP against actual outcomes (Part 45.4) would apply whatever discount its own historical accuracy analysis suggests, rather than taking the GMP figure at face value alongside the more rigorously-sourced subscription data.

---

## PART 46 — SUPPORT & SERVICE EXPERIENCE RESEARCH FOR BROKING & WEALTH PLATFORMS

### 46.1 A distinct research subject from onboarding UX and product trust

Part 40 covered onboarding/KYC funnel and in-product trust research. **Support and service experience research** is a related but distinct stream, focused specifically on what happens when a user's experience *breaks down* and they need help — a failed fund transfer, a disputed trade, a margin-call query, a KYC-document rejection — a moment of genuine stress for a financial-services user given the real money involved, making service-recovery quality disproportionately influential on retention relative to its low frequency of occurrence.

### 46.2 Customer Effort Score (CES) — a distinct metric from NPS, purpose-built for support interactions

While NPS (Part 6) measures overall relationship sentiment and CSAT measures satisfaction with a single interaction, **Customer Effort Score (CES)** — typically framed as "how easy was it to get your issue resolved" on a numeric scale — is purpose-built for support/service contexts specifically, based on the research finding that *effort reduction*, more than delight, is what drives loyalty after a service interaction: a user whose problem was resolved with minimal friction is more likely to stay than one who received an unusually delightful but effortful resolution. A TRA/market-research team supporting a broking platform should track CES specifically alongside NPS, since the two metrics measure genuinely different things and can diverge (high NPS overall, poor CES on support specifically, indicating a service-quality blind spot in an otherwise well-regarded product).

### 46.3 First Contact Resolution (FCR) — the operational metric with a real research-diagnostic use

**First Contact Resolution (FCR)** — whether an issue was fully resolved in the user's first support interaction, without requiring a follow-up contact — is both an operational efficiency metric and a genuine researchable driver of the CES/NPS relationship (Part 46.2): low FCR (issues requiring multiple contacts to resolve) is one of the most consistent, well-evidenced drivers of high customer effort and subsequent dissatisfaction, making FCR analysis (broken down by issue type — KYC queries, trade disputes, fund-transfer issues) a useful diagnostic layer beneath an aggregate CES score, similar in spirit to this handbook's driver-analysis-beneath-NPS discipline (Part 6.5).

### 46.4 Complaint-handling research and regulatory context — a heightened-stakes research area

Financial-services complaint handling operates under closer regulatory scrutiny than most consumer categories (formal complaint-resolution timelines, escalation-to-regulator/ombudsman pathways specific to broking and financial products) — research into complaint-handling experience should therefore track not just satisfaction with the resolution but whether resolution occurred within regulatory timelines and how effectively complaints were categorised/routed internally, since a pattern of complaints escalating to a regulator or ombudsman (rather than being resolved satisfactorily in-house) carries real regulatory and reputational risk beyond the individual affected customer, a genuinely higher-stakes framing than most categories' complaint research needs to carry.

## 46.5 Worked example — diagnosing a CES/NPS divergence on trade-dispute resolution

*A broking platform's overall NPS is strong (+42), but CES specifically on trade-dispute resolution (contesting an erroneous or unauthorised trade) is poor, and FCR data shows only 35% of trade disputes are resolved in the first contact, frequently requiring 3+ follow-ups.*

**Model answer.** The strong overall NPS alongside poor trade-dispute CES (Part 46.2's divergence pattern) indicates the broader product experience is well-regarded, but this specific, high-stakes support journey is a genuine blind spot the aggregate NPS number is masking — exactly why tracking CES separately, specifically for support interactions, matters. The low FCR (35%, requiring 3+ follow-ups) is the most direct diagnostic lead (Part 46.3) — investigating *why* trade disputes specifically require multiple contacts (unclear initial documentation requirements from the user, an internal routing/escalation process that passes the case between teams, an approval step requiring sign-off not available to the first-contact agent) would identify the specific process fix needed, rather than a generic "improve support" recommendation. Given trade disputes are also a complaint-handling category likely to escalate to a regulator/ombudsman if handled poorly (Part 46.4), this specific finding should be flagged as a priority fix beyond its individual CES/NPS impact — a service-quality gap in this particular journey carries disproportionate downside risk relative to most other support categories.

## PART 47 — EMPLOYEE-REVIEW MINING AS AN EQUITY-RESEARCH-ADJACENT SIGNAL

### 47.1 A distinct evidence source at the intersection of market research and equity research

Aggregated, anonymised employee reviews (Glassdoor-style platforms, where current/former employees rate and comment on their employer) have become a genuine research input for analysts covering a company's stock — a distinct evidence source from this handbook's other alternative-data material (Part 43's satellite/app-download signals), specific to organisational and workforce-health questions that neither traditional financial statements nor customer-facing research can directly answer, and squarely at the intersection of market research methodology and equity research use cases.

### 47.2 What employee-review data can plausibly signal — and the self-selection bias that limits it

Systematic patterns in employee-review data (a sustained decline in average rating, a rising share of reviews mentioning specific themes like "leadership" or "compensation," an unusual spike in review volume around a specific date) can plausibly precede or accompany real organisational stress — elevated attrition, a leadership change, a strategy shift — before it's visible in public disclosures. The critical limitation (extending this handbook's self-selection-bias caution, Part 4's B2B-research context, applied here) is that review platforms disproportionately attract employees with a strong reason to post — typically more negative or more recently-departed — meaning raw average ratings are not a representative sample of the full employee base and should never be read as if they were a properly-sampled survey.

### 47.3 Trend and theme analysis over point-in-time snapshots — the more defensible use case

Given the self-selection limitation (Part 47.2), the more defensible research application is **trend analysis over time and relative comparison to sector peers**, rather than treating any single company's absolute rating as a standalone finding — a sustained multi-quarter decline in review sentiment, or a rating meaningfully below a company's direct competitors on the same platform (holding the platform's own selection bias roughly constant across companies, a useful implicit control), is more defensible evidence of a genuine organisational trend than a single quarter's raw average, mirroring this handbook's general preference for tracked, comparative data over one-off point estimates (Part 10's tracker-consistency discipline).

### 47.4 Text-mining review content for specific, actionable themes beyond the star rating

Beyond the aggregate rating, applying this handbook's qualitative coding discipline (Part 25) to the free-text review content itself — coding recurring themes (leadership confidence, compensation competitiveness relative to peers, workload/burnout mentions, specific product or strategy concerns raised by employees closest to the work) — surfaces more specific, actionable signal than the star rating alone, similar in spirit to this handbook's broader theme of pairing a headline metric with driver analysis (Part 6.5's NPS-plus-verbatim discipline) rather than reporting the number in isolation.

## 47.5 Worked example — reading a sustained employee-sentiment decline ahead of a product launch

*A company's employee-review rating has declined steadily over four consecutive quarters, with a rising share of recent reviews specifically mentioning "unrealistic launch timelines" and "leadership out of touch with engineering realities," while the company's public messaging remains confidently on-schedule for a major upcoming product launch.*

**Model answer.** The sustained four-quarter trend (Part 47.3's trend-over-snapshot discipline) is more defensible evidence than a single quarter's dip, and the specific, recurring thematic content (Part 47.4) — "unrealistic launch timelines" and "leadership out of touch with engineering realities" — is more actionable and harder to dismiss as noise than the declining star rating alone, since it points to a specific, plausible mechanism (execution risk on the very launch the company is publicly confident about) rather than generic dissatisfaction. The correct research response is not to treat this as confirmed evidence the launch will slip (Part 47.2's self-selection caution still applies — the reviewers raising these concerns are not necessarily representative of the full engineering team's view), but to flag it as a specific, evidence-based question worth probing through other channels available to an equity/market researcher — management Q&A on the next earnings call, expert-network conversations with people close to the launch, or channel-check-style outreach (Part 4's B2B-research parallel) — treating the employee-review signal as a lead to investigate further, not a standalone conclusion to report as fact.

---

## PART 48 — REGULATORY & POLICY-CHANGE IMPACT PERCEPTION RESEARCH

### 48.1 A distinct research trigger — a proposed rule, not a product or market event

This handbook has covered market-entry research (Part 9), competitive intelligence (Part 3), and investor perception around company-specific events (Part 32's investor-perception material). **Regulatory and policy-change impact research** is triggered by something different: a regulator's proposed or newly-issued rule change (a SEBI circular altering margin requirements, disclosure norms, or product-eligibility rules) — a distinct research object because the "product" being researched is a rule, not a company's own choice, and the relevant "customers" (market participants, distributors, investors) are reacting to an externally-imposed change rather than evaluating something the research-commissioning firm designed.

### 48.2 Stakeholder-impact mapping — the necessary first step before any perception survey

Before fielding any perception research on a proposed regulatory change, the research team must first map **which distinct stakeholder groups are actually affected and how** — a margin-requirement change, for instance, affects retail F&O traders, brokers' capital and compliance operations, and clearing corporations each in genuinely different ways, meaning a single generic "how do you feel about this rule" question asked across an undifferentiated sample would blend fundamentally different impact experiences into one uninterpretable number — extending this handbook's segmentation discipline (Part 8) specifically to the regulatory-impact context, where the relevant segments are defined by *how the rule affects them operationally*, not by demographics or usage behaviour.

### 48.3 Comment-period and consultation-response research — a formal, evidence-gathering research use case

When a regulator opens a formal consultation or comment period on a proposed rule, firms and industry bodies often commission structured research specifically to inform their formal submitted response — quantifying likely operational impact (e.g. surveying member brokers on expected compliance-cost impact of a proposed change) and gathering documented member sentiment to support the industry body's formal position — a genuinely distinct research use case from typical commercial market research, since the research output feeds a regulatory submission process with its own formal evidentiary expectations, closer in spirit to this handbook's commercial due diligence material (Part 32's CDD framework) in its rigor requirements, applied to a policy-advocacy context instead of an M&A context.

### 48.4 Post-implementation impact tracking — measuring the actual effect, not just the anticipated one

Once a rule takes effect, a distinct follow-up research stream tracks the **actual** measured impact against what was anticipated pre-implementation (trading volume changes, participant behaviour shifts, unintended consequences not anticipated in the original stakeholder-impact mapping) — extending this handbook's pre/post evaluation-design discipline (Part 41.3) from a product-education context to a regulatory-change context, and a genuinely valuable research product in its own right since regulators themselves, industry bodies, and affected firms all have a stake in understanding whether a rule change

achieved its intended effect or produced meaningful unintended consequences worth flagging for a future revision.

### 48.5 Worked example — researching the impact of a proposed F&O margin-requirement increase

*SEBI proposes increasing margin requirements for retail F&O trading, open for public comment. A brokerage's research team is asked to assess likely impact ahead of submitting an industry-body comment.*

**Model answer.** Following the stakeholder-impact-mapping discipline (Part 48.2), the research should first separate at least three distinct affected groups — retail traders (who would face higher capital requirements to maintain existing position sizes, likely research question: at what margin-increase threshold do retail traders report they would meaningfully reduce trading activity or exit F&O entirely), the brokerage's own compliance/capital operations (a distinct, internal-facing impact-assessment question), and potentially smaller/regional brokers with a different capital base than a large national brokerage — each requiring a differently-designed research instrument rather than one blended survey. The output feeding the formal comment-period submission (Part 48.3) should quantify these impacts with the same evidentiary rigor this handbook applies elsewhere (documented methodology, representative sampling within each stakeholder group, clearly-labelled directional-vs-precise findings per this handbook's general precision-caution discipline, Part 1), since a weakly-evidenced submission carries less influence in a formal regulatory process than one demonstrating genuine, well-sampled stakeholder research — and the brokerage should flag internally that a post-implementation tracking study (Part 48.4) will be needed regardless of the rule's final form, to validate which anticipated impacts actually materialised.

---



## PART 49 — GIFT CITY, CROSS-BORDER & NRI INVESTOR RESEARCH

### 49.1 A distinct investor segment shaped by regulation and geography, not just wealth

Non-Resident Indian (NRI) investors, and the broader cross-border investor base increasingly served through India's GIFT City International Financial Services Centre (IFSC), constitute a genuinely distinct research segment — not primarily defined by income or risk tolerance (this handbook's usual segmentation axes, Part 8), but by a specific regulatory and geographic reality: a different applicable tax treatment, repatriation rules, and often a different regulatory investment route (e.g. NRI-specific demat/trading account structures) than a resident Indian investor, meaning research questions and product-fit assessments for this segment must be scoped around these structural differences from the outset, not treated as a wealth-tier variant of resident-investor research.

### 49.2 Why NRI investor research needs geography-specific sub-segmentation, not one blended NRI persona

A single "NRI investor" research persona would blend genuinely different sub-populations — NRIs in the Gulf/Middle East (often higher remittance volumes, different tax-treaty considerations, a large and well-established population with mature product awareness), NRIs in the US/UK/Canada (different tax-treaty and reporting complexities, e.g. US-specific PFIC considerations affecting Indian mutual fund attractiveness for US-resident NRIs specifically), and NRIs elsewhere — each facing different regulatory frictions that directly shape product fit and messaging, extending this handbook's segmentation discipline (Part 8) to a geography-and-regulation axis this handbook hasn't previously required as a primary segmentation variable.

### 49.3 GIFT City / IFSC as a distinct product-and-research context — dollar-denominated, not rupee

GIFT City's IFSC framework enables dollar-denominated investment products and trading (distinct from the rupee-denominated products this handbook's other financial-services material, Part 23, generally assumes) — meaning research into GIFT City-linked products (dollar-denominated derivatives, IFSC-listed funds) needs to account for a fundamentally different investor calculus than domestic products: currency-risk considerations the investor is implicitly taking (or avoiding) by choosing a dollar-denominated product over a rupee one, and a research population likely to include a meaningful share of NRI and foreign-institutional participants specifically drawn to the IFSC framework's distinct regulatory treatment, rather than assuming the same investor base as a comparable domestic product.

### 49.4 The remittance/relationship-triggered research moment — a distinct customer journey structure

Unlike a domestic investor's typically self-initiated account-opening journey, a meaningful share of NRI investment relationships are triggered by a specific life event with a strong existing-relationship component — a family member's recommendation, an existing NRI banking relationship being cross-sold into investment products, a specific remittance event (selling property, an inheritance) creating a lump sum needing a home for it — meaning research into the NRI investor journey should map this different trigger structure explicitly rather than assuming the same self-directed discovery funnel (Part 40.2's KYC/activation

funnel research) this handbook's domestic-investor material generally assumes, since the research questions and friction points differ meaningfully at each stage of a relationship-triggered versus self-initiated journey.

## 49.5 Worked example — scoping research for a new NRI-focused investment product

*A brokerage is considering a new NRI-focused investment product bundle and wants market research to inform the launch. The initial brief treats "NRI investors" as a single target segment.*

**Model answer.** The brief's single-segment framing should be challenged before research design begins (Part 49.2) — at minimum, splitting the research design by major geography (Gulf, US/UK/Canada, and a residual "other" category) is necessary given the genuinely different tax-treaty and regulatory considerations each faces, which directly affect which product features and messaging will actually resonate versus which will trigger disqualifying friction (e.g. US-resident NRI-specific tax complexities that could make an otherwise-attractive product a poor fit for that sub-segment specifically). The research should also explicitly map the relationship-triggered journey structure (Part 49.4) rather than assuming a self-directed discovery funnel — understanding what specific life events and existing-relationship touchpoints typically precede an NRI's decision to open a new investment relationship is likely more actionable for shaping the go-to-market approach than a generic product-feature preference survey alone, since it identifies *when and through what channel* to actually reach this segment, not just *what* to offer them once reached.

---

## PART 50 — DISCLOSURE COMPREHENSION & READABILITY TESTING

### 50.1 A distinct, narrower discipline from broad financial literacy research

Part 41 covered financial literacy at the level of general market/investing knowledge. **Disclosure comprehension and readability testing** is a narrower, distinct discipline: evaluating whether a *specific piece of mandatory disclosure text* (a risk warning, a terms-and-conditions clause, a margin-call notification, a "past performance is not indicative of future results" style statutory disclaimer) is actually understood by a realistic reader — a document-level research question, not a population-knowledge-level one, and directly relevant given how much regulated financial-services communication consists of exactly this kind of mandatory disclosure text.

### 50.2 Readability metrics — a necessary but insufficient first screen

Automated readability metrics (Flesch-Kincaid grade level, and similar formulas scoring sentence length and word complexity) give a fast, objective first screen on a disclosure's likely accessibility — but they are a **necessary, not sufficient**, check: a sentence can score as readable by a formula (short sentences, common words) while still being genuinely confusing due to dense legal/financial concept density, double negatives, or unfamiliar jargon a formula can't detect. A research team should treat a good readability score as a minimum bar cleared, not evidence the disclosure is actually understood — the same "necessary but not sufficient" caution this handbook applies to satisfaction scores as a proxy for comprehension (Part 41.3).

### 50.3 Comprehension testing — the actual test of understanding, not just readability

Genuine comprehension testing goes beyond readability scoring to directly test whether representative readers can correctly answer specific factual questions about the disclosure's content after reading it (not while looking at it) — e.g., after reading a margin-call disclosure, can the respondent correctly state what triggers a margin call and what the consequence of not responding is — a design directly analogous to Part 41.3's pre/post comprehension methodology, here applied to a single document rather than an education module, and similarly requiring actual test questions rather than a self-reported "I understood this" rating, which is prone to the same stated-vs-revealed overconfidence gap this handbook flags in its intent-vs-behaviour material (Part 22.1).

### 50.4 Plain-language rewriting and iterative testing — closing the loop on a failed comprehension test

When comprehension testing reveals a disclosure is poorly understood, the research process doesn't end at the diagnosis — the standard next step is **plain-language rewriting** (replacing jargon with everyday equivalents, restructuring dense clauses into shorter, more direct statements, adding a concrete example where a rule is stated abstractly) followed by **re-testing the rewritten version** against the same comprehension questions, iterating until the revised disclosure clears an agreed comprehension threshold — the same iterative fix-and-retest logic this handbook applies to survey instrument piloting (Part 5.8) and investor-education modules (Part 41.3), here closing the loop specifically on regulatory-disclosure clarity.

## 50.5 Worked example — testing and revising a poorly-understood margin-call disclosure

*Comprehension testing on a broker's standard margin-call disclosure text finds only 40% of tested respondents can correctly state what specifically triggers a margin call, despite the text scoring at an acceptable Flesch-Kincaid grade level.*

**Model answer.** The acceptable readability score alongside poor actual comprehension (Part 50.2's "necessary but not sufficient" distinction) confirms readability metrics alone were an inadequate check here — the disclosure is grammatically simple but conceptually unclear, likely due to dense financial-concept density (e.g. defining a margin call via a formula-based maintenance-margin threshold without a concrete worked example) rather than difficult vocabulary or sentence length, which is exactly what a Flesch-Kincaid-style metric cannot detect. The correct next step (Part 50.4) is a plain-language rewrite that specifically addresses the diagnosed comprehension gap — likely adding a concrete numeric example of a margin-call trigger rather than only the abstract rule — followed by re-testing the revised version against the same comprehension questions before considering the disclosure fixed, since a rewrite that "feels clearer" without being re-tested against the same objective comprehension bar risks repeating the same readability-score-without-comprehension-testing mistake in a new form.

---

## PART 51 — REFERRAL & WORD-OF-MOUTH GROWTH RESEARCH FOR BROKING PLATFORMS

### 51.1 A distinct growth-driver research stream from acquisition and retention

This handbook has covered CLV/CAC and retention research (Part 18) as measures of a customer relationship's ongoing value, and onboarding funnel research (Part 40) as the entry-point journey. **Referral and word-of-mouth research** studies a distinct growth mechanism entirely: how, and how much, existing customers drive *new* customer acquisition through their own advocacy — genuinely different from paid-acquisition research, since a referred customer arrives with a pre-existing trust signal (a recommendation from someone they know) rather than responding to marketing content directly.

### 51.2 NPS as an imperfect proxy for actual referral behaviour — and why to measure the real thing

NPS (Part 6) is often used as a *proxy* for referral likelihood, since its underlying question ("how likely are you to recommend") is framed around exactly this behaviour — but the well-documented gap between stated recommendation likelihood and actual referral behaviour (this handbook's recurring intent-vs-behaviour theme, Part 22.1) means NPS should not substitute for directly measuring **actual referral behaviour**: tracked referral-link usage, referral-program participation rates, and — most rigorously — the share of new customers who report an existing customer as their acquisition source, each a behavioural measure NPS alone cannot provide.

### 51.3 Referral-incentive research — does the incentive change who refers and why

When a referral program includes a financial incentive (cash, brokerage credits, both referrer and referee rewarded), research should investigate whether the incentive is **expanding** genuine advocacy (prompting satisfied customers who wouldn't have otherwise thought to refer) or primarily **substituting** for advocacy that would have happened anyway (a pure incentive-arbitrage dynamic, where referrals are transactionally motivated rather than genuine product enthusiasm) — a distinction with real implications for referred-customer quality, since research consistently shows purely incentive-motivated referrals tend to produce lower-quality, less-engaged new customers than organic, enthusiasm-driven referrals, extending this handbook's incrementality discipline (Part 19.2, Part 39.4) to the referral-growth context specifically.

### 51.4 Segmenting referral behaviour by trigger moment — when do satisfied customers actually refer

Rather than assuming referral propensity is constant over a customer's tenure, research mapping the specific **trigger moments** that precede an actual referral (a particularly positive support interaction, Part 46's service-experience material; a notable investment gain; a specific feature the customer found unexpectedly useful) identifies where in the customer journey referral-prompting touchpoints (an in-app referral prompt, a post-resolution support survey with a referral ask) would be best placed for maximum relevance and minimum intrusiveness — a more targeted approach than a generic, always-on referral prompt shown to the entire user base regardless of where they are in their own experience.

## 51.5 Worked example — diagnosing why a referral program underperforms despite high NPS

*A broking platform's NPS is strong (+38), but its referral program shows low participation and the referred customers who do sign up have below-average 90-day retention compared to non-referred customers.*

**Model answer.** The strong NPS alongside weak referral behaviour (Part 51.2's stated-vs-actual gap) confirms NPS alone shouldn't have been relied on to predict referral-program performance — satisfied customers exist, but something specific about the referral mechanism or its trigger placement (Part 51.4) may not be converting that satisfaction into actual referral action. The below-average retention among referred customers who *do* sign up (Part 51.3's incentive-quality concern) is a separate, important finding worth investigating specifically — if the program's financial incentive is attracting incentive-motivated referrers passing the link to less-genuinely-interested contacts rather than satisfied customers naturally recommending the platform to people with real investing interest, the program may be optimising for referral volume at the cost of referred-customer quality, a finding that should prompt reconsidering the incentive structure or targeting the referral ask to specific trigger moments (Part 51.4) rather than simply increasing the incentive to drive more volume of the same lower-quality referrals.

---

## PART 52 — EXPERIMENTATION METHODOLOGY: BEYOND THE SIMPLE A/B TEST

### 52.1 Why product-team experimentation deserves its own methodological depth

The glossary of this handbook has referenced A/B testing throughout as a familiar concept, but running a credible, statistically sound experimentation program for a live product (a broking app's onboarding flow, a pricing page, a feature rollout) requires methodological rigor well beyond "show two versions, see which wins" — this Part covers that deeper layer, directly relevant to a market researcher embedded in or partnering with a product team running continuous experimentation rather than only periodic standalone studies.

### 52.2 Peeking and sequential testing — the most common way a simple A/B test goes wrong

The single most common experimentation mistake is **peeking** — checking a test's results repeatedly before it reaches its pre-calculated sample size and stopping as soon as a result looks statistically significant — which dramatically inflates the true false-positive rate beyond the nominal significance threshold (a well-documented statistical problem: checking a test 10 times over its run, stopping at the first "significant" result, can produce a false-positive rate several times higher than the stated 5%). The correct disciplines are either committing to a pre-calculated sample size and duration before looking at results (extending this handbook's pre-registration-adjacent rigor from Part 5's survey-design discipline), or using **sequential testing methods** specifically designed to allow valid early stopping without inflating the false-positive rate — a genuinely different statistical framework from a standard fixed-sample-size significance test.

### 52.3 Multi-armed bandits — a distinct experimentation paradigm optimising during the test, not just after

A **multi-armed bandit** approach is a fundamentally different paradigm from a standard A/B test: rather than splitting traffic evenly and only reallocating after the test concludes, a bandit algorithm continuously shifts more traffic toward the better-performing variant *during* the test itself, reducing the "cost" of testing an underperforming variant (fewer users see the losing version for the full test duration) at the expense of a slower, noisier read on the *magnitude* of the difference between variants — a meaningful tradeoff: bandits are well-suited to situations prioritising cumulative outcome optimisation (e.g. maximising conversions during the test period itself) over precisely quantifying the size of an effect, while a standard fixed-allocation A/B test remains better suited when the *precise size* of an effect needs to be known with confidence (e.g. to build a business case for a costly follow-on investment).

### 52.4 Guardrail metrics — protecting against a "winning" test that quietly damages something else

A rigorous experimentation program tracks **guardrail metrics** alongside the primary metric being optimised — secondary metrics that shouldn't get meaningfully worse even if the primary metric improves (e.g. a checkout-flow change that lifts conversion rate but should not be allowed to also increase customer-support contact volume or reduce average order value) — extending this handbook's driver-analysis-beneath-a-headline-metric discipline (Part 6.5) specifically to the experimentation context: a test declared a

"win" on its primary metric alone, without checking guardrails, risks shipping a change that solves one problem while quietly creating another the primary metric wasn't designed to catch.

## 52.5 Worked example — diagnosing a fast "winning" bandit-based experiment

*A product team runs a bandit-based experiment on a new onboarding flow. Within days, the bandit algorithm has shifted the large majority of new-user traffic to the new flow, since it's showing a strong early lift in same-day account activation. A guardrail metric — 30-day retention — hasn't yet had time to mature for most of the cohort exposed to the new flow.*

**Model answer.** The bandit's fast reallocation toward the apparent early winner (Part 52.3's core mechanism) is working as designed for the primary same-day-activation metric, but the guardrail-metric caution (Part 52.4) applies directly here: a flow that lifts fast, same-day activation could plausibly do so by front-loading users into the funnel with weaker downstream fit (e.g. skipping an explanatory step that, while slowing initial activation, better-qualified users for sustained engagement) — meaning the bandit's fast traffic-shifting is optimising for a metric that matures quickly, potentially before the guardrail metric that matters more for long-term business health has had time to reveal a problem. The correct response is not to fully trust the bandit's early reallocation as final proof of a genuine win — holding out a smaller comparison cohort on the original flow long enough for the 30-day retention guardrail to mature, even while the bandit continues optimising the majority of traffic toward the apparent short-term winner, is the disciplined way to avoid shipping a change that wins on a fast metric while quietly damaging a slower, more consequential one.

---



## PART 53 — BROKERAGE-FEE & PRICING-MODEL PERCEPTION RESEARCH

### 53.1 A distinct pricing-research category shaped by a genuinely unusual competitive dynamic

This handbook's general pricing-research methods (Part 12's Van Westendorp, Gabor-Granger) assume a relatively conventional price-value tradeoff. Broking-platform pricing research operates in a genuinely unusual competitive category: the widespread emergence of **zero-brokerage (or near-zero) equity delivery trading** has made "free" the prevailing category norm for a core transaction type, meaning research here is less about finding an optimal price point on a demand curve and more about understanding how a platform's *overall* monetisation model (where fees are charged elsewhere — F&O brokerage, margin/interest charges, subscription tiers, payment-for-order-flow-adjacent revenue) is perceived, trusted, and understood by users who've been conditioned to expect the headline transaction itself to be free.

### 53.2 Fee-structure comprehension research — distinct from disclosure comprehension, but related

Related to Part 50's disclosure comprehension testing but distinct in scope, **fee-structure comprehension research** specifically tests whether users understand a broking platform's *complete* cost structure — not just the advertised "zero brokerage" headline, but the full picture including F&O per-lot charges, overnight margin funding interest rates, and any account-maintenance or platform fees — since research in this category consistently finds a meaningful comprehension gap between a platform's simple, prominently-marketed headline price and its more complex full cost structure, a gap that becomes a real source of user dissatisfaction and support-contact volume (Part 46) when a user's actual bill diverges from their expectation formed around the free-delivery-trading headline.

### 53.3 Cross-platform fee comparison research — the shopping-behaviour question specific to this category

Because brokerage switching costs are relatively low (opening a second demat/trading account is a modest but real hurdle, not a major one) compared to many other financial-services categories, **cross-platform fee comparison** — how actively users compare and switch platforms specifically based on fee structure, and which specific fee components (not the headline "free" claim, which has become table-stakes and largely commoditised across major platforms) actually drive comparison and switching decisions — is a genuinely distinct and more actionable research question than a generic price-satisfaction survey, since it identifies the specific cost components (often F&O charges for active traders, given the fee is charged per-lot and matters disproportionately to high-frequency users) that actually influence competitive positioning.

### 53.4 Perceived fairness research — a distinct dimension from perceived cost level

Beyond the objective cost level, **perceived fairness** of a fee structure is a distinct research dimension worth measuring separately — a fee a user perceives as transparently and predictably applied (clearly shown before an order is placed, consistent across similar transactions) can be rated favourably even at a comparable objective cost level to a fee perceived as opaque, inconsistently applied, or "hidden" (only discovered after the fact in a statement) — extending this handbook's trust-signal research theme (Part 40.4)

specifically to the pricing dimension, since a platform's fee *transparency*, not just its fee *level*, is a meaningfully separate driver of user trust and satisfaction worth its own research attention.

### 53.5 Worked example — diagnosing rising support contacts about "unexpected" F&O charges

*A broking platform sees a rising trend in support contacts specifically about "unexpected" or "hidden" F&O charges, even though the fee structure is fully disclosed in the platform's published fee schedule and shown at order placement. New F&O traders (within their first month of options trading) are disproportionately represented in these contacts.*

**Model answer.** The disclosure being technically complete (published fee schedule, shown at order time) doesn't resolve the actual problem — this is a fee-structure comprehension gap (Part 53.2) concentrated specifically among new F&O traders, most plausibly because these users formed their cost expectations around the platform's prominent zero-brokerage-delivery-trading positioning (Part 53.1) and didn't fully register the separate F&O fee structure shown at the point of a specific order, a comprehension failure at the moment of formation of overall cost expectations rather than at any single disclosure point. The research-informed fix isn't necessarily more disclosure (which is already technically complete) but addressing the perceived-fairness and expectation-setting gap (Part 53.4) — for instance, a clearer upfront cost-structure explanation specifically during F&O onboarding (tying to Part 41's pre/post comprehension-testing methodology, applied here to fee-structure understanding rather than risk education) that sets expectations before the first F&O trade, rather than relying on order-time disclosure alone to correct an expectation already formed around a different product's pricing.

---

## PART 54 — INVESTOR RELATIONS EFFECTIVENESS PERCEPTION RESEARCH

### 54.1 A distinct research subject — the IR function itself, not just its output

Earlier material in this handbook (Part 32's investor-perception research, Part 45's pre-IPO messaging testing) covered how investors perceive a company's story or a specific offering. This Part covers a related but distinct research subject: how effectively a company's **Investor Relations (IR) function** itself operates as an ongoing communication channel — earnings-call quality, disclosure timeliness and clarity, analyst-day/investor-day effectiveness, and management accessibility — a genuinely different research object from the underlying business story, since even a strong business can be perceived poorly by the investment community if its IR communication is inconsistent, unclear, or hard to access.

### 54.2 Analyst and institutional-investor perception surveys — a distinct respondent population

IR-effectiveness research typically surveys or interviews **sell-side analysts covering the stock** and **institutional/buy-side investors holding or considering the stock** — a genuinely different respondent population from the retail-investor-facing research covered elsewhere in this handbook (Part 23), with different research access considerations (this population is small, professionally sophisticated, and often reachable through existing broker/analyst relationships rather than a standard survey panel) and different content priorities (comparative benchmarking against sector peers' IR quality is often more relevant to this audience than to a retail investor).

### 54.3 Specific, benchmarkable dimensions of IR effectiveness worth researching separately

Rather than a single generic "IR quality" rating, credible IR-effectiveness research breaks the assessment into specific, separately benchmarkable dimensions: **earnings-call quality** (clarity of prepared remarks, how directly and substantively management answers analyst questions during Q&A, whether guidance is given with useful specificity), **disclosure timeliness and clarity relative to peers**, **management accessibility** (availability for follow-up questions, one-on-one investor meetings, conference participation), and **consistency of messaging over time** (whether the narrative and stated priorities shift confusingly quarter to quarter without clear explanation) — each dimension can point to a different, specific improvement area, similar in spirit to this handbook's driver-analysis-beneath-a-headline-score discipline (Part 6.5) applied to the IR function specifically.

### 54.4 Why IR-effectiveness perception has a real, if hard-to-isolate, valuation-adjacent consequence

A company's IR effectiveness plausibly affects its **cost of capital and valuation multiple indirectly**, through channels like analyst coverage breadth (well-served analysts more willing to initiate and maintain coverage, widening the pool of investors aware of the story), reduced information-asymmetry-driven valuation discounts, and lower event-driven volatility around results (clearer guidance reducing the frequency and size of surprise-driven price reactions) — a genuinely harder-to-isolate causal link than most research findings in this handbook, given the many other factors affecting valuation simultaneously, but a

link worth explicitly acknowledging as a motivating business case for IR-effectiveness research even though attributing a precise valuation impact to IR quality specifically is rarely possible with confidence.

### 54.5 Worked example — diagnosing declining analyst coverage despite unchanged fundamentals

*A mid-cap company's sell-side analyst coverage has declined from 12 covering analysts to 7 over 18 months, despite no meaningful deterioration in the underlying business. IR-effectiveness research finds analysts frequently cite inconsistent guidance (specific numeric targets given one quarter, withdrawn or vaguely restated the next) and difficulty securing timely follow-up access to management after earnings calls.*

**Model answer.** With fundamentals roughly unchanged, the coverage decline is more plausibly explained by the specific IR-quality issues surfaced in the research (Part 54.3's dimension-specific diagnosis) than by any change in the underlying investment case — inconsistent guidance (Part 54.3) makes it harder for analysts to build and defend confident forecasts, directly undermining their ability to publish and maintain conviction-backed coverage, while limited follow-up access compounds the problem by making it harder for analysts to resolve their own uncertainty through direct engagement. This diagnosis (Part 54.4's valuation-adjacent consequence made concrete) gives management and the IR function specific, actionable levers — more disciplined, consistent guidance practices and a more accessible post-earnings-call follow-up process — rather than only a general sense that "the market doesn't understand the story," a vague complaint the specific research findings here replace with two concrete, addressable root causes.

---

## PART 55 — FINFLUENCER MARKETING RESEARCH: TRUST, COMPLIANCE & EFFECTIVENESS

### 55.1 A distinct research problem from Part 34's generic influencer-campaign effectiveness

Part 34.5 covered diagnosing why an influencer campaign shows strong engagement but weak conversion — a generic effectiveness question applicable to any product category. Financial-content creators ("finfluencers") raise a genuinely distinct additional research dimension specific to this category: **regulatory compliance risk and audience trust calibration**, since financial content carries real potential for consumer harm if misunderstood or if it crosses from general education into what regulators consider unlicensed investment advice — a compliance and trust dimension a snack-food or apparel influencer campaign simply doesn't carry.

### 55.2 Disclosure-compliance research — a distinct, higher-stakes version of Part 34's transparency theme

Financial-content creator partnerships in India operate under specific disclosure requirements (clear identification of sponsored/paid content, and increasingly specific rules around registered-vs-unregistered individuals giving investment-advice-adjacent content) — research into whether a brand's finfluencer partnerships are achieving genuine, clearly-noticed disclosure compliance (not just technically present but genuinely *seen and understood* by the audience as sponsored content, extending this handbook's comprehension-testing discipline, Part 50, from mandatory product disclosures to mandatory sponsorship disclosures) is a materially higher-stakes research question here than in most other influencer-marketing categories, given the direct regulatory exposure for both the brand and the creator.

### 55.3 Credibility and expertise-perception research — distinct from generic engagement metrics

Beyond engagement metrics (views, likes, comments — Part 34's standard influencer-effectiveness measures), **credibility and expertise-perception research** specifically asks whether the audience perceives a given financial-content creator as a genuine, trustworthy source of financial information versus purely entertaining content with a financial veneer — a distinction with real downstream implications, since research in this space consistently shows audience trust and perceived expertise (not raw reach or engagement) is the stronger predictor of whether financial-content marketing actually shifts product consideration, extending this handbook's audience-brand-fit concept (Part 34.5's diagnostic framework) with an added credibility dimension specific to information-sensitive categories like finance.

### 55.4 Content-accuracy and risk-messaging audit — a research function beyond standard campaign tracking

Given the potential-harm dimension unique to financial content, a rigorous brand partnership program includes an ongoing **content-accuracy and risk-messaging audit** — systematically reviewing whether a creator's actual published content (not just the pre-approved brief) maintains balanced risk messaging over time and doesn't drift toward overpromising returns or minimising risk in pursuit of higher engagement — a research/compliance function distinct from campaign-performance tracking, closer in spirit to this handbook's regulatory-impact research material (Part 48) than to conventional marketing-effectiveness

measurement, since the downside risk here (regulatory action, reputational harm from a creator's content drifting into misleading territory) is asymmetric and not adequately captured by engagement or conversion metrics alone.

### 55.5 Worked example — evaluating a finfluencer partnership showing strong metrics but a compliance flag

*A brand's finfluencer partnership shows strong engagement and above-average click-through to the product page. A content-accuracy audit flags that several recent posts from the creator understate investment risk relative to the brand's approved messaging guidelines, though technically within the disclosed-sponsorship requirement.*

**Model answer.** Strong engagement and click-through metrics alone (Part 55.3's caution against relying on engagement as the primary success measure) should not be read as confirming this is a successful, low-risk partnership — the content-accuracy audit finding (Part 55.4) is a distinct and, in this category, potentially more consequential signal than the performance metrics, since risk-understating content creates real regulatory and reputational exposure regardless of how well it's converting. The correct response is not simply weighing the compliance flag against the strong metrics as competing considerations to balance, but recognising this as a different *category* of concern (Part 55.2's higher-stakes compliance framing) that needs direct remediation — engaging the creator on the specific risk-messaging drift, tightening review processes for ongoing (not just pre-approved) content, and treating strong short-term performance metrics as no substitute for verifying the partnership continues to meet accuracy and disclosure standards over its full duration, not just at initial brief approval.

---

## PART 56 — GAMIFICATION RESEARCH: ENGAGEMENT DESIGN & TRADING BEHAVIOUR RISK

### 56.1 A distinct research tension unique to this category — engagement design meeting real financial risk

Trading and broking apps commonly use gamification elements (streaks, badges, celebratory animations on trade execution, leaderboards, progress bars toward account milestones) drawn from the same engagement-design playbook used across consumer apps generally — but in this category specifically, these design patterns interact with genuine financial risk-taking behaviour, creating a research tension this handbook hasn't needed to address elsewhere: the same design elements that improve engagement metrics can plausibly also encourage excessive or impulsive trading behaviour, a dual-use-design problem distinct from gamification research in a typical consumer or fitness app context.

### 56.2 Engagement-metric research versus behaviour-quality research — two distinct, sometimes-conflicting measurement layers

Standard product research measures whether a gamification element improves standard engagement metrics (session frequency, feature adoption, retention) — but a rigorous research program for this specific category must run a **second, distinct measurement layer** alongside it: whether the same element correlates with lower-quality trading behaviour (increased trading frequency without a corresponding change in account performance, larger position sizes taken shortly after a gamified celebratory moment, evidence of impulsive re-entry after a loss) — extending this handbook's guardrail-metrics discipline (Part 52.4) specifically to the ethical dimension of engagement design in a financial-risk context, where the "guardrail" isn't a business metric like support-contact volume but user financial wellbeing itself.

### 56.3 Regulatory attention as a research-scoping input, not just a compliance afterthought

Regulators globally, including in India, have specifically flagged gamification elements in trading apps as a supervisory concern given the potential to encourage behaviour inconsistent with prudent investing — meaning gamification research for a broking platform should be scoped from the outset with this regulatory attention in mind (similar to Part 55.2's influencer disclosure-compliance framing), rather than treating regulatory risk as a separate compliance review that happens only after a design has already been built and tested purely on engagement metrics.

### 56.4 Segment-specific vulnerability research — new/inexperienced traders as a distinct research priority

Gamification-behaviour research should specifically prioritise **new and inexperienced traders** as a distinct segment (extending this handbook's trading-frequency segmentation, Part 40.3) rather than only analysing pooled, all-user data — since a design element's behavioural effect on a highly experienced, disciplined trader may look completely different from its effect on a first-time F&O trader with no established risk framework of their own, and pooled analysis risks averaging away a genuine, segment-specific harm signal that only shows up clearly when the newest, most vulnerable user segment is analysed separately.



## 56.5 Worked example — evaluating a "trading streak" feature ahead of a broader rollout

*A broking app is piloting a "trading streak" feature (visual recognition for consecutive days with at least one trade) that shows a strong lift in daily active trading during the pilot. Behaviour-quality analysis on the pilot cohort, segmented by experience level, shows the streak feature correlates with meaningfully increased trade frequency among first-90-day traders specifically, with no corresponding improvement in their trading outcomes, while experienced traders show negligible frequency change.*

**Model answer.** The strong pooled engagement lift alone (Part 56.2's first measurement layer) would be an incomplete and potentially misleading basis for a rollout decision — the segment-specific finding (Part 56.4) reveals the feature's actual behavioural effect is concentrated specifically among the most vulnerable segment (first-90-day traders), correlating with increased trading frequency but no corresponding performance improvement, a pattern consistent with the encouraged-impulsive-trading risk this category's gamification research specifically needs to screen for (Part 56.1). Given the regulatory attention this specific design pattern draws (Part 56.3), the research-informed recommendation is not a straightforward "ship it, engagement is up" — it should flag the segment-specific behaviour-quality finding explicitly to product leadership and compliance, and consider either redesigning the feature to avoid concentrating its effect on new-trader impulsivity (e.g. reframing the streak around account-health or learning-module completion rather than raw trade count) or excluding first-90-day traders from the feature's exposure entirely, rather than treating a strong pooled engagement metric as sufficient justification on its own.



## PART 57 — HIRING-SIGNAL & JOB-POSTING ANALYSIS, DEEPENED

### 57.1 From a one-line CI checklist item to a systematic methodology

Part 3's competitive-intelligence sources list mentioned job postings as one signal among several, in a single line. This Part deepens hiring-signal analysis into an actual methodology: systematically tracking a company's job-posting volume, functional mix, and seniority mix over time as a leading indicator of strategic direction — distinct from, and often available earlier than, the alternative-data sources (app-download rankings, satellite/web-traffic signals) covered in Part 43, since hiring intent frequently precedes a company's own public announcement of a new initiative by months.

### 57.2 Volume trend and its limitations — distinguishing genuine expansion from replacement hiring

Raw job-posting volume trend (rising, flat, or falling over successive months) is a useful first screen, but requires the same self-selection-style caution this handbook applies to other aggregate metrics (Part 47.2's review-platform caution) — a rising posting count can reflect genuine headcount expansion, or simply elevated **replacement hiring** driven by higher-than-normal attrition, two very different underlying stories that raw volume alone cannot distinguish. Cross-referencing posting volume against independent attrition-adjacent signals (Part 47's employee-review-mining material, specifically watching for a rising volume of departure-related review content) helps disambiguate which story is more likely driving an observed posting-volume increase.

### 57.3 Functional and seniority mix — the more diagnostic layer beneath aggregate volume

More diagnostic than aggregate volume is the **functional mix** of new postings (a surge specifically in sales/business-development roles suggests a go-to-market push; a surge in a specific engineering specialisation suggests a new product-line investment; a surge in a new geographic market's postings suggests market-entry activity, tying directly to this handbook's market-entry research material, Part 28) and the **seniority mix** (a new senior leadership posting in a function the company previously ran without a dedicated leader — e.g. a first-ever "Head of International Expansion" role — is a particularly strong, specific signal of strategic intent, since it represents organisational investment in a function, not just incremental headcount).

### 57.4 LinkedIn headcount-by-function trend analysis — a complementary, lagging-but-confirming data source

Beyond live job postings (a forward-looking intent signal), tracking a company's **LinkedIn headcount-by-function trend** over time provides a complementary, somewhat lagging but confirming data source — showing whether posted openings actually converted into realised headcount growth in that function, and at what pace, distinguishing a company that posts aggressively but converts slowly (potentially indicating hiring difficulty or a slower-than-planned initiative) from one whose headcount visibly tracks its posting activity closely (a cleaner confirmation the initiative is proceeding as the posting activity suggested).

## 57.5 Worked example — reading a competitor's hiring-signal pattern ahead of a market entry

*A fintech competitor's job postings show a sustained three-month increase specifically in roles based in a new geographic city the company doesn't currently have a visible presence in, spanning sales, compliance, and a senior "Regional Head" position — with LinkedIn headcount data showing early but growing realised headcount in that city over the same period, and no public announcement yet made.*

**Model answer.** The functional and seniority mix here (Part 57.3) — sales, compliance, and a senior regional leadership role, specifically concentrated in one new city — is a considerably stronger, more specific signal than a generic rise in overall posting volume would be, since this combination is characteristic of a deliberate market-entry preparation rather than routine hiring; the compliance-role inclusion specifically suggests awareness of local regulatory requirements being built out ahead of launch. The corroborating LinkedIn headcount growth in that city (Part 57.4) strengthens confidence this reflects genuine, converting hiring activity rather than postings that may not materialise into actual headcount. The appropriate research output is flagging this as a well-evidenced, pre-announcement early-warning signal of a likely market-entry move into that city — valuable specifically because it's available well before any public confirmation — while still applying this handbook's alternative-data discipline (Part 43.3's ground-truthing caution) by continuing to monitor for corroborating signals rather than treating the hiring pattern alone as definitive proof of the company's plans.

---

## PART 58 — COHORT ANALYSIS & RETENTION-CURVE METHODOLOGY

### 58.1 A distinct analytical technique from CLV's single-number output

Part 21's CLV material produces a single aggregate number describing customer value over time. **Cohort analysis** takes a related but methodologically distinct approach: grouping customers by a shared starting point (typically the month or week they first signed up) and tracking each cohort's behaviour — most commonly retention — separately over time, producing a table or curve per cohort rather than one blended aggregate figure, revealing patterns an aggregate metric averages away entirely.

### 58.2 The cohort retention table — reading rows, columns, and diagonals

A standard cohort retention table has **rows** representing acquisition cohorts (e.g. each month's new sign-ups) and **columns** representing time since acquisition (Month 1, Month 2, Month 3...), with each cell showing what percentage of that cohort remained active at that point — reading **across a row** shows how a single cohort's retention decayed over its lifetime (the retention curve shape), reading **down a column** compares different cohorts at the same point in their lifecycle (whether newer cohorts are retaining better or worse than older ones at the equivalent stage), and reading the **diagonal** reveals calendar-time effects (a specific external event, like a market crash or a product outage, hitting all active cohorts simultaneously regardless of their individual tenure) — three genuinely different questions the same table can answer depending on which direction it's read.

### 58.3 Retention-curve shape — why the curve's flattening point matters more than any single-month number

A retention curve's shape carries more diagnostic information than any single month's retention percentage: a curve that declines steeply early but **flattens** into a stable long-term plateau (a "smile curve" in retention terminology) indicates the surviving cohort has found genuine, durable value and identifies roughly where a stable core user base settles — versus a curve that continues declining without flattening, suggesting the product hasn't yet found its durable-value core and the true long-term retention rate remains unknown from the available data window. This mirrors a general research principle already applied elsewhere in this handbook to a different curve (Part 21.4's CLV-vs-churn-rate relationship): the shape of a trajectory, not a single data point along it, carries the real signal.

### 58.4 Cohort-vs-cohort comparison for diagnosing whether a product change actually helped

Comparing retention curves across cohorts acquired before versus after a specific product change (a new onboarding flow, Part 40's KYC-funnel redesign, a new feature launch) is a genuinely more rigorous way to assess whether the change improved actual retained value than looking only at an aggregate, blended retention number across the whole user base — since the aggregate number mixes old and new cohorts together and can mask or dilute a real effect visible only when cohorts are compared like-for-like at the equivalent point in their own lifecycle, extending this handbook's experimentation-methodology discipline (Part 52) to a non-experimental, observational cohort-comparison context where a formal randomised test wasn't run.

## 58.5 Worked example — reading a cohort table to evaluate a redesigned onboarding flow

*A broking platform redesigned its onboarding flow six months ago. The cohort retention table shows cohorts acquired after the redesign have noticeably higher Month-1 retention than pre-redesign cohorts, but by Month-4, both old and new cohorts have converged to a similar retention percentage.*

**Model answer.** Reading down the Month-1 column (Part 58.2) shows a genuine early improvement from the redesign — but the convergence by Month-4 (reading further along each row, Part 58.3's curve-shape discipline) reveals the improvement was concentrated in early retention specifically, not sustained long-term value: the redesign appears to be successfully retaining more users through their first month (plausibly better onboarding reducing early drop-off) without changing the underlying determinants of longer-term engagement, which apparently converge to a similar equilibrium regardless of the onboarding experience. This is a materially more precise and actionable finding than either "retention improved" (true only for Month-1) or "retention didn't really change" (true only from Month-4 onward) would be in isolation — the correct research conclusion is that the onboarding redesign solved an early-funnel problem specifically, and that if longer-term retention is the actual business priority, a different intervention (rather than further onboarding-flow iteration) is likely needed to move the Month-4+ plateau itself.

---

## PART 59 — POST-PURCHASE DISSONANCE & INVESTMENT-REGRET RESEARCH

### 59.1 A distinct post-decision research window from pre-decision behavioural finance material

This handbook's earlier behavioural-research material (Part 22's implicit-measurement toolkit, Part 40.4's risk-perception calibration) primarily addresses biases and comprehension gaps shaping a decision *before or during* the choice itself. **Post-purchase cognitive dissonance research** studies a distinct, later window: the psychological discomfort that can arise *after* a high-consideration financial commitment (opening a large position, choosing a long-term investment product, taking on significant leverage) as the investor reconciles the decision against subsequent information or outcomes — a genuinely different research object from pre-decision bias research, since it studies how people cope with a choice already made rather than how the choice itself was shaped.

### 59.2 Why financial decisions are a particularly high-dissonance-risk category

Cognitive dissonance theory holds that discomfort is greatest for decisions that are important, difficult to reverse, and involve conflicting information — a profile that describes many financial-product decisions unusually well (a long-lock-in investment product, a large position taken on margin, a pension/retirement product commitment) compared to lower-stakes, easily-reversible everyday purchases, making financial-services categories a particularly high-dissonance-risk research area worth its own dedicated attention rather than assuming generic consumer post-purchase research frameworks transfer unchanged.

### 59.3 Dissonance-reduction behaviours as a distinct, observable research signal

Rather than measuring dissonance directly (a difficult, abstract internal state), research in this space typically observes the **behaviours investors use to reduce dissonance** — selectively seeking out confirming information after a purchase (increased consumption of content validating the decision already made), avoiding information that would call the decision into question, or rationalising a suboptimal outcome after the fact — behavioural proxies for an underlying discomfort state, extending this handbook's general preference for behavioural/revealed measures over difficult-to-self-report internal states (Part 22's stated-vs-revealed theme) into the post-purchase dissonance context specifically.

### 59.4 The churn/complaint-timing signal — dissonance research's most actionable business application

A specific, highly actionable research application: tracking **when** post-purchase regret-related contact (a complaint, a request to exit a product, a support query specifically framed around "did I make the right choice") tends to occur relative to the original purchase date — research consistently finds regret-driven contact clusters at specific, predictable windows (shortly after purchase, before dissonance-reduction behaviours have had time to take hold; or at the next major market/statement-review touchpoint, when performance becomes newly salient again) — identifying these windows allows a business to proactively intervene (reassurance communication, a check-in touchpoint) at the specific moments regret is most likely to surface and potentially convert into a support contact or exit request, rather than only reacting after a complaint has already been filed.

## 59.5 Worked example — diagnosing a spike in early-tenure exit requests for a long-lock-in product

*A wealth platform's long-lock-in investment product shows a notable spike in exit/withdrawal requests specifically in the first 30 days after purchase, disproportionate to requests at any later point in the product's multi-year term, despite the product's actual documented performance being in line with its stated expectations.*

**Model answer.** The concentration of exit requests specifically in the first 30 days (Part 59.4's timing-signal discipline) is consistent with a dissonance-driven pattern rather than a genuine performance-driven grievance, since the documented performance itself is in line with expectations — the early window is precisely when dissonance-reduction behaviours (Part 59.3) haven't yet had time to settle the investor's post-decision discomfort about a large, hard-to-reverse commitment (Part 59.2's high-dissonance-risk profile). The research-informed response isn't simply reinforcing why the product's performance is fine (which the exit-requesters already technically know, given performance is in line) — it's designing a proactive touchpoint specifically timed to this identified early window (a brief, reassuring check-in reiterating the product's original purpose and time horizon, delivered before the dissonance peak rather than only in reactive response to an already-filed exit request), directly informed by the timing pattern the research surfaced rather than a generic "improve satisfaction" response that doesn't target the specific window where intervention would actually be most effective.

---

## PART 60 — RECOMMENDATION ENGINE & CONTENT-PERSONALIZATION TESTING

### 60.1 A distinct research object from robo-advisory portfolio recommendations

Part 42 covered research on algorithm-driven **portfolio allocation** recommendations specifically. This Part covers a broader, distinct category: **content and discovery personalization** — a personalised market-news feed, algorithmically-suggested watchlist additions, "stocks similar to ones you follow" recommendations, personalised fund-discovery browsing — engagement-and-discovery-oriented recommendation surfaces rather than portfolio-construction advice, raising a genuinely different research question (is this specific piece of content/suggestion relevant and useful to *this* user right now) rather than Part 42's trust-in-financial-advice question.

### 60.2 Relevance evaluation — beyond click-through rate as the sole success metric

The most common but incomplete way to evaluate a recommendation engine is click-through rate on suggested items — an incomplete measure since a user might click a sensationalist or clickbait-adjacent suggestion (a volatile penny stock, an alarmist headline) at a high rate without it being genuinely *useful* to their actual information needs or investing goals; rigorous relevance research supplements click-through with direct relevance ratings (a brief in-context "was this suggestion useful to you" prompt) and, more rigorously, downstream behaviour tracking (did the user actually engage further with the recommended item's underlying content, or bounce immediately) — extending this handbook's guardrail-metrics discipline (Part 52.4) to the personalization-quality context specifically.

### 60.3 The filter-bubble and diversity research question — a distinct quality dimension from relevance alone

Beyond individual-item relevance, a recommendation system optimised purely to maximise engagement can systematically narrow what a user sees over time (a **filter bubble**) — for a financial-content platform specifically, this raises a genuine research question distinct from generic media filter-bubble concerns: whether a user's stock/fund recommendations are becoming so narrowly tailored to their existing interests that they're no longer being exposed to genuinely relevant diversification opportunities or balanced risk perspectives, a research question with real product-design implications (deliberately preserving some exploration/diversity in recommendations, not just pure relevance-maximisation) given the category's stakes.

### 60.4 Cold-start research — a distinct methodological problem for new users with no behavioural history

A recommendation engine has no behavioural history to personalise against for a **new user** (the "cold-start problem"), meaning research into what a platform shows a first-time or low-history user is a distinct, separate research question from steady-state personalization quality for an established user — commonly relying on explicit onboarding preferences (stated interests, self-reported risk profile) or broad segment-level defaults until sufficient behavioural data accumulates, and research should specifically evaluate whether these cold-start defaults are reasonably relevant on their own merits, since a poor cold-start experience can drive early churn (Part 40.2's activation-funnel material) before the system has enough data to personalise meaningfully.

## 60.5 Worked example — diagnosing a recommendation engine's engagement-vs-quality tradeoff

*A stock-discovery feed's algorithm update lifts overall click-through rate on recommended stocks by 15%. Relevance-rating research on the updated algorithm shows in-context relevance ratings actually declined slightly, and downstream engagement (time spent on the recommended stock's detail page after clicking) dropped notably, concentrated on higher-volatility, more sensationally-framed recommendations that increased in frequency under the new algorithm.*

**Model answer.** The click-through lift alone (Part 60.2's incomplete-metric caution) would be a misleading basis for declaring the update successful — the relevance-rating decline and dropped downstream engagement reveal the algorithm is likely optimising toward clickbait-adjacent, high-volatility suggestions that generate an initial click without delivering genuine value, precisely the gap between click-through and true relevance this handbook's methodology is designed to catch. This also raises the filter-bubble/diversity question (Part 60.3) — if the algorithm is systematically shifting toward more sensational recommendation types because they generate more clicks, it may be narrowing the *kind* of content shown in a way that serves the metric being optimised rather than the user's actual investing information needs. The research-informed recommendation is reverting or further tuning the update to weight the fuller relevance-and-downstream-engagement picture rather than click-through alone, treating this as a direct, concrete illustration of why a single engagement metric is an insufficient success criterion for a recommendation system in this specific, information-sensitive category.



## PART 61 — LONGITUDINAL PANEL DESIGN FOR TRACKING SENTIMENT ACROSS MARKET CYCLES

### 61.1 A distinct methodological layer beneath a brand tracker's frozen instrument

Part 10.4 established why a survey instrument stays methodologically frozen wave over wave to isolate genuine market change from measurement artifacts, and Part 23.5 introduced investor confidence/sentiment indices as a financial-services-specific tracker application. This Part addresses a related but distinct methodological question specific to tracking studies that run for years, spanning genuinely different market conditions (a bull market, a correction, a full bear cycle): the **panel itself** — the actual group of respondents being repeatedly surveyed — has its own dynamics that can bias results independent of anything happening in the instrument or the market.

### 61.2 Panel attrition and its correlation with the very outcome being tracked

Respondents drop out of a long-running panel over time (panel attrition) — the critical risk specific to investor-sentiment tracking is that attrition is unlikely to be random with respect to the outcome being measured: investors who become disillusioned or suffer significant losses may be disproportionately likely to stop responding (from disengagement, embarrassment, or simply exiting the market and no longer feeling the survey is relevant to them), meaning a panel's remaining respondents can skew progressively more positive over a market downturn purely through differential attrition, even if no individual respondent's actual sentiment changed — a distinct, panel-specific bias mechanism beyond ordinary survey non-response bias (Part 5).

### 61.3 Panel conditioning — the survey itself changing the respondent over repeated exposure

**Panel conditioning** is a second distinct risk: respondents who are repeatedly asked the same questions can become more sophisticated or self-aware about the topic purely through repeated exposure to the survey instrument itself (a novice investor asked quarterly about their risk tolerance and market outlook may become measurably more financially literate over successive waves simply from the act of repeatedly considering these questions) — meaning observed trends in a long-running panel can partly reflect this conditioning effect rather than genuine population-level change, a confound distinct from and in addition to Part 61.2's attrition-bias risk.

### 61.4 Refreshment sampling — the standard defence against both risks simultaneously

The standard methodological defence against both attrition bias and conditioning effects is **refreshment sampling**: periodically adding newly-recruited respondents to the panel (recruited fresh, with no prior exposure to the survey instrument) alongside the continuing longitudinal respondents — comparing the fresh refreshment sample's responses against the continuing panel's responses at the same wave surfaces both problems simultaneously (a large gap suggests either significant attrition bias, conditioning effects, or both), and blending appropriately-weighted fresh and continuing respondents in the reported topline numbers keeps the tracked trend more representative of the true, current population than relying purely on an aging, self-selected, increasingly-conditioned continuing panel.

## 61.5 Worked example — diagnosing a suspiciously resilient investor-confidence tracker during a downturn

*A financial-services firm's quarterly investor-confidence tracker, running continuously for three years with no refreshment sampling, shows confidence holding up unusually well through a six-month market correction — notably more resilient than independent market-wide sentiment indicators suggest is plausible.*

**Model answer.** The absence of refreshment sampling (Part 61.4) is the first thing to flag — without a comparison against freshly-recruited respondents, it's impossible to distinguish whether the panel's resilient confidence reflects genuine, real investor sentiment or is an artifact of Part 61.2's attrition-bias mechanism (the most anxious, loss-affected respondents disproportionately dropping out of the panel during the correction, leaving a residual sample skewed toward more resilient, still-engaged investors) compounded by Part 61.3's conditioning effect (repeated respondents having become desensitised to market-volatility questions after answering them many times). The correct diagnostic step is retrospectively pulling response-rate data by wave to check whether attrition increased during the correction period specifically (a telltale sign of the Part 61.2 mechanism), and going forward, implementing refreshment sampling before trusting this tracker's readings through future market cycles — a tracking programme that held up well through one downturn purely by coincidental attrition patterns cannot be assumed to do so reliably again without this structural fix in place.

---

## APPENDIX A — GLOSSARY

TAM/SAM/SOM · NPS (Net Promoter Score) · CATI/CAPI/CAWI (interview modes) · Conjoint analysis · MaxDiff · Cross-tab · Base size · Significance (p-value) · Segmentation (a-priori vs post-hoc) · Persona · Win-loss analysis · CI (Competitive Intelligence) · Syndicated research · Panel · Likert scale · Triangulation · Van Westendorp Price Sensitivity Meter · Gabor-Granger (pricing method) · Aaker's Brand Equity model · Keller's CBBE pyramid · Unaided/aided awareness · Top-of-mind (TOM) · Brand linkage · Perceptual mapping · Factor analysis / MDS · Firmographics · Technographics · Account-based research (ABR) · DPDP Act 2023 · Sugging/frugging · Incidence rate · Straight-lining · Standardised beta (regression) · R<sup>2</sup> (variance explained) · Part-worth utility · Share of preference · Social listening · Share of voice (SoV) · A/B testing · MROC (Market Research Online Community) · Mobile ethnography · Passive metering · Shopper vs consumer research · Share of shelf (SOS) · Planogram testing · Retail audit · Mystery shopping · Share of search · Back-translation · Acquiescence bias · Extreme response style · K-means clustering · Elbow method · Brand-Price Trade-Off (BPTO) · Price elasticity of demand (PED) · Revenue management · Hypothetical bias · Retail audit panel vs household panel · Numeric vs weighted distribution · Penetration · Buying rate/frequency · Source of growth · Concept screening vs concept testing · Top-box/top-two-box · Norms database · Monadic / sequential-monadic / comparative testing · Blind vs branded testing · JAR (Just-About-Right) scale · Category management · Incrementality analysis · Baseline vs incremental sales · Pull-forward/pantry-loading effect · Promotion ROI · Marketing Mix Modeling (MMM) · Adstock/carryover · Diminishing returns (saturation) · Marginal ROI · Churn rate · Contractual vs non-contractual churn · Voluntary vs involuntary churn · Predictive churn modelling · Customer Lifetime Value (CLV) · CAC (Customer Acquisition Cost) · CLV:CAC ratio · Eye-tracking · Facial coding · Implicit Association Test (IAT) · Galvanic skin response (GSR) · Risk tolerance segmentation · Self-directed vs advised investor · Investor confidence/sentiment index · Unmoderated remote testing · Continuous/always-on research · Rapid-research decision framework · Codebook · Inductive vs deductive coding · Cohen's Kappa · Inter-coder reliability · Weighting (RIM/raking) · Post-stratification · Design effect · Progressive disclosure · BLUF (Bottom Line Up Front) · Market-entry research · Barriers to entry · Entry-mode recommendation · Brand-extension-fit study · Structural Equation Modeling (SEM) · Latent variable · Path model · Model fit indices (CFI/RMSEA) · Research roadmap · Decision-influence tracking · Cost-avoidance framing · Build-vs-buy (research capability) · Commercial due diligence (CDD) · Customer reference calls · Gross revenue retention · Attitude-behaviour gap (sustainability) · Greenwashing perception research · Investor perception research · Pre-IPO messaging testing · Audience-brand fit · Brand lift study · Sponsored-content disclosure research · Voice of Customer (VoC) · Transactional survey · Closing the loop · Retail media network · ROAS (Return on Ad Spend) · Holdout/incrementality test · In-app micro-survey · App-store review mining · Own-device usability testing · Skip logic / routing · Soft launch (survey) · Full-path testing · Trade area · Gravity model (site selection) · Site-selection research · Footfall/mobility data · Cannibalisation analysis (retail/branch) · Catchment overlap · KYC activation funnel · Trading-frequency segmentation · Trust-signal research · Risk-perception calibration · Financial literacy assessment · Pre/post comprehension design · Behavioural follow-through (education efficacy) · Risk-profiling questionnaire validation · Explainability research · Recommendation override analysis · Alternative data · Ground-truthing (alternative data) · Early-warning signal (research) · Distribution-channel research · Advisor/RM needs research · Double-blind-spot risk (channel vs end-investor) · Grey-market premium (GMP) ·

RII/NII/QIB subscription categories · Listing-day price action · Customer Effort Score (CES) · First Contact Resolution (FCR) · Complaint-handling research (regulated financial services) · Employee-review mining · Self-selection bias (review platforms) · Trend-vs-snapshot analysis · Stakeholder-impact mapping · Consultation-response research · Post-implementation impact tracking · GIFT City / IFSC · NRI geography sub-segmentation · Relationship-triggered customer journey · Readability metrics (Flesch-Kincaid) · Disclosure comprehension testing · Plain-language rewriting · Referral incentive research · Referral trigger-moment mapping · Incentive-motivated vs organic referral quality · Peeking (experimentation) · Sequential testing · Multi-armed bandit · Guardrail metrics · Fee-structure comprehension research · Cross-platform fee comparison · Perceived fee fairness · IR effectiveness research · Analyst/institutional-investor perception surveys · Earnings-call quality assessment · Finfluencer disclosure-compliance research · Credibility/expertise-perception research · Content-accuracy audit · Behaviour-quality research (gamification) · Segment-specific vulnerability research · Trading-streak / engagement-design risk · Hiring-signal analysis · Functional/seniority posting mix · LinkedIn headcount-by-function tracking · Cohort retention table · Retention-curve flattening ("smile curve") · Cohort-vs-cohort comparison · Post-purchase cognitive dissonance · Dissonance-reduction behaviours · Regret-timing window analysis · Filter bubble (recommendations) · Cold-start research · Downstream engagement tracking · Panel attrition bias · Panel conditioning effect · Refreshment sampling

## APPENDIX B — KEY FORMULAS

---

- Sample size (proportion, 95% CI):  $n = 1.9604 \times p(1-p) / e^2$
- Finite population correction:  $n_{adj} = n / (1 + (n-1)/N)$
- NPS = %Promoters (9-10) – %Detractors (0-6)
- Market size (bottom-up): # potential customers × price × purchase frequency
- Margin of error at 95% CI (given n, p):  $e = 1.96 \times \sqrt{p(1-p)/n}$

## APPENDIX C — PRACTICE / INTERVIEW Q&A

1. **Q: Why triangulate top-down and bottom-up sizing instead of picking one?** A: Each method has independent failure modes (top-down compounds percentage-assumption errors; bottom-up depends on complete unit-economics data) — agreement between the two is the actual evidence of a credible number, not either method alone.
2. **Q: A survey shows 70% of respondents "would definitely buy" a new product. Should the client trust this for a launch decision?** A: No — stated purchase intent systematically overstates actual purchase behaviour (a well-known market research finding); use intent as a directional signal, and where possible validate with a smaller real-world pilot/test market before a full launch.
3. **Q: When would you use qualitative research instead of a survey, even though surveys are more "statistical"?** A: When the goal is to discover *why* or generate hypotheses (exploratory stage), when the topic is sensitive/hard to pre-code into closed questions, or when sample access is limited (e.g. B2B decision-makers) — depth beats breadth in these cases.
4. **Q: What's wrong with reporting NPS alone as the health metric for a product?** A: It compresses a lot of information into one number and hides *why* scores are what they are; always pair with a verbatim follow-up and driver analysis (Part 6.5) so it's actionable, not just a scoreboard number.
5. **Q: How would you segment the customer base of a digital lending app?** A: Start a-priori (income band, loan purpose, city tier) for quick actionability, then validate/refine with post-hoc clustering on behavioural data (repayment behaviour, app usage frequency, cross-sell uptake) to check whether the a-priori segments actually predict the outcomes that matter (default risk, LTV) — if not, re-cut using the behavioural clusters instead.
6. **Q: A regression shows "feature completeness" is not a statistically significant driver of NPS ( $p=0.31$ ). Does that mean features don't matter to customers?** A: No — it means this sample/model found no evidence features drive *NPS specifically*; features could still matter for acquisition, for a different outcome (e.g. willingness to pay), or the measure used may not have captured the right feature dimension. Absence of significance is not evidence of absence of any effect, only of this specific effect on this specific outcome.
7. **Q: Why use Van Westendorp's four-question price meter instead of just asking "how much would you pay for this"?** A: A single willingness-to-pay question is easily gamed (respondents anchor low to seem price-savvy) and gives one number with no sense of an acceptable range; the four Van Westendorp questions triangulate a defensible price *range* (between the "too cheap" and "too expensive" crossover points) and let you cross-check internal consistency of each respondent's answers.
8. **Q: What's the difference between unaided and aided brand awareness, and why report both?** A: Unaided asks respondents to name brands with no prompt (harder, more diagnostic of true salience); aided shows a brand name and asks if they've heard of it (easier, inflated by suggestion). Reporting aided alone overstates real-world brand salience; the gap between the two is itself informative about how much a brand relies on prompted recognition versus genuine top-of-mind presence.
9. **Q: A B2B client wants a market-size number for a niche enterprise-software category, but says they only have budget/time for 40 survey responses. Is that usable?** A: In a B2B context with a small, hard-to-reach buying-committee population,  $n=40$  can be a legitimate primary-research sample for directional insight (especially when combined with qualitative depth and firmographic secondary data for the sizing itself), unlike in B2C where  $n=40$  would rarely be presented as statistically representative — the right response is to be explicit about what the sample size can and can't support (directional patterns, not a precise percentage with a tight margin of error) rather than refusing the brief outright.



10. **Q: How is a brand tracker different from a one-off brand study?** A: A tracker is methodologically frozen (fixed questions, fixed quotas, fixed timing) and run repeatedly specifically so that wave-over-wave *change* is measurable and attributable to real market movement rather than a change in how the data was collected — a one-off study optimises for depth on a single point in time, a tracker optimises for comparability over time.
11. **Q: You find your top-down and bottom-up market-size estimates differ by 40%. What do you do?** A: Do not average them or quietly pick the more favourable one — investigate which specific assumption is driving the gap (usually the weakest-evidenced percentage in the top-down chain, or an incomplete unit-economics input in the bottom-up build), report the investigated range with the reconciled explanation, and flag which number has the stronger evidentiary basis rather than presenting false precision.
12. **Q: What is "sugging" and why does it matter to a market researcher specifically?** A: Selling under the guise of research — contacting someone under a research pretext and then pitching them a product. It matters because it erodes public trust in legitimate research (lower response rates, more suspicion of survey calls generally) and is treated as a serious ethical/professional violation by industry bodies — a market researcher must keep a hard line between research contact and any sales activity.
13. **Q: Why does B2B research typically use much smaller sample sizes than B2C research, and is that a methodological weakness?** A: It reflects the actual population size, not a shortcut — a market of 500 relevant companies globally cannot statistically support a sample of thousands. It is not a weakness as long as it's disclosed honestly (a smaller, well-recruited B2B sample of the *right* stakeholders is more valuable than a larger sample of the wrong ones), and B2B research leans more on qualitative depth for exactly this reason.
14. **Q: A vendor delivers survey data with an unusually short average completion time compared to your pilot. What's your first move before analysing the data?** A: Run the data-quality checks (Part 5.8) — speeding, straight-lining, attention-check failures, open-end gibberish — before doing any substantive analysis; accepting unvalidated vendor data at face value is a common, costly mistake that surfaces only after a flawed finding has already reached a stakeholder deck.
15. **Q: Social listening shows overwhelmingly positive sentiment about a new product launch. Should the client conclude the launch is a success?** A: Not on sentiment alone — social-media populations skew younger, more urban, and more vocal than the general customer base (Part 13.1), so positive buzz doesn't guarantee broad-based adoption or purchase intent among the full target market; social listening is a fast, cheap early-signal tool best triangulated with structured survey data or actual sales/usage metrics before declaring success.
16. **Q: Why is "peeking" at A/B test results before the pre-planned sample size is reached a statistical problem, not just an impatience issue?** A: Checking significance repeatedly as data accumulates and stopping as soon as it looks significant inflates the false-positive rate in exactly the same way as the multiple-comparisons problem (Part 6.9) — each additional peek is effectively another "test," and the probability that at least one of them shows a spurious significant result by chance alone rises well above the nominal 5% threshold; the fix is committing to a sample size or test duration in advance and not stopping early based on interim results.
17. **Q: A brand's market share is 18%, but its share of shelf in a retail audit is only 11%. What's the actionable implication?** A: The brand is meaningfully under-shelved relative to its actual market performance — this is a quantified, retailer-facing negotiating point (the brand is "earning" more sales per unit of shelf space than its competitors, an efficiency argument for more facings) rather than a vague complaint about shelf placement, and it's exactly the kind of concrete, numbers-backed finding that gives a shopper-research function real influence with trade/retail relationship teams.

18. **Q: Why can't a global survey instrument simply be translated once and fielded identically across all markets in a multi-country study?** A: Literal translation doesn't guarantee equivalent meaning (an idiom or attribute term may not map cleanly across languages), and cultural response-style differences (acquiescence bias, extreme response style — Part 15.3) mean raw cross-country score comparisons can be misleading even with a perfectly translated instrument; rigorous multi-country research requires back-translation validation and often statistical adjustment for response-style differences before comparing results across countries.
19. **Q: In the worked segmentation case study (Part 6.10), why is validating that clusters predict a real business outcome described as more important than the elegance of the clustering solution itself?** A: A statistically clean, distinct, well-separated cluster solution is not automatically useful — the entire point of segmentation is to enable differentiated business action (different products, messaging, or pricing per segment), so if the clusters don't actually differ on outcomes the business cares about (repayment behaviour, LTV, conversion), the exercise has produced an interesting but non-actionable typology rather than a genuine strategic tool.
20. **Q: What's the difference between shopper research and consumer research, and why might a company need both even for the same product?** A: Consumer research explains why people choose a brand generally (need states, brand perception); shopper research studies the specific in-store/on-page purchase moment, which can diverge from stated brand preference due to shelf placement, promotions, or convenience. A company might find via consumer research that shoppers prefer Brand A, yet shopper research reveals Brand B wins at the shelf due to superior positioning — pointing to a trade-marketing/execution gap rather than a brand-perception problem, a distinction that determines whether the fix is a marketing campaign or a retail-execution intervention.
21. **Q: A channel's average ROI from an MMM decomposition looks strong, but the finance team wants to cut its budget. Why might this be the right call?** A: Average ROI (total contribution ÷ total spend) can look healthy even when marginal ROI (the return on the *next* rupee) has fallen well below breakeven once a channel is deep into its diminishing-returns saturation curve (Part 20.3-20.4) — past spend already delivered its contribution, but that doesn't justify continuing to spend at the same level if the channel is saturated; the finance team's instinct may be correctly reading the marginal picture even if the average-ROI headline number looks fine.
22. **Q: Why is CLV described as "highly non-linear" in churn rate, and what does that imply for prioritising retention investment?** A:  $CLV = (ARPU \times \text{margin}) / \text{churn rate}$  is an inverse relationship, so the CLV gain from a given percentage-point churn reduction is much larger at low churn rates than the same-sized cut delivers at high churn rates (Part 21.4's worked chart) — this implies retention investment aimed at an already-low-churn, high-value segment can generate disproportionately large CLV gains from even a small further improvement, sometimes a better-ROI investment than the same effort applied to a high-churn segment where the CLV curve has already flattened.
23. **Q: An eye-tracking study shows respondents look at a product's price for an unusually long time before deciding. What does this actually tell you, and what doesn't it tell you?** A: It tells you attention/dwell time is concentrated on price — a genuine, hard-to-self-report signal about what's driving deliberation (Part 22.2). It does *not* tell you whether that attention resulted in a positive or negative reaction (dwell time alone doesn't carry emotional valence) — that requires pairing with facial coding, GSR, or a follow-up self-report question, exactly the multi-method triangulation Part 22 emphasises rather than over-relying on any single implicit measure alone.
24. **Q: Why does researching a financial product's concept require more attention to regulatory disclosure constraints than researching, say, a snack food concept?** A: Financial-product advertising and concept communication in India is subject to SEBI (for mutual funds/broking)



and IRDAI (for insurance) disclosure requirements — a concept-test stimulus that omits mandatory risk disclosures a real launched product would have to include produces an artificially rosier reaction than the actual, compliant product would get in market (Part 23.2), a realism gap a researcher must account for when interpreting concept-test scores for regulated financial products.

25. **Q: A stakeholder asks for a rapid, 3-day turnaround study on a decision that will be very expensive and hard to reverse if wrong. What should you tell them?** A: Apply the rapid-research decision framework (Part 24.3) explicitly: rapid/agile methods trade depth for speed and are well-suited to low-stakes, easily-reversible decisions, not high-stakes, hard-to-reverse ones — the right response is to flag this mismatch directly to the stakeholder, proposing either a properly-scoped timeline for a decision this consequential, or a narrower, lower-stakes sub-question that genuinely can be answered rapidly and reliably within 3 days, rather than silently delivering an under-powered rapid study that creates false confidence in a big decision.
26. **Q: Two coders independently coding the same set of interview transcripts produce noticeably different themes. Does this mean the research has failed?** A: Not necessarily — disagreement is exactly what inter-coder reliability checks (Part 25.4) are designed to surface, and resolving it through discussion often improves the codebook itself by catching ambiguous code definitions. It only signals a real problem if agreement (measured via Cohen's Kappa) stays low even after the codebook is refined and coders are properly briefed — a single round of initial disagreement, especially before the codebook is finalised, is a normal, expected part of the process, not a failure.
27. **Q: A survey achieves 60% female / 40% male respondents against a true population split of 51/49. Why not just discard the excess female responses to fix the balance instead of weighting?** A: Discarding valid responses wastes real, already-collected data and reduces the effective sample size, worsening precision — weighting (Part 26.2) achieves the same balance by adjusting how much each response *counts* toward the topline rather than deleting information, preserving the full sample's statistical power while still correcting the demographic imbalance; discarding data to fix a skew is rarely the right call when weighting is available.
28. **Q: A stakeholder pushes back live on a finding, saying it contradicts what their sales team is telling them. How should a presenter respond in the moment?** A: Per Part 27.3: acknowledge the tension without becoming defensive, distinguish the finding (what the data actually shows) from any interpretation layered on top of it, and be explicit about what the research did and didn't cover ("we didn't specifically test that account segment the sales team is describing — that's a legitimate gap worth a follow-up") rather than either dismissing the sales team's input or abandoning a well-supported finding under pressure. Handling the challenge credibly matters more to the research's standing than avoiding it ever coming up.
29. **Q: A company wants to enter Tier-2/3 cities with a product that succeeds in metro markets. What's the key mistake to avoid in scoping this research?** A: Treating it as "the same market, smaller cities" rather than validating the assumptions that don't automatically transfer (Part 28.3) — customer need states, distribution/retail structure, and price sensitivity can all differ meaningfully in Tier-2/3 markets, and a rigorous entry study runs a real pilot to validate these before a full rollout commitment rather than assuming metro learnings apply unchanged.
30. **Q: Why is SEM described as "overkill" for many research questions despite being more statistically rigorous than simple regression?** A: SEM requires meaningfully larger sample sizes (often several hundred respondents even for a modest model), specialised software, and advanced statistical training (Part 29.4) — worth the investment for testing a multi-step causal theory with latent constructs on a reusable tracking programme, but unnecessary complexity and cost for a one-off question a standard regression driver analysis would answer just as well.

31. **Q: What's the difference between a research function demonstrating "decision-influence" versus a hard revenue-ROI number, and why does the handbook favour the former?** A: A clean, single ROI number for a research function as a whole is genuinely difficult to construct honestly, since research informs decisions several steps removed from a final business outcome (Part 30.3) — decision-influence tracking (logging which specific decisions each study informed, and the estimated cost of getting that decision wrong) and cost-avoidance framing are more defensible, achievable approaches than forcing an artificial revenue-attribution number.
32. **Q: A commercial due diligence engagement finds that customer reference calls contradict a target company's own historical retention metrics. Which should the deal team trust more?** A: Neither blindly — the historical metric (Part 31.4) describes past behaviour, while reference calls are a forward-looking signal about future risk; the rigorous approach treats the reference-call finding as a reason to stress-test the deal thesis (e.g. building a downside retention scenario into the valuation model) rather than dismissing it because it conflicts with a cleaner historical number, since CDD research specifically exists to surface exactly this kind of forward-looking risk historical data can't capture.
33. **Q: Why is stated willingness-to-pay for a sustainable product reformulation considered especially unreliable, even more than typical stated purchase intent?** A: Sustainability research has one of the most extensively documented attitude-behaviour gaps of any category (Part 32.2) — consumers consistently overstate their willingness to pay a premium for sustainable options in surveys relative to their actual purchase behaviour, making a real-money pilot (charging the actual premium and measuring real behaviour) a far more reliable validation method than trusting the stated survey figure.
34. **Q: An investor perception study ahead of a follow-on offering reveals investors are confused about how two business segments relate. What's the appropriate output of this finding?** A: Not just "investors are confused" (a finding) but a specific insight and fix (Part 33.4) — e.g. a clearer segment-relationship narrative or diagram addressing the exact confusion directly in the investor presentation, ideally tested again with a follow-up round of interviews before the actual roadshow, turning the finding into a concrete presentation-design change rather than a vague caution.
35. **Q: An influencer campaign shows strong engagement but weak tracked-link conversion. What are the three distinct possible root causes, and why does distinguishing them matter?** A: Per Part 34.5: poor audience-brand fit (engagement reflects the creator's general appeal, not genuine category interest), execution friction (a poorly-placed or unclear call-to-action/link), or weak creative (engaging content that doesn't make a clear case for the product). Each requires a different fix — better influencer selection, better campaign mechanics, or better creative briefing — so correctly diagnosing which cause is at play, rather than concluding "influencer marketing doesn't work," determines what should actually change next time.
36. **Q: A VoC dashboard shows a spike in support tickets about a redesigned checkout flow. Is this enough to identify and fix the underlying usability problem?** A: No (Part 35.4) — VoC data is well-suited to fast *detection* of a problem but less suited to fully *diagnosing* the specific failure; the appropriate next step is pulling the behavioural funnel data for the flagged flow and running targeted usability sessions on the specific step, using VoC as the early-warning trigger and a focused follow-up study as the diagnostic tool.
37. **Q: A new bank-branch site has a strong standalone trade-area profile, but sits 1.2km from an existing branch of the same bank. Why isn't the standalone trade-area number enough to greenlight the site?** A: Per Part 39.4-39.5 — a strong catchment population and income profile only establishes total addressable demand in the area, not how much of the new branch's projected

volume would be genuinely incremental versus diverted from the existing nearby branch within the overlapping catchment. Cannibalisation analysis must explicitly estimate that diverted share before the business case can be trusted, the same discipline applied to retail-media incrementality (Part 36.2) and assortment decisions (Part 19.2) elsewhere in this handbook.

38. **Q: A discount broker sees most first-time F&O traders stop trading entirely within 30 days. Is this necessarily an app-UX problem to fix?** A: Not necessarily (Part 40.5) — it could be an execution-friction UX problem, a risk-comprehension problem where new traders exit after losses exceeded what they understood the leveraged risk to be, or a natural pattern where a share of first-time options traders were always going to be one-off "curiosity" trades. Combining loss-size behavioural data with direct qualitative outreach to churned users is necessary to identify which cause dominates, since a UX fix only helps if the first hypothesis is actually true.
39. **Q: A broker's mandatory pre-F&O education module gets a 4.3/5 satisfaction score. Is that sufficient evidence the module is "working"?** A: No (Part 41.5) — satisfaction measures how pleasant or clear the module felt, not what was retained or whether behaviour changed. A credible evaluation needs a pre/post knowledge assessment using validated risk/leverage-comprehension questions (Part 41.2-41.3) to establish genuine comprehension gain, and ideally a behavioural follow-up comparing completers against a matched comparison group on downstream outcomes like early account closure or outsized first-month losses (Part 41.4).
40. **Q: A robo-advisory platform sees a spike in "moderate risk" users overriding their recommended allocation during a market drawdown, with no prior override history. What's the most likely root cause, and what should the research response be?** A: Per Part 42.5 — this most likely reflects a gap between the stated risk tolerance the questionnaire captured and the user's actual, felt risk tolerance under real market stress, a stated-vs-revealed-preference gap (Part 22.1) exposed by the drawdown itself. Rather than defaulting to an explanation-copy fix (Part 42.3), the response should check the risk-profiling questionnaire's behavioural predictive validity for this cohort (Part 42.2) — if it's failing to capture genuine risk tolerance for a meaningful subset, the questionnaire itself (e.g. adding a scenario-based drawdown-simulation question) needs redesigning, not just the explanation layer.
41. **Q: App-download-ranking data shows a fintech competitor's app jumped from #45 to #8 with no corresponding news coverage. Should this be reported as evidence of a genuine competitive shift?** A: Not yet (Part 43.5) — the ground-truthing discipline requires first checking whether this data provider's rank movements have historically correlated with real user-growth outcomes, and ruling out a paid-acquisition spike or app-store algorithm change, since a rank jump with no news is at least as consistent with temporary paid UA as organic demand growth. If the signal survives that check, the appropriate response is triggering targeted primary research (a mystery-shop, an app-store-review scan) to establish *why*, not reporting the alternative-data signal alone as a finding.
42. **Q: Advisor research on a new thematic fund is strongly positive (easy to pitch), but end-investor research shows clients misunderstand the fund's risk profile. Should this be reported as a successful fund?** A: No (Part 44.5) — this is a genuine double-blind-spot divergence (Part 44.4), not a contradiction to average away: the narrative that makes the fund easy to pitch may be the same narrative causing end-investor risk misunderstanding. The correct response is testing revised advisor-facing materials that foreground the risk profile alongside the thematic narrative, then re-running the end-investor comprehension check (Part 41's pre/post design) to confirm the fix closes the understanding gap rather than just improving the advisor-research score.
43. **Q: An IPO shows QIB subscription at 45x, NII at 60x, but RII at only 0.8x, alongside a grey-market premium suggesting a strong listing gain. How should these signals be**

**weighed together?** A: Per Part 45.5 — the divergence between overwhelming institutional demand and retail undersubscription is itself the most informative signal, suggesting the story may not be resonating with individual investors as it has with institutions applying more rigorous demand estimation. The GMP figure (Part 45.3) should be weighted cautiously given its thin, informal-market basis and weighed against this specific demand composition, not treated as an independent confirming signal, ideally discounted using a firm's own track record of validating GMP against actual outcomes (Part 45.4).

44. **Q: A broking platform's overall NPS is strong (+42), but CES on trade-dispute resolution is poor, with only 35% first-contact resolution requiring 3+ follow-ups. What does the NPS/CES divergence tell you, and where should the research look next?** A: Per Part 46.5 — the divergence shows the broader product is well-regarded, but this specific high-stakes support journey is a genuine blind spot the aggregate NPS masks, exactly why CES needs tracking separately for support interactions (Part 46.2). The low FCR is the direct diagnostic lead (Part 46.3) — investigating why disputes require multiple contacts (unclear documentation requirements, internal routing/escalation friction, an approval step unavailable to the first-contact agent) identifies the specific fix, and given trade disputes risk escalating to a regulator/ombudsman if handled poorly (Part 46.4), this finding should be prioritised beyond its CES/NPS impact alone.
45. **Q: A company's employee-review rating has declined for four straight quarters, with a rising share of reviews mentioning "unrealistic launch timelines," while public messaging remains confidently on-schedule. Should this be reported as evidence the launch will slip?** A: No (Part 47.5) — the sustained multi-quarter trend and specific recurring theme are more defensible and actionable than a single-quarter dip (Part 47.3-47.4), but the self-selection bias inherent to review platforms (Part 47.2) means the reviewers raising these concerns aren't necessarily representative of the full team's view. The correct response is flagging it as a specific, evidence-based lead to probe through other channels (earnings-call Q&A, expert-network conversations, channel checks), not reporting it as a standalone conclusion.
46. **Q: A brokerage is asked to research the likely impact of a proposed SEBI margin-requirement increase ahead of an industry-body comment submission. Why shouldn't one blended survey across all affected parties be sufficient?** A: Per Part 48.5 — stakeholder-impact mapping (Part 48.2) must come first, since retail traders, the brokerage's own compliance/capital operations, and smaller regional brokers each experience the change in genuinely different ways requiring differently-designed research instruments, and a single blended question would produce an uninterpretable average across fundamentally different impacts. The output feeding a formal comment-period submission (Part 48.3) also needs the same evidentiary rigor as any other research in this handbook — documented methodology and representative sampling within each stakeholder group — since a weakly-evidenced submission carries less regulatory influence.
47. **Q: A brief for a new NRI-focused investment product treats "NRI investors" as a single target segment. What should be challenged before research design begins?** A: Per Part 49.5 — the single-segment framing itself, since Gulf, US/UK/Canada, and other NRI geographies face genuinely different tax-treaty and regulatory considerations (Part 49.2) that directly affect product fit and messaging, not just wealth-tier variants of the same persona. The research should also map the relationship-triggered journey structure (Part 49.4) rather than assuming a self-directed discovery funnel, since understanding when and through what channel to reach this segment is often more actionable than a generic feature-preference survey alone.
48. **Q: A margin-call disclosure scores at an acceptable Flesch-Kincaid grade level, but comprehension testing finds only 40% of respondents can correctly state what triggers a margin call. What does this tell you, and what's the correct next step?** A: Per Part 50.5 — this



confirms readability metrics alone are an inadequate check (Part 50.2's "necessary but not sufficient" distinction); the disclosure is grammatically simple but conceptually unclear, likely due to dense financial-concept density a Flesch-Kincaid-style formula cannot detect. The correct next step is a plain-language rewrite targeting the diagnosed gap (e.g. adding a concrete numeric example), followed by re-testing the revised version against the same comprehension questions (Part 50.4) — a rewrite that "feels clearer" without being re-tested risks the same mistake in a new form.

49. **Q: A broking platform's NPS is strong (+38), but referral-program participation is low, and referred customers who do sign up show below-average 90-day retention. What does this combination of findings tell you?** A: Per Part 51.5 — the strong NPS alongside weak referral behaviour confirms NPS shouldn't be relied on as a proxy for actual referral behaviour (Part 51.2), since satisfaction existing doesn't mean it's converting into referral action. The below-average retention among referred sign-ups (Part 51.3) separately suggests the program's financial incentive may be attracting incentive-motivated referrers passing the link to less-genuinely-interested contacts rather than satisfied customers organically recommending the platform — a finding that should prompt reconsidering the incentive structure and referral-ask placement (Part 51.4) rather than simply increasing the incentive to drive more volume of the same lower-quality referrals.
50. **Q: A bandit-based onboarding-flow experiment shifts most traffic to a new flow within days based on a strong same-day activation lift, but the 30-day retention guardrail metric hasn't matured yet. Should the bandit's reallocation be trusted as proof of a genuine win?** A: Not yet (Part 52.5) — the bandit is working as designed for the fast-maturing primary metric (Part 52.3), but a flow that front-loads users into the funnel could plausibly lift same-day activation while degrading longer-term fit, exactly what the guardrail-metric discipline (Part 52.4) exists to catch. The disciplined response is holding out a smaller comparison cohort on the original flow long enough for the 30-day guardrail to mature, rather than treating the bandit's fast early reallocation as final proof.
51. **Q: Support contacts about "unexpected" F&O charges are rising, concentrated among first-month F&O traders, even though the fee structure is fully disclosed in the published schedule and shown at order placement. Is more disclosure the right fix?** A: Not necessarily (Part 53.5) — this is a fee-structure comprehension gap (Part 53.2) formed at the point users built cost expectations around the platform's prominent zero-brokerage-delivery positioning (Part 53.1), not a failure of any single disclosure point. The more effective fix addresses the perceived-fairness and expectation-setting gap (Part 53.4) — a clearer upfront cost explanation during F&O onboarding specifically, tested with pre/post comprehension methodology (Part 41.3) — rather than relying on order-time disclosure alone to correct an expectation already formed around a different product's pricing.
52. **Q: A mid-cap's sell-side analyst coverage has fallen from 12 to 7 analysts over 18 months with no meaningful deterioration in fundamentals. IR-effectiveness research finds analysts cite inconsistent guidance and difficulty securing timely management follow-up. What's the right diagnosis and fix?** A: Per Part 54.5 — with fundamentals unchanged, the specific IR-quality issues surfaced (Part 54.3) are the more plausible explanation: inconsistent guidance undermines analysts' ability to build and defend confident forecasts, and limited follow-up access compounds their difficulty resolving uncertainty directly. This gives management concrete, actionable levers — more disciplined guidance practices and better post-earnings-call access — rather than only a vague sense that "the market doesn't understand the story."
53. **Q: A influencer partnership shows strong engagement and click-through, but a content-accuracy audit finds recent posts understate investment risk relative to approved guidelines, though technically compliant on sponsorship disclosure. Should the strong performance metrics offset the audit finding?** A: No (Part 55.5) — the compliance flag is a

different category of concern, not a competing consideration to balance against performance metrics (Part 55.2's higher-stakes framing), since risk-understating content creates real regulatory and reputational exposure regardless of how well it converts. The correct response is direct remediation — engaging the creator on the risk-messaging drift and tightening ongoing content review, not just pre-approval — treating strong short-term metrics as no substitute for verifying accuracy standards hold across the partnership's full duration.

54. **Q: A "trading streak" pilot feature shows a strong pooled engagement lift, but segmented analysis shows it correlates with increased trade frequency among first-90-day traders specifically, with no corresponding performance improvement, while experienced traders show negligible change. Is the pooled engagement lift sufficient justification to roll the feature out?** A: No (Part 56.5) — the pooled metric alone is incomplete and potentially misleading; the segment-specific finding reveals the feature's real behavioural effect is concentrated in the most vulnerable segment (Part 56.4), correlating with increased frequency but no performance gain, a pattern consistent with encouraged-impulsive-trading risk this category's gamification research must screen for (Part 56.1-56.2). Given the regulatory attention this design pattern draws (Part 56.3), the finding should be flagged to product leadership and compliance, with redesign or exclusion of new traders considered rather than shipping on engagement alone.
55. **Q: A fintech competitor's job postings show a three-month rise specifically in sales, compliance, and a senior "Regional Head" role concentrated in one new city, with corroborating LinkedIn headcount growth there and no public announcement yet. Why is this a stronger signal than a generic rise in overall posting volume?** A: Per Part 57.5 — the specific functional and seniority mix (Part 57.3) is characteristic of deliberate market-entry preparation rather than routine hiring, with the compliance role specifically suggesting local regulatory groundwork; a generic volume rise alone couldn't distinguish this from replacement hiring (Part 57.2). The corroborating LinkedIn headcount growth (Part 57.4) confirms this is converting into real headcount, not just unfilled postings, making it a well-evidenced, pre-announcement early-warning signal — though still one to corroborate further (Part 43.3's ground-truthing discipline), not treat as definitive on its own.
56. **Q: A cohort table shows post-redesign onboarding cohorts have noticeably higher Month-1 retention than pre-redesign cohorts, but both converge to a similar retention rate by Month-4. What's the correct conclusion, and why is neither "retention improved" nor "retention didn't change" complete on its own?** A: Per Part 58.5 — reading down the Month-1 column shows a genuine early improvement, but reading further along each cohort's row (Part 58.3's curve-shape discipline) shows that improvement doesn't persist past early retention. The precise, actionable conclusion is that the redesign solved an early-funnel problem specifically without changing longer-term engagement determinants — if long-term retention is the actual priority, a different intervention targeting the Month-4+ plateau is needed, not further onboarding-flow iteration.
57. **Q: A long-lock-in investment product shows a spike in exit requests concentrated in the first 30 days after purchase, despite documented performance being in line with stated expectations. What's the more likely explanation, and what should the response be?** A: Per Part 59.5 — the early-window concentration, alongside in-line performance, points to a dissonance-driven pattern (Part 59.2-59.3) rather than a genuine performance grievance, since this is precisely when dissonance-reduction behaviours haven't yet had time to settle the discomfort of a large, hard-to-reverse commitment. The response should be a proactive touchpoint specifically timed to this identified early window (Part 59.4) — a reassuring check-in delivered before the dissonance peak — rather than a generic post-hoc satisfaction fix or simply restating performance figures the exit-requesters already know are in line.

58. **Q: A stock-discovery feed algorithm update lifts click-through rate by 15%, but relevance ratings decline slightly and downstream time-on-page drops, concentrated on higher-volatility, more sensationally-framed recommendations that increased in frequency. Should the update be declared a success based on the click-through lift?** A: No (Part 60.5) — click-through alone is an incomplete metric (Part 60.2); the relevance-rating decline and dropped downstream engagement reveal the algorithm is likely optimising toward clickbait-adjacent suggestions that generate clicks without delivering value. This also raises a filter-bubble/diversity concern (Part 60.3) if sensational recommendation types are being systematically favoured — the recommendation is reverting or tuning the update to weight the fuller relevance-and-engagement picture, not click-through alone.
59. **Q: A quarterly investor-confidence tracker with no refreshment sampling shows confidence holding up unusually well through a six-month market correction, more resilient than independent sentiment indicators suggest is plausible. What should be checked before trusting this reading?** A: Per Part 61.5 — the absence of refreshment sampling (Part 61.4) means it's impossible to distinguish genuine resilient sentiment from panel attrition bias (Part 61.2, the most anxious respondents disproportionately dropping out) compounded by panel conditioning (Part 61.3, repeated respondents becoming desensitised to volatility questions). The diagnostic step is checking whether response rates dropped during the correction specifically, and implementing refreshment sampling going forward before trusting this tracker through future cycles.